

SUSE Linux

10.1

April 07, 2006

Reference

www.novell.com



Reference

List of Authors: Jörg Arndt, Stefan Behlert, Frank Bodammer, James Branam, Volker Buzek, Klara Cihlarova, Stefan Dirsch, Olaf Donjak, Roman Drahtmüller, Thorsten Dubiel, Torsten Duwe, Thomas Fehr, Stefan Fent, Werner Fink, Jakub Friedl, Kurt Garloff, Joachim Gleißner, Carsten Groß, Andreas Grünbacher, Berthold Gunreben, Franz Hassels, Andreas Jaeger, Jana Jaeger, Klaus Kämpf, Andi Kleen, Hubert Mantel, Lars Marowsky-Bree, Chris Mason, Johannes Meixner, Lars Müller, Matthias Nagorni, Anas Nashif, Siegfried Olschner, Edith Parzefall, Peter Pöml, Thomas Renninger, Hannes Reinecke, Scott Rhoades, Thomas Rölz, Heiko Rommel, Tanja Roth, Marcus Schäfer, Thomas Schraitle, Klaus Singvogel, Frank Sundermeyer, Elisabeth Tobiasson, Hendrik Vogelsang, Klaus G. Wagner, Rebecca Walter, Christian Zoz

This publication is intellectual property of Novell Inc.

Its contents can be duplicated, either in part or in whole, provided that a copyright label is visibly located on each copy.

All information found in this book has been compiled with utmost attention to detail. However, this does not guarantee complete accuracy. Neither SUSE LINUX GmbH, the authors, nor the translators shall be held liable for possible errors or the consequences thereof.

Novell, the Novell logo, the N logo and SUSE are registered trademarks of Novell, Inc. in the United States and other countries. * Linux is a registered trademark of Linus Torvalds. All other third party trademarks are the property of their respective owners.

Contents

About This Guide	xi
Part I Advanced Deployment Scenarios	15
1 Remote Installation	17
1.1 Installation Scenarios for Remote Installation	17
1.2 Setting Up the Server Holding the Installation Sources	26
1.3 Preparing the Boot of the Target System	35
1.4 Booting the Target System for Installation	45
1.5 Monitoring the Installation Process	49
2 Advanced Disk Setup	53
2.1 LVM Configuration	53
2.2 Soft RAID Configuration	60
3 Updating the System and Package Management	65
3.1 Updating SUSE Linux	65
3.2 Software Changes from Version to Version	68
3.3 RPM—the Package Manager	86
Part II Administration	99
4 Security in Linux	101
4.1 Masquerading and Firewalls	101
4.2 SSH: Secure Network Operations	110

4.3	Encrypting Partitions and Files	116
4.4	Confining Privileges with AppArmor	119
4.5	Security and Confidentiality	128
5	Access Control Lists in Linux	141
5.1	Traditional File Permissions	141
5.2	Advantages of ACLs	143
5.3	Definitions	143
5.4	Handling ACLs	144
5.5	ACL Support in Applications	152
5.6	For More Information	152
6	System Monitoring Utilities	153
6.1	List of Open Files: lsof	153
6.2	User Accessing Files: fuser	155
6.3	File Properties: stat	155
6.4	USB Devices: lsusb	156
6.5	Information about a SCSI Device: scsiinfo	156
6.6	Processes: top	157
6.7	Process List: ps	158
6.8	Process Tree: pstree	159
6.9	Who Is Doing What: w	160
6.10	Memory Usage: free	160
6.11	Kernel Ring Buffer: dmesg	161
6.12	File Systems and Their Usage: mount, df, and du	161
6.13	The /proc File System	162
6.14	PCI Resources: lspci	165
6.15	System Calls of a Program Run: strace	166
6.16	Library Calls of a Program Run: ltrace	167
6.17	Specifying the Required Library: ldd	168
6.18	Additional Information about ELF Binaries	168
6.19	Interprocess Communication: ipcs	169
6.20	Time Measurement with time	169
Part III	System	171
7	32-Bit and 64-Bit Applications in a 64-Bit System Environment	173
7.1	Runtime Support	173
7.2	Software Development	174
7.3	Software Compilation on Biarch Platforms	174
7.4	Kernel Specifications	176

8 Booting and Configuring a Linux System	177
8.1 The Linux Boot Process	177
8.2 The init Process	181
8.3 System Configuration via /etc/sysconfig	189
9 The Boot Loader	193
9.1 Selecting a Boot Loader	194
9.2 Booting with GRUB	194
9.3 Configuring the Boot Loader with YaST	203
9.4 Uninstalling the Linux Boot Loader	208
9.5 Creating Boot CDs	208
9.6 The Graphical SUSE Screen	209
9.7 Troubleshooting	210
9.8 For More Information	212
10 Special Features of SUSE Linux	213
10.1 Information about Special Software Packages	213
10.2 Virtual Consoles	220
10.3 Keyboard Mapping	220
10.4 Language and Country-Specific Settings	221
11 Printer Operation	227
11.1 Workflow of the Printing System	228
11.2 Methods and Protocols for Connecting Printers	229
11.3 Installing the Software	229
11.4 Configuring the Printer	230
11.5 Configuration for Applications	237
11.6 Special Features in SUSE Linux	237
11.7 Troubleshooting	243
12 Dynamic Kernel Device Management with udev	251
12.1 The /dev Directory	251
12.2 Kernel uevents and udev	252
12.3 Drivers, Kernel Modules, and Devices	252
12.4 Booting and Initial Device Setup	253
12.5 Debugging udev Events	253
12.6 Influencing Kernel Device Event Handling with udev Rules	254
12.7 Persistent Device Naming	255
12.8 The Replaced hotplug Package	256
12.9 For More Information	257

13 File Systems in Linux	259
13.1 Terminology	259
13.2 Major File Systems in Linux	260
13.3 Some Other Supported File Systems	266
13.4 Large File Support in Linux	267
13.5 For More Information	268
14 The X Window System	271
14.1 X11 Setup with SaX2	271
14.2 Optimizing the X Configuration	273
14.3 Installing and Configuring Fonts	279
14.4 OpenGL—3D Configuration	284
15 FreeNX: Remotely Controlling Another Computer	289
15.1 Getting Started with NX	289
15.2 Advanced FreeNX Configuration	291
15.3 Troubleshooting	298
15.4 For More Information	300
16 Authentication with PAM	301
16.1 Structure of a PAM Configuration File	302
16.2 The PAM Configuration of sshd	303
16.3 Configuration of PAM Modules	306
16.4 For More Information	308
17 Virtualization with Xen	309
17.1 Xen Installation	310
17.2 Domain Installation	311
17.3 Starting and Controlling Xen Domains with xm	312
17.4 Troubleshooting	313
17.5 For More Information	314
Part IV Services	315
18 Basic Networking	317
18.1 IP Addresses and Routing	321
18.2 IPv6—The Next Generation Internet	323
18.3 Name Resolution	333
18.4 Configuring a Network Connection with YaST	334

18.5	Managing Network Connections with NetworkManager	345
18.6	Configuring a Network Connection Manually	348
18.7	smpppd as Dial-up Assistant	359
19	SLP Services in the Network	363
19.1	Registering Your Own Services	363
19.2	SLP Front-Ends in SUSE Linux	364
19.3	Activating SLP	365
19.4	For More Information	365
20	The Domain Name System	367
20.1	DNS Terminology	367
20.2	Configuration with YaST	368
20.3	Starting the Name Server BIND	376
20.4	The Configuration File /etc/named.conf	378
20.5	Zone Files	382
20.6	Dynamic Update of Zone Data	387
20.7	Secure Transactions	387
20.8	DNS Security	388
20.9	For More Information	389
21	Using NIS	391
21.1	Configuring NIS Servers	391
21.2	Configuring NIS Clients	397
22	Sharing File Systems with NFS	399
22.1	Importing File Systems with YaST	399
22.2	Importing File Systems Manually	400
22.3	Exporting File Systems with YaST	401
22.4	Exporting File Systems Manually	402
22.5	For More Information	404
23	DHCP	405
23.1	Configuring a DHCP Server with YaST	406
23.2	DHCP Software Packages	410
23.3	The DHCP Server dhcpcd	411
23.4	For More Information	414

24 Time Synchronization with NTP	415
24.1 Configuring an NTP Client with YaST	415
24.2 Configuring xntp in the Network	418
24.3 Setting Up a Local Reference Clock	419
25 LDAP—A Directory Service	421
25.1 LDAP versus NIS	423
25.2 Structure of an LDAP Directory Tree	424
25.3 Server Configuration with slapd.conf	427
25.4 Data Handling in the LDAP Directory	432
25.5 The YaST LDAP Client	436
25.6 Configuring LDAP Users and Groups in YaST	444
25.7 For More Information	445
26 The Apache HTTP Server	447
26.1 Quick Start	447
26.2 Configuring Apache	449
26.3 Starting and Stopping Apache	463
26.4 Installing, Activating, and Configuring Modules	465
26.5 Getting CGI Scripts to Work	472
26.6 Setting Up a Secure Web Server with SSL	475
26.7 Avoiding Security Problems	480
26.8 Troubleshooting	482
26.9 For More Information	484
27 File Synchronization	487
27.1 Available Data Synchronization Software	487
27.2 Determining Factors for Selecting a Program	491
27.3 Introduction to Unison	495
27.4 Introduction to CVS	497
27.5 Introduction to Subversion	499
27.6 Introduction to rsync	502
27.7 Introduction to mailsync	504
28 Samba	509
28.1 Terminology	509
28.2 Starting and Stopping Samba	510
28.3 Configuring a Samba Server	511
28.4 Configuring Clients	516
28.5 Samba as Login Server	517
28.6 For More Information	518

29 The Proxy Server Squid	521
29.1 Some Facts about Proxy Caches	522
29.2 System Requirements	523
29.3 Starting Squid	525
29.4 The Configuration File /etc/squid/squid.conf	527
29.5 Configuring a Transparent Proxy	533
29.6 cachemgr.cgi	536
29.7 squidGuard	537
29.8 Cache Report Generation with Calamaris	539
29.9 For More Information	540
Part V Mobility	541
30 Mobile Computing with Linux	543
30.1 Laptops	543
30.2 Mobile Hardware	551
30.3 Cellular Phones and PDAs	552
30.4 For More Information	552
31 PCMCIA	555
31.1 Controlling PCMCIA Cards Using pccardctl	556
31.2 PCMCIA in Detail	556
31.3 Troubleshooting	559
32 System Configuration Profile Management	563
32.1 Terminology	564
32.2 Setting up SCPM	564
32.3 Configuring SCPM Using a Graphical User Interface	565
32.4 Configuring SCPM Using the Command Line	572
32.5 Troubleshooting	575
32.6 For More Information	576
33 Power Management	577
33.1 Power Saving Functions	578
33.2 APM	579
33.3 ACPI	580
33.4 Rest for the Hard Disk	588
33.5 The powersave Package	589
33.6 The YaST Power Management Module	597

34 Wireless Communication	603
34.1 Wireless LAN	603
34.2 Bluetooth	614
34.3 Infrared Data Transmission	625
Index	629

About This Guide

This manual gives you a general understanding of SUSE Linux. It is intended mainly for system administrators and home users with basic system administration knowledge. This manual presents a selection of applications needed in everyday life and provides in-depth descriptions of advanced installation and configuration scenarios.

Advanced Deployment Scenarios

Learn how to deploy SUSE Linux in complex environments.

Administration

Learn how to make your SUSE Linux system secure, how to deal with file system access controls, and get to know some important utilities for Linux administrators.

System

Get an introduction to the components of your Linux system and a deeper understanding of their interaction.

Services

Learn how to configure the various network and file services that come with SUSE Linux.

Mobility

Get an introduction to mobile computing with SUSE Linux and learn how to configure the various options for wireless computing, power management, and profile management.

1 Feedback

We want to hear your comments and suggestions about this manual and the other documentation included with this product. Please use the User Comments feature at the bottom of each page of the online documentation and enter your comments there.

2 Additional Documentation

There are other manuals available on this SUSE Linux product, either online at <http://www.novell.com/documentation/> or in your installed system under /usr/share/doc/manual/:

SUSE Linux *Start-Up*

This guide introduces you to the installation procedure of SUSE Linux and the basic use of your desktop environment. An online version of this document can be found at <http://www.novell.com/documentation/suse101/>

SUSE Linux *Applications*

This guide features a selection of the most important tools offered by SUSE Linux. An online version of this document can be found at <http://www.novell.com/documentation/suse101/>.

Novell AppArmor 2.0 Administration Guide

This guide contains in-depth information on the use of *AppArmor* in your environment. An online version of this document can be found at <http://www.novell.com/documentation/apparmor/>.

3 Documentation Conventions

The following typographical conventions are used in this manual:

- /etc/passwd: filenames and directory names
- *placeholder*: replace *placeholder* with the actual value
- PATH: the environment variable PATH
- ls, --help: commands, options, and parameters
- user: users or groups
- [Alt], [Alt] + [F1]: a key to press or a key combination; keys are shown in uppercase as on a keyboard
- *File*, *File* → *Save As*: menu items, buttons

- *Dancing Penguins* (Chapter Penguins, ↑*Reference*): This is a reference to a chapter in another book.

4 About the Making of This Manual

This book is written in Novdoc, a subset of DocBook (see <http://www.docbook.org>). The XML source files were validated by `xmllint`, processed by `xsltproc`, and converted into HTML using a customized version of Norman Walsh's stylesheets.

5 Acknowledgment

With a lot of voluntary commitment, the developers of Linux cooperate on a global scale to promote the development of Linux. We thank them for their efforts—this distribution would not exist without them. Furthermore, we thank Frank Zappa and Pawar. Special thanks, of course, go to Linus Torvalds.

Have a lot of fun!

Your SUSE Team

Part I. Advanced Deployment Scenarios

Remote Installation

SUSE Linux can be installed in several different ways. As well as the usual CD or DVD installation covered in Chapter *Installation with YaST* (↑Start-Up), you can choose from various network-based approaches or even take a completely hands-off approach to the installation of SUSE Linux.

Each method is introduced by means of two short check lists: one listing the prerequisites for this method and the other illustrating the basic procedure. More detail is then provided for all the techniques used in these installation scenarios.

NOTE

In the following sections, the system to hold your new SUSE Linux installation is referred to as *target system* or *installation target*. The term *installation source* is used for all sources of installation data. This includes physical media, such as CD and DVD, and network servers distributing the installation data in your network.

1.1 Installation Scenarios for Remote Installation

This section introduces the most common installation scenarios for remote installations. For each scenario, carefully check the list of prerequisites and follow the procedure outlined for this scenario. If in need of detailed instructions for a particular step, follow the links provided for each one of them.

IMPORTANT

The configuration of the X Window System is not part of any remote installation process. After the installation has finished, log in to the target system as root, enter `telinit 3`, and start SaX2 to configure the graphics hardware as described in [Section 14.1, “X11 Setup with SaX2”](#) (page 271).

1.1.1 Simple Remote Installation via VNC—Static Network Configuration

This type of installation still requires some degree of physical access to the target system to boot for installation. The installation itself is entirely controlled by a remote workstation using VNC to connect to the installation program. User interaction is required as with the manual installation in Chapter *Installation with YaST* (↑Start-Up).

For this type of installation, make sure that the following requirements are met:

- Remote installation source: NFS, HTTP, FTP, or SMB with working network connection
- Target system with working network connection
- Controlling system with working network connection and VNC viewer software or Java-enabled browser (Firefox, Konqueror, Internet Explorer, or Opera)
- Physical boot medium (CD, or DVD) for booting the target system
- Valid static IP addresses already assigned to the installation source and the controlling system
- Valid static IP address to assign to the target system

To perform this kind of installation, proceed as follows:

- 1 Set up the installation source as described in [Section 1.2, “Setting Up the Server Holding the Installation Sources”](#) (page 26).
- 2 Boot the target system using the first CD or DVD of the SUSE Linux media kit.

- When the boot screen of the target system appears, use the boot options prompt to set the appropriate VNC options and the address of the installation source. This is described in detail in [Section 1.4, “Booting the Target System for Installation”](#) (page 45).

The target system boots to a text-based environment, giving the network address and display number under which the graphical installation environment can be addressed by any VNC viewer application or browser. VNC installations announce themselves over OpenSLP and can be found using Konqueror in `service://` or `sdp://` mode.

- On the controlling workstation, open a VNC viewing application or Web browser and connect to the target system as described in [Section 1.5.1, “VNC Installation”](#) (page 49).
- Perform the installation as described in Chapter *Installation with YaST* (↑Start-Up) .

You will need to reconnect to the target system after it reboots for the final part of the installation.

- Finish the installation.

1.1.2 Simple Remote Installation via VNC—Dynamic Network Configuration via DHCP

This type of installation still requires some degree of physical access to the target system to boot for installation. The network configuration is made with DHCP. The installation itself is entirely controlled from a remote workstation using VNC to connect to the installer, but still requires user interaction for the actual configuration efforts.

For this type of installation, make sure that the following requirements are met:

- Remote installation source: NFS, HTTP, FTP, or SMB with working network connection
- Target system with working network connection

- Controlling system with working network connection and VNC viewer software or Java-enabled browser (Firefox, Konqueror, Internet Explorer, or Opera)
- Physical boot medium (CD, DVD, custom boot disk) for booting the target system
- Running DHCP server providing IP addresses

To perform this kind of installation, proceed as follows:

- 1 Set up the installation source as described in [Section 1.2, “Setting Up the Server Holding the Installation Sources”](#) (page 26). Choose an NFS, HTTP, or FTP network server. For a SMB installation source, refer to [Section 1.2.5, “Managing a SMB Installation Source”](#) (page 34).
- 2 Boot the target system using the first CD or DVD of the SUSE Linux media kit.
- 3 When the boot screen of the target system appears, use the boot options prompt to set the appropriate VNC options and the address of the installation source. This is described in detail in [Section 1.4, “Booting the Target System for Installation”](#) (page 45).

The target system boots to a text-based environment, giving the network address and display number under which the graphical installation environment can be addressed by any VNC viewer application or browser. VNC installations announce themselves over OpenSLP and can be found using Konqueror in `service://` or `sdp://` mode.

- 4 On the controlling workstation, open a VNC viewing application or Web browser and connect to the target system as described in [Section 1.5.1, “VNC Installation”](#) (page 49).
 - 5 Perform the installation as described in Chapter *Installation with YaST* (\uparrow Start-Up).
- You will need to reconnect to the target system after it reboots for the final part of the installation.
- 6 Finish the installation.

1.1.3 Remote Installation via VNC—PXE Boot and Wake on LAN

This type of installation is completely hands-off. The target machine is started and booted remotely. User interaction is only needed for the actual installation. This approach is suitable for cross-site deployments.

To perform this type of installation, make sure that the following requirements are met:

- Remote installation source: NFS, HTTP, FTP, or SMB with working network connection
- TFTP server
- Running DHCP server for your network
- Target system capable of PXE boot, networking, and Wake on LAN, plugged in and connected to the network
- Controlling system with working network connection and VNC viewer software or Java-enabled browser (Firefox, Konqueror, Internet Explorer, or Opera)

To perform this type of installation, proceed as follows:

- 1 Set up the installation source as described in [Section 1.2, “Setting Up the Server Holding the Installation Sources”](#) (page 26). Choose an NFS, HTTP, or FTP network server or configure a SMB installation source as described in [Section 1.2.5, “Managing a SMB Installation Source”](#) (page 34).
- 2 Set up a TFTP server to hold a boot image that can be pulled by the target system. This is described in [Section 1.3.2, “Setting Up a TFTP Server”](#) (page 37).
- 3 Set up a DHCP server to provide IP addresses to all machines and reveal the location of the TFTP server to the target system. This is described in [Section 1.3.1, “Setting Up a DHCP Server”](#) (page 36).
- 4 Prepare the target system for PXE boot. This is described in further detail in [Section 1.3.5, “Preparing the Target System for PXE Boot”](#) (page 43).

- 5** Initiate the boot process of the target system using Wake on LAN. This is described in [Section 1.3.7, “Wake on LAN”](#) (page 44).
- 6** On the controlling workstation, open a VNC viewing application or Web browser and connect to the target system as described in [Section 1.5.1, “VNC Installation”](#) (page 49).
- 7** Perform the installation as described in Chapter *Installation with YaST* ([↑Start-Up](#)) .

You will need to reconnect to the target system after it reboots for the final part of the installation.
- 8** Finish the installation.

1.1.4 Simple Remote Installation via SSH—Static Network Configuration

This type of installation still requires some degree of physical access to the target system to boot for installation and to determine the IP address of the installation target. The installation itself is entirely controlled from a remote workstation using SSH to connect to the installer. User interaction is required as with the regular installation described in Chapter *Installation with YaST* ([↑Start-Up](#)).

For this type of installation, make sure that the following requirements are met:

- Remote installation source: NFS, HTTP, FTP, or SMB with working network connection
- Target system with working network connection
- Controlling system with working network connection and working SSH client software.
- Physical boot medium (CD, DVD, custom boot disk) for the target system
- Valid static IP addresses already assigned to the installation source and the controlling system

- Valid static IP address to assign to the target system

To perform this kind of installation, proceed as follows:

- 1 Set up the installation source as described in [Section 1.2, “Setting Up the Server Holding the Installation Sources”](#) (page 26).
- 2 Boot the target system using the first CD or DVD of the SUSE Linux media kit.
- 3 When the boot screen of the target system appears, use the boot options prompt to set the appropriate parameters for network connection, address of the installation source, and SSH enablement. This is described in detail in [Section 1.4.3, “Using Custom Boot Options”](#) (page 47).

The target system boots to a text-based environment, giving the network address under which the graphical installation environment can be addressed by any SSH client.

- 4 On the controlling workstation, open a terminal window and connect to the target system as described in [Section “Connecting to the Installation Program”](#) (page 52).
- 5 Perform the installation as described in Chapter *Installation with YaST* (↑Start-Up).

You will need to reconnect to the target system after it reboots for the final part of the installation.

- 6 Finish the installation.

1.1.5 Simple Remote Installation via SSH—Dynamic Network Configuration via DHCP

This type of installation still requires some degree of physical access to the target system to boot for installation and determine the IP address of the installation target. The installation itself is entirely controlled from a remote workstation using VNC to connect to the installer, but still requires user interaction for the actual configuration efforts.

For this type of installation, make sure that the following requirements are met:

- Remote installation source: NFS, HTTP, FTP, or SMB with working network connection
- Target system with working network connection
- Controlling system with working network connection and working SSH client software.
- Physical boot medium (CD or DVD) for booting the target system
- Running DHCP server providing IP addresses

To perform this kind of installation, proceed as follows:

- 1 Set up the installation source as described in [Section 1.2, “Setting Up the Server Holding the Installation Sources”](#) (page 26). Choose an NFS, HTTP, or FTP network server. For a SMB installation source, refer to [Section 1.2.5, “Managing a SMB Installation Source”](#) (page 34).
- 2 Boot the target system using the first CD or DVD of the SUSE Linux media kit.
- 3 When the boot screen of the target system appears, use the boot options prompt to pass the appropriate parameters for network connection, location of the installation source, and SSH enablement. See [Section 1.4.3, “Using Custom Boot Options”](#) (page 47) for detailed instructions on the use of these parameters.

The target system boots to a text-based environment, giving you the network address under which the graphical installation environment can be addressed by any SSH client.

- 4 On the controlling workstation, open a terminal window and connect to the target system as described in [Section “Connecting to the Installation Program”](#) (page 52).
- 5 Perform the installation as described in Chapter *Installation with YaST* (\uparrow Start-Up).

You will need to reconnect to the target system after it reboots for the final part of the installation.

6 Finish the installation.

1.1.6 Remote Installation via SSH—PXE Boot and Wake on LAN

This type of installation is completely hands-off. The target machine is started and booted remotely.

To perform this type of installation, make sure that the following requirements are met:

- Remote installation source: NFS, HTTP, FTP, or SMB with working network connection
- TFTP server
- Running DHCP server for your network, providing a static IP to the host to install
- Target system capable of PXE boot, networking, and Wake on LAN, plugged in and connected to the network
- Controlling system with working network connection and SSH client software

To perform this type of installation, proceed as follows:

- 1** Set up the installation source as described in [Section 1.2, “Setting Up the Server Holding the Installation Sources”](#) (page 26). Choose an NFS, HTTP, or FTP network server. For the configuration of a SMB installation source, refer to [Section 1.2.5, “Managing a SMB Installation Source”](#) (page 34).
- 2** Set up a TFTP server to hold a boot image that can be pulled by the target system. This is described in [Section 1.3.2, “Setting Up a TFTP Server”](#) (page 37).
- 3** Set up a DHCP server to provide IP addresses to all machines and reveal the location of the TFTP server to the target system. This is described in [Section 1.3.1, “Setting Up a DHCP Server”](#) (page 36).
- 4** Prepare the target system for PXE boot. This is described in further detail in [Section 1.3.5, “Preparing the Target System for PXE Boot”](#) (page 43).

- 5** Initiate the boot process of the target system using Wake on LAN. This is described in [Section 1.3.7, “Wake on LAN”](#) (page 44).
 - 6** On the controlling workstation, start an SSH client and connect to the target system as described in [Section 1.5.2, “SSH Installation”](#) (page 51).
 - 7** Perform the installation as described in Chapter *Installation with YaST* (\uparrow Start-Up).
- You will need to reconnect to the target system it reboots for the final part of the installation.
- 8** Finish the installation.

1.2 Setting Up the Server Holding the Installation Sources

Depending on the operating system running on the machine to use as network installation source for SUSE Linux, there are several options for the server configuration. The easiest way to set up an installation server is to use YaST on SUSE LINUX Enterprise Server 9 or SUSE Linux 9.3 and higher. On other versions of SUSE LINUX Enterprise Server or SUSE Linux, set up the installation source manually.

TIP

You can even use a Microsoft Windows machine as installation server for your Linux deployment. See [Section 1.2.5, “Managing a SMB Installation Source”](#) (page 34) for details.

1.2.1 Setting Up an Installation Server Using YaST

YaST offers a graphical tool for creating network installation sources. It supports HTTP, FTP, and NFS network installation servers.

- 1** Log in as root to the machine that should act as installation server.

2 Start *YaST* → *Miscellaneous* → *Installation Server*.

3 Select *Server Configuration*.

4 Select the server type (HTTP, FTP, or NFS).

The selected server service is started automatically every time the system starts. If a service of the selected type is already running on your system and you want to configure it manually for the server, deactivate the automatic configuration of the server service with *Do Not Configure Any Network Services*. In both cases, define the directory in which the installation data should be made available on the server.

5 Configure the required server type.

This step relates to the automatic configuration of server services. It is skipped when automatic configuration is deactivated. Define an alias for the root directory of the FTP or HTTP server on which the installation data should be found. The installation source will later be located under

`ftp://Server-IP/Alias/Name` (FTP) or under

`http://Server-IP/Alias/Name` (HTTP). *Name* stands for the name of the installation source, which is defined in the following step. If you selected NFS in the previous step, define wild cards and exports options. The NFS server will be accessible under `nfs://Server-IP/Name`. Details of NFS and exports can be found in [Chapter 22, Sharing File Systems with NFS](#) (page 399).

6 Configure the installation source.

Before the installation media are copied to their destination, define the name of the installation source (ideally, an easily remembered abbreviation of the product and version). YaST allows providing ISO images of the media instead of copies of the installation CDs. If you want this, activate the relevant check box and specify the directory path under which the ISO files can be found locally. Depending on the product to distribute using this installation server, it might be that more add-on CDs or service pack CDs are required to install the product completely. If you activate *Prompt for Additional CDs*, YaST automatically reminds you to supply these media. To announce your installation server in the network via OpenSLP, activate the appropriate option.

TIP

Consider announcing your installation source via OpenSLP if your network setup supports this option. This saves you from entering the network installation path on every target machine. The target systems are just booted using the SLP boot option and will find the network installation source without any further configuration. For details on this option, refer to [Section 1.4, “Booting the Target System for Installation”](#) (page 45).

7 Upload the installation data.

The most lengthy step in configuring an installation server is copying the actual installation CDs. Insert the media in the sequence requested by YaST and wait for the copying procedure to end. When the sources have been fully copied, return to the overview of existing information sources and close the configuration by selecting *Finish*.

Your installation server is now fully configured and ready for service. It is automatically started every time the system is started. No further intervention is required. You only need to configure and start this service correctly by hand if you have deactivated the automatic configuration of the selected network service with YaST as an initial step.

To deactivate an installation source, select *Change* in the overview to reach a list of all available installation sources. Choose the entry to remove then select *Delete*. This delete procedure only relates to the deactivation of the server service. The installation data itself remains in the directory chosen. However, you can remove it manually.

If your installation server should provide the installation data for more than one product of product version, start the YaST installation server module and select *Configure* in the overview of existing installation sources to configure the new installation source.

1.2.2 Manual Setup of an NFS Installation Source

Setting up an NFS source for installation is basically done in two steps. In the first step, create the directory structure holding the installation data and copy the installation

media over to this structure. Second, export the directory holding the installation data to the network.

To create a directory holding the installation data, proceed as follows:

- 1** Log in as `root`.
- 2** Create a directory that should later hold all installation data and change into this directory. For example:

```
mkdir install/product/productversion  
cd install/product/productversion
```

Replace `product` with an abbreviation of the product name (in this case SUSE Linux) and `productversion` with a string that contains the product name and version.

- 3** For each CD contained in the media kit execute the following commands:

- a** Copy the entire content of the installation CD into the installation server directory:

```
cp -a /media/path_to_your_CD-ROM_drive .
```

Replace `path_to_your_CD-ROM_drive` with the actual path under which your CD or DVD drive is addressed. Depending on the type of drive used in your system, this can be `cdrom`, `cdrecorder`, `dvd`, or `dvdrecorder`.

- b** Rename the directory to the CD number:

```
mv path_to_your_CD-ROM_drive CDx
```

Replace `x` with the actual number of your CD.

To export the installation sources via NFS using YaST, proceed as follows:

- 1** Log in as `root`.
- 2** Start *YaST* → *Network Services* → *NFS Server*.
- 3** Select *Start NFS Server* and *Open Port in Firewall* and click *Next*.

- 4** Select *Add Directory* and enter the path to the directory holding the installation data. In this case, it is */productversion*.
- 5** Select *Add Host* and enter the hostnames of the machines to which to export the installation data. Instead of specifying hostnames here, you could also use wild cards, ranges of network addresses, or just the domain name of your network. Enter the appropriate export options or leave the default, which works fine in most setups. For more information about the syntax used in exporting NFS shares, read the `exports` man page.

- 6** Click *Finish*.

The NFS server holding the SUSE Linux installation sources is automatically started and integrated into the boot process.

If you prefer to manually export the installation sources via NFS instead of using the YaST NFS Server module, proceed as follows:

- 1** Log in as `root`.
- 2** Open the file `/etc(exports` and enter the following line:

```
/productversion *(ro,root_squash,sync)
```

This exports the directory `/productversion` to any host that is part of this network or to any host that can connect to this server. To limit the access to this server, use netmasks or domain names instead of the general wild card `*`. Refer to the `export` man page for details. Save and exit this configuration file.

- 3** To add the NFS service to the list of servers started during system boot, execute the following commands:

```
insserv /etc/init.d/nfsserver  
insserv /etc/init.d/portmap
```

- 4** Start the NFS server using the following command:

```
rcnfsserver start
```

If you need to change the configuration of your NFS server later, modify the configuration file and restart the NFS daemon with `rcnfsserver restart`.

Announcing the NFS server via OpenSLP makes its address known to all clients in your network.

1 Log in as `root`.

2 Enter the directory `/etc/slpxreg.d/`.

3 Create a configuration file called `install.suse.nfs.reg` containing the following lines:

```
# Register the NFS Installation Server
service:install.suse:nfs://$HOSTNAME/path_instsource/CD1,en,65535
description=NFS Installation Source
```

Replace `path_instsource` with the actual path to the installation source on your server.

4 Save this configuration file and start the OpenSLP daemon using the following command:

```
rcslpd start
```

For more information about OpenSLP, refer to the package documentation located under `/usr/share/doc/packages/openslp/` or refer to [Chapter 19, SLP Services in the Network](#) (page 363).

1.2.3 Manual Setup of an FTP Installation Source

Creating an FTP installation source is very similar to creating an NFS installation source. FTP installation sources can be announced over the network using OpenSLP as well.

1 Create a directory holding the installation sources as described in [Section 1.2.2, “Manual Setup of an NFS Installation Source”](#) (page 28).

2 Configure the FTP server to distribute the contents of your installation directory:

a Log in as `root` and install the package `pure-ftpd` (a lean FTP server) using the YaST package manager.

b Enter the FTP server root directory:

```
cd/srv/ftp
```

- c** Create a subdirectory holding the installation sources in the FTP root directory:

```
mkdir instsource
```

Replace *instsource* with the product name.

- d** Copy the contents of all installation CDs into the FTP server's root directory (similar to the procedure described in [Section 1.2.2, “Manual Setup of an NFS Installation Source”](#) (page 28), [Step 3](#) (page 29)).

Alternatively, mount the contents of the already existing installation repository into the change root environment of the FTP server:

```
mount --bind path_to_instsource /srv/ftp/instsource
```

Replace *path_to_instsource* and *instsource* with values matching your setup. If you need to make this permanent, add it to */etc/fstab*.

- e** Start pure-ftpd:

```
pure-ftpd &
```

- 3** Announce the installation source via OpenSLP, if this is supported by your network setup:

- a** Create a configuration file called *install.suse.ftp.reg* under */etc/slp/reg.d/* that contains the following lines:

```
# Register the FTP Installation Server
service:install.suse:ftp://$HOSTNAME/srv/ftp/instsource/CD1,en,65535
description=FTP Installation Source
```

Replace *instsource* with the actual name to the installation source directory on your server. The *service:* line should be entered as one continuous line.

- b** Save this configuration file and start the OpenSLP daemon using the following command:

```
rcslpd start
```

1.2.4 Manual Setup of an HTTP Installation Source

Creating an HTTP installation source is very similar to creating an NFS installation source. HTTP installation sources can be announced over the network using OpenSLP as well.

- 1** Create a directory holding the installation sources as described in [Section 1.2.2, “Manual Setup of an NFS Installation Source”](#) (page 28).
- 2** Configure the HTTP server to distribute the contents of your installation directory:

- a** Install the Web server Apache as described in [Section 26.1.2, “Installation”](#) (page 448).
- b** Enter the root directory of the HTTP server (`/srv/www/htdocs`) and create a subdirectory that will hold the installation sources:

```
mkdir instsource
```

Replace `instsource` with the product name.

- c** Create a symbolic link from the location of the installation sources to the root directory of the Web server (`/srv/www/htdocs`):

```
ln -s /path_instsource /srv/www/htdocs/instsource
```

- d** Modify the configuration file of the HTTP server (`/etc/apache2/default-server.conf`) to make it follow symbolic links. Replace the following line:

```
Options None
```

with

```
Options Indexes FollowSymLinks
```

- e** Reload the HTTP server configuration using `rcapache2 reload`.

- 3** Announce the installation source via OpenSLP, if this is supported by your network setup:

- a** Create a configuration file called `install.suse.http.reg` under `/etc/slp/reg.d/` that contains the following lines:

```
# Register the HTTP Installation Server  
service:install.suse:http://$HOSTNAME/srv/www/htdocs/instsource/CD1/,en,65535  
description=HTTP Installation Source
```

Replace `path_to_instsource` with the actual path to the installation source on your server. The `service :` line should be entered as one continuous line.

- b** Save this configuration file and start the OpenSLP daemon using `rcslpd restart`.

1.2.5 Managing a SMB Installation Source

Using SMB (Samba), you can import the installation sources from a Microsoft Windows server and start your Linux deployment even with no Linux machine around.

To set up an exported Windows Share holding your SUSE Linux installation sources, proceed as follows:

- 1** Log in to your Windows machine.
- 2** Start Explorer and create a new folder that will hold the entire installation tree and name it `INSTALL`, for example.
- 3** Export this share according the procedure outlined in your Windows documentation.
- 4** Enter this share and create a subfolder, called `product`. `product` needs to be replaced with the actual product name (SUSE Linux in this case).
- 5** Copy each SUSE Linux CD into a separate folder and name these folders `CD1`, `CD2`, `CD3`, etc.
- 6** Enter the top directory of the exported share (`INSTALL`, in this example) and copy the following files and folders from `product/CD1` to this folder: `content`, `media.1`, `control.xml`, and `boot`.

7 Create a new folder under `INSTALL` and name it `yast`.

Enter the `yast` folder and create the files `order` and `instorder`.

8 Open the `order` file and enter the following line:

```
/NLD/CD1 smb://user:password@hostname/productCD1
```

Replace `user` with the username you use on the Windows machine or use `Guest` to enable guest login to this share. `password` should be replaced either with your login password or any other string for guest login. `hostname` should be replaced with the network name of your Windows machine.

9 Open the `instorder` file and add the following line:

```
/product/CD1
```

To use a SMB mounted share as installation source, proceed as follows:

1 Boot the installation target.

2 Select *Installation*.

3 Press `[F3]` and `[F4]` for a selection of installation sources.

4 Choose SMB and enter the Windows machine's name or IP address, the share name (`INSTALL`, in this example), username, and password.

After you hit `[Enter]`, YaST starts and you can perform the installation.

1.3 Preparing the Boot of the Target System

This section covers the configuration tasks needed in complex boot scenarios. It contains ready-to-apply configuration examples for DHCP, PXE boot, TFTP, and Wake on LAN.

1.3.1 Setting Up a DHCP Server

The setup of a DHCP server on SUSE Linux is done by manually editing the appropriate configuration files. This section covers extending an existing DHCP server configuration to provide the data needed to serve in a TFTP, PXE, and WOL environment.

Manual Setup of a DHCP Server

All the DHCP server needs to do, apart from providing automatic address allocation to your network clients, is to announce the IP address of the TFTP server and the file that should be pulled in by the installation routines on the target machine.

- 1** Log in as `root` to the machine hosting the DHCP server.
- 2** Append the following lines to your DHCP server's configuration file located under `/etc/dhcpd.conf`:

```
group {  
    # PXE related stuff  
    #  
    # "next server" defines the tftp server that will be used  
    next server ip_tftp_server:  
    #  
    # "filename" specifies the pxelinux image on the tftp server  
    # the server runs in chroot under /srv/tftpboot  
    filename "pxelinux.0";  
}
```

Replace `ip_of_the_tftp_server` with the actual IP address of the TFTP server.

For more information about the options available in `dhcpd.conf`, refer to the `dhcpd.conf` manual page.

- 3** Restart the DHCP server by executing `rcdhcpd restart`.

If you plan on using SSH for the remote control of a PXE and Wake on LAN installation, explicitly specify the IP address DHCP should provide to the installation target. To achieve this, modify the above mentioned DHCP configuration according to the following example:

```
group {  
    # PXE related stuff
```

```

#
# "next server" defines the tftp server that will be used
next server ip_tftp_server;
#
# "filename" specifies the pxelinux image on the tftp server
# the server runs in chroot under /srv/tftpboot
filename "pxelinux.0";
host test { hardware ethernet mac_address;
            fixed-address some_ip_address; }
}

```

The host statement introduces the hostname of the installation target. To bind the hostname and IP address to a specific host, you have to know and specify the system's hardware (MAC) address. Replace all the variables used in this example with the actual values that match your environment.

After restarting the DHCP server, it provides a static IP to the host specified, enabling you to connect to the system via SSH.

1.3.2 Setting Up a TFTP Server

Set up a TFTP server with YaST or manually on any other Linux operating system that supports xinetd and tftp. The TFTP server delivers the boot image to the target system once it boots and sends a request for it.

Setting Up a TFTP Server Using YaST

- 1** Log in as root.
- 2** Start *YaST* → *Network Services* → *TFTP Server* and install the requested package.
- 3** Click *Enable* to make sure that the server is started and included in the boot routines. No further action from your side is required to secure this. xinetd starts tftpd at boot time.
- 4** Click *Open Port in Firewall* to open the appropriate port in the firewall running on your machine. If there is no firewall running on your server, this option is not available.
- 5** Click *Browse* to browse for the boot image directory.

The default directory `/tftpboot` is created and selected automatically.

- 6 Click *Finish* to apply your settings and start the server.

Manual Setup of a TFTP Server

- 1 Log in as `root` and install the packages `tftp` and `xinetd`.
- 2 If unavailable, create `/srv/tftpboot` and `/srv/tftpboot/pxelinux.cfg` directories.
- 3 Add the appropriate files needed for the boot image as described in [Section 1.3.3, “PXE Boot”](#) (page 38).
- 4 Modify the configuration of `xinetd` located under `/etc/xinetd.d/` to make sure that the `tftp` server is started on boot:

- a If it does not exist, create a file called `tftp` under this directory with `touch tftp`. Then run `chmod 755 tftp`.
- b Open the file `tftp` and add the following lines:

```
service tftp
{
    socket_type      = dgram
    protocol         = udp
    wait             = yes
    user             = root
    server           = /usr/sbin/in.tftpd
    server_args      = -s /tftpboot
    disable          = no
}
```

- c Save the file and restart `xinetd` with `rcxinetd restart`.

1.3.3 PXE Boot

Some technical background information as well as PXE's complete specifications are available in the Preboot Execution Environment (PXE) Specification ([ftp://download.intel.com/labs/manage/wfm/download/pxespec.pdf](http://download.intel.com/labs/manage/wfm/download/pxespec.pdf)).

- 1 Change to the directory of your installation repository and copy the `linux`, `initrd`, `message`, and `memtest` files to the `/srv/tftpboot` directory by entering the following:

```
cp -a boot/loader/linux boot/loader/initrd  
boot/loader/message boot/loader/memtest /srv/tftpboot
```

- 2 Install the `syslinux` package directly from your installation CDs or DVDs with YaST.

- 3 Copy the `/usr/share/syslinux/pxelinux.0` file to the `/srv/tftpboot` directory by entering the following:

```
cp -a /usr/share/syslinux/pxelinux.0 /srv/tftpboot
```

- 4 Change to the directory of your installation repository and copy the `isolinux.cfg` file to `/srv/tftpboot/pxelinux.cfg/default` by entering the following:

```
cp -a boot/loader/isolinux.cfg /srv/tftpboot/pxelinux.cfg/default
```

- 5 Edit the `/srv/tftpboot/pxelinux.cfg/default` file and remove the lines beginning with `gfxboot`, `readinfo`, and `framebuffer`.

- 6 Insert the following entries in the append lines of the default `failsafe` and `apic` labels:

`insmod=e100`

By means of this entry, the kernel module for an Intel 100MBit/s network card is loaded on the PXE clients. This entry depends on the client's hardware and must be adapted accordingly. In the case of a Broadcom GigaBit network card, this entry should read `insmod=bcm5700`.

`netdevice=eth0`

This entry defines the client's network interface that must be used for the network installation. It is only necessary if the client is equipped with several network cards and must be adapted accordingly. In case of a single network card, this entry can be omitted.

`install=nfs://ip_instserver/path_instsource/CD1`

This entry defines the NFS server and the installation source for the client installation. Replace `ip_instserver` with the actual IP address of your

installation server. *path_instsource* should be replaced with the actual path to the installation sources. HTTP, FTP, or SMB sources are addressed in a similar manner, except for the protocol prefix, which should read http, ftp, or smb.

IMPORTANT

If you need to pass other boot options to the installation routines, such as SSH or VNC boot parameters, append them to the `install` entry. An overview of parameters and some examples are given in [Section 1.4, “Booting the Target System for Installation” \(page 45\)](#).

An example `/srv/tftpboot/pxelinux.cfg/default` file follows. Adjust the protocol prefix for the installation source to match your network setup and specify your preferred method of connecting to the installer by adding the `vnc` and `vncpassword` or the `ssh` and `sshpassword` options to the `install` entry. The lines separated by `\` must be entered as one continuous line without a line break and without the `\`.

```
default linux

# default
label linux
kernel linux
append initrd=initrd ramdisk_size=65536 insmod=e100 \
install=nfs://ip_instserver/path_instsource/product

# failsafe
label failsafe
kernel linux
append initrd=initrd ramdisk_size=65536 ide=nodma apm=off acpi=off \
insmod=e100 install=nfs://ip_instserver/path_instsource/product

# apic
label apic
kernel linux
append initrd=initrd ramdisk_size=65536 apic insmod=e100 \
install=nfs://ip_instserver/path_instsource/product

# manual
label manual
kernel linux
append initrd=initrd ramdisk_size=65536 manual=1

# rescue
label rescue
kernel linux
```

```

append initrd=initrd ramdisk_size=65536 rescue=1

# memory test
label memtest
kernel memtest

# hard disk
label harddisk
kernel
linux append SLX=0x202

implicit      0
display       message
prompt        1
timeout       100

```

Replace *ip_instserver* and *path_instsource* with the values used in your setup.

The following section serves as a short reference to the PXELINUX options used in this setup. More information about the options available can be found in the documentation of the `syslinux` package located under `/usr/share/doc/packages/syslinux/`.

1.3.4 PXELINUX Configuration Options

The options listed here are a subset of all the options available for the PXELINUX configuration file.

DEFAULT *kernel options...*

Sets the default kernel command line. If PXELINUX boots automatically, it acts as if the entries after DEFAULT had been typed in at the boot prompt, except the auto option is automatically added, indicating an automatic boot.

If no configuration file is present or no DEFAULT entry is present in the configuration file, the default is the kernel name “linux” with no options.

APPEND *options...*

Add one or more options to the kernel command line. These are added for both automatic and manual boots. The options are added at the very beginning of the kernel command line, usually permitting explicitly entered kernel options to override them.

```
LABEL label KERNEL image APPEND options...
```

Indicates that if *label* is entered as the kernel to boot, PXELINUX should instead boot *image* and the specified APPEND options should be used instead of the ones specified in the global section of the file (before the first LABEL command). The default for *image* is the same as *label* and, if no APPEND is given, the default is to use the global entry (if any). Up to 128 LABEL entries are permitted.

Note that GRUB uses the following syntax:

```
title mytitle  
kernel my_kernel my_kernel_options  
initrd myinitrd
```

while PXELINUX uses the following syntax:

```
label mylabel  
kernel mykernel  
append myoptions
```

Labels are mangled as if they were filenames and they must be unique after mangling. For example, the two labels “v2.1.30” and “v2.1.31” would not be distinguishable under PXELINUX because both mangle to the same DOS filename.

The kernel does not have to be a Linux kernel; it can be a boot sector or a COM-BOOT file.

APPEND -

Append nothing. APPEND with a single hyphen as argument in a LABEL section can be used to override a global APPEND.

```
LOCALBOOT type
```

On PXELINUX, specifying LOCALBOOT 0 instead of a KERNEL option means invoking this particular label and causes a local disk boot instead of a kernel boot.

Argument	Description
0	Perform a normal boot
4	Performs a local boot with the Universal Network Driver Interface (UNDI) driver still resident in memory

Argument	Description
5	Performs a local boot with the entire PXE stack, including the UNDI driver, still resident in memory

All other values are undefined. If you do not know what the UNDI or PXE stacks are, specify 0.

TIMEOUT *time-out*

Indicates how long to wait at the boot prompt until booting automatically, in units of 1/10 second. The time-out is cancelled as soon as the user types anything on the keyboard, the assumption being that the user completes the command begun. A time-out of zero disables the time-out completely (this is also the default).

The maximum possible time-out value is 35996 (just less than one hour).

PROMPT *flag_val*

If *flag_val* is 0, displays the boot prompt only if [Shift] or [Alt] is pressed or [Caps Lock] or [Scroll lock] is set (this is the default). If *flag_val* is 1, always displays the boot prompt.

```
F2  filename
F1  filename
...etc...
F9  filename
F10filename
```

Displays the indicated file on the screen when a function key is pressed at the boot prompt. This can be used to implement preboot online help (presumably for the kernel command line options.) For backward compatibility with earlier releases, [F10] can be also entered as [F0]. Note that there is currently no way to bind filenames to [F11] and [F12].

1.3.5 Preparing the Target System for PXE Boot

Prepare the system's BIOS for PXE boot by including the PXE option in the BIOS boot order.

WARNING

Do not place the PXE option ahead of the hard disk boot option in the BIOS. Otherwise this system would try to reinstall itself every time you boot it.

1.3.6 Preparing the Target System for Wake on LAN

Wake on LAN (WOL) requires the appropriate BIOS option to be enabled prior to the installation. Also, note down the MAC address of the target system. This data is needed to initiate Wake on LAN.

1.3.7 Wake on LAN

Wake on LAN allows a machine to be powered on via a special network packet that is sent containing the machine's MAC address. Because every machine in the world has a unique MAC identifier, you do not need to worry about accidentally powering on the wrong machine.

IMPORTANT

If the controlling machine is not located in the same network segment as the installation target that should be awakened, either configure the WOL requests to be sent as multicasts or remotely control a machine on that network segment to act as the sender of these requests.

1.3.8 Manual Wake on LAN

- 1** Log in as `root`.
- 2** Start *YaST → Software Management* and install the package `netdiag`.

- 3** Open a terminal and enter the following command as `root` to wake the target:

```
ether-wakemac_of_target
```

Replace `mac_of_target` with the actual MAC address of the target.

1.4 Booting the Target System for Installation

Basically, there are two different ways to customize the boot process for installation apart from those mentioned under [Section 1.3.7, “Wake on LAN”](#) (page 44) and [Section 1.3.3, “PXE Boot”](#) (page 38). You can either use the default boot options and F keys or use the boot options prompt of the installation boot screen to pass any boot options that the installation kernel might need on this particular hardware.

1.4.1 Using the Default Boot Options

The boot options have already been described in detail in Chapter *Installation with YaST* (↑Start-Up).

Generally, just selecting *Installation* starts the installation boot process. If problems occur, use *Installation—ACPI Disabled* or *Installation—Safe Settings*.

For more information about troubleshooting the installation process, refer to Section “*Installation Problems*” (Chapter 9, *Common Problems and Their Solutions*, ↑Start-Up).

1.4.2 Using the F Keys

The menu bar at the bottom screen offers some advanced functionality needed in some setups. Using the F keys, you can specify additional options to pass to the installation routines without having to know the detailed syntax of these parameters you would need if you entered them as boot options (see [Section 1.4.3, “Using Custom Boot Options”](#) (page 47)).

See the table below for a complete set of the options available.

Table 1.1 F Keys During Installation

Key	Purpose	Available Options	Default Value
[F1]	Provide help	None	None
[F2]	Select the installation language	All supported languages	English
[F3]	Change screen resolution for installation	<ul style="list-style-type: none">• Text mode• VESA• resolution #1• resolution #2• ...	<ul style="list-style-type: none">• Default value depends on your graphics hardware
[F4]	Select the installation source	<ul style="list-style-type: none">• CD-ROM/DVD• SLP• FTP• HTTP• NFS• SMB• Hard Disk	CD-ROM/DVD
[F5]	Apply driver update disk	Driver	None

1.4.3 Using Custom Boot Options

Using the appropriate set of boot options helps facilitate your installation procedure. Many parameters can also be configured later using the `linuxrc` routines, but using the boot options is easier. In some automated setups, the boot options can be provided with `initrd` or an `info` file.

The following table lists all installation scenarios mentioned in this chapter with the required parameters for booting and the corresponding boot options. Just append all of them in the order they appear in this table to get one boot option string that is handed to the installation routines. For example (all in one line):

```
install=... netdevice=... hostip=... netmask=... vnc=... vncpassword=...
```

Replace all the values (...) in this string with the values appropriate for your setup.

Table 1.2 Installation (Boot) Scenarios Used in This Chapter

Installation Scenario	Parameters Needed for Booting	Boot Options
Chapter <i>Installation with YaST</i> (↑Start-Up)	None: system boots automatically	None needed
Section 1.1.1, “Simple Remote Installation via VNC—Static Network Configuration” (page 18)	<ul style="list-style-type: none">• Location of the installation server• Network device• IP address• Netmask• Gateway• VNC enablement• VNC password	<ul style="list-style-type: none">• <code>install=(nfs,http,ftp,smb)://path_to_instmedia</code>• <code>netdevice=some_netdevice</code> (only needed if several network devices are available)• <code>hostip=some_ip</code>• <code>netmask=some_netmask</code>• <code>gateway=ip_gateway</code>• <code>vnc=1</code>• <code>vncpassword=some_password</code>

Installation Scenario	Parameters Needed for Booting	Boot Options
Section 1.1.2, “Simple Remote Installation via VNC—Dynamic Network Configuration via DHCP” (page 19)	<ul style="list-style-type: none"> • Location of the installation server • VNC enablement • VNC password 	<ul style="list-style-type: none"> • <code>install=(nfs,http,ftp,smb)://path_to_instmedia</code> • <code>vnc=1</code> • <code>vncpassword=some_password</code>
Section 1.1.3, “Remote Installation via VNC—PXE Boot and Wake on LAN” (page 21)	<ul style="list-style-type: none"> • Location of the installation server • Location of the TFTP server • VNC enablement • VNC password 	Not applicable; process managed through PXE and DHCP
Section 1.1.4, “Simple Remote Installation via SSH—Static Network Configuration” (page 22)	<ul style="list-style-type: none"> • Location of the installation server • Network device • IP address • Netmask • Gateway • SSH enablement • SSH password 	<ul style="list-style-type: none"> • <code>install=(nfs,http,ftp,smb)://path_to_instmedia</code> • <code>netdevice=some_netdevice</code>(only needed if several network devices are available) • <code>hostip=some_ip</code> • <code>netmask=some_netmask</code> • <code>gateway=ip_gateway</code> • <code>usessh=1</code> • <code>sshpassword=some_password</code>
Section 1.1.5, “Simple Remote Installation via	<ul style="list-style-type: none"> • Location of the installation server 	<ul style="list-style-type: none"> • <code>install=(nfs,http,ftp,smb)://path_to_instmedia</code>

Installation Scenario	Parameters Needed for Booting	Boot Options
SSH—Dynamic Network Configuration via DHCP (page 23)	<ul style="list-style-type: none">• SSH enablement• SSH password	<ul style="list-style-type: none">• <code>usessh=1</code>• <code>sshpassword=some_password</code>
Section 1.1.6, “Remote Installation via SSH—PXE Boot and Wake on LAN” (page 25)	<ul style="list-style-type: none">• Location of the installation server• Location of the TFTP server• SSH enablement• SSH password	Not applicable; process managed through PXE and DHCP

TIP

Find more information about the `linuxrc` boot options used for booting a Linux system in `/usr/share/doc/packages/linuxrc/linuxrc.html`.

1.5 Monitoring the Installation Process

There are several options for remotely monitoring the installation process. If the proper boot options have been specified while booting for installation, either VNC or SSH can be used to control the installation and system configuration from a remote workstation.

1.5.1 VNC Installation

Using any VNC viewer software, you can remotely control the installation of SUSE Linux from virtually any operating system. This section introduces the setup using a VNC viewer application or a Web browser.

Preparing for VNC Installation

All you need to do on the installation target to prepare for a VNC installation is to provide the appropriate boot options at the initial boot for installation (see [Section 1.4.3, “Using Custom Boot Options”](#) (page 47)). The target system boots into a text-based environment and waits for a VNC client to connect to the installation program.

The installation program announces the IP address and display number needed to connect for installation. If you have physical access to the target system, this information is provided right after the system booted for installation. Enter this data when your VNC client software prompts for it and provide your VNC password.

Because the installation target announces itself via OpenSLP, you can retrieve the address information of the installation target via an SLP browser without the need for any physical contact to the installation itself provided your network setup and all machines support OpenSLP:

- 1 Start the KDE file and Web browser Konqueror.
- 2 Enter `service://yast.installation.suse` in the location bar.

The target system then appears as an icon in the Konqueror screen. Clicking this icon launches the KDE VNC viewer in which to perform the installation. Alternatively, run your VNC viewer software with the IP address provided and add `:1` at the end of the IP address for the display the installation is running on.

Connecting to the Installation Program

Basically, there are two ways to connect to a VNC server (the installation target in this case). You can either start an independent VNC viewer application on any operating system or connect using a Java-enabled Web browser.

Using VNC, you can control the installation of a Linux system from any other operating system, including other Linux flavors, Windows, or Mac OS.

On a Linux machine, make sure that the package `tightvnc` is installed. On a Windows machine, install the Windows port of this application, which can be obtained at the TightVNC home page (<http://www.tightvnc.com/download.html>).

To connect to the installation program running on the target machine, proceed as follows:

- 1** Start the VNC viewer.
- 2** Enter the IP address and display number of the installation target as provided by the SLP browser or the installation program itself:

ip_address:display_number

A window opens on your desktop displaying the YaST screens as in a normal local installation.

Using a Web browser to connect to the installation program makes you totally independent of any VNC software or the underlying operating system. As long as the browser application has Java support enabled, you can use any browser (Firefox, Internet Explorer, Konqueror, Opera, etc.) to perform the installation of your Linux system.

To perform a VNC installation, proceed as follows:

- 1** Launch your preferred Web browser.
- 2** Enter the following at the address prompt:

`http://ip_address_of_target:5801`

- 3** Enter your VNC password when prompted to do so. The browser window now displays the YaST screens as in a normal local installation.

1.5.2 SSH Installation

Using SSH, you can remotely control the installation of your Linux machine using any SSH client software.

Preparing for SSH Installation

Apart from installing the appropriate software package (OpenSSH for Linux and PuTTY for Windows), you just need to pass the appropriate boot options to enable SSH for installation. See [Section 1.4.3, “Using Custom Boot Options”](#) (page 47) for details. OpenSSH is installed by default on any SUSE Linux based operating system.

Connecting to the Installation Program

- 1 Retrieve the installation target's IP address.

If you have physical access to the target machine, just take the IP address the installation routine provides at the console after the initial boot. Otherwise take the IP address that has been assigned to this particular host in the DHCP server configuration.

- 2 At a command line enter the following command:

```
ssh -X root@ip_address_of_target
```

Replace *ip_address_of_target* with the actual IP address of the installation target.

- 3 When prompted for a username, enter `root`.
- 4 When prompted for password, enter the password that has been set via the SSH boot option.

After you have successfully authenticated, a command line prompt for the installation target appears.

- 5 Enter `yast` to launch the installation program.

A window opens showing the normal YaST screens as described in Chapter *Installation with YaST* (↑Start-Up).

Advanced Disk Setup

Sophisticated system configurations require particular disk setups. To get persistent device naming with SCSI devices, use a specific start-up script or udev. Logical Volume Management (LVM) is a disk partitioning scheme that is designed to be much more flexible than the physical partitioning used in standard setups. Its snapshot functionality enables you to create data backups easily. Redundant Array of Independent Disks (RAID) offers increased data integrity, performance, and fault tolerance.

2.1 LVM Configuration

This section briefly describes the principles behind LVM and its basic features that make it useful under many circumstances. In [Section 2.1.2, “LVM Configuration with YaST”](#) (page 55), learn how to set up LVM with YaST.

WARNING

Using LVM might be associated with increased risk, such as data loss. Risks also include application crashes, power failures, and faulty commands. Save your data before implementing LVM or reconfiguring volumes. Never work without a backup.

2.1.1 The Logical Volume Manager

The Logical Volume Manager (LVM) enables flexible distribution of hard disk space over several file systems. It was developed because sometimes the need to change the

segmentation of hard disk space arises only after the initial partitioning during installation has already been done. Because it is difficult to modify partitions on a running system, LVM provides a virtual pool (volume group, VG for short) of memory space from which logical volumes (LVs) can be created as needed. The operating system accesses these LVs instead of the physical partitions. Volume groups can span more than only one disk so that several disks or parts of them may constitute one single VG. This way, LVM provides a kind of abstraction from the physical disk space that allows its segmentation to be changed in a much easier and safer way than physical repartitioning does. Background information regarding physical partitioning can be found in Section “Partition Types” (Chapter 1, *Installation with YaST, ↑Start-Up*) and Section “Partitioner” (Chapter 2, *System Configuration with YaST, ↑Start-Up*).

Figure 2.1 Physical Partitioning versus LVM

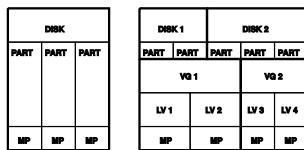


Figure 2.1, “Physical Partitioning versus LVM” (page 54) compares physical partitioning (left) with LVM segmentation (right). On the left side, one single disk has been divided into three physical partitions (PART), each with a mount point (MP) assigned so that the operating system can access them. On the right side, two disks have been divided into two and three physical partitions each. Two LVM volume groups (VG 1 and VG 2) have been defined. VG 1 contains two partitions from DISK 1 and one from DISK 2. VG 2 contains the remaining two partitions from DISK 2. In LVM, the physical disk partitions that are incorporated in a volume group are called physical volumes (PVs). Within the volume groups, four logical volumes (LV 1 through LV 4) have been defined, which can be used by the operating system via the associated mount points. The border

between different logical volumes need not be aligned with any partition border. See the border between LV 1 and LV 2 in this example.

LVM features:

- Several hard disks or partitions can be combined in a large logical volume.
- Provided the configuration is suitable, an LV (such as `/usr`) can be enlarged when the free space is exhausted.
- Using LVM, even add hard disks or LVs in a running system. However, this requires hot-swappable hardware that is capable of such actions.
- It is possible to activate a "striping mode" that distributes the data stream of a logical volume over several physical volumes. If these physical volumes reside on different disks, this can improve the reading and writing performance just like RAID 0.
- The snapshot feature enables consistent backups (especially for servers) in the running system.

With these features, using LVM already makes sense for heavily used home PCs or small servers. If you have a growing data stock, as in the case of databases, music archives, or user directories, LVM is just the right thing for you. This would allow file systems that are larger than the physical hard disk. Another advantage of LVM is that up to 256 LVs can be added. However, keep in mind that working with LVM is different from working with conventional partitions. Instructions and further information about configuring LVM is available in the official LVM HOWTO at <http://tldp.org/HOWTO/LVM-HOWTO/>.

Starting from kernel version 2.6, LVM version 2 is available, which is downward-compatible with the previous LVM and enables the continued management of old volume groups. When creating new volume groups, decide whether to use the new format or the downward-compatible version. LVM 2 does not require any kernel patches. It makes use of the device mapper integrated in kernel 2.6. This kernel only supports LVM version 2. Therefore, when talking about LVM, this section always refers to LVM version 2.

2.1.2 LVM Configuration with YaST

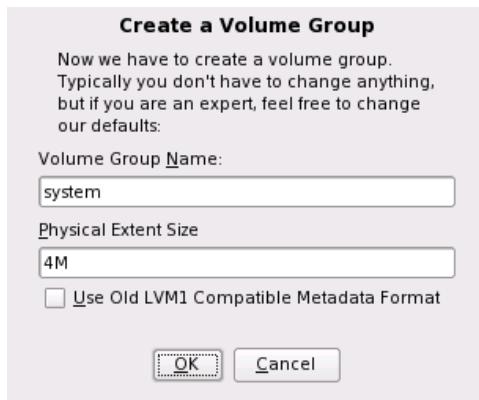
The YaST LVM configuration can be reached from the YaST Expert Partitioner (see Section “Partitioner” (Chapter 2, *System Configuration with YaST*, ↑Start-Up)). This

partitioning tool enables you to edit and delete existing partitions and create new ones that should be used with LVM. There, create an LVM partition by first clicking *Create* → *Do not format* then selecting *0x8E Linux LVM* as the partition identifier. After creating all the partitions to use with LVM, click *LVM* to start the LVM configuration.

Creating Volume Groups

If no volume group exists on your system yet, you are prompted to add one (see [Figure 2.2, “Creating a Volume Group”](#) (page 56)). It is possible to create additional groups with *Add group*, but usually one single volume group is sufficient. `system` is suggested as a name for the volume group in which the SUSE Linux system files are located. The physical extent size defines the size of a physical block in the volume group. All the disk space in a volume group is handled in chunks of this size. This value is normally set to 4 MB and allows for a maximum size of 256 GB for physical and logical volumes. The physical extent size should only be increased, for example, to 8, 16, or 32 MB, if you need logical volumes larger than 256 GB.

Figure 2.2 Creating a Volume Group

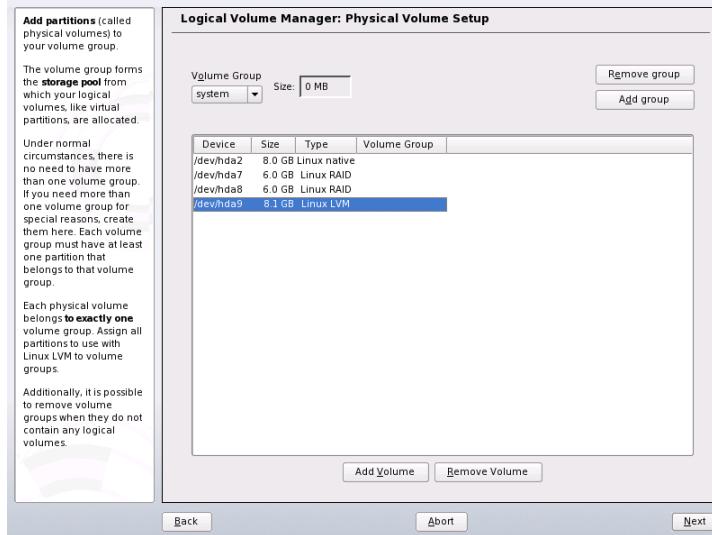


Configuring Physical Volumes

Once a volume group has been created, the following dialog lists all partitions with either the “Linux LVM” or “Linux native” type. No swap or DOS partitions are shown. If a partition is already assigned to a volume group, the name of the volume group is shown in the list. Unassigned partitions are indicated with “--”.

If there are several volume groups, set the current volume group in the selection box to the upper left. The buttons in the upper right enable creation of additional volume groups and deletion of existing volume groups. Only volume groups that do not have any partitions assigned can be deleted. All partitions that are assigned to a volume group are also referred to as a physical volumes (PV).

Figure 2.3 Physical Volume Setup



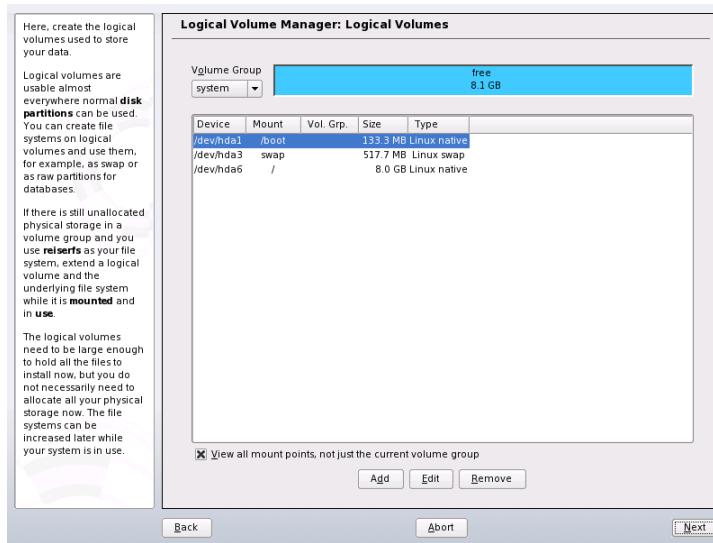
To add a previously unassigned partition to the selected volume group, first click the partition then *Add Volume*. At this point, the name of the volume group is entered next to the selected partition. Assign all partitions reserved for LVM to a volume group. Otherwise, the space on the partition remains unused. Before exiting the dialog, every volume group must be assigned at least one physical volume. After assigning all physical volumes, click *Next* to proceed to the configuration of logical volumes.

Configuring Logical Volumes

After the volume group has been filled with physical volumes, define the logical volumes the operating system should use in the next dialog. Set the current volume group in a selection box to the upper left. Next to it, the free space in the current volume group is shown. The list below contains all logical volumes in that volume group. All normal Linux partitions to which a mount point is assigned, all swap partitions, and all already

existing logical volumes are listed here. *Add*, *Edit*, and *Remove* logical volumes as needed until all space in the volume group has been exhausted. Assign at least one logical volume to each volume group.

Figure 2.4 Logical Volume Management

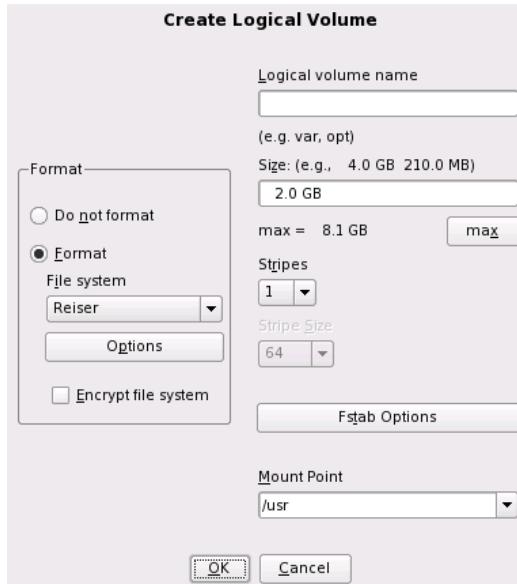


To create a new logical volume, click *Add* and fill out the pop-up that opens. As for partitioning, enter the size, file system, and mount point. Normally, a file system, such as reiserfs or ext2, is created on a logical volume and is then designated a mount point. The files stored on this logical volume can be found at this mount point on the installed system. Additionally it is possible to distribute the data stream in the logical volume among several physical volumes (striping). If these physical volumes reside on different hard disks, this generally results in a better reading and writing performance (like RAID 0). However, a striping LV with n stripes can only be created correctly if the hard disk space required by the LV can be distributed evenly to n physical volumes. If, for example, only two physical volumes are available, a logical volume with three stripes is impossible.

WARNING: Striping

YaST has no chance at this point to verify the correctness of your entries concerning striping. Any mistake made here is apparent only later when the LVM is implemented on disk.

Figure 2.5 Creating Logical Volumes



If you have already configured LVM on your system, the existing logical volumes can be entered now. Before continuing, assign appropriate mount points to these logical volumes too. With *Next*, return to the YaST Expert Partitioner and finish your work there.

Direct LVM Management

If you already have configured LVM and only want to change something, there is an alternative way to do that. In the YaST Control Center, select *System → LVM*. Basically this dialog allows the same actions as described above with the exception of physical partitioning. It shows the existing physical volumes and logical volumes in two lists and you can manage your LVM system using the methods already described.

2.2 Soft RAID Configuration

The purpose of RAID (redundant array of independent disks) is to combine several hard disk partitions into one large *virtual* hard disk to optimize performance, data security, or both. Most RAID controllers use the SCSI protocol because it can address a larger number of hard disks in a more effective way than the IDE protocol and is more suitable for parallel processing of commands. There are some RAID controllers that support IDE or SATA hard disks. Soft RAID provides the advantages of RAID systems without the additional cost of hardware RAID controllers. However, this requires some CPU time and has memory requirements that make it unsuitable for real high performance computers.

2.2.1 RAID Levels

SUSE Linux offers the option of combining several hard disks into one soft RAID system with the help of YaST—a very reasonable alternative to hardware RAID. RAID implies several strategies for combining several hard disks in a RAID system, each of them having different goals, advantages and characteristics. These variations are commonly known as *RAID levels*.

Common RAID levels are:

RAID 0

This level improves the performance of your data access by spreading out blocks of each file across multiple disk drives. Actually, this is not really a RAID, because it does not provide data backup, but the name *RAID 0* for this type of system has become the norm. With RAID 0, two or more hard disks are pooled together. The performance is very good, but the RAID system is destroyed and your data lost if even one hard disk fails.

RAID 1

This level provides adequate security for your data, because the data is copied to another hard disk 1:1. This is known as *hard disk mirroring*. If a disk is destroyed, a copy of its contents is available on another one. All of them except one could be damaged without endangering your data. However, if damage is not detected, it also may happen that damaged data is mirrored to the correct disk and data corruption happens that way. The writing performance suffers a little in the copying process compared to when using single disk access (ten to twenty percent slower), but read

access is significantly faster in comparison to any one of the normal physical hard disks, because the data is duplicated so can be parallel scanned. Generally it can be said that Level 1 provides nearly twice the read transaction rate of single disks and almost the same write transaction rate as single disks.

RAID 2 and RAID 3

These are not typical RAID implementations. Level 2 stripes data at the bit level rather than the block level. Level 3 provides byte-level striping with a dedicated parity disk and cannot service simultaneous multiple requests. Both levels are only rarely used.

RAID 4

Level 4 provides block-level striping just like Level 0 combined with a dedicated parity disk. In the case of a data disk failure, the parity data is used to create a replacement disk. However, the parity disk may create a bottleneck for write access. Nevertheless, Level 4 is sometimes used.

RAID 5

RAID 5 is an optimized compromise between Level 0 and Level 1 in terms of performance and redundancy. The hard disk space equals the number of disks used minus one. The data is distributed over the hard disks as with RAID 0. *Parity blocks*, created on one of the partitions, are there for security reasons. They are linked to each other with XOR, enabling the contents to be reconstructed by the corresponding parity block in case of system failure. With RAID 5, no more than one hard disk can fail at the same time. If one hard disk fails, it must be replaced as soon as possible to avoid the risk of losing data.

Other RAID Levels

Several other RAID levels have been developed (RAIDn, RAID 10, RAID 0+1, RAID 30, RAID 50, etc.), some of them being proprietary implementations created by hardware vendors. These levels are not very widespread, so are not explained here.

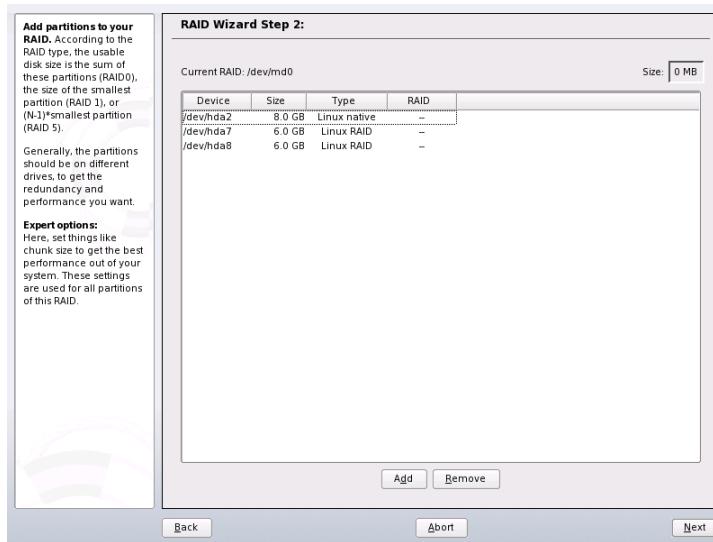
2.2.2 Soft RAID Configuration with YaST

The YaST soft RAID configuration can be reached from the YaST Expert Partitioner, described in Section “Partitioner” (Chapter 2, *System Configuration with YaST*, ↑Start-Up). This partitioning tool enables you to edit and delete existing partitions and create new ones that should be used with soft RAID. There, create RAID partitions by first

clicking *Create* → *Do not format* then selecting *0xFD Linux RAID* as the partition identifier. For RAID 0 and RAID 1, at least two partitions are needed—for RAID 1, usually exactly two and no more. If RAID 5 is used, at least three partitions are required. It is recommended to take only partitions of the same size. The RAID partitions should be stored on different hard disks to decrease the risk of losing data if one is defective (RAID 1 and 5) and to optimize the performance of RAID 0. After creating all the partitions to use with RAID, click *RAID* → *Create RAID* to start the RAID configuration.

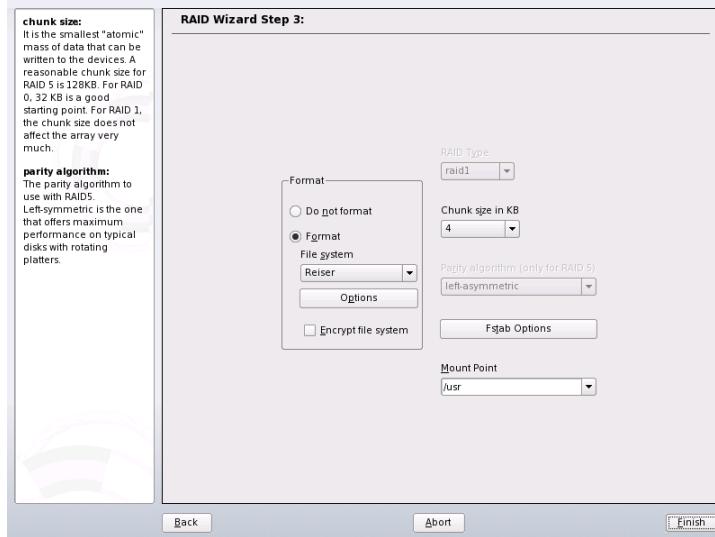
In the next dialog, choose between RAID levels 0, 1, and 5 (see [Section 2.2.1, “RAID Levels”](#) (page 60) for details). After *Next* is clicked, the following dialog lists all partitions with either the “Linux RAID” or “Linux native” type (see [Figure 2.6, “RAID Partitions”](#) (page 62)). No swap or DOS partitions are shown. If a partition is already assigned to a RAID volume, the name of the RAID device (e.g., `/dev/md0`) is shown in the list. Unassigned partitions are indicated with “--”.

Figure 2.6 RAID Partitions



To add a previously unassigned partition to the selected RAID volume, first click the partition then *Add*. At this point, the name of the RAID device is entered next to the selected partition. Assign all partitions reserved for RAID. Otherwise, the space on the partition remains unused. After assigning all partitions, click *Next* to proceed to the settings dialog where you can fine-tune the performance (see [Figure 2.7, “File System Settings”](#) (page 63)).

Figure 2.7 File System Settings



As with conventional partitioning, set the file system to use as well as encryption and the mount point for the RAID volume. Checking *Persistent Superblock* ensures that the RAID partitions are recognized as such when booting. After completing the configuration with *Finish*, see the `/dev/md0` device and others indicated with *RAID* in the expert partitioner.

2.2.3 Troubleshooting

Check the file `/proc/mdstats` to find out whether a RAID partition has been destroyed. In the event of a system failure, shut down your Linux system and replace the defective hard disk with a new one partitioned the same way. Then restart your system and enter the command `mdadm /dev/mdX --add /dev/sdX`. Replace 'X' with your particular device identifiers. This integrates the hard disk automatically into the RAID system and fully reconstructs it.

2.2.4 For More Information

Configuration instructions and more details for soft RAID can be found in the HOWTOs at:

- /usr/share/doc/packages/raidtools/Software-RAID.HOWTO.html
- <http://en.tldp.org/HOWTO/Software-RAID-HOWTO.html>

Linux RAID mailing lists are also available, such as <http://marc.theaimsgroup.com/?l=linux-raid&r=1&w=2>.

3

Updating the System and Package Management

SUSE Linux provides the option of updating an existing system without completely reinstalling it. There are two types of updates: *updating individual software packages* and *updating the entire system*. Packages can also be installed by hand using the package manager RPM.

3.1 Updating SUSE Linux

Software tends to “grow” from version to version. Therefore, take a look at the available partition space with `df` before updating. If you suspect you are running short of disk space, secure your data before updating and repartition your system. There is no general rule of thumb regarding how much space each partition should have. Space requirements depend on your particular partitioning profile, the software selected, and the version numbers of SUSE Linux.

3.1.1 Preparations

Before updating, copy the old configuration files to a separate medium, such as streamer, removable hard disk, USB stick, or ZIP drive, to secure the data. This primarily applies to files stored in `/etc` as well as some of the directories and files in `/var` and `/opt`. You may also want to write the user data in `/home` (the `HOME` directories) to a backup medium. Back up this data as `root`. Only `root` has read permission for all local files.

Before starting your update, make note of the root partition. The command `df -h` lists the device name of the root partition. In [Example 3.1, “List with `df -h`”](#) (page 66), the root partition to write down is `/dev/hda3` (mounted as `/`).

Example 3.1 *List with `df -h`*

Filesystem	Size	Used	Avail	Use%	Mounted on
<code>/dev/hda3</code>	74G	22G	53G	29%	<code>/</code>
<code>tmpfs</code>	506M	0	506M	0%	<code>/dev/shm</code>
<code>/dev/hda5</code>	116G	5.8G	111G	5%	<code>/home</code>
<code>/dev/hda1</code>	39G	1.6G	37G	4%	<code>/windows/C</code>
<code>/dev/hda2</code>	4.6G	2.6G	2.1G	57%	<code>/windows/D</code>

3.1.2 Possible Problems

If you update a default system from the previous version to this version, YaST works out necessary changes and performs them. Depending on your customizations, some steps or the entire update procedure may fail and you must resort to copying back your backup data. Here, we point out more issues to check before starting the system update.

Checking `passwd` and `group` in `/etc`

Before updating the system, make sure that `/etc/passwd` and `/etc/group` do not contain any syntax errors. For this purpose, start the verification utilities `pwck` and `grpck` as `root` and eliminate any reported errors.

PostgreSQL

Before updating PostgreSQL (`postgres`), dump the databases. See the manual page of `pg_dump`. This is only necessary if you actually used PostgreSQL prior to your update.

3.1.3 Updating with YaST

Following the preparation procedure outlined in [Section 3.1.1, “Preparations”](#) (page 65), you can now update your system:

- 1 Boot the system as for the installation, described in Section “System Start-Up for Installation” (Chapter 1, *Installation with YaST*, ↑Start-Up). In YaST, choose

a language and select *Update* in the *Installation Mode* dialog. Do not select *New Installation*.

- 2 YaST determines whether there are multiple root partitions. If there is only one, continue with the next step. If there are several, select the right partition and confirm with *Next* (`/dev/hda3` was selected in the example in [Section 3.1.1, “Preparations”](#) (page 65)). YaST reads the old `fstab` on this partition to analyze and mount the file systems listed there.
- 3 In the *Installation Settings* dialog, adjust the settings according to your requirements. Normally, you can leave the default settings untouched, but if you intend to enhance your system, check the packages offered in the *Software Selection* submenus or add support for additional languages.

You also have the possibility to make backups of various system components. Selecting backups slows down the update process. Use this option if you do not have a recent system backup.

- 4 In the following dialog, choose to update only the software that is already installed or to add new software components to the system (upgrade mode). It is advisable to accept the suggested composition, for example, *Update Based on Selection "Standard System with KDE"* or *"Standard System with GNOME"*. Adjustments can be made later with YaST.

3.1.4 Updating Individual Packages

Regardless of your overall updated environment, you can always update individual packages. From this point on, however, it is your responsibility to ensure that your system remains consistent. Update advice can be found at <http://www.novell.com/linux/download/updates/>.

Select components from the YaST package selection list according to your needs. If you select a package essential for the overall operation of the system, YaST issues a warning. Such packages should be updated only in the update mode. For example, many packages contain *shared libraries*. If you update these programs and applications in the running system, things might malfunction.

3.2 Software Changes from Version to Version

The individual aspects changed from version to version are outlined in the following in detail. This summary indicates, for example, whether basic settings have been completely reconfigured, whether configuration files have been moved to other places, or whether common applications have been significantly changed. Significant modifications that affect the daily use of the system at either the user level or the administrator level are mentioned here.

Problems and special issues of the respective versions are published online as they are identified. See the links listed below. Important updates of individual packages can be accessed at <http://www.novell.com/products/linuxprofessional/downloads/> using the YaST Online Update (YOU)—see Section “Updating Software Online” (Chapter 2, *System Configuration with YaST, ↑Start-Up*).

3.2.1 From 9.0 to 9.1

Refer to the article “Known Problems and Special Features in SUSE Linux 9.1” in the SUSE Support Database at <http://portal.suse.com> under the keyword *special features*. These articles are published for every SUSE Linux version.

Upgrading to Kernel 2.6

SUSE Linux is now based entirely on kernel 2.6. The predecessor version 2.4 cannot be used any longer, because the enclosed applications do not work with kernel 2.4. Note the following details:

- The loading of modules is configured by means of the file `/etc/modprobe.conf`. The file `/etc/modules.conf` is obsolete. YaST tries to convert the file (also see script `/sbin/generate-modprobe.conf`).
- Modules have the suffix `.ko`.
- The module `ide-scsi` is no longer needed for burning CDs.
- The prefix `snd_` has been removed from the ALSA sound module options.

- `sysfs` now complements the `/proc` file system.
- Power management (especially ACPI) has been improved and can be configured by means of a YaST module.

Mounting VFAT Partitions

When mounting VFAT partitions, the parameter `code` must be changed to `codepage`. If you have difficulties mounting a VFAT partition, check if the file `/etc/fstab` contains the old parameter name.

Standby and Suspend with ACPI

The kernel 2.6 supports standby and suspend with ACPI. This function is still in an experimental stage and may not be supported by some hardware components. To use this function, you need the `powersave` package. Information about this package is available in `/usr/share/doc/packages/powersave`. A graphical front-end is available in the `kpowersave` package.

Input Devices

Regarding the changes in connection with the input devices, refer to the already-mentioned portal article “Known Problems and Special Features in SUSE LINUX 9.1” in the Support Database at <http://portal.suse.com> under the keyword *special features*.

Native POSIX Thread Library and glibc 2.3.x

Applications linked against NGPT (Next Generation POSIX Threading) do not work with glibc 2.3.x. All affected applications that are not shipped with SUSE Linux must be compiled with linuxthreads or with NPTL (Native POSIX Thread Library). NPTL is preferred, because this is the standard for the future.

If NPTL causes difficulties, the older linuxthreads implementation can be used by setting the following environment variable (replace `kernel-version` with the version number of the respective kernel):

```
LD_ASSUME_KERNEL=kernel-version
```

The following version numbers are possible:

2.2.5 (i386, i586):

linuxthreads without floating stacks

2.4.1 (AMD64, IPF, s390x, i586, i686):, 2.4.1 (AMD64, i586, i686):

linuxthread with floating stacks

Notes regarding the kernel and linuxthreads with floating stacks: Applications using `errno`, `h_errno`, and `_res` must include the header files (`errno.h`, `netdb.h`, and `resolv.h`) with `#include`. For C++ programs with multithread support that use *thread cancellation*, the environment variable `LD_ASSUME_KERNEL=2.4.1` must be used to prompt the use of the linuxthreads library.

Adaptations for Native POSIX Thread Library

NPTL is included in SUSE Linux 9.1 as the thread package. NPTL is binary-compatible with the older linuxthreads library. However, areas in which linuxthreads violates the POSIX standard require NPTL adaptions. This includes the following: signal handling, `getpid` returning the same value in all threads, and thread handlers registered with `pthread_atfork` not working if `vfork` is used.

Network Interface Configuration

The configuration of the network interface has changed. Formerly, the hardware was initialized following the configuration of a nonexistent interface. Now, the system searches for new hardware and initializes it immediately, enabling the configuration of the new network interface.

New names have been introduced for the configuration files. Because the name of a network interface is generated dynamically and the use of hotplug devices is increasing steadily, a name like `eth0` or `eth1` is no longer suitable for configuration purposes. For this reason, unique designations, like the MAC address or the PCI slot, are used for naming interface configurations. You can use interface names as soon as they appear. Commands like `ifup eth0` or `ifdown eth0` are still possible.

The device configurations are located in `/etc/sysconfig/hardware`. The interfaces provided by these devices are usually located in `/etc/sysconfig/network`

(with different names). See the detailed description in `/usr/share/doc/packages/sysconfig/README`.

Sound Configuration

Following an update, the sound cards must be reconfigured. This can be done with the YaST sound module. As `root`, enter `/sbin/yast2 sound`.

Top-Level Domain `.local` as “link-local” Domain

The resolver library treats the top-level domain `.local` as “link-local” domain and sends multicast DNS queries to the multicast address `224.0.0.251`, port `5353`, instead of normal DNS queries. This is an incompatible change. If the domain `.local` is already used in the name server configuration, use a different domain name. For more information about multicast DNS, see <http://www.multicastedns.org>.

Systemwide UTF-8 Encoding

The default encoding for the system is UTF-8. Thus, when performing a standard installation, a locale is set with UTF-8 encoding, such as `en_US.UTF-8`. For more information, see <http://www.suse.de/~mfabian/suse-cjk/locales.html>.

Converting Filenames to UTF-8

Files in previously created file systems do not use UTF-8 encoding for the filenames (unless specified otherwise). If these filenames contain non-ASCII characters, they will be garbled. To correct this, use the `convmv` script, which converts the encoding of filenames to UTF-8.

Shell Tools Compatible with POSIX Standard of 2001

In the default setting, shell tools from the `coreutils` package (`tail`, `chown`, `head`, `sort`, etc.) no longer comply with the POSIX standard of 1992 but with the POSIX standard of 2001 (*Single UNIX Specification, version 3 == IEEE Std 1003.1-2001 == ISO/IEC 9945:2002*). The old behavior can be forced with an environment variable:

`_POSIX2_VERSION=199209`

The new value is 200112 and is used as the default for `_POSIX2_VERSION`. The SUS standard can be reviewed (free of charge, but registration is required) at <http://www.unix.org>.

TIP

Third-party software may not yet comply with the new standard. In this case, set the environment variable as described above.

/etc/gshadow Obsolete

`/etc/gshadow` has been abandoned and removed, because this file is superfluous for the following reasons:

- It is not supported by glibc.
- There is no official interface for this file. Even the shadow suite does not contain such an interface.
- Most tools that check the group password do not support the file and ignore it for the said reasons.

OpenLDAP

Because the database format has changed, the databases must be regenerated. During the update, the system attempts to perform this conversion automatically. However, there will certainly be cases in which the conversion fails.

The schema check has undergone substantial improvement. Therefore, a number of standard-noncompliant operations that were possible with the former LDAP server are no longer possible.

The syntax of the configuration file has partly changed with a view to ACLs. Following the installation, information regarding the update is available in the file `/usr/share/doc/packages/openldap2/README.update`.

Apache 1.3 Replaced with Apache 2

The Apache Web server (version 1.3) has been replaced with Apache 2. Detailed documentation for version 2.0 is available at the Web page <http://httpd.apache.org/docs-2.0/en/>. On a system with an HTTP server installation, an update removes the Apache package and installs Apache 2. Subsequently, the system must be adapted with YaST or manually. The configuration files in `/etc/httpd` are now located in `/etc/apache2`.

Either threads or processes can be selected for handling multiple concurrent queries. The process management has been moved to an independent module, the multiprocessing module (MPM). Accordingly, Apache 2 needs the `apache2-prefork` package (recommended for stability) or the `apache2-worker` package. Depending on the MPM, Apache 2 reacts differently to queries. This affects the performance as well as the use of modules. These characteristics are discussed in detail in [Section 26.4.4, “Multiprocessing Modules”](#) (page 469).

Apache 2 now supports the next-generation Internet protocol IPv6.

A mechanism has been implemented that enables module programmers to specify the desired loading sequence of the modules, relieving users of this task. The sequence in which modules are executed is often important. In earlier versions, it was determined by means of the loading sequence. For instance, a module that only gives authenticated users access to certain resources must be loaded first to prevent users without access permissions from seeing the pages.

Queries to and responses from Apache can be processed with filters.

From Samba 2.x to Samba 3.x

Following the update from Samba 2.x to Samba 3.x, `winbind` authentication is no longer available. The other authentication methods can still be used. For this reason, the following programs have been removed:

```
/usr/sbin/wb_auth  
/usr/sbin/wb_ntlmauth  
/usr/sbin/wb_info_group.pl
```

See also <http://www.squid-cache.org/Doc/FAQ/FAQ-23.html#ss23.5>

OpenSSH Update (Version 3.8p1)

`gssapi` support has been replaced with `gssapi-with-mic` to prevent potential MITM attacks. These two versions are not compatible. This means that you cannot authenticate with Kerberos tickets from older distributions, because other authentication methods are used.

SSH and Terminal Applications

When establishing a connection from a remote host (especially via SSH, telnet, and RSH) between version 9 (standard configuration with activated UTF-8) and older systems (SUSE Linux 9.0 and earlier versions in which UTF-8 is not activated by default or not supported), terminal applications may display faulty characters.

This is because OpenSSH does not forward local settings. Therefore, the default system settings that may not match the remote terminal settings are used. This affects YaST in text mode and applications executed from a remote host as a normal user (not `root`). The applications started by `root` are only affected if the user changes the standard locales for `root` (only `LC_CTYPE` is set by default).

libiodbc Discarded

Users of FreeRADIUS must now link against unixODBC, because libiodbc has been discarded.

XML Resources in `/usr/share/xml`

XML resources (DTDs, stylesheets, etc.) are installed in `/usr/share/xml`. Therefore, some directories are no longer available in `/usr/share/sgml`. If you encounter problems, modify your scripts and makefiles or use the official catalogs (especially `/etc/xml/catalog` or `/etc/sgml/catalog`).

Removable Media with subfs

Removable media are now integrated with subfs. Media no longer need to be mounted manually with `mount`. To mount the medium, simply change to the respective device directory in `/media`. Media cannot be ejected as long as they are accessed by a program.

3.2.2 From 9.1 to 9.2

Refer to the article “Known Problems and Special Features in SUSE LINUX 9.2” in the SUSE Support Database at <http://portal.suse.com> under the keyword *special features*.

Activation of the Firewall in the Proposal Dialog During the Installation

To increase the security, the enclosed firewall solution SuSEFirewall2 is activated at the end of the installation in the proposal dialog. This means that all ports are closed initially and can be opened in the proposal dialog if necessary. By default, you cannot log in from remote systems. It also interferes with network browsing and multicast applications, such as SLP, Samba ("Network Neighborhood"), and some games. You can fine-tune the firewall settings using YaST.

If network access is required during the installation or configuration of a service, the respective YaST module opens the needed TCP and UDP ports of all internal and external interfaces. If this is not desired, the user can close the ports in the YaST module or specify other detailed firewall settings.

KDE and IPv6 Support

By default, IPv6 support is not enabled for KDE. You can enable it using the `/etc/sysconfig` editor of YaST. The reason for disabling this feature is that IPv6 addresses are not properly supported by all Internet service providers and, as a consequence, this would lead to error messages while browsing the Web and delays while displaying Web pages.

YaST Online Update and Delta Packages

The YaST Online Update now supports a special kind of RPM package that only stores the binary difference from a given base package. This technique significantly reduces the package size and download time at the expense of higher CPU load for reassembling the final package. In `/etc/sysconfig/onlineupdate`, configure whether YOU should use these delta packages. See `/usr/share/doc/packages/deltarpm/README` for technical details.

Print System Configuration

At the end of the installation (proposal dialog), the ports needed for the print system must be open in the firewall configuration. Port 631/TCP and port 631/UDP are needed for CUPS and should not be closed for normal operation. Port 515/TCP (for the old LPD protocol) and the ports used by Samba must also be open for printing via LPD or SMB.

Change to X.Org

The change from XFree86 to X.Org is facilitated by compatibility links that enable access to important files and commands with the old names.

Table 3.1 Commands

XFree86	X.Org
XFree86	Xorg
xf86config	xorgconfig
xf86cfg	xorgcfg

Table 3.2 Log Files in /var/log

XFree86	X.Org
XFree86.0.log	Xorg.0.log
XFree86.0.log.old	Xorg.0.log.old

In the course of the change to X.Org, the packages were renamed from XFree86* to xorg-x11*.

Terminal Emulators for X11

We have removed a number of terminal emulators because they are either no longer maintained or do not work in the default environment, especially by not supporting

UTF-8. SUSE Linux offers standard terminals, such as xterm, the KDE and GNOME terminals, and mlterm (Multilingual Terminal Emulator for X), which might be a replacement for aterm and eterm.

Changes in the powersave Package

The configuration files in `/etc/sysconfig/powersave` have changed:

Table 3.3 *Split Configuration Files in /etc/sysconfig/powersave*

Old	Now split into
<code>/etc/sysconfig/powersave/</code>	<code>common</code>
<code>common</code>	
	<code>cpufreq</code>
	<code>events</code>
	<code>battery</code>
	<code>sleep</code>
	<code>thermal</code>

`/etc/powersave.conf` has become obsolete. Existing variables have been moved to the files listed in [Table 3.3, “Split Configuration Files in /etc/sysconfig/powersave”](#) (page 77). If you changed the “event” variables in `/etc/powersave.conf`, these must now be adapted in `/etc/sysconfig/powersave/events`.

The names of sleep states have changed from:

- suspend (ACPI S4, APM suspend)
- standby (ACPI S3, APM standby)

To:

- suspend to disk (ACPI S4, APM suspend)

- suspend to ram (ACPI S3, APM suspend)
- standby (ACPI S1, APM standby)

OpenOffice.org (OOo)

Directories:

OOo is now installed in `/usr/lib/ooo-1.1` instead of `/opt/OpenOffice.org`. The default directory for user settings is now `~/.ooo-1.1` instead of `~/OpenOffice.org1.1`.

Wrapper:

There are some new wrappers for starting the OOo components. The new names are shown [Table 3.4, “Wrapper”](#) (page 78).

Table 3.4 *Wrapper*

Old	New
<code>/usr/X11R6/bin/ooo-calc</code>	<code>/usr/bin/oocalc</code>
<code>/usr/X11R6/bin/ooo-draw</code>	<code>/usr/bin/oodraw</code>
<code>/usr/X11R6/bin/ooo-impress</code>	<code>/usr/bin/ooimpress</code>
<code>/usr/X11R6/bin/ooo-math</code>	<code>/usr/bin/oomath</code>
<code>/usr/X11R6/bin/ooo-padmin</code>	<code>/usr/sbin/oopadmin</code>
<code>/usr/X11R6/bin/ooo-setup</code>	—
<code>/usr/X11R6/bin/ooo-template</code>	<code>/usr/bin/oofromtemplate</code>
<code>/usr/X11R6/bin/ooo-web</code>	<code>/usr/bin/ooweb</code>
<code>/usr/X11R6/bin/ooo-writer</code>	<code>/usr/bin/oowriter</code>
<code>/usr/X11R6/bin/ooo</code>	<code>/usr/bin/ooffice</code>

Old	New
/usr/X11R6/bin/OOo-wrapper	/usr/bin/ooo-wrapper

The wrapper now supports the option `--icons-set` for switching between KDE and GNOME icons. The following options are no longer supported:

`--default-configuration`, `--gui`, `--java-path`, `--skip-check`, `--lang` (the language is now determined by means of locales), `--messages-in-window`, and `--quiet`.

KDE and GNOME Support:

KDE and GNOME extensions are available in the `OpenOffice_org-kde` and `OpenOffice_org-gnome` packages.

Sound Mixer kmix

The sound mixer kmix is preset as the default. For high-end hardware, there are other mixers, like QAMix, KAMix, envy24control (only ICE1712), or hdspmixer (only RME Hammerfall).

DVD Burning

In the past, a patch was applied to the `cdrecord` binary from the `cdrecord` package to support burning DVDs. Instead, a new binary `cdrecord-dvd` is installed that has this patch.

The `growisofs` program from the `dvd+rw-tools` package can now burn all DVD media (DVD+R, DVD-R, DVD+RW, DVD-RW, DVD+RL). Try using that one instead of the patched `cdrecord-dvd`.

Multiple Kernels

It is possible to install multiple kernels side by side. This feature is meant to allow administrators to upgrade from one kernel to another by installing the new kernel, verifying that the new kernel works as expected, then uninstalling the old kernel. While YaST does not yet support this feature, kernels can easily be installed and uninstalled from the shell using `rpm -i package.rpm`.

The default boot loader menus contain one kernel entry. Before installing multiple kernels, it is useful to add an entry for the extra kernels, so that they can easily be selected. The kernel that was active before installing a new kernel can be accessed as `vmlinuz.previous` and `initrd.previous`. By creating a boot loader entry similar to the default entry and having this entry refer to `vmlinuz.previous` and `initrd.previous` instead of `vmlinuz` and `initrd`, the previously active kernel can be accessed. Alternatively, GRUB and LILO support wild card boot loader entries. Refer to the GRUB info pages (`info grub`) and to the `lilo.conf` (5) manual page for details.

3.2.3 From 9.2 to 9.3

Refer to the article “Known Problems and Special Features in SUSE Linux 9.3” in the SUSE Support Database at <http://portal.suse.com> under the keyword *special features*.

Starting Manual Installation at the Kernel Prompt

The *Manual Installation* mode is gone from the boot loader screen. You can still get `linuxrc` into manual mode using `manual=1` at the boot prompt. Normally this is not necessary because you can set installation options at the kernel prompt directly, such as `textmode=1` or a URL as the installation source.

Kerberos for Network Authentication

`Kerberos` is the default for network authentication instead of `heimdal`. Converting an existing `heimdal` configuration automatically is not possible. During a system update, backup copies of configuration files are created as shown in [Table 3.5, “Backup Files”](#) (page 80).

Table 3.5 Backup Files

Old File	Backup File
<code>/etc/krb5.conf</code>	<code>/etc/krb5.conf.heimdal</code>
<code>/etc/krb5.keytab</code>	<code>/etc/krb5.keytab.heimdal</code>

The client configuration (`/etc/krb5.conf`) is very similar to the one of heimdal. If nothing special was configured, it is enough to replace the parameter `kpasswd_server` with `admin_server`.

It is not possible to copy the server-related (kdc and kadmind) data. After the system update, the old heimdal database is still available under `/var/heimdal`. MIT kerberos maintains the database under `/var/lib/kerberos/krb5kdc`.

JFS: Not Supported Anymore

Due to technical problems with JFS, it is no longer supported. The kernel file system driver is still there, but YaST does not offer partitioning with JFS.

AIDE as a Tripwire Replacement

As an intrusion detection system, use AIDE (package name `aide`), which is released under the GPL. Tripwire is no longer available on SUSE Linux.

X.Org Configuration File

The configuration tool SaX2 writes the X.Org configuration settings into `/etc/X11/xorg.conf`. During an installation from scratch, no compatibility link from `XF86Config` to `xorg.conf` is created.

XView and OpenLook Support Dropped

The packages `xview`, `xview-devel`, `xview-devel-examples`, `olvwm`, and `xtoopl` were dropped. In the past, we just provided the XView (OpenLook) base system. The XView libraries are no longer provided after the system update. Even more important, OLVWM (OpenLook Virtual Window Manager) is no longer available.

PAM Configuration

New Configuration Files (containing comments for more information)

`common-auth`

Default PAM configuration for auth section

```
common-account
```

Default PAM configuration for account section

```
common-password
```

Default PAM configuration for password changing

```
common-session
```

Default PAM configuration for session management

You should include these default configuration files from within your application-specific configuration file, because it is easier to modify and maintain one file instead of the approximately forty files that used to exist on the system. If you install an application later, it inherits the already applied changes and the administrator is not required to remember to adjust the configuration.

The changes are simple. If you have the following configuration file (which should be the default for most applications):

```
#%PAM-1.0
auth    required      pam_unix2.so
account required      pam_unix2.so
password required    pam_pwcheck.so
password required    pam_unix2.so      use_first_pass use_authok
#password required  pam_make.so    /var/yp
session required     pam_unix2.so
```

you can change it to:

```
#%PAM-1.0
auth    include      common-auth
account include      common-account
password include    common-password
session include     common-session
```

Stricter tar Syntax

The `tar` usage syntax is stricter now. The `tar` options must come before the file or directory specifications. Appending options, like `--atime-preserve` or `--numeric-owner`, after the file or directory specification makes `tar` fail. Check your backup scripts. Commands such as the following no longer work:

```
tar czf etc.tar.gz /etc --atime-preserve
```

See the `tar` info pages for more information.

3.2.4 From 9.3 to 10.0

Refer to the article “Known Problems and Special Features in SUSE Linux 10” in the SUSE Support Database at <http://portal.suse.com> under the keyword *special features*.

Becoming the Superuser Using su

By default, calling `su` to become `root` does not set the `PATH` for `root`. Either call `su -` to start a login shell with the complete environment for `root` or set `ALWAYS_SET_PATH` to `yes` in `/etc/default/su` if you want to change the default behavior of `su`.

Powersave Configuration Variables

Names of the powersave configuration variables are changed for consistency, but the sysconfig files are still the same. Find more information in [Section 33.5.1, “Configuring the powersave Package”](#) (page 589).

PCMCIA

`cardmgr` no longer manages PC cards. Instead, as with Cardbus cards and other subsystems, a kernel module manages them. All necessary actions are executed by `hotplug`. The `pcmcia` start script has been removed and `cardctl` is replaced by `pccardctl`. For more information, see `/usr/share/doc/packages/pcmciautils/README.SUSE`.

Setting Up D-BUS for Interprocess Communication in .xinitrc

Many applications now rely on D-BUS for interprocess communication (IPC). Calling `dbus-launch` starts `dbus-daemon`. The systemwide `/etc/X11/xinit/xinitrc` uses `dbus-launch` to start the window manager.

If you have a local `~/.xinitrc` file, you must change it accordingly. Otherwise applications like f-spot, banshee, tomboy, or Network Manager banshee might fail. Save

your old `~/.xinitrc`. Then copy the new template file into your home directory with:

```
cp /etc/skel/.xinitrc.template ~/.xinitrc
```

Finally, add your customizations from the saved `.xinitrc`.

NTP-Related Files Renamed

For reasons of compatibility with LSB (Linux Standard Base), most configuration files and the init script were renamed from `xntp` to `ntp`. The new filenames are:

```
/etc/slpl.reg.d/ntp.reg  
/etc/init.d/ntp  
/etc/logrotate.d/ntp  
/usr/sbin/rcntp  
/etc/sysconfig/ntp
```

Hotplug Events Handled by the udev Daemon

Hotplug events are now completely handled by the udev daemon (`udevd`). We do not use the event multiplexer system in `/etc/hotplug.d` and `/etc/dev.d` anymore. Instead `udevd` calls all hotplug helper tools directly, according to its rules. Udev rules and helper tools are provided by udev and various other packages.

TEI XSL Stylesheets

Find the TEI XSL stylesheets (`tei-xsl-stylesheets`) with a new directory layout at `/usr/share/xml/tei/stylesheet/rahtz/current`. From there, for example, use `base/p4/html/tei.xsl` to produce HTML output. For more information, see <http://www.tei-c.org/Stylesheets/teic/>

File System Change Notification for GNOME Applications

For proper functionality, GNOME applications depend on file system change notification support. For local-only file systems, install the gamin package (preferred) or run the FAM daemon. For remote file systems, run FAM on both the server and client and open the firewall for RPC calls by FAM.

GNOME (gnome-vfs2 and libgda) contains a wrapper that picks gamin or fam to provide file system change notification:

- If the FAM daemon is not running, gamin is preferred. (Rationale: Inotify is supported only by gamin and it is more efficient for local file systems).
- If the FAM daemon is running, FAM is preferred (Rationale: If FAM is running, you probably want remote notification, which is supported only by FAM).

3.2.5 From 10.0 to 10.1

Refer to the article “Known Problems and Special Features in SUSE Linux 10” in the SUSE Support Database at <http://portal.suse.com> under the keyword *special features*.

Apache 2.2

For Apache version 2.2, [Chapter 26, The Apache HTTP Server](#) (page 447) was completely reworked. In addition, find generic upgrade information at <http://httpd.apache.org/docs/2.2/upgrading.html> and the description of new features at http://httpd.apache.org/docs/2.2/new_features_2_2.html.

Starting an FTP Server (vsftpd)

By default, `xinetd` no longer starts the `vsftpd` FTP server. It is now a stand-alone daemon and you must configure it with the YaST runtime editor.

Firefox 1.5: The URL Open Command

With Firefox 1.5, the method for applications to open a Firefox instance or window has changed. The new method was already partly available in former versions where the behavior was implemented in the wrapper script.

If your application does not use `mozilla-xremote-client` or `firefox-remote`, you do not have to change anything. Otherwise the new command to open a URL is `firefox url` and it does not matter whether Firefox is already running or not. If it is already running, it follows the preference configured in *Open links from other applications in*.

From the command line, you can influence the behavior by using `firefox-new-window url` or `firefox -new-tab url`.

3.3 RPM—the Package Manager

In SUSE Linux, RPM (RPM Package Manager) is used for managing the software packages. Its main programs are `rpm` and `rpmbuild`. The powerful RPM database can be queried by the users, system administrators, and package builders for detailed information about the installed software.

Essentially, `rpm` has five modes: installing, uninstalling, or updating software packages; rebuilding the RPM database; querying RPM bases or individual RPM archives; integrity checking of packages; and signing packages. `rpmbuild` can be used to build installable packages from pristine sources.

Installable RPM archives are packed in a special binary format. These archives consist of the program files to install and certain meta information used during the installation by `rpm` to configure the software package or stored in the RPM database for documentation purposes. RPM archives normally have the extension `.rpm`.

TIP: Software Development Packages

For a number of packages, the components needed for software development (libraries, headers, include files, etc.) have been put into separate packages. These development packages are only needed if you want to compile software yourself, for example, the most recent GNOME packages. They can be identified by the name extension `-devel`, such as the packages `alsa-devel`, `gimp-devel`, and `kdelibs3-devel`.

3.3.1 Verifying Package Authenticity

SUSE Linux RPM packages have a GnuPG signature. The key including the fingerprint is:

```
1024D/9C800ACA 2000-10-19 SuSE Package Signing Key <build@suse.de>
Key fingerprint = 79C1 79B2 E1C8 20C1 890F 9994 A84E DAE8 9C80 0ACA
```

The command `rpm --checksig package-1.2.3.rpm` can be used to verify the signature of an RPM package to determine whether it really originates from SUSE Linux or from another trustworthy facility. This is especially recommended for update packages from the Internet. The SUSE Linux public package signature key normally resides in `/root/.gnupg/`. The key is additionally located in the directory `/usr/lib/rpm/gnupg/` to enable normal users to verify the signature of RPM packages.

3.3.2 Managing Packages: Install, Update, and Uninstall

Normally, the installation of an RPM archive is quite simple: `rpm -i package.rpm`. With this command, the package is installed, but only if its dependencies are fulfilled and there are no conflicts with other packages. With an error message, `rpm` requests those packages that need to be installed to meet dependency requirements. In the background, the RPM database ensures that no conflicts arise—a specific file can only belong to one package. By choosing different options, you can force `rpm` to ignore these defaults, but this is only for experts. Otherwise, risk compromising the integrity of the system and possibly jeopardize the ability to update the system.

The options `-U` or `--upgrade` and `-F` or `--frefresh` can be used to update a package, for example, `rpm -F package.rpm`. This command removes the files of the old version and immediately installs the new files. The difference between the two versions is that `-U` installs packages that previously did not exist in the system, but `-F` merely updates previously installed packages. When updating, `rpm` updates configuration files carefully using the following strategy:

- If a configuration file was not changed by the system administrator, `rpm` installs the new version of the appropriate file. No action by the system administrator is required.
- If a configuration file was changed by the system administrator before the update, `rpm` saves the changed file with the extension `.rpmodif` or `.rpmsave` (backup file) and installs the version from the new package, but only if the originally installed file and the newer version are different. If this is the case, compare the backup file (`.rpmodif` or `.rpmsave`) with the newly installed file and make your changes again in the new file. Afterwards, be sure to delete all `.rpmodif` and `.rpmsave` files to avoid problems with future updates.
- `.rpmsave` files appear if the configuration file already exists *and* if the `noreplace` label was specified in the `.spec` file.

Following an update, `.rpmsave` and `.rpmsave` files should be removed after comparing them, so they do not obstruct future updates. The `.rpmodif` extension is assigned if the file has not previously been recognized by the RPM database.

Otherwise, `.rpmsave` is used. In other words, `.rpmodif` results from updating from a foreign format to RPM. `.rpmsave` results from updating from an older RPM to a newer RPM. `.rpmsave` does not disclose any information as to whether the system administrator has made any changes to the configuration file. A list of these files is available in `/var/adm/rpmconfigcheck`. Some configuration files (like `/etc/httpd/httpd.conf`) are not overwritten to allow continued operation.

The `-U` switch is *not* just an equivalent to uninstalling with the `-e` option and installing with the `-i` option. Use `-U` whenever possible.

To remove a package, enter `rpm -e package`. `rpm` only deletes the package if there are no unresolved dependencies. It is theoretically impossible to delete Tcl/Tk, for example, as long as another application requires it. Even in this case, RPM calls for assistance from the database. If such a deletion is—for whatever reason and under unusual

circumstances—impossible, even if *no* additional dependencies exist, it may be helpful to rebuild the RPM database using the option `--rebu1ddb`.

3.3.3 RPM and Patches

To guarantee the operational security of a system, update packages must be installed in the system from time to time. Previously, a bug in a package could only be eliminated by replacing the entire package. Large packages with bugs in small files could easily result in large amounts of data. However the SUSE RPM offers a feature enabling the installation of patches in packages.

The most important considerations are demonstrated using pine as an example:

Is the patch RPM suitable for my system?

To check this, first query the installed version of the package. For pine, this can be done with

```
rpm -q pine  
pine-4.44-188
```

Then check if the patch RPM is suitable for this version of pine:

```
rpm -qp --basedon pine-4.44-224.i586.patch.rpm  
pine = 4.44-188  
pine = 4.44-195  
pine = 4.44-207
```

This patch is suitable for three different versions of pine. The installed version in the example is also listed, so the patch can be installed.

Which files are replaced by the patch?

The files affected by a patch can easily be seen in the patch RPM. The `rpm` parameter `-P` allows selection of special patch features. Display the list of files with the following command:

```
rpm -qpPl pine-4.44-224.i586.patch.rpm  
/etc/pine.conf  
/etc/pine.conf.fixed  
/usr/bin/pine
```

or, if the patch is already installed, with the following command:

```
rpm -qPl pine  
/etc/pine.conf
```

```
/etc/pine.conf.fixed  
/usr/bin/pine
```

How can a patch RPM be installed in the system?

Patch RPMs are used just like normal RPMs. The only difference is that a suitable RPM must already be installed.

Which patches are already installed in the system and for which package versions?

A list of all patches installed in the system can be displayed with the command `rpm -qPa`. If only one patch is installed in a new system (as in this example), the list appears as follows:

```
rpm -qPa  
pine-4.44-224
```

If, at a later date, you want to know which package version was originally installed, this information is also available in the RPM database. For `pine`, this information can be displayed with the following command:

```
rpm -q --basedon pine  
pine = 4.44-188
```

More information, including information about the patch feature of RPM, is available in the man pages of `rpm` and `rpmbuild`.

3.3.4 Delta RPM Packages

Delta RPM packages contain the difference between an old and a new version of an RPM package. Applying a delta RPM on an old RPM results in the complete new RPM. It is not necessary to have a copy of the old RPM, because a delta RPM can also work with an installed RPM. The delta RPM packages are even smaller in size than patch RPMs, which is an advantage when transferring update packages over the Internet. The drawback is that update operations with delta RPMs involved consume considerably more CPU cycles than plain or patch RPMs. To make YaST use delta RPM packages during YOU sessions, set `YOU_USE_DELTAS` to `yes` in `/etc/sysconfig/onlineupdate`. In this case, be prepared to have your installation media available. If `YOU_USE_DELTAS` is empty or set to `filesystem`, YOU tries to download delta packages that apply to installed files. You do not need the installation media in this case but the download time could be longer. If set to `no`, YOU only uses patch RPMs and normal RPMs.

The `prepdeltarpm`, `writedeltarpm`, and `applydeltarpm` binaries are part of the delta RPM suite (package `deltarpm`) and help you create and apply delta RPM packages. With the following commands, create a delta RPM called `new.delta.rpm`. The following command assumes that `old.rpm` and `new.rpm` are present:

```
prepdeltarpm -s seq -i info old.rpm > old.cpio  
prepdeltarpm -f new.rpm > new.cpio  
  
xdelta delta -0 old.cpio new.cpio delta  
  
writedeltarpm new.rpm delta info new.delta.rpm  
rm old.cpio new.cpio delta
```

Using `applydeltarpm`, you can reconstruct the new RPM from the file system if the old package is already installed:

```
applydeltarpm new.delta.rpm new.rpm
```

To derive it from the old RPM without accessing the file system, use the `-r` option:

```
applydeltarpm -r old.rpm new.delta.rpm new.rpm
```

See `/usr/share/doc/packages/deltarpm/README`" for technical details.

3.3.5 RPM Queries

With the `-q` option, `rpm` initiates queries, making it possible to inspect an RPM archive (by adding the option `-p`) and also to query the RPM database of installed packages. Several switches are available to specify the type of information required. See [Table 3.6, “The Most Important RPM Query Options”](#) (page 91).

Table 3.6 *The Most Important RPM Query Options*

<code>-i</code>	Package information
<code>-l</code>	File list
<code>-f FILE</code>	Query the package that contains the file <i>FILE</i> (the full path must be specified with <i>FILE</i>)
<code>-s</code>	File list with status information (implies <code>-l</code>)

-d	List only documentation files (implies -l)
-c	List only configuration files (implies -l)
--dump	File list with complete details (to be used with -l, -c, or -d)
--provides	List features of the package that another package can request with --requires
--requires, -R	Capabilities the package requires
--scripts	Installation scripts (preinstall, postinstall, uninstall)

For example, the command `rpm -q -i wget` displays the information shown in [Example 3.2, “rpm -q -i wget”](#) (page 92).

Example 3.2 *rpm -q -i wget*

```
Name      : wget                                Relocations: (not relocatable)
Version   : 1.9.1                               Vendor: SUSE LINUX AG,
Nuernberg, Germany
Release   : 50                                 Build Date: Sat 02 Oct 2004
03:49:13 AM CEST
Install date: Mon 11 Oct 2004 10:24:56 AM CEST      Build Host: f53.suse.de
Group     : Productivity/Networking/Web/Utilities    Source RPM:
wget-1.9.1-50.src.rpm
Size      : 1637514                             License: GPL
Signature : DSA/SHA1, Sat 02 Oct 2004 03:59:56 AM CEST, Key ID
a84edae89c800aca
Packager  : http://www.suse.de/feedback
URL       : http://wget.sunsite.dk/
Summary   : A tool for mirroring FTP and HTTP servers
Description :
Wget enables you to retrieve WWW documents or FTP files from a server.
This can be done in script files or via the command line.
[...]
```

The option `-f` only works if you specify the complete filename with its full path. Provide as many filenames as desired. For example, the following command

```
rpm -q -f /bin/rpm /usr/bin/wget
```

results in:

```
rpm-4.1.1-191  
wget-1.9.1-50
```

If only part of the filename is known, use a shell script as shown in [Example 3.3, “Script to Search for Packages”](#) (page 93). Pass the partial filename to the script shown as a parameter when running it.

Example 3.3 *Script to Search for Packages*

```
#!/bin/sh  
for i in $(rpm -q -a -l | grep $1); do  
    echo "\"$i\" is in package:"  
    rpm -q -f $i  
    echo ""  
done
```

The command `rpm -q --changelog rpm` displays a detailed list of change information about a specific package, sorted by date. This example shows information about the package `rpm`.

With the help of the installed RPM database, verification checks can be made. Initiate these with `-V`, `-y`, or `--verify`. With this option, `rpm` shows all files in a package that have been changed since installation. `rpm` uses eight character symbols to give some hints about the following changes:

Table 3.7 *RPM Verify Options*

5	MD5 check sum
S	File size
L	Symbolic link
T	Modification time
D	Major and minor device numbers
U	Owner
G	Group
M	Mode (permissions and file type)

In the case of configuration files, the letter `c` is printed. For example, for changes to `/etc/wgetrc` (`wget`):

```
rpm -V wget  
S.5....T c /etc/wgetrc
```

The files of the RPM database are placed in `/var/lib/rpm`. If the partition `/usr` has a size of 1 GB, this database can occupy nearly 30 MB, especially after a complete update. If the database is much larger than expected, it is useful to rebuild the database with the option `--rebuilddb`. Before doing this, make a backup of the old database. The `cron` script `cron.daily` makes daily copies of the database (packed with `gzip`) and stores them in `/var/adm/backup/rpmdb`. The number of copies is controlled by the variable `MAX_RPMDB_BACKUPS` (default: 5) in `/etc/sysconfig/backup`. The size of a single backup is approximately 1 MB for 1 GB in `/usr`.

3.3.6 Installing and Compiling Source Packages

All source packages of SUSE Linux carry a `.src.rpm` extension (source RPM).

TIP

Source packages can be copied from the installation medium to the hard disk and unpacked with YaST. They are not, however, marked as installed ([i]) in the package manager. This is because the source packages are not entered in the RPM database. Only *installed* operating system software is listed in the RPM database. When you “install” a source package, only the source code is added to the system.

The following directories must be available for `rpm` and `rpmbuild` in `/usr/src` packages (unless you specified custom settings in a file like `/etc/rpmrc`):

SOURCES

for the original sources (`.tar.bz2` or `.tar.gz` files, etc.) and for distribution-specific adjustments (mostly `.diff` or `.patch` files)

SPECS

for the `.spec` files, similar to a meta Makefile, which control the *build* process

BUILD

all the sources are unpacked, patched, and compiled in this directory

RPMS

where the completed binary packages are stored

SRPMS

here are the source RPMs

When you install a source package with YaST, all the necessary components are installed in `/usr/src/packages`: the sources and the adjustments in `SOURCES` and the relevant `.spec` file in `SPECS`.

WARNING

Do not experiment with system components (`glibc`, `rpm`, `sysvinit`, etc.), because this endangers the operability of your system.

The following example uses the `wget.src.rpm` package. After installing the package with YaST, you should have files similar to the following listing:

```
/usr/src/packages/SOURCES/nops_doc.diff  
/usr/src/packages/SOURCES/toplev_destdir.diff  
/usr/src/packages/SOURCES/wget-1.9.1+ipvmisc.patch  
/usr/src/packages/SOURCES/wget-1.9.1-brokentime.patch  
/usr/src/packages/SOURCES/wget-1.9.1-passive_ftp.diff  
/usr/src/packages/SOURCES/wget-LFS-20040909.tar.bz2  
/usr/src/packages/SOURCES/wget-wrong_charset.patch  
/usr/src/packages/SPECS/wget.spec
```

`rpmbuild -b X /usr/src/packages/SPECS/wget.spec` starts the compilation. `X` is a wild card for various stages of the build process (see the output of `--help` or the RPM documentation for details). The following is merely a brief explanation:

`-bp`

Prepare sources in `/usr/src/packages/BUILD`: unpack and patch.

`-bc`

Do the same as `-bp`, but with additional compilation.

`-bi`

Do the same as `-bp`, but with additional installation of the built software. Caution: if the package does not support the BuildRoot feature, you might overwrite configuration files.

`-bb`

Do the same as `-bi`, but with the additional creation of the binary package. If the compile was successful, the binary should be in `/usr/src/packages/RPMS`.

`-ba`

Do the same as `-bb`, but with the additional creation of the source RPM. If the compilation was successful, the binary should be in `/usr/src/packages/SRPMs`.

`--short-circuit`

Skip some steps.

The binary RPM created can now be installed with `rpm -i` or, preferably, with `rpm -U`. Installation with `rpm` makes it appear in the RPM database.

3.3.7 Compiling RPM Packages with build

The danger with many packages is that unwanted files are added to the running system during the build process. To prevent this, use `build`, which creates a defined environment in which the package is built. To establish this chroot environment, the `build` script must be provided with a complete package tree. This tree can be made available on the hard disk, via NFS, or from DVD. Set the position with `build --rpms directory`. Unlike `rpm`, the `build` command looks for the SPEC file in the source directory. To build `wget` (like in the above example) with the DVD mounted in the system under `/media/dvd`, use the following commands as `root`:

```
cd /usr/src/packages/SOURCES/
mv ../SPECS/wget.spec .
build --rpms /media/dvd/suse/ wget.spec
```

Subsequently, a minimum environment is established at `/var/tmp/build-root`. The package is built in this environment. Upon completion, the resulting packages are located in `/var/tmp/build-root/usr/src/packages/RPMS`.

The `build` script offers a number of additional options. For example, cause the script to prefer your own RPMs, omit the initialization of the build environment, or limit the `rpm` command to one of the above-mentioned stages. Access additional information with `build --help` and by reading the `build` man page.

3.3.8 Tools for RPM Archives and the RPM Database

Midnight Commander (`mc`) can display the contents of RPM archives and copy parts of them. It represents archives as virtual file systems, offering all usual menu options of Midnight Commander. Display the HEADER with `[F3]`. View the archive structure with the cursor keys and `[Enter]`. Copy archive components with `[F5]`.

KDE offers the `kpackage` tool as a front-end for `rpm`. A full-featured package manager is available as a YaST module (see Section “Installing and Removing Software” (Chapter 2, *System Configuration with YaST, ↑Start-Up*)).

Part II. Administration

Security in Linux

Masquerading and a firewall ensure a controlled data flow and data exchange. SSH (secure shell) enables you to log in to remote hosts over an encrypted connection. The encryption of files or entire partitions protects your data in the event that third parties gain access to your system. Along with technical instructions, find information about security aspects of Linux networks.

4.1 Masquerading and Firewalls

Whenever Linux is used in a networked environment, you can use the kernel functions that allow the manipulation of network packets to maintain a separation between internal and external network areas. The Linux netfilter framework provides the means to establish an effective firewall that keeps different networks apart. With the help of iptables—a generic table structure for the definition of rule sets—precisely control the packets allowed to pass a network interface. Such a packet filter can be set up quite easily with the help of SuSEfirewall2 and the corresponding YaST module.

4.1.1 Packet Filtering with iptables

The components netfilter and iptables are responsible for the filtering and manipulation of network packets as well as for network address translation (NAT). The filtering criteria and any actions associated with them are stored in chains, which must be matched one after another by individual network packets as they arrive. The chains to match are stored in tables. The `iptables` command allows you to alter these tables and rule sets.

The Linux kernel maintains three tables, each for a particular category of functions of the packet filter:

filter

This table holds the bulk of the filter rules, because it implements the *packet filtering* mechanism in the stricter sense, which determines whether packets are let through (ACCEPT) or discarded (DROP), for example.

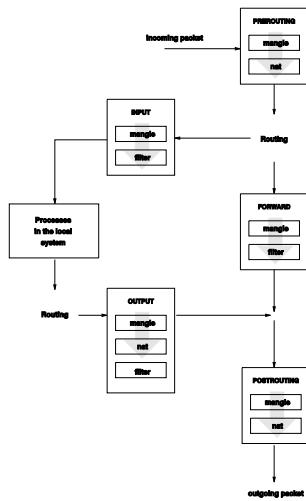
nat

This table defines any changes to the source and target addresses of packets. Using these functions also allows you to implement *masquerading*, which is a special case of NAT used to link a private network with the Internet.

mangle

The rules held in this table make it possible to manipulate values stored in IP headers (such as the type of service).

Figure 4.1 *iptables: A Packet's Possible Paths*



These tables contain several predefined chains to match packets:

PREROUTING

This chain is applied to incoming packets.

INPUT

This chain is applied to packets destined for the system's internal processes.

FORWARD

This chain is applied to packets that are only routed through the system.

OUTPUT

This chain is applied to packets originating from the system itself.

POSTROUTING

This chain is applied to all outgoing packets.

[Figure 4.1, “iptables: A Packet's Possible Paths”](#) (page 102) illustrates the paths along which a network packet may travel on a given system. For the sake of simplicity, the figure lists tables as parts of chains, but in reality these chains are held within the tables themselves.

In the simplest of all possible cases, an incoming packet destined for the system itself arrives at the `eth0` interface. The packet is first referred to the `PREROUTING` chain of the `mangle` table then to the `PREROUTING` chain of the `nat` table. The following step, concerning the routing of the packet, determines that the actual target of the packet is a process of the system itself. After passing the `INPUT` chains of the `mangle` and the `filter` table, the packet finally reaches its target, provided that the rules of the `filter` table are actually matched.

4.1.2 Masquerading Basics

Masquerading is the Linux-specific form of NAT (network address translation). It can be used to connect a small LAN (where hosts use IP addresses from the private range—see [Section 18.1.2, “Netmasks and Routing”](#) (page 321)) with the Internet (where official IP addresses are used). For the LAN hosts to be able to connect to the Internet, their private addresses are translated to an official one. This is done on the router, which acts as the gateway between the LAN and the Internet. The underlying principle is a simple one: The router has more than one network interface, typically a network card and a separate interface connecting with the Internet. While the latter links the router with the outside world, one or several others link it with the LAN hosts. With these hosts in the local network connected to the network card (such as `eth0`) of the router, they can send any packets not destined for the local network to their default gateway or router.

IMPORTANT: Using the Correct Network Mask

When configuring your network, make sure both the broadcast address and the netmask are the same for all local hosts. Failing to do so prevents packets from being routed properly.

As mentioned, whenever one of the LAN hosts sends a packet destined for an Internet address, it goes to the default router. However, the router must be configured before it can forward such packets. For security reasons, SUSE Linux does not enable this in a default installation. To enable it, set the variable `IP_FORWARD` in the file `/etc/sysconfig/sysctl` to `IP_FORWARD=yes`.

The target host of the connection can see your router, but knows nothing about the host in your internal network where the packets originated. This is why the technique is called masquerading. Because of the address translation, the router is the first destination of any reply packets. The router must identify these incoming packets and translate their target addresses, so packets can be forwarded to the correct host in the local network.

With the routing of inbound traffic depending on the masquerading table, there is no way to open a connection to an internal host from the outside. For such a connection, there would be no entry in the table. In addition, any connection already established has a status entry assigned to it in the table, so the entry cannot be used by another connection.

As a consequence of all this, you might experience some problems with a number of application protocols, such as ICQ, cucme, IRC (DCC, CTCP), and FTP (in PORT mode). Netscape, the standard FTP program, and many others use the PASV mode. This passive mode is much less problematic as far as packet filtering and masquerading are concerned.

4.1.3 Firewalling Basics

Firewall is probably the term most widely used to describe a mechanism that provides and manages a link between networks while also controlling the data flow between them. Strictly speaking, the mechanism described in this section is called a *packet filter*. A packet filter regulates the data flow according to certain criteria, such as protocols, ports, and IP addresses. This allows you to block packets that, according to their addresses, are not supposed to reach your network. To allow public access to your Web

server, for example, explicitly open the corresponding port. However, a packet filter does not scan the contents of packets with legitimate addresses, such as those directed to your Web server. For example, if incoming packets were intended to compromise a CGI program on your Web server, the packet filter would still let them through.

A more effective but more complex mechanism is the combination of several types of systems, such as a packet filter interacting with an application gateway or proxy. In this case, the packet filter rejects any packets destined for disabled ports. Only packets directed to the application gateway are accepted. This gateway or proxy pretends to be the actual client of the server. In a sense, such a proxy could be considered a masquerading host on the protocol level used by the application. One example for such a proxy is Squid, an HTTP proxy server. To use Squid, the browser must be configured to communicate via the proxy. Any HTTP pages requested are served from the proxy cache and pages not found in the cache are fetched from the Internet by the proxy. As another example, the SUSE proxy-suite (`proxy-suite`) provides a proxy for the FTP protocol.

The following section focuses on the packet filter that comes with SUSE Linux. For further information about packet filtering and firewalling, read the Firewall HOWTO included in the `howto` package. If this package is installed, read the HOWTO with `less /usr/share/doc/howto/en/txt/Firewall-HOWTO.gz`.

4.1.4 SuSEfirewall2

SuSEfirewall2 is a script that reads the variables set in `/etc/sysconfig/SuSEfirewall2` to generate a set of iptables rules. It defines three security zones, although only the first and the second one are considered in the following sample configuration:

External Zone

Given that there is no way to control what is happening on the external network, the host needs to be protected from it. In most cases, the external network is the Internet, but it could be another insecure network, such as a WLAN.

Internal Zone

This refers to the private network, in most cases the LAN. If the hosts on this network use IP addresses from the private range (see [Section 18.1.2, “Netmasks and Routing”](#) (page 321)), enable network address translation (NAT), so hosts on the internal network can access the external one.

Demilitarized Zone (DMZ)

While hosts located in this zone can be reached both from the external and the internal network, they cannot access the internal network themselves. This setup can be used to put an additional line of defense in front of the internal network, because the DMZ systems are isolated from the internal network.

Any kind of network traffic not explicitly allowed by the filtering rule set is suppressed by iptables. Therefore, each of the interfaces with incoming traffic must be placed into one of the three zones. For each of the zones, define the services or protocols allowed. The rule set is only applied to packets originating from remote hosts. Locally generated packets are not captured by the firewall.

The configuration can be performed with YaST (see [Section “Configuring with YaST”](#) (page 106)). It can also be made manually in the file `/etc/sysconfig/SuSEfirewall2`, which is well commented. Additionally, a number of example scenarios are available in `/usr/share/doc/packages/SuSEfirewall2/EXAMPLES`.

Configuring with YaST

IMPORTANT: Automatic Firewall Configuration

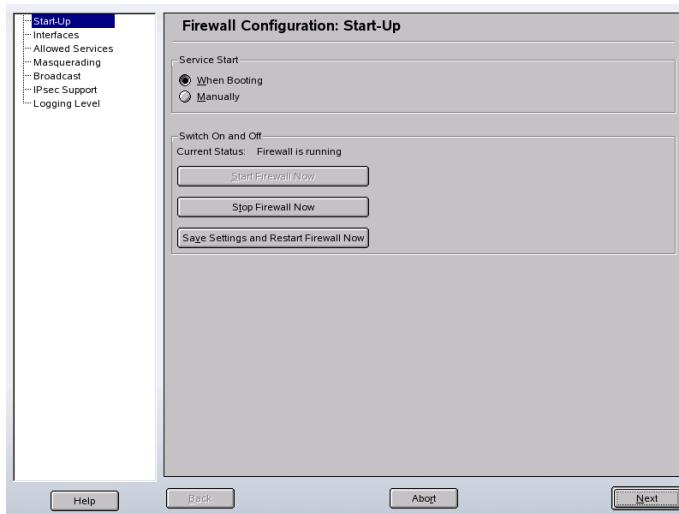
After the installation, YaST automatically starts a firewall on all configured interfaces. If a server is configured and activated on the system, YaST can modify the automatically-generated firewall configuration with the options *Open Ports on Selected Interface in Firewall* or *Open Ports on Firewall* in the server configuration modules. Some server module dialogs include a *Firewall Details* button for activating additional services and ports. The YaST firewall configuration module can be used to activate, deactivate, or reconfigure the firewall.

The YaST dialogs for the graphical configuration can be accessed from the YaST Control Center. Select *Security and Users* → *Firewall*. The configuration is divided into seven sections that can be accessed directly from the tree structure on the left side.

Start-Up

Set the start-up behavior in this dialog. In a default installation, SuSEfirewall2 is started automatically. You can also start and stop the firewall here. To implement your new settings in a running firewall, use *Save Settings and Restart Firewall Now*.

Figure 4.2 The YaST Firewall Configuration



Interfaces

All known network interfaces are listed here. To remove an interface from a zone, select the interface, press *Change*, and choose *No Zone Assigned*. To add an interface to a zone, select the interface, press *Change* and choose any of the available zones. You may also create a special interface with your own settings by using *Custom*.

Allowed Services

You need this option to offer services from your system to a zone from which it is protected. By default, the system is only protected from external zones. Explicitly allow the services that should be available to external hosts. Activate the services after selecting the desired zone in *Allowed Services for Selected Zone*.

Masquerading

Masquerading hides your internal network from external networks, such as the Internet, while enabling hosts in the internal network to access the external network transparently. Requests from the external network to the internal one are blocked and requests from the internal network seem to be issued by the masquerading server when seen externally. If special services of an internal machine need to be available to the external network, add special redirect rules for the service.

Broadcast

In this dialog, configure the UDP ports that allow broadcasts. Add the required port numbers or services to the appropriate zone, separated by spaces. See also the file `/etc/services`.

The logging of broadcasts that are not accepted can be enabled here. This may be problematic, because Windows hosts use broadcasts to know about each other and so generate many packets that are not accepted.

IPsec Support

Configure whether the IPsec service should be available to the external network in this dialog. Configure which packets are trusted under *Details*.

Logging Level

There are two rules for the logging: accepted and not accepted packets. Packets that are not accepted are DROPPED or REJECTED. Select from *Log All*, *Log Critical*, or *Do Not Log Any* for both of them.

When completed with the firewall configuration, exit this dialog with *Next*. A zone-oriented summary of your firewall configuration then opens. In it, check all settings. All services, ports, and protocols that have been allowed are listed in this summary. To modify the configuration, use *Back*. Press *Accept* to save your configuration.

Configuring Manually

The following paragraphs provide step-by-step instructions for a successful configuration. Each configuration item is marked as to whether it is relevant to firewalling or masquerading. Aspects related to the DMZ (demilitarized zone) as mentioned in the configuration file are not covered here. They are applicable only to a more complex network infrastructure found in larger organizations (corporate networks), which require extensive configuration and in-depth knowledge about the subject.

First, use the YaST module System Services (Runlevel) to enable SuSEfirewall2 in your runlevel (3 or 5 most likely). It sets the symlinks for the `SuSEfirewall2_*` scripts in the `/etc/init.d/rc?.d/` directories.

`FW_DEV_EXT` (firewall, masquerading)

The device linked to the Internet. For a modem connection, enter `ppp0`. For an ISDN link, use `ippp0`. DSL connections use `ds10`. Specify `auto` to use the interface that corresponds to the default route.

FW_DEV_INT (firewall, masquerading)

The device linked to the internal, private network (such as `eth0`). Leave this blank if there is no internal network and the firewall protects only the host on which it runs.

FW_ROUTE (firewall, masquerading)

If you need the masquerading function, set this to `yes`. Your internal hosts will not be visible to the outside, because their private network addresses (e.g., `192.168.x.x`) are ignored by Internet routers.

For a firewall without masquerading, only set this to `yes` if you want to allow access to the internal network. Your internal hosts need to use officially registered IPs in this case. Normally, however, you should *not* allow access to your internal network from the outside.

FW_MASQUERADE (masquerading)

Set this to `yes` if you need the masquerading function. This provides a virtually direct connection to the Internet for the internal hosts. It is more secure to have a proxy server between the hosts of the internal network and the Internet. Masquerading is not needed for services a proxy server provides.

FW_MASQ_NETS (masquerading)

Specify the hosts or networks to masquerade, leaving a space between the individual entries. For example:

```
FW_MASQ_NETS="192.168.0.0/24 192.168.10.1"
```

FW_PROTECT_FROM_INT (firewall)

Set this to `yes` to protect your firewall host from attacks originating in your internal network. Services are only available to the internal network if explicitly enabled.

Also see `FW_SERVICES_INT_TCP` and `FW_SERVICES_INT_UDP`.

FW_SERVICES_EXT_TCP (firewall)

Enter the TCP ports that should be made available. Leave this blank for a normal workstation at home that should not offer any services.

FW_SERVICES_EXT_UDP (firewall)

Leave this blank unless you run a UDP service and want to make it available to the outside. The services that use UDP include DNS servers, IPSec, TFTP, DHCP and others. In that case, enter the UDP ports to use.

`FW_SERVICES_INT_TCP` (firewall)

With this variable, define the services available for the internal network. The notation is the same as for `FW_SERVICES_EXT_TCP`, but the settings are applied to the *internal* network. The variable only needs to be set if `FW_PROTECT_FROM_INT` is set to `yes`.

`FW_SERVICES_INT_UDP` (firewall)

See `FW_SERVICES_INT_TCP`.

After configuring the firewall, test your setup. The firewall rule sets are created by entering `SuSEfirewall2 start` as root. Then use `telnet`, for example, from an external host to see whether the connection is actually denied. After that, review `/var/log/messages`, where you should see something like this:

```
Mar 15 13:21:38 linux kernel: SFW2-INext-DROP-DEFLT IN=eth0
OUT= MAC=00:80:c8:94:c3:e7:00:a0:c9:4d:27:56:08:00 SRC=192.168.10.0
DST=192.168.10.1 LEN=60 TOS=0x10 PREC=0x00 TTL=64 ID=15330 DF PROTO=TCP
SPT=48091 DPT=23 WINDOW=5840 RES=0x00 SYN URGP=0
OPT (020405B40402080A061AFEBCC0000000001030300)
```

Other packages to test your firewall setup are `nmap` or `nessus`. The documentation of `nmap` is found at `/usr/share/doc/packages/nmap` and the documentation of `nessus` resides in the directory `/usr/share/doc/packages/nessus-core` after installing the respective package.

4.1.5 For More Information

The most up-to-date information and other documentation about the `SuSEfirewall2` package is found in `/usr/share/doc/packages/SuSEfirewall2`. The home page of the netfilter and iptables project, <http://www.netfilter.org>, provides a large collection of documents in many languages.

4.2 SSH: Secure Network Operations

With more and more computers installed in networked environments, it often becomes necessary to access hosts from a remote location. This normally means that a user sends login and password strings for authentication purposes. As long as these strings are transmitted as plain text, they could be intercepted and misused to gain access to that user account without the authorized user even knowing about it. Apart from the fact

that this would open all the user's files to an attacker, the illegal account could be used to obtain administrator or `root` access or to penetrate other systems. In the past, remote connections were established with telnet, which offers no guards against eavesdropping in the form of encryption or other security mechanisms. There are other unprotected communication channels, like the traditional FTP protocol and some remote copying programs.

The SSH suite provides the necessary protection by encrypting the authentication strings (usually a login name and a password) and all the other data exchanged between the hosts. With SSH, the data flow could still be recorded by a third party, but the contents are encrypted and cannot be reverted to plain text unless the encryption key is known. So SSH enables secure communication over insecure networks, such as the Internet. The SSH flavor that comes with SUSE Linux is OpenSSH.

4.2.1 The OpenSSH Package

SUSE Linux installs the package OpenSSH by default. The programs `ssh`, `scp`, and `sftp` are then available as alternatives to `telnet`, `rlogin`, `rsh`, `rcp`, and `ftp`. In the default configuration, system access of a SUSE Linux system is only possible with the OpenSSH utilities and only if the firewall permits access.

4.2.2 The ssh Program

Using the `ssh` program, it is possible to log in to remote systems and work interactively. It replaces both `telnet` and `rlogin`. The `slogin` program is just a symbolic link pointing to `ssh`. For example, log in to the host `sun` with the command `ssh sun`. The host then prompts for the password on `sun`.

After successful authentication, you can work on the remote command line or use interactive applications, such as YaST. If the local username is different from the remote username, you can log in using a different login name with `ssh -l augustine sun` or `ssh augustine@sun`.

Furthermore, `ssh` offers the possibility to run commands on remote systems, as known from `rsh`. In the following example, run the command `uptime` on the host `sun` and create a directory with the name `tmp`. The program output is displayed on the local terminal of the host `earth`.

```
ssh otherplanet "uptime; mkdir tmp"
tux@otherplanet's password:
1:21pm up 2:17, 9 users, load average: 0.15, 0.04, 0.02
```

Quotation marks are necessary here to send both instructions with one command. It is only by doing this that the second command is executed on sun.

4.2.3 scp—Secure Copy

scp copies files to a remote machine. It is a secure and encrypted substitute for rcp. For example, `scp MyLetter.tex sun:` copies the file `MyLetter.tex` from the host earth to the host sun. If the username on earth is different than the username on sun, specify the latter using the `username@host` format. There is no `-l` option for this command.

After the correct password is entered, scp starts the data transfer and shows a growing row of asterisks to simulate a progress bar. In addition, the program displays the estimated time of arrival to the right of the progress bar. Suppress all output by giving the option `-q`.

scp also provides a recursive copying feature for entire directories. The command `scp -r src/ sun:` copies the entire contents of the directory `src` including all subdirectories to the `backup` directory on the host sun. If this subdirectory does not exist yet, it is created automatically.

The option `-p` tells scp to leave the time stamp of files unchanged. `-C` compresses the data transfer. This minimizes the data volume to transfer, but creates a heavier burden on the processor.

4.2.4 sftp—Secure File Transfer

The sftp program can be used instead of scp for secure file transfer. During an sftp session, you can use many of the commands known from ftp. The sftp program may be a better choice than scp, especially when transferring data for which the filenames are unknown.

4.2.5 The SSH Daemon (`sshd`)—Server-Side

To work with the SSH client programs `ssh` and `scp`, a server, the SSH daemon, must be running in the background, listening for connections on TCP/IP port 22. The daemon generates three key pairs when starting for the first time. Each key pair consist of a private and a public key. Therefore, this procedure is referred to as public key-based. To guarantee the security of the communication via SSH, access to the private key files must be restricted to the system administrator. The file permissions are set accordingly by the default installation. The private keys are only required locally by the SSH daemon and must not be given to anyone else. The public key components (recognizable by the name extension `.pub`) are sent to the client requesting the connection. They are readable for all users.

A connection is initiated by the SSH client. The waiting SSH daemon and the requesting SSH client exchange identification data to compare the protocol and software versions and to prevent connections through the wrong port. Because a child process of the original SSH daemon replies to the request, several SSH connections can be made simultaneously.

For the communication between SSH server and SSH client, OpenSSH supports versions 1 and 2 of the SSH protocol. A newly installed SUSE Linux system defaults to version 2. To continue using version 1 after an update, follow the instructions in `/usr/share/doc/packages/openssh/README.SuSE`. This document also describes how an SSH 1 environment can be transformed into a working SSH 2 environment with just a few steps.

When using version 1 of SSH, the server sends its public host key and a server key, which is regenerated by the SSH daemon every hour. Both allow the SSH client to encrypt a freely chosen session key, which is sent to the SSH server. The SSH client also tells the server which encryption method (cipher) to use.

Version 2 of the SSH protocol does not require a server key. Both sides use an algorithm according to Diffie-Helman to exchange their keys.

The private host and server keys are absolutely required to decrypt the session key and cannot be derived from the public parts. Only the SSH daemon contacted can decrypt the session key using its private keys (see `man /usr/share/doc/packages/openssh/RFC.nr off`). This initial connection phase can be watched closely by turning on the verbose debugging option `-v` of the SSH client.

Version 2 of the SSH protocol is used by default. Override this to use version 1 of the protocol with the `-1` switch. The client stores all public host keys in `~/.ssh/known_hosts` after its first contact with a remote host. This prevents any man-in-the-middle attacks—attempts by foreign SSH servers to use spoofed names and IP addresses. Such attacks are detected either by a host key that is not included in `~/.ssh/known_hosts` or by the server's inability to decrypt the session key in the absence of an appropriate private counterpart.

It is recommended to back up the private and public keys stored in `/etc/ssh/` in a secure, external location. In this way, key modifications can be detected and the old ones can be used again after a reinstallation. This spares users any unsettling warnings. If it is verified that, despite the warning, it is indeed the correct SSH server, the existing entry for the system must be removed from `~/.ssh/known_hosts`.

4.2.6 SSH Authentication Mechanisms

Now the actual authentication takes place, which, in its simplest form, consists of entering a password as mentioned above. The goal of SSH was to introduce a secure software that is also easy to use. Because it is meant to replace `rsh` and `rlogin`, SSH must also be able to provide an authentication method appropriate for daily use. SSH accomplishes this by way of another key pair, which is generated by the user. The SSH package provides a helper program for this: `ssh-keygen`. After entering `ssh-keygen -t rsa` or `ssh-keygen -t dsa`, the key pair is generated and you are prompted for the base filename in which to store the keys.

Confirm the default setting and answer the request for a passphrase. Even if the software suggests an empty passphrase, a text from 10 to 30 characters is recommended for the procedure described here. Do not use short and simple words or phrases. Confirm by repeating the passphrase. Subsequently, you will see where the private and public keys are stored, in this example, the files `id_rsa` and `id_rsa.pub`.

Use `ssh-keygen -p -t rsa` or `ssh-keygen -p -t dsa` to change your old passphrase. Copy the public key component (`id_rsa.pub` in the example) to the remote machine and save it to `~/.ssh/authorized_keys`. You will be asked to authenticate yourself with your passphrase the next time you establish a connection. If this does not occur, verify the location and contents of these files.

In the long run, this procedure is more troublesome than giving your password each time. Therefore, the SSH package provides another tool, `ssh-agent`, which retains the

private keys for the duration of an X session. The entire X session is started as a child process of ssh-agent. The easiest way to do this is to set the variable `usessh` at the beginning of the `.xsession` file to `yes` and log in via a display manager, such as KDM or XDM. Alternatively, enter `ssh-agent startx`.

Now you can use ssh or scp as usual. If you have distributed your public key as described above, you are no longer prompted for your password. Take care of terminating your X session or locking it with a password protection application, such as `xlock`.

All the relevant changes that resulted from the introduction of version 2 of the SSH protocol are also documented in the file `/usr/share/doc/packages/openssh/README.SuSE`.

4.2.7 X, Authentication, and Forwarding Mechanisms

Beyond the previously described security-related improvements, SSH also simplifies the use of remote X applications. If you run `ssh` with the option `-X`, the `DISPLAY` variable is automatically set on the remote machine and all X output is exported to the remote machine over the existing SSH connection. At the same time, X applications started remotely and locally viewed with this method cannot be intercepted by unauthorized individuals.

By adding the option `-A`, the ssh-agent authentication mechanism is carried over to the next machine. This way, you can work from different machines without having to enter a password, but only if you have distributed your public key to the destination hosts and properly saved it there.

Both mechanisms are deactivated in the default settings, but can be permanently activated at any time in the systemwide configuration file `/etc/ssh/sshd_config` or the user's `~/.ssh/config`.

ssh can also be used to redirect TCP/IP connections. In the examples below, SSH is told to redirect the SMTP and the POP3 port, respectively:

```
ssh -L 25:sun:25 earth
```

With this command, any connection directed to `earth` port 25 (SMTP) is redirected to the SMTP port on `sun` via an encrypted channel. This is especially useful for those using

SMTP servers without SMTP-AUTH or POP-before-SMTP features. From any arbitrary location connected to a network, e-mail can be transferred to the “home” mail server for delivery. Similarly, all POP3 requests (port 110) on earth can be forwarded to the POP3 port of sun with this command:

```
ssh -L 110:sun:110 earth
```

Both commands must be executed as `root`, because the connection is made to privileged local ports. E-mail is sent and retrieved by normal users in an existing SSH connection. The SMTP and POP3 host must be set to `localhost` for this to work. Additional information can be found in the manual pages for each of the programs described above and also in the files under `/usr/share/doc/packages/openssh`.

4.3 Encrypting Partitions and Files

Every user has some confidential data that third parties should not be able to access. The more connected and mobile you are, the more carefully you should handle your data. The encryption of files or entire partitions is recommended if others have access over a network connection or direct physical access.

WARNING: Encrypted Media Is Limited Protection

Be aware that with the methods described in this section, you cannot protect your running system from being compromised. After the encrypted media is successfully mounted, everybody with appropriate permissions has access to it. Encrypted media makes sense if you lose your computer or it is stolen and unauthorized individuals want to read your confidential data.

The following list features a number of imaginable usage scenarios.

Laptops

If you travel with your laptop, it is a good idea to encrypt hard disk partitions containing confidential data. If you lose your laptop or if it is stolen, your data will be out of reach if it resides in an encrypted file system or a single encrypted file.

Removable Media

USB flash drives or external hard disks are as prone to being stolen as laptops. An encrypted file system provides protection against third-party access.

Workstations

In companies where almost everyone has access to your computer, it can make sense to encrypt partition or single files.

4.3.1 Setting Up a Crypto File System with YaST

YaST offers the encryption of files or partitions during installation as well as in an already installed system. An encrypted file can be created at any time, because it fits nicely in an existing partition layout. To encrypt an entire partition, dedicate a partition for encryption in the partition layout. The standard partitioning proposal as suggested by YaST does not, by default, include an encrypted partition. Add it manually in the partitioning dialog.

Creating an Encrypted Partition during Installation

WARNING: Password Input

Observe the warnings about password security when setting the password for encrypted partitions and memorize it well. Without the password, the encrypted data cannot be accessed or restored.

The YaST expert dialog for partitioning, described in Section “Partitioner” (Chapter 2, *System Configuration with YaST*, ↑Start-Up), offers the options needed for creating an encrypted partition. Click *Create* like when creating a regular partition. In the dialog that opens, enter the partitioning parameters for the new partition, such as the desired formatting and the mount point. Complete the process by clicking *Encrypt File System*. In the following dialog, enter the password twice. The new encrypted partition is created after the partitioning dialog is closed by clicking *OK*. While booting, the operating system requests the password before mounting the partition.

If you do not want to mount the encrypted partition during start-up, click `Enter` when prompted for the password. Then decline the offer to enter the password again. In this case, the encrypted file system is not mounted and the operating system continues booting, blocking access to your data. The partition is available to all users once it has been mounted.

If the encrypted file system should only be mounted when necessary, enable *Do Not Mount During Booting* in the *fstab Options* dialog. The respective partition will not be mounted when the system is booted. To make it available afterwards, mount it manually with `mount name_of_partition mount_point`. Enter the password when prompted to do so. After finishing your work with the partition, unmount it with `umount name_of_partition` to protect it from access by other users.

Creating an Encrypted Partition on a Running System

WARNING: Activating Encryption in a Running System

It is also possible to create encrypted partitions on a running system like during installation. However, encrypting an existing partition destroys all data on it.

On a running system, select *System → Partitioning* in the YaST control center. Click *Yes* to proceed. Instead of selecting *Create* as mentioned above, click *Edit*. The rest of the procedure is the same.

Installing Encrypted Files

Instead of using a partition, it is possible to create encrypted file systems within single files for holding confidential data. These are created from the same YaST dialog. Select *Crypt File* and enter the path to the file to create along with its intended size. Accept the proposed formatting settings and the file system type. Then specify the mount point and decide whether the encrypted file system should be mounted when the system is booted.

The advantage of encrypted files is that they can be added without repartitioning the hard disk. They are mounted with the help of a loop device and behave just like normal partitions.

Using vi to Encrypt Files

The disadvantage of using encrypted partitions is that while the partition is mounted, at least `root` can access the data. To prevent this, `vi` can be used in encrypted mode.

Use `vi -x filename` to edit a new file. vi prompts you to set a password, after which it encrypts the content of the file. Whenever you access this file, vi requests the correct password.

For even more security, you can place the encrypted text file in an encrypted partition. This is recommended because the encryption used in vi is not very strong.

4.3.2 Encrypting the Content of Removable Media

YaST treats removable media like external hard disks or USB flash drives like any other hard disk. Files or partitions on such media can be encrypted as described above. However, do not select to mount these media when the system is booted, because they are usually only connected while the system is running.

4.4 Confining Privileges with AppArmor

Many security vulnerabilities result from bugs in *trusted* programs. A *trusted* program runs with privilege that some attacker would like to have and the program fails to keep that trust if there is a bug in the program that allows the attacker to acquire that privilege.

Novell® AppArmor is an application security solution designed specifically to provide least privilege confinement to suspect programs. AppArmor allows the administrator to specify the domain of activities the program can perform by developing a security *profile* for that application—a listing of files that the program may access and the operations the program may perform.

Effective hardening of a computer system requires minimizing the number of programs that mediate privilege then securing the programs as much as possible. With Novell AppArmor, you only need to profile the programs that are exposed to attack in your environment, which drastically reduces the amount of work required to harden your computer. AppArmor profiles enforce policies to make sure that programs do what they are supposed to do, but nothing else.

Administrators only need to care about the applications that are vulnerable to attacks and generate profiles for these. Hardening a system thus comes down to building and maintaining the AppArmor profile set and monitoring any policy violations or exceptions logged by AppArmor's reporting facility.

Building AppArmor profiles to confine an application is very straightforward and intuitive. AppArmor ships with several tools that assist in profile creation. AppArmor does not require you to do any programming or script handling. The only task that is required from the administrator is to determine a policy of strictest access and execute permissions for each application that needs to be hardened.

Updates or modifications to the application profiles are only required if the software configuration or the desired range of activities changes. AppArmor offers intuitive tools to handle profile updates or modifications.

Users will not notice AppArmor at all. It runs “behind the scenes” and does not require any user interaction. Performance will not be affected noticeably by AppArmor. If some activity of the application is not covered by an AppArmor profile or if some activity of the application is prevented by AppArmor, the administrator needs to adjust the profile of this application to cover this kind of behavior.

This guide outlines the basic tasks that need to be performed with AppArmor to effectively harden a system. For more in-depth information, refer to *Novell AppArmor 2.0 Administration Guide*.

4.4.1 Installing Novell AppArmor

Users installing either a GNOME or the KDE desktop selection can skip this section, because Novell AppArmor is installed by default as part of these selections.

If installing neither of these desktops or even going for an entirely text-based environment, do the following to install the required packages using the YaST package manager.

- 1** Log in as `root` and start YaST.
- 2** In the YaST Control Center, select *Software* → *Software Management*.
- 3** Use the search functionality of YaST (keywords “AppArmor”) to install the following packages:

- apparmor-parser
- libapparmor
- apparmor-docs
- yast2-apparmor
- apparmor-profiles
- apparmor-utils

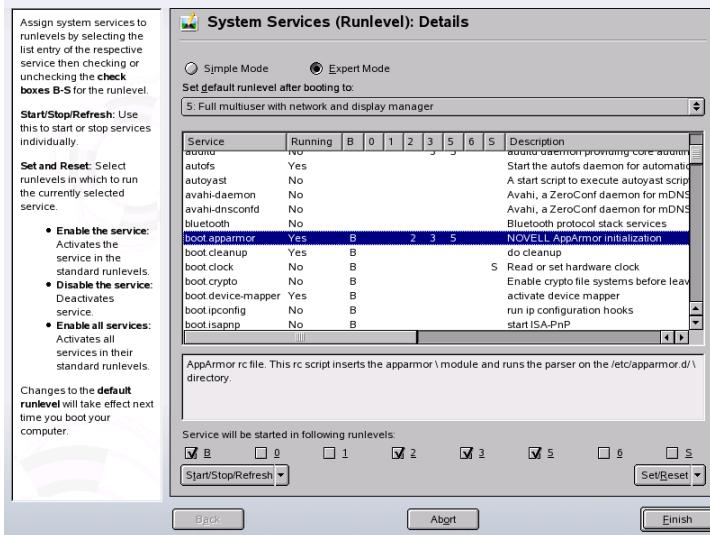
- 4 Select all these packages for installation then select *Accept*. YaST resolves any dependencies and installs all the packages for you.
- 5 After YaST has finished updating the system configuration, select *Finish* to leave the package manager.

4.4.2 Enabling Novell AppArmor

After Novell AppArmor has been installed, explicitly enable it to make sure that it is started whenever your system boots. Use the YaST System Services (Runlevel) module for this task:

- 1 Log in as `root` and start YaST.
- 2 Start *System* → *System Services (Runlevel)*.
- 3 In the list of services displayed, select `apparmor`. See [Figure 4.3, “Enabling Novell AppArmor Using YaST”](#) (page 122).
- 4 Click *Enable* to enable AppArmor permanently.
- 5 Click *Finish* to accept your settings.

Figure 4.3 Enabling Novell AppArmor Using YaST



Using the YaST Runlevel tool, you enable services permanently—these settings survive a reboot of your system. To enable AppArmor temporarily—for the duration of one session only—proceed as follows:

- 1 Log in as `root` and start YaST.
- 2 Start *Novell AppArmor* → *AppArmor Control Panel*.
- 3 Set the *AppArmor Status* to *AppArmor is enabled* by clicking *Configure* → *Enable* → *OK*.
- 4 Apply your settings with *Done*.

4.4.3 Getting Started with Profiling Applications

Prepare a successful deployment of Novell AppArmor on your system by carefully considering the following items:

- 1 Determine the applications to profile. Read more on this in [Section “Choosing the Applications to Profile”](#) (page 123).
- 2 Build the needed profiles as roughly outlined in [Section “Building and Modifying Profiles”](#) (page 124). Check the results and adjust the profiles when necessary.
- 3 Keep track of what is happening on your system by running AppArmor reports and dealing with security events. Refer to [Section “Configuring Novell AppArmor Event Notification and Reports”](#) (page 126).
- 4 Update your profiles whenever your environment changes or you need to react to security events logged by AppArmor's reporting tool. Refer to [Section “Updating Your Profiles”](#) (page 127).

Choosing the Applications to Profile

You only need to protect the programs that are exposed to attacks in your particular setup, so only use profiles for those applications you really run. Use the following list to determine the most likely candidates:

Network Agents

Programs (servers and clients) have open network ports and network agents are server programs that respond to those network ports. User clients (such as mail clients and Web browsers) also have open network ports and mediate privilege.

Web Applications

CGI Perl scripts, PHP pages, and more complex Web applications can be invoked through a Web browser.

Cron Jobs

Programs that the cron daemon periodically runs read input from a variety of sources.

To find out which processes are currently running with open network ports and might need a profile to confine them, run `unconfined` as root.

Example 4.1 Output of unconfined

```
19848 /usr/sbin/cupsd not confined
19887 /usr/sbin/sshd not confined
19947 /usr/lib/postfix/master not confined
29205 /usr/sbin/sshd confined by '/usr/sbin/sshd (enforce)'
```

Each of the processes in the above example labeled `not confined` might need a custom profile to confine it. Those labeled `confined by` are already protected by AppArmor.

TIP: For More Information

For more information about choosing the the right applications to profile, refer to Chapter *Selecting Programs to Immunize (Novell AppArmor 2.0 Administration Guide)*.

Building and Modifying Profiles

Novell AppArmor on SUSE Linux ships with a preconfigured set of profiles for the most important applications. In addition to that, you can use AppArmor to create your own profiles for a set of applications defined in `/etc/apparmor/README.profiles`.

There are two ways of managing profiles. One is to use the graphical front-end provided by the YaST Novell AppArmor modules and the other is to use the command line tools provided by the AppArmor suite itself. Both methods basically work the same way.

Running `unconfined` as described in [Section “Choosing the Applications to Profile”](#) (page 123) identifies a list of applications that may need a profile to run in a safe mode.

For each application, perform the following steps to create a profile:

- 1 As `root`, let AppArmor create a rough outline of the application's profile by running `genprof programname`

or

running *YaST → Novell AppArmor → Add Profile Wizard* and specifying the complete path of the application to profile.

A basic profile is outlined and AppArmor is put into learning mode, which means that it logs any activity of the program you are executing but does not restrict it, yet.

- 2 Run the full range of the application's actions to let AppArmor get a very specific picture of its activities.

- 3** Let AppArmor analyze the log files generated in [Step 2](#) (page 124). Do this either by running typing **[S]** in genprof

or

clicking *Scan system log for AppArmor events* in the *Add Profile Wizard* and follow the instructions given in the wizard until the profile is completed.

AppArmor scans the logs it recorded during the application's run and asks you to set the access rights for each event that was logged. Either set them for each file or use globbing.

- 4** Once all access permissions are set, your profile is set to enforce mode mode. The profile is applied and AppArmor restricts the application according to the profile just created.

If you started genprof against an application that had an existing profile that was in complain mode, this profile will remain in learning mode upon exit of this learning cycle. For more information on changing the mode of a profile, refer to Section *Complain or Learning Mode* ([Chapter 3, Building Novell AppArmor Profiles, Novell AppArmor 2.0 Administration Guide](#)) and to Section *Enforce Mode* ([Chapter 3, Building Novell AppArmor Profiles, Novell AppArmor 2.0 Administration Guide](#)).

Test your profile settings by performing every task you need with the application you just confined. Normally, the confined program runs smoothly and you do not notice AppArmor activities at all. However, if you notice certain misbehavior with your application, check the system logs and see if AppArmor is too closely constricting your application. Find the appropriate logs in `/var/log/messages` or run `dmesg`.

Any output resembling the following example hints at AppArmor too closely confining your application:

```
AppArmor: REJECTING w access to /var/run/nscd/socket (traceroute(2050) profile  
/usr/sbin/traceroute active /usr/sbin/traceroute)
```

To adjust the profile, run the *Add Profile Wizard* again as described above and let it analyze the log messages relating this particular application. Determine the access rights or restrictions when prompted by YaST.

TIP: For More Information

For more information about profile building and modification, refer to Chapter *Building Novell AppArmor Profiles (Novell AppArmor 2.0 Administration Guide)*.

Configuring Novell AppArmor Event Notification and Reports

Set up event notification in Novell AppArmor so you can review security events. Event Notification is an Novell AppArmor feature that informs a specified e-mail recipient when systemic Novell AppArmor activity occurs under the chosen severity level. This feature is currently available via the YaST interface.

To set up event notification in YaST, proceed as follows:

- 1** Make sure that a mail server is running on your system to deliver the event notifications.
- 2** Log in as `root` and start YaST. Then select *Novell AppArmor → AppArmor Control Panel*.
- 3** In *Enable Security Event Notification* section, select *Configure*.
- 4** For each record type (*Terse*, *Summary*, and *Verbose*) set a report frequency, enter the e-mail address to receive the reports and determine the severity of events to log. If you want to include unknown events in the event reports, check *Include Unknown Severity Events*.

NOTE

Unless you are familiar with AppArmor's event categorization, choose to be notified about events for all security levels.

- 5** Leave this dialog with *OK → Finish* to apply your settings.

Configure Novell AppArmor reports. Using reports, you can read important Novell AppArmor security events reported in the log files without manually sifting through

the cumbersome messages only useful to the logprof tool. You can decrease the size of the report by filtering by date range or program name.

To configure the AppArmor reports, proceed as follows:

- 1 Log in as `root` and start YaST. Select *Novell AppArmor* → *AppArmor Reports*.
- 2 Select the type of report you want to examine or configure from *Executive Security Summary*, *Applications Audit*, and *Security Incident Report*.
- 3 Edit the report generation frequency, e-mail address, export format, and the location of the reports by selecting *Edit* and providing the requested data.
- 4 To run a report of the selected type, click *Run Now*.
- 5 Browse through the archived reports of a given type by selecting *View Archive* and specifying the report type.

or

Delete unneeded reports or add new ones.

TIP: For More Information

For more information about configuring event notification in Novell AppArmor, refer to Section “Setting Up Event Notification” (Chapter 4, *Managing Profiled Applications, Novell AppArmor 2.0 Administration Guide*). More information about report configuration can be found in Section “Reports” (Chapter 4, *Managing Profiled Applications, Novell AppArmor 2.0 Administration Guide*) .

Updating Your Profiles

Software and system configurations change over time. As a result of that your profile setup for AppArmor might need some fine-tuning from time to time. AppArmor checks your system log for policy violations or other AppArmor events and lets you adjust your profile set accordingly. Any application behavior that is outside of any profile definition can also be addressed using the *Update Profile Wizard*.

To update your profile set, proceed as follows:

- 1 Log in as `root` and start YaST.
- 2 Start *Novell AppArmor* → *Update Profile Wizard*.
- 3 Adjust access or execute rights to any resource or for any executable that has been logged when prompted.
- 4 Leave YaST after you answered all questions. Your changes are applied to the respective profiles.

TIP: For More Information

For more information about updating your profiles from the system logs, refer to Section “Updating Profiles from Syslog Entries” (Chapter 3, *Building Novell AppArmor Profiles, Novell AppArmor 2.0 Administration Guide*).

4.5 Security and Confidentiality

One of the main characteristics of a Linux or UNIX system is its ability to handle several users at the same time (multiuser) and to allow these users to perform several tasks (multitasking) on the same computer simultaneously. Moreover, the operating system is network transparent. The users often do not know whether the data and applications they are using are provided locally from their machine or made available over the network.

With the multiuser capability, the data of different users must be stored separately. Security and privacy need to be guaranteed. Data security was already an important issue, even before computers could be linked through networks. Just like today, the most important concern was the ability to keep data available in spite of a lost or otherwise damaged data medium, a hard disk in most cases.

This section is primarily focused on confidentiality issues and on ways to protect the privacy of users, but it cannot be stressed enough that a comprehensive security concept should always include procedures to have a regularly updated, workable, and tested backup in place. Without this, you could have a very hard time getting your data back—not only in the case of some hardware defect, but also if the suspicion arises that someone has gained unauthorized access and tampered with files.

4.5.1 Local Security and Network Security

There are several ways of accessing data:

- personal communication with people who have the desired information or access to the data on a computer
- directly from the console of a computer (physical access)
- over a serial line
- using a network link

In all these cases, a user should be authenticated before accessing the resources or data in question. A Web server might be less restrictive in this respect, but you still would not want it to disclose all your personal data to any surfer.

In the list above, the first case is the one where the highest amount of human interaction is involved, such as when you are contacting a bank employee and are required to prove that you are the person owning that bank account. Then you are asked to provide a signature, a PIN, or a password to prove that you are the person you claim to be. In some cases, it might be possible to elicit some intelligence from an informed person just by mentioning known bits and pieces to win the confidence of that person by using clever rhetoric. The victim could be led to reveal gradually more information, maybe without even becoming aware of it. Among hackers, this is called *social engineering*. You can only guard against this by educating people and by dealing with language and information in a conscious way. Before breaking into computer systems, attackers often try to target receptionists, service people working with the company, or even family members. In many cases, such an attack based on social engineering is only discovered at a much later time.

A person wanting to obtain unauthorized access to your data could also use the traditional way and try to get at your hardware directly. Therefore, the machine should be protected against any tampering so that no one can remove, replace, or cripple its components. This also applies to backups and even any network cable or the power cord. Also secure the boot procedure, because there are some well-known key combinations that might provoke unusual behavior. Protect yourself against this by setting passwords for the BIOS and the boot loader.

Serial terminals connected to serial ports are still used in many places. Unlike network interfaces, they do not rely on a network protocol to communicate with the host. A simple cable or an infrared port is used to send plain characters back and forth between the devices. The cable itself is the weakest point of such a system: with an older printer connected to it, it is easy to record anything that runs over the wires. What can be achieved with a printer can also be accomplished in other ways, depending on the effort that goes into the attack.

Reading a file locally on a host requires other access rules than opening a network connection with a server on a different host. There is a distinction between local security and network security. The line is drawn where data must be put into packets to be sent somewhere else.

Local Security

Local security starts with the physical environment in the location where the computer is running. Set up your machine in a place where security is in line with your expectations and needs. The main goal of local security is to keep users separate from each other, so no user can assume the permissions or the identity of another. This is a general rule to be observed, but it is especially true for the user `root`, who holds the supreme power on the system. `root` can take on the identity of any other local user without being prompted for the password and read any locally stored file.

Passwords

On a Linux system, passwords are not stored as plain text and the text string entered is not simply matched with the saved pattern. If this were the case, all accounts on your system would be compromised as soon as someone got access to the corresponding file. Instead, the stored password is encrypted and, each time it is entered, is encrypted again and the two encrypted strings are compared. This only provides more security if the encrypted password cannot be reverse-computed into the original text string.

This is actually achieved by a special kind of algorithm, also called *trapdoor algorithm*, because it only works in one direction. An attacker who has obtained the encrypted string is not able to get your password by simply applying the same algorithm again. Instead, it would be necessary to test all the possible character combinations until a combination is found that looks like your password when encrypted. With passwords eight characters long, there are quite a number of possible combinations to calculate.

In the seventies, it was argued that this method would be more secure than others due to the relative slowness of the algorithm used, which took a few seconds to encrypt just one password. In the meantime, however, PCs have become powerful enough to do several hundred thousand or even millions of encryptions per second. Because of this, encrypted passwords should not be visible to regular users (`/etc/shadow` cannot be read by normal users). It is even more important that passwords are not easy to guess, in case the password file becomes visible due to some error. Consequently, it is not really useful to “translate” a password like “tantalize” into “t@nt@1lz3”.

Replacing some letters of a word with similar looking numbers is not safe enough. Password cracking programs that use dictionaries to guess words also play with substitutions like that. A better way is to make up a word with no common meaning, something that only makes sense to you personally, like the first letters of the words of a sentence or the title of a book, such as “The Name of the Rose” by Umberto Eco. This would give the following safe password: “TNotRbUE9”. In contrast, passwords like “beerbuddy” or “jasmine76” are easily guessed even by someone who has only some casual knowledge about you.

The Boot Procedure

Configure your system so it cannot be booted from a floppy or from CD, either by removing the drives entirely or by setting a BIOS password and configuring the BIOS to allow booting from a hard disk only. Normally, a Linux system is started by a boot loader, allowing you to pass additional options to the booted kernel. Prevent others from using such parameters during boot by setting an additional password in `/boot/grub/menu.lst` (see [Chapter 9, The Boot Loader](#) (page 193)). This is crucial to your system's security. Not only does the kernel itself run with `root` permissions, but it is also the first authority to grant `root` permissions at system start-up.

File Permissions

As a general rule, always work with the most restrictive privileges possible for a given task. For example, it is definitely not necessary to be `root` to read or write e-mail. If the mail program has a bug, this bug could be exploited for an attack that acts with exactly the permissions of the program when it was started. By following the above rule, minimize the possible damage.

The permissions of the more than 200,000 files included in a SUSE distribution are carefully chosen. A system administrator who installs additional software or other files

should take great care when doing so, especially when setting the permission bits. Experienced and security-conscious system administrators always use the `-l` option with the command `ls` to get an extensive file list, which allows them to detect any incorrect file permissions immediately. An incorrect file attribute does not only mean that files could be changed or deleted. These modified files could be executed by `root` or, in the case of configuration files, programs could use such files with the permissions of `root`. This significantly increases the possibilities of an attacker. Attacks like this are called cuckoo eggs, because the program (the egg) is executed (hatched) by a different user (bird), just like a cuckoo tricks other birds into hatching its eggs.

A SUSE Linux system includes the files `permissions.permissions.easy`, `permissions.secure`, and `permissions.paranoid`, all in the directory `/etc`. The purpose of these files is to define special permissions, such as world-writable directories or, for files, the setuser ID bit (programs with the setuser ID bit set do not run with the permissions of the user that has launched it, but with the permissions of the file owner, in most cases `root`). An administrator can use the file `/etc/permissions.local` to add his own settings.

To define which of the above files is used by SUSE's configuration programs to set permissions accordingly, select *Security* in YaST. To learn more about the topic, read the comments in `/etc/permissions` or consult the manual page of `chmod` (`man chmod`).

Buffer Overflows and Format String Bugs

Special care must be taken whenever a program is supposed to process data that can or could be changed by a user, but this is more of an issue for the programmer of an application than for regular users. The programmer must make sure that his application interprets data in the correct way, without writing it into memory areas that are too small to hold it. Also, the program should hand over data in a consistent manner, using the interfaces defined for that purpose.

A *buffer overflow* can happen if the actual size of a memory buffer is not taken into account when writing to that buffer. There are cases where this data (as generated by the user) uses up some more space than what is available in the buffer. As a result, data is written beyond the end of that buffer area, which, under certain circumstances, makes it possible for a program to execute program sequences influenced by the user (and not by the programmer), rather than just processing user data. A bug of this kind may have

serious consequences, especially if the program is being executed with special privileges (see [Section “File Permissions”](#) (page 131)).

Format string bugs work in a slightly different way, but again it is the user input that could lead the program astray. In most cases, these programming errors are exploited with programs executed with special permissions—setuid and setgid programs—which also means that you can protect your data and your system from such bugs by removing the corresponding execution privileges from programs. Again, the best way is to apply a policy of using the lowest possible privileges (see [Section “File Permissions”](#) (page 131)).

Given that buffer overflows and format string bugs are bugs related to the handling of user data, they are not only exploitable if access has been given to a local account. Many of the bugs that have been reported can also be exploited over a network link. Accordingly, buffer overflows and format string bugs should be classified as being relevant for both local and network security.

Viruses

Contrary to what some people say, there are viruses that run on Linux. However, the viruses that are known were released by their authors as a *proof of concept* to prove that the technique works as intended. None of these viruses have been spotted *in the wild* so far.

Viruses cannot survive and spread without a host on which to live. In this case, the host would be a program or an important storage area of the system, such as the master boot record, which needs to be writable for the program code of the virus. Owing to its multiuser capability, Linux can restrict write access to certain files, especially important with system files. Therefore, if you did your normal work with `root` permissions, you would increase the chance of the system being infected by a virus. In contrast, if you follow the principle of using the lowest possible privileges as mentioned above, chances of getting a virus are slim.

Apart from that, you should never rush into executing a program from some Internet site that you do not really know. SUSE's RPM packages carry a cryptographic signature as a digital label that the necessary care was taken to build them. Viruses are a typical sign that the administrator or the user lacks the required security awareness, putting at risk even a system that should be highly secure by its very design.

Viruses should not be confused with worms, which belong to the world of networks entirely. Worms do not need a host to spread.

Network Security

Network security is important for protecting from an attack that is started outside. The typical login procedure requiring a username and a password for user authentication is still a local security issue. In the particular case of logging in over a network, differentiate between the two security aspects. What happens until the actual authentication is network security and anything that happens afterwards is local security.

X Window System and X Authentication

As mentioned at the beginning, network transparency is one of the central characteristics of a UNIX system. X, the windowing system of UNIX operating systems, can make use of this feature in an impressive way. With X, it is basically no problem to log in at a remote host and start a graphical program that is then sent over the network to be displayed on your computer.

When an X client should be displayed remotely using an X server, the latter should protect the resource managed by it (the display) from unauthorized access. In more concrete terms, certain permissions must be given to the client program. With the X Window System, there are two ways to do this, called host-based access control and cookie-based access control. The former relies on the IP address of the host where the client should run. The program to control this is `xhost`. `xhost` enters the IP address of a legitimate client into a tiny database belonging to the X server. However, relying on IP addresses for authentication is not very secure. For example, if there were a second user working on the host sending the client program, that user would have access to the X server as well—just like someone stealing the IP address. Because of these shortcomings, this authentication method is not described in more detail here, but you can learn about it with `man xhost`.

In the case of cookie-based access control, a character string is generated that is only known to the X server and to the legitimate user, just like an ID card of some kind. This cookie (the word goes back not to ordinary cookies, but to Chinese fortune cookies, which contain an epigram) is stored on login in the file `.Xauthority` in the user's home directory and is available to any X client wanting to use the X server to display a window. The file `.Xauthority` can be examined by the user with the tool `xauth`. If you were to rename `.Xauthority` or if you deleted the file from your home direc-

tory by accident, you would not be able to open any new windows or X clients. Read more about X Window System security mechanisms in the man page of `Xsecurity` (`man Xsecurity`).

SSH (secure shell) can be used to encrypt a network connection completely and forward it to an X server transparently without the encryption mechanism being perceived by the user. This is also called X forwarding. X forwarding is achieved by simulating an X server on the server side and setting a `DISPLAY` variable for the shell on the remote host. Further details about SSH can be found in [Section 4.2, “SSH: Secure Network Operations”](#) (page 110).

WARNING

If you do not consider the host where you log in to be a secure host, do not use X forwarding. With X forwarding enabled, an attacker could authenticate via your SSH connection to intrude on your X server and sniff your keyboard input, for instance.

Buffer Overflows and Format String Bugs

As discussed in [Section “Buffer Overflows and Format String Bugs”](#) (page 132), buffer overflows and format string bugs should be classified as issues concerning both local and network security. As with the local variants of such bugs, buffer overflows in network programs, when successfully exploited, are mostly used to obtain `root` permissions. Even if that is not the case, an attacker could use the bug to gain access to an unprivileged local account to exploit any other vulnerabilities that might exist on the system.

Buffer overflows and format string bugs exploitable over a network link are certainly the most frequent form of remote attacks in general. Exploits for these—programs to exploit these newly-found security holes—are often posted on the security mailing lists. They can be used to target the vulnerability without knowing the details of the code. Over the years, experience has shown that the availability of exploit codes has contributed to more secure operating systems, obviously due to the fact that operating system makers were forced to fix the problems in their software. With free software, anyone has access to the source code (SUSE Linux comes with all available source codes) and anyone who finds a vulnerability and its exploit code can submit a patch to fix the corresponding bug.

Denial of Service

The purpose of a denial of service (DoS) attack is to block a server program or even an entire system, something that could be achieved by various means: overloading the server, keeping it busy with garbage packets, or exploiting a remote buffer overflow. Often a DoS attack is made with the sole purpose of making the service disappear. However, once a given service has become unavailable, communications could become vulnerable to *man-in-the-middle attacks* (sniffing, TCP connection hijacking, spoofing) and DNS poisoning.

Man in the Middle: Sniffing, Hijacking, Spoofing

In general, any remote attack performed by an attacker who puts himself between the communicating hosts is called a *man-in-the-middle attack*. What almost all types of man-in-the-middle attacks have in common is that the victim is usually not aware that there is something happening. There are many possible variants, for example, the attacker could pick up a connection request and forward that to the target machine. Now the victim has unwittingly established a connection with the wrong host, because the other end is posing as the legitimate destination machine.

The simplest form of a man-in-the-middle attack is called *sniffer*—the attacker is “just” listening to the network traffic passing by. As a more complex attack, the “man in the middle” could try to take over an already established connection (hijacking). To do so, the attacker would need to analyze the packets for some time to be able to predict the TCP sequence numbers belonging to the connection. When the attacker finally seizes the role of the target host, the victims notice this, because they get an error message saying the connection was terminated due to a failure. The fact that there are protocols not secured against hijacking through encryption, which only perform a simple authentication procedure upon establishing the connection, makes it easier for attackers.

Spoofing is an attack where packets are modified to contain counterfeit source data, usually the IP address. Most active forms of attack rely on sending out such fake packets—something that, on a Linux machine, can only be done by the superuser (`root`).

Many of the attacks mentioned are carried out in combination with a DoS. If an attacker sees an opportunity to bring down a certain host abruptly, even if only for a short time, it makes it easier for him to push the active attack, because the host will not be able to interfere with the attack for some time.

DNS Poisoning

DNS poisoning means that the attacker corrupts the cache of a DNS server by replying to it with spoofed DNS reply packets, trying to get the server to send certain data to a victim who is requesting information from that server. Many servers maintain a trust relationship with other hosts, based on IP addresses or hostnames. The attacker needs a good understanding of the actual structure of the trust relationships among hosts to disguise itself as one of the trusted hosts. Usually, the attacker analyzes some packets received from the server to get the necessary information. The attacker often needs to target a well-timed DoS attack at the name server as well. Protect yourself by using encrypted connections that are able to verify the identity of the hosts to which to connect.

Worms

Worms are often confused with viruses, but there is a clear difference between the two. Unlike viruses, worms do not need to infect a host program to live. Instead, they are specialized to spread as quickly as possible on network structures. The worms that appeared in the past, such as Ramen, Lion, or Adore, make use of well-known security holes in server programs like bind8 or lprNG. Protection against worms is relatively easy. Given that some time elapses between the discovery of a security hole and the moment the worm hits your server, there is a good chance that an updated version of the affected program is available on time. That is only useful if the administrator actually installs the security updates on the systems in question.

4.5.2 Some General Security Tips and Tricks

To handle security competently, it is important to keep up with new developments and stay informed about the latest security issues. One very good way to protect your systems against problems of all kinds is to get and install the updated packages recommended by security announcements as quickly as possible. SUSE security announcements are published on a mailing list to which you can subscribe by following the link <http://www.novell.com/linux/security/securitysupport.html>. The list suse-security-announce@suse.de is a first-hand source of information regarding updated packages and includes members of SUSE's security team among its active contributors.

The mailing list suse-security@suse.de is a good place to discuss any security issues of interest. Subscribe to it on the same Web page.

bugtraq@securityfocus.com is one of the best-known security mailing lists worldwide. Reading this list, which receives between 15 and 20 postings per day, is recommended. More information can be found at <http://www.securityfocus.com>.

The following is a list of rules you may find useful in dealing with basic security concerns:

- According to the rule of using the most restrictive set of permissions possible for every job, avoid doing your regular jobs as `root`. This reduces the risk of getting a cuckoo egg or a virus and protects you from your own mistakes.
- If possible, always try to use encrypted connections to work on a remote machine. Using `ssh` (secure shell) to replace `telnet`, `ftp`, `rsh`, and `rlogin` should be standard practice.
- Avoid using authentication methods based on IP addresses alone.
- Try to keep the most important network-related packages up-to-date and subscribe to the corresponding mailing lists to receive announcements on new versions of such programs (`bind`, `sendmail`, `ssh`, etc.). The same should apply to software relevant to local security.
- Change the `/etc/permissions` file to optimize the permissions of files crucial to your system's security. If you remove the setuid bit from a program, it might well be that it cannot do its job anymore in the intended way. On the other hand, consider that, in most cases, the program will also have ceased to be a potential security risk. You might take a similar approach with world-writable directories and files.
- Disable any network services you do not absolutely require for your server to work properly. This makes your system safer. Open ports, with the socket state LISTEN, can be found with the program `netstat`. As for the options, it is recommended to use `netstat -ap` or `netstat -anp`. The `-p` option allows you to see which process is occupying a port under which name.

Compare the `netstat` results with those of a thorough port scan done from outside your host. An excellent program for this job is `nmap`, which not only checks out the ports of your machine, but also draws some conclusions as to which services are waiting behind them. However, port scanning may be interpreted as an aggressive act, so do not do this on a host without the explicit approval of the administrator.

Finally, remember that it is important not only to scan TCP ports, but also UDP ports (options `-sS` and `-sU`).

- To monitor the integrity of the files of your system in a reliable way, use the program AIDE (Advanced Intrusion Detection Environment), available on SUSE Linux. Encrypt the database created by AIDE to prevent someone from tampering with it. Furthermore, keep a backup of this database available outside your machine, stored on an external data medium not connected to it by a network link.
- Take proper care when installing any third-party software. There have been cases where a hacker had built a trojan horse into the tar archive of a security software package, which was fortunately discovered very quickly. If you install a binary package, have no doubts about the site from which you downloaded it.

SUSE's RPM packages are gpg-signed. The key used by SUSE for signing is:

ID:9C800ACA 2000-10-19 SUSE Package Signing Key <build@suse.de>

Key fingerprint = 79C1 79B2 E1C8 20C1 890F 9994 A84E DAE8 9C80 0ACA

The command `rpm --checksig package.rpm` shows whether the checksum and the signature of an uninstalled package are correct. Find the key on the first CD of the distribution and on most key servers worldwide.

- Check your backups of user and system files regularly. Consider that if you do not test whether the backup works, it might actually be worthless.
- Check your log files. Whenever possible, write a small script to search for suspicious entries. Admittedly, this is not exactly a trivial task. In the end, only you can know which entries are unusual and which are not.
- Use `tcp_wrapper` to restrict access to the individual services running on your machine, so you have explicit control over which IP addresses can connect to a service. For further information regarding `tcp_wrapper`, consult the manual pages of `tcpd` and `hosts_access` (`man 8 tcpd`, `man hosts_access`).
- Use SuSEfirewall to enhance the security provided by `tcpd` (`tcp_wrapper`).
- Design your security measures to be redundant: a message seen twice is much better than no message at all.

4.5.3 Using the Central Security Reporting Address

If you discover a security-related problem (please check the available update packages first), write an e-mail to security@suse.de. Please include a detailed description of the problem and the version number of the package concerned. SUSE will try to send a reply as soon as possible. You are encouraged to pgp encrypt your e-mail messages. SUSE's pgp key is:

```
ID:3D25D3D9 1999-03-06 SUSE Security Team <security@suse.de>
Key fingerprint = 73 5F 2E 99 DF DB 94 C4 8F 5A A3 AE AF 22 F2 D5
```

This key is also available for download from <http://www.novell.com/linux/security/securitysupport.html>.

5

Access Control Lists in Linux

POSIX ACLs (access control lists) can be used as an expansion of the traditional permission concept for file system objects. With ACLs, permissions can be defined more flexibly than the traditional permission concept allows.

The term *POSIX ACL* suggests that this is a true POSIX (*portable operating system interface*) standard. The respective draft standards POSIX 1003.1e and POSIX 1003.2c have been withdrawn for several reasons. Nevertheless, ACLs as found on many systems belonging to the UNIX family are based on these drafts and the implementation of file system ACLs as described in this chapter follows these two standards as well. They can be viewed at <http://wt.xpilot.org/publications posix.1e/>.

5.1 Traditional File Permissions

The basics of traditional Linux file permissions are explained in Section “Users and Access Permissions” (Chapter 3, *Working with the Shell*, ↑Start-Up). More advanced features are the setuid, setgid, and sticky bit.

5.1.1 The setuid Bit

In certain situations, the access permissions may be too restrictive. Therefore, Linux has additional settings that enable the temporary change of the current user and group identity for a specific action. For example, the `passwd` program normally requires root permissions to access `/etc/passwd`. This file contains some important information, like the home directories of users and user and group IDs. Thus, a normal user

would not be able to change `passwd`, because it would be too dangerous to grant all users direct access to this file. A possible solution to this problem is the *setuid* mechanism. *setuid* (set user ID) is a special file attribute that instructs the system to execute programs marked accordingly under a specific user ID. Consider the `passwd` command:

```
-rwsr-xr-x 1 root shadow 80036 2004-10-02 11:08 /usr/bin/passwd
```

You can see the `s` that denotes that the setuid bit is set for the user permission. By means of the setuid bit, all users starting the `passwd` command execute it as `root`.

5.1.2 The `setgid` Bit

The setuid bit applies to users. However, there is also an equivalent property for groups: the *setgid* bit. A program for which this bit was set runs under the group ID under which it was saved, no matter which user starts it. Therefore, in a directory with the setgid bit, all newly created files and subdirectories are assigned to the group to which the directory belongs. Consider the following example directory:

```
drwxrws--- 2 tux archive 48 Nov 19 17:12 backup
```

You can see the `s` that denotes that the setgid bit is set for the group permission. The owner of the directory and members of the group `archive` may access this directory. Users that are not members of this group are “mapped” to the respective group. The effective group ID of all written files will be `archive`. For example, a backup program that runs with the group ID `archive` is able to access this directory even without root privileges.

5.1.3 The Sticky Bit

There is also the *sticky bit*. It makes a difference whether it belongs to an executable program or a directory. If it belongs to a program, a file marked in this way is loaded to RAM to avoid needing to get it from the hard disk each time it is used. This attribute is used rarely, because modern hard disks are fast enough. If this bit is assigned to a directory, it prevents users from deleting each other's files. Typical examples include the `/tmp` and `/var/tmp` directories:

```
drwxrwxrwt 2 root root 1160 2002-11-19 17:15 /tmp
```

5.2 Advantages of ACLs

Traditionally, three permission sets are defined for each file object on a Linux system. These sets include the read (r), write (w), and execute (x) permissions for each of three types of users—the file owner, the group, and other users. In addition to that, it is possible to set the *set user id*, the *set group id*, and the *sticky bit*. This lean concept is fully adequate for most practical cases. However, for more complex scenarios or advanced applications, system administrators formerly had to use a number of tricks to circumvent the limitations of the traditional permission concept.

ACLs can be used as an extension of the traditional file permission concept. They allow assignment of permissions to individual users or groups even if these do not correspond to the original owner or the owning group. Access control lists are a feature of the Linux kernel and are currently supported by ReiserFS, Ext2, Ext3, JFS, and XFS. Using ACLs, complex scenarios can be realized without implementing complex permission models on the application level.

The advantages of ACLs are evident if you want to replace a Windows server with a Linux server. Some of the connected workstations may continue to run under Windows even after the migration. The Linux system offers file and print services to the Windows clients with Samba. With Samba supporting access control lists, user permissions can be configured both on the Linux server and in Windows with a graphical user interface (only Windows NT and later). With `winbindd`, part of the `samba` suite, it is even possible to assign permissions to users only existing in the Windows domain without any account on the Linux server.

5.3 Definitions

user class

The conventional POSIX permission concept uses three *classes* of users for assigning permissions in the file system: the owner, the owning group, and other users.

Three permission bits can be set for each user class, giving permission to read (r), write (w), and execute (x).

access ACL

The user and group access permissions for all kinds of file system objects (files and directories) are determined by means of access ACLs.

default ACL

Default ACLs can only be applied to directories. They determine the permissions a file system object inherits from its parent directory when it is created.

ACL entry

Each ACL consists of a set of ACL entries. An ACL entry contains a type, a qualifier for the user or group to which the entry refers, and a set of permissions. For some entry types, the qualifier for the group or users is undefined.

5.4 Handling ACLs

[Table 5.1, “ACL Entry Types”](#) (page 145) summarizes the six possible types of ACL entries, each defining permissions for a user or a group of users. The *owner* entry defines the permissions of the user owning the file or directory. The *owning group* entry defines the permissions of the file's owning group. The superuser can change the owner or owning group with `chown` or `chgrp`, in which case the owner and owning group entries refer to the new owner and owning group. Each *named user* entry defines the permissions of the user specified in the entry's qualifier field. Each *named group* entry defines the permissions of the group specified in the entry's qualifier field. Only the named user and named group entries have a qualifier field that is not empty. The *other* entry defines the permissions of all other users.

The *mask* entry further limits the permissions granted by named user, named group, and owning group entries by defining which of the permissions in those entries are effective and which are masked. If permissions exist in one of the mentioned entries as well as in the mask, they are effective. Permissions contained only in the mask or only in the actual entry are not effective—meaning the permissions are not granted. All permissions defined in the owner and owning group entries are always effective. The example in [Table 5.2, “Masking Access Permissions”](#) (page 145) demonstrates this mechanism.

There are two basic classes of ACLs: A *minimum* ACL contains only the entries for the types owner, owning group, and other, which correspond to the conventional permission bits for files and directories. An *extended* ACL goes beyond this. It must contain a mask entry and may contain several entries of the named user and named group types.

Table 5.1 ACL Entry Types

Type	Text Form
owner	user::rwx
named user	user:name:rwx
owning group	group::rwx
named group	group:name:rwx
mask	mask::rwx
other	other::rwx

Table 5.2 Masking Access Permissions

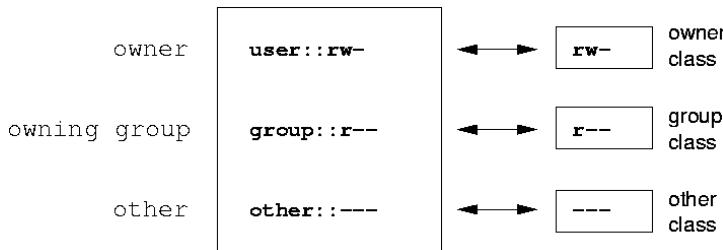
Entry Type	Text Form	Permissions
named user	user:geek0:r-x	r-x
mask	mask::rw-	rw-
	effective permissions:	r--

5.4.1 ACL Entries and File Mode Permission Bits

Figure 5.1, “Minimum ACL: ACL Entries Compared to Permission Bits” (page 146) and Figure 5.2, “Extended ACL: ACL Entries Compared to Permission Bits” (page 146) illustrate the two cases of a minimum ACL and an extended ACL. The figures are structured in three blocks—the left block shows the type specifications of the ACL entries, the center block displays an example ACL, and the right block shows the respective permission bits according to the conventional permission concept, for example, as displayed by `ls -l`. In both cases, the *owner class* permissions are mapped to the

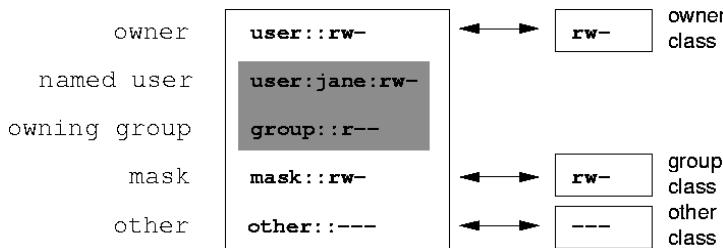
ACL entry owner. *Other class* permissions are mapped to the respective ACL entry. However, the mapping of the *group class* permissions is different in the two cases.

Figure 5.1 Minimum ACL: ACL Entries Compared to Permission Bits



In the case of a minimum ACL—without mask—the group class permissions are mapped to the ACL entry owning group. This is shown in [Figure 5.1, “Minimum ACL: ACL Entries Compared to Permission Bits”](#) (page 146). In the case of an extended ACL—with mask—the group class permissions are mapped to the mask entry. This is shown in [Figure 5.2, “Extended ACL: ACL Entries Compared to Permission Bits”](#) (page 146).

Figure 5.2 Extended ACL: ACL Entries Compared to Permission Bits



This mapping approach ensures the smooth interaction of applications, regardless of whether they have ACL support. The access permissions that were assigned by means of the permission bits represent the upper limit for all other “fine adjustments” made with an ACL. Changes made to the permission bits are reflected by the ACL and vice versa.

5.4.2 A Directory with an Access ACL

With `getfacl` and `setfacl` on the command line, you can access ACLs. The usage of these commands is demonstrated in the following example.

Before creating the directory, use the `umask` command to define which access permissions should be masked each time a file object is created. The command `umask 027` sets the default permissions by giving the owner the full range of permissions (0), denying the group write access (2), and giving other users no permissions at all (7). `umask` actually masks the corresponding permission bits or turns them off. For details, consult the `umask` man page.

`mkdir mydir` creates the `mydir` directory with the default permissions as set by `umask`. Use `ls -dl mydir` to check whether all permissions were assigned correctly. The output for this example is:

```
drwxr-x--- ... tux project3 ... mydir
```

With `getfacl mydir`, check the initial state of the ACL. This gives information like:

```
# file: mydir
# owner: tux
# group: project3
user::rwx
group::r-x
other::---
```

The first three output lines display the name, owner, and owning group of the directory. The next three lines contain the three ACL entries owner, owning group, and other. In fact, in the case of this minimum ACL, the `getfacl` command does not produce any information you could not have obtained with `ls`.

Modify the ACL to assign read, write, and execute permissions to an additional user `geeko` and an additional group `mascots` with:

```
setfacl -m user:geeko:rwx,group:mascots:rwx mydir
```

The option `-m` prompts `setfacl` to modify the existing ACL. The following argument indicates the ACL entries to modify (multiple entries are separated by commas). The final part specifies the name of the directory to which these modifications should be applied. Use the `getfacl` command to take a look at the resulting ACL.

```
# file: mydir
# owner: tux
# group: project3
user::rwx
user:geeko:rwx
group::r-x
group:mascots:rwx
```

```
mask::rwx
other::---
```

In addition to the entries initiated for the user `geeko` and the group `mascots`, a mask entry has been generated. This mask entry is set automatically so that all permissions are effective. `setfacl` automatically adapts existing mask entries to the settings modified, unless you deactivate this feature with `-n`. `mask` defines the maximum effective access permissions for all entries in the group class. This includes named user, named group, and owning group. The group class permission bits displayed by `ls -dl mydir` now correspond to the mask entry.

```
drwxrwx---+ ... tux project3 ... mydir
```

The first column of the output contains an additional `+` to indicate that there is an *extended* ACL for this item.

According to the output of the `ls` command, the permissions for the mask entry include write access. Traditionally, such permission bits would mean that the owning group (here `project3`) also has write access to the directory `mydir`. However, the effective access permissions for the owning group correspond to the overlapping portion of the permissions defined for the owning group and for the mask—which is `r-x` in our example (see [Table 5.2, “Masking Access Permissions”](#) (page 145)). As far as the effective permissions of the owning group in this example are concerned, nothing has changed even after the addition of the ACL entries.

Edit the mask entry with `setfacl` or `chmod`. For example, use `chmod g-w mydir`. `ls -dl mydir` then shows:

```
drwxr-x---+ ... tux project3 ... mydir
```

`getfacl mydir` provides the following output:

```
# file: mydir
# owner: tux
# group: project3
user::rwx
user:geeko:rwx          # effective: r-x
group::r-x
group:mascots:rwx       # effective: r-x
mask::r-x
other::---
```

After executing the `chmod` command to remove the write permission from the group class bits, the output of the `ls` command is sufficient to see that the mask bits must have changed accordingly: write permission is again limited to the owner of `mydir`.

The output of the `getfacl` confirms this. This output includes a comment for all those entries in which the effective permission bits do not correspond to the original permissions, because they are filtered according to the mask entry. The original permissions can be restored at any time with `chmod g+w mydir`.

5.4.3 A Directory with a Default ACL

Directories can have a default ACL, which is a special kind of ACL defining the access permissions that objects in the directory inherit when they are created. A default ACL affects both subdirectories and files.

Effects of a Default ACL

There are two ways in which the permissions of a directory's default ACL are passed to the files and subdirectories:

- A subdirectory inherits the default ACL of the parent directory both as its default ACL and as an access ACL.
- A file inherits the default ACL as its access ACL.

All system calls that create file system objects use a `mode` parameter that defines the access permissions for the newly created file system object. If the parent directory does not have a default ACL, the permission bits as defined by the `umask` are subtracted from the permissions as passed by the `mode` parameter, with the result being assigned to the new object. If a default ACL exists for the parent directory, the permission bits assigned to the new object correspond to the overlapping portion of the permissions of the `mode` parameter and those that are defined in the default ACL. The `umask` is disregarded in this case.

Application of Default ACLs

The following three examples show the main operations for directories and default ACLs:

1. Add a default ACL to the existing directory `mydir` with:

```
setfacl -d -m group:mascots:r-x mydir
```

The option `-d` of the `setfacl` command prompts `setfacl` to perform the following modifications (option `-m`) in the default ACL.

Take a closer look at the result of this command:

```
getfacl mydir

# file: mydir
# owner: tux
# group: project3
user::rwx
user:geeko:rwx
group::r-x
group:mascots:rwx
mask::rwx
other::---
default:user::rwx
default:group::r-x
default:group:mascots:r-x
default:mask::r-x
default:other::---
```

`getfacl` returns both the access ACL and the default ACL. The default ACL is formed by all lines that start with `default`. Although you merely executed the `setfacl` command with an entry for the `mascots` group for the default ACL, `setfacl` automatically copied all other entries from the access ACL to create a valid default ACL. Default ACLs do not have an immediate effect on access permissions. They only come into play when file system objects are created. These new objects inherit permissions only from the default ACL of their parent directory.

2. In the next example, use `mkdir` to create a subdirectory in `mydir`, which inherits the default ACL.

```
mkdir mydir/mysubdir

getfacl mydir/mysubdir

# file: mydir/mysubdir
# owner: tux
# group: project3
user::rwx
group::r-x
group:mascots:r-x
mask::r-x
other::---
default:user::rwx
default:group::r-x
default:group:mascots:r-x
```

```
default:mask::r-x
default:other::---
```

As expected, the newly-created subdirectory `mysubdir` has the permissions from the default ACL of the parent directory. The access ACL of `mysubdir` is an exact reflection of the default ACL of `mydir`. The default ACL that this directory will hand down to its subordinate objects is also the same.

3. Use `touch` to create a file in the `mydir` directory, for example, `touch mydir/myfile`. `ls -l mydir/myfile` then shows:

```
-rw-r-----+ ... tux project3 ... mydir/myfile
```

The output of `getfacl mydir/myfile` is:

```
# file: mydir/myfile
# owner: tux
# group: project3
user::rw-
group::r-x      # effective:r--
group:mascots:r-x # effective:r--
mask::r--
other::---
```

`touch` uses a mode with the value `0666` when creating new files, which means that the files are created with read and write permissions for all user classes, provided no other restrictions exist in `umask` or in the default ACL (see [Section “Effects of a Default ACL” \(page 149\)](#)). In effect, this means that all access permissions not contained in the `mode` value are removed from the respective ACL entries. Although no permissions were removed from the ACL entry of the group class, the mask entry was modified to mask permissions not set in `mode`.

This approach ensures the smooth interaction of applications, such as compilers, with ACLs. You can create files with restricted access permissions and subsequently mark them as executable. The `mask` mechanism guarantees that the right users and groups can execute them as desired.

5.4.4 The ACL Check Algorithm

A check algorithm is applied before any process or application is granted access to an ACL-protected file system object. As a basic rule, the ACL entries are examined in the following sequence: owner, named user, owning group or named group, and other. The

access is handled in accordance with the entry that best suits the process. Permissions do not accumulate.

Things are more complicated if a process belongs to more than one group and would potentially suit several group entries. An entry is randomly selected from the suitable entries with the required permissions. It is irrelevant which of the entries triggers the final result “access granted”. Likewise, if none of the suitable group entries contains the required permissions, a randomly selected entry triggers the final result “access denied”.

5.5 ACL Support in Applications

ACLs can be used to implement very complex permission scenarios that meet the requirements of modern applications. The traditional permission concept and ACLs can be combined in a smart manner. The basic file commands (`cp`, `mv`, `ls`, etc.) support ACLs, as do Samba and Konqueror.

Unfortunately, many editors and file managers still lack ACL support. When copying files with Emacs, for instance, the ACLs of these files are lost. When modifying files with an editor, the ACLs of files are sometimes preserved and sometimes not, depending on the backup mode of the editor used. If the editor writes the changes to the original file, the access ACL is preserved. If the editor saves the updated contents to a new file that is subsequently renamed to the old filename, the ACLs may be lost, unless the editor supports ACLs. Except for the star archiver, there are currently no backup applications that preserve ACLs.

5.6 For More Information

Detailed information about ACLs is available at <http://acl.bestbits.at/>. Also see the man pages for `getfacl(1)`, `acl(5)`, and `setfacl(1)`.

6

System Monitoring Utilities

A number of programs and mechanisms, some of which are presented here, can be used to examine the status of your system. Also described are some utilities that are useful for routine work, along with their most important parameters.

For each of the commands introduced, examples of the relevant outputs are presented. In these examples, the first line is the command itself (after the > or # sign prompt). Omissions are indicated with square brackets ([. . .]) and long lines are wrapped where necessary. Line breaks for long lines are indicated by a backslash (\).

```
# command -x -y
output line 1
output line 2
output line 3 is annoyingly long, so long that \
    we have to break it
output line 3
[...]
output line 98
output line 99
```

The descriptions have been kept short to allow as many utilities as possible to be mentioned. Further information for all the commands can be found in the man pages. Most of the commands also understand the parameter --help, which produces a brief list of the possible parameters.

6.1 List of Open Files: lsof

To view a list of all the files open for the process with process ID *PID*, use -p. For example, to view all the files used by the current shell, enter:

```
tester@linux:~> lsof -p $$
```

COMMAND	PID	USER	FD	TYPE	DEVICE	SIZE	NODE	NAME
bash	5552	tester	cwd	DIR	3,3	1512	117619	/home/tester
bash	5552	tester	rtd	DIR	3,3	584	2	/
bash	5552	tester	txt	REG	3,3	498816	13047	/bin/bash
bash	5552	tester	mem	REG	0,0		0	[heap] (stat: No such \ file or directory)
bash	5552	tester	mem	REG	3,3	217016	115687	/var/run/nscd/passwd
bash	5552	tester	mem	REG	3,3	208464	11867	\ /usr/lib/locale/en_GB.utf8/LC_CTYPE
bash	5552	tester	mem	REG	3,3	882134	11868	\ /usr/lib/locale/en_GB.utf8/LC_COLLATE
bash	5552	tester	mem	REG	3,3	1386997	8837	/lib/libc-2.3.6.so
bash	5552	tester	mem	REG	3,3	13836	8843	/lib/libdl-2.3.6.so
bash	5552	tester	mem	REG	3,3	290856	12204	/lib/libcurses.so.5.5
bash	5552	tester	mem	REG	3,3	26936	13004	/lib/libhistory.so.5.1
bash	5552	tester	mem	REG	3,3	190200	13006	/lib/libreadline.so.5.1
bash	5552	tester	mem	REG	3,3	54	11842	\ /usr/lib/locale/en_GB.utf8/LC_NUMERIC
bash	5552	tester	mem	REG	3,3	2375	11663	\ /usr/lib/locale/en_GB.utf8/LC_TIME
bash	5552	tester	mem	REG	3,3	290	11736	\ /usr/lib/locale/en_GB.utf8/LC_MONETARY
bash	5552	tester	mem	REG	3,3	52	11831	\ /usr/lib/locale/en_GB.utf8/LC_MESSAGES/SYS_LC_MESSAGES
bash	5552	tester	mem	REG	3,3	34	11862	\ /usr/lib/locale/en_GB.utf8/LC_PAPER
bash	5552	tester	mem	REG	3,3	62	11839	\ /usr/lib/locale/en_GB.utf8/LC_NAME
bash	5552	tester	mem	REG	3,3	127	11664	\ /usr/lib/locale/en_GB.utf8/LC_ADDRESS
bash	5552	tester	mem	REG	3,3	56	11735	\ /usr/lib/locale/en_GB.utf8/LC_TELEPHONE
bash	5552	tester	mem	REG	3,3	23	11866	\ /usr/lib/locale/en_GB.utf8/LC_MEASUREMENT
bash	5552	tester	mem	REG	3,3	21544	9109	\ /usr/lib/gconv/gconv-modules.cache
bash	5552	tester	mem	REG	3,3	366	9720	\ /usr/lib/locale/en_GB.utf8/LC_IDENTIFICATION
bash	5552	tester	mem	REG	3,3	97165	8828	/lib/ld-2.3.6.so
bash	5552	tester	Ou	CHR	136,5		2	/dev/pts/5
bash	5552	tester	lu	CHR	136,5		2	/dev/pts/5
bash	5552	tester	2u	CHR	136,5		2	/dev/pts/0
bash	3838	tester	255u	CHR	136,0		2	/dev/pts/0
bash	5552	tester	Ou	CHR	136,5		7	/dev/pts/5
bash	5552	tester	1u	CHR	136,5		7	/dev/pts/5
bash	5552	tester	2u	CHR	136,5		7	/dev/pts/5
bash	5552	tester	255u	CHR	136,5		7	/dev/pts/5
X	5646	root	mem	CHR	1,1		1006	/dev/mem
lsof	5673	tester	Ou	CHR	136,5		7	/dev/pts/5
lsof	5673	tester	2u	CHR	136,5		7	/dev/pts/5

The special shell variable `$$`, whose value is the process ID of the shell, has been used.

The command `lsof` lists all the files currently open when used without any parameters. Because there are often thousands of open files, listing all of them is rarely useful. However, the list of all files can be combined with search functions to generate useful lists. For example, list all used character devices:

```
tester@linux:~> lsof | grep CHR
```

bash	3838	tester	Ou	CHR	136,0		2	/dev/pts/0
bash	3838	tester	1u	CHR	136,0		2	/dev/pts/0
bash	3838	tester	2u	CHR	136,0		2	/dev/pts/0
bash	3838	tester	255u	CHR	136,0		2	/dev/pts/0
bash	5552	tester	Ou	CHR	136,5		7	/dev/pts/5
bash	5552	tester	1u	CHR	136,5		7	/dev/pts/5
bash	5552	tester	2u	CHR	136,5		7	/dev/pts/5
bash	5552	tester	255u	CHR	136,5		7	/dev/pts/5
X	5646	root	mem	CHR	1,1		1006	/dev/mem
lsof	5673	tester	Ou	CHR	136,5		7	/dev/pts/5
lsof	5673	tester	2u	CHR	136,5		7	/dev/pts/5

```
grep      5674    tester   1u      CHR  136,5          7 /dev/pts/5
grep      5674    tester   2u      CHR  136,5          7 /dev/pts/5
```

6.2 User Accessing Files: fuser

It can be useful to determine what processes or users are currently accessing certain files. Suppose, for example, you want to unmount a file system mounted at `/mnt`. `umount` returns "device is busy." The command `fuser` can then be used to determine what processes are accessing the device:

```
tester@linux:~> fuser -v /mnt/*
```

```
USER          PID ACCESS COMMAND
/mnt/notes.txt  26597 f.... less
```

Following termination of the `less` process, which was running on another terminal, the file system can successfully be unmounted.

6.3 File Properties: stat

The command `stat` displays file properties:

```
tester@linux:~> stat /etc/profile
  File: `/etc/profile'
  Size: 7930          Blocks: 16          IO Block: 4096   regular file
Device: 303h/771d      Inode: 40657        Links: 1
Access: (0644/-rw-r--r--)  Uid: (    0/    root)    Gid: (    0/    root)
Access: 2006-01-06 16:45:43.000000000 +0100
Modify: 2005-11-21 14:54:35.000000000 +0100
Change: 2005-12-19 09:51:04.000000000 +0100
```

The parameter `--filesystem` produces details of the properties of the file system in which the specified file is located:

```
tester@linux:~> stat /etc/profile --filesystem
  File: "/etc/profile"
  ID: 0          Namelen: 255       Type: reiserfs
Block size: 4096      Fundamental block size: 4096
Blocks: Total: 2622526     Free: 1809771     Available: 1809771
Inodes: Total: 0         Free: 0
```

6.4 USB Devices: `lsusb`

The command `lsusb` lists all USB devices. With the option `-v`, print a more detailed list. The detailed information is read from the directory `/proc/bus/usb/`. The following is the output of `lsusb` with the following USB devices attached: hub, memory stick, hard disk, and a mouse.

```
linux:/ # lsusb
Bus 004 Device 007: ID 0ea0:2168 Ours Technology, Inc. Transcend JetFlash \
2.0 / Astone USB Drive
Bus 004 Device 006: ID 04b4:6830 Cypress Semiconductor Corp. USB-2.0 IDE \
Adapter
Bus 004 Device 005: ID 05e3:0605 Genesys Logic, Inc.
Bus 004 Device 001: ID 0000:0000
Bus 003 Device 001: ID 0000:0000
Bus 002 Device 001: ID 0000:0000
Bus 001 Device 005: ID 046d:c012 Logitech, Inc. Optical Mouse
Bus 001 Device 001: ID 0000:0000
```

6.5 Information about a SCSI Device: `scsiinfo`

The command `scsiinfo` lists information about a SCSI device. With the option `-l`, list all SCSI devices known to the system (similar information is obtained via the command `lsscsi`). The following is the output of `scsiinfo -i /dev/sda`, which gives information about a hard disk. The option `-a` gives even more information.

```
linux:/ # scsiinfo -i /dev/sda
Inquiry command
-----
Relative Address          0
Wide bus 32               0
Wide bus 16               1
Synchronous neg.          1
Linked Commands           1
Command Queueing          1
SftRe                      0
Device Type                0
Peripheral Qualifier      0
Removable?                 0
Device Type Modifier       0
ISO Version                 0
ECMA Version                 0
ANSI Version                3
```

```

AENC                      0
TrmIOP                     0
Response Data Format      2
Vendor:                   FUJITSU
Product:                  MAS3367NP
Revision level:           0104A0K7P43002BE

```

The option `-d` puts out a defects list with two tables of bad blocks of a hard disk: first the one supplied by the vendor (manufacturer table) and second the list of bad blocks that appeared in operation (grown table). If the number of entries in the grown table increases, it might be a good idea to replace the hard disk.

6.6 Processes: `top`

The command `top`, which stands for "table of processes," displays a list of processes that is refreshed every two seconds. To terminate the program, press `[Q]`. The parameter `-n 1` terminates the program after a single display of the process list. The following is an example output of the command `top -n 1`:

```

tester@linux:~> top -n 1
top - 17:06:28 up 2:10,  5 users,  load average: 0.00, 0.00, 0.00
Tasks: 85 total,   1 running, 83 sleeping,   1 stopped,   0 zombie
Cpu(s):  5.5% us,  0.8% sy,  0.8% ni, 91.9% id,  1.0% wa,  0.0% hi,  0.0% si
Mem: 515584k total, 506468k used,   9116k free, 66324k buffers
Swap: 658656k total,     0k used, 658656k free, 353328k cached

```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1	root	16	0	700	272	236	S	0.0	0.1	0:01.33	init
2	root	34	19	0	0	0	S	0.0	0.0	0:00.00	ksoftirqd/0
3	root	10	-5	0	0	0	S	0.0	0.0	0:00.27	events/0
4	root	10	-5	0	0	0	S	0.0	0.0	0:00.01	khelper
5	root	10	-5	0	0	0	S	0.0	0.0	0:00.00	kthread
11	root	10	-5	0	0	0	S	0.0	0.0	0:00.05	kblockd/0
12	root	20	-5	0	0	0	S	0.0	0.0	0:00.00	kacpid
472	root	20	0	0	0	0	S	0.0	0.0	0:00.00	pdflush
473	root	15	0	0	0	0	S	0.0	0.0	0:00.06	pdflush
475	root	11	-5	0	0	0	S	0.0	0.0	0:00.00	aio/0
474	root	15	0	0	0	0	S	0.0	0.0	0:00.07	kswapd0
681	root	10	-5	0	0	0	S	0.0	0.0	0:00.01	kseriod
839	root	10	-5	0	0	0	S	0.0	0.0	0:00.02	reiserfs/0
923	root	13	-4	1712	552	344	S	0.0	0.1	0:00.67	udevd
1343	root	10	-5	0	0	0	S	0.0	0.0	0:00.00	khubd
1587	root	20	0	0	0	0	S	0.0	0.0	0:00.00	shpchpd_event
1746	root	15	0	0	0	0	S	0.0	0.0	0:00.00	wl_control
1752	root	15	0	0	0	0	S	0.0	0.0	0:00.00	wl_bus_master1
2151	root	16	0	1464	496	416	S	0.0	0.1	0:00.00	acpid
2165	messageb	16	0	3340	1048	792	S	0.0	0.2	0:00.64	dbus-daemon

```

2166 root      15   0 1840  752  556 S  0.0  0.1  0:00.01 syslog-ng
2171 root      16   0 1600  516  320 S  0.0  0.1  0:00.00 klogd
2235 root      15   0 1736  800  652 S  0.0  0.2  0:00.10 resmgrd
2289 root      16   0 4192 2852 1444 S  0.0  0.6  0:02.05 haldd
2403 root      23   0 1756  600  524 S  0.0  0.1  0:00.00 haldd-addon-acpi
2709 root      19   0 2668 1076  944 S  0.0  0.2  0:00.00 NetworkManagerD
2714 root      16   0 1756  648  564 S  0.0  0.1  0:00.56 haldd-addon-stor

```

If you press **F** while `top` is running, a menu opens with which to make extensive changes to the format of the output.

The parameter `-U UID` monitors only the processes associated with a particular user. Replace *UID* with the user ID of the user. `top -U `id -u`` returns the UID of the user on the basis of the username and displays his processes.

6.7 Process List: ps

The command `ps` produces a list of processes. Most parameters must be written without a minus sign. Refer to `ps --help` for a brief help or to the man page for extensive help.

To list all processes with user and command line information, use `ps axu`:

```

tester@linux:~> ps axu
USER        PID %CPU %MEM      VSZ      RSS TTY      STAT START      TIME COMMAND
root         1  0.0  0.0      696     272 ?          S    12:59   0:01 init [5]
root         2  0.0  0.0       0      0 ?          SN   12:59   0:00 [ksoftirqd/0]
root         3  0.0  0.0       0      0 ?          S<  12:59   0:00 [events/0]
[...]
tester      4047  0.0  6.0 158548 31400 ?          Ssl  13:02   0:06 mono-best \
--debug /usr/lib/beagle/Best.exe --autostarted
tester      4057  0.0  0.7  9036  3684 ?          S    13:02   0:00 \
/opt/gnome/sbin/gnome-vfs-daemon
--oaf-activate-iid=OAFIID:GNOME_VFS_Daemon_Factory --oa
tester      4067  0.0  0.1  2204   636 ?          S    13:02   0:00 \
/opt/gnome/lib/nutilus/mapping-daemon
tester      4072  0.0  1.0 15996  5160 ?          Ss   13:02   0:00 \
gnome-screensaver
tester      4114  0.0  3.7 130988 19172 ?          SLL  13:06   0:04 sound-juicer
tester      4818  0.0  0.3   4192  1812 pts/0      Ss   15:59   0:00 -bash
tester      4959  0.0  0.1  2324   816 pts/0      R+   16:17   0:00 ps axu

```

To check how many `sshd` processes are running, use the option `-p` together with the command `pidof`, which lists the process IDs of the given processes.

```
tester@linux:~> ps -p `pidof sshd`  
PID TTY      STAT   TIME COMMAND  
3524 ?        Ss      0:00 /usr/sbin/sshd -o PidFile=/var/run/sshd.init.pid  
4813 ?        Ss      0:00 sshd: tester [priv]  
4817 ?        R       0:00 sshd: tester@pts/0
```

The process list can be formatted according to your needs. The option **-L** returns a list of all keywords. Enter the following command to issue a list of all processes sorted by memory usage:

```
tester@linux:~> ps ax --format pid,rss,cmd --sort rss  
PID  RSS CMD  
 2    0  [ksoftirqd/0]  
 3    0  [events/0]  
 4    0  [khelper]  
 5    0  [kthread]  
11    0  [kblockd/0]  
12    0  [kacpid]  
472   0  [pdfflush]  
473   0  [pdfflush]  
[...]  
4028 17556 nautilus --no-default-window --sm-client-id default2  
4118 17800 ksnapshot  
4114 19172 sound-juicer  
4023 25144 gnome-panel --sm-client-id default1  
4047 31400 mono-best --debug /usr/lib/beagle/Best.exe --autostarted  
3973 31520 mono-beagled --debug /usr/lib/beagle/BeagleDaemon.exe \  
--bg --autostarted
```

6.8 Process Tree: `pstree`

The command `pstree` produces a list of processes in the form of a tree:

```
tester@linux:~> pstree  
init-+-NetworkManagerD  
|  |-acpid  
|  |-3*[automount]  
|  |-cron  
|  |-cupsd  
|  |-2*[dbus-daemon]  
|  |-dbus-launch  
|  |-dcopserver  
|  |-dhcpcd  
|  |-events/0  
|  |-gpg-agent  
|  |-hald-+-hald-addon-acpi  
|  |  `--hald-addon-stor  
|  |-kded  
|  |-kdeinit-+-kdesu---su---kdesu_stub---yast2---y2controlcenter
```

```

|      | -kio_file
|      | -klauncher
|      | -konqueror
|      | -konsole---bash---su---bash
|      |           `--bash
|      `--kwin
|-kdesktop---kdesktop_lock---xmatrix
|-kdesud
|-kdm-+-X
|   `--kdm---startkde---kwrapper
[...]

```

The parameter `-p` adds the process ID to a given name. To have the command lines displayed as well, use the `-a` parameter:

6.9 Who Is Doing What: w

With the command `w`, find out who is logged onto the system and what each user is doing. For example:

```

tester@linux:~> w
16:33:03 up 3:33, 2 users, load average: 0.14, 0.06, 0.02
USER     TTY      LOGIN@     IDLE     JCPU    PCPU WHAT
tester    :0      16:33 ?xdm?    9.42s  0.15s /bin/sh /opt/kde3/bin/startkde
tester    pts/0    15:59    0.00s  0.19s  0.00s w

```

If any users of other systems have logged in remotely, the parameter `-f` shows the computers from which they have established the connection.

6.10 Memory Usage: free

The utility `free` examines RAM usage. Details of both free and used memory (and swap areas) are shown:

```

ester@linux:~> free
              total        used         free       shared      buffers      cached
Mem:      515584      501704      13880          0      73040     334592
-/+ buffers/cache:      94072      421512
Swap:      658656          0      658656

```

The options `-b`, `-k`, `-m`, `-g` show output in bytes, KB, MB, or GB, respectively. The parameter `-d delay` ensures that the display is refreshed every `delay` seconds. For example, `free -d 1.5` produces an update every 1.5 seconds.

6.11 Kernel Ring Buffer: dmesg

The Linux kernel keeps certain messages in a ring buffer. To view these messages, enter the command `dmesg`:

```
$ dmesg
[...]
end_request: I/O error, dev fd0, sector 0
subfs: unsuccessful attempt to mount media (256)
e100: eth0: e100_watchdog: link up, 100Mbps, half-duplex
NET: Registered protocol family 17
IA-32 Microcode Update Driver: v1.14 <tigran@veritas.com>
microcode: CPU0 updated from revision 0xe to 0x2e, date = 08112004
IA-32 Microcode Update Driver v1.14 unregistered
bootsplash: status on console 0 changed to on
NET: Registered protocol family 10
Disabled Privacy Extensions on device c0326ea0(lo)
IPv6 over IPv4 tunneling driver
powernow: This module only works with AMD K7 CPUs
bootsplash: status on console 0 changed to on
```

Older events are logged in the files `/var/log/messages` and `/var/log/warn`.

6.12 File Systems and Their Usage: mount, df, and du

The command `mount` shows which file system (device and type) is mounted at which mount point:

```
tester@linux:~> mount
/dev/hda3 on / type reiserfs (rw,acl,user_xattr)
proc on /proc type proc (rw)
sysfs on /sys type sysfs (rw)
udev on /dev type tmpfs (rw)
devpts on /dev/pts type devpts (rw,mode=0620,gid=5)
/dev/hd1 on /boot type ext2 (rw,acl,user_xattr)
/dev/hda4 on /local type reiserfs (rw,acl,user_xattr)
/dev/fd0 on /media/floppy type subfs
(rw,nosuid,nodev,noatime,fs=floppyfss,procuid)
```

Obtain information about total usage of the file systems with the command `df`. The parameter `-h` (or `--human-readable`) transforms the output into a form understandable for common users.

```
tester@linux:~> df -h
Filesystem           Size   Used  Avail Use% Mounted on
/dev/hda3            11G   3.2G  6.9G  32% /
udev                 252M  104K  252M   1% /dev
/dev/hda1            16M   6.6M  7.8M  46% /boot
/dev/hda4            27G   34M   27G   1% /local
```

Display the total size of all the files in a given directory and its subdirectories with the command `du`. The parameter `-s` suppresses the output of detailed information. `-h` again transforms the data into a human-readable form:

```
tester@linux:~> du -sh /local
1.7M    /local
```

6.13 The /proc File System

The `/proc` file system is a pseudo file system in which the kernel reserves important information in the form of virtual files. For example, display the CPU type with this command:

```
ester@linux:~> cat /proc/cpuinfo
processor      : 0
vendor_id     : AuthenticAMD
cpu family    : 6
model         : 8
model name    : AMD Athlon(tm) XP 2400+
stepping       : 1
cpu MHz       : 2009.343
cache size    : 256 KB
fdiv_bug      : no
[...]
```

The allocation and use of interrupts can be queried with the following command:

```
tester@linux:~> cat /proc/interrupts
          CPU0
0:    3577519      XT-PIC  timer
1:      130      XT-PIC  i8042
2:        0      XT-PIC  cascade
5:    564535      XT-PIC  Intel 82801DB-ICH4
7:        1      XT-PIC  parport0
8:        2      XT-PIC  rtc
9:        1      XT-PIC  acpi, uhci_hcd:usb1, ehci_hcd:usb4
10:       0      XT-PIC  uhci_hcd:usb3
11:    71772      XT-PIC  uhci_hcd:usb2, eth0
12:   101150      XT-PIC  i8042
14:    33146      XT-PIC  ide0
15:   149202      XT-PIC  ide1
```

```
NMI:          0  
LOC:          0  
ERR:          0  
MIS:          0
```

Some of the important files and their contents are:

```
/proc/devices  
available devices
```

```
/proc/modules  
kernel modules loaded
```

```
/proc/cmdline  
kernel command line
```

```
/proc/meminfo  
detailed information about memory usage
```

```
/proc/config.gz  
gzip-compressed configuration file of the kernel currently running
```

Further information is available in the text file /usr/src/linux/
Documentation/filesystems/proc.txt. Information about processes cur-
rently running can be found in the /proc/NNN directories, where NNN is the process
ID (PID) of the relevant process. Every process can find its own characteristics in
/proc/self/:

```
tester@linux:~> ls -l /proc/self  
lrwxrwxrwx 1 root root 64 2006-01-09 13:03 /proc/self -> 5356  
tester@linux:~> ls -l /proc/self/  
total 0  
dr-xr-xr-x 2 tester users 0 2006-01-09 17:04 attr  
-r----- 1 tester users 0 2006-01-09 17:04 auxv  
-r--r--r-- 1 tester users 0 2006-01-09 17:04 cmdline  
lrwxrwxrwx 1 tester users 0 2006-01-09 17:04 cwd -> /home/tester  
-r----- 1 tester users 0 2006-01-09 17:04 environ  
lrwxrwxrwx 1 tester users 0 2006-01-09 17:04 exe -> /bin/ls  
dr-x----- 2 tester users 0 2006-01-09 17:04 fd  
-rw-r--r-- 1 tester users 0 2006-01-09 17:04 loginuid  
-r--r--r-- 1 tester users 0 2006-01-09 17:04 maps  
-rw----- 1 tester users 0 2006-01-09 17:04 mem  
-r--r--r-- 1 tester users 0 2006-01-09 17:04 mounts  
-rw-r--r-- 1 tester users 0 2006-01-09 17:04 oom_adj  
-r--r--r-- 1 tester users 0 2006-01-09 17:04 oom_score  
lrwxrwxrwx 1 tester users 0 2006-01-09 17:04 root -> /  
-rw----- 1 tester users 0 2006-01-09 17:04 seccomp
```

```
-r--r--r-- 1 tester users 0 2006-01-09 17:04 smaps
-r--r--r-- 1 tester users 0 2006-01-09 17:04 stat
-r--r--r-- 1 tester users 0 2006-01-09 17:04 statm
-r--r--r-- 1 tester users 0 2006-01-09 17:04 status
dr-xr-xr-x 3 tester users 0 2006-01-09 17:04 task
-r--r--r-- 1 tester users 0 2006-01-09 17:04 wchan
```

The address assignment of executables and libraries is contained in the maps file:

```
tester@linux:~> cat /proc/self/maps
08048000-0804c000 r-xp 00000000 03:03 17753      /bin/cat
0804c000-0804d000 rw-p 00004000 03:03 17753      /bin/cat
0804d000-0806e000 rw-p 0804d000 00:00 0          [heap]
b7d27000-b7d5a000 r--p 00000000 03:03 11867      \
    /usr/lib/locale/en_GB.utf8/LC_CTYPE
b7d5a000-b7e32000 r--p 00000000 03:03 11868      \
    /usr/lib/locale/en_GB.utf8/LC_COLLATE
b7e32000-b7e33000 rw-p b7e32000 00:00 0
b7e33000-b7f45000 r-xp 00000000 03:03 8837      /lib/libc-2.3.6.so
b7f45000-b7f46000 r--p 00112000 03:03 8837      /lib/libc-2.3.6.so
b7f46000-b7f48000 rw-p 00113000 03:03 8837      /lib/libc-2.3.6.so
b7f48000-b7f4c000 rw-p b7f48000 00:00 0
b7f52000-b7f53000 r--p 00000000 03:03 11842      \
    /usr/lib/locale/en_GB.utf8/LC_NUMERIC
[...]
b7f5b000-b7f61000 r-s 00000000 03:03 9109      \
    /usr/lib/gconv/gconv-modules.cache
b7f61000-b7f62000 r--p 00000000 03:03 9720      \
    /usr/lib/locale/en_GB.utf8/LC_IDENTIFICATION
b7f62000-b7f76000 r-xp 00000000 03:03 8828      /lib/ld-2.3.6.so
b7f76000-b7f78000 rw-p 00013000 03:03 8828      /lib/ld-2.3.6.so
bfd61000-bfd76000 rw-p bfd61000 00:00 0        [stack]
fffffe000-fffff000 ---p 00000000 00:00 0        [vdso]
```

6.13.1 procinfo

Important information from the /proc file system is summarized by the command procinfo:

```
tester@linux:~> procinfo
Linux 2.6.15-rc5-git3-2-default (geeko@buildhost) (gcc 4.1.0 20051129) #1 Wed
Dec 14 13:10:38 UTC 2005 1CPU [linux.suse.de]

Memory:      Total        Used        Free        Shared       Buffers
Mem:         515584      509472      6112          0        73024
Swap:        658656          0        658656

Bootup: Mon Jan  9 12:59:08 2006     Load average: 0.10 0.04 0.05 1/86 5406

user :      0:02:07.98   0.8% page in :      442638 disk 1:      20125r
```

```

13476w
nice : 0:02:20.91 0.9% page out: 134950
system: 0:00:42.93 0.3% page act: 70577
IOWait: 0:01:25.40 0.6% page dea: 11696
hw irq: 0:00:08.94 0.1% page flt: 1423622
sw irq: 0:00:01.29 0.0% swap in : 0
idle : 4:06:30.54 97.3% swap out: 0
uptime: 4:13:20.72 context : 3813145

irq 0: 3799268 timer irq 8: 2 rtc
irq 1: 130 i8042 irq 9: 1 acpi, uhci_hcd:usb1,
irq 2: 0 cascade [4] irq 10: 0 uhci_hcd:usb3
irq 3: 8 irq 11: 75905 uhci_hcd:usb2, eth0
irq 4: 8 irq 12: 101150 i8042
irq 5: 564535 Intel 82801DB-ICH4 irq 14: 33733 ide0
irq 6: 9 irq 15: 157045 ide1
irq 7: 1 parport0 [3]

```

To see all the information, use the parameter `-a`. The parameter `-nN` produces updates of the information every N seconds. In this case, terminate the program by pressing `Q`.

By default, the cumulative values are displayed. The parameter `-d` produces the differential values. `procinfo -dn5` displays the values that have changed in the last five seconds:

6.14 PCI Resources: `lspci`

The command `lspci` lists the PCI resources:

```

linux:~ # lspci
00:00.0 Host bridge: Intel Corporation 82845G/GL[Brookdale-G]/GE/PE \
    DRAM Controller/Host-Hub Interface (rev 01)
00:01.0 PCI bridge: Intel Corporation 82845G/GL[Brookdale-G]/GE/PE \
    Host-to-AGP Bridge (rev 01)
00:1d.0 USB Controller: Intel Corporation 82801DB/DBL/DBM \
    (ICH4/ICH4-L/ICH4-M) USB UHCI Controller #1 (rev 01)
00:1d.1 USB Controller: Intel Corporation 82801DB/DBL/DBM \
    (ICH4/ICH4-L/ICH4-M) USB UHCI Controller #2 (rev 01)
00:1d.2 USB Controller: Intel Corporation 82801DB/DBL/DBM \
    (ICH4/ICH4-L/ICH4-M) USB UHCI Controller #3 (rev 01)
00:1d.7 USB Controller: Intel Corporation 82801DB/DBM \
    (ICH4/ICH4-M) USB2 EHCI Controller (rev 01)
00:1e.0 PCI bridge: Intel Corporation 82801 PCI Bridge (rev 81)
00:1f.0 ISA bridge: Intel Corporation 82801DB/DBL (ICH4/ICH4-L) \
    LPC Interface Bridge (rev 01)
00:1f.1 IDE interface: Intel Corporation 82801DB (ICH4) IDE \

```

```
Controller (rev 01)
00:1f.3 SMBus: Intel Corporation 82801DB/DBL/DBM (ICH4/ICH4-L/ICH4-M) \
SMBus Controller (rev 01)
00:1f.5 Multimedia audio controller: Intel Corporation 82801DB/DBL/DBM \
(ICH4/ICH4-L/ICH4-M) AC'97 Audio Controller (rev 01)
01:00.0 VGA compatible controller: Matrox Graphics, Inc. G400/G450 (rev 85)
02:08.0 Ethernet controller: Intel Corporation 82801DB PRO/100 VE (LOM) \
Ethernet Controller (rev 81)
```

Using `-v` results in a more detailed listing:

```
linux:~ # lspci  
[...]  
02:08.0 Ethernet controller: Intel Corporation 82801DB PRO/100 VE (LOM)\\  
Ethernet Controller (rev 81)  
    Subsystem: Fujitsu Siemens Computer GmbH: Unknown device 1001  
    Flags: bus master, medium devsel, latency 66, IRQ 11  
    Memory at d1000000 (32-bit, non-prefetchable) [size=4K]  
    I/O ports at 3000 [size=64]  
    Capabilities: [dc] Power Management version 2
```

Information about device name resolution is obtained from file /usr/share/pci.ids. PCI IDs not listed in this file are marked “Unknown device”.

The parameter `-vv` produces all the information that could be queried by the program. To view the pure numeric values, you should use the parameter `-n`.

6.15 System Calls of a Program Run: `strace`

The utility `strace` enables you to trace all the system calls of a process currently running. Enter the command in the normal way, adding `strace` at the beginning of the line:

```

fstat64(3, {st_mode=S_IFREG|0755, st_size=36659, ...}) = 0
[...]
stat64(1, {st_mode=S_IFCHR|0620, st_rdev=makedev(136, 0), ...}) = 0
mmap2(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) \
= 0xb7ca7000
write(1, "bin Desktop Documents music\tM"..., 55) = 55
\ music Music public_html tmp
) = 55
close(1) = 0
munmap(0xb7ca7000, 4096) = 0
exit_group(0) = ?

```

For example, to trace all attempts to open a particular file, use the following:

```

tester@linux:~> strace -e open ls .bashrc
open("/etc/ld.so.cache", O_RDONLY)      = 3
open("/lib/librt.so.1", O_RDONLY)        = 3
open("/lib/libacl.so.1", O_RDONLY)       = 3
open("/lib/libc.so.6", O_RDONLY)         = 3
open("/lib/libpthread.so.0", O_RDONLY)    = 3
open("/lib/libattr.so.1", O_RDONLY)       = 3
[...]

```

To trace all the child processes, use the parameter **-f**. The behavior and output format of strace can be largely controlled. For information, see **man strace**.

6.16 Library Calls of a Program Run: **ltrace**

The command **ltrace** enables you to trace the library calls of a process. This command is used in a similar fashion to **strace**. The parameter **-c** outputs the number and duration of the library calls that have occurred:

```

tester@linux:~> ltrace -c find ~
% time     seconds   usecs/call   calls     function
----- -----
34.37    6.758937    245    27554 __errno_location
33.53    6.593562    788    8358 __fprintf_chk
12.67    2.490392   144    17212 strlen
11.97    2.353302   239    9845 readdir64
  2.37    0.466754    27   16716 __ctype_get_mb_cur_max
  1.18    0.231189    27    8531 strcpy
  1.17    0.230765    27    8358 memcpy
[...]
  0.00    0.000036    36          1 textdomain

```

```
-----  
100.00 19.662715 105717 total
```

6.17 Specifying the Required Library: `ldd`

The command `ldd` can be used to find out which libraries would load the dynamic executable specified as argument:

```
tester@linux:~> ldd /bin/ls  
 linux-gate.so.1 => (0xfffffe000)  
 librt.so.1 => /lib/librt.so.1 (0xb7f97000)  
 libacl.so.1 => /lib/libacl.so.1 (0xb7f91000)  
 libc.so.6 => /lib/libc.so.6 (0xb7e79000)  
 libpthread.so.0 => /lib/libpthread.so.0 (0xb7e67000)  
 /lib/ld-linux.so.2 (0xb7fb6000)  
 libattr.so.1 => /lib/libattr.so.1 (0xb7e63000)
```

Static binaries do not need any dynamic libraries:

```
tester@linux:~> ldd /bin/sash  
 not a dynamic executable  
tester@linux:~> file /bin/sash  
/bin/sash: ELF 32-bit LSB executable, Intel 80386, version 1 (SYSV), for  
GNU/Linux 2.6.4, statically linked, for GNU/Linux 2.6.4, stripped
```

6.18 Additional Information about ELF Binaries

The content of binaries can be read with the `readelf` utility. This even works with ELF files that were built for other hardware architectures:

```
tester@linux:~> readelf --file-header /bin/ls  
ELF Header:  
 Magic: 7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00  
 Class: ELF32  
 Data: 2's complement, little endian  
 Version: 1 (current)  
 OS/ABI: UNIX - System V  
 ABI Version: 0  
 Type: EXEC (Executable file)  
 Machine: Intel 80386  
 Version: 0x1
```

```

Entry point address:          0x8049b60
Start of program headers:    52 (bytes into file)
Start of section headers:    81112 (bytes into file)
Flags:                      0x0
Size of this header:         52 (bytes)
Size of program headers:    32 (bytes)
Number of program headers:  9
Size of section headers:   40 (bytes)
Number of section headers: 30
Section header string table index: 29

```

6.19 Interprocess Communication: ipcs

The command `ipcs` produces a list of the IPC resources currently in use:

```

----- Shared Memory Segments -----
key      shmid   owner   perms      bytes      nattch     status
0x00000000 58261504  tester   600        393216     2          dest
0x00000000 58294273  tester   600        196608     2          dest
0x00000000 83886083  tester   666        43264      2
0x00000000 83951622  tester   666        192000     2
0x00000000 83984391  tester   666        282464     2
0x00000000 84738056  root     644        151552     2          dest

----- Semaphore Arrays -----
key      semid   owner   perms      nsems
0x4d038abf 0          tester   600        8

----- Message Queues -----
key      msqid   owner   perms      used-bytes   messages

```

6.20 Time Measurement with time

The time spent by commands can be determined with the `time` utility. This utility is available in two versions: as a shell built-in and as a program (`/usr/bin/time`).

```

tester@linux:~> time find . > /dev/null
real    0m4.051s
user    0m0.042s
sys     0m0.205s

```


Part III. System

32-Bit and 64-Bit Applications in a 64-Bit System Environment

SUSE Linux is available for several 64-bit platforms. This does not necessarily mean that all the applications included have already been ported to 64-bit platforms. SUSE Linux supports the use of 32-bit applications in a 64-bit system environment. This chapter offers a brief overview of how this support is implemented on 64-bit SUSE Linux platforms. It explains how 32-bit applications are executed (runtime support) and how 32-bit applications should be compiled to enable them to run both in 32-bit and 64-bit system environments. Additionally, find information about the Kernel API and an explanation of how 32-bit applications can run under a 64-bit kernel.

SUSE Linux for the 64-bit platforms AMD64 and EM64T is designed so that existing 32-bit applications run in the 64-bit environment “out-of-the-box.” This support means that you can continue to use your preferred 32-bit applications without waiting for a corresponding 64-bit port to become available.

7.1 Runtime Support

IMPORTANT: Conflicts between Application Versions

If an application is available both for 32-bit and 64-bit environments, the parallel installation of both versions is bound to lead to problems. In such cases, decide on one of the two versions and install and use this.

To be executed correctly, every application requires a range of libraries. Unfortunately, the names for the 32-bit and 64-bit versions of these libraries are identical. They must be differentiated from each other in another way.

To retain compatibility with the 32-bit version, the libraries are stored at the same place in the system as in the 32-bit environment. The 32-bit version of `libc.so.6` is located under `/lib/libc.so.6` in both the 32-bit and 64-bit environments.

All 64-bit libraries and object files are located in directories called `lib64`. The 64-bit object files you would normally expect to find under `/lib`, `/usr/lib`, and `/usr/X11R6/lib` are now found under `/lib64`, `/usr/lib64`, and `/usr/X11R6/lib64`. This means that there is space for the 32-bit libraries under `/lib`, `/usr/lib` and `/usr/X11R6/lib`, so the filename for both versions can remain unchanged.

No subdirectories of the object directories whose data content does not depend on the word size are moved. For example, the X11 fonts are still found in the usual location under `/usr/X11R6/lib/X11/fonts`. This scheme conforms to the LSB (Linux Standards Base) and the FHS (File System Hierarchy Standard).

7.2 Software Development

A biarch development toolchain allows generation of 32-bit and 64-bit objects. The default is to compile 64-bit objects. It is possible to generate 32-bit objects by using special flags. For GCC, this special flag is `-m32`.

All header files must be written in an architecture-independent form. The installed 32-bit and 64-bit libraries must have an API (application programming interface) that matches the installed header files. The normal SUSE environment is designed according to this principle. In the case of manually updated libraries, resolve these issues yourself.

7.3 Software Compilation on Biarch Platforms

To develop binaries for the other architecture on a biarch architecture, the respective libraries for the second architecture must additionally be installed. These packages are called `rpmname-32bit`. You also need the respective headers and libraries from the `rpmname-devel` packages and the development libraries for the second architecture from `rpmname-devel-32bit`.

Most Open Source programs use an `autoconf`-based program configuration. To use `autoconf` for configuring a program for the second architecture, overwrite the normal compiler and linker settings of `autoconf` by running the `configure` script with additional environment variables.

The following example refers to an AMD64 or EM64T system with x86 as the second architecture:

1. Set `autoconf` to use the 32-bit compiler:

```
CC="gcc -m32"
```

2. Instruct the linker to process 32-bit objects:

```
LD="ld -m elf64_i386"
```

3. Set the assembler to generate 32-bit objects:

```
AS="gcc -c -m32"
```

4. Determine that the libraries for libtool and so on come from `/usr/lib`:

```
LDFLAGS="-L/usr/lib"
```

5. Determine that the libraries are stored in the `lib` subdirectory:

```
--libdir=/usr/lib
```

6. Determine that the 32-bit X libraries are used:

```
--x-libraries=/usr/X11R6/lib/
```

Not all of these variables are needed for every program. Adapt them to the respective program.

```
CC="gcc -m64" \
LDFLAGS="-L/usr/lib64;" \
.configure \
--prefix=/usr \
--libdir=/usr/lib64
make
make install
```

7.4 Kernel Specifications

The 64-bit kernels for AMD64 and EM64T offer both a 64-bit and a 32-bit kernel ABI (application binary interface). The latter is identical with the ABI for the corresponding 32-bit kernel. This means that the 32-bit application can communicate with the 64-bit kernel in the same way as with the 32-bit kernel.

The 32-bit emulation of system calls for a 64-bit kernel does not support a number of APIs used by system programs. This depends on the platform. For this reason, a small number of applications, like `lspci` or the LVM administration programs, must be compiled as 64-bit programs to function properly.

A 64-bit kernel can only load 64-bit kernel modules that have been specially compiled for this kernel. It is not possible to use 32-bit kernel modules.

TIP

Some applications require separate kernel-loadable modules. If you intend to use such a 32-bit application in a 64-bit system environment, contact the provider of this application and SUSE to make sure that the 64-bit version of the kernel-loadable module and the 32-bit compiled version of the kernel API are available for this module.

8

Booting and Configuring a Linux System

Booting a Linux system involves various different components. This chapter outlines the underlying principles and highlights the components involved. The concept of runlevels and SUSE's system configuration with `sysconfig` are also discussed in this chapter.

8.1 The Linux Boot Process

The Linux boot process consists of several stages each represented by another component. The following list briefly summarizes the boot process and features all the major components involved.

1. **BIOS** After the computer has been turned on, the BIOS initializes the screen and keyboard and tests the main memory. Up to this stage, the machine does not access any mass storage media. Subsequently, the information about the current date, time, and the most important peripherals are loaded from the CMOS values. When the first hard disk and its geometry are recognized, the system control passes from the BIOS to the boot loader.
2. **Boot Loader** The first physical 512-byte data sector of the first hard disk is loaded into the main memory and the *boot loader* that resides at the beginning of this sector takes over. The commands executed by the boot loader determine the remaining part of the boot process. Therefore, the first 512 bytes on the first hard disk are referred to as the *Master Boot Record* (MBR). The boot loader then passes control to the actual operating system, in this case, the Linux kernel.

More information about GRUB, the Linux boot loader, can be found in [Chapter 9, “The Boot Loader”](#) (page 193).

3. **Kernel and initramfs** To pass system control, the boot loader loads both the kernel and an initial RAM-based file system (initramfs) into memory. The contents of the initial ramfs can be used by the kernel directly. The init ramfs contains a small executable called init that handles the mounting of the real root file system. In former versions of SUSE Linux, these tasks were handled by initrd and linuxrc, respectively. For more information about initramfs, refer to [Section 8.1.1, “initramfs”](#) (page 178).
4. **init on initramfs** This program performs all actions needed to mount the proper root file system, like providing kernel functionality for the needed file system and device drivers for mass storage controllers with udev. After the root file system has been found, it is checked for errors and mounted. If this has been successful, the initramfs is cleaned and the init program on the root file system is executed. For more information about init, refer to [Section 8.1.2, “init on initramfs”](#) (page 179). Find more information about udev in [Chapter 12, “Dynamic Kernel Device Management with udev”](#) (page 251).
5. **init** init handles the actual booting of the system through several different levels providing different functionality. init is described in [Section 8.2, “The init Process”](#) (page 181).

8.1.1 initramfs

initramfs is a small cpio archive that the kernel can load to a RAM disk. It provides a minimal Linux environment that enables the execution of programs before the actual root file system is mounted. This minimal Linux environment is loaded into memory by BIOS routines and does not have specific hardware requirements other than sufficient memory. initramfs must always provide an executable named init that should execute the actual init program on the root file system for the boot process to proceed.

Before the actual root file system can be mounted and the actual operating system can be started, the kernel needs the corresponding drivers to access the device on which the root file system is located. These drivers may include special drivers for certain kinds of hard drives or even network drivers to access a network file system. The needed modules for the root file system may be loaded by init on initramfs. After the modules are loaded, udev provides the initramfs with the needed devices. initramfs is available

during the entire boot process. This makes it possible to handle all device events generated during boot.

If you need to change hardware (hard disks) in an installed system and this hardware requires different drivers to be present in the kernel at boot time, you must update initramfs. This is done in the same way as with initramfs' predecessor, initrd, by calling `mkinitrd`. Calling `mkinitrd` without any argument creates an initramfs. Calling `mkinitrd -R` creates an initrd. In SUSE Linux, the modules to load are specified by the variable `INITRD_MODULES` in `/etc/sysconfig/kernel`. After installation, this variable is automatically set to the correct value. The modules are loaded in exactly the order in which they appear in `INITRD_MODULES`. This is especially important if several SCSI drivers are used, because otherwise the names of the hard disks would change. Strictly speaking, it would be sufficient just to load those drivers needed to access the root file system. However, all SCSI drivers needed for installation are loaded by means of initramfs or initrd because later loading could be problematic.

IMPORTANT: Updating initramfs or initrd

The boot loader loads initramfs or initrd in the same way as the kernel. It is not necessary to reinstall GRUB after updating initramfs or initrd, because GRUB searches the directory for the right file when booting.

8.1.2 init on initramfs

The main purpose of init on initramfs is to prepare the mounting of and access to the real root file system. Depending on your actual system configuration, init is responsible for the following tasks.

Loading Kernel Modules

Depending on your hardware configuration, special drivers may be needed to access the hardware components of your computer (the most important component being your hard drive). To access the final root file system, the kernel needs to load the proper file system drivers.

Providing Special Block Files

For each loaded module, the kernel generates device events. udev handles these events and generates the needed device special files on a RAM file system in `/dev`. Without those special files, the file system would not be accessible.

Managing RAID and LVM Setups

If you configured your system to hold the root file system under RAID or LVM, init sets up LVM or RAID to enable access to the root file system later. Information about RAID can be found in [Section 2.2, “Soft RAID Configuration”](#) (page 60). Information about LVM can be found in [Section 2.1, “LVM Configuration”](#) (page 53).

Managing Network Configuration

If you configured your system to use a network-mounted root file system (mounted via NFS), init must make sure that the proper network drivers are loaded and that they are set up to allow access to the root file system.

When init is called during the initial boot as part of the installation process, its tasks differ from those mentioned earlier:

Finding the Installation Medium

As you start the installation process, your machine loads an installation kernel and a special initrd with the YaST installer from the installation medium. The YaST installer, which is run in a RAM file system, needs to have information about the actual location of the installation medium to access it and install the operating system.

Initiating Hardware Recognition and Loading Appropriate Kernel Modules

As mentioned in [Section 8.1.1, “initramfs”](#) (page 178), the boot process starts with a minimum set of drivers that can be used with most hardware configurations. init starts an initial hardware scanning process that determines the set of drivers suitable for your hardware configuration. These values are later written to `INITRD_MODULES` in `/etc/sysconfig/kernel` to enable any subsequent boot process to use a custom initrd or in a `/etc/sysconfig/hardware/hwconfig-*` file if the device is not needed during the boot process. During the installation process, init loads this set of modules.

Loading the Installation System or Rescue System

As soon as the hardware has been properly recognized, the appropriate drivers have been loaded, and udev has created the device special files, init starts the installation system, which contains the actual YaST installer, or the rescue system.

Starting YaST

Finally, init starts YaST, which starts package installation and system configuration.

8.2 The init Process

The program init is the process with process ID 1. It is responsible for initializing the system in the required way. init takes a special role. It is started directly by the kernel and resists signal 9, which normally kills processes. All other programs are either started directly by init or by one of its child processes.

init is centrally configured in the `/etc/inittab` file where the *runlevels* are defined (see [Section 8.2.1, “Runlevels”](#) (page 181)). The file also specifies which services and daemons are available in each of the levels. Depending on the entries in `/etc/inittab`, several scripts are run by init. For reasons of clarity, these scripts, called *init scripts*, all reside in the directory `/etc/init.d` (see [Section 8.2.2, “Init Scripts”](#) (page 184)).

The entire process of starting the system and shutting it down is maintained by init. From this point of view, the kernel can be considered a background process whose task is to maintain all other processes and adjust CPU time and hardware access according to requests from other programs.

8.2.1 Runlevels

In Linux, *runlevels* define how the system is started and what services are available in the running system. After booting, the system starts as defined in `/etc/inittab` in the line `initdefault`. Usually this is 3 or 5. See [Table 8.1, “Available Runlevels”](#) (page 181). As an alternative, the runlevel can be specified at boot time (at the boot prompt, for instance). Any parameters that are not directly evaluated by the kernel itself are passed to init.

Table 8.1 Available Runlevels

Runlevel	Description
0	System halt
S	Single user mode; from the boot prompt, only with US keyboard mapping
1	Single user mode

Runlevel	Description
2	Local multiuser mode without remote network (NFS, etc.)
3	Full multiuser mode with network
4	Not used
5	Full multiuser mode with network and X display manager—KDM, GDM, or XDM
6	System reboot

IMPORTANT: Avoid Runlevel 2 with a Partition Mounted via NFS

You should not use runlevel 2 if your system mounts a partition like `/usr` via NFS. The system might behave unexpectedly if program files or libraries are missing because the NFS service is not available in runlevel 2 (local multiuser mode without remote network).

To change runlevels while the system is running, enter `telinit` and the corresponding number as an argument. Only the system administrator is allowed to do this. The following list summarizes the most important commands in the runlevel area.

`telinit 1` or `shutdown now`

The system changes to *single user mode*. This mode is used for system maintenance and administration tasks.

`telinit 3`

All essential programs and services (including network) are started and regular users are allowed to log in and work with the system without a graphical environment.

`telinit 5`

The graphical environment is enabled. Usually a display manager like XDM, GDM, or KDM is started. If autologin is enabled, the local user is logged in to the preselected window manager (GNOME or KDE or any other window manager).

```
telinit 0 or shutdown -h now
```

The system halts.

```
telinit 6 or shutdown -r now
```

The system halts then reboots.

Runlevel 5 is the default runlevel in all SUSE Linux standard installations. Users are prompted for login with a graphical interface or the default user is logged in automatically. If the default runlevel is 3, the X Window System must be configured properly, as described in [Chapter 14, *The X Window System*](#) (page 271), before the runlevel can be switched to 5. If this is done, check whether the system works in the desired way by entering `telinit 5`. If everything turns out as expected, you can use YaST to set the default runlevel to 5.

Generally, two things happen when you change runlevels. First, stop scripts of the current runlevel are launched, closing down some programs essential for the current runlevel. Then start scripts of the new runlevel are started. Here, in most cases, a number of programs are started. For example, the following occurs when changing from runlevel 3 to 5:

1. The administrator (`root`) requests init to change to a different runlevel by entering `telinit 5`.
2. init consults its configuration file (`/etc/inittab`) and determines it should start `/etc/init.d/rc` with the new runlevel as a parameter.
3. Now `rc` calls all the stop scripts of the current runlevel, but only those for which there is no start script in the new runlevel. In this example, these are all the scripts that reside in `/etc/init.d/rc3.d` (old runlevel was 3) and start with a K. The number following K specifies the order to start, because there are some dependencies to consider.
4. The last things to start are the start scripts of the new runlevel. These are, in this example, in `/etc/init.d/rc5.d` and begin with an S. The same procedure regarding the order in which they are started is applied here.

When changing into the same runlevel as the current runlevel, init only checks `/etc/inittab` for changes and starts the appropriate steps, for example, for starting a `getty` on another interface. The same functionality may be achieved with the command `telinit q`.

8.2.2 Init Scripts

There are two types of scripts in `/etc/init.d`:

Scripts Executed Directly by init

This is the case only during the boot process or if an immediate system shutdown is initiated (power failure or a user pressing `[Ctrl] + [Alt] + [Del]`). The execution of these scripts is defined in `/etc/inittab`.

Scripts Executed Indirectly by init

These are run when changing the runlevel and always call the master script `/etc/init.d/rc`, which guarantees the correct order of the relevant scripts.

All scripts are located in `/etc/init.d`. Scripts that are run at boot time are called through symbolic links from `/etc/init.d/boot.d`. Scripts for changing the runlevel are called through symbolic links from one of the subdirectories (`/etc/init.d/rc0.d` to `/etc/init.d/rc6.d`). This is just for clarity reasons and avoids duplicate scripts if they are used in several runlevels. Because every script can be executed as both a start and a stop script, these scripts must understand the parameters `start` and `stop`. The scripts also understand the `restart`, `reload`, `force-reload`, and `status` options. These different options are explained in [Table 8.2, “Possible init Script Options”](#) (page 184). Scripts that are run directly by init do not have these links. They are run independently from the runlevel when needed.

Table 8.2 Possible init Script Options

Option	Description
<code>start</code>	Start service.
<code>stop</code>	Stop service.
<code>restart</code>	If the service is running, stop it then restart it. If it is not running, start it.
<code>reload</code>	Reload the configuration without stopping and restarting the service.

Option	Description
force-reload	Reload the configuration if the service supports this. Otherwise, do the same as if <code>restart</code> had been given.
status	Show the current status of service.

Links in each runlevel-specific subdirectory make it possible to associate scripts with different runlevels. When installing or uninstalling packages, these links are added and removed with the help of the program `insserv` (or using `/usr/lib/lsb/install_initd`, which is a script calling this program). See the `insserv(8)` man page for details.

A short introduction to the boot and stop scripts launched first or last, respectively, follows as well as an explanation of the maintaining script.

`boot`

Executed while starting the system directly using `init`. It is independent of the chosen runlevel and is only executed once. Here, the `proc` and `pts` file systems are mounted and `blogd` (boot logging daemon) is activated. If the system is booted for the first time after an update or an installation, the initial system configuration is started.

The `blogd` daemon is a service started by `boot` and `rc` before any other one. It is stopped after the actions triggered by the above scripts (running a number of sub-scripts, for example) are completed. `blogd` writes any screen output to the log file `/var/log/boot.msg`, but only if and when `/var` is mounted read-write. Otherwise, `blogd` buffers all screen data until `/var` becomes available. Get further information about `blogd` on the `blogd(8)` man page.

The script `boot` is also responsible for starting all the scripts in `/etc/init.d/` `boot.d` with a name that starts with `S`. There, the file systems are checked and loop devices are configured if needed. The system time is also set. If an error occurs while automatically checking and repairing the file system, the system administrator can intervene after first entering the root password. Last executed is the script `boot.local`.

`boot.local`

Here, enter additional commands to execute at boot before changing into a runlevel. It can be compared to `AUTOEXEC.BAT` on DOS systems.

`boot.setup`

This script is executed when changing from single user mode to any other runlevel and is responsible for a number of basic settings, such as the keyboard layout and initialization of the virtual consoles.

`halt`

This script is only executed while changing into runlevel 0 or 6. Here, it is executed either as `halt` or as `reboot`. Whether the system shuts down or reboots depends on how `halt` is called.

`rc`

This script calls the appropriate stop scripts of the current runlevel and the start scripts of the newly selected runlevel.

You can create your own scripts and easily integrate them into the scheme described above. For instructions about formatting, naming, and organizing custom scripts, refer to the specifications of the LSB and to the man pages of `init`, `init.d`, and `insserv`. Additionally consult the man pages of `startproc` and `killproc`.

WARNING: Faulty init Scripts May Halt Your System

Faulty init scripts may hang your machine. Edit such scripts with great care and, if possible, subject them to heavy testing in the multiuser environment. Some useful information about init scripts can be found in [Section 8.2.1, “Runlevels” \(page 181\)](#).

To create a custom init script for a given program or service, use the file `/etc/init.d/skeleton` as a template. Save a copy of this file under the new name and edit the relevant program and filenames, paths, and other details as needed. You may also need to enhance the script with your own parts, so the correct actions are triggered by the init procedure.

The `INIT INFO` block at the top is a required part of the script and must be edited. See [Example 8.1, “A Minimal INIT INFO Block” \(page 187\)](#).

Example 8.1 A Minimal INIT INFO Block

```
### BEGIN INIT INFO
# Provides:          FOO
# Required-Start:    $syslog $remote_fs
# Required-Stop:     $syslog $remote_fs
# Default-Start:    3 5
# Default-Stop:     0 1 2 6
# Description:      Start FOO to allow XY and provide YZ
### END INIT INFO
```

In the first line of the `INFO` block, after `Provides:`, specify the name of the program or service controlled by this init script. In the `Required-Start:` and `Required-Stop:` lines, specify all services that need to be started or stopped before the service itself is started or stopped. This information is used later to generate the numbering of script names, as found in the runlevel directories. After `Default-Start:` and `Default-Stop:`, specify the runlevels in which the service should automatically be started or stopped. Finally, for `Description:`, provide a short description of the service in question.

To create the links from the runlevel directories (`/etc/init.d/rc?.d/`) to the corresponding scripts in `/etc/init.d/`, enter the command `insserv new-script-name`. The `insserv` program evaluates the `INIT INFO` header to create the necessary links for start and stop scripts in the runlevel directories (`/etc/init.d/rc?.d/`). The program also takes care of the correct start and stop order for each runlevel by including the necessary numbers in the names of these links. If you prefer a graphical tool to create such links, use the runlevel editor provided by YaST, as described in [Section 8.2.3, “Configuring System Services \(Runlevel\) with YaST”](#) (page 188).

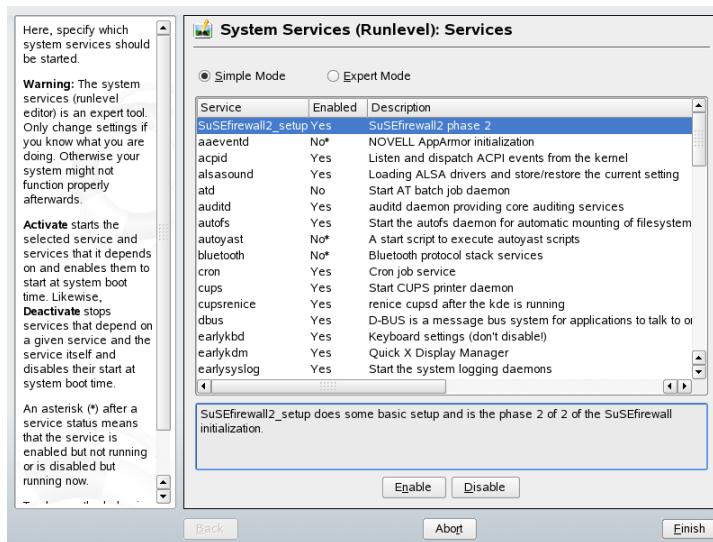
If a script already present in `/etc/init.d/` should be integrated into the existing runlevel scheme, create the links in the runlevel directories right away with `insserv` or by enabling the corresponding service in the runlevel editor of YaST. Your changes are applied during the next reboot—the new service is started automatically.

Do not set these links manually. If something is wrong in the `INFO` block, problems will arise when `insserv` is run later for some other service. The manually-added service will be removed with the next run of `insserv`.

8.2.3 Configuring System Services (Runlevel) with YaST

After starting this YaST module with *YaST* → *System* → *System Services (Runlevel)*, it displays an overview listing all the available services and the current status of each service (disabled or enabled). Decide whether to use the module in *Simple Mode* or in *Expert Mode*. The default *Simple Mode* should be sufficient for most purposes. The left column shows the name of the service, the center column indicates its current status, and the right column gives a short description. For the selected service, a more detailed description is provided in the lower part of the window. To enable a service, select it in the table then select *Enable*. The same steps apply to disable a service.

Figure 8.1 System Services (Runlevel)



For detailed control over the runlevels in which a service is started or stopped or to change the default runlevel, first select *Expert Mode*. The current default runlevel or “initdefault” (the runlevel into which the system boots by default) is displayed at the top. Normally, the default runlevel of a SUSE Linux system is runlevel 5 (full multiuser mode with network and X). A suitable alternative might be runlevel 3 (full multiuser mode with network).

This YaST dialog allows the selection of one of the runlevels (as listed in [Table 8.1, “Available Runlevels”](#) (page 181)) as the new default. Additionally use the table in this window to enable or disable individual services and daemons. The table lists the services and daemons available, shows whether they are currently enabled on your system, and, if so, for which runlevels. After selecting one of the rows with the mouse, click the check boxes representing the runlevels (*B*, *0*, *1*, *2*, *3*, *5*, *6*, and *S*) to define the runlevels in which the selected service or daemon should be running. Runlevel *4* is initially undefined to allow creation of a custom runlevel. A brief description of the currently selected service or daemon is provided below the table overview.

With *Start*, *Stop*, or *Refresh*, decide whether a service should be activated. *Refresh status* checks the current status. *Set* or *Reset* lets you select whether to apply your changes to the system or to restore the settings that existed before starting the runlevel editor. Selecting *Finish* saves the changed settings to disk.

WARNING: Faulty Runlevel Settings May Damage Your System

Faulty runlevel settings may render a system unusable. Before applying your changes, make absolutely sure that you know their consequences.

8.3 System Configuration via /etc/sysconfig

The main configuration of SUSE Linux is controlled by the configuration files in `/etc/sysconfig`. The individual files in `/etc/sysconfig` are only read by the scripts to which they are relevant. This ensures that network settings, for example, only need to be parsed by network-related scripts. Many other system configuration files are generated according to the settings in `/etc/sysconfig`. This task is performed by SuSEconfig. For example, if you change the network configuration, SuSEconfig might make changes to the file `/etc/host.conf` as well, because this is one of the files relevant for the network configuration. This concept enables you to make basic changes to your configuration without needing to reboot the system.

There are two ways to edit the system configuration. Either use the YaST sysconfig Editor or edit the configuration files manually.

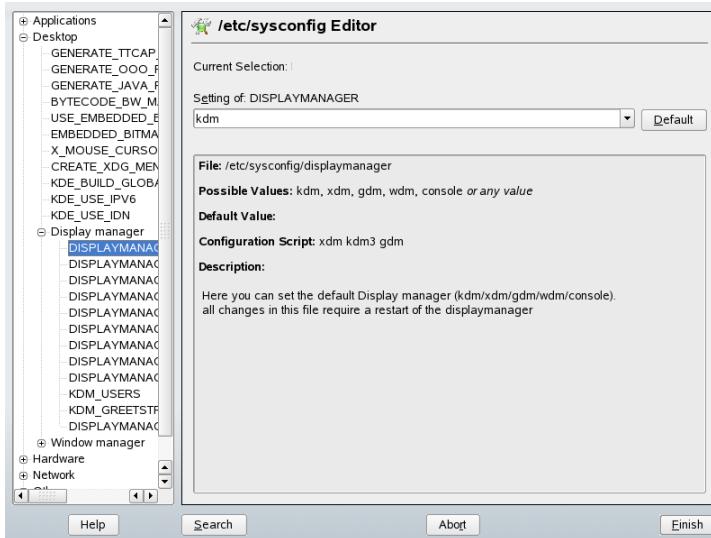
8.3.1 Changing the System Configuration Using the YaST sysconfig Editor

The YaST sysconfig editor provides an easy-to-use front-end to system configuration. Without any knowledge of the actual location of the configuration variable you need to change, you can just use the built-in search function of this module, change the value of the configuration variable as needed, and let YaST take care of applying these changes, updating configurations that depend on the values set in `sysconfig` and restarting services.

WARNING: Modifying `/etc/sysconfig/*` Files Can Damage Your Installation

Do not modify the `/etc/sysconfig` files if you lack previous experience and knowledge. It could do considerable damage to your system. The files in `/etc/sysconfig` include a short comment for each variable to explain what effect they actually have.

Figure 8.2 System Configuration Using the sysconfig Editor



The YaST sysconfig dialog is split into three parts. The left part of the dialog shows a tree view of all configurable variables. When you select a variable, the right part displays a

both the current selection and the current setting of this variable. Below, a third window displays a short description of the variable's purpose, possible values, the default value, and the actual configuration file from which this variable originates. The dialog also provides information about which configuration script is executed after changing the variable and which new service is started as a result of the change. YaST prompts you to confirm your changes and informs you which scripts will be executed after you leave the dialog by selecting *Finish*. Also select the services and scripts to skip for now, so they are started later. YaST applies all changes automatically and restarts any services involved for your changes to take an effect.

8.3.2 Changing the System Configuration Manually

To manually change the system configuration, proceed as follows

- 1** Become `root`.
- 2** Bring the system into single user mode (runlevel 1) with `init 1`.
- 3** Change the configuration files as needed with an editor of your choice.

If you do not use YaST to change the configuration files in `/etc/sysconfig`, make sure that empty variable values are represented by two quotation marks (`KEYTABLE=""`) and that values with blanks in them are enclosed in quotation marks. Values consisting of one word only do not need to be quoted.

- 4** Execute `SuSEconfig` to make sure that the changes take effect.
- 5** Bring your system back to the previous runlevel with a command like `init default_runlevel`. Replace `default_runlevel` with the default runlevel of the system. Choose 5 if you want to return to full multiuser with network and X or choose 3 if you prefer to work in full multiuser with network.

This procedure is mainly relevant when changing systemwide settings, such as the network configuration. Small changes should not require going into single user mode, but you may still do so to make absolutely sure that all the programs concerned are correctly restarted.

TIP: Configuring Automated System Configuration

To disable the automated system configuration by SuSEconfig, set the variable `ENABLE_SUSECONFIG` in `/etc/sysconfig/suseconfig` to `no`. Do not disable SuSEconfig if you want to use the SUSE installation support. It is also possible to disable the autoconfiguration partially.

The Boot Loader

This chapter describes how to configure GRUB, the boot loader used in SUSE Linux. A special YaST module is available for performing all settings. If you are not familiar with the subject of booting in Linux, read the following sections to acquire some background information. This chapter also describes some of the problems frequently encountered when booting with GRUB and their solutions.

This chapter focuses on boot management and the configuration of the boot loader GRUB. The boot procedure as a whole is outlined in [Chapter 8, Booting and Configuring a Linux System](#) (page 177). A boot loader represents the interface between machine (BIOS) and the operating system (SUSE Linux). The configuration of the boot loader directly impacts the start of the operating system.

The following terms appear frequently in this chapter and might need some explanation:

Master Boot Record

The structure of the MBR is defined by an operating system-independent convention. The first 446 bytes are reserved for the program code. They typically hold the boot loader program, in this case, GRUB stage 1. The next 64 bytes provide space for a partition table of up to four entries (see Section “Partition Types” (Chapter 1, *Installation with YaST*, ↑Start-Up)). The partition table contains information about the partitioning of the hard disk and the file system type. The operating system needs this table for handling the hard disk. With GRUB stage 1 in the MBR, exactly one partition must be marked *active*. The last two bytes of the MBR must contain a static “magic number” (AA55). An MBR containing a different value is regarded as invalid by BIOSs.

Boot Sectors

Boot sectors are the first sectors of hard disk partitions with the exception of the extended partition, which merely serves as a “container” for other partitions. These boot sectors have 512 bytes of space for code used to boot an operating system installed in the respective partition. This applies to boot sectors of formatted DOS, Windows, and OS/2 partitions, which also contain some important basic data of the file system. In contrast, the boot sectors of Linux partitions are initially empty after setting up a file system other than XFS. Therefore, a Linux partition is not bootable by itself, even if it contains a kernel and a valid root file system. A boot sector with valid code for booting the system has the same magic number as the MBR in its last two bytes (AA55).

9.1 Selecting a Boot Loader

By default, the boot loader GRUB is used in SUSE Linux. However, in some cases and for special hardware and software constellations, LILO may be necessary. If you update from an older SUSE Linux version that uses LILO, LILO is installed.

Information about the installation and configuration of LILO is available in the Support Database under the keyword LILO and in `/usr/share/doc/packages/lilo`.

9.2 Booting with GRUB

GRUB (Grand Unified Bootloader) comprises two stages. stage1 consists of 512 bytes and its only task is to load the second stage of the boot loader. Subsequently, stage2 is loaded. This stage contains the main part of the boot loader.

In some configurations, an intermediate stage 1.5 can be used, which locates and loads stage 2 from an appropriate file system. If possible, this method is chosen by default on installation or when initially setting up GRUB with YaST.

stage2 is able to access many file systems. Currently, Ext2, Ext3, ReiserFS, Minix, and the DOS FAT file system used by Windows are supported. To a certain extent, JFS, XFS, and UFS and FFS used by BSD systems are also supported. Since version 0.95, GRUB is also able to boot from a CD or DVD containing an ISO 9660 standard file system pursuant to the “El Torito” specification. Even before the system is booted, GRUB can access file systems of supported BIOS disk devices (floppy disks or hard

disks, CD drives, and DVD drives detected by the BIOS). Therefore, changes to the GRUB configuration file (`menu.lst`) do not require a reinstallation of the boot manager. When the system is booted, GRUB reloads the menu file with the valid paths and partition data of the kernel or the initial RAM disk (`initrd`) and locates these files.

The actual configuration of GRUB is based on three files that are described below:

`/boot/grub/menu.lst`

This file contains all information about partitions or operating systems that can be booted with GRUB. Without this information, the GRUB command line prompts the user for how to proceed (see [Section “Editing Menu Entries during the Boot Procedure”](#) (page 199) for details).

`/boot/grub/device.map`

This file translates device names from the GRUB and BIOS notation to Linux device names.

`/etc/grub.conf`

This file contains the commands, parameters, and options the GRUB shell needs for installing the boot loader correctly.

GRUB can be controlled in various ways. Boot entries from an existing configuration can be selected from the graphical menu (splash screen). The configuration is loaded from the file `menu.lst`.

In GRUB, all boot parameters can be changed prior to booting. For example, errors made when editing the menu file can be corrected in this way. Boot commands can also be entered interactively at a kind of input prompt (see [Section “Editing Menu Entries during the Boot Procedure”](#) (page 199)). GRUB offers the possibility of determining the location of the kernel and the `initrd` prior to booting. In this way, you can even boot an installed operating system for which no entry exists in the boot loader configuration.

GRUB actually exists in two versions: as a boot loader and as a normal Linux program in `/usr/sbin/grub`. This program is referred to as the *GRUB shell*. It provides an emulation of GRUB in the installed system and can be used to install GRUB or test new settings before applying them. The functionality to install GRUB as the boot loader on a hard disk or floppy disk is integrated in GRUB in the form of the commands `install` and `setup`. This is available in the GRUB shell when Linux is loaded.

9.2.1 The GRUB Boot Menu

The graphical splash screen with the boot menu is based on the GRUB configuration file `/boot/grub/menu.lst`, which contains all information about all partitions or operating systems that can be booted by the menu.

Every time the system is booted, GRUB loads the menu file from the file system. For this reason, GRUB does not need to be reinstalled after every change to the file. Use the YaST boot loader to modify the GRUB configuration as described in [Section 9.3, “Configuring the Boot Loader with YaST”](#) (page 203).

The menu file contains commands. The syntax is very simple. Every line contains a command followed by optional parameters separated by spaces like in the shell. For historical reasons, some commands permit an `=` in front of the first parameter. Comments are introduced by a hash (`#`).

To identify the menu items in the menu overview, set a `title` for every entry. The text (including any spaces) following the keyword `title` is displayed as a selectable option in the menu. All commands up to the next `title` are executed when this menu item is selected.

The simplest case is the redirection to boot loaders of other operating systems. The command is `chainloader` and the argument is usually the boot block of another partition in GRUB block notation. For example:

```
chainloader (hd0,3)+1
```

The device names in GRUB are explained in [Section “Naming Conventions for Hard Disks and Partitions”](#) (page 197). This example specifies the first block of the fourth partition of the first hard disk.

Use the command `kernel` to specify a kernel image. The first argument is the path to the kernel image in a partition. The other arguments are passed to the kernel on its command line.

If the kernel does not have built-in drivers for access to the root partition or a recent Linux system with advanced hotplug features is used, `initrd` must be specified with a separate GRUB command whose only argument is the path to the `initrd` file. Because the loading address of the `initrd` is written to the loaded kernel image, the command `initrd` must follow immediately after the `kernel` command.

The command `root` simplifies the specification of kernel and `initrd` files. The only argument of `root` is a device or a partition. This device is used for all kernel, `initrd`, or other file paths for which no device is explicitly specified until the next `root` command.

The `boot` command is implied at the end of every menu entry, so it does not need to be written into the menu file. However, if you use GRUB interactively for booting, you must enter the `boot` command at the end. The command itself has no arguments. It merely boots the loaded kernel image or the specified chain loader.

After writing all menu entries, define one of them as the `default` entry. Otherwise, the first one (entry 0) is used. You can also specify a time-out in seconds after which the default entry should boot. `timeout` and `default` usually precede the menu entries. An example file is described in [Section “An Example Menu File” \(page 198\)](#).

Naming Conventions for Hard Disks and Partitions

The naming conventions GRUB uses for hard disks and partitions differ from those used for normal Linux devices. In GRUB, the numbering of the partitions starts with zero. This means that `(hd0, 0)` is the first partition of the first hard disk. On a common desktop machine with a hard disk connected as primary master, the corresponding Linux device name is `/dev/hda1`.

The four possible primary partitions are assigned the partition numbers 0 to 3. The logical partitions are numbered from 4:

```
(hd0,0)    first primary partition of the first hard disk  
(hd0,1)    second primary partition  
(hd0,2)    third primary partition  
(hd0,3)    fourth primary partition (usually an extended partition)  
(hd0,4)    first logical partition  
(hd0,5)    second logical partition
```

Being dependent on BIOS devices, GRUB does not distinguish between IDE, SATA, SCSI, and hardware RAID devices. All hard disks recognized by the BIOS or other controllers are numbered according to the boot sequence preset in the BIOS.

Unfortunately, it is often not possible to map the Linux device names to BIOS device names exactly. It generates this mapping with the help of an algorithm and saves it to the file `device.map`, which can be edited if necessary. Information about the file `device.map` is available in [Section 9.2.2, “The File device.map” \(page 200\)](#).

A complete GRUB path consists of a device name written in parentheses and the path to the file in the file system in the specified partition. The path begins with a slash. For example, the bootable kernel could be specified as follows on a system with a single IDE hard disk containing Linux in its first partition:

```
(hd0,0)/boot/vmlinuz
```

An Example Menu File

The following example shows the structure of a GRUB menu file. The example installation has a Linux boot partition under /dev/hda5, a root partition under /dev/hda7, and a Windows installation under /dev/hda1.

```
gfxmenu (hd0,4)/message
color white/blue black/light-gray
default 0
timeout 8

title linux
  kernel (hd0,4)/vmlinuz root=/dev/hda7 vga=791
  initrd (hd0,4)/initrd

title windows
  chainloader (hd0,0)+1

title floppy
  chainloader (fd0)+1

title failsafe
  kernel (hd0,4)/vmlinuz.shipped root=/dev/hda7 ide=nodma \
  apm=off acpi=off vga=normal nosmp maxcpus=0 3
  initrd (hd0,4)/initrd.shipped
```

The first block defines the configuration of the splash screen:

```
gfxmenu (hd0,4)/message
```

The background image message is located in /dev/hda5.

```
color white/blue black/light-gray
```

Color scheme: white (foreground), blue (background), black (selection), and light gray (background of the selection). The color scheme has no effect on the splash screen, only on the customizable GRUB menu that you can access by exiting the splash screen with **Esc**.

```
default 0
```

The first menu entry `title linux` is the one to boot by default.

```
timeout 8
```

After eight seconds without any user input, GRUB automatically boots the default entry. To deactivate automatic boot, delete the `timeout` line. If you set `timeout 0`, GRUB boots the default entry immediately.

The second and largest block lists the various bootable operating systems. The sections for the individual operating systems are introduced by `title`.

- The first entry (`title linux`) is responsible for booting SUSE Linux. The kernel (`vmlinuz`) is located in the first logical partition (the boot partition) of the first hard disk. Kernel parameters, such as the root partition and VGA mode, are appended here. The root partition is specified according to the Linux naming convention (`/dev/hda7/`), because this information is read by the kernel and has nothing to do with GRUB. The `initrd` is also located in the first logical partition of the first hard disk.
- The second entry is responsible for loading Windows. Windows is booted from the first partition of the first hard disk (`hd0, 0`). The command `chainloader +1` causes GRUB to read and execute the first sector of the specified partition.
- The next entry enables booting from floppy disk without modifying the BIOS settings.
- The boot option `failsafe` starts Linux with a selection of kernel parameters that enables Linux to boot even on problematic systems.

The menu file can be changed whenever necessary. GRUB then uses the modified settings during the next boot. Edit the file permanently using YaST or an editor of your choice. Alternatively, make temporary changes interactively using the `edit` function of GRUB. See [Section “Editing Menu Entries during the Boot Procedure”](#) (page 199).

Editing Menu Entries during the Boot Procedure

In the graphical boot menu, select the operating system to boot with the arrow keys. If you select a Linux system, you can enter additional boot parameters at the boot prompt. To edit individual menu entries directly, press `[Esc]` to exit the splash screen and get to the GRUB text-based menu then press `[E]`. Changes made in this way only apply to the current boot and are not adopted permanently.

IMPORTANT: Keyboard Layout during the Boot Procedure

The US keyboard layout is the only one available when booting.

Editing menu entries facilitates the repair of a defective system that can no longer be booted, because the faulty configuration file of the boot loader can be circumvented by manually entering parameters. Manually entering parameters during the boot procedure is also useful for testing new settings without impairing the native system.

After activating the editing mode, use the arrow keys to select the menu entry of the configuration to edit. To make the configuration editable, press **E** again. In this way, edit incorrect partitions or path specifications before they have a negative effect on the boot process. Press **Enter** to exit the editing mode and return to the menu. Then press **B** to boot this entry. Further possible actions are displayed in the help text at the bottom.

To enter changed boot options permanently and pass them to the kernel, open the file `menu.lst` as the user `root` and append the respective kernel parameters to the existing line, separated by spaces:

```
title linux
  kernel (hd0,0)/vmlinuz root=/dev/hda3 additional parameter
  initrd (hd0,0)/initrd
```

GRUB automatically adopts the new parameters the next time the system is booted. Alternatively, this change can also be made with the YaST boot loader module. Append the new parameters to the existing line, separated by spaces.

9.2.2 The File `device.map`

The file `device.map` maps GRUB and BIOS device names to Linux device names. In a mixed system containing IDE and SCSI hard disks, GRUB must try to determine the boot sequence by a special procedure, because GRUB may not have access to the BIOS information on the boot sequence. GRUB saves the result of this analysis in the file `/boot/grub/device.map`. For a system on which the boot sequence in the BIOS is set to IDE before SCSI, the file `device.map` could appear as follows:

```
(fd0)  /dev/fd0
(hd0)  /dev/hda
(hd1)  /dev/sda
```

Because the order of IDE, SCSI, and other hard disks depends on various factors and Linux is not able to identify the mapping, the sequence in the file `device.map` can be set manually. If you encounter problems when booting, check if the sequence in this file corresponds to the sequence in the BIOS and use the GRUB prompt to modify it temporarily if necessary. After the Linux system has booted, the file `device.map` can be edited permanently with the YaST boot loader module or an editor of your choice.

IMPORTANT: SATA Disks

Depending on the controller, SATA disks are either recognized as IDE (`/dev/hd x`) or SCSI (`/dev/sd x`) devices.

After manually changing `device.map`, execute the following command to reinstall GRUB. This command causes the file `device.map` to be reloaded and the commands listed in `grub.conf` to be executed:

```
grub --batch < /etc/grub.conf
```

9.2.3 The File `/etc/grub.conf`

The third most important GRUB configuration file after `menu.lst` and `device.map` is `/etc/grub.conf`. This file contains the commands, parameters, and options the GRUB shell needs for installing the boot loader correctly:

```
root (hd0,4)
      install /grub/stage1 (hd0,3) /grub/stage2 0x8000 (hd0,4)/grub/menu.lst
      quit
```

Meaning of the individual entries:

`root (hd0,4)`

This command tells GRUB to apply the following commands to the first logical partition of the first hard disk (the location of the boot files).

`install` parameter

The command `grub` should be run with the parameter `install`. `stage1` of the boot loader should be installed in the the extended partition container `(/grub/stage1 (hd0,3))`. `stage2` should be loaded to the memory address `0x8000` `(/grub/stage2 0x8000)`. The last entry `((hd0,4)/grub/menu.lst)` tells GRUB where to look for the menu file.

9.2.4 Setting a Boot Password

Even before the operating system is booted, GRUB enables access to file systems. Users without root permissions can access files in your Linux system to which they have no access once the system is booted. To block this kind of access or prevent users from booting certain operating systems, set a boot password.

IMPORTANT: Boot Password and Splash Screen

If you use a boot password for GRUB, the usual splash screen is not displayed.

As the user `root`, proceed as follows to set a boot password:

- 1 At the root prompt, enter `grub`.
- 2 Encrypt the password in the GRUB shell:

```
grub> md5crypt  
Password: ****  
Encrypted: $1$ls2dv/$JOYcdxIn7CJk9xShzzJVw/
```

- 3 Paste the encrypted string into the global section of the file `menu.lst`:

```
gfxmenu (hd0,4)/message  
color white/blue black/light-gray  
default 0  
timeout 8  
password --md5 $1$ls2dv/$JOYcdxIn7CJk9xShzzJVw/
```

Now GRUB commands can only be executed at the boot prompt after pressing **P** and entering the password. However, users can still boot all operating systems from the boot menu.

- 4 To prevent one or several operating systems from being booted from the boot menu, add the entry `lock` to every section in `menu.lst` that should not be bootable without entering a password. For example:

```
title linux  
kernel (hd0,4)/vmlinuz root=/dev/hda7 vga=791  
initrd (hd0,4)/initrd  
lock
```

After rebooting the system and selecting the Linux entry from the boot menu, the following error message is displayed:

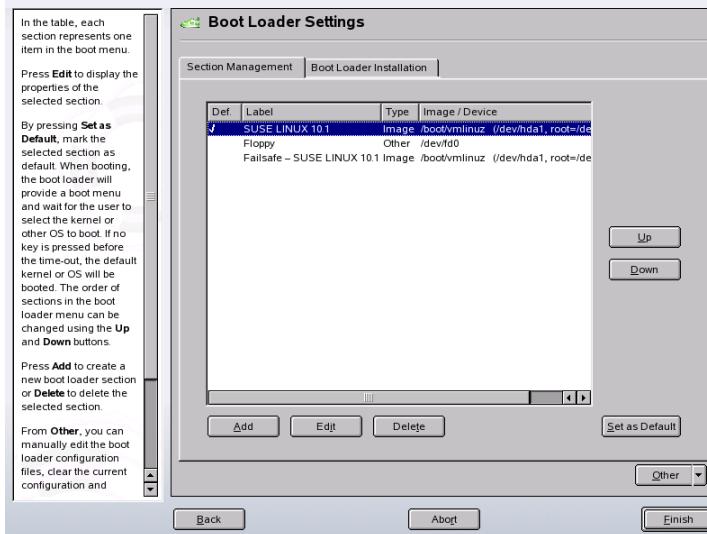
```
Error 32: Must be authenticated
```

Press **[Enter]** to enter the menu. Then press **[P]** to get a password prompt. After entering the password and pressing **[Enter]**, the selected operating system (Linux in this case) should boot.

9.3 Configuring the Boot Loader with YaST

The easiest way to configure the boot loader in your SUSE Linux system is to use the YaST module. In the YaST Control Center, select *System → Boot Loader Configuration*. As in [Figure 9.1, “Configuring the Boot Loader with YaST”](#) (page 203), this shows the current boot loader configuration of your system and allows you to make changes.

Figure 9.1 Configuring the Boot Loader with YaST



Use the *Section Management* tab to edit, change, and delete boot loader sections for the individual operating systems. To add an option, click *Add*. To change the value of an existing option, select it with the mouse and click *Edit*. If you do not want to use an existing option at all, select it and click *Delete*. If you are not familiar with boot loader options, read [Section 9.2, “Booting with GRUB”](#) (page 194) first.

Use the *Boot Loader Installation* tab to view and change settings related to type, location, and advanced loader settings.

9.3.1 Boot Loader Type

Set the boot loader type in *Boot Loader Installation*. The default boot loader in SUSE Linux is GRUB. To use LILO, proceed as follows:

Procedure 9.1 *Changing the Boot Loader Type*

1 Select the *Boot Loader Installation* tab.

2 For *Boot Loader*, select *LILO*.

3 In the dialog box that opens, select one of the following actions:

Propose New Configuration

Have YaST propose a new configuration.

Convert Current Configuration

Have YaST convert the current configuration. When converting the configuration, some settings may be lost.

Start New Configuration from Scratch

Write a custom configuration. This action is not available during the installation of SUSE Linux.

Read Configuration Saved on Disk

Load your own `/etc/lilo.conf`. This action is not available during the installation of SUSE Linux.

4 Click *OK* to save the changes

5 Click *Finish* in the main dialog to apply the changes.

During the conversion, the old GRUB configuration is saved to disk. To use it, simply change the boot loader type back to GRUB and choose *Restore Configuration Saved before Conversion*. This action is available only on an installed system.

NOTE: Custom Boot Loader

If you want use a boot loader other than GRUB or LILO, select *Do Not Install Any Boot Loader*. Read the documentation of your boot loader carefully before choosing this option.

9.3.2 Boot Loader Location

To change the location of the boot loader, follow these steps:

Procedure 9.2 *Changing the Boot Loader Location*

- 1 Select the *Boot Loader Installation* tab then select one of the following options for *Boot Loader Location*:

Master Boot Record of /dev/hdX

This installs the boot loader in the MBR of a disk. X identifies the hard disk, for example, a, b, c, or d:

```
hda => ide0 master  
hdb => ide0 slave  
hdc => ide1 master  
hdd => ide1 slave
```

Boot Sector of Boot Partition /dev/hdXY

The boot sector of the /boot partition. This option is the default if you have several operating systems installed on your hard drive. The Y stands for the partition (1, 2, 3, 4, 5, etc.) as in:

```
/dev/hda1
```

Boot Sector of Root Partition /dev/hdXY

The boot sector of the / (root) partition. Unless a /boot partition is necessary or the MBR needs to be used, this is the preferred default.

Other

Use this option to specify the location of the boot loader manually.

- 2 Click *Finish* to apply your changes.

9.3.3 Default System

To change the system that is booted by default, proceed as follows:

Procedure 9.3 *Setting the Default System*

- 1 Open the *Section Management* tab.
- 2 Select the desired system from the list.
- 3 Click *Set as Default*.
- 4 Click *Finish* to activate these changes.

9.3.4 Boot Loader Time-Out

The boot loader does not boot the default system immediately. During the time-out, you can select the system to boot or write some kernel parameters. To set the boot loader time-out, proceed as follows:

Procedure 9.4 *Changing the Boot Loader Time-Out*

- 1 Open the *Boot Loader Installation* tab.
- 2 Click *Boot Loader Options*.
- 3 Check *Show Boot Menu*.
- 4 In *Boot Menu*, change the value of *Boot Menu Time-Out* by typing in a new value, clicking the appropriate arrow key with your mouse, or by using the arrow keys on the keyboard.
- 5 Click *OK*.
- 6 Click *Finish* to save the changes.

Set for the boot menu should be displayed permanently without timing out by disabling *Continue Booting after a Time-Out*.

9.3.5 Security Settings

Using this YaST module, you can also set a password to protect booting. This gives you an additional level of security.

Procedure 9.5 Setting a Boot Loader Password

- 1** Open the *Boot Loader Installation* tab.
- 2** Click *Boot Loader Options*.
- 3** In *Password Protection*, check *Protect Boot Loader with Password* and set your password.
- 4** Click *OK*.
- 5** Click *Finish* to save the changes.

9.3.6 Disk Order

If your computer has more than one hard disk, you can specify the boot sequence of the disks to match the BIOS setup of the machine (see [Section 9.2.2, “The File device.map”](#) (page 200)). To do so, proceed as follows:

Procedure 9.6 Setting the Disk Order

- 1** Open the *Boot Loader Installation* tab.
- 2** Click *Boot Loader Installation Details*.
- 3** If more than one disk is listed, select a disk and click *Up* or *Down* to reorder the displayed disks.
- 4** Click *OK* to save the changes.
- 5** Click *Finish* to save the changes.

Using this module, you can also replace the master boot record with generic code, which boots the active partition. Click *Replace MBR with Generic Code* in *Disk System Area Update*. Enable *Activate Boot Loader Partition* to activate the partition that contains the boot loader. Click *Finish* to save the changes.

9.4 Uninstalling the Linux Boot Loader

YaST can be used to uninstall the Linux boot loader and restore the MBR to the state it had prior to the installation of Linux. During the installation, YaST automatically creates a backup copy of the original MBR and restores it on request.

To uninstall GRUB, start the YaST boot loader module (*System → Boot Loader Configuration*). In the first dialog, select *Reset → Restore MBR of Hard Disk* and exit the dialog with *Finish*.

9.5 Creating Boot CDs

If problems occur booting your system using a boot manager or if the boot manager cannot be installed on the MBR of your hard disk or a floppy disk, it is also possible to create a bootable CD with all the necessary start-up files for Linux. This requires a CD writer installed in your system.

Creating a bootable CD-ROM with GRUB merely requires a special form of *stage2* called *stage2_eltorito* and, optionally, a customized *menu.lst*. The classic files *stage1* and *stage2* are not required.

Procedure 9.7 Creating Boot CDs

- 1 Create a directory in which to create the ISO image, for example:

```
cd /tmp  
mkdir iso
```

- 2 Create a subdirectory for GRUB:

```
mkdir -p iso/boot/grub
```

- 3** Copy the kernel, the files `stage2_eltorito`, `initrd`, `menu.lst`, and `/boot/message` to `iso/boot/`:

```
cp /boot/vmlinuz iso/boot/
cp /boot/initrd iso/boot/
cp /boot/message iso/boot/
cp /boot/grub/menu.lst iso/boot/grub
```

- 4** Adjust the path entries in `iso/boot/menu.lst` to make them point to a CD-ROM device. Do this by replacing the device name of the hard disks, listed in the format `(hd*)`, in the pathnames with the device name of the CD-ROM drive, which is `(cd)`:

```
gfxmenu (cd)/boot/message
timeout 8
default 0

title Linux
kernel (cd)/boot/vmlinuz root=/dev/hda5 vga=794 resume=/dev/hda1 \
splash=verbose showopts
initrd (cd)/boot/initrd
```

- 5** Create the ISO image with the following command:

```
mkisofs -R -b boot/grub/stage2_eltorito -no-emul-boot \
-boot-load-size 4 -boot-info-table -o grub.iso iso
```

- 6** Write the resulting file `grub.iso` to a CD using your preferred utility.

9.6 The Graphical SUSE Screen

Since SUSE Linux 7.2, the graphical SUSE screen is displayed on the first console if the option “`vga=<value>`” is used as a kernel parameter. If you install using YaST, this option is automatically activated in accordance with the selected resolution and the graphics card. There are three ways to disable the SUSE screen, if desired:

Disabling the SUSE Screen When Necessary

Enter the command `echo 0 >/proc/splash` on the command line to disable the graphical screen. To activate it again, enter `echo 1 >/proc/splash`.

Disabling the SUSE screen by default.

Add the kernel parameter `splash=0` to your boot loader configuration. Chapter 9, *The Boot Loader* (page 193) provides more information about this. However, if you prefer the text mode, which was the default in earlier versions, set `vga=normal`.

Completely Disabling the SUSE Screen

Compile a new kernel and disable the option *Use splash screen instead of boot logo in framebuffer support*.

TIP

Disabling framebuffer support in the kernel automatically disables the splash screen as well. SUSE cannot provide any support for your system if you run it with a custom kernel.

9.7 Troubleshooting

This section lists some of the problems frequently encountered when booting with GRUB and a short description of possible solutions. Some of the problems are covered in articles in the Support Database at <http://portal.suse.de/sdb/en/index.html>. If your specific problem is not included in this list, use the search dialog of the Support Database at <https://portal.suse.com/PM/page/search.pm> to search for keywords like *GRUB*, *boot*, and *boot loader*.

GRUB and XFS

XFS leaves no room for `stage1` in the partition boot block. Therefore, do not specify an XFS partition as the location of the boot loader. This problem can be solved by creating a separate boot partition that is not formatted with XFS.

GRUB and JFS

Although technically possible, the combination of GRUB with JFS is problematic. In this case, create a separate boot partition (`/boot`) and format it with Ext2. Install GRUB in this partition.

GRUB Reports GRUB Geom Error

GRUB checks the geometry of connected hard disks when the system is booted. Sometimes, the BIOS returns inconsistent information and GRUB reports a GRUB Geom Error. If this is the case, use LILO or update the BIOS. Detailed information

about the installation, configuration, and maintenance of LILO is available in the Support Database under the keyword LILO.

GRUB also returns this error message if Linux was installed on an additional hard disk that is not registered in the BIOS. *stage1* of the boot loader is found and loaded correctly, but *stage2* is not found. This problem can be remedied by registering the new hard disk in the BIOS.

System Containing IDE and SCSI Hard Disks Does Not Boot

During the installation, YaST may have incorrectly determined the boot sequence of the hard disks. For example, GRUB may regard /dev/hda as hd0 and /dev/sda as hd1, although the boot sequence in the BIOS is reversed (*SCSI before IDE*).

In this case, correct the hard disks during the boot process with the help of the GRUB command line. After the system has booted, edit `device.map` to apply the new mapping permanently. Then check the GRUB device names in the files `/boot/grub/menu.lst` and `/boot/grub/device.map` and reinstall the boot loader with the following command:

```
grub --batch < /etc/grub.conf
```

Booting Windows from the Second Hard Disk

Some operating systems, such as Windows, can only boot from the first hard disk. If such an operating system is installed on a hard disk other than the first hard disk, you can effect a logical change for the respective menu entry.

```
...
title windows
map (hd0) (hd1)
map (hd1) (hd0)
chainloader (hd1,0)+1
...
```

In this example, Windows is started from the second hard disk. For this purpose, the logical order of the hard disks is changed with `map`. This change does not affect the logic within the GRUB menu file. Therefore, the second hard disk must be specified for `chainloader`.

9.8 For More Information

Extensive information about GRUB is available at <http://www.gnu.org/software/grub/>. Also refer to the grub info page. You can also search for the keyword “GRUB” in the Support Database at <http://portal.suse.de/sdb/en/index.html> to get information about special issues.

10

Special Features of SUSE Linux

This chapter starts with information about various software packages, the virtual consoles, and the keyboard layout. We talk about software components like `bash`, `cron`, and `logrotate`, because they were changed or enhanced during the last release cycles. Even if they are small or considered of minor importance, users may want to change their default behavior, because these components are often closely coupled with the system. The chapter is finished by a section about language and country-specific settings (I18N and L10N).

10.1 Information about Special Software Packages

The programs `bash`, `cron`, `logrotate`, `locate`, `ulimit`, and `free`, and the file `resolv.conf` are very important for system administrators and many users. Man pages and info pages are two useful sources of information about commands, but both are not always available. GNU Emacs is a popular and very configurable text editor.

10.1.1 The Package `bash` and `/etc/profile`

Bash is the default shell in SUSE Linux. When used as a login shell, it reads several initialization files. Bash processes them in the order they appear in this list.

1. /etc/profile
2. ~/.profile
3. /etc/bash.bashrc
4. ~/.bashrc

Custom settings can be made in `~/.profile` or in `~/.bashrc`. To ensure the correct processing of these files, it is necessary to copy the basic settings from `/etc/skel/.profile` or `/etc/skel/.bashrc` into the home directory of the user. It is recommended to copy the settings from `/etc/skel` following an update. Execute the following shell commands to prevent the loss of personal adjustments:

```
mv ~/.bashrc ~/.bashrc.old  
cp /etc/skel/.bashrc ~/.bashrc  
mv ~/.profile ~/.profile.old  
cp /etc/skel/.profile ~/.profile
```

Then copy personal adjustments back from the `*.old` files.

10.1.2 The cron Package

If you want to run commands regularly and automatically in the background at predefined times, cron is the traditional tool to use. cron is driven by specially formatted time tables. Some of them come with the system and users can write their own tables if needed.

The cron tables are located in `/var/spool/cron/tabs`. `/etc/crontab` serves as a systemwide cron table. Enter the username to run the command directly after the time table and before the command. In [Example 10.1, “Entry in /etc/crontab”](#) (page 214), `root` is entered. Package-specific tables, located in `/etc/cron.d`, have the same format. See the `cron` man page (`man cron`).

Example 10.1 Entry in `/etc/crontab`

```
1-59/5 * * * *    root    test -x /usr/sbin/atrun && /usr/sbin/atrun
```

You cannot edit `/etc/crontab` by calling the command `crontab -e`. This file must be loaded directly into an editor, modified, then saved.

A number of packages install shell scripts to the directories `/etc/cron.hourly`, `/etc/cron.daily`, `/etc/cron.weekly`, and `/etc/cron.monthly`, whose execution is controlled by `/usr/lib/cron/run-crons`. `/usr/lib/cron/run-crons` is run every 15 minutes from the main table (`/etc/crontab`). This guarantees that processes that may have been neglected can be run at the proper time.

To run the `hourly`, `daily`, or other periodic maintenance scripts at custom times, remove the time stamp files regularly using `/etc/crontab` entries (see [Example 10.2, “/etc/crontab: Remove Time Stamp Files”](#) (page 215), which removes the `hourly` one before every full hour, the `daily` one once a day at 2:14 a.m., etc.).

Example 10.2 /etc/crontab: Remove Time Stamp Files

```
59 * * * *      root  rm -f /var/spool/cron/lastrun/cron.hourly
14 2 * * *      root  rm -f /var/spool/cron/lastrun/cron.daily
29 2 * * 6      root  rm -f /var/spool/cron/lastrun/cron.weekly
44 2 1 * *      root  rm -f /var/spool/cron/lastrun/cron.monthly
```

The daily system maintenance jobs have been distributed to various scripts for reasons of clarity. They are contained in the package `aaa_base`. `/etc/cron.daily` contains, for example, the components `suse.de-backup-rpmdb`, `suse.de-clean-tmp`, or `suse.de-cron-local`.

10.1.3 Log Files: Package logrotate

There are a number of system services (*daemons*) that, along with the kernel itself, regularly record the system status and specific events to log files. This way, the administrator can regularly check the status of the system at a certain point in time, recognize errors or faulty functions, and troubleshoot them with pinpoint precision. These log files are normally stored in `/var/log` as specified by FHS and grow on a daily basis. The `logrotate` package helps control the growth of these files.

Configure `logrotate` with the file `/etc/logrotate.conf`. In particular, the `include` specification primarily configures the additional files to read. SUSE Linux ensures that programs that produce log files install individual configuration files in `/etc/logrotate.d`. For example, such programs come with the packages `apache2` (`/etc/logrotate.d/apache2`) and `syslogd` (`/etc/logrotate.d/syslog`).

Example 10.3 Example for /etc/logrotate.conf

```
# see "man logrotate" for details
# rotate log files weekly
weekly

# keep 4 weeks worth of backlogs
rotate 4

# create new (empty) log files after rotating old ones
create

# uncomment this if you want your log files compressed
#compress

# RPM packages drop log rotation information into this directory
include /etc/logrotate.d

# no packages own lastlog or wtmp - we'll rotate them here
#/var/log/wtmp {
#    monthly
#    create 0664 root utmp
#    rotate 1
#}

# system-specific logs may be also be configured here.
```

logrotate is controlled through cron and is called daily by /etc/cron.daily/logrotate.

IMPORTANT

The `create` option reads all settings made by the administrator in /etc/permissions*. Ensure that no conflicts arise from any personal modifications.

10.1.4 The Command locate

locate, a command for quickly finding files, is not included in the standard scope of installed software. If desired, install the package `find-locate`. The `updatedb` process is started automatically every night or about 15 minutes after booting the system.

10.1.5 The Command ulimit

With the `ulimit` (*user limits*) command, it is possible to set limits for the use of system resources and to have these displayed. `ulimit` is especially useful for limiting the memory available for applications. With this, an application can be prevented from using too much memory on its own, which could bring the system to a standstill.

`ulimit` can be used with various options. To limit memory usage, use the options listed in [Table 10.1, “`ulimit`: Setting Resources for the User”](#) (page 217).

Table 10.1 *ulimit: Setting Resources for the User*

-m	Maximum size of physical memory
-v	Maximum size of virtual memory
-s	Maximum size of the stack
-c	Maximum size of the core files
-a	Display of limits set

Systemwide entries can be made in `/etc/profile`. There, enable creation of core files, needed by programmers for *debugging*. A normal user cannot increase the values specified in `/etc/profile` by the system administrator, but can make special entries in `~/ .bashrc`.

Example 10.4 *ulimit: Settings in `~/bashrc`*

```
# Limits of physical memory:  
ulimit -m 98304  
  
# Limits of virtual memory:  
ulimit -v 98304
```

Memory amounts must be specified in KB. For more detailed information, see `man bash`.

IMPORTANT

Not all shells support `ulimit` directives. PAM (for instance, `pam_limits`) offers comprehensive adjustment possibilities if you depend on encompassing settings for these restrictions.

10.1.6 The free Command

The `free` command is somewhat misleading if your goal is to find out how much RAM is currently being used. That information can be found in `/proc/meminfo`. These days, users with access to a modern operating system, such as Linux, should not really need to worry much about memory. The concept of *available RAM* dates back to before the days of unified memory management. The slogan *free memory is bad memory* applies well to Linux. As a result, Linux has always made the effort to balance out caches without actually allowing free or unused memory.

Basically, the kernel does not have direct knowledge of any applications or user data. Instead, it manages applications and user data in a *page cache*. If memory runs short, parts of it are written to the swap partition or to files, from which they can initially be read with the help of the `mmap` command (see `man mmap`).

The kernel also contains other caches, such as the *slab cache*, where the caches used for network access are stored. This may explain differences between the counters in `/proc/meminfo`. Most, but not all of them, can be accessed via `/proc/slabinfo`.

10.1.7 The File /etc/resolv.conf

Domain name resolution is handled through the file `/etc/resolv.conf`. Refer to Chapter 20, *The Domain Name System* (page 367).

This file is updated by the script `/sbin/modify_resolvconf` exclusively, with no other program having permission to modify `/etc/resolv.conf` directly. Enforcing this rule is the only way to guarantee that the system's network configuration and the relevant files are kept in a consistent state.

10.1.8 Man Pages and Info Pages

For some GNU applications (such as tar), the man pages are no longer maintained. For these commands, use the `--help` option to get a quick overview of the info pages, which provide more in-depth instructions. info is GNU's hypertext system. Read an introduction to this system by entering `info info`. Info pages can be viewed with Emacs by entering `emacs -f info` or directly in a console with `info`. You can also use `tkinfo`, `xinfo`, or the SUSE help system to view info pages.

10.1.9 Settings for GNU Emacs

GNU Emacs is a complex work environment. The following sections cover the configuration files processed when GNU Emacs is started. More information is available at <http://www.gnu.org/software/emacs/>.

On start-up, Emacs reads several files containing the settings of the user, system administrator, and distributor for customization or preconfiguration. The initialization file `~/.emacs` is installed to the home directories of the individual users from `/etc/skel/.emacs`, in turn, reads the file `/etc/skel/.gnu-emacs`. To customize the program, copy `.gnu-emacs` to the home directory (with `cp /etc/skel/.gnu-emacs ~/.gnu-emacs`) and make the desired settings there.

`.gnu-emacs` defines the file `~/.gnu-emacs-custom` as `custom-file`. If users make settings with the `customize` options in Emacs, the settings are saved to `~/.gnu-emacs-custom`.

With SUSE Linux, the `emacs` package installs the file `site-start.el` in the directory `/usr/share/emacs/site-lisp`. The file `site-start.el` is loaded before the initialization file `~/.emacs`. Among other things, `site-start.el` ensures that special configuration files distributed with Emacs add-on packages, such as `psgml`, are loaded automatically. Configuration files of this type are located in `/usr/share/emacs/site-lisp`, too, and always begin with `suse-start-`. The local system administrator can specify systemwide settings in `default.el`.

More information about these files is available in the Emacs info file under *Init File: info:/emacs/InitFile*. Information about how to disable loading these files (if necessary) is also provided at this location.

The components of Emacs are divided into several packages:

- The base package `emacs`.
- `emacs-x11` (usually installed): the program *with* X11 support.
- `emacs-nox`: the program *without* X11 support.
- `emacs-info`: online documentation in info format.
- `emacs-el`: the uncompiled library files in Emacs Lisp. These are not required at runtime.
- Numerous add-on packages can be installed if needed: `emacs-auctex` (for LaTeX), `psgml` (for SGML and XML), `gnuserv` (for client and server operation), and others.

10.2 Virtual Consoles

Linux is a multiuser and multitasking system. The advantages of these features can be appreciated even on a stand-alone PC system. In text mode, there are six virtual consoles available. Switch between them using `[Alt] + [F1]` to `[Alt] + [F6]`. The seventh console is reserved for X and the tenth console shows kernel messages. More or fewer consoles can be assigned by modifying the file `/etc/inittab`.

To switch to a console from X without shutting it down, use `[Ctrl] + [Alt] + [F1]` to `[Ctrl] + [Alt] + [F6]`. To return to X, press `[Alt] + [F7]`.

10.3 Keyboard Mapping

To standardize the keyboard mapping of programs, changes were made to the following files:

```
/etc/inputrc  
/usr/X11R6/lib/X11/Xmodmap  
/etc/skel/.Xmodmap  
/etc/skel/.exrc  
/etc/skel/.less  
/etc/skel/.lesskey  
/etc/csh.cshrc
```

```
/etc/termcap  
/usr/lib/terminfo/x/xterm  
/usr/X11R6/lib/X11/app-defaults/XTerm  
/usr/share/emacs/<VERSION>/site-lisp/term/*.el
```

These changes only affect applications that use `terminfo` entries or whose configuration files are changed directly (`vi`, `less`, etc.). Applications not shipped with SUSE Linux should be adapted to these defaults.

Under X, the compose key (multikey) can be accessed using `[Ctrl] + [Shift]` (right). Also see the corresponding entry in `/usr/X11R6/lib/X11/Xmodmap`.

Further settings are possible using the X Keyboard Extension (XKB). This extension is also used by the desktop environments GNOME (`gswitchit`) and KDE (`kxkb`).

TIP: For More Information

Information about XKB is available in `/etc/X11/xkb/README` and the documents listed there.

Detailed information about the input of Chinese, Japanese, and Korean (CJK) is available at Mike Fabian's page: <http://www.suse.de/~mfabian/suse-cjk/input.html>.

10.4 Language and Country-Specific Settings

SUSE Linux is, to a very large extent, internationalized and can be modified for local needs in a flexible manner. In other words, internationalization (*I18N*) allows specific localizations (*L10N*). The abbreviations I18N and L10N are derived from the first and last letters of the words and, in between, the number of letters omitted.

Settings are made with `LC_` variables defined in the file `/etc/sysconfig/language`. This refers not only to *native language support*, but also to the categories *Messages (Language)*, *Character Set*, *Sort Order*, *Time and Date*, *Numbers*, and *Money*. Each of these categories can be defined directly with its own variable or indirectly with a master variable in the file `language` (see the `locale` man page).

RC_LC_MESSAGES, RC_LC_CTYPE, RC_LC_COLLATE, RC_LC_TIME,
RC_LC_NUMERIC, RC_LC_MONETARY

These variables are passed to the shell without the RC_ prefix and represent the listed categories. The shell profiles concerned are listed below. The current setting can be shown with the command `locale`.

RC_LC_ALL

This variable, if set, overwrites the values of the variables already mentioned.

RC_LANG

If none of the previous variables are set, this is the fallback. By default, SUSE Linux only sets RC_LANG. This makes it easier for users to enter their own values.

ROOT_USES_LANG

A yes or no variable. If it is set to no, root always works in the POSIX environment.

The variables can be set with the YaST sysconfig editor (see [Section 8.3.1, “Changing the System Configuration Using the YaST sysconfig Editor”](#) (page 190)). The value of such a variable contains the language code, country code, encoding, and modifier. The individual components are connected by special characters:

```
LANG=<language>[[_<COUNTRY>].<Encoding>[@<Modifier>]]
```

10.4.1 Some Examples

You should always set the language and country codes together. Language settings follow the standard ISO 639 available at <http://www.evertype.com/standards/iso639/iso639-en.html> and <http://www.loc.gov/standards/iso639-2/>. Country codes are listed in ISO 3166 available at http://www.din.de/gremien/nas/nabd/iso3166ma/codlstp1/en_listp1.html.

It only makes sense to set values for which usable description files can be found in /usr/lib/locale. Additional description files can be created from the files in /usr/share/i18n using the command `localeddef`. The description files are part of the glibc-i18n package. A description file for en_US.UTF-8 (for English and United States) can be created with:

```
localedef -i en_US -f UTF-8 en_US.UTF-8
```

```
LANG=en_US.UTF-8
```

This is the default setting if American English is selected during installation. If you selected another language, that language is enabled but still with UTF-8 as the character encoding.

```
LANG=en_US.ISO-8859-1
```

This sets the language to English, country to United States, and the character set to ISO-8859-1. This character set does not support the Euro sign, but it can be useful sometimes for programs that have not been updated to support UTF-8. The string defining the charset (ISO-8859-1 in this case) is then evaluated by programs like Emacs.

```
LANG=en_IE@euro
```

The above example explicitly includes the Euro sign in a language setting. Strictly speaking, this setting is obsolete now, because UTF-8 also covers the Euro symbol. It is only useful if an application does not support UTF-8, but ISO-8859-15.

SuSEconfig reads the variables in /etc/sysconfig/language and writes the necessary changes to /etc/SuSEconfig/profile and /etc/SuSEconfig/csh.cshrc. /etc/SuSEconfig/profile is read or sourced by /etc/profile. /etc/SuSEconfig/csh.cshrc is sourced by /etc/csh.cshrc. This makes the settings available systemwide.

Users can override the system defaults by editing their `~/.bashrc` accordingly. For instance, if you do not want to use the systemwide `en_US` for program messages, include `LC_MESSAGES=es_ES` so messages are displayed in Spanish instead.

10.4.2 Locale Settings in `~/.i18n`

If you are not satisfied with locale system defaults, change the settings in `~/.i18n`. Entries in `~/.i18n` override system defaults from /etc/sysconfig/language. Use the same variable names but without the `RC_` namespace prefixes, for example, use `LANG` instead of `RC_LANG`.

10.4.3 Settings for Language Support

Files in the category *Messages* are, as a rule, only stored in the corresponding language directory (like `en`) to have a fallback. If you set `LANG` to `en_US` and the message file

in `/usr/share/locale/en_US/LC_MESSAGES` does not exist, it falls back to `/usr/share/locale/en/LC_MESSAGES`.

A fallback chain can also be defined, for example, for Breton to French or for Galician to Spanish to Portuguese:

```
LANGUAGE="br_FR:fr_FR"
```

```
LANGUAGE="gl_ES:es_ES:pt_PT"
```

If desired, use the Norwegian variants Nynorsk and Bokmål instead (with additional fallback to no):

```
LANG="nn_NO"
```

```
LANGUAGE="nn_NO:nb_NO:no"
```

or

```
LANG="nb_NO"
```

```
LANGUAGE="nb_NO:nn_NO:no"
```

Note that in Norwegian, `LC_TIME` is also treated differently.

One problem that can arise is a separator used to delimit groups of digits not being recognized properly. This occurs if `LANG` is set to only a two-letter language code like `de`, but the definition file glibc uses is located in `/usr/share/lib/de_DE/LC_NUMERIC`. Thus `LC_NUMERIC` must be set to `de_DE` to make the separator definition visible to the system.

10.4.4 For More Information

- *The GNU C Library Reference Manual*, Chapter “Locales and Internationalization”. It is included in `glibc-info`.
- Markus Kuhn, *UTF-8 and Unicode FAQ for Unix/Linux*, currently at <http://www.cl.cam.ac.uk/~mgk25/unicode.html>.

- *Unicode-Howto*, by Bruno Haible: /usr/share/doc/howto/en/html/Unicode-HOWTO.html.

Printer Operation

CUPS is the standard print system in SUSE Linux. CUPS is highly user-oriented. In many cases, it is compatible with LPRng or can be adapted with relatively little effort. LPRng is included in SUSE Linux only for reasons of compatibility.

Printers can be distinguished by interface, such as USB or network, and printer language. When buying a printer, make sure that the printer has an interface that is supported by the hardware and a suitable printer language. Printers can be categorized on the basis of the following three classes of printer languages:

PostScript Printers

PostScript is the printer language in which most print jobs in Linux and Unix are generated and processed by the internal print system. This language is already quite old and very efficient. If PostScript documents can be processed directly by the printer and do not need to be converted in additional stages in the print system, the number of potential error sources is reduced. Because PostScript printers are subject to substantial license costs, these printers usually cost more than printers without a PostScript interpreter.

Standard Printer (languages like PCL and ESC/P)

Although these printer languages are quite old, they are still undergoing expansion to address new features in printers. In the case of known printer languages, the print system can convert PostScript jobs to the respective printer language with the help of Ghostscript. This processing stage is referred to as interpreting. The best-known languages are PCL, which is mostly used by HP printers and their clones, and ESC/P, which is used by Epson printers. These printer languages are usually supported by Linux and produce a decent print result. Linux may not be able to address some functions of extremely new and fancy printers, because the open

source developers may still be working on these features. Except for the `hpijs` drivers developed by HP, there are currently no printer manufacturers who develop Linux drivers and make them available to Linux distributors under an open source license. Most of these printers are in the medium price range.

Proprietary Printers (usually GDI printers)

Usually only one or several Windows drivers are available for proprietary printers. These printers do not support any of the common printer languages and the printer languages they use are subject to change when a new edition of a model is released. See [Section 11.7.1, “Printers without Standard Printer Language Support”](#) (page 243) for more information.

Before you buy a new printer, refer to the following sources to check how well the printer you intend to buy is supported:

- <http://cdb.suse.de/>—the SUSE Linux printer database
- <http://www.linuxprinting.org/>—the LinuxPrinting.org printer database
- <http://www.cs.wisc.edu/~ghost/>—the Ghostscript Web page
- `/usr/share/doc/packages/ghostscript/catalog.devices`—list of included drivers

The online databases always show the latest Linux support status. However, a Linux distribution can only integrate the drivers available at production time. Accordingly, a printer currently rated as “perfectly supported” may not have had this status when the latest SUSE Linux version was released. Thus, the databases may not necessarily indicate the correct status, but only provide an approximation.

11.1 Workflow of the Printing System

The user creates a print job. The print job consists of the data to print plus information for the spooler, such as the name of the printer or the name of the printer queue, and, optionally, the information for the filter, such as printer-specific options.

A dedicated printer queue exists for every printer. The spooler holds the print job in the queue until the desired printer is ready to receive data. When the printer is ready, the spooler sends the data through the filter and back-end to the printer.

The filter converts the data the user wants to print (ASCII, PostScript, PDF, JPEG, etc.) into printer-specific data (PostScript, PCL, ESC/P, etc.). The features of the printer are described in the PPD files. A PPD file contains printer-specific options with the parameters needed to enable them on the printer. The filter system makes sure that options selected by the user are enabled.

If you use a PostScript printer, the filter system converts the data into printer-specific PostScript. This does not require a printer driver. If you use a non-PostScript printer, the filter system converts the data into printer-specific data using Ghostscript. This requires a Ghostscript printer driver suitable for your printer. The back-end receives the printer-specific data from the filter passes it to the printer.

11.2 Methods and Protocols for Connecting Printers

There are various possibilities for connecting a printer to the system. The configuration of the CUPS print system does not distinguish between a local printer and a printer connected to the system over the network. In Linux, local printers must be connected as described in the manual of the printer manufacturer. CUPS supports serial, USB, parallel, and SCSI connections. For more information about the printer connection, read the article *CUPS in a Nutshell* in the Support Database at <http://portal.suse.com>. Find the article by entering *cups* in the search dialog.

WARNING: Cable Connection to the Machine

When connecting the printer to the machine, do not forget that only USB devices can be plugged in or unplugged during operation. The system should be shut down before changing other kinds of connections.

11.3 Installing the Software

PPD (PostScript printer description) is the computer language that describes the properties, like resolution, and options, such as the availability of a duplex unit. These descriptions are required for using various printer options in CUPS. Without a PPD file, the print data would be forwarded to the printer in a “raw” state, which is usually not

desired. During the installation of SUSE Linux, many PPD files are preinstalled to enable even printers without PostScript support to be used.

To configure a PostScript printer, the best approach is to get a suitable PPD file. Many PPD files are available in the package `manufacturer-PPDs`, which is automatically installed within the scope of the standard installation. See [Section 11.6.3, “PPD Files in Various Packages”](#) (page 240) and [Section 11.7.2, “No Suitable PPD File Available for a PostScript Printer”](#) (page 243).

New PPD files can be stored in the directory `/usr/share/cups/model/` or added to the print system with YaST (see [Section “Manual Configuration”](#) (page 231)). Subsequently, the PPD file can be selected during the installation.

Be careful if a printer manufacturer wants you to install entire software packages in addition to modifying configuration files. First, this kind of installation would result in the loss of the support provided by SUSE Linux and, second, print commands may work differently and the system may no longer be able to address devices of other manufacturers. For this reason, the installation of manufacturer software is not recommended.

11.4 Configuring the Printer

After connecting the printer to the computer and installing the software, install the printer in the system. This should be done with the tools delivered with SUSE Linux. Because SUSE Linux puts great emphasis on security, third-party tools often have difficulties with the security restrictions and cause more complications than benefits. See [Section 11.6.1, “CUPS Server and Firewall”](#) (page 238) and [Section 11.6.2, “Changes in the CUPS Print Service”](#) (page 239) for more information about troubleshooting.

11.4.1 Local Printers

If an unconfigured local printer is detected when you log in, YaST starts for configuring it. This uses the same dialogs as the following description of configuration.

To configure the printer, select *Hardware → Printer* in the YaST control center. This opens the main printer configuration window, where the detected devices are listed in the upper part. The lower part lists any queues configured so far. If your printer was not detected, configure it manually.

IMPORTANT

If the *Printer* entry is not available in the YaST control center, the `yast2-printer` package probably is not installed. To solve this problem, install the `yast2-printer` package and restart YaST.

Automatic Configuration

YaST is able to configure the printer automatically if the parallel or USB port can be set up automatically and the connected printer can be detected. The printer database must also contain the ID string of the printer that YaST retrieves during the automatic hardware detection. If the hardware ID differs from the model designation, select the model manually.

To make sure that everything works properly, each configuration should be checked with the print test function of YaST. The test page also provides important information about the configuration tested.

Manual Configuration

If the requirements for automatic configuration are not met or if you want a custom setup, configure the printer manually. Depending on how successful the autodetection is and how much information about the printer model is found in the database, YaST may be able to determine the right settings automatically or at least make a reasonable preselection.

The following parameters must be configured:

Hardware Connection (Port)

The configuration of the hardware connection depends on whether YaST has been able to find the printer during hardware autodetection. If YaST is able to detect the printer model automatically, it can be assumed that the printer connection works on the hardware level and no settings need to be changed in this respect. If YaST is unable to autodetect the printer model, there may be some problem with the connection on the hardware level. In this case, some manual intervention is required to configure the connection.

In the *Printer Configuration* dialog, press *Add* to start the manual configuration workflow. Here, select your *Printer Type* (for example `USB printer`) and, with *Next*, enter the *Printer Connection* and select the device.

Name of the Queue

The queue name is used when issuing print commands. The name should be relatively short and consist of lowercase letters and numbers only. Enter the *Name for printing* in the next dialog (*Queue name*).

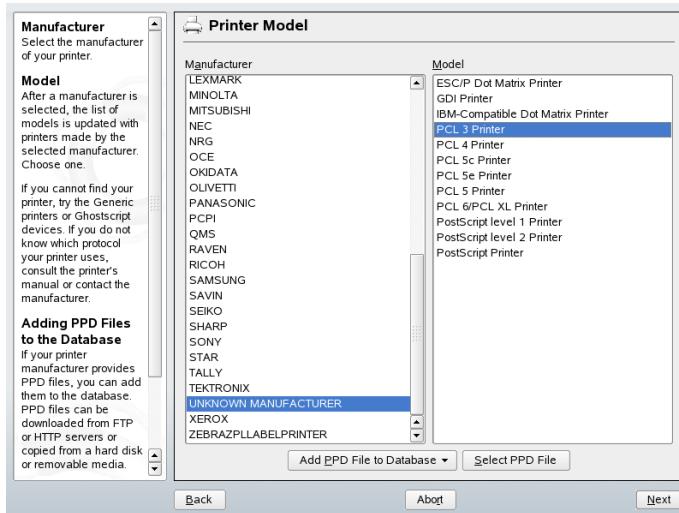
Printer Model and PPD File

All printer-specific parameters, such as the Ghostscript driver to use and the printer filter parameters for the driver, are stored in a PPD (PostScript Printer Description) file. See [Section 11.3, “Installing the Software”](#) (page 229) for more information about PPD files.

For many printer models, several PPD files are available, for example, if several Ghostscript drivers work with the given model. When you select a manufacturer and a model in the next dialog (*Printer model*), YaST selects the PPD file that corresponds to the printer. If several PPD files are available for the model, YaST defaults to one of them (normally the one marked `recommended`). You can change the chosen PPD file in the next dialog with *Edit*.

For non-PostScript models, all printer-specific data is produced by the Ghostscript driver. For this reason, the driver configuration is the single most important factor determining the output quality. The printout is affected both by the kind of Ghostscript driver (PPD file) selected and the options specified for it. If necessary, change additional options (as made available by the PPD file) after selecting *Edit*.

Figure 11.1 Selecting the Printer Model



Always check whether your settings work as expected by printing the test page. If the output is garbled, for example, with several pages almost empty, you should be able to stop the printer by first removing all paper then stopping the test from YaST.

If the printer database does not include an entry for your model, you can either add a new PPD file by selecting *Add PPD File to Database*, or use a collection of generic PPD files to make the printer work with one of the standard printer languages. To do so, select *UNKNOWN MANUFACTURER* as your printer manufacturer.

Advanced Settings

Normally, you do not need to change any of these settings.

11.4.2 Network Printers

A network printer can support various protocols, some of them even concurrently. Although most of the supported protocols are standardized, some manufacturers expand (modify) the standard because they test systems that have not implemented the standard correctly or because they want to provide certain functions that are not available in the standard. Manufacturers then provide drivers for only a few operating systems, eliminating difficulties with those systems. Unfortunately, Linux drivers are rarely provided. The current situation is such that you cannot act on the assumption that every protocol works smoothly in Linux. Therefore, you may have to experiment with various options to achieve a functional configuration.

CUPS supports the `socket`, `LPD`, `IPP`, and `smb` protocols. Here is some detailed information about these protocols:

socket

Socket refers to a connection in which the data is sent to an Internet socket without first performing a data handshake. Some of the socket port numbers that are commonly used are 9100 or 35. An example device URI is

```
socket://host-printer:9100/.
```

LPD (line printer daemon)

The proven LPD protocol is described in RFC 1179. Under this protocol, some job-related data, such as the ID of the printer queue, is sent before the actual print data is sent. Therefore, a printer queue must be specified when configuring the LPD protocol for the data transmission. The implementations of diverse printer manufacturers are flexible enough to accept any name as the printer queue. If necessary, the printer manual should indicate what name to use. LPT, LPT1, LP1, or similar names are often used. An LPD queue can also be configured on a different Linux or Unix host in the CUPS system. The port number for an LPD service is 515. An example device URI is `lpd://host-printer/LPT1`.

IPP (Internet printing protocol)

IPP is a relatively new (1999) protocol based on the HTTP protocol. With IPP, more job-related data is transmitted than with the other protocols. CUPS uses IPP for internal data transmission. This is the preferred protocol for a forwarding queue between two CUPS servers. The name of the print queue is necessary to configure IPP correctly. The port number for IPP is 631. Example device URIs are

```
ipp://host-printer/ps and  
ipp://host-cupsserver/printers/ps.
```

SMB (Windows share)

CUPS also supports printing on printers connected to Windows shares. The protocol used for this purpose is SMB. SMB uses the port numbers 137, 138, and 139.

Example device URIs are

```
smb://user:password@workgroup/server/printer,  
smb://user:password@host/printer, and smb://server/printer.
```

The protocol supported by the printer must be determined before configuration. If the manufacturer does not provide the needed information, the command `nmap`, which comes with the `nmap` package, can be used to guess the protocol. `nmap` checks a host for open ports. For example:

```
nmap -p 35,137-139,515,631,9100-10000 printerIP
```

Configuring CUPS in the Network Using YaST

Network printers should be configured with YaST. YaST facilitates the configuration and is best equipped to handle the security restrictions in CUPS (see [Section 11.6.2, “Changes in the CUPS Print Service” \(page 239\)](#)). For guidelines for installation of CUPS in the network, read the article *CUPS in a Nutshell* in the Support Database at <http://portal.suse.com>.

Start the printer configuration then click *Add*. If not told otherwise by the network administrator try the option *Print Directly to a Network Printer* and proceed according to your local requirements.

Configuring with Command Line Tools

Alternatively, CUPS can be configured with command-line tools like `lpadmin` and `lpoptions`. You need a device URI (uniform resource identifier) consisting of a back-end, such as `usb`, and parameters, like `/dev/usb/lp0`. For example, the full URI could be `parallel:/dev/lp0` (printer connected to the first parallel port) or `usb:/dev/usb/lp0` (first detected printer connected to the USB port).

With `lpadmin`, the CUPS server administrator can add, remove, or manage class and print queues. To add a printer queue use the following syntax:

```
lpadmin -p queue -v device-URI \
-P PPD-file -E
```

Then the device (`-v`) will be available as `queue` (`-p`), using the specified PPD file (`-P`). This means that you must know the PPD file and the name of the device if you want to configure the printer manually.

Do not use `-E` as the first option. For all CUPS commands, `-E` as the first argument sets use of an encrypted connection. To enable the printer, `-E` must be used as shown in the following example:

```
lpadmin -p ps -v parallel:/dev/lp0 -P \
/usr/share/cups/model/Postscript.ppd.gz -E
```

The following example configures a network printer:

```
lpadmin -p ps -v socket://192.168.1.0:9100/ -P \
/usr/share/cups/model/Postscript-level1.ppd.gz -E
```

For more options of `lpadmin`, see the `lpadmin(1)` man page.

During printer setup, certain options are set as default. These options can be modified for every print job (depending on the print tool used). Changing these default options with YaST is also possible. Using command line tools, set default options as follows:

1 First, list all options:

```
lpoptions -p queue -l
```

Example:

```
Resolution/Output Resolution: 150dpi *300dpi 600dpi
```

The activated default option is evident from the preceding asterisk (*).

2 Change the option with `lpadmin`:

```
lpadmin -p queue -o Resolution=600dpi
```

3 Check the new setting:

```
lpoptions -p queue -l
```

```
Resolution/Output Resolution: 150dpi 300dpi *600dpi
```

Settings are written to `~/.lpoptions` when a normal user runs `lpoptions`. root settings are written to `/etc/cups/lpoptions`.

11.5 Configuration for Applications

Applications rely on the existing printer queues in the same way as command line tools do. There is usually no need to reconfigure the printer for a particular application, because you should be able to print from applications using the available queues.

To print from the command line, enter `lp -d queuename filename`, substituting the corresponding names for `queuename` and `filename`.

Some applications rely on the `lp` command for printing. In this case, enter the correct command in the application's print dialog, usually without specifying `filename`, for example, `lp -d queuename`. To make this work with KDE programs, enable *Print through an external program*. Otherwise you cannot enter the print command.

Tools such as xpp and the KDE program kprinter provide a graphical interface for choosing among queues and setting both CUPS standard options and printer-specific options made available through the PPD file. You can use kprinter as the standard printing interface of non-KDE applications by specifying `kprinter` or `kprinter --stdin` as the print command in the print dialogs of these applications. The behavior of the application itself determines which of these two commands to choose. If set up correctly, the application should call the kprinter dialog whenever a print job is issued from it, so you can use the dialog to select a queue and set other printing options. This requires that the application's own print setup does not conflict with that of kprinter and that printing options are only changed through kprinter after it has been enabled.

11.6 Special Features in SUSE Linux

A number of CUPS features have been adapted for SUSE Linux. Some of the most important changes are covered here.

11.6.1 CUPS Server and Firewall

There are several ways to configure CUPS as the client of a network server.

1. For every queue on the network server, you can configure a local queue through which to forward all jobs to the corresponding network server (forwarding queue). Usually, this approach is not recommended, because all client machines must be reconfigured whenever the configuration of the network server changes.
2. Print jobs can also be forwarded directly to one network server. For this type of configuration, do not run a local CUPS daemon. `lp` or corresponding library calls of other programs can send jobs directly to the network server. However, this configuration does not work if you also want to print on a local printer.
3. The CUPS daemon can listen to IPP broadcast packets that other network servers send to announce available queues.

This is the best CUPS configuration for printing over remote CUPS servers. However, there is a risk that an attacker sends IPP broadcasts with queues and the local daemon accesses a counterfeit queue. If it then displays the queue with the same name as another queue on the local server, the owner of the job may believe the job is sent to a local server, while in reality it is sent to the attacker's server.

YaST can find CUPS servers by either scanning local network hosts to see if they offer the IPP service or by listening to IPP broadcasts. This requires the firewall to let incoming packets on port 631/UDP (service IPP client) pass through. This is automatically enabled when you have configured your machine to be in the internal firewall zone. Opening a port to configure access to remote queues in the external zone can be a security risk because an attacker could broadcast a server that might be accepted by users. By default IPP broadcasts are rejected in the external zone. See [Section “Configuring with YaST”](#) (page 106) for details on firewall configuration.

Alternatively, the user can detect CUPS servers by actively scanning the local network hosts or configure all queues manually. However, because of the reasons mentioned in the beginning of this section, this method is not recommended.

11.6.2 Changes in the CUPS Print Service

These changes were initially applied for SUSE Linux 9.1.

cupsd Runs as the User lp

On start-up, cupsd changes from the user `root` to the user `lp`. This provides a much higher level of security, because the CUPS print service does not run with unrestricted permissions, only with the permissions needed for the print service.

However, the authentication (the password check) cannot be performed via `/etc/shadow`, because `lp` has no access to `/etc/shadow`. Instead, the CUPS-specific authentication via `/etc/cups/passwd.md5` must be used. For this purpose, a CUPS administrator with the CUPS administration group `sys` and a CUPS password must be entered in `/etc/cups/passwd.md5`. To do this, enter the following as `root`:

```
lppasswd -g sys -a CUPS-admin-name
```

This setting is also essential if you want to use the CUPS administration Web front-end or the KDE printer administration tool.

When cupsd runs as `lp`, `/etc/printcap` cannot be generated, because `lp` is not permitted to create files in `/etc/`. Therefore, cupsd generates `/etc/cups/printcap`. To ensure that applications that can only read queue names from `/etc/printcap` continue to work properly, `/etc/printcap` is a symbolic link pointing to `/etc/cups/printcap`.

When cupsd runs as `lp`, port 631 cannot be opened. Therefore, cupsd cannot be reloaded with `rccups reload`. Use `rccups restart` instead.

Generalized Functionality for `BrowseAllow` and `BrowseDeny`

The access permissions set for `BrowseAllow` and `BrowseDeny` apply to all kinds of packages sent to cupsd. The default settings in `/etc/cups/cupsd.conf` are as follows:

```
BrowseAllow @LOCAL  
BrowseDeny All
```

and

```
<Location />
  Order Deny,Allow
  Deny From All
  Allow From 127.0.0.1
  Allow From 127.0.0.2
  Allow From @LOCAL
</Location>
```

In this way, only LOCAL hosts can access `cupsd` on a CUPS server. LOCAL hosts are hosts whose IP addresses belong to a non-PPP interface (interfaces whose `IFF_POINTOPOINT` flags are not set) and whose IP addresses belong to the same network as the CUPS server. Packets from all other hosts are rejected immediately.

cupsd Activated by Default

In a standard installation, `cupsd` is activated automatically, enabling comfortable access to the queues of CUPS network servers without any additional manual actions. The items in [Section “cupsd Runs as the User lp”](#) (page 239) and [Section “Generalized Functionality for `BrowseAllow` and `BrowseDeny`”](#) (page 239) are vital preconditions for this feature, because otherwise the security would not be sufficient for an automatic activation of `cupsd`.

11.6.3 PPD Files in Various Packages

The YaST printer configuration sets up the queues for CUPS using only the PPD files installed in `/usr/share/cups/model/` on the system. To find the suitable PPD files for the printer model, YaST compares the vendor and model determined during hardware detection with the vendors and models in all PPD files available in `/usr/share/cups/model/` on the system. For this purpose, the YaST printer configuration generates a database from the vendor and model information extracted from the PPD files. When you select a printer from the list of vendors and models, receive the PPD files matching the vendor and model.

The configuration using only PPD files and no other information sources has the advantage that the PPD files in `/usr/share/cups/model/` can be modified freely. The YaST printer configuration recognizes changes and regenerates the vendor and model database. For example, if you only have PostScript printers, normally you do not need the Foomatic PPD files in the `cups-drivers` package or the Gimp-Print PPD files

in the `cups-drivers-stp` package. Instead, the PPD files for your PostScript printers can be copied directly to `/usr/share/cups/model/` (if they do not already exist in the manufacturer-PPDs package) to achieve an optimum configuration for your printers.

CUPS PPD Files in the `cups` Package

The generic PPD files in the `cups` package have been complemented with adapted Foomatic PPD files for PostScript level 1 and level 2 printers:

- `/usr/share/cups/model/Postscript-level1.ppd.gz`
- `/usr/share/cups/model/Postscript-level2.ppd.gz`

PPD Files in the `cups-drivers` Package

Normally, the Foomatic printer filter `foomatic-rip` is used together with Ghostscript for non-PostScript printers. Suitable Foomatic PPD files have the entries `*NickName: ... Foomatic/Ghostscript driver and *cupsFilter: ... foomatic-rip`. These PPD files are located in the `cups-drivers` package.

YaST prefers a Foomatic PPD file if a Foomatic PPD file with the entry `*NickName: ... Foomatic ... (recommended)` matches the printer model and the manufacturer-PPDs package does not contain a more suitable PPD file.

Gimp-Print PPD Files in the `cups-drivers-stp` Package

Instead of `foomatic-rip`, the CUPS filter `rastertoprinter` from Gimp-Print can be used for many non-PostScript printers. This filter and suitable Gimp-Print PPD files are available in the `cups-drivers-stp` package. The Gimp-Print PPD files are located in `/usr/share/cups/model/stp/` and have the entries `*NickName: ... CUPS+Gimp-Print and *cupsFilter: ... rastertoprinter`.

PPD Files from Printer Manufacturers in the manufacturer-PPDs Package

The manufacturer-PPDs package contains PPD files from printer manufacturers that are released under a sufficiently liberal license. PostScript printers should be configured with the suitable PPD file of the printer manufacturer, because this file enables the use of all functions of the PostScript printer. YaST prefers a PPD file from the manufacturer-PPDs package if the following conditions are met:

- The vendor and model determined during the hardware detection match the vendor and model in a PPD file from the manufacturer-PPDs package.
- The PPD file from the manufacturer-PPDs package is the only suitable PPD file for the printer model or there is a Foomatic PPD file with a *NickName : . . . Foomatic/Postscript (recommended) entry that also matches the printer model.

Accordingly, YaST does not use any PPD file from the manufacturer-PPDs package in the following cases:

- The PPD file from the manufacturer-PPDs package does not match the vendor and model. This may happen if the manufacturer-PPDs package contains only one PPD file for similar models, for example, if there is no separate PPD file for the individual models of a model series, but the model name is specified in a form like Funprinter 1000 series in the PPD file.
- The Foomatic PostScript PPD file is not recommended. This may be because the printer model does not operate efficiently enough in PostScript mode, for example, the printer may be unreliable in this mode because it has too little memory or the printer is too slow because its processor is too weak. Furthermore, the printer may not support PostScript by default, for example, because PostScript support is only available as an optional module.

If a PPD file from the manufacturer-PPDs package is suitable for a PostScript printer, but YaST cannot configure it for these reasons, select the respective printer model manually in YaST.

11.7 Troubleshooting

The following sections cover some of the most frequently encountered printer hardware and software problems and ways to solve or circumvent these problems.

11.7.1 Printers without Standard Printer Language Support

Printers that do not support any common printer language and can only be addressed with special control sequences are called *GDI printers*. These printers only work with the operating system versions for which the manufacturer delivers a driver. *GDI* is a programming interface developed by Microsoft for graphics devices. The actual problem is not the programming interface, but the fact that GDI printers can only be addressed with the proprietary printer language of the respective printer model.

Some printers can be switched to operate either in GDI mode or one of the standard printer languages. Some manufacturers provide proprietary drivers for their GDI printers. The disadvantage of proprietary printer drivers is that there is no guarantee that these work with the installed print system and that they are suitable for the various hardware platforms. In contrast, printers that support a standard printer language do not depend on a special print system version or a special hardware platform.

Instead of spending time trying to make a proprietary Linux driver work, it may be more cost-effective to purchase a supported printer. This would solve the driver problem once and for all, eliminating the need to install and configure special driver software and obtain driver updates that may be required due to new developments in the print system.

11.7.2 No Suitable PPD File Available for a PostScript Printer

If the manufacturer-PPDs package does not contain any suitable PPD file for a PostScript printer, it should be possible to use the PPD file from the driver CD of the printer manufacturer or download a suitable PPD file from the Web page of the printer manufacturer.

If the PPD file is provided as a zip archive (.zip) or a self-extracting zip archive (.exe), unpack it with `unzip`. First, review the license terms of the PPD file. Then use the `cupstestppd` utility to check if the PPD file complies with “Adobe PostScript Printer Description File Format Specification, version 4.3.” If the utility returns “FAIL,” the errors in the PPD files are serious and are likely to cause major problems. The problem spots reported by `cupstestppd` should be eliminated. If necessary, ask the printer manufacturer for a suitable PPD file.

11.7.3 Parallel Ports

The safest approach is to connect the printer directly to the first parallel port and to select the following parallel port settings in the BIOS:

- I/O address: 378 (hexadecimal)
- Interrupt: irrelevant
- Mode: Normal, SPP, or Output Only
- DMA: disabled

If the printer cannot be addressed on the parallel port despite these settings, enter the I/O address explicitly in accordance with the setting in the BIOS in the form `0x378` in `/etc/modprobe.conf`. If there are two parallel ports that are set to the I/O addresses 378 and 278 (hexadecimal), enter these in the form `0x378, 0x278`.

If interrupt 7 is free, it can be activated with the entry shown in [Example 11.1, “/etc/modprobe.conf: Interrupt Mode for the First Parallel Port”](#) (page 244). Before activating the interrupt mode, check the file `/proc/interrupts` to see which interrupts are already in use. Only the interrupts currently being used are displayed. This may change depending on which hardware components are active. The interrupt for the parallel port must not be used by any other device. If you are not sure, use the polling mode with `irq=none`.

Example 11.1 /etc/modprobe.conf: Interrupt Mode for the First Parallel Port

```
alias parport_lowlevel parport_pc
options parport_pc io=0x378 irq=7
```

11.7.4 Network Printer Connections

Identifying Network Problems

Connect the printer directly to the computer. For test purposes, configure the printer as a local printer. If this works, the problems are related to the network.

Checking the TCP/IP Network

The TCP/IP network and name resolution must be functional.

Checking a Remote lpd

Use the following command to test if a TCP connection can be established to lpd (port 515) on *host*:

```
netcat -z host 515 && echo ok || echo failed
```

If the connection to lpd cannot be established, lpd may not be active or there may be basic network problems.

As the user `root`, use the following command to query a (possibly very long) status report for *queue* on remote *host*, provided the respective lpd is active and the host accepts queries:

```
echo -e "\004queue" \
| netcat -w 2 -p 722 host 515
```

If lpd does not respond, it may not be active or there may be basic network problems. If lpd responds, the response should show why printing is not possible on the queue on *host*. If you receive a response like that in [Example 11.2, “Error Message from the lpd”](#) (page 245), the problem is caused by the remote lpd.

Example 11.2 Error Message from the lpd

```
lpd: your host does not have line printer access
lpd: queue does not exist
printer: spooling disabled
printer: printing disabled
```

Checking a Remote cupsd

By default, the CUPS network server should broadcast its queues every 30 seconds on UDP port 631. Accordingly, the following command can be used to test whether there is a CUPS network server in the network.

```
netcat -u -l -p 631 & PID=$! ; sleep 40 ; kill $PID
```

If a broadcasting CUPS network server exists, the output appears as shown in [Example 11.3, “Broadcast from the CUPS Network Server”](#) (page 246).

Example 11.3 Broadcast from the CUPS Network Server

```
ipp://host.domain:631/printers/queue
```

The following command can be used to test if a TCP connection can be established to `cupsd` (port 631) on `host`:

```
netcat -z host 631 && echo ok || echo failed
```

If the connection to `cupsd` cannot be established, `cupsd` may not be active or there may be basic network problems. `lpstat -h host -l -t` returns a (possibly very long) status report for all queues on `host`, provided the respective `cupsd` is active and the host accepts queries.

The next command can be used to test if the `queue` on `host` accepts a print job consisting of a single carriage-return character. Nothing should be printed. Possibly, a blank page may be ejected.

```
echo -en "\r" \
| lp -d queue -h host
```

Troubleshooting a Network Printer or Print Server Box

Spoolers running in a print server box sometimes cause problems when they have to deal with a lot of print jobs. Because this is caused by the spooler in the print server box, there is nothing you can do about it. As a work-around, circumvent the spooler in the print server box by addressing the printer connected to the print server box directly via TCP socket. See [Section 11.4.2, “Network Printers”](#) (page 234).

In this way, the print server box is reduced to a converter between the various forms of data transfer (TCP/IP network and local printer connection). To use this method, you need to know the TCP port on the print server box. If the printer is connected to the print server box and powered on, this TCP port can usually be determined with the `nmap` utility from the `nmap` package some time after the print server box is powered on. For example, `nmap IP-address` may deliver the following output for a print server box:

Port	State	Service
23/tcp	open	telnet
80/tcp	open	http
515/tcp	open	printer

```
631/tcp      open        cups
9100/tcp     open        jetdirect
```

This output indicates that the printer connected to the print server box can be addressed via TCP socket on port 9100. By default, nmap only checks a number of commonly known ports listed in /usr/share/nmap/nmap-services. To check all possible ports, use the command nmap -p *from_port-to_port* *IP-address*. This may take some time. For further information, refer to the nmap man page.

Enter a command like

```
echo -en "\rHello\r\f" | netcat -w 1 IP-address port
cat file | netcat -w 1 IP-address port
```

to send character strings or files directly to the respective port to test if the printer can be addressed on this port.

11.7.5 Defective Printouts without Error Message

For the print system, the print job is completed when the CUPS back-end completes the data transfer to the recipient (printer). If the further processing on the recipient fails, for example, if the printer is not able to print the printer-specific data, the print system does not notice this. If the printer is not able to print the printer-specific data, select a different PPD file that is more suitable for the printer.

11.7.6 Disabled Queues

If the data transfer to the recipient fails entirely after several attempts, the CUPS back-end, such as `usb` or `socket`, reports an error to the print system (to `cupsd`). The back-end decides whether and how many attempts make sense until the data transfer is reported as impossible. Because further attempts would be in vain, `cupsd` disables printing for the respective queue. After eliminating the cause of the problem, the system administrator must reenable printing with the command `/usr/bin/enable`.

11.7.7 CUPS Browsing: Deleting Print Jobs

If a CUPS network server broadcasts its queues to the client hosts via browsing and a suitable local `cupsd` is active on the client hosts, the client `cupsd` accepts print jobs from applications and forwards them to the `cupsd` on the server. When `cupsd` accepts a print job, it is assigned a new job number. Therefore, the job number on the client host is different from the job number on the server. Because a print job is usually forwarded immediately, it cannot be deleted with the job number on the client host, because the client `cupsd` regards the print job as completed as soon as it has been forwarded to the server `cupsd`.

To delete the print job on the server, use a command such as `lpstat -h print-server -o` to determine the job number on the server, provided the server has not already completed the print job (that is, sent it to the printer). Using this job number, the print job on the server can be deleted:

```
cancel -h print-server queue-jobnumber
```

11.7.8 Defective Print Jobs and Data Transfer Errors

Print jobs remain in the queues and printing resumes if you switch the printer off and on or shut down and reboot the computer during the printing process. Defective print jobs must be removed from the queue with `cancel`.

If a print job is defective or an error occurs in the communication between the host and the printer, the printer prints numerous sheets of paper with unintelligible characters, because it is unable to process the data correctly. To deal with this, follow these steps:

- 1** To stop printing, remove all paper from ink jet printers or open the paper trays of laser printers. High-quality printers have a button for canceling the current printout.
- 2** The print job may still be in the queue, because jobs are only removed after they are sent completely to the printer. Use `lpstat -o` or `lpstat -h print-server -o` to check which queue is currently printing. Delete the print job with `cancel queue-jobnumber` or `cancel -h print-server queue-jobnumber`.

- 3** Some data may still be transferred to the printer even though the print job has been deleted from the queue. Check if a CUPS back-end process is still running for the respective queue and terminate it. For example, for a printer connected to the parallel port, the command `fuser -k /dev/lp0` can be used to terminate all processes that are still accessing the printer (more precisely: the parallel port).
- 4** Reset the printer completely by switching it off for some time. Then insert the paper and turn on the printer.

11.7.9 Debugging the CUPS Print System

Use the following generic procedure to locate problems in the CUPS print system:

- 1** Set `LogLevel debug` in `/etc/cups/cupsd.conf`.
- 2** Stop `cupsd`.
- 3** Remove `/var/log/cups/error_log*` to avoid having to search through very large log files.
- 4** Start `cupsd`.
- 5** Repeat the action that led to the problem.
- 6** Check the messages in `/var/log/cups/error_log*` to identify the cause of the problem.

11.7.10 For More Information

Solutions to many specific problems are presented in the SUSE Support Database (<http://portal.suse.com/>). Locate the relevant articles with keyword searches.

12

Dynamic Kernel Device Management with udev

Since version 2.6, the kernel is capable of adding or removing almost any device in the running system. Changes in device state (whether a device is plugged in or removed) need to be propagated to userspace. Devices need to be configured as soon as they are plugged in and discovered. Users of a certain device need to be informed about any state changes of this device. udev provides the needed infrastructure to dynamically maintain the device node files and symbolic links in the `/dev` directory. udev rules provide a way to plug external tools into the kernel device event processing. This enables you to customize udev device handling, for example, by adding certain scripts to execute as part of kernel device handling, or request and import additional data to evaluate during device handling.

12.1 The `/dev` Directory

The device nodes in the `/dev` directory provide access to the corresponding kernel devices. With udev, the `/dev` directory reflects the current state of the kernel. Every kernel device has one corresponding device file. If a device is disconnected from the system, the device node is removed.

The content of the `/dev` directory is kept on a temporary file system and all files are created from scratch at every system start-up. Manually created or changed files intentionally do not survive a reboot. Static files and directories that should always be present in the `/dev` directory regardless of the state of the corresponding kernel device can be placed in the `/lib/udev/devices` directory. At system start-up, the contents of that directory is copied to the `/dev` directory with the same ownership and permissions as the files in `/lib/udev/devices`.

12.2 Kernel uevents and udev

The required device information is exported by the sysfs file system. For every device the kernel has detected and initialized, a directory with the device name is created. It contains attribute files with device-specific properties. Every time a device is added or removed, the kernel sends a uevent to notify udev of the change.

The udev daemon reads and parses all provided rules from the `/etc/udev/rules.d/*.rules` files once at start-up and keeps them in memory. If rules files are changed, added, or removed, the daemon receives an event and updates the in-memory representation of the rules.

Every received event is matched against the set of provides rules. The rules can add or change event environment keys, request a specific name for the device node to create, add symlinks pointing to the node, or add programs to run after the device node is created. The driver core uevents are received from a kernel netlink socket.

12.3 Drivers, Kernel Modules, and Devices

The kernel bus drivers probe for devices. For every detected device, the kernel creates an internal device structure and the driver core sends a uevent to the udev daemon. Bus devices identify themselves by a specially-formatted ID, which tells what kind of device it is. Usually these IDs consist of vendor and product ID and other subsystem-specific values. Every bus has its own scheme for these IDs, called MODALIAS. The kernel takes the device information, composes a MODALIAS ID string from it, and sends that string along with the event. For a USB mouse, it looks like this:

```
MODALIAS=usb:v046DpC03Ed2000dc00dsc00dp00ic03isc01ip02
```

Every device driver carries a list of known aliases for devices it can handle. The list is contained in the kernel module file itself. The program depmod reads the ID lists and creates the file `modules.alias` in the kernel's `/lib/modules` directory for all currently available modules. With this infrastructure, module loading is as easy as calling modprobe for every event that carries a MODALIAS key. If modprobe `$MODALIAS` is called, it matches the device alias composed for the device with the

aliases provided by the modules. If a matching entry is found, that module is loaded. All this is triggered by udev and happens automatically.

12.4 Booting and Initial Device Setup

All device events happening during the boot process before the udev daemon is running are lost, because the infrastructure to handle these events lives on the root file system and is not available at that time. To cover that loss, the kernel provides a `uevent` file for every device in the sysfs file system. By writing `add` to that file, the kernel resends the same event as the one lost during boot. A simple loop over all `uevent` files in `/sys` triggers all events again to create the device nodes and perform device setup.

As an example, a USB mouse present during boot may not be initialized by the early boot logic, because the driver is not available that time. The event for the device discovery was lost and failed to find a kernel module for the device. Instead of manually searching for possibly connected devices, udev just requests all device events from the kernel after the root file system is available, so the event for the USB mouse device just runs again. Now it finds the kernel module on the mounted root file system and the USB mouse can be initialized.

From userspace, there is no visible difference between a device coldplug sequence and a device discovery during runtime. In both cases, the same rules are used to match and the same configured programs are run.

12.5 Debugging udev Events

The program `udevmonitor` can be used to visualize the driver core events and the timing of the udev event processes.

```
UEVENT [1132632714.285362] add@/devices/pci0000:00/0000:00:1d.1/usb2/2-2
UEVENT [1132632714.288166] add@/devices/pci0000:00/0000:00:1d.1/usb2/2-2/2-2:1.0
UEVENT [1132632714.309485] add@/class/input/input6
UEVENT [1132632714.309511] add@/class/input/input6/mouse2
UEVENT [1132632714.309524] add@/class/usb_device/usbdev2.12
UDEV   [1132632714.348966] add@/devices/pci0000:00/0000:00:1d.1/usb2/2-2
UDEV   [1132632714.420947] add@/devices/pci0000:00/0000:00:1d.1/usb2/2-2/2-2:1.0
UDEV   [1132632714.427298] add@/class/input/input6
UDEV   [1132632714.434223] add@/class/usb_device/usbdev2.12
UDEV   [1132632714.439934] add@/class/input/input6/mouse2
```

The UEVENT lines show the events the kernel has sent over netlink. The UDEV lines show the finished udev event handlers. The timing is printed in microseconds. The time between UEVENT and UDEV is the time udev took to process this event or the udev daemon has delayed its execution to synchronize this event with related and already running events. For example, events for hard disk partitions always wait for the main disk device event to finish, because the partition events may rely on the data the main disk event has queried from the hardware.

`udevmonitor --env` shows the complete event environment:

```
UDEV [1132633002.937243] add@/class/input/input7
UDEV_LOG=3
ACTION=add
DEVPATH=/class/input/input7
SUBSYSTEM=input
SEQNUM=1043
PHYSDEVPATH=/devices/pci0000:00/0000:00:1d.1/usb2/2-2/2-2:1.0
PHYSDEVBUS=usb
PHYSDEVDRIVER=usbhid
PRODUCT=3/46d/c03e/2000
NAME="Logitech USB-PS/2 Optical Mouse"
PHYS="usb-0000:00:1d.1-2/input0"
UNIQ=""
EV=7
KEY=70000 0 0 0 0 0 0 0 0
REL=103
```

udev also sends messages to syslog. The default syslog priority that controls which messages are sent to syslog is specified in the udev configuration file `/etc/udev/udev.conf`. The log priority of the running daemon can be changed with `udevcontrol log_priority=level/number`.

12.6 Influencing Kernel Device Event Handling with udev Rules

A udev rule can match with any property the kernel adds to the event itself or any information that the kernel exports to `sysfs`. The rule can also request additional information from external programs. Every event is matched against all provided rules. All rules are located in the `/etc/udev/rules.d` directory.

Every line in the rules file contains at least one key value pair. There are two kinds of keys, match and assignment keys. If all match keys match their values, the rule is applied

and the assignment keys are assigned the specified value. A matching rule may specify the name of the device node, add symlinks pointing to the node, or run a specified program as part of the event handling. If no matching rule is found, the default device node name is used to create the device node. The rule syntax and the provided keys to match or import data are described in the udev man page.

12.7 Persistent Device Naming

The dynamic device directory and the udev rules infrastructure make it possible to provide stable names for all disk devices—regardless of their order of recognition or the connection used to plug the device. Every appropriate block device the kernel creates is examined by tools with special knowledge about certain buses, drive types, or file systems. Along with the dynamic kernel-provided device node name, udev maintains classes of persistent symbolic links pointing to the device:

```
/dev/disk
|-- by-id
|   |-- scsi-SATA_HTS726060M9AT00_MRH453M4HWHG7B -> ../../sda
|   |-- scsi-SATA_HTS726060M9AT00_MRH453M4HWHG7B-part1 -> ../../sda1
|   |-- scsi-SATA_HTS726060M9AT00_MRH453M4HWHG7B-part6 -> ../../sda6
|   |-- scsi-SATA_HTS726060M9AT00_MRH453M4HWHG7B-part7 -> ../../sda7
|   |-- usb-Generic_STORAGE_DEVICE_02773 -> ../../sdd
|   `-- usb-Generic_STORAGE_DEVICE_02773-part1 -> ../../sdd1
|-- by-label
|   |-- Photos -> ../../sdd1
|   |-- SUSE10 -> ../../sda7
|   `-- devel -> ../../sda6
|-- by-path
|   |-- pci-0000:00:1f.2-scsi-0:0:0:0 -> ../../sda
|   |-- pci-0000:00:1f.2-scsi-0:0:0:0-part1 -> ../../sda1
|   |-- pci-0000:00:1f.2-scsi-0:0:0:0-part6 -> ../../sda6
|   |-- pci-0000:00:1f.2-scsi-0:0:0:0-part7 -> ../../sda7
|   |-- pci-0000:00:1f.2-scsi-1:0:0:0 -> ../../sr0
|   |-- usb-02773:0:0:2 -> ../../sdd
|   |-- usb-02773:0:0:2-part1 -> ../../sdd1
`-- by-uuid
    |-- 159a47a4-e6e6-40be-a757-a629991479ae -> ../../sda7
    |-- 3e999973-00c9-4917-9442-b7633bd95b9e -> ../../sda6
    '-- 4210-8F8C -> ../../sdd1
```

12.8 The Replaced hotplug Package

The formerly used hotplug package is entirely replaced by the udev and the udev-related kernel infrastructure. The following parts of the former hotplug infrastructure have been made obsolete or had their functionality taken over by udev:

/etc/hotplug/*.agent
No longer needed or moved to /lib/udev

/etc/hotplug/*.rc
Replaced by the /sys/*/uevent trigger

/etc/hotplug/blacklist
Replaced by the blacklist option in modprobe.conf

/etc/dev.d/*
Replaced by the udev rule RUN key

/etc/hotplug.d/*
Replaced by the udev rule RUN key

/sbin/hotplug
Replaced by udevd listening to netlink; only used in the initial RAM file system until the root file system can be mounted, then it is disabled

/dev/*
Replaced by dynamic udev and static content in /lib/udev/devices/*

The following files and directories contain the crucial elements of the udev infrastructure:

/etc/udev/udev.conf
Main udev configuration file

/etc/udev/rules.d/*
udev event matching rules

/lib/udev/devices/*
Static /dev content

```
/lib/udev/*
    Helper programs called from udev rules
```

12.9 For More Information

For more information about the udev infrastructure, refer to the following man pages:

udev

General information about udev, keys, rules, and other important configuration issues.

udevinfo

udevinfo can be used to query device information from the udev database.

udevd

Information about the udev event managing daemon.

udevmonitor

udevmonitor prints the kernel and udev event sequence to the console. This tool is mainly used for debugging purposes.

File Systems in Linux

Linux supports a number of different file systems. This chapter presents a brief overview of the most popular Linux file systems, elaborating on their design concepts, advantages, and fields of application. Some additional information about LFS (large file support) in Linux is also provided.

13.1 Terminology

metadata

A file system–internal data structure that assures all the data on disk is properly organized and accessible. Essentially, it is “data about the data.” Almost every file system has its own structure of metadata, which is part of why the file systems show different performance characteristics. It is extremely important to maintain metadata intact, because otherwise all data on the file system could become inaccessible.

inode

Inodes contain various information about a file, including size, number of links, pointers to the disk blocks where the file contents are actually stored, and date and time of creation, modification, and access.

journal

In the context of a file system, a journal is an on-disk structure containing a kind of log in which the file system stores what it is about to change in the file system's metadata. Journaling greatly reduces the recovery time of a Linux system because it obsoletes the lengthy search process that checks the entire file system at system start-up. Instead, only the journal is replayed.

13.2 Major File Systems in Linux

Unlike two or three years ago, choosing a file system for a Linux system is no longer a matter of a few seconds (Ext2 or ReiserFS?). Kernels starting from 2.4 offer a variety of file systems from which to choose. The following is an overview of how these file systems basically work and which advantages they offer.

It is very important to bear in mind that there may be no file system that best suits all kinds of applications. Each file system has its particular strengths and weaknesses, which must be taken into account. Even the most sophisticated file system cannot replace a reasonable backup strategy, however.

The terms *data integrity* and *data consistency*, when used in this chapter, do not refer to the consistency of the user space data (the data your application writes to its files). Whether this data is consistent must be controlled by the application itself.

IMPORTANT: Setting Up File Systems

Unless stated otherwise in this chapter, all the steps required to set up or change partitions and file systems can be performed using YaST.

13.2.1 ReiserFS

Officially one of the key features of the 2.4 kernel release, ReiserFS has been available as a kernel patch for 2.2.x SUSE kernels since SUSE Linux version 6.4. ReiserFS was designed by Hans Reiser and the Namesys development team. It has proven itself to be a powerful alternative to Ext2. Its key assets are better disk space utilization, better disk access performance, and faster crash recovery.

ReiserFS's strengths, in more detail, are:

Better Disk Space Utilization

In ReiserFS, all data is organized in a structure called B^* -balanced tree. The tree structure contributes to better disk space utilization because small files can be stored directly in the B^* tree leaf nodes instead of being stored elsewhere and just maintaining a pointer to the actual disk location. In addition to that, storage is not allocated in chunks of 1 or 4 kB, but in portions of the exact size needed. Another benefit lies in the dynamic allocation of inodes. This keeps the file system more flexible than traditional file systems, like Ext2, where the inode density must be specified at file system creation time.

Better Disk Access Performance

For small files, file data and “stat_data” (inode) information are often stored next to each other. They can be read with a single disk I/O operation, meaning that only one access to disk is required to retrieve all the information needed.

Fast Crash Recovery

Using a journal to keep track of recent metadata changes makes a file system check a matter of seconds, even for huge file systems.

Reliability through Data Journaling

ReiserFS also supports data journaling and ordered data modes similar to the concepts outlined in the Ext3 section, [Section 13.2.3, “Ext3”](#) (page 262). The default mode is `data=ordered`, which ensures both data and metadata integrity, but uses journaling only for metadata.

13.2.2 Ext2

The origins of Ext2 go back to the early days of Linux history. Its predecessor, the Extended File System, was implemented in April 1992 and integrated in Linux 0.96c. The Extended File System underwent a number of modifications and, as Ext2, became the most popular Linux file system for years. With the creation of journaling file systems and their astonishingly short recovery times, Ext2 became less important.

A brief summary of Ext2's strengths might help understand why it was—and in some areas still is—the favorite Linux file system of many Linux users.

Solidity

Being quite an “old-timer,” Ext2 underwent many improvements and was heavily tested. This may be the reason why people often refer to it as rock-solid. After a system outage when the file system could not be cleanly unmounted, e2fsck starts to analyze the file system data. Metadata is brought into a consistent state and pending files or data blocks are written to a designated directory (called `lost +found`). In contrast to journaling file systems, e2fsck analyzes the entire file system and not just the recently modified bits of metadata. This takes significantly longer than checking the log data of a journaling file system. Depending on file system size, this procedure can take half an hour or more. Therefore, it is not desirable to choose Ext2 for any server that needs high availability. However, because Ext2 does not maintain a journal and uses significantly less memory, it is sometimes faster than other file systems.

Easy Upgradability

The code for Ext2 is the strong foundation on which Ext3 could become a highly-acclaimed next-generation file system. Its reliability and solidity were elegantly combined with the advantages of a journaling file system.

13.2.3 Ext3

Ext3 was designed by Stephen Tweedie. Unlike all other next-generation file systems, Ext3 does not follow a completely new design principle. It is based on Ext2. These two file systems are very closely related to each other. An Ext3 file system can be easily built on top of an Ext2 file system. The most important difference between Ext2 and Ext3 is that Ext3 supports journaling. In summary, Ext3 has three major advantages to offer:

Easy and Highly Reliable Upgrades from Ext2

Because Ext3 is based on the Ext2 code and shares its on-disk format as well as its metadata format, upgrades from Ext2 to Ext3 are incredibly easy. Unlike transitions to other journaling file systems, such as ReiserFS or XFS, which can be quite tedious (making backups of the entire file system and recreating it from scratch), a transition to Ext3 is a matter of minutes. It is also very safe, because recreating an entire file system from scratch might not work flawlessly. Considering the number of existing Ext2 systems that await an upgrade to a journaling file system, you can easily figure out why Ext3 might be of some importance to many system administrators.

Downgrading from Ext3 to Ext2 is as easy as the upgrade. Just perform a clean unmount of the Ext3 file system and remount it as an Ext2 file system.

Reliability and Performance

Some other journaling file systems follow the “metadata-only” journaling approach. This means your metadata is always kept in a consistent state, but the same cannot be automatically guaranteed for the file system data itself. Ext3 is designed to take care of both metadata and data. The degree of “care” can be customized. Enabling Ext3 in the `data=journal` mode offers maximum security (data integrity), but can slow down the system because both metadata and data are journaled. A relatively new approach is to use the `data=ordered` mode, which ensures both data and metadata integrity, but uses journaling only for metadata. The file system driver collects all data blocks that correspond to one metadata update. These data blocks are written to disk before the metadata is updated. As a result, consistency is achieved for metadata and data without sacrificing performance. A third option to use is `data=writeback`, which allows data to be written into the main file system after its metadata has been committed to the journal. This option is often considered the best in performance. It can, however, allow old data to reappear in files after crash and recovery while internal file system integrity is maintained. Unless you specify something else, Ext3 is run with the `data=ordered` default.

13.2.4 Converting an Ext2 File System into Ext3

To convert an Ext2 file system to Ext3, proceed as follows:

- 1 Create an Ext3 journal by running `tune2fs -j` as root. This creates an Ext3 journal with the default parameters.

To decide yourself how large the journal should be and on which device it should reside, run `tune2fs -J` instead together with the desired journal options `size=` and `device=`. More information about the `tune2fs` program is available in the `tune2fs` manual page.

- 2 To ensure that the Ext3 file system is recognized as such, edit the file `/etc/fstab` as root, changing the file system type specified for the corresponding partition from `ext2` to `ext3`. The change takes effect after the next reboot.
- 3 To boot a root file system set up as an Ext3 partition, include the modules `ext3` and `jbd` in the `initrd`. To do this, edit `/etc/sysconfig/kernel` as root, adding `ext3` and `jbd` to the `INITRD_MODULES` variable. After saving the

changes, run the `mkinitrd` command. This builds a new initrd and prepares it for use.

13.2.5 Reiser4

Right after kernel 2.6 had been released, the family of journaling file systems was joined by another member: Reiser4. Reiser4 is fundamentally different from its predecessor ReiserFS (version 3.6). It introduces the concept of plug-ins to tweak the file system functionality and a finer grained security concept.

Fine Grained Security Concept

In designing Reiser4, its developers put an emphasis on the implementation of security-relevant features. Reiser4 therefore comes with a set of dedicated security plug-ins. The most important one introduces the concept of file “items.” Currently, file access controls are defined per file. If there is a large file containing information relevant to several users, groups, or applications, the access rights had be fairly imprecise to include all parties involved. In Reiser4, you can split those files into smaller portions (the “items”). Access rights can then be set for each item and each user separately, allowing a much more precise file security management. A perfect example would be `/etc/passwd`. To date, only `root` can read and edit the file while non-`root` users only get read access to this file. Using the item concept of Reiser4, you could split this file in a set of items (one item per user) and allow users or applications to modify their own data but not access other users' data. This concept adds both to security and flexibility.

Extensibility through Plug-Ins

Many file system functions and external functions normally used by a file system are implemented as plug-ins in Reiser4. These plug-ins can easily be added to the base system. You no longer need to recompile the kernel or reformat the hard disk to add new functionalities to your file system.

Better File System Layout through Delayed Allocation

Like XFS, Reiser4 supports delayed allocation. See [Section 13.2.6, “XFS”](#) (page 265). Using delayed allocation even for metadata can result in better overall layout.

13.2.6 XFS

Originally intended as the file system for their IRIX OS, SGI started XFS development in the early 1990s. The idea behind XFS was to create a high-performance 64-bit journaling file system to meet the extreme computing challenges of today. XFS is very good at manipulating large files and performs well on high-end hardware. However, even XFS has a drawback. Like ReiserFS, XFS takes great care of metadata integrity, but less of data integrity.

A quick review of XFS's key features explains why it may prove a strong competitor for other journaling file systems in high-end computing.

High Scalability through the Use of Allocation Groups

At the creation time of an XFS file system, the block device underlying the file system is divided into eight or more linear regions of equal size. Those are referred to as *allocation groups*. Each allocation group manages its own inodes and free disk space. Practically, allocation groups can be seen as file systems in a file system. Because allocation groups are rather independent of each other, more than one of them can be addressed by the kernel simultaneously. This feature is the key to XFS's great scalability. Naturally, the concept of independent allocation groups suits the needs of multiprocessor systems.

High Performance through Efficient Management of Disk Space

Free space and inodes are handled by B^+ trees inside the allocation groups. The use of B^+ trees greatly contributes to XFS's performance and scalability. XFS uses *delayed allocation*. It handles allocation by breaking the process into two pieces. A pending transaction is stored in RAM and the appropriate amount of space is reserved. XFS still does not decide where exactly (speaking of file system blocks) the data should be stored. This decision is delayed until the last possible moment. Some short-lived temporary data may never make its way to disk, because it may be obsolete by the time XFS decides where actually to save it. Thus XFS increases write performance and reduces file system fragmentation. Because delayed allocation results in less frequent write events than in other file systems, it is likely that data loss after a crash during a write is more severe.

Preallocation to Avoid File System Fragmentation

Before writing the data to the file system, XFS *reserves* (preallocates) the free space needed for a file. Thus, file system fragmentation is greatly reduced. Performance is increased because the contents of a file are not distributed all over the file system.

13.3 Some Other Supported File Systems

Table 13.1, “File System Types in Linux” (page 266) summarizes some other file systems supported by Linux. They are supported mainly to ensure compatibility and interchange of data with different kinds of media or foreign operating systems.

Table 13.1 File System Types in Linux

cramfs	<i>Compressed ROM file system:</i> A compressed read-only file system for ROMs.
hpfs	<i>High Performance File System:</i> The IBM OS/2 standard file system—only supported in read-only mode.
iso9660	Standard file system on CD-ROMs.
minix	This file system originated from academic projects on operating systems and was the first file system used in Linux. Today, it is used as a file system for floppy disks.
msdos	<i>fat</i> , the file system originally used by DOS, is today used by various operating systems.
ncpfs	File system for mounting Novell volumes over networks.
nfs	<i>Network File System:</i> Here, data can be stored on any machine in a network and access may be granted via a network.
smbfs	<i>Server Message Block</i> is used by products such as Windows to enable file access over a network.
sysv	Used on SCO UNIX, Xenix, and Coherent (commercial UNIX systems for PCs).
ufs	Used by BSD, SunOS, and NeXTstep. Only supported in read-only mode.

umsdos	<i>UNIX on MSDOS</i> : Applied on top of a normal <code>fat</code> file system, achieves UNIX functionality (permissions, links, long filenames) by creating special files.
vfat	<i>Virtual FAT</i> : Extension of the <code>fat</code> file system (supports long filenames).
ntfs	<i>Windows NT file system</i> , read-only.

13.4 Large File Support in Linux

Originally, Linux supported a maximum file size of 2 GB. This was enough before the explosion of multimedia and as long as no one tried to manipulate huge databases on Linux. Becoming more and more important for server computing, the kernel and C library were modified to support file sizes larger than 2 GB when using a new set of interfaces that applications must use. Today, almost all major file systems offer LFS support, allowing you to perform high-end computing. [Table 13.2, “Maximum Sizes of File Systems \(On-Disk Format\)”](#) (page 267) offers an overview of the current limitations of Linux files and file systems.

Table 13.2 Maximum Sizes of File Systems (On-Disk Format)

File System	File Size (Bytes)	File System Size (Bytes)
Ext2 or Ext3 (1 kB block size)	2^{34} (16 GB)	2^{41} (2 TB)
Ext2 or Ext3 (2 kB block size)	2^{38} (256 GB)	2^{43} (8 TB)
Ext2 or Ext3 (4 kB block size)	2^{41} (2 TB)	2^{43} -4096 (16 TB-4096 Bytes)
Ext2 or Ext3 (8 kB block size) (systems with 8 kB pages, like Alpha)	2^{46} (64 TB)	2^{45} (32 TB)
ReiserFS v3	2^{46} (64 TB)	2^{45} (32 TB)

File System	File Size (Bytes)	File System Size (Bytes)
XFS	2^{63} (8 EB)	2^{63} (8 EB)
NFSv2 (client side)	2^{31} (2 GB)	2^{63} (8 EB)
NFSv3 (client side)	2^{63} (8 EB)	2^{63} (8 EB)

IMPORTANT: Linux Kernel Limits

Table 13.2, “Maximum Sizes of File Systems (On-Disk Format)” (page 267) describes the limitations regarding the on-disk format. The 2.6 kernel imposes its own limits on the size of files and file systems handled by it. These are as follows:

File Size

On 32-bit systems, files may not exceed the size of 2 TB (2^{41} bytes).

File System Size

File systems may be up to 2^{73} bytes large. However, this limit is still out of reach for the currently available hardware.

13.5 For More Information

Each of the file system projects described above maintains its own home page on which to find mailing list information, further documentation, and FAQs.

- <http://e2fsprogs.sourceforge.net/>
- <http://www.zipworld.com.au/~akpm/linux/ext3/>
- <http://www.namesys.com/>
- <http://oss.software.ibm.com/developerworksopensource/jfs/>
- <http://oss.sgi.com/projects/xfs/>

A comprehensive multipart tutorial about Linux file systems can be found at *IBM developerWorks*: <http://www-106.ibm.com/developerworks/library/l-fs.html>. For a comparison of the different journaling file systems in Linux, look at Juan I. Santos Florido's article at *Linuxgazette*: <http://www.linuxgazette.com/issue55/florido.html>. Those interested in an in-depth analysis of LFS in Linux should try Andreas Jaeger's LFS site: http://www.suse.de/~aj/linux_lfs.html.

14

The X Window System

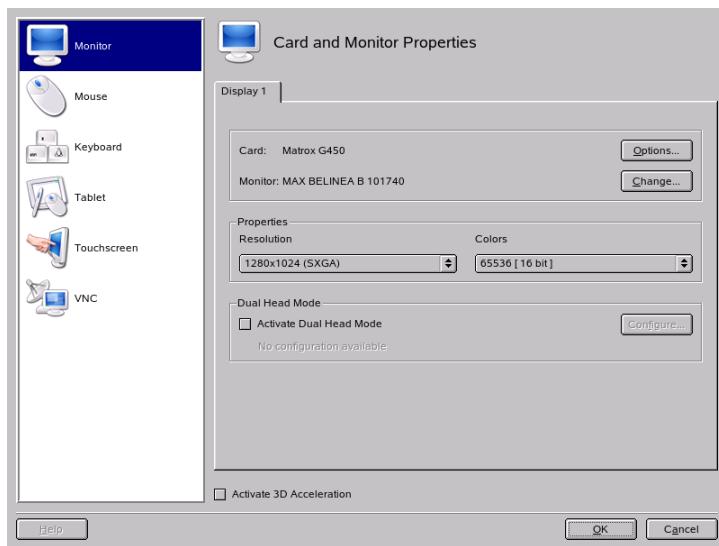
The X Window System (X11) is the de facto standard for graphical user interfaces in UNIX. X is network-based, enabling applications started on one host to be displayed on another host connected over any kind of network (LAN or Internet). This chapter describes the setup and optimization of the X Window System environment, provides background information about the use of fonts in SUSE Linux, and explains the configuration of OpenGL and 3D.

The following text contains several references to documentation that can be found below `/usr/share/doc/packages/Xorg` and `/usr/share/doc/howto/en`. This material along with respective manual pages is available only if the appropriate documentation packages are installed (`xorg-x11-doc`, `xorg-x11-man`, and `howtoen`).

14.1 X11 Setup with SaX2

The graphical user interface, or X server, handles the communication between hardware and software. Desktops, like KDE and GNOME, and the wide variety of window managers, use the X server for interaction with the user. The graphical user interface is initially configured during installation. To change the settings afterwards, use the respective module from the YaST control center or run SaX2 manually from the command line with the command `sax2`. The SaX2 main window provides a common interface for the individual modules from the YaST control center.

Figure 14.1 The Main Window of SaX2



In the left navigation bar, there are six items, each of them showing the respective configuration dialog from the YaST control center. Find the sections mentioned below in Chapter *System Configuration with YaST* (↑Start-Up).

Monitor

For a description of the monitor and graphics card configuration, see Section “Card and Monitor Properties” (Chapter 2, *System Configuration with YaST*, ↑Start-Up).

Mouse

For a description of the mouse configuration in the graphical environment, see Section “Mouse Properties” (Chapter 2, *System Configuration with YaST*, ↑Start-Up).

Keyboard

For a description of the keyboard configuration in the graphical environment, see Section “Keyboard Properties” (Chapter 2, *System Configuration with YaST*, ↑Start-Up).

Tablet

For a description of the graphics tablet configuration, see Section “Tablet Properties” (Chapter 2, *System Configuration with YaST*, ↑Start-Up).

Touchscreen

For a description of the touchscreen configuration, see Section “Touchscreen Properties” (Chapter 2, *System Configuration with YaST*, ↑Start-Up).

VNC

For a description of the VNC configuration, see Section “Remote Access Properties” (Chapter 2, *System Configuration with YaST*, ↑Start-Up).

14.2 Optimizing the X Configuration

X.Org is an open source implementation of the X Window System. It is further developed by the X.Org Foundation, which is also responsible for the development of new technologies and standards of the X Window System.

To use the available hardware, including mouse, graphics card, monitor, and keyboard, in the best way possible, the configuration can be optimized manually. Some aspects of this optimization are explained below. For detailed information about configuring the X Window System, review the various files in the directory `/usr/share/doc/packages/Xorg` and `man xorg.conf`.

WARNING

Be very careful when configuring your X Window System. Never start the X Window System until the configuration is finished. A wrongly configured system can cause irreparable damage to your hardware (this applies especially to fixed-frequency monitors). The authors of this book and SUSE Linux cannot be held responsible for damage. This information has been carefully researched, but this does not guarantee that all methods presented here are correct and will not damage your hardware.

The programs SaX2 and xorgconfig create the file `xorg.conf`, by default in `/etc/X11`. This is the primary configuration file for the X Window System. Find all the settings here concerning your graphics card, mouse, and monitor.

The following sections describe the structure of the configuration file `/etc/X11/xorg.conf`. It consists of several sections, each one dealing with a certain aspect of the configuration. Each section starts with the keyword `Section <designation>` and ends with `EndSection`. The sections have the form:

```

Section designation
entry 1
entry 2
entry n
EndSection

```

The available section types are listed in [Table 14.1, “Sections in /etc/X11/xorg.conf”](#) (page 274).

Table 14.1 *Sections in /etc/X11/xorg.conf*

Type	Meaning
Files	This section describes the paths used for fonts and the RGB color table.
ServerFlags	General switches are set here.
InputDevice	Input devices, like keyboards and special input devices (touch-pads, joysticks, etc.), are configured in this section. Important parameters in this section are <code>Driver</code> and the options defining the <code>Protocol</code> and <code>Device</code> .
Monitor	Describes the monitor used. The individual elements of this section are the name, which is referred to later in the <code>Screen</code> definition, the <code>bandwidth</code> , and the synchronization frequency limits (<code>HorizSync</code> and <code>VertRefresh</code>). Settings are given in MHz, kHz, and Hz. Normally, the server refuses any modeline that does not correspond with the specification of the monitor. This prevents too high frequencies from being sent to the monitor by accident.
Modes	The modeline parameters are stored here for the specific screen resolutions. These parameters can be calculated by SaX2 on the basis of the values given by the user and normally do not need to be changed. Intervene manually at this point if, for example, you want to connect a fixed frequency monitor. Find details of the meaning of individual number values in the HOWTO files in <code>/usr/share/doc/howto/en/html/XFree86-Video-Timings-HOWTO</code> .

Type	Meaning
Device	This section defines a specific graphics card. It is referenced by its descriptive name.
Screen	This section puts together a Monitor and a Device to form all the necessary settings for X.Org. In the Display subsection, specify the size of the virtual screen (Virtual), the ViewPort, and the Modes used with this screen.
ServerLayout	This section defines the layout of a single or multihead configuration. This section binds the input devices InputDevice and the display devices Screen.

Monitor, Device, and Screen are explained in more detail below. Further information about the other sections can be found in the manual pages of `X.Org` and `xorg.conf`.

There can be several different Monitor and Device sections in `xorg.conf`. Even multiple Screen sections are possible. The following ServerLayout section determines which one is used.

14.2.1 Screen Section

The screen section combines a monitor with a device section and determines the resolution and color depth to use. A screen section might resemble [Example 14.1, “Screen Section of the File /etc/X11/xorg.conf”](#) (page 276).

Example 14.1 Screen Section of the File /etc/X11/xorg.conf

```
Section "Screen"
    DefaultDepth 16
    SubSection "Display"
        Depth       16
        Modes      "1152x864" "1024x768" "800x600"
        Virtual     1152x864
    EndSubSection
    SubSection "Display"
        Depth       24
        Modes      "1280x1024"
    EndSubSection
    SubSection "Display"
        Depth       32
        Modes      "640x480"
    EndSubSection
    SubSection "Display"
        Depth       8
        Modes      "1280x1024"
    EndSubSection
    Device      "Device[0]"
    Identifier   "Screen[0]"
    Monitor     "Monitor[0]"
EndSection
```

The Identifier line (here Screen[0]) gives this section a defined name with which it can be uniquely referenced in the following ServerLayout section. The lines Device and Monitor specify the graphics card and the monitor that belong to this definition. These are just links to the Device and Monitor sections with their corresponding names or *identifiers*. These sections are discussed in detail below.

Use the DefaultDepth setting to select the color depth the server should use unless it is started with a specific color depth. There is a Display subsection for each color depth. The keyword Depth assigns the color depth valid for this subsection. Possible values for Depth are 8, 15, 16, and 24. Not all X server modules support all these values.

After the color depth, a list of resolutions is set in the Modes section. This list is checked by the X server from left to right. For each resolution, the X server searches for a suitable Modeline in the Modes section. The Modeline depends on the capability of both the monitor and the graphics card. The Monitor settings determine the resulting Modeline.

The first resolution found is the `Default` mode. With `Ctrl` + `Alt` + `+` (on the number pad), switch to the next resolution in the list to the right. With `Ctrl` + `Alt` + `-` (on the number pad), switch to the left. This enables you to vary the resolution while X is running.

The last line of the `Display` subsection with `Depth 16` refers to the size of the virtual screen. The maximum possible size of a virtual screen depends on the amount of memory installed on the graphics card and the desired color depth, not on the maximum resolution of the monitor. Because modern graphics cards have a large amount of video memory, you can create very large virtual desktops. However, you may no longer be able to use 3D functionality if you fill most of the video memory with a virtual desktop. If the card has 16 MB video RAM, for example, the virtual screen can be up to 4096x4096 pixels in size at 8-bit color depth. Especially for accelerated cards, however, it is not recommended to use all your memory for the virtual screen, because this memory on the card is also used for several font and graphics caches.

14.2.2 Device Section

A device section describes a specific graphics card. You can have as many device entries in `xorg.conf` as you like, provided their names are differentiated using the keyword `Identifier`. As a rule—if you have more than one graphics card installed—the sections are simply numbered in order. The first one is called `Device[0]`, the second one `Device[1]`, and so on. The following file shows an excerpt from the `Device` section of a computer with a Matrox Millennium PCI graphics card:

```
Section "Device"
    BoardName      "MGA2064W"
    BusID         "0:19:0"
    Driver        "mga"
    Identifier    "Device[0]"
    VendorName   "Matrox"
    Option       "sw_cursor"
EndSection
```

If you use SaX2 for configuration, the device section should look something like the above example. Both the `Driver` and `BusID` are dependent on the hardware installed in your computer and are detected by SaX2 automatically. The `BusID` defines the PCI or AGP slot in which the graphics card is installed. This matches the ID displayed by the command `lspci`. The X server needs details in decimal form, but `lspci` displays these in hexadecimal form.

With the `Driver` parameter, specify the driver to use for this graphics card. If the card is a Matrox Millennium, the driver module is called `mga`. The X server then searches through the `ModulePath` defined in the `Files` section in the `drivers` subdirectory. In a standard installation, this is the directory `/usr/X11R6/lib/modules/drivers`. `_drv.o` is added to the name, so, in the case of the `mga` driver, the driver file `mga_drv.o` is loaded.

The behavior of the X server or of the driver can also be influenced through additional options. An example of this is the option `sw_cursor`, which is set in the `device` section. This deactivates the hardware mouse cursor and depicts the mouse cursor using software. Depending on the driver module, there are various options available, which can be found in the description files of the driver modules in the directory `/usr/X11R6/lib/X11/doc`. Generally valid options can also be found in the manual pages (`man xorg.conf` and `man X.Org`).

14.2.3 Monitor and Modes Section

Like the `Device` sections, the `Monitor` and `Modes` sections describe one monitor each. The configuration file `/etc/X11/xorg.conf` can contain as many `Monitor` sections as desired. The `server layout` section specifies which `Monitor` section is relevant.

`Monitor` definitions should only be set by experienced users. The modelines constitute an important part of the `Monitor` sections. Modelines set horizontal and vertical timings for the respective resolution. The monitor properties, especially the allowed frequencies, are stored in the `Monitor` section.

WARNING

Unless you have an in-depth knowledge of monitor and graphics card functions, nothing should be changed in the modelines, because this could cause severe damage to your monitor.

Those who try to develop their own monitor descriptions should be very familiar with the documentation in `/usr/X11/lib/X11/doc`. The section covering the video modes deserves a special mention. It describes, in detail, how the hardware functions and how to create modelines.

Manual specification of modelines is rarely required today. If you are using a modern multisync monitor, the allowed frequencies and optimal resolutions can, as a rule, be read directly from the monitor by the X server via DDC, as described in the SaX2 configuration section. If this is not possible for some reason, use one of the VESA modes included in the X server. This will function with practically all graphics card and monitor combinations.

14.3 Installing and Configuring Fonts

The installation of additional fonts in SUSE Linux is very easy. Simply copy the fonts to any directory located in the X11 font path (see [Section 14.3.1, “X11 Core Fonts”](#) (page 280)). To enable use of the fonts, the installation directory should be a subdirectory of the directories configured in `/etc/fonts/fonts.conf` (see [Section 14.3.2, “Xft”](#) (page 281)).

The font files can be copied manually (as `root`) to a suitable directory, such as `/usr/X11R6/lib/X11/fonts/truetype`. Alternatively, the task can be performed with the KDE font installer in the KDE Control Center. The result is the same.

Instead of copying the actual fonts, you can also create symbolic links. For example, you may want to do this if you have licensed fonts on a mounted Windows partition and want to use them. Subsequently, run `SuSEconfig --module fonts`.

`SuSEconfig --module fonts` executes the script `/usr/sbin/fonts-config`, which handles the configuration of the fonts. To see what this script does, refer to the manual page of the script (`man fonts-config`).

The procedure is the same for bitmap fonts, TrueType and OpenType fonts, and Type1 (PostScript) fonts. All these font types can be installed in any directory. Only CID-keyed fonts require a slightly different procedure. For this, see [Section 14.3.3, “CID-Keyed Fonts”](#) (page 284).

X.Org contains two completely different font systems: the old *X11 core font system* and the newly designed *Xft and fontconfig* system. The following sections briefly describe these two systems.

14.3.1 X11 Core Fonts

Today, the X11 core font system supports not only bitmap fonts but also scalable fonts, like Type1 fonts, TrueType and OpenType fonts, and CID-keyed fonts. Scalable fonts are only supported without antialiasing and subpixel rendering and the loading of large scalable fonts with glyphs for many languages may take a long time. Unicode fonts are also supported, but their use may be slow and require more memory.

The X11 core font system has a few inherent weaknesses. It is outdated and can no longer be extended in a meaningful fashion. Although it must be retained for reasons of backward compatibility, the more modern Xft and fontconfig system should be used if at all possible.

For its operation, the X server needs to know what fonts it has available and where in the system it can find them. This is handled by a `FontPath` variable, which contains the path to all valid system font directories. In each of these directories, a file named `fonts.dir` lists the available fonts in this directory. The `FontPath` is generated by the X server at start-up. It searches for a valid `fonts.dir` file in each of the `FontPath` entries in the configuration file `/etc/X11/xorg.conf`. These entries are found in the `Files` section. Display the actual `FontPath` with `xset q`. This path may also be changed at runtime with `xset`. To add an additional path, use `xset +fp <path>`. To remove an unwanted path, use `xset -fp <path>`.

If the X server is already active, newly installed fonts in mounted directories can be made available with the command `xset fp rehash`. This command is executed by `SuSEconfig --module fonts`. Because the command `xset` needs access to the running X server, this only works if `SuSEconfig --module fonts` is started from a shell that has access to the running X server. The easiest way to achieve this is to assume `root` permissions by entering `su` and the root password. `su` transfers the access permissions of the user who started the X server to the root shell. To check if the fonts were installed correctly and are available by way of the X11 core font system, use the command `xlsfonts` to list all available fonts.

By default, SUSE Linux uses UTF-8 locales. Therefore, Unicode fonts should be preferred (font names ending with `iso10646-1` in `xlsfonts` output). All available Unicode fonts can be listed with `xlsfonts | grep iso10646-1`. Nearly all Unicode fonts available in SUSE Linux contain at least the glyphs needed for European languages (formerly encoded as `iso-8859-*`).

14.3.2 Xft

From the outset, the programmers of Xft made sure that scalable fonts including antialiasing are supported well. If Xft is used, the fonts are rendered by the application using the fonts, not by the X server as in the X11 core font system. In this way, the respective application has access to the actual font files and full control of how the glyphs are rendered. This constitutes the basis for the correct display of text in a number of languages. Direct access to the font files is very useful for embedding fonts for printing to make sure that the printout looks the same as the screen output.

In SUSE Linux, the two desktop environments KDE and GNOME, Mozilla, and many other applications already use Xft by default. Xft is already used by more applications than the old X11 core font system.

Xft uses the fontconfig library for finding fonts and influencing how they are rendered. The properties of fontconfig are controlled by the global configuration file `/etc/fonts/fonts.conf` and the user-specific configuration file `~/.fonts.conf`. Each of these fontconfig configuration files must begin with

```
<?xml version="1.0"?>
<!DOCTYPE fontconfig SYSTEM "fonts.dtd">
<fontconfig>
```

and end with

```
</fontconfig>
```

To add directories to search for fonts, append lines such as the following:

```
<dir>/usr/local/share/fonts/</dir>
```

However, this is usually not necessary. By default, the user-specific directory `~/.fonts` is already entered in `/etc/fonts/fonts.conf`. Accordingly, all you need to do to install additional fonts is to copy them to `~/.fonts`.

You can also insert rules that influence the appearance of the fonts. For example, enter

```
<match target="font">
  <edit name="antialias" mode="assign">
    <bool>false</bool>
  </edit>
</match>
```

to disable antialiasing for all fonts or

```
<match target="font">
  <test name="family">
    <string>Luxi Mono</string>
    <string>Luxi Sans</string>
  </test>
  <edit name="antialias" mode="assign">
    <bool>false</bool>
  </edit>
</match>
```

to disable antialiasing for specific fonts.

By default, most applications use the font names `sans-serif` (or the equivalent `sans`), `serif`, or `monospace`. These are not real fonts but only aliases that are resolved to a suitable font, depending on the language setting.

Users can easily add rules to `~/.fonts.conf` to resolve these aliases to their favorite fonts:

```
<alias>
  <family>sans-serif</family>
  <prefer>
    <family>FreeSans</family>
  </prefer>
</alias>
<alias>
  <family>serif</family>
  <prefer>
    <family>FreeSerif</family>
  </prefer>
</alias>
<alias>
  <family>monospace</family>
  <prefer>
    <family>FreeMono</family>
  </prefer>
</alias>
```

Because nearly all applications use these aliases by default, this affects almost the entire system. Thus, you can easily use your favorite fonts almost everywhere without having to modify the font settings in the individual applications.

Use the command `fc-list` to find out which fonts are installed and available for use. For instance, the command `fc-list` returns a list of all fonts. To find out which of the available scalable fonts (`:scalable=true`) contain all glyphs required for Hebrew (`:lang=he`), their font names (`family`), their style (`style`), their weight (`weight`), and the name of the files containing the fonts, enter the following command:

```
fc-list ":lang=he:scalable=true" family style weight
```

The output of this command could look like the following:

```
FreeSansBold.ttf: FreeSans:style=Bold:weight=200
FreeMonoBoldOblique.ttf: FreeMono:style=BoldOblique:weight=200
FreeSerif.ttf: FreeSerif:style=Medium:weight=80
FreeSerifBoldItalic.ttf: FreeSerif:style=BoldItalic:weight=200
FreeSansOblique.ttf: FreeSans:style=Oblique:weight=80
FreeSerifItalic.ttf: FreeSerif:style=Italic:weight=80
FreeMonoOblique.ttf: FreeMono:style=Oblique:weight=80
FreeMono.ttf: FreeMono:style=Medium:weight=80
FreeSans.ttf: FreeSans:style=Medium:weight=80
FreeSerifBold.ttf: FreeSerif:style=Bold:weight=200
FreeSansBoldOblique.ttf: FreeSans:style=BoldOblique:weight=200
FreeMonoBold.ttf: FreeMono:style=Bold:weight=200
```

Important parameters that can be queried with `fc-list`:

Table 14.2 Parameters of `fc-list`

Parameter	Meaning and Possible Values
family	Name of the font family, for example, <code>FreeSans</code> .
foundry	The manufacturer of the font, for example, <code>urw</code> .
style	The font style, such as <code>Medium</code> , <code>Regular</code> , <code>Bold</code> , <code>Italic</code> , or <code>Heavy</code> .
lang	The language that the font supports, for example, <code>de</code> for German, <code>ja</code> for Japanese, <code>zh-TW</code> for traditional Chinese, or <code>zh-CN</code> for simplified Chinese.
weight	The font weight, such as <code>80</code> for regular or <code>200</code> for bold.
slant	The slant, usually <code>0</code> for none and <code>100</code> for italic.
file	The name of the file containing the font.
outline	<code>true</code> for outline fonts or <code>false</code> for other fonts.

Parameter	Meaning and Possible Values
scalable	true for scalable fonts or false for other fonts.
bitmap	true for bitmap fonts or false for other fonts.
pixelsize	Font size in pixels. In connection with fc-list, this option only makes sense for bitmap fonts.

14.3.3 CID-Keyed Fonts

In contrast to the other font types, you cannot simply install CID-keyed fonts in just any directory. CID-keyed fonts must be installed in `/usr/share/ghostscript/Resource/CIDFont`. This is not relevant for Xft and fontconfig, but it is necessary for Ghostscript and the X11 core font system.

TIP

See <http://www.xfree86.org/current/fonts.html> for more information about fonts under X11.

14.4 OpenGL—3D Configuration

14.4.1 Hardware Support

SUSE Linux includes several OpenGL drivers for 3D hardware support. [Table 14.3, “Supported 3D Hardware”](#) (page 284) provides an overview.

Table 14.3 *Supported 3D Hardware*

OpenGL Driver	Supported Hardware
nVidia	nVidia Chips: all except some legacy chipsets (GeForce2 and older)

OpenGL Driver	Supported Hardware
DRI	Intel i810/i815/i830M, Intel 845G/852GM/855GM/865G/915G,915GM/945G Matrox G200/G400/G450/G550, ATI Rage 128(Pro)/Radeon (up to 9250)

If you are installing with YaST for the first time, 3D acceleration can be activated during installation, provided YaST detects 3D support. For nVidia graphics chips, the nVidia driver must be installed first. To do this, select the nVidia driver patch in YOU (YaST Online Update). Due to license restrictions, the nVidia driver is not included in the distribution.

If you update your system instead, the procedure for configuring 3D hardware support is different. This depends on which OpenGL driver is used. Further details are provided in the following section.

14.4.2 OpenGL Drivers

The OpenGL drivers nVidia and DRI can be configured easily with SaX2. For nVidia adapters, the nVidia driver must be installed first. Enter the command `3Ddiag` to check if the configuration for nVidia or DRI is correct.

For security reasons, only users belonging to the group `video` are permitted to access the 3D hardware. Therefore, make sure that all local users are members of this group. Otherwise, the slow *software rendering fallback* of the OpenGL driver is used for OpenGL applications. Use the command `id` to check whether the current user belongs to the group `video`. If this is not the case, use YaST to add the user to the group.

14.4.3 The Diagnosis Tool `3Ddiag`

The diagnosis tool `3Ddiag` allows verification of the 3D configuration in SUSE Linux. This is a command line tool that must be started in a terminal. Enter `3Ddiag -h` to list possible options for `3Ddiag`.

To verify the X.Org configuration, the tool checks if the packages needed for 3D support are installed and if the correct OpenGL library and GLX extension are used. Follow the instructions of 3Ddiag if you receive failed messages. If everything is correct, you only see done messages on the screen.

14.4.4 OpenGL Test Utilities

For testing OpenGL, the program `glxgears` and games like `tuxracer` and `armagetron` (packages have the same names) can be useful. If 3D support has been activated, it should be possible to play these smoothly on a fairly new computer. Without 3D support, these games would run very slowly (slideshow effect). Use the `glxinfo` command to verify that 3D is active, in which case the output contains a line with `direct rendering: Yes`.

14.4.5 Troubleshooting

If the OpenGL 3D test results are negative (the games cannot be smoothly played), use 3Ddiag to make sure no errors exist in the configuration (failed messages). If correcting these does not help or if failed messages have not appeared, take a look at the X.Org log files.

Often, you will find the line `DRI is disabled` in the X.Org file `/var/log/Xorg.0.log`. The exact cause can only be discovered by closely examining the log file—a task requiring some experience.

In such cases, no configuration error exists, because this would have already been detected by 3Ddiag. Consequently, at this point, the only choice is to use the software rendering fallback of the DRI driver, which does not provide 3D hardware support. You should also go without 3D support if you get OpenGL representation errors or instability. Use SaX2 to disable 3D support completely.

14.4.6 Installation Support

Apart from the software rendering fallback of the DRI driver, some OpenGL drivers in Linux are still in developmental phases and are therefore considered experimental. The drivers are included in the distribution because of the high demand for 3D hardware acceleration in Linux. Considering the experimental status of some OpenGL

drivers, SUSE cannot offer any installation support for configuring 3D hardware acceleration or provide any further assistance with related problems. The basic configuration of the graphical user interface (X Window System) does not include 3D hardware acceleration configuration. If you experience problems with 3D hardware acceleration, it is recommended to disable 3D support completely.

14.4.7 For More Information

For information, refer to the README files in `/usr/X11R6/lib/X11/doc`. Find more information about nvidia driver installation at <http://www.suse.de/~sndirsch/nvidia-installer-HOWTO.html>.

FreeNX: Remotely Controlling Another Computer

15

FreeNX is a GPL implementation of the NX Server, used for remote access and display of another computer. It provides near-local speed application responsiveness over high-latency, low-bandwidth links.

15.1 Getting Started with NX

The following steps outline the basic procedure for establishing a working NX setup that allows up to 10 clients to connect to the NX server:

- 1 Install the following packages on the server and client machine using the YaST Software Management module:

Server Machine	Client Machine
<ul style="list-style-type: none">• NX	<ul style="list-style-type: none">• NX
<ul style="list-style-type: none">• FreeNX	<ul style="list-style-type: none">• knx (for KDE sessions)• NoMachine nxclient (for non-KDE sessions)

- 2 Set up the NX server by issuing the following command as root:

```
nxsetup --install --clean --purge --setup-nomachine-key
```

The server runs and runs according to the default settings in `/etc/nxserver/node.conf`. Any user can remotely connect from another workstation. For advanced NX server configuration, refer to [Section 15.2, “Advanced FreeNX Configuration”](#) (page 291).

If you prefer a more secure setup with private keys distributed to each client, refer to the instructions given in [Section 15.2.1, “Configuring SSH Authentication Using Client Keys”](#) (page 292).

- 3 Configure the firewall on the machine hosting the NX server to allow NX connections.
 - a Log in to the server machine as root and start the YaST Firewall module.
 - b Select *Allowed Services* to enter the service configuration dialog and select *External Zone*.
 - c Select *Advanced* to enter the port details for NX.
 - d Open ports 22 (SSH), 5000 to 5009, and 7000 to 7009 to allow NX traffic. Do this by entering the following in *TCP ports*:

```
22 5000:5009 7000:7009
```
 - e Store your settings and restart the firewall by selecting *OK* → *Next* → *Accept*.

TIP

For detailed information about firewall configuration for NX, refer to `/usr/share/doc/packages/FreeNX/NX-Firewall.txt`.

To remotely connect to another workstation and use KDE as your desktop choice, proceed as follows:

- 1 Start KNX from the main menu.
- 2 The first time you log in, you need to create a new connection. To create a connection, do the following:

- a** In *KNX Client Login*, click *Connection Settings*.
 - b** Enter a name for the connection, such as the name of the server.
 - c** Enter the host information, the port number, and the bandwidth for your connection.
 - d** From *Session type*, select *UNIX/KDE* to start a KDE session.
 - e** Select a display resolution.
 - f** Click *OK*.
- 3** Once you are connected and the remote connection appears on your screen, you can access applications and use the remote computer as though you were sitting at that machine.

To remotely connect to another machine running GNOME as your desktop choice, proceed as follows:

- 1** Download and install the nxclient package from NoMachine via http://www.nomachine.com/download_client_linux.php.
- 2** Start *NX Connection Wizard* from the main menu.
- 3** In three steps, enter the name of the connection, port and host details, and connection type, select the session type *Unix/Gnome*, decide whether to create a shortcut on your desktop, and click *Finish*.
- 4** To connect to the remote desktop, click the NX shortcut on your desktop and provide username and password and click *OK*.

The remote desktop appears on your screen.

15.2 Advanced FreeNX Configuration

The following sections introduce some advanced features mainly needed in more complex NX scenarios.

15.2.1 Configuring SSH Authentication Using Client Keys

The authentication configured in [Section 15.1, “Getting Started with NX”](#) (page 289) solely relies on username and password credentials. For a more secure authentication, NX can be configured to generate a pair of SSH keys. The client key is then copied from the server machine to any client that should be able to connect to the NX server. Clients that do not present this key cannot authenticate at the NX server. This feature is only supported for the FreeNX server/knx client combination.

To configure the NX server to use this authentication method and generate the appropriate key pair, proceed as follows:

- 1** Log in as root to the server machine.
- 2** Open the server's configuration file `/etc/nxserver/node.conf` and make sure that `ENABLE_SSH_AUTHENTICATION` is set to 1 (which should be default).
- 3** Install the server with the following command:

```
nxsetup --install --clean --purge
```

- 4** Adjust the access permissions to `/var/lib/nxserver/home/.ssh/authorized_keys2`:

```
chmod 640 /var/lib/nxserver/home/.ssh/authorized_keys2
```

- 5** Log out.

To configure knx to use this key, proceed as follows:

- 1** At the server machine, log in as root.
- 2** Copy the key file to the location on the client machine where knx needs it, replacing `client` with the client's address.

```
scp /var/lib/nxserver/home/.ssh/client.id_dsa.key client:/usr/share/knx/
```
- 3** Log in to the client machine as root.

- 4 Adjust the access permissions as follows:

```
chmod 644 /usr/share/knx/client.id_dsa.key
```

- 5 Log out.

15.2.2 Configuring PAM Authentication

By default, FreeNX allows anyone to open an NX session provided this user is present in the user database of your server (locally or via LDAP, NIS, etc.). This behavior is toggled by the `ENABLE_PAM_AUTHENTICATION` variable in `/usr/bin/nxserver` on the server machine. The default value here is 1. Setting it to 0 disables the PAM-mediated user authentication (`PAM_AUTH`) for FreeNX.

If `ENABLE_PAM_AUTHENTICATION` is set to 0, you need to add users and passwords manually. To add local NX users on the server, proceed as follows:

- 1 Log in to the server machine as root.
- 2 Make sure that any user to add exists in the system's database of local users by checking the contents of `/etc/passwd` or using the YaST User Management module.
- 3 For each user to add, add the username with the command `nxserver --adduser`. Then add the user's password with `nxserver --passwd`.
- 4 Restart the server with `nxserver --restart` and log out.

15.2.3 Using Systemwide and User-Specific Configuration Files

The FreeNX server's behavior is controlled via `/etc/node.conf`. You can either run a global NX server configuration or run the server with user-specific configurations. This comes into play, if you have different users running NX on one machine with different requirements.

In the following example, assume user `joe` wants NX automatically to start with a certain application as soon as he opens an NX session. To enable this behavior for this user only, proceed as follows:

- 1** Log in as root.
- 2** Enter the `/etc/nxserver` directory:

```
cd /etc/nxserver
```

- 3** Save a copy of the NX server's configuration file (`node.conf`) under `joe.node.conf` in the same directory.
- 4** Edit the appropriate parameters (`NODE_AUTOSTART` and `ENABLE_AUTORECONNECT`) in `joe.node.conf`. For details on these features, refer to [Section 15.2.5, “Configuring Autostart Tasks and Exporting Configurations”](#) (page 295) and [Section 15.2.4, “Suspending and Resuming NX Sessions”](#) (page 294).
- 5** Reinstall the NX server to activate the new configuration:

```
nxsetup --install --clean --purge --setup-nomachine-key
```

The user-specific configuration overrides the global configuration.
- 6** Log out.

15.2.4 Suspending and Resuming NX Sessions

As with sessions on a mobile computer, it is equally possible to configure NX to allow suspend and resume of user sessions. A suspended session reopens exactly in the same state as you left it.

To configure suspend and resume for NX sessions, proceed as follows:

- 1** Log in as root.
- 2** Open the server's configuration file, `/etc/nxserver/node.conf`, and edit it as follows:

```
ENABLE_PASSDB_AUTHENTICATION="0"
ENABLE_USER_DB="0"
ENABLE_AUTORECONNECT="1"
```

- 3 Save and exit the configuration file and restart the server with `nxserver --restart`.

- 4 Log out.

To suspend a session on exit, click the *X* in the top right corner of your NX window and select *Suspend* to suspend your session and to exit the client. Upon reconnect, you are asked whether to resume the old session or start a fresh one.

15.2.5 Configuring Autostart Tasks and Exporting Configurations

FreeNX offers an autostart functionality that allows you to launch certain tasks when starting or resuming an NX session, provided the underlying application supports the `start` and `resume` properties. For example, you can automatically clean up the desktop or do other autostart tasks when you start FreeNX. This is especially useful when you reconnect a session, even from a different NX client (where you cannot use the standard KDE or GNOME mechanisms).

To configure autostart features, proceed as follows:

- 1 Log in as root on the server machine.
- 2 Open the server's configuration file `/etc/nxserver/node.conf` and edit the `NODE_AUTOSTART` variable to the following, replacing `myprogram` with the program that should be executed on start or resume of an NX session:

```
NODE_AUTOSTART=myprogram
```

- 3 Save and exit the configuration file.
- 4 Restart the server with the `nxserver --restart` command and log out.

The program specified now starts every time a session is started or resumed.

You can also export the variables NX_USERIP and NX_SESSIONID to make them accessible in the user's environment. This allows, for example, putting an icon onto the desktop with the generic content and accessing a Samba server running on the user's thin client. To make the contents of a floppy on the thin client's floppy drive available to the user, proceed as follows:

- 1** Enable the export of the NX_USERIP and NX_SESSIONID variables on the server side:

- a** Log in as root on the server.
 - b** Open the server's configuration /etc/nxserver/node.conf and set the following variables:

```
EXPORT_USERIP="1"  
EXPORT_SESSIONID="1"
```

- c** Save and exit the server's configuration and restart the server using the nxserver --restart command.

- d** Log out.

- 2** On the client side, open a session, export the floppy drive via SMB and create an icon on the desktop:

- a** Export the contents of your floppy drive through Samba using your file manager (Nautilus or Konqueror).
 - b** Create a floppy.desktop file in the Desktop directory and enter the following line:

```
Exec=smb://$NX_USERIP/floppy
```

The server exports the thin client's IP address, allowing you to access the thin client's floppy drive with the floppy icon in the NX session.

15.2.6 Creating a Chain of NX Servers

A chain of NX servers enables you to traverse firewalls and cope with IP masquerading. An external “gateway” server can be used to forward incoming connections to an internal server hidden behind a (masquerading) firewall.

To configure a chain of NX servers, proceed as follows:

- 1 Configure the internal server as described in [Section 15.2.1, “Configuring SSH Authentication Using Client Keys”](#) (page 292) and distribute the server's private key (`client.id_dsa.key`) to `/usr/NX/share/` on the gateway.
- 2 On the gateway server, proceed as follows:

- a Log in as root.

- b Set the following variables in `/etc/nxserver/node.conf`, replacing `myinternalhost` with the IP address of the internal NX server:

```
ENABLE_SERVER_FORWARD="1"  
SERVER_FORWARD_HOST="myinternalhost"  
SERVER_FORWARD_KEY="/usr/NX/share/client.id_dsa.key"
```

- c Restart the external server to apply the altered configuration with the command `nxserver --restart` and log out.

Any incoming connection is forwarded to the internal server.

15.2.7 Installing and Running FreeNX and NoMachine on the Same Server

You can install and run FreeNX and the commercial NoMachine NX server on the same machine without interference. This is implemented in FreeNX by forwarding the connection to the NoMachine installed on the same machine.

To enable this feature, proceed as follows:

- 1 Log in as root on the server machine.

- 2 Open the server's configuration file for FreeNX under `/etc/nxserver/node.conf` and set the following variable:

```
ENABLE_NOMACHINE_FORWARD="1"
```

- 3 Save this file and restart the FreeNX server with the `nserver --restart` command.

- 4 Log out.

To connect to the NoMachine server, use the standard username and password credentials. To connect to the FreeNX server, prepend `freenx.` to the normal username (for example, `freenx.joedoe`) and use the usual password.

15.3 Troubleshooting

The following sections hint at some of the most frequent problems encountered when using FreeNX and provide possible solutions to these problems.

15.3.1 knx Hangs When Trying to Establish a Connection

You are trying to establish a connection to your NX server using knx. When initiating the connection, knx fails to authenticate the user and no remote session is ever started.

To determine why this happens and how the problem can be solved, proceed as follows:

- 1 Check whether Novell AppArmor is running on the server machine and proceed as described in [Section 15.3.2, “Connection to NX Server Cannot Be Established”](#) (page 299).
- 2 Retry establishing a connection between knx and the server.
- 3 Check whether the firewall on the client side allows SSH traffic by starting the YaST Firewall module and checking whether SSH is listed among *Allowed Services* for the *External Zone*. Enable SSH if it has not already been enabled.

- 4** Check the firewall on the server side for SSH and for the NX ports listed in [Section 15.1, “Getting Started with NX”](#) (page 289). Open these ports if they had previously been closed.
- 5** Retry establishing a connection between knx and the server.
- 6** Log in to the server as root and proceed as follows:
 - a** Enter the /tmp directory and check for lock files of the NX server:

```
cd /
ls -ltr .nx*
```
 - b** If there are any of these old lock files, remove them.
 - c** Log out.

- 7** Retry establishing a connection between knx and the server.
- 8** On the client machine, delete and reinstall the knx client using the YaST Software Management module.

You should be able to connect to the server now, provided you followed all the instructions above.

15.3.2 Connection to NX Server Cannot Be Established

After starting knx and initiating the connection, you get the following error message:

```
Connection to NX server could not be established. Connection timed out.
```

To determine the source of this problem, proceed as follows:

- 1** Log in to the server as root.
- 2** Check the output of the dmesg command for an entry like the following:

```
SubDomain: REJECTING r access to
/var/lib/nxserver/home/.ssh/authorized_keys2 (sshd(31247) profile
/usr/sbin/sshd active /usr/sbin/sshd)
```

This entry tells you that Novell AppArmor running on the server does not allow the ssh daemon to access some NX-specific files.

- 3 Stop AppArmor on the server machine

or

Put the sshd profile into learning mode and add permissions for accessing the NX files to the existing profile. This is described in more detail in the *Novell AppArmor 2.0 Administration Guide*.

- 4 Reconnect to the server.

15.3.3 User Authentication Succeeds, Remote Connection Not Established

After you have run knx and initiated the session, knx succeeds in authenticating the user, but instead of a terminal window opening with the new session, you get an error message like the following:

```
Could not yet establish the connection to the remote proxy. Do you  
want to terminate the current session?
```

The connection failed due to the higher ports used in negotiating the NX remote session not being opened on the server's firewall. To adjust the firewall settings on the server, proceed as described in [Section 15.1, “Getting Started with NX”](#) (page 289).

15.4 For More Information

For the latest information about the current FreeNX package, refer to the README at `/usr/share/doc/packages/FreeNX/README.SUSE`. For additional information about NX Server, use the command `nxserver --help`.

16

Authentication with PAM

Linux uses PAM (pluggable authentication modules) in the authentication process as a layer that mediates between user and application. PAM modules are available on a systemwide basis, so they can be requested by any application. This chapter describes how the modular authentication mechanism works and how it is configured.

System administrators and programmers often want to restrict access to certain parts of the system or to limit the use of certain functions of an application. Without PAM, applications must be adapted every time a new authentication mechanism, such as LDAP or SAMBA, is introduced. This process, however, is rather time-consuming and error-prone. One way to avoid these drawbacks is to separate applications from the authentication mechanism and delegate authentication to centrally managed modules. Whenever a newly required authentication scheme is needed, it is sufficient to adapt or write a suitable PAM module for use by the program in question.

Every program that relies on the PAM mechanism has its own configuration file in the directory `/etc/pam.d/programname`. These files define the PAM modules used for authentication. In addition, there are global configuration files for most PAM modules under `/etc/security`, which define the exact behavior of these modules (examples include `pam_env.conf`, `pam_pwcheck.conf`, `pam_unix2.conf`, and `time.conf`). Every application that uses a PAM module actually calls a set of PAM functions, which then process the information in the various configuration files and return the result to the calling application.

16.1 Structure of a PAM Configuration File

Each line in a PAM configuration file contains a maximum of four columns:

```
<Type of module> <Control flag> <Module path> <Options>
```

PAM modules are processed as stacks. Different types of modules have different purposes, for example, one module checks the password, another one verifies the location from which the system is accessed, and yet another one reads user-specific settings. PAM knows about four different types of modules:

`auth`

The purpose of this type of module is to check the user's authenticity. This is traditionally done by querying a password, but it can also be achieved with the help of a chip card or through biometrics (fingerprints or iris scan).

`account`

Modules of this type check whether the user has general permission to use the requested service. As an example, such a check should be performed to ensure that no one can log in under the username of an expired account.

`password`

The purpose of this type of module is to enable the change of an authentication token. In most cases, this is a password.

`session`

Modules of this type are responsible for managing and configuring user sessions. They are started before and after authentication to register login attempts in system logs and configure the user's specific environment (mail accounts, home directory, system limits, etc.).

The second column contains control flags to influence the behavior of the modules started:

`required`

A module with this flag must be successfully processed before the authentication may proceed. After the failure of a module with the `required` flag, all other

modules with the same flag are processed before the user receives a message about the failure of the authentication attempt.

requisite

Modules having this flag must also be processed successfully, in much the same way as a module with the `required` flag. However, in case of failure a module with this flag gives immediate feedback to the user and no further modules are processed. In case of success, other modules are subsequently processed, just like any modules with the `required` flag. The `requisite` flag can be used as a basic filter checking for the existence of certain conditions that are essential for a correct authentication.

sufficient

After a module with this flag has been successfully processed, the calling application receives an immediate message about the success and no further modules are processed, provided there was no preceding failure of a module with the `required` flag. The failure of a module with the `sufficient` flag has no direct consequences, in the sense that any subsequent modules are processed in their respective order.

optional

The failure or success of a module with this flag does not have any direct consequences. This can be useful for modules that are only intended to display a message (for example, to tell the user that mail has arrived) without taking any further action.

include

If this flag is given, the file specified as argument is inserted at this place.

The module path does not need to be specified explicitly, as long as the module is located in the default directory `/lib/security` (for all 64-bit platforms supported by SUSE Linux, the directory is `/lib64/security`). The fourth column may contain an option for the given module, such as `debug` (enables debugging) or `nullok` (allows the use of empty passwords).

16.2 The PAM Configuration of sshd

To show how the theory behind PAM works, consider the PAM configuration of `sshd` as a practical example:

Example 16.1 PAM Configuration for sshd

```
%PAM-1.0
auth    include      common-auth
auth    required     pam_nologin.so
account include    common-account
password include   common-password
session include    common-session
# Enable the following line to get resmgr support for
# ssh sessions (see /usr/share/doc/packages/resmgr/README.SuSE)
#session optional   pam_resmgr.so fake_ttyname
```

The typical PAM configuration of an application (sshd, in this case) contains four include statements referring to the configuration files of four module types: `common-auth`, `common-account`, `common-password`, and `common-session`. These four files hold the default configuration for each module type. By including them instead of calling each module separately for each PAM application, automatically get an updated PAM configuration if the administrator changes the defaults. In former times, you had to adjust all configuration files manually for all applications when changes to PAM occurred or a new application was installed. Now the PAM configuration is made with central configuration files and all changes are automatically inherited by the PAM configuration of each service.

The first include file (`common-auth`) calls two modules of the `auth` type: `pam_env` and `pam_unix2`. See [Example 16.2, “Default Configuration for the auth Section”](#) (page 304).

Example 16.2 Default Configuration for the auth Section

```
auth    required     pam_env.so
auth    required     pam_unix2.so
```

The first one, `pam_env`, loads the file `/etc/security/pam_env.conf` to set the environment variables as specified in this file. This can be used to set the `DISPLAY` variable to the correct value, because the `pam_env` module knows about the location from which the login is taking place. The second one, `pam_unix2`, checks the user's login and password against `/etc/passwd` and `/etc/shadow`.

After the modules specified in `common-auth` have been successfully called, a third module called `pam_nologin` checks whether the file `/etc/nologin` exists. If it does, no user other than `root` may log in. The whole stack of `auth` modules is processed before `sshd` gets any feedback about whether the login has succeeded. Given that all modules of the stack have the `required` control flag, they must all be processed successfully before `sshd` receives a message about the positive result. If one of the

modules is not successful, the entire module stack is still processed and only then is sshd notified about the negative result.

As soon as all modules of the auth type have been successfully processed, another include statement is processed, in this case, that in [Example 16.3, “Default Configuration for the account Section”](#) (page 305). common-account contains just one module, pam_unix2. If pam_unix2 returns the result that the user exists, sshd receives a message announcing this success and the next stack of modules (password) is processed, shown in [Example 16.4, “Default Configuration for the password Section”](#) (page 305).

Example 16.3 Default Configuration for the account Section

```
account required      pam_unix2.so
```

Example 16.4 Default Configuration for the password Section

```
password required    pam_pwcheck.so  nullok
password required    pam_unix2.so     nullok use_first_pass use_authok
#password required   pam_make.so    /var/yp
```

Again, the PAM configuration of sshd involves just an include statement referring to the default configuration for password modules located in common-password. These modules must successfully be completed (control flag required) whenever the application requests the change of an authentication token. Changing a password or another authentication token requires a security check. This is achieved with the pam_pwcheck module. The pam_unix2 module used afterwards carries over any old and new passwords from pam_pwcheck, so the user does not need to authenticate again. This also makes it impossible to circumvent the checks carried out by pam_pwcheck. The modules of the password type should be used wherever the preceding modules of the account or the auth type are configured to complain about an expired password.

Example 16.5 Default Configuration for the session Section

```
session required     pam_limits.so
session required     pam_unix2.so
```

As the final step, the modules of the session type, bundled in the common-session file are called to configure the session according to the settings for the user in question. Although pam_unix2 is processed again, it has no practical consequences due to its none option specified in the respective configuration file of this module, pam_unix2.conf. The pam_limits module loads the file /etc/security/limits.conf,

which may define limits on the use of certain system resources. The `session` modules are called a second time when user logs out.

16.3 Configuration of PAM Modules

Some of the PAM modules are configurable. The corresponding configuration files are located in `/etc/security`. This section briefly describes the configuration files relevant to the `sshd` example—`pam_unix2.conf`, `pam_env.conf`, `pam_pwcheck.conf`, and `limits.conf`.

16.3.1 pam_unix2.conf

The traditional password-based authentication method is controlled by the PAM module `pam_unix2`. It can read the necessary data from `/etc/passwd`, `/etc/shadow`, NIS maps, NIS+ tables, or an LDAP database. The behavior of this module can be influenced by configuring the PAM options of the individual application itself or globally by editing `/etc/security/pam_unix2.conf`. A very basic configuration file for the module is shown in [Example 16.6, “pam_unix2.conf”](#) (page 306).

Example 16.6 `pam_unix2.conf`

```
auth:    nullok
account:
password:      nullok
session:       none
```

The `nullok` option for module types `auth` and `password` specifies that empty passwords are permitted for the corresponding type of account. Users are also allowed to change passwords for their accounts. The `none` option for the module type `session` specifies that no messages are logged on its behalf (this is the default). Learn about additional configuration options from the comments in the file itself and from the manual page `pam_unix2(8)`.

16.3.2 pam_env.conf

This file can be used to define a standardized environment for users that is set whenever the `pam_env` module is called. With it, preset environment variables using the following syntax:

```
VARIABLE [DEFAULT=[value]] [ OVERRIDE=[value]]
```

VARIABLE

Name of the environment variable to set.

```
[DEFAULT=[value]]
```

Default value the administrator wants set.

```
[ OVERRIDE=[value]]
```

Values that may be queried and set by pam_env, overriding the default value.

A typical example of how pam_env can be used is the adaptation of the DISPLAY variable, which is changed whenever a remote login takes place. This is shown in [Example 16.7, “pam_env.conf”](#) (page 307).

Example 16.7 pam_env.conf

```
REMOTEHOST      DEFAULT=localhost OVERRIDE=@{PAM_RHOST}
DISPLAY        DEFAULT=${REMOTEHOST}:0.0 OVERRIDE=${DISPLAY}
```

The first line sets the value of the REMOTEHOST variable to localhost, which is used whenever pam_env cannot determine any other value. The DISPLAY variable in turn contains the value of REMOTEHOST. Find more information in the comments in the file /etc/security/pam_env.conf.

16.3.3 pam_pwcheck.conf

This configuration file is for the pam_pwcheck module, which reads options from it for all password type modules. Settings stored in this file take precedence over the PAM settings of an individual application. If application-specific settings have not been defined, the application uses the global settings. [Example 16.8, “pam_pwcheck.conf”](#) (page 307) tells pam_pwcheck to allow empty passwords and modification of passwords. More options for the module are mentioned in the file /etc/security/pam_pwcheck.conf.

Example 16.8 pam_pwcheck.conf

```
password:    nullok
```

16.3.4 limits.conf

System limits can be set on a user or group basis in the file `limits.conf`, which is read by the `pam_limits` module. The file allows you to set hard limits, which may not be exceeded at all, and soft limits, which may be exceeded temporarily. To learn about the syntax and the available options, read the comments included in the file.

16.4 For More Information

In the directory `/usr/share/doc/packages/pam` of your installed system, find the following additional documentation:

READMEs

In the top level of this directory, there are some general README files. The sub-directory `modules` holds README files about the available PAM modules.

The Linux-PAM System Administrators' Guide

This document includes everything that a system administrator should know about PAM. It discusses a range of topics, from the syntax of configuration files to the security aspects of PAM. The document is available as a PDF file, in HTML format, and as plain text.

The Linux-PAM Module Writers' Manual

This document summarizes the topic from the developer's point of view, with information about how to write standard-compliant PAM modules. It is available as a PDF file, in HTML format, and as plain text.

The Linux-PAM Application Developers' Guide

This document includes everything needed by an application developer who wants to use the PAM libraries. It is available as a PDF file, in HTML format, and as plain text.

Thorsten Kukuk has developed a number of PAM modules for SUSE Linux and made some information available about them at <http://www.suse.de/~kukuk/pam/>.

Virtualization with Xen

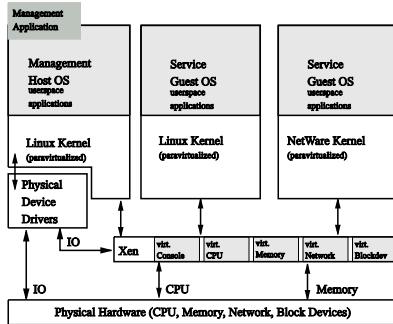
Xen makes it possible to run several Linux systems on one physical machine. The hardware for the different systems is provided virtually. This chapter gives an overview of the possibilities and limitations of this technology. Sections about installing, configuring, and running Xen complete this introduction.

Virtual machines commonly need to emulate the hardware a system needs. The disadvantage is that the emulated hardware is much slower than the real silicon. Xen has a different approach. It restricts emulation to as few parts as possible. To achieve this, Xen uses *paravirtualization*. This is a technique that presents virtual machines similarly, but not identically to the underlying hardware. Therefore, host and guest operating systems are adapted on kernel level. The user space remains unchanged. Xen controls the hardware with a hypervisor and a controlling guest, also called Domain-0. These provide all needed virtualized block and network devices. The guest systems use these virtual block and network devices to run the system and connect to other guests or the local network. When several physical machines running Xen are configured in a way that the virtual block and network devices are available, it is also possible to migrate a guest system from one piece of hardware to another while running. Originally, Xen was developed to run up to 100 guest systems on one computer, but this number depends strongly on the system requirements of the running guest systems, especially the memory consumption.

To limit the CPU utilization, the Xen hypervisor offers three different schedulers. The scheduler also may be changed while running the guest system, making it is possible to change the priority of the running guest system. On a higher level, migrating a guest may also be used to adjust the available CPU power.

The Xen virtualization system also has some drawbacks regarding the supported hardware. Several closed source drivers, like those from Nvidia or ATI, do not work as expected. In these cases, you must use the open source drivers if available, even if they do not support the full capabilities of the chips. Also several WLAN chips and Cardbus bridges are not supported when using Xen. In version 2, Xen does not support PAE (physical address extension), which means that it does not support more than 4 GB of memory. ACPI is not supported. Power management and other modes that depend on ACPI do not work. Another limitation of Xen is that it is currently not possible to just boot a block device. To boot, it is always necessary to have the correct kernel and initrd available in Domain-0.

Figure 17.1 Xen Overview



17.1 Xen Installation

The installation procedure of Xen involves the setup of a Domain-0 domain and the installation of Xen guests. First, make sure that the needed packages are installed. These

are `python`, `bridge-utils`, `xen`, `xen-tools`, `xen-tools-ioemu`, and a `kernel-xen` package. When selecting Xen during installation, Xen is added to the GRUB configuration. For other cases, make an entry in `boot/grub/menu.lst`. This entry should be similar to the following:

```
title Xen3
  kernel  (hd0,0)/boot/xen.gz
  module  (hd0,0)/boot/vmlinuz-xen <parameters>
  module  (hd0,0)/boot/initrd-xen
```

Replace `(hd0,0)` with the partition that holds your `/boot` directory. See also [Chapter 9, *The Boot Loader*](#) (page 193). Replace `<parameters>` with the parameters normally used to boot a Linux kernel. Then reboot into Xen mode. This boots the Xen hypervisor and a slightly changed Linux kernel as Domain-0 that runs most of the hardware. Apart from the exceptions already mentioned, everything should work as normal.

17.2 Domain Installation

Each guest domain must be installed individually. For this, run the YaST module Virtual Machine Installation (Xen) in the Software group. In the following interface, provide all information that is needed to run the installer. This information is split into four categories:

Options

Here, set the name of the guest domain, its memory resource, and boot options for the installer.

Disks

Choose to create a file system image or a real physical partition. The file system images are stored in the directory `/var/lib/xen/images`. Make sure that you have sufficient disk space in this directory.

Operating System

The operating system that should be used to install the guest domain. This system is selected in the YaST module Installation Source and cannot be set in this workflow.

Network

Only bridged networking is supported by this module. Add the number of virtual network cards you need.

The next dialog starts the installation system after doing several setup tasks. This system is identical to a standard installation in text mode as described in Section “YaST in Text Mode” (Chapter 2, *System Configuration with YaST*, ↑Start-Up).

If you need to change the configuration of a guest domain, you must do that directly in the configuration file. This file is located in `/etc/xen` and has a name identical to the name of the guest domain.

17.3 Starting and Controlling Xen Domains with `xm`

Xen automatically reduces the amount of memory of Domain-0 to meet the requirements of the newly started guest domain. If there is not enough memory available, the guest is not started. You can always check the available memory of Domain-0 with the command `free`.

It is always possible to detach a console or reattach it from another terminal. To detach, use `[Ctrl] + [1]`. To reattach, first check the ID of the needed guest with `xm list` and attach to that ID with `xm console ID`.

The `xm` tool of Xen has many possible parameters. View a list with a short explanation by entering `xm help`. [Table 17.1, “`xm Commands`”](#) (page 312) provides some of the most important commands as a starting point.

Table 17.1 *xm Commands*

<code>xm help</code>	Print a list of commands that are available for the <code>xm</code> tool.
<code>xm console <i>ID</i></code>	Connect to the first console (tty1) of the guest with ID <i>ID</i> .
<code>xm mem-set <i>ID Mem</i></code>	Set the memory size of the domain with ID <i>ID</i> to <i>Mem</i> in MB.

<code>xm create <i>domname</i> [-c]</code>	Start the domain with configuration file <i>domname</i> . The optional <code>-c</code> links the current terminal to the first tty of the new guest.
<code>xm shutdown <i>ID</i></code>	Do a normal shutdown of the guest with ID <i>ID</i> .
<code>xm destroy <i>ID</i></code>	Terminate the guest with ID <i>ID</i> immediately.
<code>xm list</code>	Print a list of all running domains with their respective ID, memory, and CPU time values.
<code>xm info</code>	Display information about the Xen host, including CPU and memory information.

17.4 Troubleshooting

This section provides some hints about how to solve common problems. It is not meant as an exhaustive step-by-step instruction, but should help get started with solving some issues.

There are networking issues in Xen3.

The concept of networking has changed considerably from Xen2 to Xen3. Domain-0 is no longer directly connected to the bridge to prevent blocking of the bridge. Unfortunately, the initialization scripts of the system cannot handle the current configuration. To restart the network, run `/etc/init.d/xend restart`.

I need to do a file system check.

If the file system check did not work automatically, you may do it manually from Domain-0. Shut down your guest and run `fsck` on the image while it is not mounted. If `fsck` complains that the file system is mounted, check your mounts with the command `mount`.

DHCP does not get IP addresses.

DHCP needs several iptables kernel modules to run. Either those have not been installed or you updated your kernel and forgot to update the kernel modules in the guest system.

There is a problem booting the hypervisor and the messages go away too quickly
Connect your Xen machine to another workstation with a serial nullmodem cable.
Then, on the Xen side, add the following parameter to the line

```
kernel (hd0,0)/boot/xen.gz com1=115200,8n1
```

Before you boot Xen, start a terminal program on your workstation. As an example, this may be

```
screen /dev/ttys0 115200
```

Change the device and speed according to your needs.

17.5 For More Information

More information about Xen can be found on the following Web sites:

- `/usr/share/doc/packages/xen/user/html/index.html`—Official information for Xen users. It requires the package `xen-doc-html`.
- `/usr/share/doc/packages/xen/interface/html/index.html`—Some more technical interface documentation. It also requires the package `xen-doc-html`.
- <http://www.cl.cam.ac.uk/Research/SRG/netos/xen/index.html>—Xen home page with many different documentation links.
- <http://lists.xensource.com/>—Several mailing lists about Xen.
- <http://wiki.xensource.com/xenwiki>—Xen wiki for the open source community.

Part IV. Services

18

Basic Networking

Linux offers the necessary networking tools and features for integration into all types of network structures. The customary Linux protocol, TCP/IP, has various services and special features, which are discussed here. Network access using a network card, modem, or other device can be configured with YaST. Manual configuration is also possible. Only the fundamental mechanisms and the relevant network configuration files are discussed in this chapter.

Linux and other Unix operating systems use the TCP/IP protocol. It is not a single network protocol, but a family of network protocols that offer various services. The protocols listed in [Table 18.1, “Several Protocols in the TCP/IP Protocol Family”](#) (page 318) are provided for the purpose of exchanging data between two machines via TCP/IP. Networks combined by TCP/IP, comprising a worldwide network are also referred to, in their entirety, as “the Internet.”

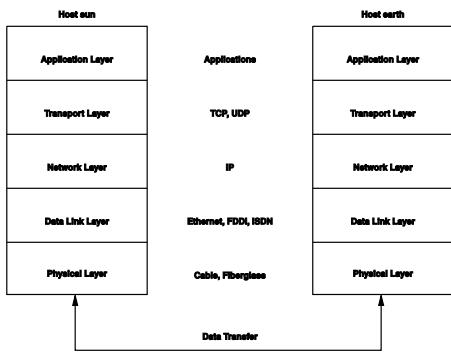
RFC stands for *Request for Comments*. RFCs are documents that describe various Internet protocols and implementation procedures for the operating system and its applications. The RFC documents describe the setup of Internet protocols. To expand your knowledge about any of the protocols, refer to the appropriate RFC documents. They are available online at <http://www.ietf.org/rfc.html>.

Table 18.1 Several Protocols in the TCP/IP Protocol Family

Protocol	Description
TCP	Transmission Control Protocol: A connection-oriented secure protocol. The data to transmit is first sent by the application as a stream of data then converted by the operating system to the appropriate format. The data arrives at the respective application on the destination host in the original data stream format in which it was initially sent. TCP determines whether any data has been lost during the transmission and that there is no mix-up. TCP is implemented wherever the data sequence matters.
UDP	User Datagram Protocol: A connectionless, insecure protocol. The data to transmit is sent in the form of packets generated by the application. The order in which the data arrives at the recipient is not guaranteed and data loss is a possibility. UDP is suitable for record-oriented applications. It features a smaller latency period than TCP.
ICMP	Internet Control Message Protocol: Essentially, this is not a protocol for the end user, but a special control protocol that issues error reports and can control the behavior of machines participating in TCP/IP data transfer. In addition, it provides a special echo mode that can be viewed using the program ping.
IGMP	Internet Group Management Protocol: This protocol controls machine behavior when implementing IP multicast.

As shown in [Figure 18.1, “Simplified Layer Model for TCP/IP”](#) (page 319), data exchange takes place in different layers. The actual network layer is the insecure data transfer via IP (Internet protocol). On top of IP, TCP (transmission control protocol) guarantees, to a certain extent, security of the data transfer. The IP layer is supported by the underlying hardware-dependent protocol, such as ethernet.

Figure 18.1 Simplified Layer Model for TCP/IP



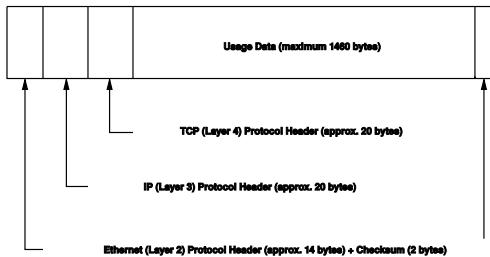
The diagram provides one or two examples for each layer. The layers are ordered according to *abstraction levels*. The lowest layer is very close to the hardware. The uppermost layer, however, is almost a complete abstraction from the hardware. Every layer has its own special function. The special functions of each layer are mostly implicit in their description. The data link and physical layers represent the physical network used, such as ethernet.

Almost all hardware protocols work on a packet-oriented basis. The data to transmit is packaged in *packets*, because it cannot be sent all at once. The maximum size of a TCP/IP packet is approximately 64 KB. Packets are normally quite a bit smaller, because the network hardware can be a limiting factor. The maximum size of a data packet on an ethernet is about fifteen hundred bytes. The size of a TCP/IP packet is limited to this amount when the data is sent over an ethernet. If more data is transferred, more data packets need to be sent by the operating system.

For the layers to serve their designated functions, additional information regarding each layer must be saved in the data packet. This takes place in the *header* of the packet. Every layer attaches a small block of data, called the protocol header, to the front of each emerging packet. A sample TCP/IP data packet traveling over an ethernet cable is illustrated in [Figure 18.2, “TCP/IP Ethernet Packet”](#) (page 320). The proof sum is

located at the end of the packet, not at the beginning. This simplifies things for the network hardware.

Figure 18.2 TCP/IP Ethernet Packet



When an application sends data over the network, the data passes through each layer, all implemented in the Linux kernel except the physical layer. Each layer is responsible for preparing the data so it can be passed to the next layer. The lowest layer is ultimately responsible for sending the data. The entire procedure is reversed when data is received. Like the layers of an onion, in each layer the protocol headers are removed from the transported data. Finally, the transport layer is responsible for making the data available for use by the applications at the destination. In this manner, one layer only communicates with the layer directly above or below it. For applications, it is irrelevant whether data is transmitted via a 100 MBit/s FDDI network or via a 56-kbit/s modem line. Likewise, it is irrelevant for the data line which kind of data is transmitted, as long as packets are in the correct format.

18.1 IP Addresses and Routing

The discussion in this section is limited to IPv4 networks. For information about IPv6 protocol, the successor to IPv4, refer to [Section 18.2, “IPv6—The Next Generation Internet”](#) (page 323).

18.1.1 IP Addresses

Every computer on the Internet has a unique 32-bit address. These 32 bits (or 4 bytes) are normally written as illustrated in the second row in [Example 18.1, “Writing IP Addresses”](#) (page 321).

Example 18.1 Writing IP Addresses

IP Address (binary):	11000000 10101000 00000000 00010100
IP Address (decimal):	192. 168. 0. 20

In decimal form, the four bytes are written in the decimal number system, separated by periods. The IP address is assigned to a host or a network interface. It cannot be used anywhere else in the world. There are exceptions to this rule, but these are not relevant in the following passages.

The points in IP addresses indicate the hierarchical system. Until the 1990s, IP addresses were strictly categorized in classes. However, this system has proven too inflexible and was discontinued. Now, *classless routing* (CIDR, classless interdomain routing) is used.

18.1.2 Netmasks and Routing

Netmasks are used to define the address range of a subnetwork. If two hosts are in the same subnetwork, they can reach each other directly, if they are not in the same subnetwork, they need the address of a gateway that handles all the traffic between the subnetwork and the rest of the world. To check if two IP addresses are in the same subnet, simply “AND” both addresses with the netmask. If the result is identical, both IP addresses are in the same local network. If there are differences, the remote IP address, and thus the remote interface, can only be reached over a gateway.

To understand how the netmask works, look at [Example 18.2, “Linking IP Addresses to the Netmask”](#) (page 322). The netmask consists of 32 bits that identify how much of

an IP address belongs to the network. All those bits that are 1 mark the corresponding bit in the IP address as belonging to the network. All bits that are 0 mark bits inside the subnetwork. This means that the more bits are 1, the smaller the subnetwork is. Because the netmask always consists of several successive 1 bits, it is also possible to just count the number of bits in the netmask. In [Example 18.2, “Linking IP Addresses to the Netmask”](#) (page 322) the first net with 24 bits could also be written as 192.168.0.0/24.

Example 18.2 *Linking IP Addresses to the Netmask*

IP address (192.168.0.20):	11000000 10101000 00000000 00010100
Netmask (255.255.255.0):	11111111 11111111 11111111 00000000

Result of the link:	11000000 10101000 00000000 00000000
In the decimal system:	192. 168. 0. 0

IP address (213.95.15.200):	11010101 10111111 00001111 11001000
Netmask (255.255.255.0):	11111111 11111111 11111111 00000000

Result of the link:	11010101 10111111 00001111 00000000
In the decimal system:	213. 95. 15. 0

To give another example: all machines connected with the same ethernet cable are usually located in the same subnetwork and are directly accessible. Even when the subnet is physically divided by switches or bridges, these hosts can still be reached directly.

IP addresses outside the local subnet can only be reached if a gateway is configured for the target network. In the most common case, there is only one gateway that handles all traffic that is external. However, it is also possible to configure several gateways for different subnets.

If a gateway has been configured, all external IP packets are sent to the appropriate gateway. This gateway then attempts to forward the packets in the same manner—from host to host—until it reaches the destination host or the packet's TTL (time to live) expires.

Table 18.2 *Specific Addresses*

Address Type	Description
Base Network Address	This is the netmask AND any address in the network, as shown in Example 18.2, “Linking IP Addresses to the Netmask”

Address Type	Description
	(page 322) under Result. This address cannot be assigned to any hosts.
Broadcast Address	This basically says, “Access all hosts in this subnetwork.” To generate this, the netmask is inverted in binary form and linked to the base network address with a logical OR. The above example therefore results in 192.168.0.255. This address cannot be assigned to any hosts.
Local Host	The address 127.0.0.1 is assigned to the “loopback device” on each host. A connection can be set up to your own machine with this address.

Because IP addresses must be unique all over the world, you cannot just select random addresses. There are three address domains to use if you want to set up a private IP-based network. These cannot get any connection from the rest of the Internet, because they cannot be transmitted over the Internet. These address domains are specified in RFC 1597 and listed in [Table 18.3, “Private IP Address Domains”](#) (page 323).

Table 18.3 Private IP Address Domains

Network/Netmask	Domain
10.0.0.0/255.0.0.0	10.x.x.x
172.16.0.0/255.240.0.0	172.16.x.x – 172.31.x.x
192.168.0.0/255.255.0.0	192.168.x.x

18.2 IPv6—The Next Generation Internet

Due to the emergence of the WWW (World Wide Web), the Internet has experienced explosive growth with an increasing number of computers communicating via TCP/IP

in the past fifteen years. Since Tim Berners-Lee at CERN (<http://public.web.cern.ch>) invented the WWW in 1990, the number of Internet hosts has grown from a few thousand to about a hundred million.

As mentioned, an IPv4 address consists of only 32 bits. Also, quite a few IP addresses are lost—they cannot be used due to the way in which networks are organized. The number of addresses available in your subnet is two to the power of the number of bits, minus two. A subnetwork has, for example, 2, 6, or 14 addresses available. To connect 128 hosts to the Internet, for example, you need a subnetwork with 256 IP addresses, from which only 254 are usable, because two IP addresses are needed for the structure of the subnetwork itself: the broadcast and the base network address.

Under the current IPv4 protocol, DHCP or NAT (network address translation) are the typical mechanisms used to circumvent the potential address shortage. Combined with the convention to keep private and public address spaces separate, these methods can certainly mitigate the shortage. The problem with them lies in their configuration, which is a chore to set up and a burden to maintain. To set up a host in an IPv4 network, you need a number of address items, such as the host's own IP address, the subnetmask, the gateway address, and maybe a name server address. All these items need to be known and cannot be derived from somewhere else.

With IPv6, both the address shortage and the complicated configuration should be a thing of the past. The following sections tell more about the improvements and benefits brought by IPv6 and about the transition from the old protocol to the new one.

18.2.1 Advantages

The most important and most visible improvement brought by the new protocol is the enormous expansion of the available address space. An IPv6 address is made up of 128 bit values instead of the traditional 32 bits. This provides for as many as several quadrillion IP addresses.

However, IPv6 addresses are not only different from their predecessors with regard to their length. They also have a different internal structure that may contain more specific information about the systems and the networks to which they belong. More details about this are found in [Section 18.2.2, “Address Types and Structure”](#) (page 326).

The following is a list of some other advantages of the new protocol:

Autoconfiguration

IPv6 makes the network “plug and play” capable, which means that a newly set up system integrates into the (local) network without any manual configuration. The new host uses its automatic configuration mechanism to derive its own address from the information made available by the neighboring routers, relying on a protocol called the *neighbor discovery* (ND) protocol. This method does not require any intervention on the administrator's part and there is no need to maintain a central server for address allocation—an additional advantage over IPv4, where automatic address allocation requires a DHCP server.

Mobility

IPv6 makes it possible to assign several addresses to one network interface at the same time. This allows users to access several networks easily, something that could be compared with the international roaming services offered by mobile phone companies: when you take your mobile phone abroad, the phone automatically logs in to a foreign service as soon as it enters the corresponding area, so you can be reached under the same number everywhere and are able to place an outgoing call just like in your home area.

Secure Communication

With IPv4, network security is an add-on function. IPv6 includes IPSec as one of its core features, allowing systems to communicate over a secure tunnel to avoid eavesdropping by outsiders on the Internet.

Backward Compatibility

Realistically, it would be impossible to switch the entire Internet from IPv4 to IPv6 at one time. Therefore, it is crucial that both protocols are able to coexist not only on the Internet, but also on one system. This is ensured by compatible addresses (IPv4 addresses can easily be translated into IPv6 addresses) and through the use of a number of tunnels. See [Section 18.2.3, “Coexistence of IPv4 and IPv6”](#) (page 330). Also, systems can rely on a *dual stack IP* technique to support both protocols at the same time, meaning that they have two network stacks that are completely separate, such that there is no interference between the two protocol versions.

Custom Tailored Services through Multicasting

With IPv4, some services, such as SMB, need to broadcast their packets to all hosts in the local network. IPv6 allows a much more fine-grained approach by enabling servers to address hosts through *multicasting*—by addressing a number of hosts as parts of a group (which is different from addressing all hosts through *broadcasting*)

or each host individually through *unicasting*). Which hosts are addressed as a group may depend on the concrete application. There are some predefined groups to address all name servers (the *all name servers multicast group*), for example, or all routers (the *all routers multicast group*).

18.2.2 Address Types and Structure

As mentioned, the current IP protocol is lacking in two important aspects: there is an increasing shortage of IP addresses and configuring the network and maintaining the routing tables is becoming a more complex and burdensome task. IPv6 solves the first problem by expanding the address space to 128 bits. The second one is countered by introducing a hierarchical address structure, combined with sophisticated techniques to allocate network addresses, as well as *multihoming* (the ability to assign several addresses to one device, giving access to several networks).

When dealing with IPv6, it is useful to know about three different types of addresses:

Unicast

Addresses of this type are associated with exactly one network interface. Packets with such an address are delivered to only one destination. Accordingly, unicast addresses are used to transfer packets to individual hosts on the local network or the Internet.

Multicast

Addresses of this type relate to a group of network interfaces. Packets with such an address are delivered to all destinations that belong to the group. Multicast addresses are mainly used by certain network services to communicate with certain groups of hosts in a well-directed manner.

Anycast

Addresses of this type are related to a group of interfaces. Packets with such an address are delivered to the member of the group that is closest to the sender, according to the principles of the underlying routing protocol. Anycast addresses are used to make it easier for hosts to find out about servers offering certain services in the given network area. All servers of the same type have the same anycast address. Whenever a host requests a service, it receives a reply from the server with the closest location, as determined by the routing protocol. If this server should fail for some reason, the protocol automatically selects the second closest server, then the third one, and so forth.

An IPv6 address is made up of eight four-digit fields, each representing 16 bits, written in hexadecimal notation. They are also separated by colons (:). Any leading zero bytes within a given field may be dropped, but zeros within the field or at its end may not. Another convention is that more than four consecutive zero bytes may be collapsed into a double colon. However, only one such :: is allowed per address. This kind of shorthand notation is shown in [Example 18.3, “Sample IPv6 Address”](#) (page 327), where all three lines represent the same address.

Example 18.3 *Sample IPv6 Address*

```
fe80 : 0000 : 0000 : 0000 : 0000 : 10 : 1000 : 1a4  
fe80 :     0 :     0 :     0 :     0 : 10 : 1000 : 1a4  
fe80 :                   : 10 : 1000 : 1a4
```

Each part of an IPv6 address has a defined function. The first bytes form the prefix and specify the type of address. The center part is the network portion of the address, but it may be unused. The end of the address forms the host part. With IPv6, the netmask is defined by indicating the length of the prefix after a slash at the end of the address. An address, as shown in [Example 18.4, “IPv6 Address Specifying the Prefix Length”](#) (page 327), contains the information that the first 64 bits form the network part of the address and the last 64 form its host part. In other words, the 64 means that the netmask is filled with 64 1-bit values from the left. Just like with IPv4, the IP address is combined with AND with the values from the netmask to determine whether the host is located in the same subnetwork or in another one.

Example 18.4 *IPv6 Address Specifying the Prefix Length*

```
fe80::10:1000:1a4/64
```

IPv6 knows about several predefined types of prefixes. Some of these are shown in [Table 18.4, “Various IPv6 Prefixes”](#) (page 327).

Table 18.4 *Various IPv6 Prefixes*

Prefix (hex)	Definition
00	IPv4 addresses and IPv4 over IPv6 compatibility addresses. These are used to maintain compatibility with IPv4. Their use still requires a router able to translate IPv6 packets into IPv4 packets. Several special addresses, such as the one for the loopback device, have this prefix as well.

Prefix (hex)	Definition
2 or 3 as the first digit	Aggregatable global unicast addresses. As is the case with IPv4, an interface can be assigned to form part of a certain subnetwork. Currently, there are the following address spaces: 2001::/16 (production quality address space) and 2002::/16 (6to4 address space).
fe80::/10	Link-local addresses. Addresses with this prefix should not be routed and should therefore only be reachable from within the same subnetwork.
fec0::/10	Site-local addresses. These may be routed, but only within the network of the organization to which they belong. In effect, they are the IPv6 equivalent of the current private network address space, such as 10.x.x.x.
ff	These are multicast addresses.

A unicast address consists of three basic components:

Public Topology

The first part (which also contains one of the prefixes mentioned above) is used to route packets through the public Internet. It includes information about the company or institution that provides the Internet access.

Site Topology

The second part contains routing information about the subnetwork to which to deliver the packet.

Interface ID

The third part identifies the interface to which to deliver the packet. This also allows for the MAC to form part of the address. Given that the MAC is a globally unique, fixed identifier coded into the device by the hardware maker, the configuration procedure is substantially simplified. In fact, the first 64 address bits are consolidated to form the EUI-64 token, with the last 48 bits taken from the MAC, and the remaining 24 bits containing special information about the token type. This also makes it possible to assign an EUI-64 token to interfaces that do not have a MAC, such as those based on PPP or ISDN.

On top of this basic structure, IPv6 distinguishes between five different types of unicast addresses:

:: (unspecified)

This address is used by the host as its source address when the interface is initialized for the first time—when the address cannot yet be determined by other means.

::1 (loopback)

The address of the loopback device.

IPv4 Compatible Addresses

The IPv6 address is formed by the IPv4 address and a prefix consisting of 96 zero bits. This type of compatibility address is used for tunneling (see [Section 18.2.3, “Coexistence of IPv4 and IPv6”](#) (page 330)) to allow IPv4 and IPv6 hosts to communicate with others operating in a pure IPv4 environment.

IPv4 Addresses Mapped to IPv6

This type of address specifies a pure IPv4 address in IPv6 notation.

Local Addresses

There are two address types for local use:

link-local

This type of address can only be used in the local subnetwork. Packets with a source or target address of this type should not be routed to the Internet or other subnetworks. These addresses contain a special prefix (`fe80 :: /10`) and the interface ID of the network card, with the middle part consisting of zero bytes. Addresses of this type are used during automatic configuration to communicate with other hosts belonging to the same subnetwork.

site-local

Packets with this type of address may be routed to other subnetworks, but not to the wider Internet—they must remain inside the organization's own network. Such addresses are used for intranets and are an equivalent of the private address space defined by IPv4. They contain a special prefix (`fec0 :: /10`), the interface ID, and a 16 bit field specifying the subnetwork ID. Again, the rest is filled with zero bytes.

As a completely new feature introduced with IPv6, each network interface normally gets several IP addresses, with the advantage that several networks can be accessed through the same interface. One of these networks can be configured completely auto-

matically using the MAC and a known prefix with the result that all hosts on the local network can be reached as soon as IPv6 is enabled (using the link-local address). With the MAC forming part of it, any IP address used in the world is unique. The only variable parts of the address are those specifying the *site topology* and the *public topology*, depending on the actual network in which the host is currently operating.

For a host to go back and forth between different networks, it needs at least two addresses. One of them, the *home address*, not only contains the interface ID but also an identifier of the home network to which it normally belongs (and the corresponding prefix). The home address is a static address and, as such, it does not normally change. Still, all packets destined to the mobile host can be delivered to it, regardless of whether it operates in the home network or somewhere outside. This is made possible by the completely new features introduced with IPv6, such as *stateless autoconfiguration* and *neighbor discovery*. In addition to its home address, a mobile host gets one or more additional addresses that belong to the foreign networks where it is roaming. These are called *care-of* addresses. The home network has a facility that forwards any packets destined to the host when it is roaming outside. In an IPv6 environment, this task is performed by the *home agent*, which takes all packets destined to the home address and relays them through a tunnel. On the other hand, those packets destined to the care-of address are directly transferred to the mobile host without any special detours.

18.2.3 Coexistence of IPv4 and IPv6

The migration of all hosts connected to the Internet from IPv4 to IPv6 is a gradual process. Both protocols will coexist for some time to come. The coexistence on one system is guaranteed where there is a *dual stack* implementation of both protocols. That still leaves the question of how an IPv6 enabled host should communicate with an IPv4 host and how IPv6 packets should be transported by the current networks, which are predominantly IPv4 based. The best solutions offer tunneling and compatibility addresses (see [Section 18.2.2, “Address Types and Structure”](#) (page 326)).

IPv6 hosts that are more or less isolated in the (worldwide) IPv4 network can communicate through tunnels: IPv6 packets are encapsulated as IPv4 packets to move them across an IPv4 network. Such a connection between two IPv4 hosts is called a *tunnel*. To achieve this, packets must include the IPv6 destination address (or the corresponding prefix) as well as the IPv4 address of the remote host at the receiving end of the tunnel. A basic tunnel can be configured manually according to an agreement between the hosts' administrators. This is also called *static tunneling*.

However, the configuration and maintenance of static tunnels is often too labor-intensive to use them for daily communication needs. Therefore, IPv6 provides for three different methods of *dynamic tunneling*:

6over4

IPv6 packets are automatically encapsulated as IPv4 packets and sent over an IPv4 network capable of multicasting. IPv6 is tricked into seeing the whole network (Internet) as a huge local area network (LAN). This makes it possible to determine the receiving end of the IPv4 tunnel automatically. However, this method does not scale very well and is also hampered by the fact that IP multicasting is far from widespread on the Internet. Therefore, it only provides a solution for smaller corporate or institutional networks where multicasting can be enabled. The specifications for this method are laid down in RFC 2529.

6to4

With this method, IPv4 addresses are automatically generated from IPv6 addresses, enabling isolated IPv6 hosts to communicate over an IPv4 network. However, a number of problems have been reported regarding the communication between those isolated IPv6 hosts and the Internet. The method is described in RFC 3056.

IPv6 Tunnel Broker

This method relies on special servers that provide dedicated tunnels for IPv6 hosts. It is described in RFC 3053.

IMPORTANT: The 6bone Initiative

In the heart of the “old-time” Internet, there is already a globally distributed network of IPv6 subnets that are connected through tunnels. This is the *6bone* network (<http://www.6bone.net>), an IPv6 test environment that may be used by programmers and Internet providers who want to develop and offer IPv6-based services to gain the experience necessary to implement the new protocol. More information can be found on the project’s Internet site.

18.2.4 Configuring IPv6

To configure IPv6, you do not normally need to make any changes on the individual workstations. However, IPv6 support must be loaded. To do this, enter `modprobe ipv6` as root.

Because of the autoconfiguration concept of IPv6, the network card is assigned an address in the *link-local* network. Normally, no routing table management takes place on a workstation. The network routers can be queried by the workstation, using the *router advertisement protocol*, for what prefix and gateways should be implemented. The radvd program can be used to set up an IPv6 router. This program informs the workstations which prefix to use for the IPv6 addresses and which routers. Alternatively, use zebra for automatic configuration of both addresses and routing.

Consult the ifup(8) man page to get information about how to set up various types of tunnels using the /etc/sysconfig/network files.

18.2.5 For More Information

The above overview does not cover the topic of IPv6 comprehensively. For a more in-depth look at the new protocol, refer to the following online documentation and books:

<http://www.ngnet.it/e/cosa-ipv6.php>

An article series providing a well-written introduction to the basics of IPv6. A good primer on the topic.

<http://www.bieringer.de/linux/IPv6/>

Here, find the Linux IPv6-HOWTO and many links related to the topic.

<http://www.6bone.net/>

Visit this site if you want to join a tunneled IPv6 network.

<http://www.ipv6.org/>

The starting point for everything about IPv6.

RFC 2640

The fundamental RFC about IPv6.

IPv6 Essentials

A book describing all the important aspects of the topic is *IPv6 Essentials* by Silvia Hagen (ISBN 0-596-00125-8).

18.3 Name Resolution

DNS assists in assigning an IP address to one or more names and assigning a name to an IP address. In Linux, this conversion is usually carried out by a special type of software known as bind. The machine that takes care of this conversion is called a *name server*. The names make up a hierarchical system in which each name component is separated by dots. The name hierarchy is, however, independent of the IP address hierarchy described above.

Consider a complete name, such as `earth.example.com`, written in the format `hostname.domain`. A full name, referred to as a *fully qualified domain name* (FQDN), consists of a hostname and a domain name (`example.com`). The latter also includes the *top level domain* or TLD (`.com`).

TLD assignment has become quite confusing for historical reasons. Traditionally, three-letter domain names are used in the USA. In the rest of the world, the two-letter ISO national codes are the standard. In addition to that, longer TLDs were introduced in 2000 that represent certain spheres of activity (for example, `.info`, `.name`, `.museum`).

In the early days of the Internet (before 1990), the file `/etc/hosts` was used to store the names of all the machines represented over the Internet. This quickly proved to be impractical in the face of the rapidly growing number of computers connected to the Internet. For this reason, a decentralized database was developed to store the hostnames in a widely distributed manner. This database, similar to the name server, does not have the data pertaining to all hosts in the Internet readily available, but can dispatch requests to other name servers.

The top of the hierarchy is occupied by *root name servers*. These root name servers manage the top level domains and are run by the Network Information Center (NIC). Each root name server knows about the name servers responsible for a given top level domain. Information about top level domain NICs is available at <http://www.internic.net>.

DNS can do more than just resolve hostnames. The name server also knows which host is receiving e-mails for an entire domain—the *mail exchanger (MX)*.

For your machine to resolve an IP address, it must know about at least one name server and its IP address. Easily specify such a name server with the help of YaST. If you have a modem dial-up connection, you may not need to configure a name server manually

at all. The dial-up protocol provides the name server address as the connection is made. The configuration of name server access with SUSE Linux is described in [Chapter 20, “The Domain Name System”](#) (page 367).

The protocol `whois` is closely related to DNS. With this program, quickly find out who is responsible for any given domain.

18.4 Configuring a Network Connection with YaST

There are many supported networking types on Linux. Most of them use different device names and the configuration files are spread over several locations in the file system. For a detailed overview of the aspects of manual network configuration, see [Section 18.6, “Configuring a Network Connection Manually”](#) (page 348).

During installation, YaST can be used to configure automatically all interfaces that have been detected. Additional hardware can be configured any time after installation in the installed system. The following sections describe the network configuration for all types of network connections supported by SUSE Linux.

18.4.1 Configuring the Network Card with YaST

After starting the module, YaST displays a general network configuration dialog. Choose whether to use YaST or NetworkManager to manage all your network devices. To use NetworkManager, check *User Controlled with NetworkManager*. Find detailed information about NetworkManager in [Section 18.5, “Managing Network Connections with NetworkManager”](#) (page 345). If you want to configure your network in the traditional way with the YaST, check *Traditional Method with ifup*.

The upper part of the traditional configuration shows a list with all the network cards available for configuration. Any card properly detected is listed with its name. Devices that could not be detected may be configured using *Add* as described in [Section “Manual Configuration of an Undetected Network Card”](#) (page 335). Configure a new network card or change an existing configuration.

Manual Configuration of an Undetected Network Card

Configuring a network card that was not detected includes the following items:

Network Configuration

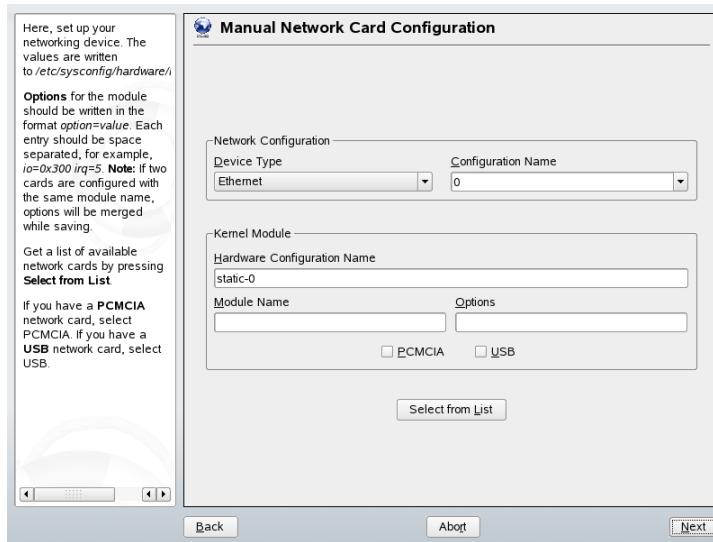
Set the device type of the interface from the available options and the configuration name. Information about the naming conventions for configuration names is available in the `getcfg(8)` man page.

Kernel Module

Hardware Configuration Name specifies the name of the `/etc/sysconfig/hardware/hwcfg-*` file containing the hardware settings of your network card. This contains the name of the suitable kernel module as well as the needed options to initialize the hardware. Usually, YaST proposes useful names for PCMCIA and USB hardware. For other hardware, `hwcfg-static-0` usually only makes sense if the card is configured with the configuration name 0.

If the network card is a PCMCIA or USB device, activate the respective check boxes and exit this dialog with *Next*. Otherwise, select your network card model from *Select from List*. YaST then automatically selects the suitable kernel module for the card. Exit this dialog with *Next*.

Figure 18.3 Configuration of the Network Card



Setting the Network Address

Set the device type of the interface and the configuration name. Select the device type from those provided. Specify a configuration name according to your needs. Usually, the default settings are useful and can be accepted. Information about the naming conventions for configuration names is available in the `getcfg(8)` man page .

If you selected *Wireless* as the device type of the interface, configure the operating mode, the network name (ESSID), and the encryption in the next dialog, *Wireless Network Card Configuration*. Click *OK* to complete the configuration of your card. A detailed description of the configuration of WLAN cards is provided in [Section 34.1.3, “Configuration with YaST” \(page 608\)](#). For all other interface types, proceed with the network address setup:

Automatic Address Setup (via DHCP)

If your network includes a DHCP server, you can rely on it to set up your network address automatically. The option should also be used if you are using a DSL line but with no static IP assigned by the ISP. If you decide to use DHCP, configure the details after selecting *DHCP Client Options*. Specify whether the DHCP server should always honor broadcast requests and any identifier to use. By default, DHCP

servers use the card's hardware address to identify an interface. If you have a virtual host setup where different hosts communicate through the same interface, an identifier is necessary to distinguish them.

Static Address Setup

If you have a static address, enable that option. Then enter the address and subnet mask for your network. The preset subnet mask should match the requirements of a typical home network.

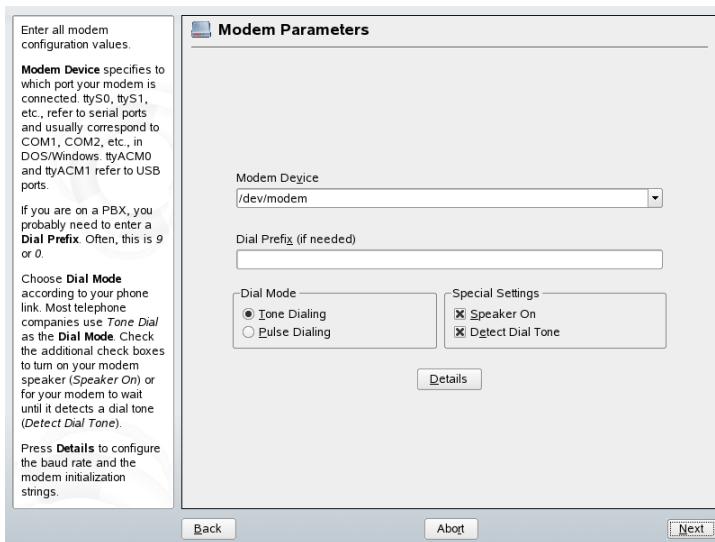
Leave this dialog by selecting *Next* or proceed to configure the hostname, name server, and routing details (see the sections on DNS Server (↑Start-Up) and Routing (↑Start-Up)).

Advanced enables you to specify more complex settings. Under *Detailed Settings*, use *User Controlled* to delegate the control over the network card from the administrator (`root`) to the normal user. For mobile operation, this allows the user to adapt changing network connections in a more flexible way, because he can control the activation or deactivation of the interface. The MTU (maximum transmission unit) and the type of *Device Activation* can also be set in this dialog.

18.4.2 Modem

In the YaST Control Center, access the modem configuration under *Network Devices*. If your modem was not automatically detected, open the dialog for manual configuration. In the dialog that opens, enter the interface to which the modem is connected under *Modem*.

Figure 18.4 Modem Configuration



If you are behind a private branch exchange (PBX), you may need to enter a dial prefix. This is often a zero. Consult the instructions that came with the PBX to find out. Also select whether to use tone or pulse dialing, whether the speaker should be on, and whether the modem should wait until it detects a dial tone. The last option should not be enabled if the modem is connected to an exchange.

Under *Details*, set the baud rate and the modem initialization strings. Only change these settings if your modem was not autodetected or if it requires special settings for data transmission to work. This is mainly the case with ISDN terminal adapters. Leave this dialog by clicking *OK*. To delegate control over the modem to the normal user without root permissions, activate *User Controlled*. In this way, a user without administrator permissions can activate or deactivate an interface. Under *Dial Prefix Regular Expression*, specify a regular expression. The *Dial Prefix* in KInternet, which can be modified by the normal user, must match this regular expression. If this field is left empty, the user cannot set a different *Dial Prefix* without administrator permissions.

In the next dialog, select the ISP (Internet service provider). To choose from a predefined list of ISPs operating in your country, select *Country*. Alternatively, click *New* to open a dialog in which to provide the data for your ISP. This includes a name for the dial-up connection and ISP as well as the login and password provided by your ISP. Enable *Always Ask for Password* to be prompted for the password each time you connect.

In the last dialog, specify additional connection options:

Dial on Demand

If you enable dial on demand, set at least one name server.

Modify DNS when Connected

This option is enabled by default, with the effect that the name server address is updated each time you connect to the Internet.

Automatically Retrieve DNS

If the provider does not transmit its domain name server after connecting, disable this option and enter the DNS data manually.

Stupid Mode

This option is enabled by default. With it, input prompts sent by the ISP's server are ignored to prevent them from interfering with the connection process.

External Firewall Interface and Restart Firewall

Selecting these options enables the SUSEfirewall2, which protects you from outside attacks for the duration of your Internet connection.

Idle Time-Out (seconds)

With this option, specify a period of network inactivity after which the modem disconnects automatically.

IP Details

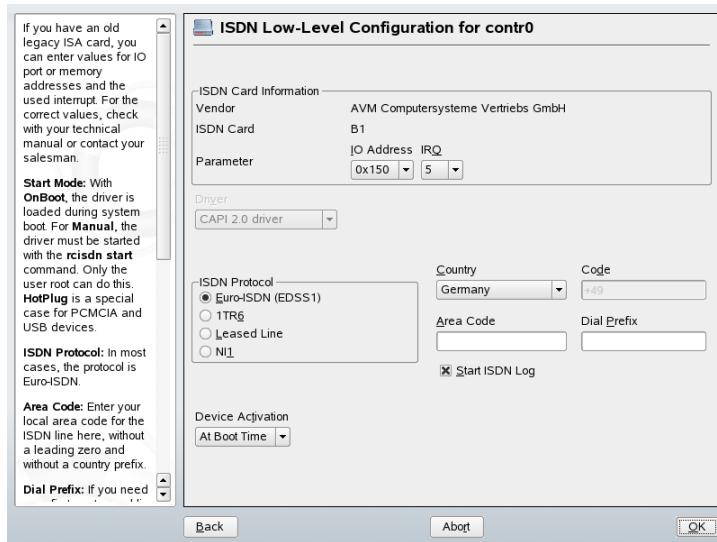
This opens the address configuration dialog. If your ISP does not assign a dynamic IP address to your host, disable *Dynamic IP Address* then enter your host's local IP address and the remote IP address. Ask your ISP for this information. Leave *Default Route* enabled and close the dialog by selecting *OK*.

Selecting *Next* returns to the original dialog, which displays a summary of the modem configuration. Close this dialog with *Finish*.

18.4.3 ISDN

Use this module to configure one or several ISDN cards for your system. If YaST did not detect your ISDN card, manually select it. Multiple interfaces are possible, but several ISPs can be configured for one interface. In the subsequent dialogs, set the ISDN options necessary for the proper functioning of the card.

Figure 18.5 ISDN Configuration

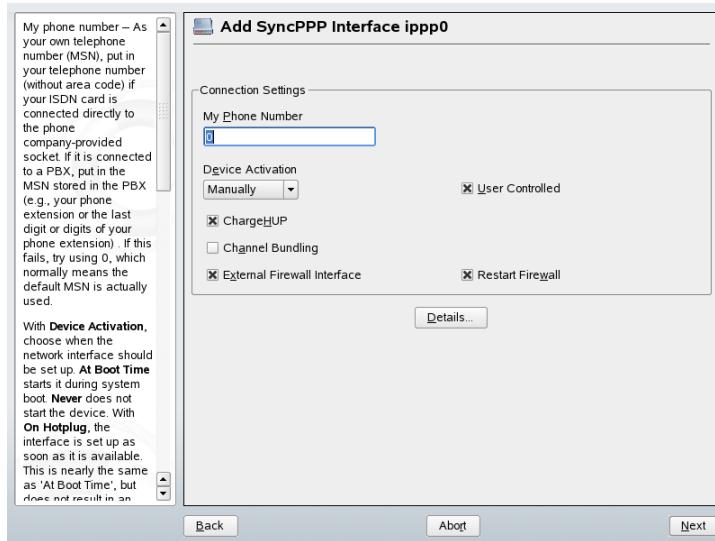


In the next dialog, shown in [Figure 18.5, “ISDN Configuration”](#) (page 340), select the protocol to use. The default is *Euro-ISDN (EDSS1)*, but for older or larger exchanges, select *ITR6*. If you are in the US, select *N1*. Select your country in the relevant field. The corresponding country code then appears in the field next to it. Finally, provide your *Area Code* and the *Dial Prefix* if necessary.

Start Mode defines how the ISDN interface should be started: *At Boot Time* causes the ISDN driver to be initialized each time the system boots. *Manually* requires you to load the ISDN driver as `root` with the command `rcisdn start`. *On Hotplug*, used for PCMCIA or USB devices, loads the driver after the device is plugged in. When finished with these settings, select *OK*.

In the next dialog, specify the interface type for your ISDN card and add ISPs to an existing interface. Interfaces may be either the SyncPPP or the RawIP type, but most ISPs operate in the SyncPPP mode, which is described below.

Figure 18.6 ISDN Interface Configuration



The number to enter for *My Phone Number* depends on your particular setup:

ISDN Card Directly Connected to Phone Outlet

A standard ISDN line provides three phone numbers (called multiple subscriber numbers, or MSNs). If the subscriber asked for more, there may be up to 10. One of these MSNs must be entered here, but without your area code. If you enter the wrong number, your phone operator automatically falls back to the first MSN assigned to your ISDN line.

ISDN Card Connected to a Private Branch Exchange

Again, the configuration may vary depending on the equipment installed:

1. Smaller private branch exchanges (PBX) built for home purposes mostly use the Euro-ISDN (EDSS1) protocol for internal calls. These exchanges have an internal S0 bus and use internal numbers for the equipment connected to them.

Use one of the internal numbers as your MSN. You should be able to use at least one of the exchange's MSNs that have been enabled for direct outward dialing. If this does not work, try a single zero. For further information, consult the documentation that came with your phone exchange.

2. Larger phone exchanges designed for businesses normally use the 1TR6 protocol for internal calls. Their MSN is called EAZ and usually corresponds to the direct-dial number. For the configuration under Linux, it should be sufficient to enter the last digit of the EAZ. As a last resort, try each of the digits from 1 to 9.

For the connection to be terminated just before the next charge unit is due, enable *ChargeHUP*. However, remember that may not work with every ISP. You can also enable channel bundling (multilink PPP) by selecting the corresponding option. Finally, you can enable SuSEfirewall2 for your link by selecting *External Firewall Interface* and *Restart Firewall*. To enable the normal user without administrator permissions to activate or deactivate the interface, select the *User Controlled*.

Details opens a dialog in which to implement more complex connection schemes, which are not relevant for normal home users. Leave the *Details* dialog by selecting *OK*.

In the next dialog, make IP address settings. If you have not been given a static IP by your provider, select *Dynamic IP Address*. Otherwise, use the fields provided to enter your host's local IP address and the remote IP address according to the specifications of your ISP. If the interface should be the default route to the Internet, select *Default Route*. Each host can only have one interface configured as the default route. Leave this dialog by selecting *Next*.

The following dialog allows you to set your country and select an ISP. The ISPs included in the list are call-by-call providers only. If your ISP is not in the list, select *New*. This opens the *Provider Parameters* dialog in which to enter all the details for your ISP. When entering the phone number, do not include any blanks or commas among the digits. Finally, enter your login and the password as provided by the ISP. When finished, select *Next*.

To use *Dial on Demand* on a stand-alone workstation, also specify the name server (DNS server). Most ISPs support dynamic DNS, which means the IP address of a name server is sent by the ISP each time you connect. For a single workstation, however, you still need to provide a placeholder address like 192.168.22.99. If your ISP does not support dynamic DNS, specify the name server IP addresses of the ISP. If desired, specify a time-out for the connection—the period of network inactivity (in seconds) after which the connection should be automatically terminated. Confirm your settings with *Next*. YaST displays a summary of the configured interfaces. To make all these settings active, select *Finish*.

18.4.4 Cable Modem

In some countries, such as Austria and the US, it is quite common to access the Internet through the TV cable network. The TV cable subscriber usually gets a modem that is connected to the TV cable outlet on one side and to a computer network card on the other (using a 10Base-TG twisted pair cable). The cable modem then provides a dedicated Internet connection with a fixed IP address.

Depending on the instructions provided by your ISP, when configuring the network card either select *Automatic Address Setup (via DHCP)* or *Static Address Setup*. Most providers today use DHCP. A static IP address often comes as part of a special business account.

18.4.5 DSL

To configure your DSL device, select the *DSL* module from the *YaST Network Devices* section. This YaST module consists of several dialogs in which to set the parameters of DSL links based on one of the following protocols:

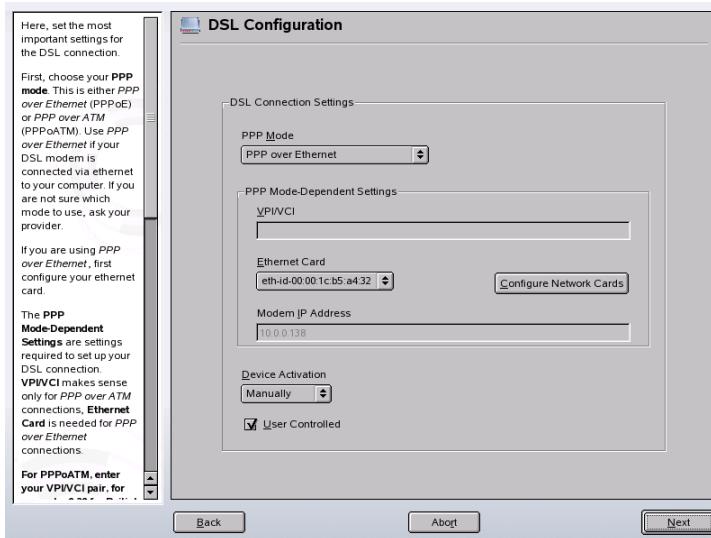
- PPP over Ethernet (PPPoE)
- PPP over ATM (PPPoATM)
- CAPI for ADSL (Fritz Cards)
- Point-to-Point Tunneling Protocol (PPTP)—Austria

The configuration of a DSL connection based on PPPoE or PPTP requires that the corresponding network card has already been set up in the correct way. If you have not done so yet, first configure the card by selecting *Configure Network Cards* (see [Section 18.4.1, “Configuring the Network Card with YaST”](#) (page 334)). In the case of a DSL link, addresses may be assigned automatically but not via DHCP, which is why you should not enable the option *Automatic address setup (via DHCP)*. Instead, enter a static dummy address for the interface, such as 192.168.22.1. In *Subnet Mask*, enter 255.255.255.0. If you are configuring a stand-alone workstation, leave *Default Gateway* empty.

TIP

Values in *IP Address* and *Subnet Mask* are only placeholders. They are only needed to initialize the network card and do not represent the DSL link as such.

Figure 18.7 *DSL Configuration*



To begin the DSL configuration (see [Figure 18.7, “DSL Configuration”](#) (page 344)), first select the PPP mode and the ethernet card to which the DSL modem is connected (in most cases, this is `eth0`). Then use *Device Activation* to specify whether the DSL link should be established during the boot process. Click *User Controlled* to authorize the normal user without root permissions to activate or deactivate the interface with KInternet. The dialog also lets you select your country and choose from a number of ISPs operating in it. The details of any subsequent dialogs of the DSL configuration depend on the options set so far, which is why they are only briefly mentioned in the following paragraphs. For details on the available options, read the detailed help available from the dialogs.

To use *Dial on Demand* on a stand-alone workstation, also specify the name server (DNS server). Most ISPs support dynamic DNS—the IP address of a name server is sent by the ISP each time you connect. For a single workstation, however, provide a placeholder address like `192.168.22.99`. If your ISP does not support dynamic DNS, enter the name server IP address provided by your ISP.

Idle Time-Out (seconds) defines a period of network inactivity after which to terminate the connection automatically. A reasonable time-out value is between 60 and 300 seconds. If *Dial on Demand* is disabled, it may be useful to set the time-out to zero to prevent automatic hang-up.

The configuration of T-DSL is very similar to the DSL setup. Just select *T-Online* as your provider and YaST opens the T-DSL configuration dialog. In this dialog, provide some additional information required for T-DSL—the line ID, the T-Online number, the user code, and your password. All of these should be included in the information you received after subscribing to T-DSL.

18.5 Managing Network Connections with NetworkManager

NetworkManager is the ideal solution for a mobile workstation. With NetworkManager, you do not need to worry about configuring network interfaces and switching between networks when you are moving. NetworkManager can automatically connect to known WLAN networks. If you have two or more connection possibilities, it can connect to the faster one.

NOTE: NetworkManager and SCPM

Do not use NetworkManager together with SCPM when SCPM profiles also change network settings. If you want to use SCPM and NetworkManager at the same time, disable the network resource in SCPM configuration.

NetworkManager is not a suitable solution in the following cases:

- Your computer has a static address.
- You want to use more than one provider for dial-up for one interface.
- You want to use WPA-EAP encryption for your WLAN connection.
- Your computer is a router for your network.
- Your computer provides network services for other computers in your network, for example, it is a DHCP or DNS server.

18.5.1 Controlling NetworkManager

To start NetworkManager, enable NetworkManager in the YaST module of your network device. Because NetworkManager does not need standard network configuration, the YaST configuration becomes inactive. NetworkManager automatically chooses the best network available, but it can automatically connect only to a known network. For the first connection to a network, use the NetworkManager applet. If the network requires additional information, like its name, password or encryption key, NetworkManager prompts for it.

Both KDE and GNOME have their own applets for NetworkManager. An appropriate applet should start automatically with the desktop environment. The applet is then shown as an icon in the system tray. Functions of both applets are similar, but their interfaces are different. They can also be used in other graphical environments with standard system tray support.

The KNetworkManager Applet

KNetworkManager is a KDE applet for controlling NetworkManager. If it is not running start it with the command `knetworkmanager`. When it is running, a blue Earth icon is shown in the system tray. Right-clicking the icon opens the KNetworkManager menu with various commands for managing network connections.

The menu contains available network connections, both wired and wireless devices. If you hold the mouse cursor over them, details are displayed. The currently used connection is checked in the menu. The signal strength of wireless networks is also shown in the menu. Encrypted wireless networks are marked with a blue lock icon. To connect to an encrypted network, choose it from the menu. In the dialog that opens, choose the type of *Encryption* the network uses and enter the appropriate *Passphrase* or *Key*.

To connect to a network that does not broadcast its service set identifier (ESSID) and therefore cannot be detected automatically, choose *Connect to Other Wireless Network*. In the dialog that opens, enter the ESSID and set encryption parameters if necessary.

To access dial-up connections, choose *Dial-Up Connections*. When dial-up connections have already been defined, start the connection by clicking the connection to use. *Configure Dial-Up Connections* opens YaST where you can define a new dial-up connection.

To disable any active network connection, choose *Options* → *Switch to Offline mode* from the KNetworkManager menu. To reenable the connection, choose *Options* → *Switch to Online mode*. To disable wireless network connections, choose *Options* → *Disable Wireless* from the KNetworkManager menu. To reenable wireless connections, choose *Options* → *Enable Wireless*. Enabling networking takes a few seconds.

The GNOME NetworkManager Applet

GNOME also has its own applet for NetworkManager. If it is not running, start it with the command `nm-applet`. When it is running, an icon is shown in the system tray. The appearance of the icon depends on state of the network connection. If you are not sure what the icon means, hold the mouse cursor over it until an explanation appears.

Left-click the applet icon to show a menu with available networks. The currently used connection is checked in the menu. The signal strength of wireless networks is also shown in the menu. Encrypted wireless networks are marked with a shield icon. To connect to an encrypted network, choose it from the menu. In the dialog that opens, choose the type of *Encryption* the network uses and enter the appropriate *Passphrase* or *Key*.

To connect to a network that does not broadcast its service set identifier (ESSID) and therefore cannot be detected automatically, left-click the icon and choose *Connect to Other Wireless Network*. In the dialog that opens, enter the ESSID and set encryption parameters if necessary.

To disable networking, right-click the applet icon and uncheck *Enable Networking*. To disable wireless networking, right-click the applet icon and uncheck *Enable Wireless*.

To get information about the current connection (including interface used, IP address, and hardware address), right-click the applet icon and choose *Connection information* from the menu.

18.5.2 For More Information

More information about NetworkManager and D-BUS can be found on the following Web sites and directories:

- <http://www.gnome.org/projects/NetworkManager/>—NetworkManager project page

- <http://www.freedesktop.org/Software/dbus>—D-BUS project page
- /usr/share/doc/packages/NetworkManager

18.6 Configuring a Network Connection Manually

Manual configuration of the network software should always be the last alternative. Using YaST is recommended. However, this background information about the network configuration can also assist your work with YaST.

All built-in network cards and hotplug network cards (PCMCIA, USB, some PCI cards) are detected and configured via hotplug. The system sees a network card in two different ways: first as a physical device and second as an interface. The insertion or detection of a device triggers a hotplug event. This hotplug event triggers the initialization of the device with the script `hwup`. When the network card is initialized as a new network interface, the kernel generates another hotplug event that triggers the setup of the interface with `ifup`.

The kernel numbers interface names according to the temporal order of their registration. The initialization sequence is decisive for the assignment of names. If one of several network card fails, the numbering of all subsequently initialized cards is shifted. For real hotpluggable cards, the order in which the devices are connected is what matters.

To achieve a flexible configuration, the configuration of the device (hardware) and the interface has been separated and the mapping of configurations to devices and interfaces is no longer managed on the basis of the interface names. The device configurations are located in `/etc/sysconfig/hardware/hwcfg-*`. The interface configurations are located in `/etc/sysconfig/network/ifcfg-*`. The names of the configurations are assigned in such a way that they describe the devices and interfaces with which they are associated. Because the former mapping of drivers to interface name required static interface names, this mapping can no longer take place in `/etc/modprobe.conf`. In the new concept, alias entries in this file would cause undesirable side effects.

The configuration names—everything after `hwcfg-` or `ifcfg-`—can describe the devices by means of the slot, a device-specific ID, or the interface name. For example, the configuration name for a PCI card could be `bus-pci-0000:02:01.0` (PCI

slot) or `vpid-0x8086-0x1014-0x0549` (vendor and product ID). The name of the associated interface could be `bus-pci-0000:02:01.0` or `wlan-id-00:05:4e:42:31:7a` (MAC address).

To assign a certain network configuration to any card of a certain type (of which only one is inserted at a time) instead of a certain card, select less specific configuration names. For example, `bus-pcmcia` would be used for all PCMCIA cards. On the other hand, the names can be limited by a preceding interface type. For example, `wlan-bus-usb` would be assigned to WLAN cards connected to a USB port.

The system always uses the configuration that best describes an interface or the device providing the interface. The search for the most suitable configuration is handled by `getcfg`. The output of `getcfg` delivers all information that can be used for describing a device. Details regarding the specification of configuration names are available in the manual page of `getcfg`.

With the described method, a network interface is configured with the correct configuration even if the network devices are not always initialized in the same order. However, the name of the interface still depends on the initialization sequence. There are two ways to ensure reliable access to the interface of a certain network card:

- `getcfg-interface configuration name` returns the name of the associated network interface. Therefore, the configuration name, such as `firewall`, `dhcpd`, `routing`, or various virtual network interfaces (tunnels), can be entered in some configuration files instead of the interface name, which is not persistent.
- Persistent interface names are assigned to each interface automatically. You may adjust them to suit your needs. When creating interface names, proceed as outlined in `/etc/udev/rules.d/30-net_persistent_names.rules`. However, the persistent name `pname` should not be the same as the name that would automatically be assigned by the kernel. Therefore, `eth*`, `tr*`, `wlan*`, and so on are not permitted. Instead, use `net*` or descriptive names like `external`, `internal`, or `dmz`. Make sure that the same interface name is not used twice. Allowed characters in interface names are restricted to `[a-zA-Z0-9]`. A persistent name can only be assigned to an interface immediately after its registration, which means that the driver of the network card must be reloaded or `hwup device description` must be executed. The command `rcnetwork restart` is not sufficient for this purpose.

IMPORTANT: Using Persistent Interface Names

The use of persistent interface names has not been tested in all areas. Therefore, some applications may not be able to handle freely selected interface names.

`ifup` requires an existing interface, because it does not initialize the hardware. The initialization of the hardware is handled by the command `hwup` (executed by `hotplug` or `coldplug`). When a device is initialized, `ifup` is automatically executed for the new interface via `hotplug` and the interface is set up if the start mode is `onboot`, `hotplug`, or `auto` and the network service was started. Formerly, the command `ifup interfaceName` triggered the hardware initialization. Now the procedure has been reversed. First, a hardware component is initialized then all other actions follow. In this way, a varying number of devices can always be configured in the best way possible with an existing set of configurations.

[Table 18.5, “Manual Network Configuration Scripts”](#) (page 350) summarizes the most important scripts involved in the network configuration. Where possible, the scripts are distinguished by hardware and interface.

Table 18.5 *Manual Network Configuration Scripts*

Configura- tion Stage	Command	Function
Hardware	<code>hw{up,down,status}</code>	The <code>hw*</code> scripts are executed by the hotplug subsystem to initialize a device, undo the initialization, or query the status of a device. More information is available in the manual page of <code>hwup</code> .
Interface	<code>getcfg</code>	<code>getcfg</code> can be used to query the interface name associated with a configuration name or a hardware description. More information is available in the manual page of <code>getcfg</code> .
Interface	<code>if{up,down,status}</code>	The <code>if*</code> scripts start existing network interfaces or return the status of the

Configuration Stage	Command	Function
	ifup	specified interface. More information is available in the manual page of <code>ifup</code> .

More information about hotplug and persistent device names is available in [Chapter 12, *Dynamic Kernel Device Management with udev*](#) (page 251).

18.6.1 Configuration Files

This section provides an overview of the network configuration files and explains their purpose and the format used.

/etc/sysconfig/hardware/hwcfg-*

These files contain the hardware configurations of network cards and other devices. They contain the needed parameters, such as the kernel module, start mode, and script associations. Refer to the manual page of `hwup` for details. Regardless of the existing hardware, the `hwcfg-static-*` configurations are applied when coldplug is started.

/etc/sysconfig/network/ifcfg-*

These files contain the configurations for network interface. They include information such as the start mode and the IP address. Possible parameters are described in the manual page of `ifup`. Additionally, all variables from the files `dhcp`, `wireless`, and `config` can be used in the `ifcfg-*` files if a general setting should be used for only one interface.

/etc/sysconfig/network/config, dhcp, wireless

The file `config` contains general settings for the behavior of `ifup`, `ifdown`, and `ifstatus`. `dhcp` contains settings for DHCP and `wireless` for wireless LAN cards. The variables in all three configuration files are commented and can also be used in `ifcfg-*` files, where they are treated with higher priority.

/etc/sysconfig/network/routes, ifroute-*

The static routing of TCP/IP packets is determined here. All the static routes required by the various system tasks can be entered in the /etc/sysconfig/network/routes file: routes to a host, routes to a host via a gateway, and routes to a network. For each interface that needs individual routing, define an additional configuration file: /etc/sysconfig/network/ifroute-*. Replace * with the name of the interface. The entries in the routing configuration files look like this:

#	Destination	Dummy/Gateway	Netmask	Device
	127.0.0.0	0.0.0.0	255.255.255.0	lo
	204.127.235.0	0.0.0.0	255.255.255.0	eth0
	default	204.127.235.41	0.0.0.0	eth0
	207.68.156.51	207.68.145.45	255.255.255.255	eth1
	192.168.0.0	207.68.156.51	255.255.0.0	eth1

The route's destination is in the first column. This column may contain the IP address of a network or host or, in the case of *reachable* name servers, the fully qualified network or hostname.

The second column contains the default gateway or a gateway through which a host or network can be accessed. The third column contains the netmask for networks or hosts behind a gateway. For example, the mask is 255.255.255.255 for a host behind a gateway.

The fourth column is only relevant for networks connected to the local host such as loopback, Ethernet, ISDN, PPP, and dummy device. The device name must be entered here.

An (optional) fifth column can be used to specify the type of a route. Columns that are not needed should contain a minus sign – to ensure that the parser correctly interprets the command. For details, refer to the `routes(5)` man page.

/etc/resolv.conf

The domain to which the host belongs is specified in this file (keyword `search`). Also listed is the status of the name server address to access (keyword `nameserver`). Multiple domain names can be specified. When resolving a name that is not fully qualified, an attempt is made to generate one by attaching the individual search entries. Use multiple name servers by entering several lines, each beginning with `nameserver`.

Precede comments with # signs. YaST enters the specified name server in this file.

[Example 18.5, “/etc/resolv.conf”](#) (page 353) shows what /etc/resolv.conf could look like.

Example 18.5 /etc/resolv.conf

```
# Our domain
search example.com
#
# We use sun (192.168.0.20) as nameserver
nameserver 192.168.0.20
```

Some services, like pppd (wvdial), ipppd (isdn), dhcp (dhcpcd and dhclient), pcmcia, and hotplug, modify the file /etc/resolv.conf by means of the script `modify_resolvconf`. If the file /etc/resolv.conf has been temporarily modified by this script, it contains a predefined comment giving information about the service that modified it, the location where the original file has been backed up, and how to turn off the automatic modification mechanism. If /etc/resolv.conf is modified several times, the file includes modifications in a nested form. These can be reverted in a clean way even if this reversal takes place in an order different from the order in which modifications were introduced. Services that may need this flexibility include isdn, pcmcia, and hotplug.

If a service was not terminated in a normal, clean way, `modify_resolvconf` can be used to restore the original file. Also, on system boot, a check is performed to see whether there is an uncleaned, modified `resolv.conf`, for example, after a system crash, in which case the original (unmodified) `resolv.conf` is restored.

YaST uses the command `modify_resolvconf` check to find out whether `resolv.conf` has been modified and subsequently warns the user that changes will be lost after restoring the file. Apart from this, YaST does not rely on `modify_resolvconf`, which means that the impact of changing `resolv.conf` through YaST is the same as that of any manual change. In both cases, changes have a permanent effect. Modifications requested by the mentioned services are only temporary.

/etc/hosts

In this file, shown in [Example 18.6, “/etc/hosts”](#) (page 354), IP addresses are assigned to hostnames. If no name server is implemented, all hosts to which an IP connection will be set up must be listed here. For each host, enter a line consisting of the IP address, the fully qualified hostname, and the hostname into the file. The IP address

must be at the beginning of the line and the entries separated by blanks and tabs. Comments are always preceded by the # sign.

Example 18.6 /etc/hosts

```
127.0.0.1 localhost
192.168.0.20 sun.example.com sun
192.168.0.0 earth.example.com earth
```

/etc/networks

Here, network names are converted to network addresses. The format is similar to that of the hosts file, except the network names precede the addresses. See [Example 18.7, “/etc/networks” \(page 354\)](#).

Example 18.7 /etc/networks

```
loopback      127.0.0.0
localnet      192.168.0.0
```

/etc/host.conf

Name resolution—the translation of host and network names via the *resolver* library—is controlled by this file. This file is only used for programs linked to libc4 or libc5. For current glibc programs, refer to the settings in /etc/nsswitch.conf. A parameter must always stand alone in its own line. Comments are preceded by a # sign. [Table 18.6, “Parameters for /etc/host.conf” \(page 354\)](#) shows the parameters available. A sample /etc/host.conf is shown in [Example 18.8, “/etc/host.conf” \(page 355\)](#).

Table 18.6 Parameters for /etc/host.conf

order <i>hosts</i> , <i>bind</i>	Specifies in which order the services are accessed for the name resolution. Available arguments are (separated by blank spaces or commas):
	<i>hosts</i> : Searches the /etc/hosts file
	<i>bind</i> : Accesses a name server
	<i>nis</i> : Uses NIS

multi <i>on/off</i>	Defines if a host entered in <code>/etc/hosts</code> can have multiple IP addresses.
nospoof <i>on</i> spoofalert <i>on/off</i>	These parameters influence the name server <i>spoofing</i> , but, apart from that, do not exert any influence on the network configuration.
trim <i>domainname</i>	The specified domain name is separated from the hostname after hostname resolution (as long as the hostname includes the domain name). This option is useful if only names from the local domain are in the <code>/etc/hosts</code> file, but should still be recognized with the attached domain names.

Example 18.8 `/etc/host.conf`

```
# We have named running
order hosts bind
# Allow multiple addrs
multi on
```

/etc/nsswitch.conf

The introduction of the GNU C Library 2.0 was accompanied by the introduction of the *Name Service Switch* (NSS). Refer to the `nsswitch.conf` (5) man page and *The GNU C Library Reference Manual* for details.

The order for queries is defined in the file `/etc/nsswitch.conf`. A sample `nsswitch.conf` is shown in [Example 18.9, “/etc/nsswitch.conf”](#) (page 356). Comments are introduced by # signs. In this example, the entry under the `hosts` database means that a request is sent to `/etc/hosts` (files) via DNS (see [Chapter 20, The Domain Name System](#) (page 367)).

Example 18.9 /etc/nsswitch.conf

```
passwd:      compat
group:       compat

hosts:        files dns
networks:    files dns

services:     db files
protocols:   db files

netgroup:    files
automount:  files nis
```

The “databases” available over NSS are listed in [Table 18.7, “Databases Available via /etc/nsswitch.conf”](#) (page 356). In addition, automount, bootparams, netmasks, and publickey are expected in the near future. The configuration options for NSS databases are listed in [Table 18.8, “Configuration Options for NSS “Databases””](#) (page 357).

Table 18.7 Databases Available via /etc/nsswitch.conf

aliases	Mail aliases implemented by sendmail; see man 5 aliases.
ethers	Ethernet addresses.
group	For user groups, used by getrent. See also the man page for group.
hosts	For hostnames and IP addresses, used by gethostbyname and similar functions.
netgroup	Valid host and user lists in the network for the purpose of controlling access permissions; see the netgroup (5) man page.
networks	Network names and addresses, used by getnetent.
passwd	User passwords, used by getpwent; see the passwd (5) man page.

protocols	Network protocols, used by <code>getprotoent</code> ; see the <code>protocols(5)</code> man page.
rpc	Remote procedure call names and addresses, used by <code>getrpcbyname</code> and similar functions.
services	Network services, used by <code>getservent</code> .
shadow	Shadow passwords of users, used by <code>getspnam</code> ; see the <code>shadow(5)</code> man page.

Table 18.8 Configuration Options for NSS “Databases”

files	directly access files, for example, <code>/etc/aliases</code>
db	access via a database
nis, nisplus	NIS, see also Chapter 21, Using NIS (page 391)
dns	can only be used as an extension for <code>hosts</code> and <code>networks</code>
compat	can only be used as an extension for <code>passwd</code> , <code>shadow</code> , and <code>group</code>

/etc/nscd.conf

This file is used to configure nscd (name service cache daemon). See the `nscd(8)` and `nscd.conf(5)` man pages. By default, the system entries of `passwd` and `groups` are cached by nscd. This is important for the performance of directory services, like NIS and LDAP, because otherwise the network connection needs to be used for every access to names or groups. `hosts` is not cached by default, because the mechanism in nscd to cache hosts makes the local system unable to trust forward and reverse lookup checks. Instead of asking nscd to cache names, set up a caching DNS server.

If the caching for `passwd` is activated, it usually takes about fifteen seconds until a newly added local user is recognized. Reduce this waiting time by restarting nscd with the command `rcnscd restart`.

/etc/hostname

This contains the hostname without the domain name attached. This file is read by several scripts while the machine is booting. It may only contain one line in which the hostname is set.

18.6.2 Start-Up Scripts

Apart from the configuration files described above, there are also various scripts that load the network programs while the machine is booting. These are started as soon as the system is switched to one of the *multiuser runlevels*. Some of these scripts are described in [Table 18.9, “Some Start-Up Scripts for Network Programs”](#) (page 358).

Table 18.9 Some Start-Up Scripts for Network Programs

/etc/init.d/network	This script handles the configuration of the network interfaces. The hardware must already have been initialized by /etc/init.d/coldplug (via hotplug). If the network service was not started, no network interfaces are implemented when they are inserted via hotplug.
/etc/init.d/inetd	Starts xinetd. xinetd can be used to make server services available on the system. For example, it can start vsftpd whenever an FTP connection is initiated.
/etc/init.d/portmap	Starts the portmapper needed for the RPC server, such as an NFS server.
/etc/init.d/nfsserver	Starts the NFS server.
/etc/init.d/sendmail	Controls the sendmail process.
/etc/init.d/ypserv	Starts the NIS server.
/etc/init.d/ypbind	Starts the NIS client.

18.7 smpppd as Dial-up Assistant

Most home users do not have a dedicated line connecting them to the Internet. Instead, they use dial-up connections. Depending on the dial-up method (ISDN or DSL), the connection is controlled by ipppd or pppd. Basically, all that needs to be done to go online is to start these programs correctly.

If you have a flat-rate connection that does not generate any additional costs for the dial-up connection, simply start the respective daemon. Control the dial-up connection with a KDE applet or a command-line interface. If the Internet gateway is not the host you are using, you might want to control the dial-up connection by way of a network host.

This is where smpppd is involved. It provides a uniform interface for auxiliary programs and acts in two directions. First, it programs the required pppd or ipppd and controls its dial-up properties. Second, it makes various providers available to the user programs and transmits information about the current status of the connection. As smpppd can also be controlled by way of the network, it is suitable for controlling dial-up connections to the Internet from a workstation in a private subnetwork.

18.7.1 Configuring smpppd

The connections provided by smpppd are automatically configured by YaST. The actual dial-up programs KInternet and cinternet are also preconfigured. Manual settings are only required to configure additional features of smpppd, such as remote control.

The configuration file of smpppd is `/etc/smpppd.conf`. By default, it does not enable remote control. The most important options of this configuration file are:

`open-inet-socket = yes / no`

To control smpppd via the network, this option must be set to `yes`. The port on which smpppd listens is 3185. If this parameter is set to `yes`, the parameters `bind-address`, `host-range`, and `password` should also be set accordingly.

`bind-address = ip`

If a host has several IP addresses, use this parameter to determine at which IP address smpppd should accept connections.

`host-range = min ip max ip`

The parameter `host-range` defines a network range. Hosts whose IP addresses are within this range are granted access to smpppd. All hosts not within this range are denied access.

`password = password`

By assigning a password, limit the clients to authorized hosts. As this is a plain-text password, you should not overrate the security it provides. If no password is assigned, all clients are permitted to access smpppd.

`slp-register = yes / no`

With this parameter, the smpppd service can be announced in the network via SLP.

More information about smpppd is available in the `smpppd(8)` and `smpppd.conf(5)` man pages.

18.7.2 Configuring KInternet, cinternet, and qinternet for Remote Use

KInternet, cinternet, and qinternet can be used to control a local or remote smpppd. cinternet is the command-line counterpart of the graphical KInternet. qinternet is basically the same as KInternet, but does not use the KDE libraries, so it can be used without KDE and must be installed separately. To prepare these utilities for use with a remote smpppd, edit the configuration file `/etc/smpppd-c.conf` manually or using KInternet. This file only uses three options:

`sites = list of sites`

Here, tell the front-ends where to search for smpppd. The front-ends test the options in the order specified here. The `local` option orders the establishment of a connection to the local smpppd. `gateway` points to an smpppd on the gateway. The connection should be established as specified under `server` in `config-file`. `slp` orders the front-ends to connect to an smpppd found via SLP.

`server = server`

Here, specify the host on which smpppd runs.

`password = password`

Insert the password selected for smpppd.

If smpppd is active, you can now try to access it, for example, with `cinternet --verbose --interface-list`. If you experience difficulties at this point, refer to the `smpppd-c.conf(5)` and `cinternet(8)` man pages.

19

SLP Services in the Network

The *service location protocol* (SLP) was developed to simplify the configuration of networked clients within a local network. To configure a network client, including all required services, the administrator traditionally needs detailed knowledge of the servers available in the network. SLP makes the availability of selected services known to all clients in the local network. Applications that support SLP can use the information distributed and be configured automatically.

SUSE Linux supports installation using installation sources provided with SLP and contains many system services with integrated support for SLP. YaST and Konqueror both have appropriate front-ends for SLP. You can use SLP to provide networked clients with central functions, such as an installation server, YOU server, file server, or print server on your SUSE Linux.

19.1 Registering Your Own Services

Many applications under SUSE Linux already have integrated SLP support through the use of the `libslp` library. If a service has not been compiled with SLP support, use one of the following methods to make it available with SLP:

Static Registration with `/etc/slp.reg.d`

Create a separate registration file for each new service. The following is an example of a file for registering a scanner service:

```
## Register a saned service on this system
## en means english language
## 65535 disables the timeout, so the service registration does
## not need refreshes
```

```
service:scanner.sane://$HOSTNAME:6566,en,65535
watch-port=tcp=6566
description=SANE scanner daemon
```

The most important line in this file is the *service URL*, which begins with `service:`. This contains the service type (`scanner . sane`) and the address under which the service is available on the server. `$HOSTNAME` is automatically replaced with the full hostname. The name of the TCP port on which the relevant service can be found follows, separated by a colon. Then enter the language in which the service should appear and the duration of registration in seconds. These should be separated from the service URL by commas. Set the value for the duration of registration between 0 and 65535. 0 prevents registration. 65535 removes all restrictions.

The registration file also contains the two variables `watch-tcp-port` and `description`. `watch-tcp-port` links the SLP service announcement to whether the relevant service is active by having `slpd` check the status of the service. The second variable contains a more precise description of the service that is displayed in suitable browsers.

Static Registration with `/etc/slp.reg`

The only difference from the procedure with `/etc/slp.reg.d` is the grouping of all services within a central file.

Dynamic Registration with `slptool`

If a service should be registered for SLP from proprietary scripts, use the `slptool` command line front-end.

19.2 SLP Front-Ends in SUSE Linux

SUSE Linux contains several front-ends that enable SLP information to be checked and used by means of a network:

`slptool`

`slptool` is a simple command line program that can be used to announce SLP inquiries in the network or announce proprietary services. `slptool --help` lists all available options and functions. `slptool` can also be called from scripts that process SLP information.

YaST SLP Browser

YaST contains a separate SLP browser that lists all services in the local network announced by SLP in a tree diagram under *Network Services* → *SLP Browser*.

Konqueror

When used as a network browser, Konqueror can display all SLP services available in the local network at `slp:/`. Click the icons in the main window to obtain more detailed information about the relevant service. If you use Konqueror with `service:/`, click the relevant icon once in the browser window to set up a connection with the selected service.

19.3 Activating SLP

`sldp` must run on your system if you want to offer services. It is not necessary to start this daemon simply to make service inquiries. Like most system services in SUSE Linux, the `sldp` daemon is controlled by means of a separate init script. The daemon is inactive by default. To activate it for the duration of a session, run `rcsldp start` as `root` to start it and `rcsldp stop` to stop it. Perform a restart or status check with `restart` or `status`. If `sldp` should be active by default, run the `insserv sldp` command once as `root`. This automatically includes `sldp` in the set of services to start when a system boots.

19.4 For More Information

The following sources provide further information about SLP:

RFC 2608, 2609, 2610

RFC 2608 generally deals with the definition of SLP. RFC 2609 deals with the syntax of the service URLs used in greater detail and RFC 2610 deals with DHCP via SLP.

<http://www.openslp.com>

The home page of the OpenSLP project.

`/usr/share/doc/packages/openslp`

This directory contains all available documentation for SLP, including a `README` . SuSE containing the SUSE Linux details, the RFCs mentioned above, and two

introductory HTML documents. Programmers who want to use the SLP functions should install the `openslp-devel` package to consult its supplied *Programmers Guide*.

20

The Domain Name System

DNS (domain name system) is needed to resolve the domain names and hostnames into IP addresses. In this way, the IP address 192.168.0.0 is assigned to the hostname `earth`, for example. Before setting up your own name server, read the general information about DNS in [Section 18.3, “Name Resolution”](#) (page 333). The following configuration examples refer to BIND.

20.1 DNS Terminology

Zone

The domain namespace is divided into regions called zones. For instance, if you have `opensuse.org`, you have the `opensuse` section, or zone, of the `.org` domain.

DNS server

The DNS server is a server that maintains the name and IP information for a domain. You can have a primary DNS server for master zone, a secondary server for slave zone, or a slave server without any zones for caching.

Master zone DNS server

The master zone includes all hosts from your network and a DNS server master zone stores up-to-date records for all the hosts in your domain.

Slave zone DNS server

A slave zone is a copy of the master zone. The slave zone DNS server obtains its zone data with zone transfer operations from its master server. The slave zone DNS server responds authoritatively for the zone as long as it has valid

(not expired) zone data. If the slave cannot obtain a new copy of the zone data, it stops responding for the zone.

Forwarder

Forwarders are DNS servers to which your DNS server should send queries it cannot answer.

Record

The record is information about name and IP address. Supported records and their syntax are described in BIND documentation. Some special records are:

NS record

An NS record tells name servers which machines are in charge of a given domain zone.

MX record

The MX (mail exchange) records describe the machines to contact for directing mail across the Internet.

SOA record

SOA (Start of Authority) record is the first record in a zone file. The SOA record is used when using DNS to synchronize data between multiple computers.

20.2 Configuration with YaST

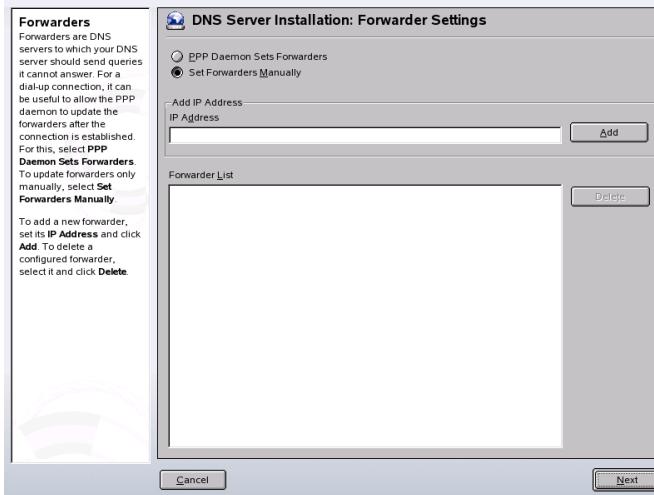
You can use the DNS module of YaST to configure a DNS server for your local network. When starting the module for the first time, a wizard starts, prompting you to make just a few basic decisions concerning administration of the server. Completing this initial setup produces a very basic server configuration that should be functioning in its essential aspects. The expert mode can be used to deal with more advanced configuration tasks.

20.2.1 Wizard Configuration

The wizard consists of three steps or dialogs. At the appropriate places in the dialogs, you are given the opportunity to enter the expert configuration mode.

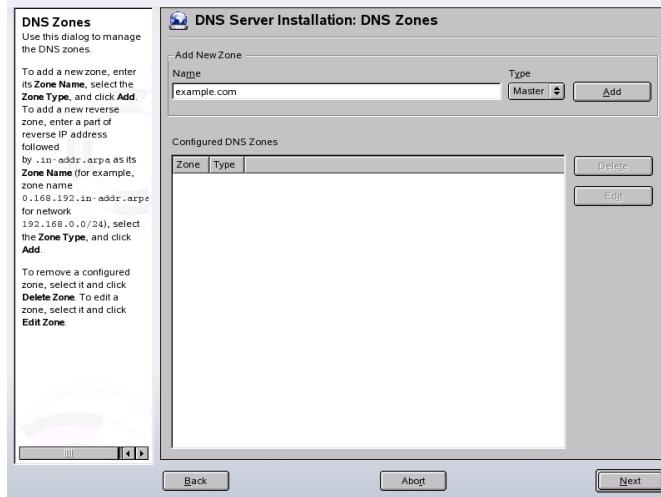
- When starting the module for the first time, the *Forwarder Settings* dialog, shown in [Figure 20.1, “DNS Server Installation: Forwarder Settings”](#) (page 369), opens. In it, decide whether the PPP daemon should provide a list of forwarders on dial-up via DSL or ISDN (*PPP Daemon Sets Forwarders*) or whether you want to supply your own list (*Set Forwarders Manually*).

Figure 20.1 DNS Server Installation: Forwarder Settings



- The *DNS Zones* dialog consists of several parts and is responsible for the management of zone files, described in [Section 20.5, “Zone Files”](#) (page 382). For a new zone, provide a name for it in *Zone Name*. To add a reverse zone, the name must end in `.in-addr.arpa`. Finally, select the *Zone Type* (master or slave). See [Figure 20.2, “DNS Server Installation: DNS Zones”](#) (page 370). Click *Edit Zone* to configure other settings of an existing zone. To remove a zone, click *Delete Zone*.

Figure 20.2 DNS Server Installation: DNS Zones



- 3** In the final dialog, you can open the ports for the DNS service in the firewall that is activated during the installation and decide whether DNS should be started. The expert configuration can also be accessed from this dialog. See [Figure 20.3, “DNS Server Installation: Finish Wizard”](#) (page 371).

Figure 20.3 DNS Server Installation: Finish Wizard



20.2.2 Expert Configuration

After starting the module, YaST opens a window displaying several configuration options. Completing it results in a DNS server configuration with the basic functions in place:

Starting the DNS Server

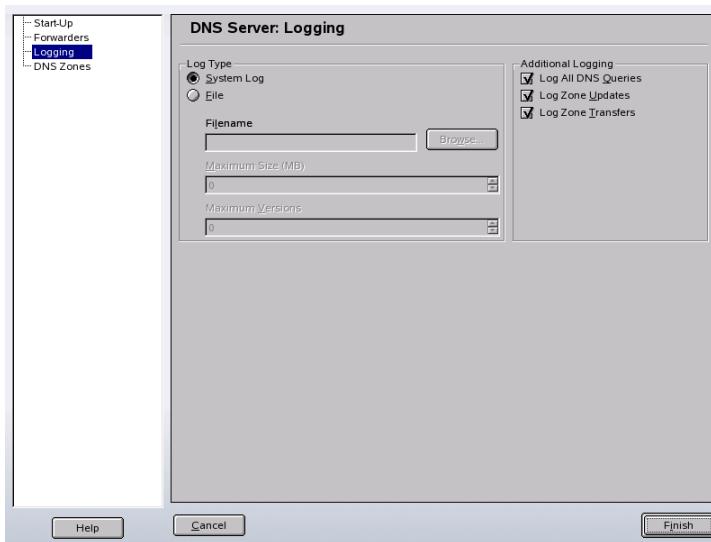
Under *Booting*, define whether the DNS server should be started when the system boots (during booting the system) or manually. To start the DNS server immediately, select *Start DNS Server Now*. To stop the DNS server, select *Stop DNS Server Now*. To save the current settings, select *Save Settings and Restart DNS Server Now*. You can open the DNS port in the firewall with *Open Port in Firewall* and modify the firewall settings with *Firewall Details*.

Logging

To set what the DNS server should log and how, select *Logging*. Under *Log Type*, specify where the DNS server should write the log data. Use the systemwide log file `/var/log/messages` by selecting *Log to System Log* or specify a different file by selecting *Log to File*. In the latter case, additionally specify the maximum file size in megabytes and the number of log files to store.

Further options are available under *Additional Logging*. Enabling *Log All DNS Queries* causes *every* query to be logged, in which case the log file could grow extremely large. For this reason, it is not a good idea to enable this option for other than debugging purposes. To log the data traffic during zone updates between DHCP and DNS server, enable *Log Zone Updates*. To log the data traffic during a zone transfer from master to slave, enable *Log Zone Transfer*. See [Figure 20.4, “DNS Server: Logging” \(page 372\)](#).

Figure 20.4 DNS Server: Logging

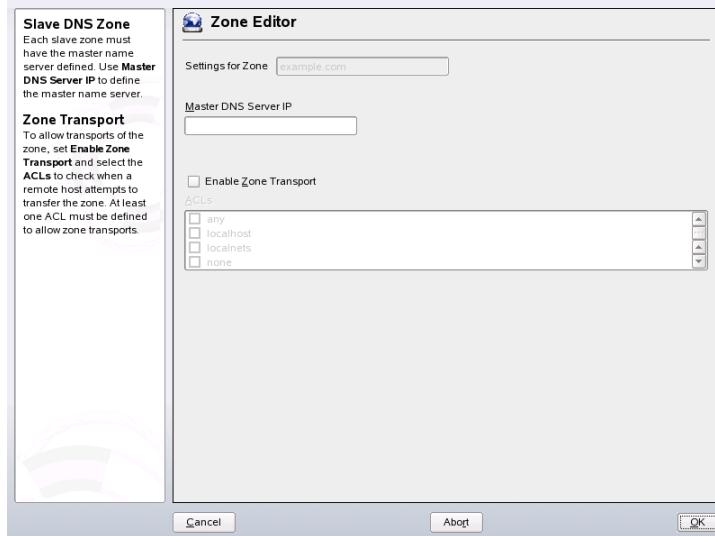


Adding a Slave Zone

To add a slave zone, select *DNS Zones*, choose the zone type *Slave*, and click *Add*.

In the *Zone Editor* under *Master DNS Server*, specify the master from which the slave should fetch its data. To limit access to the server, select one of the ACLs from the list. See [Figure 20.5, “DNS Server: Slave Zone Editor” \(page 373\)](#).

Figure 20.5 *DNS Server: Slave Zone Editor*



Adding a Master Zone

To add a master zone, select *DNS Zones*, choose the zone type *Master*, write the name of the new zone, and click *Add*.

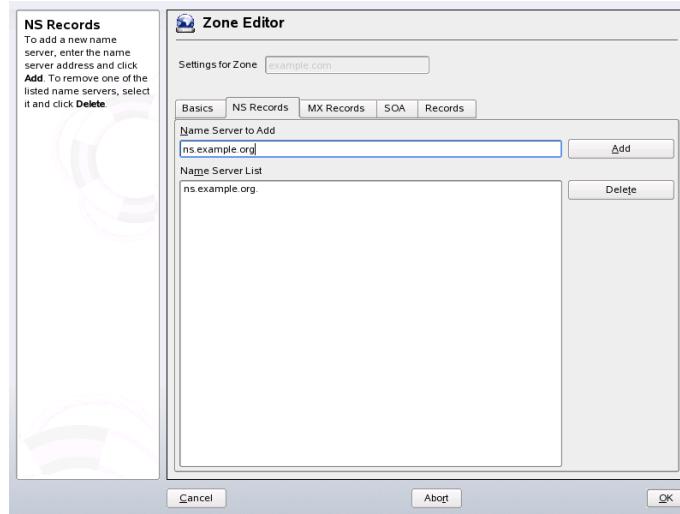
Editing a Master Zone

To edit a master zone, select *DNS Zones*, choose the zone type *Master*, select the master zone from the table, and click *Edit*. The dialog consists of several pages: *Basic* (the one opened first), *NS Records*, *MX Records*, *SOA*, and *Records*.

Zone Editor (NS Records)

This dialog allows you to define alternative name servers for the zones specified. Make sure that your own name server is included in the list. To add a record, enter its name under *Name Server to Add* then confirm with *Add*. See [Figure 20.6, “DNS Server: Zone Editor \(NS Records\)”](#) (page 374).

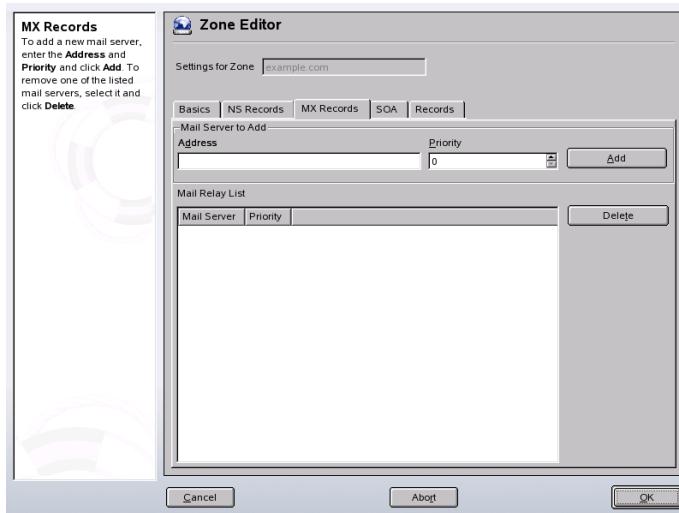
Figure 20.6 DNS Server: Zone Editor (NS Records)



Zone Editor (MX Records)

To add a mail server for the current zone to the existing list, enter the corresponding address and priority value. After doing so, confirm by selecting *Add*. See [Figure 20.7, “DNS Server: Zone Editor \(MX Records\)”](#) (page 375).

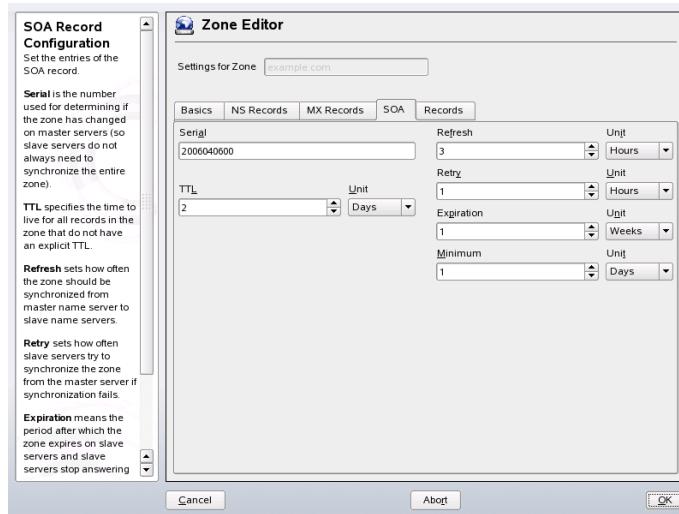
Figure 20.7 DNS Server: Zone Editor (MX Records)



Zone Editor (SOA)

This page allows you to create SOA (start of authority) records. For an explanation of the individual options, refer to [Example 20.6, “File /var/lib/named/world.zone”](#) (page 383).

Figure 20.8 DNS Server: Zone Editor (SOA)



Zone Editor (Records)

This dialog manages name resolution. In *Record Key*, enter the hostname then select its type. *A-Record* represents the main entry. The value for this should be an IP address. *CNAME* is an alias. Use the types *NS* and *MX* for detailed or partial records that expand on the information provided in the *NS Records* and *MX Records* tabs. These three types resolve to an existing A record. *PTR* is for reverse zones. It is the opposite of an A record.

20.3 Starting the Name Server BIND

On a SUSE Linux system, the name server BIND (*Berkeley Internet name domain*) comes preconfigured so it can be started right after installation without any problem. If you already have a functioning Internet connection and have entered `127.0.0.1` as the name server address for `localhost` in `/etc/resolv.conf`, you normally already have a working name resolution without needing to know the DNS of the provider. BIND carries out name resolution via the root name server, a notably slower process. Normally, the DNS of the provider should be entered with its IP address in the configuration file `/etc/named.conf` under `forwarders` to ensure effective and secure name resolution. If this works so far, the name server runs as a pure *caching-only* name server. Only when you configure its own zones will it become a proper DNS.

A simple example of this is included in the documentation in `/usr/share/doc/packages/bind/sample-config`.

TIP: Automatic Adaptation of the Name Server Information

Depending on the type of Internet connection or the network connection, the name server information can automatically be adapted to the current conditions. To do this, set the variable `MODIFY_NAMED_CONF_DYNAMICALLY` in the file `/etc/sysconfig/network/config` to yes.

However, do not set up any official domains until assigned one by the responsible institution. Even if you have your own domain and it is managed by the provider, you are better off not using it, because BIND would otherwise not forward requests for this domain. The Web server at the provider, for example, would not be accessible for this domain.

To start the name server, enter the command `rcnamed start` as root. If “done” appears to the right in green, named, as the name server process is called, has been started successfully. Test the name server immediately on the local system with the `host` or `dig` programs, which should return `localhost` as the default server with the address `127.0.0.1`. If this is not the case, `/etc/resolv.conf` probably contains an incorrect name server entry or the file does not exist at all. For the first test, enter `host 127.0.0.1`, which should always work. If you get an error message, use `rcnamed status` to see whether the server is actually running. If the name server does not start or behaves unexpectedly, you can usually find the cause in the log file `/var/log/messages`.

To use the name server of the provider or one already running on your network as the forwarder, enter the corresponding IP address or addresses in the options section under `forwarders`. The addresses included in [Example 20.1, “Forwarding Options in named.conf”](#) (page 377) are just examples. Adjust these entries to your own setup.

Example 20.1 Forwarding Options in named.conf

```
options {
    directory "/var/lib/named";
    forwarders { 10.11.12.13; 10.11.12.14; };
    listen-on { 127.0.0.1; 192.168.0.99; };
    allow-query { 127/8; 192.168.0/24; };
    notify no;
};
```

The `options` entry is followed by entries for the zone, `localhost`, and `0.0.127.in-addr.arpa`. The `type hint` entry under `."` should always be present. The corresponding files do not need to be modified and should work as they are. Also make sure that each entry is closed with a `";` and that the curly braces are in the correct places. After changing the configuration file `/etc/named.conf` or the zone files, tell BIND to reread them with `rcnamed reload`. Achieve the same by stopping and restarting the name server with `rcnamed restart`. Stop the server at any time by entering `rcnamed stop`.

20.4 The Configuration File `/etc/named.conf`

All the settings for the BIND name server itself are stored in the file `/etc/named.conf`. However, the zone data for the domains to handle, consisting of the hostnames, IP addresses, and so on, are stored in separate files in the `/var/lib/named` directory. The details of this are described later.

`/etc/named.conf` is roughly divided into two areas. One is the `options` section for general settings and the other consists of `zone` entries for the individual domains. A `logging` section and `acl` (access control list) entries are optional. Comment lines begin with a `#` sign or `//`. A minimal `/etc/named.conf` is shown in [Example 20.2, “A Basic /etc/named.conf”](#) (page 379).

Example 20.2 A Basic /etc/named.conf

```
options {
    directory "/var/lib/named";
    forwarders { 10.0.0.1; };
    notify no;
};

zone "localhost" in {
    type master;
    file "localhost.zone";
};

zone "0.0.127.in-addr.arpa" in {
    type master;
    file "127.0.0.zone";
};

zone "." in {
    type hint;
    file "root.hint";
};
```

20.4.1 Important Configuration Options

`directory "filename";`

Specifies the directory in which BIND can find the files containing the zone data. Usually, this is `/var/lib/named`.

`forwarders { ip-address; };`

Specifies the name servers (mostly of the provider) to which DNS requests should be forwarded if they cannot be resolved directly. Replace `ip-address` with an IP address like `10.0.0.1`.

`forward first;`

Causes DNS requests to be forwarded before an attempt is made to resolve them via the root name servers. Instead of `forward first`, `forward only` can be written to have all requests forwarded and none sent to the root name servers. This makes sense for firewall configurations.

`listen-on port 53 { 127.0.0.1; ip-address; };`

Tells BIND on which network interfaces and port to accept client queries. `port 53` does not need to be specified explicitly, because 53 is the default port. Enter

`127.0.0.1` to permit requests from the local host. If you omit this entry entirely, all interfaces are used by default.

`listen-on-v6 port 53 {any;};`

Tells BIND on which port it should listen for IPv6 client requests. The only alternative to `any` is `none`. As far as IPv6 is concerned, the server only accepts a wild card address.

`query-source address * port 53;`

This entry is necessary if a firewall is blocking outgoing DNS requests. This tells BIND to post requests externally from port 53 and not from any of the high ports above 1024.

`query-source-v6 address * port 53;`

Tells BIND which port to use for IPv6 queries.

`allow-query { 127.0.0.1; net;};`

Defines the networks from which clients can post DNS requests. Replace `net` with address information like `192.168.1/24`. The `/24` at the end is an abbreviated expression for the netmask, in this case, `255.255.255.0`.

`allow-transfer ! *;;`

Controls which hosts can request zone transfers. In the example, such requests are completely denied with `! *.` Without this entry, zone transfers can be requested from anywhere without restrictions.

`statistics-interval 0;`

In the absence of this entry, BIND generates several lines of statistical information per hour in `/var/log/messages`. Set it to 0 to suppress these statistics completely or set an interval in minutes.

`cleaning-interval 720;`

This option defines at which time intervals BIND clears its cache. This triggers an entry in `/var/log/messages` each time it occurs. The time specification is in minutes. The default is 60 minutes.

`interface-interval 0;`

BIND regularly searches the network interfaces for new or nonexisting interfaces. If this value is set to 0, this is not done and BIND only listens at the interfaces detected at start-up. Otherwise, the interval can be defined in minutes. The default is sixty minutes.

notify no;

no prevents other name servers from being informed when changes are made to the zone data or when the name server is restarted.

20.4.2 Logging

What, how, and where logging takes place can be extensively configured in BIND. Normally, the default settings should be sufficient. [Example 20.3, “Entry to Disable Logging”](#) (page 381) shows the simplest form of such an entry and completely suppresses any logging.

Example 20.3 Entry to Disable Logging

```
logging {  
    category default { null; };  
};
```

20.4.3 Zone Entries

Example 20.4 Zone Entry for my-domain.de

```
zone "my-domain.de" in {  
    type master;  
    file "my-domain.zone";  
    notify no;  
};
```

After zone, specify the name of the domain to administer (my-domain.de) followed by in and a block of relevant options enclosed in curly braces, as shown in [Example 20.4, “Zone Entry for my-domain.de”](#) (page 381). To define a *slave zone*, switch the type to slave and specify a name server that administers this zone as master (which, in turn, may be a slave of another master), as shown in [Example 20.5, “Zone Entry for other-domain.de”](#) (page 381).

Example 20.5 Zone Entry for other-domain.de

```
zone "other-domain.de" in {  
    type slave;  
    file "slave/other-domain.zone";  
    masters { 10.0.0.1; };  
};
```

The zone options:

`type master;`

By specifying `master`, tell BIND that the zone is handled by the local name server. This assumes that a zone file has been created in the correct format.

`type slave;`

This zone is transferred from another name server. It must be used together with `masters`.

`type hint;`

The zone `.` of the `hint` type is used to set the root name servers. This zone definition can be left as is.

`file my-domain.zone` or file “`slave/other-domain.zone`”;

This entry specifies the file where zone data for the domain is located. This file is not required for a slave, because this data is fetched from another name server. To differentiate master and slave files, use the directory `slave` for the slave files.

`masters { server-ip-address; };`

This entry is only needed for slave zones. It specifies from which name server the zone file should be transferred.

`allow-update {! *; };`

This option controls external write access, which would allow clients to make a DNS entry—something not normally desirable for security reasons. Without this entry, zone updates are not allowed at all. The above entry achieves the same because `! *` effectively bans any such activity.

20.5 Zone Files

Two types of zone files are needed. One assigns IP addresses to hostnames and the other does the reverse: it supplies a hostname for an IP address.

TIP: Using the Dot in Zone Files

The `.` has an important meaning in the zone files. If hostnames are given without a final `.`, the zone is appended. Complete hostnames specified with a full domain name must end with a `.` to avoid having the domain added to it

again. A missing or wrongly placed dot is probably the most frequent cause of name server configuration errors.

The first case to consider is the zone file `world.zone`, responsible for the domain `world.cosmos`, shown in [Example 20.6, “File /var/lib/named/world.zone”](#) (page 383).

Example 20.6 *File /var/lib/named/world.zone*

```
$TTL 2D
world.cosmos. IN SOA      gateway  root.world.cosmos. (
    2003072441 ; serial
    1D          ; refresh
    2H          ; retry
    1W          ; expiry
    2D )        ; minimum

    IN NS      gateway
    IN MX      10 sun

gateway   IN A       192.168.0.1
           IN A       192.168.1.1
sun       IN A       192.168.0.2
moon      IN A       192.168.0.3
earth     IN A       192.168.1.2
mars      IN A       192.168.1.3
www       IN CNAME   moon
```

Line 1:

`$TTL` defines the default time to live that should apply to all the entries in this file.
In this example, entries are valid for a period of two days (2 D).

Line 2:

This is where the SOA (start of authority) control record begins:

- The name of the domain to administer is `world.cosmos` in the first position. This ends with a `.`, because otherwise the zone would be appended a second time. Alternatively, `@` can be entered here, in which case the zone would be extracted from the corresponding entry in `/etc/named.conf`.
- After `IN SOA` is the name of the name server in charge as master for this zone. The name is expanded from `gateway` to `gateway.world.cosmos`, because it does not end with a `..`.
- An e-mail address of the person in charge of this name server follows. Because the `@` sign already has a special meaning, `.` is entered here instead. For

`root@world.cosmos` the entry must read `root.world.cosmos..` The `.` must be included at the end to prevent the zone from being added.

- The `(` includes all lines up to `)` into the SOA record.

Line 3:

The `serial` number is an arbitrary number that is increased each time this file is changed. It is needed to inform the secondary name servers (slave servers) of changes. For this, a 10 digit number of the date and run number, written as `YYYYMMDDNN`, has become the customary format.

Line 4:

The `refresh rate` specifies the time interval at which the secondary name servers verify the zone `serial` number. In this case, one day.

Line 5:

The `retry rate` specifies the time interval at which a secondary name server, in case of error, attempts to contact the primary server again. Here, two hours.

Line 6:

The `expiration time` specifies the time frame after which a secondary name server discards the cached data if it has not regained contact to the primary server. Here, it is a week.

Line 7:

The last entry in the SOA record specifies the `negative caching TTL`—the time for which results of unresolved DNS queries from other servers may be cached.

Line 9:

The `IN NS` specifies the name server responsible for this domain. `gateway` is extended to `gateway.world.cosmos` because it does not end with a `..`. There can be several lines like this—one for the primary and one for each secondary name server. If `notify` is not set to `no` in `/etc/named.conf`, all the name servers listed here are informed of the changes made to the zone data.

Line 10:

The `MX` record specifies the mail server that accepts, processes, and forwards e-mails for the domain `world.cosmos`. In this example, this is the host `sun.world.cosmos`. The number in front of the hostname is the preference value. If there are multiple MX entries, the mail server with the smallest value is

taken first and, if mail delivery to this server fails, an attempt is made with the next higher value.

Lines 12–17:

These are the actual address records where one or more IP addresses are assigned to hostnames. The names are listed here without a . because they do not include their domain, so `world.cosmos` is added to all of them. Two IP addresses are assigned to the host `gateway`, because it has two network cards. Wherever the host address is a traditional one (IPv4), the record is marked with A. If the address is an IPv6 address, the entry is marked with A6. The previous token for IPv6 addresses was AAAA, which is now obsolete.

NOTE: A6 Syntax

The A6 record has a slightly different syntax than AAAA. Because of the fragmentation possibility, it is necessary to provide information about missed bits before the address. You must provide this information even if you want to use a completely unfragmented address. For the old AAAA record with the syntax

```
pluto IN      AAAA 2345:00C1:CA11:0001:1234:5678:9ABC:DEF0  
pluto IN      AAAA 2345:00D2:DA11:0001:1234:5678:9ABC:DEF0
```

You need to add information about missing bits in A6 format. Because the example above is complete (does not miss any bits), the A6 format of this record is:

```
pluto IN      AAAA 0 2345:00C1:CA11:0001:1234:5678:9ABC:DEF0  
pluto IN      AAAA 0 2345:00D2:DA11:0001:1234:5678:9ABC:DEF0
```

Do not use IPv4 addresses with IPv6 mapping. If a host has an IPv4 address, it uses an A record, not an A6.

Line 18:

The alias `www` can be used to address `mond` (`CNAME` means *canonical name*).

The pseudodomain `in-addr.arpa` is used for the reverse lookup of IP addresses into hostnames. It is appended to the network part of the address in reverse notation. So `192.168.1` is resolved into `1.168.192.in-addr.arpa`. See [Example 20.7, “Reverse Lookup”](#) (page 386).

Example 20.7 Reverse Lookup

```
$TTL 2D
1.168.192.in-addr.arpa. IN SOA gateway.world.cosmos. root.world.cosmos. (
    2003072441      ; serial
    1D              ; refresh
    2H              ; retry
    1W              ; expiry
    2D )            ; minimum

    IN NS           gateway.world.cosmos.

1          IN PTR       gateway.world.cosmos.
2          IN PTR       earth.world.cosmos.
3          IN PTR       mars.world.cosmos.
```

Line 1:

\$TTL defines the standard TTL that applies to all entries here.

Line 2:

The configuration file should activate reverse lookup for the network 192.168.1.0. Given that the zone is called 1.168.192.in-addr.arpa, should not be added to the hostnames. Therefore, all hostnames are entered in their complete form—with their domain and with a . at the end. The remaining entries correspond to those described for the previous world.cosmos example.

Lines 3–7:

See the previous example for world.cosmos.

Line 9:

Again this line specifies the name server responsible for this zone. This time, however, the name is entered in its complete form with the domain and a . at the end.

Lines 11–13:

These are the pointer records hinting at the IP addresses on the respective hosts. Only the last part of the IP address is entered at the beginning of the line, without the . at the end. Appending the zone to this (without the .in-addr.arpa) results in the complete IP address in reverse order.

Normally, zone transfers between different versions of BIND should be possible without any problem.

20.6 Dynamic Update of Zone Data

The term *dynamic update* refers to operations by which entries in the zone files of a master server are added, changed, or deleted. This mechanism is described in RFC 2136. Dynamic update is configured individually for each zone entry by adding an optional `allow-update` or `update-policy` rule. Zones to update dynamically should not be edited by hand.

Transmit the entries to update to the server with the command `nsupdate`. For the exact syntax of this command, check the manual page for `nsupdate` (`man 8 nsupdate`). For security reasons, any such update should be performed using TSIG keys as described in [Section 20.7, “Secure Transactions”](#) (page 387).

20.7 Secure Transactions

Secure transactions can be made with the help of transaction signatures (TSIGs) based on shared secret keys (also called TSIG keys). This section describes how to generate and use such keys.

Secure transactions are needed for communication between different servers and for the dynamic update of zone data. Making the access control dependent on keys is much more secure than merely relying on IP addresses.

Generate a TSIG key with the following command (for details, see `man dnssec-keygen`):

```
dnssec-keygen -a hmac-md5 -b 128 -n HOST host1-host2
```

This creates two files with names similar to these:

```
Khost1-host2.+157+34265.private Khost1-host2.+157+34265.key
```

The key itself (a string like `eJIKuCyyGJwwuN3xAteKgg==`) is found in both files. To use it for transactions, the second file (`Khost1-host2.+157+34265.key`) must be transferred to the remote host, preferably in a secure way (using `scp`, for example). On the remote server, the key must be included in the file `/etc/named.conf` to enable a secure communication between `host1` and `host2`:

```
key host1-host2. {
    algorithm hmac-md5;
```

```
secret "ejIkuCyyGJwwuN3xAtEkgg==;  
};
```

WARNING: File Permissions of /etc/named.conf

Make sure that the permissions of `/etc/named.conf` are properly restricted. The default for this file is 0640, with the owner being `root` and the group `named`. As an alternative, move the keys to an extra file with specially limited permissions, which is then included from `/etc/named.conf`.

To enable the server `host1` to use the key for `host2` (which has the address 192.168.2.3 in this example), the server's `/etc/named.conf` must include the following rule:

```
server 192.168.2.3 {  
    keys { host1-host2. ; };  
};
```

Analogous entries must be included in the configuration files of `host2`.

Add TSIG keys for any ACLs (access control lists, not to be confused with file system ACLs) that are defined for IP addresses and address ranges to enable transaction security. The corresponding entry could look like this:

```
allow-update { key host1-host2. ; };
```

This topic is discussed in more detail in the *BIND Administrator Reference Manual* under `update-policy`.

20.8 DNS Security

DNSSEC, or DNS security, is described in RFC 2535. The tools available for DNSSEC are discussed in the BIND Manual.

A zone considered secure must have one or several zone keys associated with it. These are generated with `dnssec-keygen`, just like the host keys. The DSA encryption algorithm is currently used to generate these keys. The public keys generated should be included in the corresponding zone file with an `$INCLUDE` rule.

With the command `dnssec-makekeyset`, all keys generated are packaged into one set, which must then be transferred to the parent zone in a secure manner. On the parent, the set is signed with `dnssec-signkey`. The files generated by this command are then used to sign the zones with `dnssec-signzone`, which in turn generates the files to include for each zone in `/etc/named.conf`.

20.9 For More Information

For additional information, refer to the *BIND Administrator Reference Manual* from package `bind-doc`, which is installed under `/usr/share/doc/packages/bind/`. Consider additionally consulting the RFCs referenced by the manual and the manual pages included with BIND. `/usr/share/doc/packages/bind/README` SuSE contains up-to-date information about BIND in SUSE Linux.

21

Using NIS

As soon as multiple UNIX systems in a network want to access common resources, it becomes important that all user and group identities are the same for all machines in that network. The network should be transparent to users: whatever machines they use, they always find themselves in exactly the same environment. This is made possible by means of NIS and NFS services. NFS distributes file systems over a network and is discussed in [Chapter 22, *Sharing File Systems with NFS*](#) (page 399).

NIS (Network Information Service) can be described as a database-like service that provides access to the contents of `/etc/passwd`, `/etc/shadow`, and `/etc/group` across networks. NIS can also be used for other purposes (making the contents of files like `/etc/hosts` or `/etc/services` available, for example), but this is beyond the scope of this introduction. People often refer to NIS as *YP*, because it works like the network's "yellow pages."

21.1 Configuring NIS Servers

To distribute NIS information across networks, you can either have one single server (*a master*) that serves to all clients or you can have NIS slave servers request this information from the master then relay it to their respective clients.

- To configure just one NIS server for your network, proceed with [Section 21.1.1, "Configuring a NIS Master Server"](#) (page 392).

- If your NIS master server should export its data to slave servers in other subnets, set up the master server as described in [Section 21.1.1, “Configuring a NIS Master Server”](#) (page 392) and set up slave servers in the subnets as described in [Section 21.1.2, “Configuring a NIS Slave Server”](#) (page 396).

21.1.1 Configuring a NIS Master Server

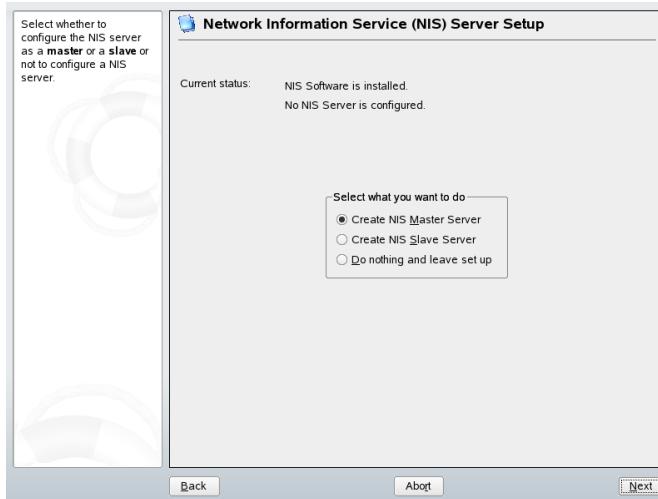
To configure a NIS master server for your network, proceed as follows:

- 1 Start *YaST* → *Network Services* → *NIS Server*.
- 2 If you need just one NIS server in your network or if this server is to act as the master for further NIS slave servers, select *Install and set up NIS Master Server*. YaST installs the required packages.

TIP

If NIS server software is already installed on your machine, initiate the creation of a NIS master server by clicking *Create NIS Master Server*.

Figure 21.1 *NIS Server Setup*



3 Determine basic NIS setup options:

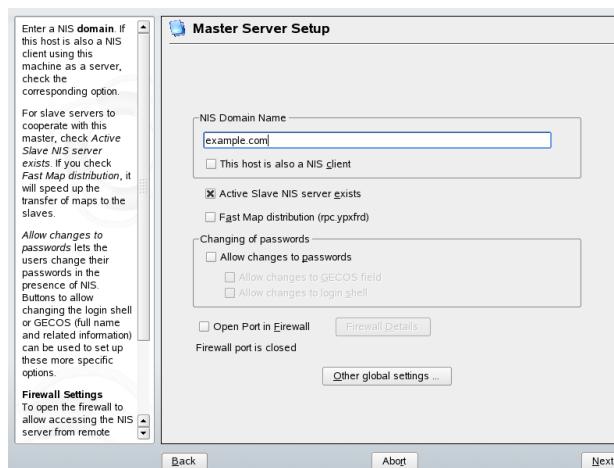
- a** Enter the NIS domain name.
- b** Define whether the host should also be a NIS client, enabling users to log in and access data from the NIS server, by selecting *This host is also a NIS client*.

Select *Changing of passwords* to allow users in your network (both local users and those managed through the NIS server) to change their passwords on the NIS server (with the command `yppasswd`).

This makes the options *Allow Changes to GECOS Field* and *Allow Changes to Login Shell* available. “GECOS” means that the users can also change their names and address settings with the command `ypchfn`. “SHELL” allows users to change their default shell with the command `ypchsh`, for example, to switch from `bash` to `sh`. The new shell must be one of the predefined entries in `/etc/shells`.

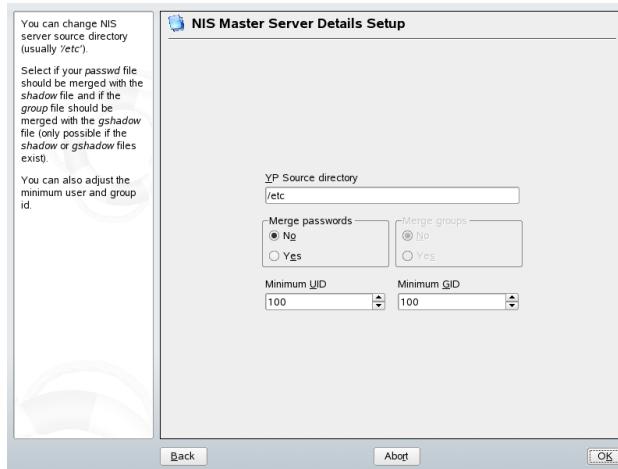
- c** If your NIS server should act as a master server to NIS slave servers in other subnets, select *Active Slave NIS Server exists*.
- d** Select *Open Ports in Firewall* to have YaST adapt the firewall settings for the NIS server.

Figure 21.2 Master Server Setup



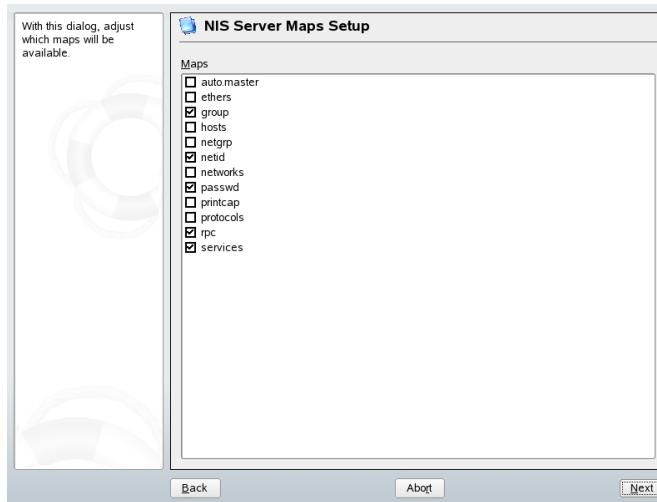
- e** Leave this dialog with *Next* or click *Other global settings* to make additional settings. *Other global settings* include changing the source directory of the NIS server (*/etc* by default). In addition, passwords can be merged here. The setting should be *Yes* so the files (*/etc/passwd*, */etc/shadow*, and */etc/group*) are used to build the user database. Also determine the smallest user and group ID that should be offered by NIS. Click *OK* to confirm your settings and return to the previous screen.

Figure 21.3 Changing the Directory and Synchronizing Files for a NIS Server



- 4 If you previously enabled *Active Slave NIS Server Exists*, enter the hostnames used as slaves and click *Next*.
- 5 If you do not use slave servers, the slave configuration is skipped and you continue directly to the dialog for the database configuration. Here, specify the *maps*, the partial databases to transfer from the NIS server to the client. The default settings are usually adequate. Leave this dialog with *Next*.
- 6 Check which maps should be available and click *Next* to continue.

Figure 21.4 NIS Server Maps Setup

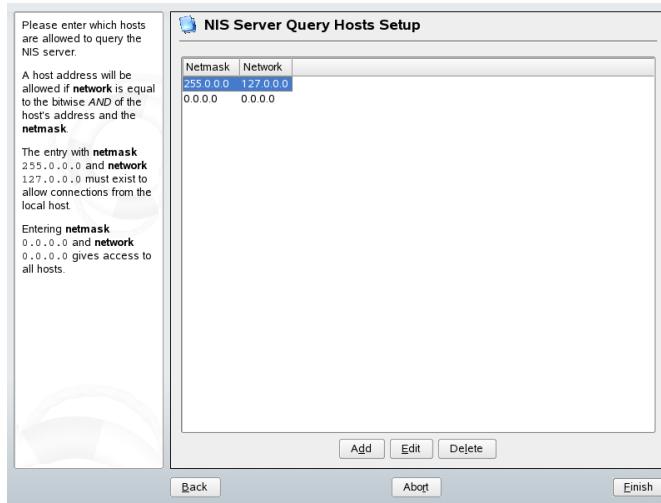


- 7 Enter the hosts that are allowed to query the NIS server. You can add, edit, or delete hosts by clicking the appropriate button. Specify from which networks requests can be sent to the NIS server. Normally, this is your internal network. In this case, there should be the following two entries:

255.0.0.0	127.0.0.0
0.0.0.0	0.0.0.0

The first entry enables connections from your own host, which is the NIS server. The second one allows all hosts to send requests to the server.

Figure 21.5 Setting Request Permissions for a NIS Server



- 8** Click *Finish* to save changes and exit the setup.

21.1.2 Configuring a NIS Slave Server

To configure additional NIS *slave servers* in your network, proceed as follows:

- 1** Start *YaST* → *Network Services* → *NIS Server*.
- 2** Select *Install and set up NIS Slave Server* and click *Next*.

TIP

If NIS server software is already installed on your machine, initiate the creation of a NIS slave server by clicking *Create NIS Slave Server*.

- 3** Complete the basic setup of your NIS slave server:
 - a** Enter the NIS domain.
 - b** Enter hostname or IP address of the master server.

- c** Set *This host is also a NIS client* if you want to enable user logins on this server.
 - d** Adapt the firewall settings with *Open Ports in Firewall*.
 - e** Click *Next*.
- 4** Enter the hosts that are allowed to query the NIS server. You can add, edit, or delete hosts by clicking the appropriate button. Specify from which networks requests can be sent to the NIS server. Normally, this is all hosts. In this case, there should be the following two entries:
- | | |
|-----------|-----------|
| 255.0.0.0 | 127.0.0.0 |
| 0.0.0.0 | 0.0.0.0 |
- The first entry enables connections from your own host, which is the NIS server. The second one allows all hosts with access to the same network to send requests to the server.
- 5** Click *Finish* to save changes and exit the setup.

21.2 Configuring NIS Clients

Use the module *NIS Client* to configure a workstation to use NIS. Select whether the host has a static IP address or receives one issued by DHCP. DHCP can also provide the NIS domain and the NIS server. For information about DHCP, see [Chapter 23, *DHCP*](#) (page 405). If a static IP address is used, specify the NIS domain and the NIS server manually. See [Figure 21.6, “Setting Domain and Address of a NIS Server”](#) (page 398). *Find* makes YaST search for an active NIS server in your whole network. Depending on the size of your local network, this may be a time-consuming process. *Broadcast* asks for a NIS server in the local network after the specified servers fail to respond.

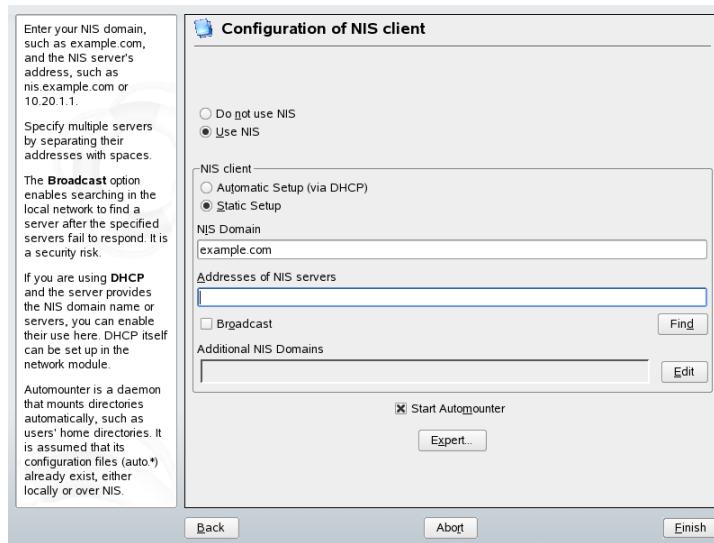
You can also specify multiple servers by entering their addresses in *Addresses of NIS servers* and separating them by spaces.

Depending on your local installation, you may also want to activate the automounter. This option also installs additional software if required.

In the expert settings, disable *Answer Remote Hosts* if you do not want other hosts to be able to query which server your client is using. By checking *Broken Server*, the client is enabled to receive replies from a server communicating through an unprivileged port. For further information, see man `ypbind`.

After you have made your settings, click *Finish* to save them and return to the YaST control center.

Figure 21.6 Setting Domain and Address of a NIS Server



Sharing File Systems with NFS

As mentioned in [Chapter 21, *Using NIS*](#) (page 391), NFS works with NIS to make a network transparent to the user. With NFS, it is possible to distribute file systems over the network. It does not matter at which terminal users are logged in. They always find themselves in the same environment.

Like NIS, NFS is a client/server system. A machine can be both—it can supply file systems over the network (export) and mount file systems from other hosts (import).

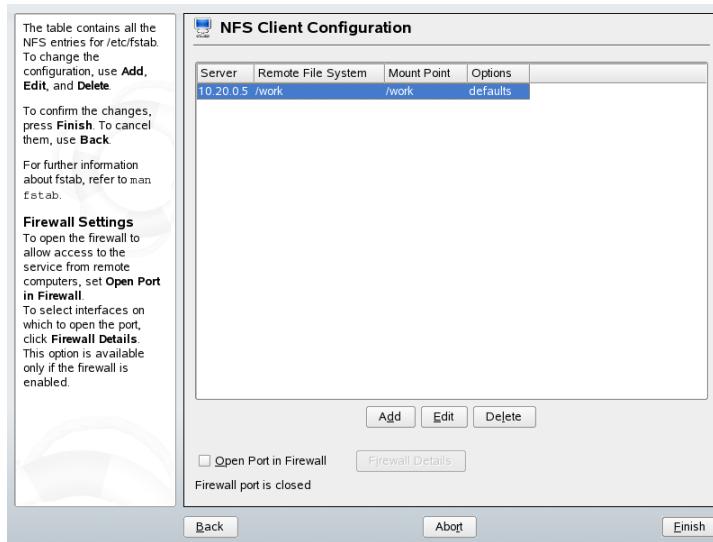
IMPORTANT: Need for DNS

In principle, all exports can be made using IP addresses only. To avoid time-outs, however, you should have a working DNS system. This is necessary at least for logging purposes, because the mountd daemon does reverse lookups.

22.1 Importing File Systems with YaST

Users authorized to do so can mount NFS directories from an NFS server into their own file trees. This can be achieved most easily using the YaST module *NFS Client*. Just enter the hostname of the NFS server, the directory to import, and the mount point at which to mount this directory locally. All this is done after *Add* is clicked in the first dialog. Click *Open Port in Firewall* to open the firewall to allow access to the service from remote computers. The firewall status is displayed next to the check box. Clicking *OK* saves your changes. See [Figure 22.1, “NFS Client Configuration with YaST”](#) (page 400).

Figure 22.1 NFS Client Configuration with YaST



22.2 Importing File Systems Manually

File systems can easily be imported manually from an NFS server. The only prerequisite is a running RPC port mapper, which can be started by entering the command `rcportmap start` as root. Once this prerequisite is met, remote file systems exported on the respective machines can be mounted in the file system just like local hard disks using the command `mount` with the following syntax:

```
mount host:remote-path local-path
```

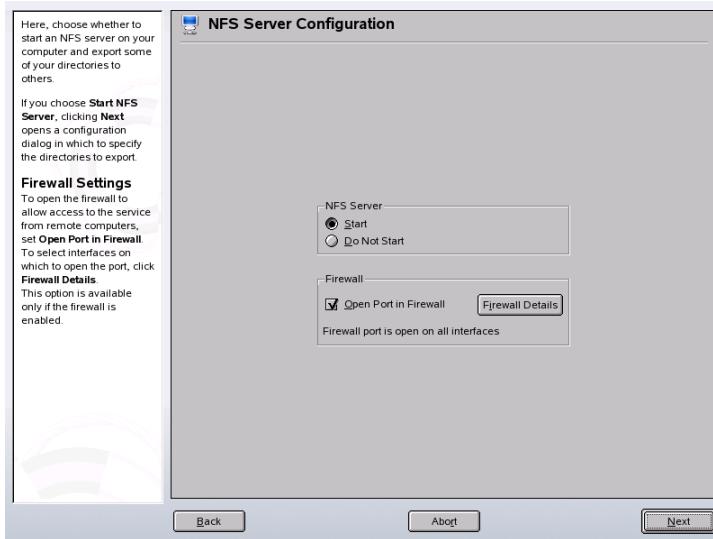
If user directories from the machine sun, for example, should be imported, use the following command:

```
mount sun:/home /home
```

22.3 Exporting File Systems with YaST

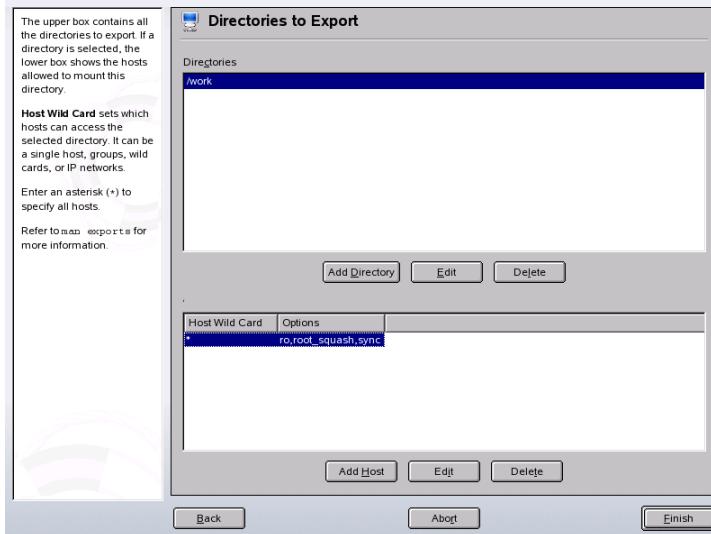
With YaST, turn a host in your network into an NFS server—a server that exports directories and files to all hosts granted access to it. This could be done to provide applications to all members of a group without installing them locally on each and every host. To install such a server, start YaST and select *Network Services* → *NFS Server*. A dialog like that in [Figure 22.2, “NFS Server Configuration Tool”](#) (page 401) opens.

Figure 22.2 NFS Server Configuration Tool



Next, activate *Start NFS Server* and click *Next*. In the upper text field, enter the directories to export. Below, enter the hosts that should have access to them. This dialog is shown in [Figure 22.3, “Configuring an NFS Server with YaST”](#) (page 402). There are four options that can be set for each host: single host, netgroups, wildcards, and IP networks. A more thorough explanation of these options is provided by *man exports*. *Exit* completes the configuration.

Figure 22.3 Configuring an NFS Server with YaST



IMPORTANT: Automatic Firewall Configuration

If a firewall is active on your system (SuSEfirewall2), YaST adapts its configuration for the NFS server by enabling the `nfs` service when *Open Ports in Firewall* is selected.

22.4 Exporting File Systems Manually

If you do not want to use YaST, make sure the following systems run on the NFS server:

- RPC portmapper (`portmap`)
- RPC mount daemon (`rpc.mountd`)
- RPC NFS daemon (`rpc.nfsd`)

For these services to be started by the scripts `/etc/init.d/portmap` and `/etc/init.d/nfsserver` when the system is booted, enter the commands `insserv /etc/init.d/nfsserver` and `insserv /etc/init.d/portmap`. Also define which file systems should be exported to which host in the configuration file `/etc(exports`.

For each directory to export, one line is needed to set which machines may access that directory with what permissions. All subdirectories of this directory are automatically exported as well. Authorized machines are usually specified with their full names (including domain name), but it is possible to use wild cards like * or ? (which expand the same way as in the Bash shell). If no machine is specified here, any machine is allowed to import this file system with the given permissions.

Set permissions for the file system to export in brackets after the machine name. The most important options are shown in [Table 22.1, “Permissions for Exported File System”](#) (page 403).

Table 22.1 *Permissions for Exported File System*

option	meaning
ro	The file system is exported with read-only permission (default).
rw	The file system is exported with read-write permission.
root_squash	This ensures that the user <code>root</code> of an importing machine does not have <code>root</code> permissions on this file system. This is achieved by assigning user ID 65534 to users with user ID 0 (<code>root</code>). This user ID should be set to <code>nobody</code> (which is the default).
no_root_squash	Does not assign user ID 0 to user ID 65534, keeping the <code>root</code> permissions valid.
link_relative	Converts absolute links (those beginning with <code>/</code>) to a sequence of <code>.. .</code> . This is only useful if the entire file system of a machine is mounted (default).

option	meaning
link_absolute	Symbolic links remain untouched.
map_identity	User IDs are exactly the same on both client and server (default).
map_daemon	Client and server do not have matching user IDs. This tells nfsd to create a conversion table for user IDs. The ugidd daemon is required for this to work.

Your `exports` file might look like [Example 22.1, “/etc\(exports”](#) (page 404). `/etc/exports` is read by `mountd` and `nfsd`. If you change anything in this file, restart `mountd` and `nfsd` for your changes to take effect. This can easily be done with `rcnfsserver restart`.

Example 22.1 `/etc/exports`

```
#  
# /etc/exports  
#  
/home          sun(rw)    venus(rw)  
/usr/X11        sun(ro)    venus(ro)  
/usr/lib/texmf  sun(ro)    venus(rw)  
/              earth(ro,root_squash)  
/home/ftp      (ro)  
# End of exports
```

22.5 For More Information

Information about configuring an NFS server is available in `/usr/share/doc/packages/nfs-utils/README` and the documents listed there. The detailed technical documentation is available online at <http://nfs.sourceforge.net/>

.

DHCP

The purpose of the *dynamic host configuration protocol* (DHCP) is to assign network settings centrally from a server rather than configuring them locally on each and every workstation. A host configured to use DHCP does not have control over its own static address. It is enabled to configure itself completely and automatically according to directions from the server. If you use the NetworkManager on the client side, you do not need to configure the client at all. This is useful if you have changing environments and only one interface active at a time. Never use NetworkManager on a machine that runs a DHCP server.

One way to configure a DHCP server is to identify each client using the hardware address of its network card (which is fixed in most cases), then supply that client with identical settings each time it connects to the server. DHCP can also be configured to assign addresses to each interested client dynamically from an address pool set up for that purpose. In the latter case, the DHCP server tries to assign the same address to the client each time it receives a request, even over longer periods. This works only if the network does not have more clients than addresses.

DHCP makes life easier for system administrators. Any changes, even bigger ones, related to addresses and the network configuration in general can be implemented centrally by editing the server's configuration file. This is much more convenient than reconfiguring numerous workstations. Also it is much easier to integrate machines, particularly new machines, into the network, because they can be given an IP address from the pool. Retrieving the appropriate network settings from a DHCP server is especially useful in the case of laptops regularly used in different networks.

A DHCP server supplies not only the IP address and the netmask, but also the hostname, domain name, gateway, and name server addresses for the client to use. In addition to that, DHCP allows a number of other parameters to be configured in a centralized way, for example, a time server from which clients may poll the current time or even a print server.

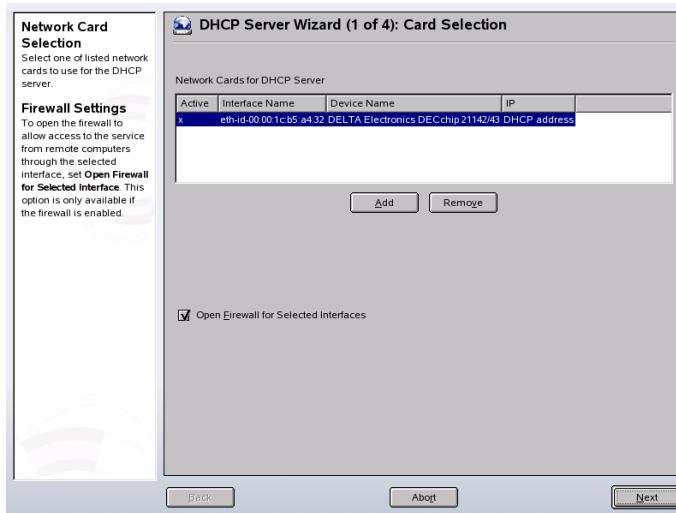
23.1 Configuring a DHCP Server with YaST

When the module is started for the first time, a wizard starts, prompting you to make a few basic decision concerning server administration. Completing this initial setup produces a very basic server configuration that should function in essential aspects. The expert mode can be used to deal with more advanced configuration tasks.

Card Selection

In the first step, YaST looks for the network interfaces available on your system then displays them in a list. From the list, select the interface on which the DHCP server should listen and click *Add*. After this, select *Open Firewall for Selected Interfaces* to open the firewall for this interface. See [Figure 23.1, “DHCP Server: Card Selection”](#) (page 407).

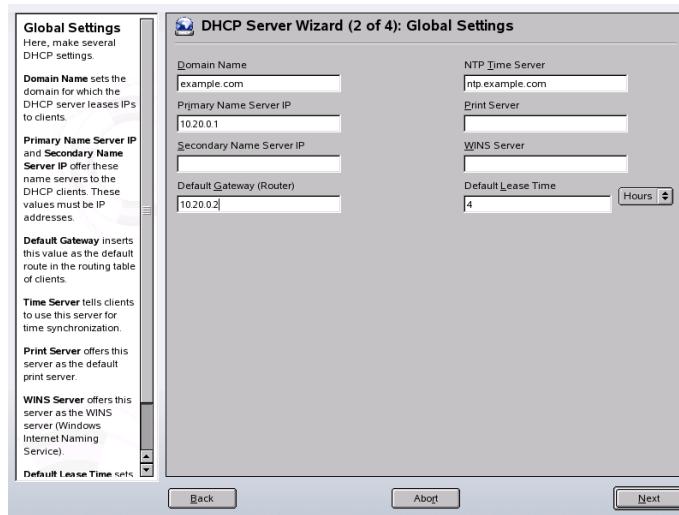
Figure 23.1 DHCP Server: Card Selection



Global Settings

In the entry fields, provide the network specifics for all clients the DHCP server should manage. These specifics are the domain name, address of a time server, addresses of the primary and secondary name server, addresses of a print and a WINS server (for a mixed network with both Windows and Linux clients), gateway address, and lease time. See [Figure 23.2, “DHCP Server: Global Settings” \(page 408\)](#).

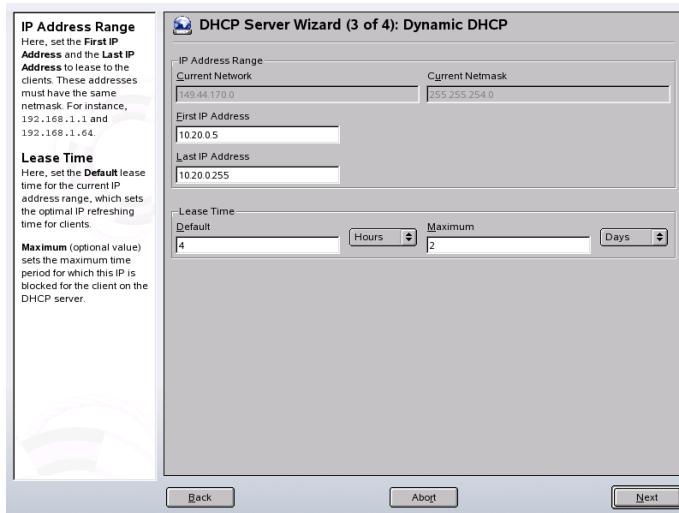
Figure 23.2 DHCP Server: Global Settings



Dynamic DHCP

In this step, configure how dynamic IP addresses should be assigned to clients. To do so, specify an IP range from which the server can assign addresses to DHCP clients. All these addresses must be covered by the same netmask. Also specify the lease time during which a client may keep its IP address without needing to request an extension of the lease. Optionally, specify the maximum lease time—the period during which the server reserves an IP address for a particular client. See [Figure 23.3, “DHCP Server: Dynamic DHCP”](#) (page 409).

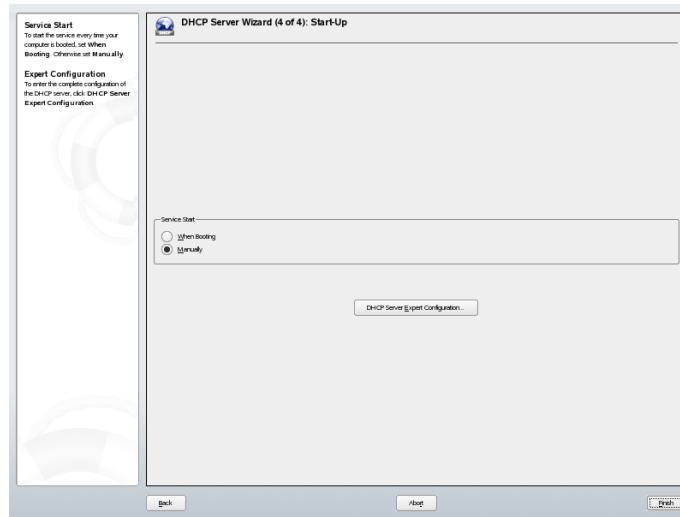
Figure 23.3 DHCP Server: Dynamic DHCP



Finishing the Configuration and Setting the Start Mode

After the third part of the configuration wizard, a last dialog is shown in which you can define how the DHCP server should be started. Here, specify whether to start the DHCP server automatically when the system is booted or manually when needed (for example, for test purposes). Click *Finish* to complete the configuration of the server. See [Figure 23.4, “DHCP Server: Start-Up” \(page 410\)](#).

Figure 23.4 DHCP Server: Start-Up



23.2 DHCP Software Packages

Both a DHCP server and DHCP clients are available for SUSE Linux. The DHCP server available is `dhcpd` (published by the Internet Software Consortium). On the client side, choose between two different DHCP client programs: `dhclient` (also from ISC) and the DHCP client daemon in the `dhpcd` package.

SUSE Linux installs `dhpcd` by default. The program is very easy to handle and is launched automatically on each system boot to watch for a DHCP server. It does not need a configuration file to do its job and works out of the box in most standard setups. For more complex situations, use the ISC `dhclient`, which is controlled by means of the configuration file `/etc/dhclient.conf`.

23.3 The DHCP Server `dhcpd`

The core of any DHCP system is the dynamic host configuration protocol daemon. This server *leases* addresses and watches how they are used, according to the settings defined in the configuration file `/etc/dhcpd.conf`. By changing the parameters and values in this file, a system administrator can influence the program's behavior in numerous ways. Look at the basic sample `/etc/dhcpd.conf` file in [Example 23.1, “The Configuration File `/etc/dhcpd.conf`”](#) (page 411).

Example 23.1 The Configuration File `/etc/dhcpd.conf`

```
default-lease-time 600;          # 10 minutes
max-lease-time 7200;           # 2 hours

option domain-name "cosmos.all";
option domain-name-servers 192.168.1.1, 192.168.1.2;
option broadcast-address 192.168.1.255;
option routers 192.168.1.254;
option subnet-mask 255.255.255.0;

subnet 192.168.1.0 netmask 255.255.255.0
{
    range 192.168.1.10 192.168.1.20;
    range 192.168.1.100 192.168.1.200;
}
```

This simple configuration file should be sufficient to get the DHCP server to assign IP addresses in the network. Make sure that a semicolon is inserted at the end of each line, because otherwise `dhcpd` is not started.

The sample file can be divided into three sections. The first one defines how many seconds an IP address is leased to a requesting client by default (`default-lease-time`) before it should apply for renewal. The section also includes a statement of the maximum period for which a machine may keep an IP address assigned by the DHCP server without applying for renewal (`max-lease-time`).

In the second part, some basic network parameters are defined on a global level:

- The line `option domain-name` defines the default domain of your network.
- With the entry `option domain-name-servers`, specify up to three values for the DNS servers used to resolve IP addresses into hostnames and vice versa. Ideally, configure a name server on your machine or somewhere else in your network before setting up DHCP. That name server should also define a hostname for each

dynamic address and vice versa. To learn how to configure your own name server, read [Chapter 20, *The Domain Name System*](#) (page 367).

- The line `option broadcast-address` defines the broadcast address the requesting client should use.
- With `option routers`, set where the server should send data packets that cannot be delivered to a host on the local network (according to the source and target host address and the subnet mask provided). In most cases, especially in smaller networks, this router is identical to the Internet gateway.
- With `option subnet-mask`, specify the netmask assigned to clients.

The last section of the file defines a network, including a subnet mask. To finish, specify the address range that the DHCP daemon should use to assign IP addresses to interested clients. In [Example 23.1, “The Configuration File /etc/dhcpd.conf”](#) (page 411), clients may be given any address between 192.168.1.10 and 192.168.1.20 as well as 192.168.1.100 and 192.168.1.200.

After editing these few lines, you should be able to activate the DHCP daemon with the command `rcdhcpd start`. It will be ready for use immediately. Use the command `rcdhcpd check-syntax` to perform a brief syntax check. If you encounter any unexpected problems with your configuration—the server aborts with an error or does not return `done` on `start`—you should be able to find out what has gone wrong by looking for information either in the main system log `/var/log/messages` or on console 10 (`Ctrl` + `Alt` + `F10`).

On a default SUSE Linux system, the DHCP daemon is started in a chroot environment for security reasons. The configuration files must be copied to the chroot environment so the daemon can find them. Normally, there is no need to worry about this because the command `rcdhcpd start` automatically copies the files.

23.3.1 Clients with Fixed IP Addresses

DHCP can also be used to assign a predefined, static address to a specific client. Addresses assigned explicitly always take priority over dynamic addresses from the pool. A static address never expires in the way a dynamic address would, for example, if there were not enough addresses available and the server needed to redistribute them among clients.

To identify a client configured with a static address, dhcpcd uses the hardware address, which is a globally unique, fixed numerical code consisting of six octet pairs for the identification of all network devices (for example, 00:00:45:12:EE:F4). If the respective lines, like the ones in [Example 23.2, “Additions to the Configuration File”](#) (page 413), are added to the configuration file of [Example 23.1, “The Configuration File /etc/dhcpcd.conf”](#) (page 411), the DHCP daemon always assigns the same set of data to the corresponding client.

Example 23.2 Additions to the Configuration File

```
host earth {  
hardware ethernet 00:00:45:12:EE:F4;  
fixed-address 192.168.1.21;  
}
```

The name of the respective client (`host hostname`, here `earth`) is entered in the first line and the MAC address in the second line. On Linux hosts, find the MAC address with the command `ip link show` followed by the network device (for example, `eth0`). The output should contain something like

```
link/ether 00:00:45:12:EE:F4
```

In the preceding example, a client with a network card having the MAC address 00:00:45:12:EE:F4 is assigned the IP address 192.168.1.21 and the hostname `earth` automatically. The type of hardware to enter is `ether` in nearly all cases, although `token-ring`, which is often found on IBM systems, is also supported.

23.3.2 The SUSE Linux Version

To improve security, the SUSE version of the ISC's DHCP server comes with the non-root/chroot patch by Ari Edelkind applied. This enables dhcpcd to run with the user ID `nobody` and run in a chroot environment (`/var/lib/dhcp`). To make this possible, the configuration file `dhcpcd.conf` must be located in `/var/lib/dhcp/etc`. The init script automatically copies the file to this directory when starting.

Control the server's behavior regarding this feature by means of entries in the file `/etc/sysconfig/dhcpcd`. To run dhcpcd without the chroot environment, set the variable `DHCPD_RUN_CHROOTED` in `/etc/sysconfig/dhcpcd` to “no”.

To enable dhcpcd to resolve hostnames even from within the chroot environment, some other configuration files must be copied as well:

- /etc/localtime
- /etc/host.conf
- /etc/hosts
- /etc/resolv.conf

These files are copied to `/var/lib/dhcp/etc/` when starting the init script. Take these copies into account for any changes that they require if they are dynamically modified by scripts like `/etc/ppp/ip-up`. However, there should be no need to worry about this if the configuration file only specifies IP addresses (instead of hostnames).

If your configuration includes additional files that should be copied into the chroot environment, set these under the variable `DHCPD_CONF_INCLUDE_FILES` in the file `/etc/sysconfig/dhcpd`. To ensure that the DHCP logging facility keeps working even after a restart of the `syslog-ng` daemon, there is an additional entry `SYSLOGD_ADDITIONAL_SOCKET_DHCP` in the file `/etc/sysconfig/syslog`.

23.4 For More Information

More information about DHCP is available at the Web site of the *Internet Software Consortium* (<http://www.isc.org/products/DHCP/>). Information is also available in the `dhcpd`, `dhcpd.conf`, `dhcpd.leases`, and `dhcp-options` `man` pages.

24

Time Synchronization with NTP

The NTP (network time protocol) mechanism is a protocol for synchronizing the system time over the network. First, a machine can obtain the time from a server that is a reliable time source. Second, a machine can itself act as a time source for other computers in the network. The goal is twofold—maintaining the absolute time and synchronizing the system time of all machines within a network.

Maintaining an exact system time is important in many situations. The built-in hardware (BIOS) clock does often not meet the requirements of applications like databases.

Manual correction of the system time would lead to severe problems because, for example, a backward leap can cause malfunction of critical applications. Within a network, it is usually necessary to synchronize the system time of all machines, but manual time adjustment is a bad approach. xntp provides a mechanism to solve these problems. It continuously adjusts the system time with the help of reliable time servers in the network. It further enables the management of local reference clocks, such as radio-controlled clocks.

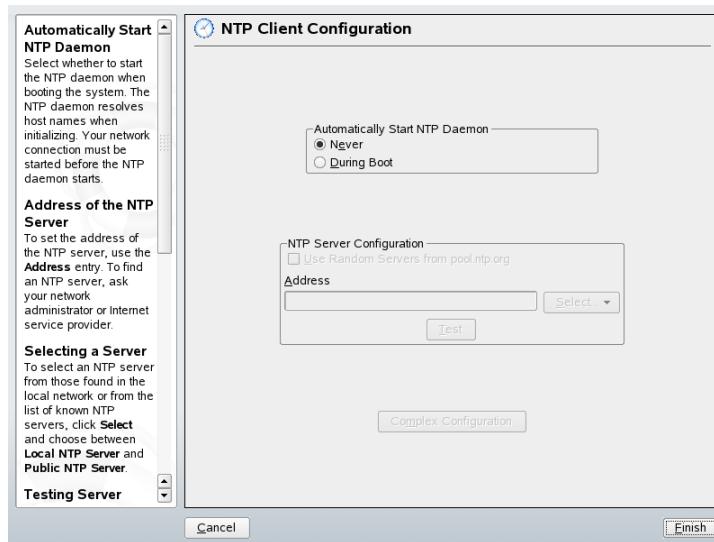
24.1 Configuring an NTP Client with YaST

xntp is preset to use the local computer clock as a time reference. Using the (BIOS) clock, however, only serves as a fallback for the case that no time source of greater precision is available. SUSE Linux facilitates the configuration of an NTP client with YaST. Use the quick or complex configuration for clients that do not run the SuSEfirewall because they are part of a protected intranet. Both are described in the following.

24.1.1 Quick NTP Client Configuration

The easy NTP client configuration (*Network Services → NTP Client*) consists of two dialogs. Set the start mode of `xntpd` and the server to query in the first dialog. To start `xntpd` automatically when the system is booted, click *During Boot*. Then specify the *NTP Server Configuration*. Either click *Use Random Server...*, if you cannot use a local time server, or click *Select* to access a second dialog in which to select a suitable time server for your network.

Figure 24.1 YaST: Configuring an NTP Client

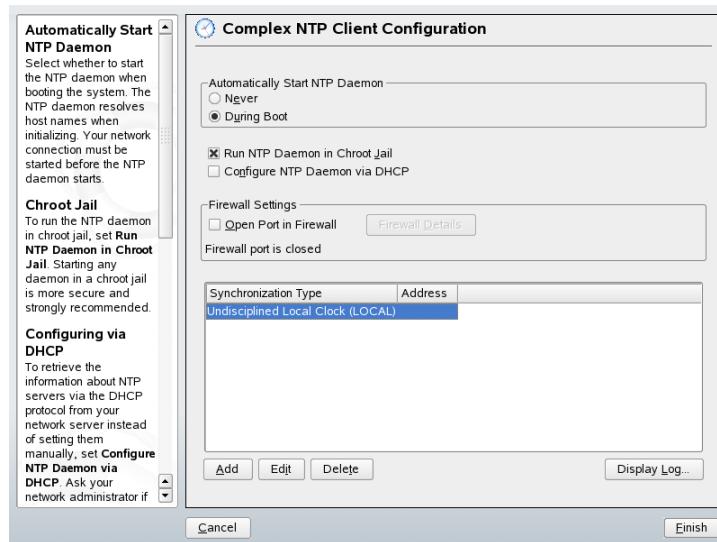


In the detailed server selection dialog, determine whether to implement time synchronization using a time server from your local network (*Local NTP Server*) or an Internet-based time server that takes care of your time zone (*Public NTP Server*). For a local time server, click *Lookup* to start an SLP query for available time servers in your network. Select the most suitable time server from the list of search results and exit the dialog with *OK*. For a public time server, select your country (time zone) and a suitable server from the list under *Public NTP Server* then exit the dialog with *OK*. In the main dialog, test the availability of the selected server with *Test* and quit the dialog with *Finish*.

24.1.2 Complex NTP Client Configuration

The complex configuration of an NTP client can be accessed under *Complex Configuration* from the main dialog of the *NTP Client* module, shown in [Figure 24.1, “YaST: Configuring an NTP Client”](#) (page 416), after selecting the start-up mode as described in the quick configuration.

Figure 24.2 YaST: Complex NTP Client Configuration



In *Complex NTP Client Configuration*, determine whether `xntpd` should be started in a chroot jail. This increases the security in the event of an attack over `xntpd`, because it prevents the attacker from compromising the entire system. *Configure NTP Daemon via DHCP* sets up the NTP client to get a list of the NTP servers available in your network via DHCP.

The servers and other time sources for the client to query are listed in the lower part. Modify this list as needed with *Add*, *Edit*, and *Delete*. *Display Log* provides the possibility to view the log files of your client.

Click *Add* to add a new source of time information. In the following dialog, select the type of source with which the time synchronization should be made. The following options are available:

Server

Another dialog enables you to select an NTP server (as described in [Section 24.1.1, “Quick NTP Client Configuration”](#) (page 416)). Activate *Use for Initial Synchronization* to trigger the synchronization of the time information between the server and the client when the system is booted. An input field allows you to specify additional options for xntpd. Refer to `/usr/share/doc/packages/xntp-doc` (part of the `xntp-doc` package) for detailed information.

Peer

A peer is a machine to which a symmetric relationship is established: it acts both as a time server and as a client. To use a peer in the same network instead of a server, enter the address of the system. The rest of the dialog is identical to the *Server* dialog.

Radio Clock

To use a radio clock in your system for the time synchronization, enter the clock type, unit number, device name, and other options in this dialog. Click *Driver Calibration* to fine-tune the driver. Detailed information about the operation of a local radio clock is available in `/usr/share/doc/packages/xntp-doc/html/refclock.htm`.

Outgoing Broadcast

Time information and queries can also be transmitted by broadcast in the network. In this dialog, enter the address to which such broadcasts should be sent. Do not activate broadcasting unless you have a reliable time source like a radio controlled clock.

Incoming Broadcast

If you want your client to receive its information via broadcast, enter the address from which the respective packets should be accepted in this fields.

24.2 Configuring xntp in the Network

The easiest way to use a time server in the network is to set server parameters. For example, if a time server called `ntp.example.com` is reachable from the network, add its name to the file `/etc/ntp.conf` by adding the line `server ntp.example.com`. To add more time servers, insert additional lines with the keyword `server`. After initializing xntpd with the command `rcntpd start`, it takes

about one hour until the time is stabilized and the drift file for correcting the local computer clock is created. With the drift file, the systematic error of the hardware clock can be computed as soon as the computer is powered on. The correction is used immediately, resulting in a higher stability of the system time.

There are two possible ways to use the NTP mechanism as a client: First, the client can query the time from a known server in regular intervals. With many clients, this approach can cause a high load on the server. Second, the client can wait for NTP broadcasts sent out by broadcast time servers in the network. This approach has the disadvantage that the quality of the server is unknown and a server sending out wrong information can cause severe problems.

If the time is obtained via broadcast, you do not need the server name. In this case, enter the line `broadcastclient` in the configuration file `/etc/ntp.conf`. To use one or more known time servers exclusively, enter their names in the line starting with `servers`.

24.3 Setting Up a Local Reference Clock

The software package `xntp` contains drivers for connecting local reference clocks. A list of supported clocks is available in the `xntp-doc` package in the file `/usr/share/doc/packages/xntp-doc/html/refclock.htm`. Every driver is associated with a number. In `xntp`, the actual configuration takes place by means of pseudo IPs. The clocks are entered in the file `/etc/ntp.conf` as though they existed in the network. For this purpose, they are assigned special IP addresses in the form `127.127.t.u`. Here, `t` stands for the type of the clock and determines which driver is used and `u` for unit, which determines the interface used.

Normally, the individual drivers have special parameters that describe configuration details. The file `/usr/share/doc/packages/xntp-doc/html/driverNN.htm` (where `NN` is the number of the driver) provides information about the particular type of clock. For example, the “type 8” clock (radio clock over serial interface) requires an additional mode that specifies the clock more precisely. The Conrad DCF77 receiver module, for example, has mode 5. To use this clock as a preferred reference, specify the keyword `prefer`. The complete server line for a Conrad DCF77 receiver module would be:

```
server 127.127.8.0 mode 5 prefer
```

Other clocks follow the same pattern. Following the installation of the `xntp-doc` package, the documentation for `xntp` is available in the directory `/usr/share/doc/packages/xntp-doc/html`. The file `/usr/share/doc/packages/xntp-doc/html/refclock.htm` provides links to the driver pages describing the driver parameters.

25

LDAP—A Directory Service

The Lightweight Directory Access Protocol (LDAP) is a set of protocols designed to access and maintain information directories. LDAP can be used for numerous purposes, like user and group management, system configuration management, or address management. This chapter provides a basic understanding of how OpenLDAP works and how to manage LDAP data with YaST. While there are several implementations of the LDAP protocol, this chapter focuses entirely on the OpenLDAP implementation.

It is crucial within a networked environment to keep important information structured and quickly available. This can be done with a directory service that, like the common yellow pages, keeps information available in a well-structured, quickly searchable form.

In the ideal case, a central server keeps the data in a directory and distributes it to all clients using a certain protocol. The data is structured in a way that allows a wide range of applications to access it. That way, it is not necessary for every single calendar tool and e-mail client to keep its own database—a central repository can be accessed instead. This notably reduces the administration effort for the information. The use of an open and standardized protocol like LDAP ensures that as many different client applications as possible can access such information.

A directory in this context is a type of database optimized for quick and effective reading and searching:

- To make numerous (concurrent) reading accesses possible, write access is limited to a small number of updates by the administrator. Conventional databases are optimized for accepting the largest possible data volume in a short time.

- Because write accesses can only be executed in a restricted fashion, a directory service is employed for administering mostly unchanging, static information. Data in a conventional database typically changes very often (*dynamic* data). Phone numbers in a company directory do not change nearly as often as, for example, the figures administered in accounting.
- When static data is administered, updates of the existing data sets are very rare. When working with dynamic data, especially when data sets like bank accounts or accounting are concerned, the consistency of the data is of primary importance. If an amount should be subtracted from one place to be added to another, both operations must happen concurrently, within a *transaction*, to ensure balance over the data stock. Databases support such transactions. Directories do not. Short-term inconsistencies of the data are quite acceptable in directories.

The design of a directory service like LDAP is not laid out to support complex update or query mechanisms. All applications accessing this service should gain access quickly and easily.

Many directory services have previously existed and still exist both in Unix and outside it. Novell NDS, Microsoft ADS, Banyan's Street Talk, and the OSI standard X.500 are just a few examples. LDAP was originally planned as a lean flavor of DAP, the directory access protocol, which was developed for accessing X.500. The X.500 standard regulates the hierarchical organization of directory entries.

LDAP is a trimmed down version of DAP. Without losing the X.500 entry hierarchy, profit from LDAP's cross-platform capabilities and save resources. The use of TCP/IP makes it substantially easier to establish interfaces between a docking application and the LDAP service.

LDAP, meanwhile, has evolved and is increasingly employed as a stand-alone solution without X.500 support. LDAP supports *referrals* with LDAPv3 (the protocol version in package `openldap2`), making it possible to have distributed databases. The usage of SASL (simple authentication and security layer) is also new.

LDAP is not limited to querying data from X.500 servers, as it was originally planned. There is an open source server `slapd`, which can store object information in a local database. There is also an extension called `slurpd`, which is responsible for replicating multiple LDAP servers.

The `openldap2` package consists of:

slapd

A stand-alone LDAPv3 server that administers object information in a BerkeleyDB-based database.

slurpd

This program enables the replication of modifications to data on the local LDAP server to other LDAP servers installed on the network.

additional tools for system maintenance

slapcat, slapadd, slapindex

25.1 LDAP versus NIS

The Unix system administrator traditionally uses the NIS service for name resolution and data distribution in a network. The configuration data contained in the files in /etc and the directories group, hosts, mail, netgroup, networks, passwd, printcap, protocols, rpc, and services are distributed by clients all over the network. These files can be maintained without major effort because they are simple text files. The handling of larger amounts of data, however, becomes increasingly difficult due to nonexistent structuring. NIS is only designed for Unix platforms. This means it is not suitable as a centralized data administration tool in heterogeneous networks.

Unlike NIS, the LDAP service is not restricted to pure Unix networks. Windows servers (from 2000) support LDAP as a directory service. Novell also offers an LDAP service. Application tasks mentioned above are additionally supported in non-Unix systems.

The LDAP principle can be applied to any data structure that should be centrally administered. A few application examples are:

- Employment as a replacement for the NIS service
- Mail routing (postfix, sendmail)
- Address books for mail clients, like Mozilla, Evolution, and Outlook
- Administration of zone descriptions for a BIND9 name server
- User authentication with Samba in heterogeneous networks

This list can be extended because LDAP is extensible, unlike NIS. The clearly-defined hierarchical structure of the data eases the administration of large amounts of data, because it can be searched better.

25.2 Structure of an LDAP Directory Tree

An LDAP directory has a tree structure. All entries (called objects) of the directory have a defined position within this hierarchy. This hierarchy is called the *directory information tree* (DIT). The complete path to the desired entry, which unambiguously identifies it, is called *distinguished name* or DN. A single node along the path to this entry is called *relative distinguished name* or RDN. Objects can generally be assigned to one of two possible types:

container

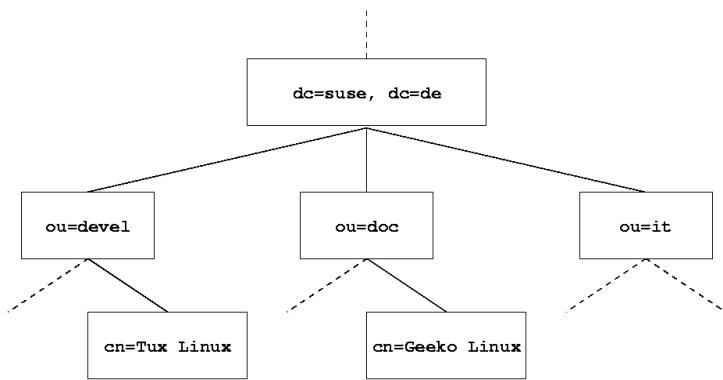
These objects can themselves contain other objects. Such object classes are `root` (the root element of the directory tree, which does not really exist), `c` (country), `ou` (organizational unit), and `dc` (domain component). This model is comparable to the directories (folders) in a file system.

leaf

These objects sit at the end of a branch and have no subordinate objects. Examples are `person`, `InetOrgPerson`, or `groupofNames`.

The top of the directory hierarchy has a root element `root`. This can contain `c` (country), `dc` (domain component), or `o` (organization) as subordinate elements. The relations within an LDAP directory tree become more evident in the following example, shown in [Figure 25.1, “Structure of an LDAP Directory”](#) (page 425).

Figure 25.1 Structure of an LDAP Directory



The complete diagram comprises a fictional directory information tree. The entries on three levels are depicted. Each entry corresponds to one box in the picture. The complete, valid *distinguished name* for the fictional SUSE employee Geeko Linux, in this case, is `cn=Geeko Linux, ou=doc, dc=suse, dc=de`. It is composed by adding the RDN `cn=Geeko Linux` to the DN of the preceding entry `ou=doc, dc=suse, dc=de`.

The global determination of which types of objects should be stored in the DIT is done following a *scheme*. The type of an object is determined by the *object class*. The object class determines what attributes the concerned object must or can be assigned. A scheme, therefore, must contain definitions of all object classes and attributes used in the desired application scenario. There are a few common schemes (see RFC 2252 and 2256). It is, however, possible to create custom schemes or to use multiple schemes complementing each other if this is required by the environment in which the LDAP server should operate.

Table 25.1, “[Commonly Used Object Classes and Attributes](#)” (page 426) offers a small overview of the object classes from `core.schema` and `inetorgperson.schema` used in the example, including required attributes and valid attribute values.

Table 25.1 Commonly Used Object Classes and Attributes

Object Class	Meaning	Example Entry	Compulsory Attributes
dcObject	<i>domainComponent</i> (name components of the domain)	suse	dc
organizationalUnit	<i>organizationalUnit</i> (organizational unit)	doc	ou
inetOrgPerson	<i>inetOrgPerson</i> (person-related data for the intranet or Internet)	Gecko Linux	sn and cn

[Example 25.1, “Excerpt from schema.core”](#) (page 426) shows an excerpt from a schema directive with explanations (line numbering for explanatory reasons).

Example 25.1 Excerpt from schema.core

```
#1 attributetype (2.5.4.11 NAME ( 'ou' 'organizationalUnitName')
#2           DESC 'RFC2256: organizational unit this object belongs to'
#3           SUP name )

...
#4 objectclass ( 2.5.6.5 NAME 'organizationalUnit'
#5           DESC 'RFC2256: an organizational unit'
#6           SUP top STRUCTURAL
#7           MUST ou
#8 MAY (userPassword $ searchGuide $ seeAlso $ businessCategory
      $ x121Address $ registeredAddress $ destinationIndicator
      $ preferredDeliveryMethod $ telexNumber
      $ teletexTerminalIdentifier $ telephoneNumber
      $ internationaliSDNNumber $ facsimileTelephoneNumber
      $ street $ postOfficeBox $ postalCode $ postalAddress
      $ physicalDeliveryOfficeName
      $ st $ l $ description) )
...
```

The attribute type `organizationalUnitName` and the corresponding object class `organizationalUnit` serve as an example here. Line 1 features the name of the attribute, its unique OID (*object identifier*) (numerical), and the abbreviation of the attribute.

Line 2 gives brief description of the attribute with DESC. The corresponding RFC on which the definition is based is also mentioned here. SUP in line 3 indicates a superordinate attribute type to which this attribute belongs.

The definition of the object class organizationalUnit begins in line 4, like in the definition of the attribute, with an OID and the name of the object class. Line 5 features a brief description of the object class. Line 6, with its entry SUP top, indicates that this object class is not subordinate to another object class. Line 7, starting with MUST, lists all attribute types that *must* be used in conjunction with an object of the type organizationalUnit. Line 8, starting with MAY, lists all attribute types that are permitted in conjunction with this object class.

A very good introduction to the use of schemes can be found in the documentation of OpenLDAP. When installed, find it in /usr/share/doc/packages/openldap2/admin-guide/index.html.

25.3 Server Configuration with slapd.conf

Your installed system contains a complete configuration file for your LDAP server at /etc/openldap/slapd.conf. The single entries are briefly described here and necessary adjustments are explained. Entries prefixed with a hash (#) are inactive. This comment character must be removed to activate them.

25.3.1 Global Directives in slapd.conf

Example 25.2 *slapd.conf: Include Directive for Schemes*

```
include      /etc/openldap/schema/core.schema
include      /etc/openldap/schema/cosine.schema
include      /etc/openldap/schema/inetorgperson.schema
include      /etc/openldap/schema/rfc2307bis.schema
include      /etc/openldap/schema/yast.schema
```

This first directive in slapd.conf, shown in [Example 25.2, “slapd.conf: Include Directive for Schemes”](#) (page 427), specifies the scheme by which the LDAP directory is organized. The entry core.schema is compulsory. Additionally required schemes

are appended to this directive. Information can be found in the included OpenLDAP documentation.

Example 25.3 *slapd.conf: pidfile and argsfile*

```
pidfile /var/run/slapd/slapd.pid
argsfile /var/run/slapd/slapd.args
```

These two files contain the PID (process ID) and some of the arguments with which the `slapd` process is started. There is no need for modifications here.

Example 25.4 *slapd.conf: Access Control*

```
# Sample Access Control
#           Allow read access of root DSE
# Allow self write access
#           Allow authenticated users read access
#           Allow anonymous users to authenticate
# access to dn="" by * read
access to * by self write
          by users read
          by anonymous auth
#
# if no access controls are present, the default is:
#           Allow read by all
#
# rootdn can always write!
```

[Example 25.4, “slapd.conf: Access Control”](#) (page 428) is the excerpt from `slapd.conf` that regulates the access permissions for the LDAP directory on the server. The settings made here in the global section of `slapd.conf` are valid as long as no custom access rules are declared in the database-specific section. These would overwrite the global declarations. As presented here, all users have read access to the directory, but only the administrator (`rootdn`) can write to this directory. Access control regulation in LDAP is a highly complex process. The following tips can help:

- Every access rule has the following structure:

```
access to <what> by <who> <access>
```

- *what* is a placeholder for the object or attribute to which access is granted. Individual directory branches can be protected explicitly with separate rules. It is also possible to process regions of the directory tree with one rule by using regular expressions. `slapd` evaluates all rules in the order in which they are listed in the configuration file. More general rules should be listed after more specific ones—the first rule `slapd` regards as valid is evaluated and all following entries are ignored.

- *who* determines who should be granted access to the areas determined with *what*. Regular expressions may be used. `slapd` again aborts the evaluation of *who* after the first match, so more specific rules should be listed before the more general ones. The entries shown in [Table 25.2, “User Groups and Their Access Grants”](#) (page 429) are possible.

Table 25.2 User Groups and Their Access Grants

Tag	Scope
*	All users without exception
anonymous	Not authenticated (“anonymous”) users
users	Authenticated users
self	Users connected with the target object
dn.regex=<regex>	All users matching the regular expression

- *access* specifies the type of access. Use the options listed in [Table 25.3, “Types of Access”](#) (page 429).

Table 25.3 Types of Access

Tag	Scope of Access
none	No access
auth	For contacting the server
compare	To objects for comparison access
search	For the employment of search filters
read	Read access
write	Write access

slapd compares the access right requested by the client with those granted in `slapd.conf`. The client is granted access if the rules allow a higher or equal right than the requested one. If the client requests higher rights than those declared in the rules, it is denied access.

Example 25.5, “`slapd.conf: Example for Access Control`” (page 430) shows an example of a simple access control that can be arbitrarily developed using regular expressions.

Example 25.5 `slapd.conf: Example for Access Control`

```
access to dn.regex="ou=([^\,]+),dc=suse,dc=de"
by dn.regex="cn=administrator,ou=$1,dc=suse,dc=de" write
by user read
by * none
```

This rule declares that only its respective administrator has write access to an individual `ou` entry. All other authenticated users have read access and the rest of the world has no access.

TIP: Establishing Access Rules

If there is no `access to` rule or no matching `by` directive, access is denied. Only explicitly declared access rights are granted. If no rules are declared at all, the default principle is write access for the administrator and read access for the rest of the world.

Find detailed information and an example configuration for LDAP access rights in the online documentation of the installed `openldap2` package.

Apart from the possibility to administer access permissions with the central server configuration file (`slapd.conf`), there is access control information (ACI). ACI allows storage of the access information for individual objects within the LDAP tree. This type of access control is not yet common and is still considered experimental by the developers. Refer to <http://www.openldap.org/faq/data/cache/758.html> for information.

25.3.2 Database-Specific Directives in slapd.conf

Example 25.6 *slapd.conf: Database-Specific Directives*

```
database bdb
checkpoint      1024      5
cachesize       10000
suffix "dc=suse,dc=de"
rootdn "cn=admin,dc=suse,dc=de"
# Cleartext passwords, especially for the rootdn, should
# be avoided. See slappasswd(8) and slapd.conf(5) for details.
# Use of strong authentication encouraged.
rootpw secret
# The database directory MUST exist prior to running slapd AND
# should only be accessible by the slapd/tools. Mode 700 recommended.
directory /var/lib/ldap
# Indices to maintain
index objectClass eq
```

The type of database, a Berkeley database in this case, is determined in the first line of this section (see [Example 25.6, “slapd.conf: Database-Specific Directives”](#) (page 431)). `checkpoint` determines the amount of data (in kb) that is kept in the transaction log before it is written to the actual database and the time (in minutes) between two write actions. `cachesize` sets the number of objects kept in the database's cache. `suffix` determines for which portion of the LDAP tree this server should be responsible. The following `rootdn` determines who owns administrator rights to this server. The user declared here does not need to have an LDAP entry or exist as regular user. The administrator password is set with `rootpw`. Instead of using `secret` here, it is possible to enter the hash of the administrator password created by `slappasswd`. The `directory` directive indicates the directory (in the file system) where the database directories are stored on the server. The last directive, `index objectClass eq`, results in the maintenance of an index of all object classes. Attributes for which users search most often can be added here according to experience. Custom Access rules defined here for the database are used instead of the global Access rules.

25.3.3 Starting and Stopping the Servers

Once the LDAP server is fully configured and all desired entries have been made according to the pattern described in [Section 25.4, “Data Handling in the LDAP Directory”](#) (page 432), start the LDAP server as `root` by entering `rcldap start`. To stop the server manually, enter the command `rcldap stop`. Request the status of the running LDAP server with `rcldap status`.

The YaST runlevel editor, described in [Section 8.2.3, “Configuring System Services \(Runlevel\) with YaST”](#) (page 188), can be used to have the server started and stopped automatically on boot and halt of the system. It is also possible to create the corresponding links to the start and stop scripts with the `insserv` command from a command prompt as described in [Section 8.2.2, “Init Scripts”](#) (page 184).

25.4 Data Handling in the LDAP Directory

OpenLDAP offers a series of tools for the administration of data in the LDAP directory. The four most important tools for adding to, deleting from, searching through, and modifying the data stock are briefly explained below.

25.4.1 Inserting Data into an LDAP Directory

Once the configuration of your LDAP server in `/etc/openldap/lsapd.conf` is correct and ready to go (it features appropriate entries for `suffix`, `directory`, `rootdn`, `rootpw`, and `index`), proceed to entering records. OpenLDAP offers the `ldapadd` command for this task. If possible, add the objects to the database in bundles for practical reasons. LDAP is able to process the LDIF format (LDAP data interchange format) for this. An LDIF file is a simple text file that can contain an arbitrary number of pairs of attribute and value. Refer to the schema files declared in `slapd.conf` for the available object classes and attributes. The LDIF file for creating a rough framework for the example in [Figure 25.1, “Structure of an LDAP Directory”](#) (page 425) would look like that in [Example 25.7, “Example for an LDIF File”](#) (page 433).

Example 25.7 Example for an LDIF File

```
# The SUSE Organization
dn: dc=suse,dc=de
objectClass: dcObject
objectClass: organization
o: SUSE AG dc: suse

# The organizational unit development (devel)
dn: ou=devel,dc=suse,dc=de
objectClass: organizationalUnit
ou: devel

# The organizational unit documentation (doc)
dn: ou=doc,dc=suse,dc=de
objectClass: organizationalUnit
ou: doc

# The organizational unit internal IT (it)
dn: ou=it,dc=suse,dc=de
objectClass: organizationalUnit
ou: it
```

IMPORTANT: Encoding of LDIF Files

LDAP works with UTF-8 (Unicode). Umlauts must be encoded correctly. Use an editor that supports UTF-8, such as Kate or recent versions of Emacs. Otherwise, avoid umlauts and other special characters or use `recode` to recode the input to UTF-8.

Save the file with the `.ldif` suffix then pass it to the server with the following command:

```
ldapadd -x -D <dn of the administrator> -W -f <file>.ldif
```

`-x` switches off the authentication with SASL in this case. `-D` declares the user that calls the operation. The valid DN of the administrator is entered here just like it has been configured in `slapd.conf`. In the current example, this is `cn=admin,dc=suse,dc=de`. `-W` circumvents entering the password on the command line (in clear text) and activates a separate password prompt. This password was previously determined in `slapd.conf` with `rootpw`. `-f` passes the filename. See the details of running `ldapadd` in [Example 25.8, “`ldapadd` with `example.ldif`”](#) (page 434).

Example 25.8 *ldapadd with example.ldif*

```
ldapadd -x -D cn=admin,dc=suse,dc=de -W -f example.ldif
```

```
Enter LDAP password:  
adding new entry "dc=suse,dc=de"  
adding new entry "ou=devel,dc=suse,dc=de"  
adding new entry "ou=doc,dc=suse,dc=de"  
adding new entry "ou=it,dc=suse,dc=de"
```

The user data of individuals can be prepared in separate LDIF files. [Example 25.9, “LDIF Data for Tux”](#) (page 434) adds Tux to the new LDAP directory.

Example 25.9 *LDIF Data for Tux*

```
# coworker Tux  
dn: cn=Tux Linux,ou=devel,dc=suse,dc=de  
objectClass: inetOrgPerson  
cn: Tux Linux  
givenName: Tux  
sn: Linux  
mail: tux@suse.de  
uid: tux  
telephoneNumber: +49 1234 567-8
```

An LDIF file can contain an arbitrary number of objects. It is possible to pass entire directory branches to the server at once or only parts of it as shown in the example of individual objects. If it is necessary to modify some data relatively often, a fine subdivision of single objects is recommended.

25.4.2 Modifying Data in the LDAP Directory

The tool `ldapmodify` is provided for modifying the data stock. The easiest way to do this is to modify the corresponding LDIF file then pass this modified file to the LDAP server. To change the telephone number of colleague Tux from +49 1234 567-8 to +49 1234 567-10, edit the LDIF file like in [Example 25.10, “Modified LDIF File tux.ldif”](#) (page 434).

Example 25.10 *Modified LDIF File tux.ldif*

```
# coworker Tux  
dn: cn=Tux Linux,ou=devel,dc=suse,dc=de  
changetype: modify  
replace: telephoneNumber  
telephoneNumber: +49 1234 567-10
```

Import the modified file into the LDAP directory with the following command:

```
ldapmodify -x -D cn=admin,dc=suse,dc=de -W -f tux.ldif
```

Alternatively, pass the attributes to change directly to `ldapmodify`. The procedure for this is described below:

1. Start `ldapmodify` and enter your password:

```
ldapmodify -x -D cn=admin,dc=suse,dc=de -W  
Enter LDAP password:
```

2. Enter the changes while carefully complying with the syntax in the order presented below:

```
dn: cn=Tux Linux,ou=devel,dc=suse,dc=de  
changetype: modify  
replace: telephoneNumber  
telephoneNumber: +49 1234 567-10
```

Find detailed information about `ldapmodify` and its syntax in the `ldapmodify(1)` man page.

25.4.3 Searching or Reading Data from an LDAP Directory

OpenLDAP provides, with `ldapsearch`, a command line tool for searching data within an LDAP directory and reading data from it. A simple query would have the following syntax:

```
ldapsearch -x -b dc=suse,dc=de "(objectClass=*)"
```

The option `-b` determines the search base—the section of the tree within which the search should be performed. In the current case, this is `dc=suse,dc=de`. To perform a more finely-grained search in specific subsections of the LDAP directory (for example, only within the `devel` department), pass this section to `ldapsearch` with `-b`. `-x` requests activation of simple authentication. `(objectClass=*)` declares that all objects contained in the directory should be read. This command option can be used after the creation of a new directory tree to verify that all entries have been recorded correctly and the server responds as desired. More information about the use of `ldapsearch` can be found in the corresponding man page (`ldapsearch(1)`).

25.4.4 Deleting Data from an LDAP Directory

Delete unwanted entries with `ldapdelete`. The syntax is similar to that of the commands described above. To delete, for example, the complete entry for Tux Linux, issue the following command:

```
ldapdelete -x -D cn=admin,dc=suse,dc=de -W cn=Tux \
Linux,ou=devel,dc=suse,dc=de
```

25.5 The YaST LDAP Client

YaST includes a module to set up LDAP-based user management. If you did not enable this feature during the installation, start the module by selecting *Network Services → LDAP Client*. YaST automatically enables any PAM and NSS related changes as required by LDAP (described below) and installs the necessary files.

25.5.1 Standard Procedure

Background knowledge of the processes acting in the background of a client machine helps you understand how the YaST LDAP client module works. If LDAP is activated for network authentication or the YaST module is called, the packages `pam_ldap` and `nss_ldap` are installed and the two corresponding configuration files are adapted. `pam_ldap` is the PAM module responsible for negotiation between login processes and the LDAP directory as the source of authentication data. The dedicated module `pam_ldap.so` is installed and the PAM configuration is adapted (see [Example 25.11, “pam_unix2.conf Adapted to LDAP” \(page 436\)](#)).

Example 25.11 pam_unix2.conf Adapted to LDAP

```
auth:      use_ldap
account:   use_ldap
password:  use_ldap
session:   none
```

When manually configuring additional services to use LDAP, include the PAM LDAP module in the PAM configuration file corresponding to the service in `/etc/pam.d`. Configuration files already adapted to individual services can be found in `/usr/share/doc/packages/pam_ldap/pam.d/`. Copy appropriate files to `/etc/pam.d`.

glibc name resolution through the `nsswitch` mechanism is adapted to the employment of LDAP with `nss_ldap`. A new, adapted file `nsswitch.conf` is created in `/etc/` with the installation of this package. More about the workings of `nsswitch.conf` can be found in [Section 18.6.1, “Configuration Files”](#) (page 351). The following lines must be present in `nsswitch.conf` for user administration and authentication with LDAP. See [Example 25.12, “Adaptations in nsswitch.conf”](#) (page 437).

Example 25.12 *Adaptations in nsswitch.conf*

```
passwd: compat  
group: compat  
  
passwd_compat: ldap  
group_compat: ldap
```

These lines order the resolver library of glibc first to evaluate the corresponding files in `/etc` and additionally access the LDAP server as sources for authentication and user data. Test this mechanism, for example, by reading the content of the user database with the command `getent passwd`. The returned set should contain a survey of the local users of your system as well as all users stored on the LDAP server.

To prevent regular users managed through LDAP from logging in to the server with `ssh` or `login`, the files `/etc/passwd` and `/etc/group` each need to include an additional line. This is the line `+::::::/sbin/nologin` in `/etc/passwd` and `+::: in /etc/group`.

25.5.2 Configuration of the LDAP Client

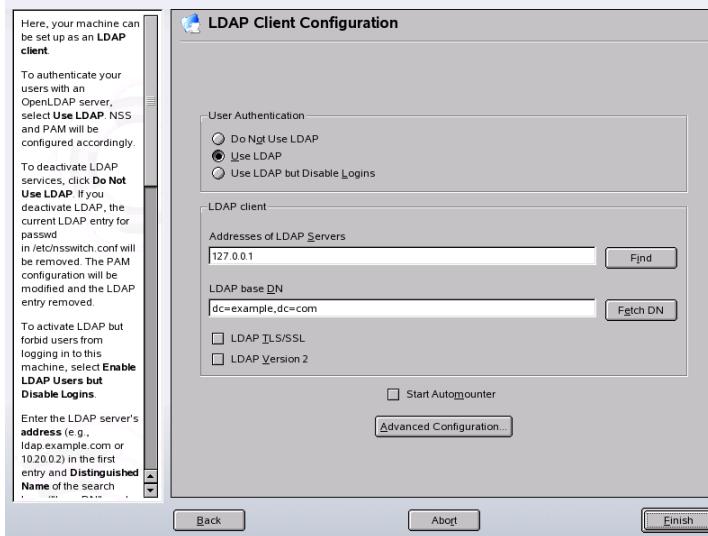
After the initial adjustments of `nss_ldap`, `pam_ldap`, `/etc/passwd`, and `/etc/group` have been taken care of by YaST, you can simply connect your client to the server and let YaST do user management via LDAP. This basic setup is described in [Section “Basic Configuration”](#) (page 438).

Use the YaST LDAP client to further configure the YaST group and user configuration modules. This includes manipulating the default settings for new users and groups and the number and nature of the attributes assigned to a user or a group. LDAP user management allows you to assign far more and different attributes to users and groups than traditional user or group management solutions. This is described in [Section “Configuring the YaST Group and User Administration Modules”](#) (page 441).

Basic Configuration

The basic LDAP client configuration dialog ([Figure 25.2, “YaST: Configuration of the LDAP Client”](#) (page 438)) opens during installation if you choose LDAP user management or when you select *Network Services → LDAP Client* in the YaST Control Center in the installed system.

Figure 25.2 YaST: Configuration of the LDAP Client

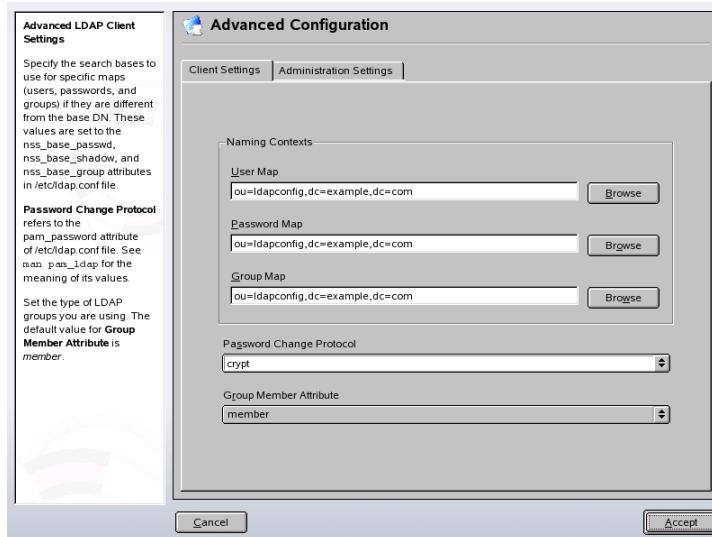


To authenticate users of your machine against an OpenLDAP server and enable user management via OpenLDAP, proceed as follows:

- 1 Click *Use LDAP* to enable the use of LDAP. Select *Use LDAP but Disable Logins* instead if you want to use LDAP for authentication, but do not want other users to log in to this client.

- 2** Enter the IP address of the LDAP server to use.
- 3** Enter the *LDAP base DN* to select the search base on the LDAP server.
If you want to retrieve the base DN automatically, click *Fetch DN*. YaST then checks for any LDAP database on the server address specified above. Choose the appropriate base DN from the search results given by YaST.
- 4** If TLS or SSL protected communication with the server is required, select *LDAP TLS/SSL*.
- 5** If the LDAP server still uses LDAPv2, explicitly enable the use of this protocol version by selecting *LDAP Version 2*.
- 6** Select *Start Automounter* to mount remote directories on your client, such as a remotely managed /home.
- 7** Click *Finish* to apply your settings.

Figure 25.3 YaST: Advanced Configuration



To modify data on the server as administrator, click *Advanced Configuration*. The following dialog is split in two tabs. See [Figure 25.3, “YaST: Advanced Configuration”](#) (page 439):

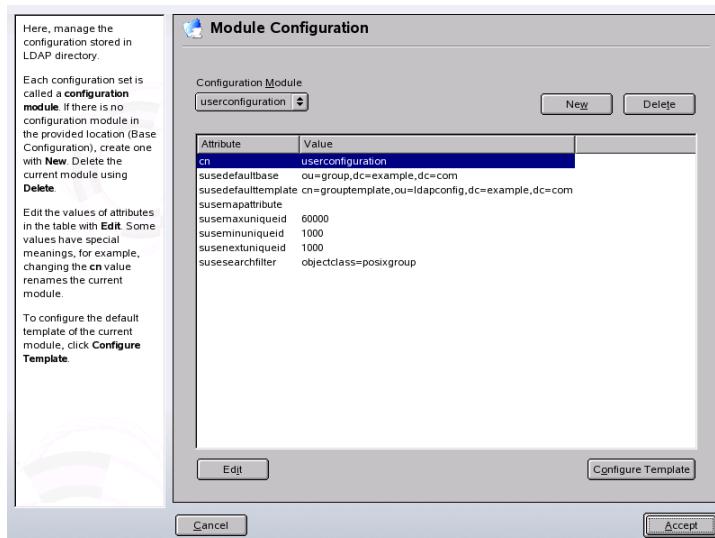
- 1** In the *Client Settings* tab, adjust the following settings:
 - a** If the search base for users, passwords, and groups differs from the global search base specified the *LDAP base DN*, enter these different naming contexts in *User Map*, *Password Map*, and *Group Map*.
 - b** Specify the password change protocol. The standard method to use whenever a password is changed is `crypt`, meaning that password hashes generated by `crypt` are used. For details on this and other options, refer to the `pam_ldap` man page.
 - c** Specify the LDAP group to use with *Group Member Attribute*. The default value for this is `member`.
- 2** In *Administration Settings*, adjust the following settings:
 - a** Set the base for storing your user management data via *Configuration Base DN*.
 - b** Enter the appropriate value for *Administrator DN*. This DN must be identical with the `rootdn` value specified in `/etc/openldap/slapd.conf` to enable this particular user to manipulate data stored on the LDAP server. Enter the full DN (such as `cn=admin,dc=suse,dc=de`) or activate *Append Base DN* to have the base DN added automatically when you enter `cn=admin`.
 - c** Check *Create Default Configuration Objects* to create the basic configuration objects on the server to enable user management via LDAP.
 - d** If your client machine should act as a file server for home directories across your network, check *Home Directories on This Machine*.
 - e** Click *Accept* to leave the *Advanced Configuration* then *Finish* to apply your settings.

Use *Configure User Management Settings* to edit entries on the LDAP server. Access to the configuration modules on the server is then granted according to the ACLs and ACIs stored on the server. Follow the procedures outlined in [Section “Configuring the YaST Group and User Administration Modules”](#) (page 441).

Configuring the YaST Group and User Administration Modules

Use the YaST LDAP client to adapt the YaST modules for user and group administration and to extend them as needed. Define templates with default values for the individual attributes to simplify the data registration. The presets created here are stored as LDAP objects in the LDAP directory. The registration of user data is still done with the regular YaST modules for user and group management. The registered data is stored as LDAP objects on the server.

Figure 25.4 YaST: Module Configuration



The dialog for module configuration (Figure 25.4, “YaST: Module Configuration” (page 441)) allows the creation of new modules, selection and modification of existing configuration modules, and design and modification of templates for such modules.

To create a new configuration module, proceed as follows:

- 1 Click *New* and select the type of module to create. For a user configuration module, select `suseuserconfiguration` and for a group configuration choose `susegroupconfiguration`.
- 2 Choose a name for the new template.

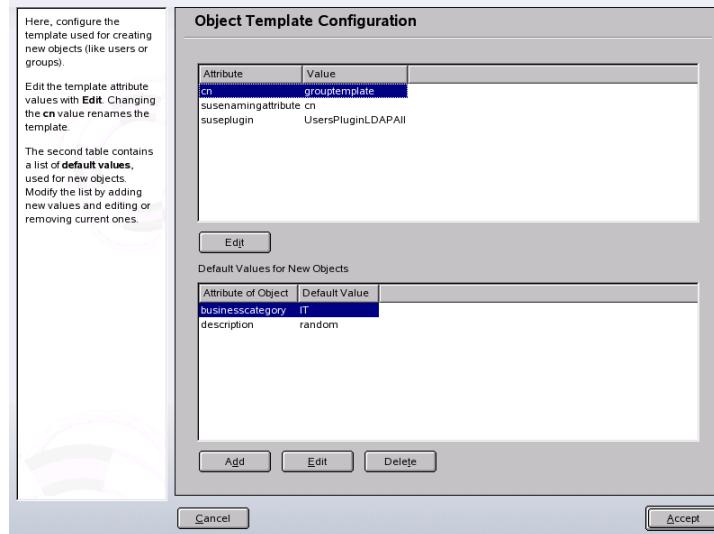
The content view then features a table listing all attributes allowed in this module with their assigned values. Apart from all set attributes, the list also contains all other attributes allowed by the current schema but currently not used.

- 3 Accept the preset values or adjust the defaults to use in group and user configuration by selecting the respective attribute, pressing *Edit*, and entering the new value. Rename a module by simply changing the `cn` attribute of the module. Clicking *Delete* deletes the currently selected module.
- 4 After you click *OK*, the new module is added to the selection menu.

The YaST modules for group and user administration embed templates with sensible standard values. To edit a template associated with a configuration module, proceed as follows:

- 1 In the *Module Configuration* dialog, click *Configure Template*.
- 2 Determine the values of the general attributes assigned to this template according to your needs or leave some of them empty. Empty attributes are deleted on the LDAP server.
- 3 Modify, delete, or add new default values for new objects (user or group configuration objects in the LDAP tree).

Figure 25.5 YaST: Configuration of an Object Template



Connect the template to its module by setting the `susedefaulttemplate` attribute value of the module to the DN of the adapted template.

TIP

The default values for an attribute can be created from other attributes by using a variable instead of an absolute value. For example, when creating a new user, `cn=%sn %givenName` is created automatically from the attribute values for `sn` and `givenName`.

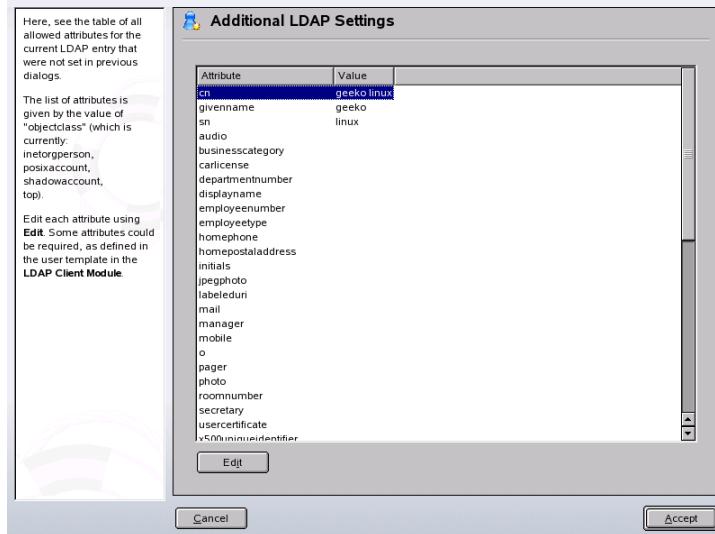
Once all modules and templates are configured correctly and ready to run, new groups and users can be registered in the usual way with YaST.

25.6 Configuring LDAP Users and Groups in YaST

The actual registration of user and group data differs only slightly from the procedure when not using LDAP. The following brief instructions relate to the administration of users. The procedure for administering groups is analogous.

- 1** Access the YaST user administration with *Security & Users* → *User Administration*.
- 2** Use *Set Filter* to limit the view of users to the LDAP users and enter the password for Root DN.
- 3** Click *Add* and enter the configuration of a new user. A dialog with four tabs opens:
 - a** Specify username, login, and password in the *User Data* tab.
 - b** Check the *Details* tab for the group membership, login shell, and home directory of the new user. If necessary, change the default to values that better suit your needs. The default values as well as those of the password settings can be defined with the procedure described in [Section “Configuring the YaST Group and User Administration Modules”](#) (page 441).
 - c** Modify or accept the default *Password Settings*.
 - d** Enter the *Plug-Ins* tab, select the LDAP plug-in, and click *Launch* to configure additional LDAP attributes assigned to the new user (see [Figure 25.6, “YaST: Additional LDAP Settings”](#) (page 445)).
- 4** Click *Accept* to apply your settings and leave the user configuration.

Figure 25.6 YaST: Additional LDAP Settings



The initial input form of user administration offers *LDAP Options*. This gives the possibility to apply LDAP search filters to the set of available users or go to the module for the configuration of LDAP users and groups by selecting *LDAP User and Group Configuration*.

25.7 For More Information

More complex subjects, like SASL configuration or establishment of a replicating LDAP server that distributes the workload among multiple slaves, were intentionally not included in this chapter. Detailed information about both subjects can be found in the *OpenLDAP 2.2 Administrator's Guide* (references follow).

The Web site of the OpenLDAP project offers exhaustive documentation for beginning and advanced LDAP users:

OpenLDAP Faq-O-Matic

A very rich question and answer collection concerning installation, configuration, and use of OpenLDAP. Find it at <http://www.openldap.org/faq/data/cache/1.html>.

Quick Start Guide

Brief step-by-step instructions for installing your first LDAP server. Find it at <http://www.openldap.org/doc/admin22/quickstart.html> or on an installed system in /usr/share/doc/packages/openldap2/admin-guide/quickstart.html.

OpenLDAP 2.2 Administrator's Guide

A detailed introduction to all important aspects of LDAP configuration, including access controls and encryption. See <http://www.openldap.org/doc/admin22/> or, on an installed system, /usr/share/doc/packages/openldap2/admin-guide/index.html.

Understanding LDAP

A detailed general introduction to the basic principles of LDAP: <http://www.redbooks.ibm.com/redbooks/pdfs/sg244986.pdf>.

Printed literature about LDAP:

- *LDAP System Administration* by Gerald Carter (ISBN 1-56592-491-6)
- *Understanding and Deploying LDAP Directory Services* by Howes, Smith, and Good (ISBN 0-672-32316-8)

The ultimate reference material for the subject of LDAP is the corresponding RFCs (request for comments), 2251 to 2256.

26

The Apache HTTP Server

With a share of more than 70%, the Apache HTTP Server (Apache) is the world's most widely-used Web server according to the November 2005 Survey from <http://www.netcraft.com/>. Apache, developed by the Apache Software Foundation (<http://www.apache.org/>), is available for most operating systems. SUSE Linux includes Apache version 2.2. In this chapter, learn how to install, configure and set up a Web server; how to use SSL, CGI, and additional modules; and how to troubleshoot Apache.

26.1 Quick Start

With the help of this section, quickly set up and start Apache. You must be `root` to install and configure Apache.

26.1.1 Requirements

Make sure that the following requirements are met before trying to set up the Apache Web server:

1. The machine's network is configured properly. For more information about this topic, refer to [Chapter 18, Basic Networking](#) (page 317).

2. The machine's exact system time is maintained by synchronizing with a time server. This is necessary because parts of the HTTP protocol depend on the correct time. See [Chapter 24, Time Synchronization with NTP](#) (page 415) to learn more about this topic.
3. The latest security updates are installed. If in doubt, run a YaST Online Update.
4. The default Web server port (port 80) is opened in the firewall. For this, configure the SUSEFirewall2 to allow the service *HTTP Server* in the external zone. This can be done using YaST. [Section “Configuring with YaST”](#) (page 106) gives details.

26.1.2 Installation

Apache on SUSE Linux is not installed by default. To install it, start YaST and select *Software* → *Software Management*. Now choose *Filters* → *Selections* and select *Simple Web Server with Apache2*. Confirm the installation of the dependent packages to finish the installation process.

Apache is installed with a standard, predefined configuration that runs “out of the box”. The installation includes the multiprocessing module `apache2-prefork` as well the PHP5 module. Refer to [Section 26.4, “Installing, Activating, and Configuring Modules”](#) (page 465) for more information about modules.

26.1.3 Start

To start Apache and make sure that it is automatically started during boot, start YaST and select *System* → *System Services (Runlevel)*. Search for `apache2` and *Enable* the service. The Web server starts immediately. By saving your changes with *Finish*, the system is configured to automatically start Apache in runlevels 3 and 5 during boot. For more information about the runlevels in SUSE Linux and a description of the YaST runlevel editor, refer to [Section 8.2.3, “Configuring System Services \(Runlevel\) with YaST”](#) (page 188).

To start Apache using the shell, run `rccache2 start`. To make sure that Apache is automatically started during boot in runlevels 3 and 5, use `chkconfig -a apache2`.

If you have not received error messages when starting Apache, the Web server should be running now. Start a browser and open <http://localhost/>. You should see an Apache test page starting with “If you can see this, it means that the installation of the Apache Web server software on this system was successful.” If you do not see this page, refer to [Section 26.8, “Troubleshooting”](#) (page 482).

Now that the Web server is running, you can add your own documents, adjust the configuration according to your needs, or add functionality by installing modules.

26.2 Configuring Apache

Apache in SUSE Linux can be configured in two different ways: with YaST or manually. Manual configuration offers a higher level of detail, but lacks the convenience of the YaST GUI.

IMPORTANT: Configuration Changes

Changes to most configuration values for Apache only take effect after Apache is restarted or reloaded. This happens automatically when using YaST and finishing the configuration with *Enabled* checked for the *HTTP Service*. Manual restart is described in [Section 26.3, “Starting and Stopping Apache”](#) (page 463). Most configuration changes only require a reload with `rcaapache2 reload`.

26.2.1 Configuring Apache Manually

Configuring Apache manually involves editing the plain text configuration files as the user `root`.

Configuration Files

Apache configuration files can be found in two different locations:

- `/etc/sysconfig/apache2`
- `/etc/apache2/`

/etc/sysconfig/apache2

/etc/sysconfig/apache2 controls some global settings of Apache, like modules to load, additional configuration files to include, flags with which the server should be started, and flags that should be added to the command line. Every configuration option in this file is extensively documented and therefore not mentioned here. For a general-purpose Web server, the settings in /etc/sysconfig/apache2 should be sufficient for any configuration needs.

IMPORTANT: No SuSEconfig Module for Apache

The SuSEconfig module for Apache has been removed from SUSE Linux. It is no longer necessary to run `SuSEconfig` after changing `/etc/sysconfig/apache2`.

/etc/apache2/

/etc/apache2/ hosts all configuration files for Apache. In the following, the purpose of each file is explained. Each file includes several configuration options (also referred to as *directives*). Every configuration option in these files is extensively documented and therefore not mentioned here.

The Apache configuration files are organized as follows:

```
/etc/apache2/
|
|- charset.conv
|- conf.d/
|   |
|   |- *.conf
|
|- default-server.conf
|- errors.conf
|- httpd.conf
|- listen.conf
|- magic
|- mime.types
|- mod_*.conf
|- server-tuning.conf
|- ssl-global.conf
|- ssl.*
|- sysconfig.d
|   |
|   |- global.conf
|   |- include.conf
```

```
|   |- loadmodule.conf . .
|
|- uid.conf
|- vhosts.d
|   |- *.conf
```

Apache Configuration Files in /etc/apache2/

`charset.conv`

Specifies which character sets to use for different languages. Do not edit.

`conf.d/*.conf`

Configuration files added by other modules. These configuration files can be included into your virtual host configuration where needed. See `vhosts.d/vhost.template` for examples. By doing so, you can provide different module sets for different virtual hosts.

`default-server.conf`

Global configuration for all virtual hosts with reasonable defaults. Instead of changing the values, overwrite them with a virtual host configuration.

`errors.conf`

Defines how Apache responds to errors. To customize these messages for all virtual hosts, edit this file. Otherwise overwrite these directives in your virtual host configurations.

`httpd.conf`

The main Apache server configuration file. Avoid changing this file. It mainly contains include statements and global settings. Overwrite global settings in the respective configuration files listed here. Change host-specific settings (such as document root) in your virtual host configuration.

`listen.conf`

Binds Apache to specific IP addresses and ports. Name-based virtual hosting (see [Section “Name-Based Virtual Hosts”](#) (page 453)) is also configured here.

`magic`

Data for the `mime_magic` module that helps Apache automatically determine the MIME type of an unknown file. Do not change.

`mime.types`

MIME types known by the system (this actually is a link to `/etc/mime.types`).

Do not edit. If you need to add MIME types not listed here, add them to `mod_mime-defaults.conf`.

`mod_*.conf`

Configuration files for the modules that are installed by default. Refer to [Section 26.4, “Installing, Activating, and Configuring Modules”](#) (page 465) for details.

Note that configuration files for optional modules reside in the directory `conf.d`.

`server-tuning.conf`

Contains configuration directives for the different MPMs (see [Section 26.4.4, “Multiprocessing Modules”](#) (page 469)) as well as general configuration options that control Apache's performance. Properly test your Web server when making changes here.

`ssl-global.conf` and `ssl.*`

Global SSL configuration and SSL certificate data. Refer to [Section 26.6, “Setting Up a Secure Web Server with SSL”](#) (page 475) for details.

`sysconfig.d/*.conf`

Configuration files automatically generated from `/etc/sysconfig/apache2`.

Do not change any of these files—edit `/etc/sysconfig/apache2` instead.

Put no other configuration files in this directory.

`uid.conf`

Specifies under which user and group ID Apache runs. Do not change.

`vhosts.d/*.conf`

Your virtual host configuration should go here. The directory contains template files for virtual hosts with and without SSL. Every file in this directory ending in `.conf` is automatically included in the Apache configuration. Refer to [Section “Virtual Host Configuration”](#) (page 452) for details.

Virtual Host Configuration

The term *virtual host* refers to Apache's ability to serve multiple URIs (universal resource identifiers) from the same physical machine. This means that several domains, such as `www.example.com` and `www.example.net`, are run by a single Web server on one physical machine.

It is common practice to use virtual hosts to save administrative effort (only a single Web server needs to be maintained) and hardware expenses (each domain does not require a dedicated server). Virtual hosts can be name based, IP based, or port based.

Virtual hosts can be configured via YaST (see [Section “Virtual Hosts”](#) (page 460)) or by manually editing a configuration file. By default, Apache in SUSE Linux is prepared for one configuration file per virtual host in `/etc/apache2/vhosts.d/`. All files in this directory with the extension `.conf` are automatically included to the configuration. A basic template for a virtual host is provided in this directory (`vhost.template` or `vhost-ssl.template` for a virtual host with SSL support).

TIP: Always Create a Virtual Host Configuration

It is recommended to always create a virtual host configuration file, even if your Web server only hosts one domain. In doing so, you not only have the domain-specific configuration in one file, but you can always fall back to a working basic configuration by simply moving, deleting, or renaming the configuration file for the virtual host. For the same reason, you should also create separate configuration files for each virtual host.

The `<VirtualHost></VirtualHost>` block holds the information that applies to a particular domain. When Apache receives a client request for a defined virtual host, it uses the directives enclosed in this section. Almost all directives can be used in a virtual host context. See <http://httpd.apache.org/docs/2.0/mod/quickreference.html> for further information about Apache's configuration directives.

Name-Based Virtual Hosts

With name-based virtual hosts, more than one Web site is served per IP address. Apache uses the host field in the HTTP header sent by the client to connect the request to a matching `ServerName` entry of one of the virtual host declarations. If no matching `ServerName` is found, the first specified virtual host is used as a default.

The directive `NameVirtualHost` tells Apache on which IP address and, optionally, which port to listen for requests by clients containing the domain name in the HTTP header. This option is configured in the configuration file `/etc/apache2/listen.conf`.

The first argument can be a fully qualified domain name, but it is recommended to use the IP address. The second argument is the port and is optional. By default, port 80 is used and is configured via the `Listen` directive.

The wild card `*` can be used for both the IP address and the port number to receive requests on all interfaces. IPv6 addresses must be enclosed in square brackets.

Example 26.1 Variations of Name-Based VirtualHost Entries

```
# NameVirtualHost IP-address[:Port]
NameVirtualHost 192.168.1.100:80
NameVirtualHost 192.168.1.100
NameVirtualHost *:80
NameVirtualHost *
NameVirtualHost [2002:c0a8:164::]:80
```

The opening `VirtualHost` tag takes the IP address (or fully qualified domain name) previously declared with the `NameVirtualHost` as an argument in a name-based virtual host configuration. A port number previously declared with the `NameVirtualHost` directive is optional.

The wild card `*` is also allowed as a substitute for the IP address. This syntax is only valid in combination with the wild card usage in `NameVirtualHost *`. When using IPv6 addresses, the address must be included in square brackets.

Example 26.2 Name-Based VirtualHost Directives

```
<VirtualHost 192.168.1.100:80>
  ...
</VirtualHost>

<VirtualHost 192.168.1.100>
  ...
</VirtualHost>

<VirtualHost *:80>
  ...
</VirtualHost>

<VirtualHost *>
  ...
</VirtualHost>

<VirtualHost [2002:c0a8:164::]>
  ...
</VirtualHost>
```

IP-Based Virtual Hosts

This alternative virtual host configuration requires the setup of multiple IPs for a machine. One instance of Apache hosts several domains, each of which is assigned a different IP.

The physical server must have one IP address for each IP-based virtual host. If the machine does not have multiple network cards, virtual network interfaces (IP aliasing) can also be used.

The following example shows Apache running on a machine with the IP 192.168.0.10, hosting two domains on the additional IPs 192.168.0.20 and 192.168.0.30. A separate `VirtualHost` block is needed for every virtual server.

Example 26.3 IP-Based VirtualHost Directives

```
<VirtualHost 192.168.0.20>
  ...
</VirtualHost>

<VirtualHost 192.168.0.30>
  ...
</VirtualHost>
```

Here, `VirtualHost` directives are only specified for interfaces other than 192.168.0.10. When a `Listen` directive is also configured for 192.168.0.10, a separate IP-based virtual host must be created to answer HTTP requests to that interface—otherwise the directives found in the default server configuration (`/etc/apache2/default-server.conf`) are applied.

Basic Virtual Host Configuration

At least the following directives should be present in each virtual host configuration in order to set up a virtual host. See `/etc/apache2/vhosts.d/vhost.template` for more options.

`ServerName`

The fully qualified domain name under which the host should be addressed.

DocumentRoot

Path to the directory from which Apache should serve files for this host. For security reasons, access to the entire file system is forbidden by default, so you must explicitly unlock this directory within a `Directory` container.

ServerAdmin

E-mail address of the server administrator. This address is, for example, shown on error pages Apache creates.

ErrorLog

The error log file for this virtual host. Although it is not necessary to create separate error log files for each virtual host, it is common practice to do so, because it makes debugging of errors much easier. `/var/log/apache2/` is the default directory where Apache's log files should be kept.

CustomLog

The access log file for this virtual host. Although it is not necessary to create separate access log files for each virtual host, it is common practice to do so, because it allows separate analysis of access statistics for each host. `/var/log/apache2/` is the default directory where Apache's log files should be kept.

As mentioned above, access to the whole file system is forbidden by default for security reasons. Therefore, explicitly unlock the `DocumentRoot` directory in which you have placed the files Apache should serve:

```
<Directory "/srv/www/example.com_htdocs">
    Order allow,deny
    Allow from all
</Directory>
```

The complete configuration file looks like this:

Example 26.4 Basic VirtualHost Configuration

```
<VirtualHost 192.168.0.10>
    ServerName www.example.com
    DocumentRoot /srv/www/example.com_htdocs
    ServerAdmin webmaster@example.com
    ErrorLog /var/log/apache2/www.example.com_log
    CustomLog /var/log/apache2/www.example.com-access_log common
    <Directory "/srv/www/example.com">
        Order allow,deny
        Allow from all
    </Directory>
</VirtualHost>
```

26.2.2 Configuring Apache with YaST

To configure your Web server with YaST, start YaST and select *Network Services → HTTP Server*. When starting the module for the first time, the *HTTP Server Wizard* starts, prompting you to make just a few basic decisions concerning administration of the server. After having finished the wizard, the dialog in [Section “HTTP Server Configuration”](#) (page 461) starts every time you call the *HTTP Server* module.

HTTP Server Wizard

The *HTTP Server Wizard* consists of five steps. In the last step of the dialog, you are given the opportunity to enter the expert configuration mode to make even more specific settings.

Network Device Selection

Here, specify the network interfaces and ports Apache uses to listen for incoming requests. You can select any combination of existing network interfaces and their respective IP addresses. Ports from all three ranges (well-known ports, registered ports, and dynamic or private ports) that are not reserved by other services can be used. The default setting is to listen on all network interfaces (IP addresses) on port 80.

Check *Open Firewall for Selected Ports* to open the ports in the firewall that the Web server listens on. This is necessary to make the Web server available on the network, which can be a LAN, WAN, or the public Internet. Keeping the port closed is only useful in test situations where no external access to the Web server is necessary.

Click *Next* to continue with configuration.

Modules

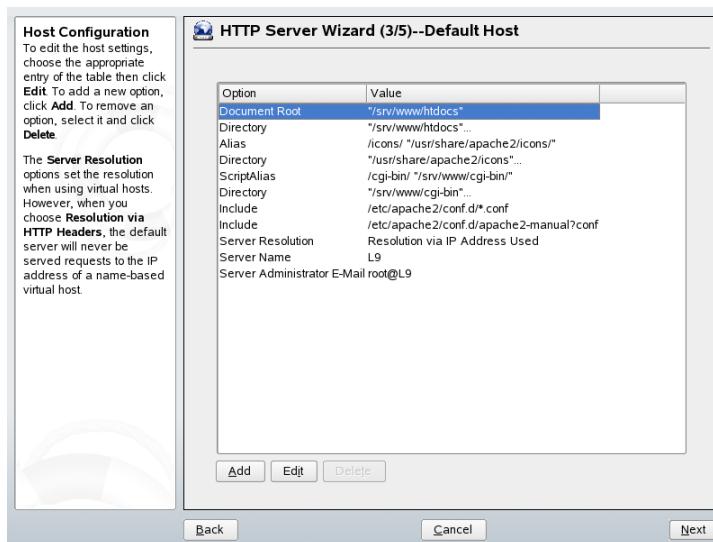
The *Modules* configuration option allows for the activation or deactivation of the script languages, the web server should support. For the activation or deactivation of other modules, refer to [Section “Server Modules”](#) (page 462). Click *Next* to advance to the next dialog.

Default Host

This option pertains to the default Web server. As explained in [Section “Virtual Host Configuration”](#) (page 452), Apache can serve multiple virtual hosts from a single physical machine. The first declared virtual host in the configuration file is commonly referred to as the *default host*. Each virtual host inherits the default host's configuration.

To edit the host settings (also called *directives*), choose the appropriate entry in the table then click *Edit*. To add new directives, click *Add*. To delete a directive, select it and click *Delete*.

Figure 26.1 HTTP Server Wizard: Default Host



Here is list of the default settings of the server:

Document Root

Path to the directory from which Apache serves files for this host. `/srv/www/htdocs` is the default location.

Alias

With the help of `Alias` directives, URLs can be mapped to physical file system locations. This means that a certain path even outside the `Document Root` in the file system can be accessed via a URL aliasing that path.

The default SUSE Linux Alias `/icons` points to `/usr/share/apache2/icons` for the Apache icons displayed in the directory index view.

ScriptAlias

Similar to the Alias directive, the ScriptAlias directive maps a URL to a file system location. The difference is that ScriptAlias designates the target directory as a CGI location, meaning that CGI scripts should be executed in that location.

Directory

With the Directory setting, you can enclose a group of configuration options that will only apply to the specified directory.

Access and display options for the directories `/usr/share/apache2/icons` and `/srv/www/cgi-bin` are configured here. It should not be necessary to change the defaults.

Include

With include, additional configuration files can be specified. `/etc/apache2/conf.d/` is the directory containing the configuration files that come with external modules. By default, all files in this directory (`*.conf`) are included. `/etc/apache2/conf.d/apache2-manual?conf` is the directory containing all apache2-manual configuration files.

Server Name

This specifies the default URL used by clients to contact the Web server. Use a fully qualified domain name (FQDN) to reach the Web server at `http://FQDN/` or its IP address. You cannot choose an arbitrary name here—the server must be “known” under this name.

Server Administrator E-Mail

E-mail address of the server administrator. This address is, for example, shown on error pages Apache creates.

Server Resolution

This option refers to [Section “Virtual Host Configuration”](#) (page 452). *Determine Request Server by HTTP Headers* lets a VirtualHost answer on a request to its server name (see [Section “Name-Based Virtual Hosts”](#) (page 453)). *Determine Request Server by Server IP Address* makes Apache select the requested host by

the HTTP header information the client sends. See [Section “IP-Based Virtual Hosts”](#) (page 455) for more details on IP-based virtual hosts.

After finishing with the *Default Host* step, click *Next* to continue with the configuration.

Virtual Hosts

In this step, the wizard displays a list of already configured virtual hosts (see [Section “Virtual Host Configuration”](#) (page 452)). If you have not made manual changes prior to starting the YaST HTTP wizard, only one virtual host is present—one identical to the default host configured in the previous step. It is marked as default with an asterisk next to the server name.

To add a host, click *Add* to open a dialog in which to enter basic information about the host. *Server Identification* includes the server name, server contents root (`DocumentRoot`), and administrator e-mail. *Server Resolution* is used to determine how a host is identified (name based or IP based). These options are explained in [Section “Default Host”](#) (page 458).

Clicking *Next* advances to the second part of the virtual host configuration dialog.

In part two of the virtual host configuration you can specify whether to enable CGI scripts and which directory to use for these scripts. It is also possible to enable SSL. If you do so, you must specify the path to the certificate as well. See [Section 26.6.2, “Configuring Apache with SSL”](#) (page 479) for details on SSL and certificates. With the *Directory Index* option, you can specify which file to display when the client requests a directory (by default, `index.html`). Add one or more filenames (space-separated) if you want to change this. With *Enable Public HTML*, the content of the users public directories (`~user/public_html/`) is made available on the server under `http://www.example.com/~user`.

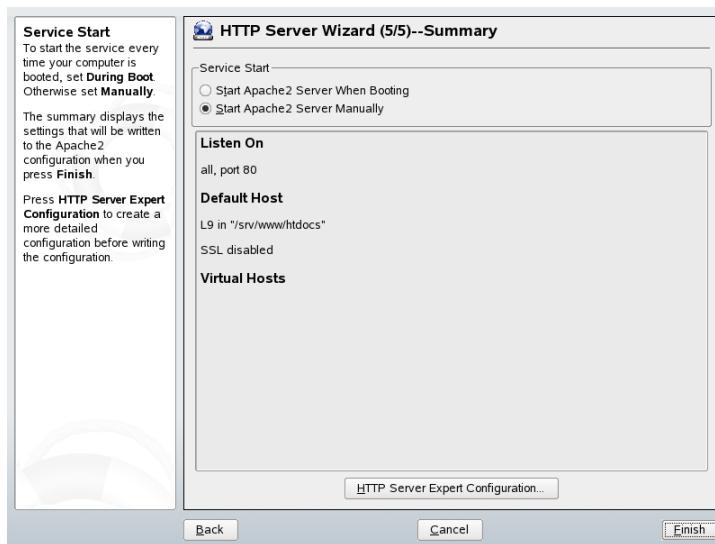
IMPORTANT: Creating Virtual Hosts

It is not possible to add virtual hosts at will. If using name-based virtual hosts, each hostname must be resolved on the network. If using IP-based virtual hosts, you can assign only one host to each IP address available.

Summary

This is the final step of the wizard. Here, determine how and when the Apache server is started: when booting or manually. Also see a short summary of the configuration made so far. If you are satisfied with your settings, click *Finish* to complete configuration. If you want to change something, click *Back* until you have reached the desired dialog. Clicking *HTTP Server Expert Configuration* opens the dialog described in Section “[HTTP Server Configuration](#)” (page 461).

Figure 26.2 *HTTP Server Wizard: Summary*



HTTP Server Configuration

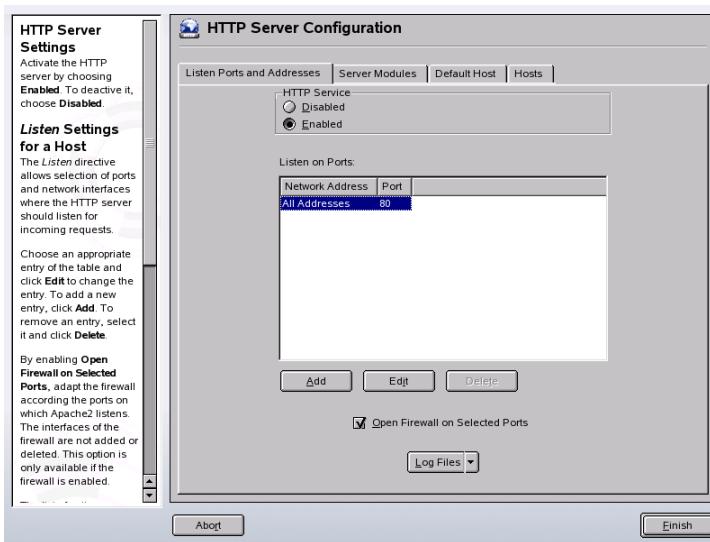
The *HTTP Server Configuration* dialog also lets you make even more adjustments to the configuration than the wizard (which only runs if you configure your Web server for the first time). It consists of four tabs described in the following. No configuration option you change here is effective immediately—you always must confirm your changes with *Finish* to make them effective. Clicking *Cancel* leaves the configuration module and discards your changes.

Listen Ports and Addresses

In *HTTP Service*, select whether Apache should be running (*Enabled*) or stopped (*Disabled*). In *Listen on Ports*, *Add*, *Edit*, or *Delete* addresses and ports on which the server should be available. The default is to listen on all interfaces on port 80. You should always check *Open Firewall on Selected Ports*, because otherwise the Web server is not reachable from the outside. Keeping the port closed is only useful in test situations where no external access to the Web server is necessary.

With *Log Files*, watch either the access log or the error log. This is useful if you want to test your configuration. The log file opens in a separate window from which you can also restart or reload the Web server (see [Section 26.3, “Starting and Stopping Apache”](#) (page 463) for details). These commands are effective immediately.

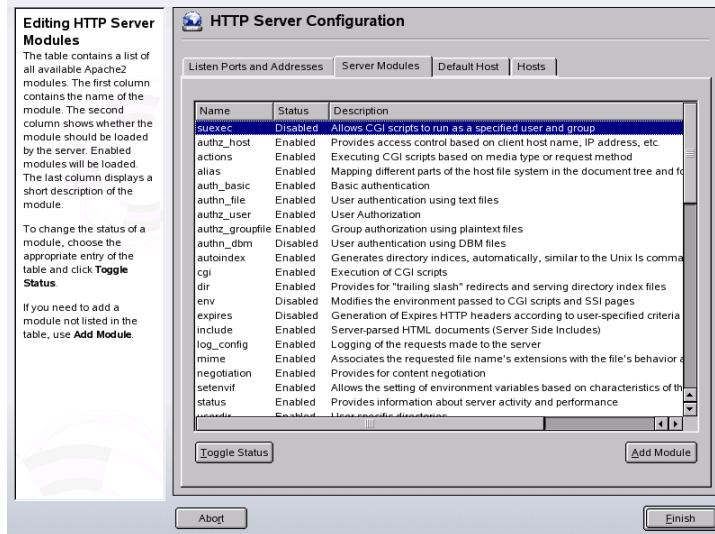
Figure 26.3 HTTP Server Configuration: Listen Ports and Addresses



Server Modules

You can change the status (enabled or disabled) of Apache2 modules by clicking *Toggle Status*. Click *Add Module* to add a new module that is already installed but not yet listed. Learn more about modules in [Section 26.4, “Installing, Activating, and Configuring Modules”](#) (page 465).

Figure 26.4 HTTP Server Configuration: Server Modules



Default Host or Hosts

These dialogs are identical to the ones already described. Refer to [Section “Default Host”](#) (page 458) and [Section “Virtual Hosts”](#) (page 460).

26.3 Starting and Stopping Apache

If configured with YaST (see [Section 26.2.2, “Configuring Apache with YaST”](#) (page 457)), Apache is started at boot time in runlevels 3 and 5 and stopped in runlevels 0, 1, 2, and 6. You can change this behavior using YaST’s runlevel editor or the command line tool `chkconfig`.

To start, stop, or manipulate Apache on a running system, use the init script `/usr/sbin/rapache2` (refer to [Section 8.2.2, “Init Scripts”](#) (page 184) for a general information about init scripts.). The `rapache2` command takes the following parameters:

```
start
Starts Apache if it is not already running.
```

`startssl`

Starts Apache with SSL support if it is not already running. For more information about SSL support, refer to [Section 26.6, “Setting Up a Secure Web Server with SSL”](#) (page 475).

`restart`

Stops then restarts Apache. Starts the Web server if it was not running before.

`try-restart`

Stops then restarts Apache only if it has been running before.

`reload` or `graceful`

Stops the Web server by advising all forked Apache processes to first finish their requests before shutting down. As each process dies, it is replaced by a newly started one, resulting in complete “restart” of Apache.

TIP

`rcaapache2 reload` is the preferred method of restarting Apache in production environments, for example, to activate a change in the configuration, because it allows all clients to be served without causing connection break-offs.

`configtest`

Checks the syntax of the configuration files without affecting a running Web server. Because this check is forced every time the server is started, reloaded, or restarted, it is usually not necessary to run the test explicitly (if a configuration error is found, the Web server is not started, reloaded, or restarted).

`probe`

Probes for the necessity of a reload (checks whether the configuration has changed) and suggests the required arguments for the `rcaapache2` command.

`server-status` and `full-server-status`

Dumps a short or full status screen, respectively. Requires either `lynx` or `w3m` installed as well as the module `mod_status` enabled. In addition to that, `status` must be added to `APACHE_SERVER_FLAGS` in the file `/etc/sysconfig/apache2`.

TIP: Additional Flags

If you specify additional flags to the `rcaapache2`, these are passed through to the Web server.

26.4 Installing, Activating, and Configuring Modules

The Apache software is built in a modular fashion: all functionality except some core tasks is handled by modules. This has progressed so far that even HTTP is processed by a module (`http_core`).

Apache modules can be compiled into the Apache binary at build time or dynamically loaded at runtime. Refer to [Section 26.4.2, “Activation and Deactivation”](#) (page 466) for details of how to load modules dynamically.

Apache modules can be divided into four different categories:

Base Modules

Base modules are compiled into Apache by default. Apache in SUSE Linux has only `mod_so` (needed to load other modules) and `http_core` compiled in. All others are available as shared objects: rather than being included in the server binary itself, they can be included at runtime.

Extension Modules

In general, modules labeled as extensions are included in the Apache software package, but are usually not compiled into the server statically. In SUSE Linux, they are available as shared objects that can be loaded into Apache at runtime.

External Modules

Modules labeled external are not included in the official Apache distribution. SUSE Linux provides several of them readily available for use.

Multiprocessing Modules

MPMs are responsible for accepting and handling requests to the Web server, representing the core of the Web server software.

26.4.1 Module Installation

If you have followed the default way of installing Apache (described in [Section 26.1.2, “Installation” \(page 448\)](#)), it is installed with all base and extension modules, the multi-processing module Prefork MPM, and the external module PHP5.

You can install additional external modules by starting YaST and choosing *Software → Software Management*. Now choose *Filter → Search* and search for *apache*. Among other packages, the result list contains all available external Apache modules.

26.4.2 Activation and Deactivation

Using YaST, you can activate or deactivate the script language modules (PHP5, Perl, Python, and Ruby) with the module configuration described in [Section “HTTP Server Wizard” \(page 457\)](#). All other modules can be enabled or disabled as described in [Section “Server Modules” \(page 462\)](#).

If you prefer to activate or deactivate the modules manually, use the commands `a2enmod mod_foo` or `a2dismod mod_foo`, respectively. `a2enmod -l` outputs a list of all currently active modules.

IMPORTANT: Including Configuration Files for External Modules

If you have activated external modules manually, make sure to load their configuration files in all virtual host configurations. Configuration files for external modules are located under `/etc/apache2/conf.d/` and are not loaded by default. If you need the same modules on each virtual host, you can include `*.conf` from this directory. Otherwise include individual files. See `/etc/apache2/vhost.d/vhost.template` for examples.

26.4.3 Base and Extension Modules

All base and extension modules are described in detail in the Apache documentation. Only a brief description of the most important modules is available here. Refer to <http://httpd.apache.org/docs/2.2/mod/> to learn details about each module.

`mod_alias`

Provides `Alias` and `Redirect` directives with which you can map a URL to a specific directory (`Alias`) or redirect a requested URL to another location. This module is enabled by default.

`mod_auth*`

The authentication modules provide different authentication methods: basic authentication with `mod_auth_basic` or digest authentication with `mod_auth_digest`. Digest authentication in Apache 2.2 is considered experimental.

`mod_auth_basic` and `mod_auth_digest` must be combined with an authentication provider module, `mod_authn_*` (for example, `mod_authn_file` for text file-based authentication) and with an authorization module `mod_authz_*` (for example, `mod_authz_user` for user authorization).

More information about this topic is available in the “Authentication howto” at <http://httpd.apache.org/docs/2.2/howto/auth.html>

`mod_autoindex`

`Autoindex` generates directory listings when no index file (for example, `index.html`) is present. The look and feel of these indexes is configurable. This module is enabled by default. However, directory listings are disabled by default via the `Options` directive—overwrite this setting in your virtual host configuration. The default configuration file for this module is located at `/etc/apache2/mod_autoindex-defaults.conf`.

`mod_cgi`

`mod_cgi` is needed to execute CGI scripts. This module is enabled by default.

`mod_deflate`

Using this module, Apache can be configured to compress given file types on the fly before delivering them.

`mod_dir`

`mod_dir` provides the `DirectoryIndex` directive with which you can configure which files are automatically delivered when a directory is requested (`index.html` by default). It also provides an automatic redirect to the correct URL when a directory request does not contain a trailing slash. This module is enabled by default.

`mod_expires`

With `mod_expires`, you can control how often proxy and browser caches refresh your documents by sending an `Expires` header. This module is enabled by default.

`mod_include`

`mod_include` lets you use Server Side Includes (SSI), which provide a basic functionality to generate HTML pages dynamically. This module is enabled by default.

`mod_info`

Provides a comprehensive overview of the server configuration under `http://localhost/server-info/`. For security reasons, you should always limit access to this URL. By default only `localhost` is allowed to access this URL. `mod_info` is configured at `/etc/apache2/mod_info.conf`

`mod_log_config`

With this module, you can configure the looks of the Apache log files. This module is enabled by default.

`mod_mime`

The `mime` module takes care that a file is delivered with the correct MIME header based on the filename's extension (for example `text/html` for HTML documents). This module is enabled by default.

`mod_negotiation`

Necessary for content negotiation. See <http://httpd.apache.org/docs/2.2/content-negotiation.html> for more information. This module is enabled by default.

`mod_rewrite`

Provides the functionality of `mod_alias`, but offers more features and flexibility. With `mod_rewrite`, you can redirect URLs based on multiple rules, request headers, and more.

`mod_speling`

`mod_speling` attempts to automatically correct typographical errors in URLs, such as capitalization errors.

`mod_ssl`

Enables encrypted connections between Web server and clients. See [Section 26.6, “Setting Up a Secure Web Server with SSL”](#) (page 475) for details. This module is enabled by default.

`mod_status`

Provides information on server activity and performance under `http://localhost/server-status/`. For security reasons, you should always limit access to this URL. By default, only `localhost` is allowed to access this URL. `mod_status` is configured at `/etc/apache2/mod_status.conf`

`mod_suexec`

`mod_suexec` lets you run CGI scripts under a different user and group. This module is enabled by default.

`mod_userdir`

Enables user-specific directories available under `~user/`. The `UserDir` directive must be specified in the configuration. This module is enabled by default.

26.4.4 Multiprocessing Modules

SUSE Linux provides two different multiprocessing modules (MPMs) for use with Apache.

Prefork MPM

The prefork MPM implements a nonthreaded, preforking Web server. It makes the Web server behave similarly to Apache version 1.x in that it isolates each request and handles it by forking a separate child process. Thus problematic requests cannot affect others, avoiding a lockup of the Web server.

While providing stability with this process-based approach, the prefork MPM consumes more system resources than its counterpart, the worker MPM. The prefork MPM is considered the default MPM for Unix-based operating systems.

IMPORTANT: MPMs in This Document

This document assumes Apache is used with the prefork MPM.

Worker MPM

The worker MPM provides a multithreaded Web server. A thread is a “lighter” form of a process. The advantage of a thread over a process is its lower resource consumption.

Instead of only forking child processes, the worker MPM serves requests by using threads with server processes. The preforked child processes are multithreaded. This approach makes Apache perform better by consuming fewer system resources than the prefork MPM.

One major disadvantage is the stability of the worker MPM: if a thread becomes corrupt, all threads of a process can be affected. In the worst case, this may result in a server crash. Especially when using the Common Gateway Interface (CGI) with Apache under heavy load, internal server errors might occur due to threads unable to communicate with system resources. Another argument against using the worker MPM with Apache is that not all available Apache modules are thread-safe and thus cannot be used in conjunction with the worker MPM.

WARNING: Using PHP Modules with MPMs

Not all available PHP modules are thread-safe. Using the worker MPM with mod_php is strongly discouraged.

26.4.5 External Modules

Find a list of all external modules shipped with SUSE Linux here. Find the module's documentation in the listed directory.

FastCGI

FastCGI is a language-independent, scalable, and open extension to CGI that provides high performance without the limitations of server-specific APIs. FastCGI applications are very fast because they are persistent. There is no per-request start-up and initialization overhead.

Package Name: apache2-mod_fastcgi

Configuration File: /etc/apache2/conf.d/mod_fastcgi.conf

More Information: /usr/share/doc/packages/apache2-mod_fastcgi

mod_perl

mod_perl enables you to run Perl scripts in an embedded interpreter. The persistent interpreter embedded in the server avoids the overhead of starting an external interpreter and the penalty of Perl start-up time.

Package Name: apache2-mod_perl

Configuration File: /etc/apache2/conf.d/mod_perl.conf
More Information: /usr/share/doc/packages/apache2-mod_perl

mod_php5

PHP is a server-side, cross-platform HTML embedded scripting language.

Package Name: apache2-mod_php5

Configuration File: /etc/apache2/conf.d/php5.conf

More Information: /usr/share/doc/packages/apache2-mod_php5

mod_python

mod_python allows embedding Python within the Apache HTTP server for a considerable boost in performance and added flexibility in designing Web-based applications.

Package Name: apache2-mod_python

More Information: /usr/share/doc/packages/apache2-mod_python

mod_ruby

mod_ruby embeds the Ruby interpreter into the Apache Web server, allowing Ruby CGI scripts to be executed natively. These scripts start much faster than without mod_ruby.

Package Name: apache2-mod_ruby

More Information: /usr/share/doc/packages/apache2-mod_ruby

mod_jk-ap20

This module provides connectors between Apache and a Tomcat Servlet Container.

Package Name: mod_jk-ap20

More Information: /usr/share/doc/packages/mod_jk-ap20

26.4.6 Compilation

Apache can be extended by advanced users by writing custom modules. To develop modules for Apache or compile third-party modules, the package apache2-devel is required along with the corresponding development tools. apache2-devel also contains the apxs2 tools, which are necessary for compiling additional modules for Apache.

`apxs2` enables the compilation and installation of modules from source code (including the required changes to the configuration files), which creates *dynamic shared objects* (DSOs) that can be loaded into Apache at runtime.

The `apxs2` binaries are located under `/usr/sbin`:

- `/usr/sbin/apxs2`—suitable for building an extension module that works with any MPM. The installation location is `/usr/lib/apache2`.
- `/usr/sbin/apxs2-prefork`—suitable for prefork MPM modules. The installation location is `/usr/lib/apache2-prefork`.
- `/usr/sbin/apxs2-worker`—suitable for worker MPM modules.

`apxs2` installs modules so they can be used for all MPMs. The other two programs install modules so they can only be used for the respective MPMs. `apxs2` installs modules in `/usr/lib/apache2`, `apxs2-prefork` and `apxs2-worker` installs modules in `/usr/lib/apache2-prefork` or `/usr/lib/apache2-worker`.

Install and activate a module from source code with the commands `cd /path/to/module/source; apxs2 -cia mod_foo.c` (`-c` compiles the module, `-i` installs it, and `-a` activates it). Other options of `apxs2` are described in the `apxs2(1)` man page.

26.5 Getting CGI Scripts to Work

Apache's Common Gateway Interface (CGI) lets you create dynamic content with programs or scripts usually referred to as CGI scripts. CGI scripts can be written in any programming language. Usually, script languages such as Perl or PHP are used.

To enable Apache to deliver content created by CGI scripts, `mod_cgi` needs to be activated. `mod_alias` is also needed. Both modules are enabled by default. Refer to [Section 26.4.2, “Activation and Deactivation”](#) (page 466) for details on activating modules.

WARNING: CGI Security

Allowing the server to execute CGI scripts is a potential security hole. Refer to [Section 26.7, “Avoiding Security Problems”](#) (page 480) for additional information.

26.5.1 Apache Configuration

In SUSE Linux, the execution of CGI scripts is only allowed in the directory `/srv/www/cgi-bin/`. This location is already configured to execute CGI scripts. If you have created a virtual host configuration (see [Section “Virtual Host Configuration”](#) (page 452)) and want to place your scripts in a host-specific directory, you must unlock and configure this directory.

Example 26.5 VirtualHost CGI Configuration

```
ScriptAlias /cgi-bin/ "/srv/www/example.com cgi-bin/"❶
<Directory "/srv/www/example.com cgi-bin/">
    Options +ExecCGI❷
    AddHandler cgi-script .cgi .pl❸
    Order allow,deny❹
    Allow from all
</Directory>
```

- ❶ Tells Apache to handle all files within this directory as CGI scripts.
- ❷ Enables CGI script execution
- ❸ Tells the server to treat files with the extensions `.pl` and `.cgi` as CGI scripts. Adjust according to your needs.
- ❹ The `Order` and `Allow` directives control the default access state and the order in which `Allow` and `Deny` directives are evaluated. In this case “`deny`” statements are evaluated before “`allow`” statements and access from everywhere is enabled.

26.5.2 Running an Example Script

CGI programming differs from "regular" programming in that the CGI programs and scripts must be preceded by a MIME-Type header such as `Content-type: text/html`. This header is sent to the client, so it understands what kind of content it receives. Secondly, the script's output must be something the client, usually a Web browser, understands—HTML in most cases or plain text or images, for example.

A simple test script available under `/usr/share/doc/packages/apache2/test-cgi` is part of the Apache package. It outputs the content of some environment variables as plain text. Copy this script to either `/srv/www/cgi-bin/` or the script directory of your virtual host (`/srv/www/example.com/cgi-bin/`) and name it `test.cgi`.

Files accessible by the Web server should be owned by the user `root` (see [Section 26.7, “Avoiding Security Problems”](#) (page 480) for additional information). Because the Web server runs with a different user, the CGI scripts must be world-executable and world-readable. Change into the CGI directory and use the command `chmod 755 test.cgi` to apply the proper permissions.

Now call `http://localhost/cgi-bin/test.cgi` or `http://example.com/cgi-bin/test.cgi`. You should see the “CGI/1.0 test script report”.

26.5.3 Troubleshooting

If you do not see the output of the test program but an error message instead, check the following:

CGI Troubleshooting

- Have you reloaded the server after having changed the configuration? Check with `rcaapache2 probe`.
- If you have configured your custom CGI directory, is it configured properly? If in doubt, try the script within the default CGI directory `/srv/www/cgi-bin/` and call it with `http://localhost/cgi-bin/test.cgi`.
- Are the file permissions correct? Change into the CGI directory and execute the `ls -l test.cgi`. Its output should start with
`-rwxr-xr-x 1 root root`
- Make sure that the script does not contain programming errors. If you have not changed `test.cgi`, this should not be the case, but if you are using your own programs, always make sure that they do not contain programming errors.

26.6 Setting Up a Secure Web Server with SSL

Whenever sensitive data, such as credit card information, is transferred between Web server and client, it would be desirable to have a secure, encrypted connection with authentication. mod_ssl provides strong encryption using the secure sockets layer (SSL) and transport layer security (TLS) protocols for HTTP communication between a client and the Web server. Using SSL/TSL, a private connection between Web server and client is established. Data integrity is ensured and client and server are able to authenticate each other.

For this purpose, the server sends an SSL certificate that holds information proving the server's valid identity before any request to a URL is answered. In turn, this guarantees that the server is the uniquely correct end point for the communication. Additionally, the certificate generates an encrypted connection between client and server that can transport information without the risk of exposing sensitive, plain-text content.

mod_ssl does not implement the SSL/TSL protocols itself, but acts as an interface between Apache and an SSL library. In SUSE Linux, the OpenSSL library is used. OpenSSL is automatically installed with Apache.

The most visible effect of using mod_ssl with Apache is that URLs are prefixed with `https://` instead of `http://`.

26.6.1 Creating an SSL Certificate

In order to use SSL/TSL with the Web server, you need to create an SSL certificate. This certificate is needed for the authorization between Web server and client, so that each party can clearly identify the other party. To ensure the integrity of the certificate, it must be signed by a party every user trusts.

There are three types of certificates you can create: a “dummy” certificate for testing purposes only, a self-signed certificate for a defined circle of users that trust you, and a certificate signed by an independent, publicly-known certificate authority (CA).

Creating a certificate is basically a two step process. First, a private key for the certificate authority is generated then the server certificate is signed with this key.

TIP: For More Information

To learn more about concepts and definitions of SSL/TSL, refer to
http://httpd.apache.org/docs/2.2/ssl/ssl_intro.html.

Creating a “Dummy” Certificate

Generating a dummy certificate is simple. Just call the script
/usr/bin/gensslcert. It creates or overwrites the following files:

- /etc/apache2/ssl.crt/ca.crt
- /etc/apache2/ssl.crt/server.crt
- /etc/apache2/ssl.key/server.key
- /etc/apache2/ssl.csr/server.csr

A copy of ca.crt is also placed at /srv/www/htdocs/CA.crt for download.

IMPORTANT

A dummy certificate should never be used on a production system. Only use it for testing purposes.

Creating a Self-Signed Certificate

If you are setting up a secure Web server for an Intranet or for a defined circle of users, it might be sufficient if you sign a certificate with your own certificate authority (CA).

Creating a self-signed certificate is an interactive nine-step process. Change into the directory /usr/share/doc/packages/apache2 and run the following command:
./mkcert.sh make --no-print-directory /usr/bin/openssl
/usr/sbin/ custom. Do not attempt to run this command from outside this directory. The program provides a series of prompts, some of which require user input.

Procedure 26.1 *Creating a Self-Signed Certificate with mkcert.sh*

- 1 Decide the signature algorithm used for certificates

Choose RSA (, the default), because some older browsers have problems with DSA.

2 Generating RSA private key for CA (1024 bit)

No interaction needed.

3 Generating X.509 certificate signing request for CA

Create the CA's distinguished name here. This requires you to answer a few questions, such as country name or organization name. Enter valid data, because everything you enter here later shows up in the certificate. You do not need to answer every question. If one does not apply to you or you want to leave it blank, use “.”. Common name is the name of the CA itself—choose a significant name, such as *My company* CA.

4 Generating X.509 certificate for CA signed by itself

Choose certificate version (the default).

5 Generating RSA private key for SERVER (1024 bit)

No interaction needed.

6 Generating X.509 certificate signing request for SERVER

Create the distinguished name for the server key here. Questions are almost identical to the ones already answered for the CA's distinguished name. The data entered here applies to the Web server and does not necessarily need to be identical to the CA's data (for example, if the server is located elsewhere).

IMPORTANT: Selecting a Common Name

The common name you enter here must be the fully qualified hostname of your secure server (for example, www.example.com). Otherwise the browser issues a warning that the certificate does not match the server when accessing the Web server.

7 Generating X.509 certificate signed by own CA

Choose certificate version (the default).

- 8** Encrypting RSA private key of CA with a pass phrase for security

It is strongly recommended to encrypt the private key of the CA with a password, so choose **[Y]** and enter a password.

- 9** Encrypting RSA private key of SERVER with a pass phrase for security

Encrypting the server key with a password requires you to enter this password every time you start the Web server. This makes it difficult to automatically start the server on boot or to restart the Web server. Therefore, it is common sense to say **[N]** to this question. Keep in mind that your key is unprotected when not encrypted with a password and make sure that only authorized persons have access to the key.

IMPORTANT: Encrypting the Server Key

If you choose to encrypt the server key with a password, increase the value for APACHE_TIMEOUT in /etc/sysconfig/apache2. Otherwise you do not have enough time to enter the passphrase before the attempt to start the server is stopped unsuccessfully.

The script's result page presents a list of certificates and keys it has generated. Contrary to what the script outputs, the files have not been generated in the local directory `conf`, but to the correct locations under `/etc/apache2/`.

The last step is to copy the CA certificate file from `/etc/apache2/ssl.crt/ca.crt` to a location where your users can access it in order to incorporate it into the list of known and trusted CAs in their Web browsers. Otherwise a browser complains that the certificate was issued by an unknown authority. The certificate is valid for one year.

IMPORTANT: Self-Signed Certificates

Only use a self-signed certificate on a Web server that is accessed by people who know and trust you as a certificate authority. It is not recommended to use such a certificate on a public shop, for example.

Getting an Officially Signed Certificate

There are a number of official certificate authorities that sign your certificates. The certificate is signed by a trustworthy third party, so can be fully trusted. Publicly operating secure Web servers usually have got an officially signed certificate.

The best-known official CAs are Thawte (<http://www.thawte.com/>) or Verisign (www.verisign.com). These and other CAs are already compiled into all browsers, so certificates signed by these certificate authorities are automatically accepted by the browser.

When requesting an officially signed certificate, you do not send a certificate to the CA. Instead, issue a Certificate Signing Request (CSR). To create a CSR, call the script `/usr/share/ssl/misc/CA.sh -newreq`.

First the script asks for a password with which the CSR should be encrypted. Then you are asked to enter a distinguished name. This requires you to answer a few questions, such as country name or organization name. Enter valid data—everything you enter here later shows up in the certificate and is checked. You do not need to answer every question. If one does not apply to you or you want to leave it blank, use “.”. Common name is the name of the CA itself—choose a significant name, such as *My company* CA. Last, a challenge password and an alternative company name must be entered.

Find the CSR in the directory from which you called the script. The file is named `newreq.pem`.

26.6.2 Configuring Apache with SSL

The default port for SSL and TLS requests on the Web server side is 443. There is no conflict between a “regular” Apache listening on port 80 and an SSL/TLS-enabled Apache listening on port 443. In fact, HTTP and HTTPS can be run with the same Apache instance. Usually separate virtual hosts are used to dispatch requests to port 80 and port 443 to separate virtual servers.

IMPORTANT: Firewall Configuration

Do not forget to open the firewall for SSL-enabled Apache on port 443. This can be done with YaST as described in [Section “Configuring with YaST”](#) (page 106).

To use SSL, it must be activated in the global server configuration. Open `/etc/sysconfig/apache2` in an editor and search for `APACHE_MODULES`. Add “ssl” to the list of modules if it is not already present (`mod_ssl` is activated by default). Next, search for `APACHE_SERVER_FLAGS` and add “SSL”. If you have chosen to encrypt your server certificate with a password, you should also increase the value for `APACHE_TIMEOUT`, so you have enough time to enter the passphrase when Apache starts. Restart the server to make these changes active. A reload is not sufficient.

The virtual host configuration directory contains a template `/etc/apache2/vhosts.d/vhost-ssl.template` with SSL-specific directives that are extensively documented. Refer to [Section “Virtual Host Configuration” \(page 452\)](#) for the general virtual host configuration.

To get started, it should be sufficient to adjust the values for the following directives:

- `DocumentRoot`
- `ServerName`
- `ServerAdmin`
- `ErrorLog`
- `TransferLog`

IMPORTANT: Name-Based Virtual Hosts and SSL

It is not possible to run multiple SSL-enabled virtual hosts on a server with only one IP address. Users connecting to such a setup receive a warning message stating that the certificate does not match the server name every time they visit the URL. A separate IP address or port is necessary for every SSL-enabled domain to achieve communication based on a valid SSL certificate.

26.7 Avoiding Security Problems

A Web server exposed to the public Internet requires an ongoing administrative effort. It is inevitable that security issues appear, both related to the software and to accidental misconfiguration. Here are some tips for how to deal with them.

26.7.1 Up-to-Date Software

If there are vulnerabilities found in the Apache software, a security advisory will be issued by SUSE. It contains instructions for fixing the vulnerabilities, which in turn should be applied soon as possible. The SUSE security announcements are available from the following locations:

- **Web Page** http://www.suse.com/us/private/support/online_help/mailinlists/
- **Mailing List** http://www.suse.com/us/private/support/online_help/mailinlists/
- **RSS Feed** http://www.novell.com/linux/security/suse_security.xml

26.7.2 DocumentRoot Permissions

By default in SUSE Linux, the `DocumentRoot` directory `/srv/www/htdocs` and the `CGI` directory `/srv/www/cgi-bin` belong to the user and group `root`. You should not change these permissions. If the directories were writable for all, any user could place files into them. These files might then be executed by Apache with the permissions of `wwwrun`, which may give the user unintended access to file system resources. Use subdirectories of `/srv/www` to place the `DocumentRoot` and `CGI` directories for your virtual hosts and make sure that directories and files belong to user and group `root`.

26.7.3 File System Access

By default, access to the whole file system is denied in `/etc/apache2/httpd.conf`. You should never overwrite these directives, but specifically enable access to all directories Apache should be able to read (see [Section “Basic Virtual Host Configuration”](#) (page 455) for details). In doing so, ensure that no critical files, such as password or system configuration files, can be read from the outside.

26.7.4 CGI Scripts

Interactive scripts in Perl, PHP, SSI, or any other programming language can essentially run arbitrary commands and therefore present a general security issue. Scripts that will be executed from the server should only be installed from sources the server administrator trusts—allowing users to run their own scripts is generally not a good idea. It is also recommended to do security audits for all scripts.

To make the administration of scripts as easy as possible, it is common practice to limit the execution of CGI scripts to specific directories instead of globally allowing them. The directives `ScriptAlias` and `Option ExecCGI` are used for configuration. The SUSE Linux default configuration does not allow execution of CGI scripts from everywhere.

All CGI scripts run as the same user, so different scripts can potentially conflict with each other. The module suEXEC lets you run CGI scripts under a different user and group.

26.7.5 User Directories

When enabling user directories (with `mod_userdir` or `mod_rewrite`) you should strongly consider not allowing `.htaccess` files, which would allow users to overwrite security settings. At least you should limit the user's engagement by using the directive `AllowOverRide`. In SUSE Linux, `.htaccess` files are enabled by default, but the user is not allowed to overwrite any `Option` directives when using `mod_userdir` (see the `/etc/apache2/mod_userdir.conf` configuration file).

26.8 Troubleshooting

If Apache does not start, the Web page is not accessible, or users cannot connect to the Web server, it is important to find the cause of the problem. Here are some typical places to look for error explanations and important things to check.

First, `rcaapache2` (described in [Section 26.3, “Starting and Stopping Apache”](#) (page 463)) is verbose about errors, so can be quite helpful if it is actually used for operating Apache. Sometimes it is tempting to use the binary `/usr/sbin/httpd2` for starting or stopping the Web server. Avoid doing this and use the `rcaapache2` script instead. `rcaapache2` even provides tips and hints for solving configuration errors.

Second, the importance of log files cannot be overemphasized. In case of both fatal and nonfatal errors, the Apache log files, mainly the error log file, are the places to look for causes. Additionally, you can control the verbosity of the logged messages with the `LogLevel` directive if more detail is needed in the log files. By default, the error log file is located at `/var/log/apache2/error_log`.

TIP: A Simple Test

Watch the Apache log messages with the command `tail -F /var/log/apache2/my_error_log`. Then run `rcaapache2 restart`. Now, try to connect with a browser and check the output.

A common mistake is not to open the ports for Apache in the firewall configuration of the server. If you configure Apache with YaST, there is a separate option available to take care of this specific issue (see [Section 26.2.2, “Configuring Apache with YaST”](#) (page 457)). If you are configuring Apache manually, open firewall ports for HTTP and HTTPS via YaST’s firewall module.

If the error cannot be tracked down with the help of any these, check the online Apache bug database at http://httpd.apache.org/bug_report.html. Additionally, the Apache user community can be reached via a mailing list available at <http://httpd.apache.org/userslist.html>. A recommended newsgroup is <comp.infosystems.www.servers.unix>.

26.9 For More Information

The package `apache2-doc` contains the complete Apache manual in various localizations for local installation and reference. It is not installed by default—the quickest way to install it is to use the command `yast -i apache2-doc`. Once installed, the Apache manual is available at <http://localhost/manual/>. You may also access it on the Web at <http://httpd.apache.org/docs-2.2/>. SUSE-specific configuration hints are available in the directory `/usr/share/doc/packages/apache2/README.*`.

26.9.1 Apache 2.2

For a list of new features in Apache 2.2, refer to http://httpd.apache.org/docs/2.2/new_features_2_2.html. Information about upgrading from version 2.0 to 2.2 is available at <http://httpd.apache.org/docs-2.2/upgrading.html>.

26.9.2 Apache Modules

More information about external Apache modules from [Section 26.4.5, “External Modules”](#) (page 470) is available at the following locations:

FastCGI

<http://www.fastcgi.com/>

mod_perl

<http://perl.apache.org/>

mod_php5

<http://www.php.net/manual/en/install.unix.apache2.php>

mod_python

<http://www.modpython.org/>

mod_ruby

<http://www.modruby.net/>

26.9.3 Development

More information about developing Apache modules or about getting involved in the Apache Web server project are available at the following locations:

Apache Developer Information

<http://httpd.apache.org/dev/>

Apache Developer Documentation

<http://httpd.apache.org/docs/2.2/developer/>

Writing Apache Modules with Perl and C

<http://www.modperl.com/>

26.9.4 Miscellaneous Sources

If you experience difficulties specific to Apache in SUSE Linux, take a look at the SUSE Support Database at <http://portal.suse.com/sdb/en/index.html>. The history of Apache is provided at http://httpd.apache.org/ABOUT_APACHE.html. This page also explains why the server is called Apache.

27

File Synchronization

Today, many people use several computers—one computer at home, one or several computers at the workplace, and possibly a laptop or PDA on the road. Many files are needed on all these computers. You may want to be able to work with all computers and modify the files and subsequently have the latest version of the data available on all computers.

27.1 Available Data Synchronization Software

Data synchronization is no problem for computers that are permanently linked by means of a fast network. In this case, use a network file system, like NFS, and store the files on a server, enabling all hosts to access the same data via the network. This approach is impossible if the network connection is poor or not permanent. When you are on the road with a laptop, copies of all needed files must be on the local hard disk. However, it is then necessary to synchronize modified files. When you modify a file on one computer, make sure a copy of the file is updated on all other computers. For occasional copies, this can be done manually with scp or rsync. However, if many files are involved, the procedure can be complicated and requires great care to avoid errors, such as overwriting a new file with an old file.

WARNING: Risk of Data Loss

Before you start managing your data with a synchronization system, you should be well acquainted with the program used and test its functionality. A backup is indispensable for important files.

The time-consuming and error-prone task of manually synchronizing data can be avoided by using one of the programs that use various methods to automate this job. The following summaries are merely intended to convey a general understanding of how these programs work and how they can be used. If you plan to use them, read the program documentation.

27.1.1 Unison

Unison is not a network file system. Instead, the files are simply saved and edited locally. The program Unison can be executed manually to synchronize files. When the synchronization is performed for the first time, a database is created on the two hosts, containing checksums, time stamps, and permissions of the selected files. The next time it is executed, Unison can recognize which files were changed and propose transmission from or to the other host. Usually all suggestions can be accepted.

27.1.2 CVS

CVS, which is mostly used for managing program source versions, offers the possibility to keep copies of the files on multiple computers. Accordingly, it is also suitable for data synchronization. CVS maintains a central repository on the server in which the files and changes to files are saved. Changes that are performed locally are committed to the repository and can be retrieved from other computers by means of an update. Both procedures must be initiated by the user.

CVS is very resilient to errors when changes occur on several computers. The changes are merged and, if changes took place in the same lines, a conflict is reported. When a conflict occurs, the database remains in a consistent state. The conflict is only visible for resolution on the client host.

27.1.3 Subversion

In contrast to CVS, which “evolved,” Subversion (SVN) is a consistently designed project. Subversion was developed as a technically improved successor to CVS.

Subversion has been improved in many respects to its predecessor. Due to its history, CVS only maintains files and is oblivious of directories. Directories also have a version history in Subversion and can be copied and renamed just like files. It is also possible to add metadata to every file and to every directory. This metadata can be fully maintained with versioning. As opposed to CVS, Subversion supports transparent network access over dedicated protocols, like WebDAV (Web-based Distributed Authoring and Versioning). WebDAV extends the functionality of the HTTP protocol to allow collaborative write access to files on remote Web servers.

Subversion was largely assembled on the basis of existing software packages. Therefore, the Apache Web server and the WebDAV extension always run in conjunction with Subversion.

27.1.4 mailsync

Unlike the synchronization tools covered in the previous sections, mailsync only synchronizes e-mails between mailboxes. The procedure can be applied to local mailbox files as well as to mailboxes on an IMAP server.

Based on the message ID contained in the e-mail header, the individual messages are either synchronized or deleted. Synchronization is possible between individual mailboxes and between mailbox hierarchies.

27.1.5 rsync

When no version control is needed but large directory structures need to be synchronized over slow network connections, the tool rsync offers well-developed mechanisms for transmitting only changes within files. This not only concerns text files, but also binary files. To detect the differences between files, rsync subdivides the files into blocks and computes checksums over them.

The effort put into the detection of the changes comes at a price. The systems to synchronize should be scaled generously for the usage of rsync. RAM is especially important.

27.1.6 Novell iFolder

Novell iFolder lets you access your files from anywhere at any time. If you place your files in your iFolder directory, it is instantly synchronized to the server. With this method, you can work from everywhere.

Novell iFolder also lets you work offline. This is very convenient if you do not have an Internet connection, for example, when working with a laptop on the road. After you successfully connect to the Internet, the changes to your work is sent to the server.

Working with iFolder involves these steps:

1. Log in before you work with iFolder.
2. Change your preferences for how often you want to synchronize.
3. Synchronize your files and view the activity between your client and the iFolder server.
4. Resolve any conflicts during the synchronization. Conflicts occur when you change the same file on two different computers. iFolder stores conflicted files in a separate directory to prevent data loss.

For more information about iFolder, see <http://www.novell.com/en-en/documentation/>. Find tips and tricks about iFolder at <http://www.novell.com/coolssolutions/ifmag/>.

27.2 Determining Factors for Selecting a Program

There are some important factors to consider when deciding which program to use.

27.2.1 Client-Server versus Peer-to-Peer

Two different models are commonly used for distributing data. In the first model, all clients synchronize their files with a central server. The server must be accessible by all clients at least occasionally. This model is used by Subversion, CVS, and WebDAV.

The other possibility is to let all networked hosts synchronize their data between each other as peers. This is the concept followed by unison. rsync actually works in client mode, but any client can also act as a server.

27.2.2 Portability

Subversion, CVS, and unison are also available for many other operating systems, including various Unix and Windows systems.

27.2.3 Interactive versus Automatic

In Subversion, CVS, WebDAV, and Unison, the data synchronization is started manually by the user. This allows fine control over the data to synchronize and easy conflict handling. However, if the synchronization intervals are too long, conflicts are more likely to occur.

27.2.4 Conflicts: Incidence and Solution

Conflicts only rarely occur in Subversion or CVS, even when several people work on one large program project. This is because the documents are merged on the basis of individual lines. When a conflict occurs, only one client is affected. Usually conflicts in Subversion or CVS can easily be resolved.

Unison reports conflicts, allowing the affected files to be excluded from the synchronization. However, changes cannot be merged as easily as in Subversion or CVS.

As opposed to Subversion or CVS, where it is possible to partially accept changes in cases of conflict, WebDAV only performs a check-in when the complete modification is considered successful.

There is no conflict handling in rsync. The user is responsible for not accidentally overwriting files and manually resolving all possible conflicts. To be on safe side, a versioning system like RCS can be additionally employed.

27.2.5 Selecting and Adding Files

In its standard configuration, Unison synchronizes an entire directory tree. New files appearing in the tree are automatically included in the synchronization.

In Subversion or CVS, new directories and files must be added explicitly using the command `svn add` or `cvs add`, respectively. This results in greater user control over the files to synchronize. On the other hand, new files are often overlooked, especially when the question marks in the output of `svn update` and `svn status` or `cvs update` are ignored due to the large number of files.

27.2.6 History

An additional feature of Subversion or CVS is that old file versions can be reconstructed. A brief editing remark can be inserted for each change and the development of the files can easily be traced later based on the content and the remarks. This is a valuable aid for theses and program texts.

27.2.7 Data Volume and Hard Disk Requirements

A sufficient amount of free space for all distributed data is required on the hard disks of all involved hosts. Subversion and CVS require additional space for the repository database on the server. The file history is also stored on the server, requiring even more space. When files in text format are changed, only the modified lines need to be saved.

Binary files require additional space amounting to the size of the file every time the file is changed.

27.2.8 GUI

Unison offers a graphical user interface that displays the synchronization procedures Unison wants to perform. Accept the proposal or exclude individual files from the synchronization. In text mode, interactively confirm the individual procedures.

Experienced users normally run Subversion or CVS from the command line. However, graphical user interfaces are available for Linux, such as cervisia, and for other operating systems, like wincvs. Many development tools, such as kdevelop, and text editors, such as emacs, provide support for CVS or Subversion. The resolution of conflicts is often much easier to perform with these front-ends.

27.2.9 User Friendliness

Unison and rsync are rather easy to use and are also suitable for newcomers. CVS and Subversion are somewhat more difficult to operate. Users should understand the interaction between the repository and local data. Changes to the data should first be merged locally with the repository. This is done with the command `cvs update` or `svn update`. Then the data must be sent back to the repository with the command `cvs commit` or `svn commit`. Once this procedure has been understood, newcomers are also able to use CVS or Subversion with ease.

27.2.10 Security against Attacks

During transmission, the data should ideally be protected against interception and manipulation. Unison, CVS, rsync, and Subversion can easily be used via ssh (secure shell), providing security against attacks of this kind. Running CVS or Unison via rsh (remote shell) should be avoided. Accessing CVS with the *pserver* mechanism in insecure networks is likewise not advisable. Subversion already provides the necessary security measures by running with Apache.

27.2.11 Protection against Data Loss

CVS has been used by developers for a long time to manage program projects and is extremely stable. Because the development history is saved, CVS even provides protection against certain user errors, such as unintentional deletion of a file. Despite Subversion not being as common as CVS, it is already being employed in productive environments, for example, by the Subversion project itself.

Unison is still relatively new, but boasts a high level of stability. However, it is more sensitive to user errors. Once the synchronization of the deletion of a file has been confirmed, there is no way to restore the file.

Table 27.1 Features of the File Synchronization Tools: -- = very poor, - = poor or not available, o = medium, + = good, ++ = excellent, x = available

	unison	CVS/SVN	rsync	mailsync
Client/Server	equal	C-S/C-S	C-S	equal
Portability	Lin,Un*x,Win	Lin,Un*x,Win	Lin,Un*x,Win	Lin,Un*x
Interactivity	x	x/x	x	-
Speed	-	o/+	+	+
Conflicts	o	++/++	o	+
File Sel.	Dir.	Sel./file, dir.	Dir.	Mailbox
History	-	x/x	-	-
Hard Disk Space	o	--	o	+
GUI	+	o/o	-	-
Difficulty	+	o/o	+	o
Attacks	+(ssh)	+/(ssh)	+(ssh)	+(SSL)

	unison	CVS/SVN	rsync	mailsync
Data Loss	+	++/++	+	+

27.3 Introduction to Unison

Unison is an excellent solution for synchronizing and transferring entire directory trees. The synchronization is performed in both directions and can be controlled by means of an intuitive graphical front-end. A console version can also be used. The synchronization can be automated so interaction with the user is not required, but experience is necessary.

27.3.1 Requirements

Unison must be installed on the client as well as on the server. In this context, the term *server* refers to a second, remote host (unlike CVS, explained in [Section 27.1.2, “CVS”](#) (page 488)).

In the following section, Unison is used together with ssh. In this case, an SSH client must be installed on the client and an SSH server must be installed on the server.

27.3.2 Using Unison

The approach used by Unison is the association of two directories (*roots*) with each other. This association is symbolic—it is not an online connection. In this example, the directory layout is as follows:

Client: /home/tux/dir1

Server: /home/geeko/dir2

You want to synchronize these two directories. The user is known as tux on the client and as geeko on the server. The first thing to do is to test if the client-server communication works:

```
unison -testserver /home/tux/dir1 ssh://geeko@server//homes/geeko/dir2
```

The most frequently encountered problems are:

- The Unison versions used on the client and server are not compatible.
- The server does not allow SSH connections.
- Neither of the two specified paths exists.

If everything works, omit the option `-testserver`. During the first synchronization, Unison does not yet know the relationship between the two directories and submits suggestions for the transfer direction of the individual files and directories. The arrows in the *Action* column indicate the transfer direction. A question mark means that Unison is not able to make a suggestion regarding the transfer direction because both versions were changed or are new.

The arrow keys can be used to set the transfer direction for the individual entries. If the transfer directions are correct for all displayed entries, simply click *Go*.

The characteristics of Unison (for example, whether to perform the synchronization automatically in clear cases) can be controlled by means of command-line parameters specified when the program is started. View the complete list of all parameters with `unison --help`.

Example 27.1 *The file `~/.unison/exampleprefs`*

```
root=/home/tux/dir1
root=ssh://wilber@server//homes/wilber/dir2
batch=true
```

For each pair, a synchronization log is maintained in the user directory `~/.unison`. Configuration sets, such as `~/.unison/exampleprefs`, can also be stored in this directory. To start the synchronization, specify this file as the command-line parameter as in `unison exampleprefs`.

27.3.3 For More Information

The official documentation of Unison is extremely useful. For this reason, this section merely provides a brief introduction. The complete manual is available at <http://www.cis.upenn.edu/~bcpierce/unison/> and in the SUSE package `unison`.

27.4 Introduction to CVS

CVS is suitable for synchronization purposes if individual files are edited frequently and are stored in a file format, such as ASCII text or program source text. The use of CVS for synchronizing data in other formats, such as JPEG files, is possible, but leads to large amounts of data, because all variants of a file are stored permanently on the CVS server. In such cases, most of the capabilities of CVS cannot be used. The use of CVS for synchronizing files is only possible if all workstations can access the same server.

27.4.1 Configuring a CVS Server

The *server* is the host on which all valid files are located, including the latest versions of all files. Any stationary workstation can be used as a server. If possible, the data of the CVS repository should be included in regular backups.

When configuring a CVS server, it might be a good idea to grant users access to the server via SSH. If the user is known to the server as `tux` and the CVS software is installed on the server as well as on the client, the following environment variables must be set on the client side:

```
CVS_RSH=ssh CVS_ROOT=tux@server:/serverdir
```

The command `cvs init` can be used to initialize the CVS server from the client side. This needs to be done only once.

Finally, the synchronization must be assigned a name. Select or create a directory on the client exclusively to contain files to manage with CVS (the directory can also be empty). The name of the directory is also the name of the synchronization. In this example, the directory is called `synchome`. Change to this directory and enter the following command to set the synchronization name to `synchome`:

```
cvs import synthome tux wilber
```

Many CVS commands require a comment. For this purpose, CVS starts an editor (the editor defined in the environment variable `$EDITOR` or `vi` if no editor was defined). The editor call can be circumvented by entering the comment in advance on the command line, such as in the following example:

```
cvs import -m 'this is a test' synthome tux wilber
```

27.4.2 Using CVS

The synchronization repository can now be checked out from all hosts with `cvs co synchome`. This creates a new subdirectory `synchome` on the client. To commit your changes to the server, change to the directory `synchome` (or one of its subdirectories) and enter `cvs commit`.

By default, all files (including subdirectories) are committed to the server. To commit only individual files or directories, specify them as in `cvs commit file1 directory1`. New files and directories must be added to the repository with a command like `cvs add file1 directory1` before they are committed to the server. Subsequently, commit the newly added files and directories with `cvs commit file1 directory1`.

If you change to another workstation, check out the synchronization repository, if this has not been done during an earlier session at the same workstation (see above).

Start the synchronization with the server with `cvs update`. Update individual files or directories as in `cvs update file1 directory1`. To see the difference between the current files and the versions stored on the server, use the command `cvs diff` or `cvs diff file1 directory1`. Use `cvs -nq update` to see which files would be affected by an update.

Here are some of the status symbols displayed during an update:

U

The local version was updated. This affects all files that are provided by the server and missing on the local system.

M

The local version was modified. If there were changes on the server, it was possible to merge the differences in the local copy.

P

The local version was patched with the version on the server.

C

The local file conflicts with current version in the repository.

?

This file does not exist in CVS.

The status `M` indicates a locally modified file. Either commit the local copy to the server or remove the local file and run the update again. In this case, the missing file is retrieved from the server. If you commit a locally modified file and the file was changed in the same line and committed, you might get a conflict, indicated with `C`.

In this case, look at the conflict marks (`>>` and `<<`) in the file and decide between the two versions. As this can be a rather unpleasant job, you might decide to abandon your changes, delete the local file, and enter `cvs up` to retrieve the current version from the server.

27.4.3 For More Information

This section merely offers a brief introduction to the many possibilities of CVS. Extensive documentation is available at the following URLs:

<http://www.cvshome.org/>
<http://www.gnu.org/manual/>

27.5 Introduction to Subversion

Subversion is a free open source versioning control system and is widely regarded as the successor to CVS, meaning that features already introduced for CVS are normally also in Subversion. It is especially recommended when the advantages of CVS are sought without having to put up with its disadvantages. Many of these features have already been briefly introduced in [Section 27.1.3, “Subversion”](#) (page 489).

27.5.1 Installing a Subversion Server

The installation of a repository database on a server is a relatively simple procedure. Subversion provides a dedicated administration tool for this purpose. The command to enter for creating a new repository is:

```
svnadmin create /path/to/repository
```

Other options can be listed with `svnadmin help`. As opposed to CVS, Subversion is not based on RCS, but rather on different types of repository. The Berkeley Database or FSFS (a repository that uses the file system directly) is commonly used. Do not install a repository on remote file systems, like NFS, AFS, or Windows SMB. The database requires POSIX locking mechanisms, which these file systems do not support.

The command `svnlook` provides information about an existing repository.

```
svnlook info /path/to/repository
```

A server must be configured to allow different users to access the repository. Either use the Apache Web server with WebDAV to do this or use `svnserve`, the server packaged with Subversion. Once `svnserve` is up and running, the repository can be accessed with a URL with `svn://` or `svn+ssh://`. Users that should authenticate themselves when calling `svn` can be set in `/etc/svnserve.conf`.

A decision for Apache or for `svnserve` depends on many factors. It is recommended to browse the Subversion book. More information about it can be found in [Section 27.5.3, “For More Information”](#) (page 502).

27.5.2 Usage and Operation

Use the command `svn` (similar to `cvs`) to access a Subversion repository. With `svn help`, obtain the description of a parameter of a command:

```
checkout (co): Check out a working copy from a repository.  
usage: checkout URL[@REV]... [PATH]
```

If specified, REV determines in which revision the URL is first looked up.

If PATH is omitted, the basename of the URL will be used as the destination. If multiple URLs are given each will be checked out into a sub-directory of PATH, with the name of the sub-directory being the basename of the URL.

...

The content provided by a correctly configured server fitted with a corresponding repository can be accessed by any client with one of the following commands:

```
svn list http://svn.example.com/path/to/project
```

or

```
svn list svn://svn.example.com/path/to/project
```

Save an existing project in the current directory (check it out) with the command `svn checkout`:

```
svn checkout http://svn.example.com/path/to/project nameofproject
```

Checking out creates a new subdirectory `nameofproject` on the client. Operations (adding, copying, renaming, deleting) can then be performed on it:

```
svn add file  
svn copy oldfile newfile  
svn move oldfile newfile  
svn delete file
```

These commands can also be used on directories. Subversion can additionally record properties of a file or directory:

```
svn propset license GPL foo.txt
```

The preceding example sets the value `GPL` for the property `license`. Display properties with `svn proplist`:

```
svn proplist --verbose foo.txt  
Properties on 'foo.txt':  
license : GPL
```

Save the changes to the server with `svn commit`. Another user can incorporate your changes in his working directory by synchronizing with the server using `svn update`.

Unlike CVS, the status of a working directory in Subversion can be displayed *without* accessing the repository with `svn status`. Local changes are displayed in five columns, with the first one being the most important one:

''

No changes.

'A'

Object is marked for addition.

'D'

Object is marked for deletion.

'M'

Object was modified.

'C'	Object is in conflict.
'T'	Object was ignored.
'?'	Object is not being maintained by versioning control.
'!'	Object is reported missing. This flag appears when the object was deleted or moved without the <code>svn</code> command.
'~'	Object was being maintained as a file but has since been replaced by a directory or the opposite has occurred.

The second column shows the status of properties. The meaning of all other columns can be read in the Subversion book.

27.5.3 For More Information

The first point of reference is the home page of the Subversion project at <http://subversion.tigris.org/>. A highly recommendable book can be found in the directory `file:///usr/share/doc/packages/subversion/html/book.html` after installation of the package `subversion-doc` and is also available online at <http://svnbook.red-bean.com/svnbook/index.html>.

27.6 Introduction to rsync

`rsync` is useful when large amounts of data need to be transmitted regularly while not changing too much. This is, for example, often the case when creating backups. Another application concerns staging servers. These are servers that store complete directory trees of Web servers that are regularly mirrored onto a Web server in a DMZ.

27.6.1 Configuration and Operation

rsync can be operated in two different modes. It can be used to archive or copy data. To accomplish this, only a remote shell, like ssh, is required on the target system. However, rsync can also be used as a daemon to provide directories to the network.

The basic mode of operation of rsync does not require any special configuration. rsync directly allows mirroring complete directories onto another system. As an example, the following command creates a backup of the home directory of tux on a backup server named sun:

```
rsync -az -e ssh /home/tux/ tux@sun:backup
```

The following command is used to play the directory back:

```
rsync -az -e ssh tux@sun:backup /home/tux/
```

Up to this point, the handling does not differ much from that of a regular copying tool, like scp.

rsync should be operated in “rsync” mode to make all its features fully available. This is done by starting the rsyncd daemon on one of the systems. Configure it in the file `/etc/rsyncd.conf`. For example, to make the directory `/srv/ftp` available with rsync, use the following configuration:

```
gid = nobody
uid = nobody
read only = true
use chroot = no
transfer logging = true
log format = %h %o %f %l %b
log file = /var/log/rsyncd.log

[FTP]
    path = /srv/ftp
    comment = An Example
```

Then start rsyncd with `rcrsyncd start`. rsyncd can also be started automatically during the boot process. Set this up by activating this service in the runlevel editor provided by YaST or by manually entering the command `insserv rsyncd`. rsyncd can alternatively be started by xinetd. This is, however, only recommended for servers that rarely use rsyncd.

The example also creates a log file listing all connections. This file is stored in `/var/log/rsyncd.log`.

It is then possible to test the transfer from a client system. Do this with the following command:

```
rsync -avz sun::FTP
```

This command lists all files present in the directory `/srv/ftp` of the server. This request is also logged in the log file `/var/log/rsyncd.log`. To start an actual transfer, provide a target directory. Use `.` for the current directory. For example:

```
rsync -avz sun::FTP .
```

By default, no files are deleted while synchronizing with rsync. If this should be forced, the additional option `--delete` must be stated. To ensure that no newer files are deleted, the option `--update` can be used instead. Any conflicts that arise must be resolved manually.

27.6.2 For More Information

Important information about rsync is provided in the man pages `man rsync` and `man rsyncd.conf`. A technical reference about the operating principles of rsync is featured in `/usr/share/doc/packages/rsync/tech_report.ps`. Find latest news about rsync on the project Web site at <http://rsync.samba.org/>.

27.7 Introduction to mailsync

mailsync is mainly suitable for the following three tasks:

- Synchronization of locally stored e-mails with mails stored on a server
- Migration of mailboxes to a different format or to a different server
- Integrity check of a mailbox or search for duplicates

27.7.1 Configuration and Use

mailsync distinguishes between the mailbox itself (the *store*) and the connection between two mailboxes (the *channel*). The definitions of the stores and channels are stored in `~/.mailsync`. The following paragraphs explain a number of store examples.

A simple definition might appear as follows:

```
store saved-messages {  
    pat Mail/saved-messages  
prefix  Mail/  
}
```

Mail/ is a subdirectory of the user's home directory that contains e-mail folders, including the folder saved-messages. If mailsync is started with mailsync -m saved-messages, it lists an index of all messages in saved-messages. If the following definition is made

```
store localdir {  
pat      Mail/*  
prefix  Mail/  
}
```

the command mailsync -m localdir lists all messages stored under Mail/. In contrast, the command mailsync localdir lists the folder names. The specifications of a store on an IMAP server appear as follows:

```
store imapinbox {  
server {mail.edu.harvard.com/user=gulliver}  
ref   {mail.edu.harvard.com}  
pat   INBOX  
}
```

The above example merely addresses the main folder on the IMAP server. A store for the subfolders would appear as follows:

```
store imapdir {  
server {mail.edu.harvard.com/user=gulliver}  
ref {mail.edu.harvard.com}  
pat INBOX.*  
prefix  INBOX.  
}
```

If the IMAP server supports encrypted connections, the server specification should be changed to

```
server {mail.edu.harvard.com/ssl/user=gulliver}
```

or, if the server certificate is not known, to

```
server {mail.edu.harvard.com/ssl/novalidate-cert/user=gulliver}
```

The prefix is explained later.

Now the folders under `Mail/` should be connected to the subdirectories on the IMAP server:

```
channel folder localdir imapdir {  
msinfo .mailsync.info  
}
```

`mailsync` uses the `msinfo` file to keep track of the messages that have already been synchronized.

The command `mailsync folder` does the following:

- Expands the mailbox pattern on both sides.
- Removes the prefix from the resulting folder names.
- Synchronizes the folders in pairs (or creates them if they do not exist).

Accordingly, the folder `INBOX.sent-mail` on the IMAP server is synchronized with the local folder `Mail/sent-mail` (provided the definitions explained above exist). The synchronization between the individual folder is performed as follows:

- If a message already exists on both sides, nothing happens.
- If the message is missing on one side and is new (not listed in the `msinfo` file), it is transmitted there.
- If the message merely exists on one side and is old (already listed in the `msinfo` file), it is deleted there (because the message that had obviously existed on the other side was deleted).

To know in advance which messages will be transmitted and which will be deleted during a synchronization, start `mailsync` with a channel *and* a store with `mailsync folder localdir`. This command produces a list of all messages that are new on the local host as well as a list of all messages that would be deleted on the IMAP side during a synchronization. Similarly, the command `mailsync folder imapdir` produces a list of all messages that are new on the IMAP side and a list of all messages that would be deleted on the local host during a synchronization.

27.7.2 Possible Problems

In the event of a data loss, the safest method is to delete the relevant channel log file `msinfo`. Accordingly, all messages that only exist on one side are viewed as new and are therefore transmitted during the next synchronization.

Only messages with a message ID are included in the synchronization. Messages lacking a message ID are simply ignored, which means they are not transmitted or deleted. A missing message ID is usually caused by faulty programs when sending or writing a message.

On certain IMAP servers, the main folder is addressed with `INBOX` and subfolders are addressed with a randomly selected name (in contrast to `INBOX` and `INBOX.name`). Therefore, for such IMAP servers, it is not possible to specify a pattern exclusively for the subfolders.

After the successful transmission of messages to an IMAP server, the mailbox drivers (c-client) used by `mailsync` set a special status flag. For this reason, some e-mail programs, like `mutt`, are not able to recognize these messages as new. Disable the setting of this special status flag with the option `-n`.

27.7.3 For More Information

The README in `/usr/share/doc/packages/mailsync/`, which is included in `mailsync`, provides additional information. In this connection, RFC 2076 “Common Internet Message Headers” is of special interest.

28

Samba

Using Samba, a Unix machine can be configured as a file and print server for DOS, Windows, and OS/2 machines. Samba has developed into a fully-fledged and rather complex product. Configure Samba with YaST, SWAT (a Web interface), or the configuration file.

28.1 Terminology

SMB protocol

Samba uses the SMB (server message block) protocol that is based on the NetBIOS services. Due to pressure from IBM, Microsoft released the protocol so other software manufacturers could establish connections to a Microsoft domain network. With Samba, the SMB protocol works on top of the TCP/IP protocol, so the TCP/IP protocol must be installed on all clients.

CIFS protocol

CIFS (common Internet file system) protocol is another protocol supported by Samba. CIFS defines a standard remote file system access protocol for use over the network, enabling groups of users to work together and share documents across the network.

NetBIOS

NetBIOS is a software interface (API) designed for communication between machines. Here, a name service is provided. It enables machines connected to the network to reserve names for themselves. After reservation, these machines can be addressed by name. There is no central process that checks names. Any machine

on the network can reserve as many names as it wants as long as the names are not already in use. The NetBIOS interface can now be implemented for different network architectures. An implementation that works relatively closely with network hardware is called NetBEUI, but this is often referred to as NetBIOS. Network protocols implemented with NetBIOS are IPX from Novell (NetBIOS via TCP/IP) and TCP/IP.

The NetBIOS names sent via TCP/IP have nothing in common with the names used in `/etc/hosts` or those defined by DNS. NetBIOS uses its own, completely independent naming convention. However, it is recommended to use names that correspond to DNS hostnames to make administration easier. This is the default used by Samba.

Samba server

Samba server is a server that provides SMB/CIFS services and NetBIOS over IP naming services to clients. For Linux, there are two daemons for Samba server: smnd for SMB/CIFS services and nmbd for naming services.

Samba client

Samba client is a system that uses Samba services from a Samba server over the SMB protocol. All common operating systems, such as Mac OS X, Windows, and OS/2, support the SMB protocol. The TCP/IP protocol must be installed on all computers. Samba provides a client for the different UNIX flavors. For Linux, there is a kernel module for SMB that allows the integration of SMB resources on the Linux system level. You do not need run any daemon for Samba client.

Shares

SMB servers provide hardware space to their clients by means of shares. Shares are printers and directories with their subdirectories on the server. It is exported by means of a name and can be accessed by its name. The share name can be set to any name—it does not have to be the name of the export directory. A printer is also assigned a name. Clients can access the printer by its name.

28.2 Starting and Stopping Samba

You can start or stop the Samba server automatically during boot or manually. Starting and stopping policy is a part of the YaST Samba server configuration described in [Section 28.3.1, “Configuring a Samba Server with YaST”](#) (page 511).

To stop or start running Samba services with YaST, use *System* → *System Services (Runlevel)*. From a command line, stop services required for Samba with `rcsmb stop && rcnmb stop` and start them with `rcnmb start && rcsmb start`.

28.3 Configuring a Samba Server

A samba server in SUSE Linux can be configured in two different ways: with YaST or manually. Manual configuration offers a higher level of detail, but lacks the convenience of the YaST GUI.

28.3.1 Configuring a Samba Server with YaST

To configure a Samba server, start YaST and select *Network Services* → *Samba Server*. When starting the module for the first time, the *Samba Server Installation* dialog starts, prompting you to make just a few basic decisions concerning administration of the server then at the end of the configuration prompts for the password of Samba root. For later starts, the *Samba Server Configuration* dialog appears.

The *Samba Server Installation* dialog consists of two steps:

Workgroup or Domain Name

Select an existing name from *Workgroup or Domain Name* or enter a new one and click *Next*.

Samba Server Type

In the next step, specify whether your server should act as PDC and click *Next*.

You can change all settings from *Samba Server Installation* later in the *Samba Server Configuration* dialog with the *Identity* tab.

During first start of Samba server module the *Samba Server Configuration* dialog appears directly after *Samba Server Installation* dialog. It consists of three tabs:

Start Up

In the *Start Up* tab, you can set starting of the Samba server. To start the service every time your system boots, select *During Boot*. To activate manual start, choose *Manually*. More information about starting a Samba server is provided in [Section 28.2, “Starting and Stopping Samba”](#) (page 510).

In this tab, you can also open ports in your firewall. To do so, select *Open Port in Firewall*. If you have multiple network interfaces, select the network interface for Samba services by clicking *Firewall Details*, selecting the interfaces, and clicking *OK*.

Shares

In this tab, determine the Samba shares to activate. There are some predefined shares, like homes and printers. Use *Toggle Status* to switch between *Active* and *Inactive*. Click *Add* to add new shares and *Delete* to delete the selected share.

Identity

In the *Identity* tab, you can determine the domain with which the host is associated (*Base Settings*) and whether to use an alternative hostname in the network (*NetBIOS Host Name*). To set expert global settings or set user authentication, click *Advanced Settings*.

Click *Finish* to close the configuration.

28.3.2 Web Administration with SWAT

An alternative tool for Samba server administration is SWAT (Samba Web Administration Tool). It provides a simple Web interface with which to configure the Samba server. To use SWAT, open <http://localhost:901> in a Web browser and log in as user `root`. If you do not have a special Samba root account, use the system `root` account.

NOTE: Activating SWAT

After Samba server installation, SWAT is not activated. To activate it, open *Network Services* → *Network Services (xinetd)* in YaST, enable the network services configuration, select `swat` from the table, and click *Toggle Status (On or Off)*.

28.3.3 Configuring the Server Manually

If you intend to use Samba as a server, install `samba`. The main configuration file of Samba is `/etc/samba/smb.conf`. This file can be divided into two logical parts. The `[global]` section contains the central and global settings. The `[share]` sections contain the individual file and printer shares. By means of this approach, details regarding the shares can be set differently or globally in the `[global]` section, which enhances the structural transparency of the configuration file.

The global Section

The following parameters of the `[global]` section need some adjustment to match the requirements of your network setup so other machines can access your Samba server via SMB in a Windows environment.

`workgroup = TUX-NET`

This line assigns the Samba server to a workgroup. Replace `TUX-NET` with an appropriate workgroup of your networking environment. Your Samba server appears under its DNS name unless this name has been assigned to any other machine in the network. If the DNS name is not available, set the server name using `netbiosname=MYNAME`. See `man smb.conf` for more details about this parameter.

`os level = 2`

This parameter triggers whether your Samba server tries to become LMB (local master browser) for its workgroup. Choose a very low value to spare the existing Windows network from any disturbances caused by a misconfigured Samba server. More information about this important topic can be found in the files `BROWSING.txt` and `BROWSING-Config.txt` under the `textdocs` subdirectory of the package documentation.

If no other SMB server is present in your network (such as a Windows NT or 2000 server) and you want the Samba server to keep a list of all systems present in the local environment, set the `os_level` to a higher value (for example, 65). Your Samba server is then chosen as LMB for your local network.

When changing this setting, consider carefully how this could affect an existing Windows network environment. First test the changes in an isolated network or at a noncritical time of day.

wins support and wins server

To integrate your Samba server into an existing Windows network with an active WINS server, enable the `wins server` option and set its value to the IP address of that WINS server.

If your Windows machines are connected to separate subnets and should still be aware of each other, you need to set up a WINS server. To turn a Samba server into such a WINS server, set the option `wins support = Yes`. Make sure that only one Samba server of the network has this setting enabled. The options `wins server` and `wins support` must never be enabled at the same time in your `smb.conf` file.

Shares

The following examples illustrate how a CD-ROM drive and the user directories (`homes`) are made available to the SMB clients.

[cdrom]

To avoid having the CD-ROM drive accidentally made available, these lines are deactivated with comment marks (semicolons in this case). Remove the semicolons in the first column to share the CD-ROM drive with Samba.

Example 28.1 A CD-ROM Share

```
; [cdrom]
;           comment = Linux CD-ROM
;           path = /media/cdrom
;           locking = No
```

[cdrom] and comment

The entry `[cdrom]` is the name of the share that can be seen by all SMB clients on the network. An additional `comment` can be added to further describe the share.

```
path = /media/cdrom
path exports the directory /media/cdrom.
```

By means of a very restrictive default configuration, this kind of share is only made available to the users present on this system. If this share should be made available to everybody, add a line `guest ok = yes` to the configuration. This setting gives read permissions to anyone on the network. It is recommended to handle this

parameter with great care. This applies even more to the use of this parameter in the [global] section.

[homes]

The [home] share is of special importance here. If the user has a valid account and password for the Linux file server and his own home directory, he can be connected to it.

Example 28.2 homes Share

```
[homes]
comment = Home Directories
valid users = %S
browseable = No
read only = No
create mask = 0640
directory mask = 0750
```

[homes]

As long as there is no other share using the share name of the user connecting to the SMB server, a share is dynamically generated using the [homes] share directives. The resulting name of the share is the username.

valid users = %S

%S is replaced with the concrete name of the share as soon as a connection has been successfully established. For a [homes] share, this is always the username. As a consequence, access rights to a user's share are restricted exclusively to the user.

browseable = No

This setting makes the share invisible in the network environment.

read only = No

By default, Samba prohibits write access to any exported share by means of the read only = Yes parameter. To make a share writable, set the value read only = No, which is synonymous with writable = Yes.

create mask = 0640

Systems that are not based on MS Windows NT do not understand the concept of UNIX permissions, so they cannot assign permissions when creating a file. The parameter create mask defines the access permissions assigned to newly created files. This only applies to writable shares. In effect, this setting means the owner has read and write permissions and the members of the

owner's primary group have read permissions. `valid users = %S` prevents read access even if the group has read permissions. For the group to have read or write access, deactivate the line `valid users = %S`.

Security Levels

To improve security, each share access can be protected with a password. SMB has three possible ways of checking the permissions:

Share Level Security (security = share)

A password is firmly assigned to a share. Everyone who knows this password has access to that share.

User Level Security (security = user)

This variation introduces the concept of the user to SMB. Each user must register with the server with his own password. After registration, the server can grant access to individual exported shares dependent on usernames.

Server Level Security (security = server):

To its clients, Samba pretends to be working in user level mode. However, it passes all password queries to another user level mode server, which takes care of authentication. This setting expects an additional parameter (`password server`).

The selection of share, user, or server level security applies to the entire server. It is not possible to offer individual shares of a server configuration with share level security and others with user level security. However, you can run a separate Samba server for each configured IP address on a system.

More information about this subject can be found in the Samba HOWTO Collection. For multiple servers on one system, pay attention to the options `interfaces` and `bind interfaces only`.

28.4 Configuring Clients

Clients can only access the Samba server via TCP/IP. NetBEUI and NetBIOS via IPX cannot be used with Samba.

28.4.1 Configuring a Samba Client with YaST

Configure a Samba client to access resources (files or printers) on the Samba server. Enter the domain or workgroup in the dialog *Samba Workgroup*. Click *Browse* to display all available groups and domains, which can be selected with the mouse. If you activate *Also Use SMB Information for Linux Authentication*, the user authentication runs over the Samba server. After completing all settings, click *Finish* to finish the configuration.

28.4.2 Windows 9x and ME

Windows 9x and ME already have built-in support for TCP/IP. However, this is not installed as the default. To add TCP/IP, go to *Control Panel* → *System* and choose *Add → Protocols* → *TCP/IP from Microsoft*. After rebooting your Windows machine, find the Samba server by double-clicking the desktop icon for the network environment.

TIP

To use a printer on the Samba server, install the standard or Apple-PostScript printer driver from the corresponding Windows version. It is best to link this to the Linux printer queue, which accepts Postscript as an input format.

28.5 Samba as Login Server

In networks where predominantly Windows clients are found, it is often preferable that users may only register with a valid account and password. In a Windows-based network, this task is handled by a Windows NT server configured as a primary domain controller (PDC) but this can also be done with the help of a Samba server. The entries that must be made in the [global] section of `smb.conf` are shown in [Example 28.3, “Global Section in smb.conf”](#) (page 517).

Example 28.3 *Global Section in smb.conf*

```
[global]
workgroup = TUX-NET
domain logons = Yes
domain master = Yes
```

If encrypted passwords are used for verification purposes—this is the default setting with well-maintained MS Windows 9x installations, MS Windows NT 4.0 from service pack 3, and all later products—the Samba server must be able to handle these. The entry `encrypt passwords = yes` in the `[global]` section enables this (with Samba version 3, this is now the default). In addition, it is necessary to prepare user accounts and passwords in an encryption format that conforms with Windows. Do this with the command `smbpasswd -a name`. Create the domain account for the computers, required by the Windows NT domain concept, with the following commands:

Example 28.4 *Setting Up a Machine Account*

```
useradd hostname\$  
    smbpasswd -a -m hostname
```

With the `useradd` command, a dollar sign is added. The command `smbpasswd` inserts this automatically when the parameter `-m` is used. The commented configuration example (`/usr/share/doc/packages/Samba/examples/smb.conf.SuSE`) contains settings that automate this task.

Example 28.5 *Automated Setup of a Machine Account*

```
add machine script = /usr/sbin/useradd -g nogroup -c "NT Machine Account" \  
    -s /bin/false %m\$
```

To make sure that Samba can execute this script correctly, choose a Samba user with the required administrator permissions. To do so, select one user and add it to the `ntadmin` group. After that, all users belonging to this Linux group can be assigned `Domain Admin` status with the command:

```
net groupmap add ntgroup="Domain Admins" unixgroup=ntadmin
```

More information about this topic is provided in Chapter 12 of the Samba HOWTO Collection, found in `/usr/share/doc/packages/samba/Samba-HOWTO-Collection.pdf`.

28.6 For More Information

Detailed Samba information is available in the digital documentation. Enter `apropos samba` at the command line to display some manual pages or just browse the `/usr/share/doc/packages/samba` directory if Samba documentation is installed for

more online documentation and examples. Find a commented example configuration (`smb.conf.SuSE`) in the `examples` subdirectory.

The Samba HOWTO Collection provided by the Samba team includes a section about troubleshooting. In addition to that, Part V of the document provides a step-by-step guide to checking your configuration. You can find Samba HOWTO Collection in `/usr/share/doc/packages/samba/Samba-HOWTO-Collection.pdf` after installing package `samba-doc`.

The Proxy Server Squid

Squid is a widely-used proxy cache for Linux and UNIX platforms. This means that it stores requested Internet objects, such as data on a Web or FTP server, on a machine that is closer to the requesting workstation than the server. This chapter discusses its configuration, the settings required to get it running, how to configure the system to do transparent proxying, how to gather statistics about using the cache with the help of programs, like Calamaris and cachemgr, and how to filter Web contents with squidGuard.

Squid acts as a proxy cache. It redirects object requests from clients (in this case, from Web browsers) to the server. When the requested objects arrive from the server, it delivers the objects to the client and keeps a copy of them in the hard disk cache. One of the advantages of caching is that several clients requesting the same object can be served from the hard disk cache. This enables clients to receive the data much faster than from the Internet. This procedure also reduces the network traffic.

Along with the actual caching, Squid offers a wide range of features such as distributing the load over intercommunicating hierarchies of proxy servers, defining strict access control lists for all clients accessing the proxy, allowing or denying access to specific Web pages with the help of other applications, and generating statistics about frequently-visited Web pages for the assessment of the users' surfing habits. Squid is not a generic proxy. It normally proxies only HTTP connections. It does also support the protocols FTP, Gopher, SSL, and WAIS, but it does not support other Internet protocols, such as Real Audio, news, or video conferencing. Because Squid only supports the UDP protocol to provide communication between different caches, many other multimedia programs are not supported.

29.1 Some Facts about Proxy Caches

As a proxy cache, Squid can be used in several ways. When combined with a firewall, it can help with security. Multiple proxies can be used together. It can also determine what types of objects should be cached and for how long.

29.1.1 Squid and Security

It is possible to use Squid together with a firewall to secure internal networks from the outside using a proxy cache. The firewall denies all clients access to external services except Squid. All Web connections must be established by the proxy. With this configuration, Squid completely controls Web access.

If the firewall configuration includes a DMZ, the proxy should operate within this zone. [Section 29.5, “Configuring a Transparent Proxy”](#) (page 533) describes how to implement a *transparent* proxy. This simplifies the configuration of the clients, because in this case they do not need any information about the proxy.

29.1.2 Multiple Caches

Several instances of Squid can be configured to exchange objects between them. This reduces the total system load and increases the chances of finding an object already existing in the local network. It is also possible to configure cache hierarchies, so a cache is able to forward object requests to sibling caches or to a parent cache—causing it to get objects from another cache in the local network or directly from the source.

Choosing the appropriate topology for the cache hierarchy is very important, because it is not desirable to increase the overall traffic on the network. For a very large network, it would make sense to configure a proxy server for every subnetwork and connect them to a parent proxy, which in turn is connected to the proxy cache of the ISP.

All this communication is handled by ICP (Internet cache protocol) running on top of the UDP protocol. Data transfers between caches are handled using HTTP (hypertext transmission protocol) based on TCP.

To find the most appropriate server from which to get the objects, one cache sends an ICP request to all sibling proxies. These answer the requests via ICP responses with a

HIT code if the object was detected or a MISS if it was not. If multiple HIT responses were found, the proxy server decides from which server to download, depending on factors such as which cache sent the fastest answer or which one is closer. If no satisfactory responses are received, the request is sent to the parent cache.

TIP

To avoid duplication of objects in different caches in the network, other ICP protocols are used, such as CARP (cache array routing protocol) or HTCP (hypertext cache protocol). The more objects maintained in the network, the greater the possibility of finding the desired one.

29.1.3 Caching Internet Objects

Not all objects available in the network are static. There are a lot of dynamically generated CGI pages, visitor counters, and encrypted SSL content documents. Objects like this are not cached because they change each time they are accessed.

The question remains as to how long all the other objects stored in the cache should stay there. To determine this, all objects in the cache are assigned one of various possible states. Web and proxy servers find out the status of an object by adding headers to these objects, such as “Last modified” or “Expires” and the corresponding date. Other headers specifying that objects must not be cached are used as well.

Objects in the cache are normally replaced, due to a lack of free hard disk space, using algorithms such as LRU (last recently used). Basically this means that the proxy expunges the objects that have not been requested for the longest time.

29.2 System Requirements

The most important thing is to determine the maximum network load the system must bear. It is, therefore, important to pay more attention to the load peaks, because these might be more than four times the day's average. When in doubt, it would be better to overestimate the system's requirements, because having Squid working close to the limit of its capabilities could lead to a severe loss in the quality of the service. The following sections point to the system factors in order of significance.

29.2.1 Hard Disks

Speed plays an important role in the caching process, so this factor deserves special attention. For hard disks, this parameter is described as *random seek time*, measured in milliseconds. Because the data blocks that Squid reads from or writes to the hard disk tend to be rather small, the seek time of the hard disk is more important than its data throughput. For the purposes of a proxy, hard disks with high rotation speeds are probably the better choice, because they allow the read-write head to be positioned in the required spot more quickly. One possibility to speed up the system is to use a number of disks concurrently or to employ striping RAID arrays.

29.2.2 Size of the Disk Cache

In a small cache, the probability of a HIT (finding the requested object already located there) is small, because the cache is easily filled and the less requested objects are replaced by newer ones. If, for example, one GB is available for the cache and the users only surf ten MB per day, it would take more than one hundred days to fill the cache.

The easiest way to determine the needed cache size is to consider the maximum transfer rate of the connection. With a 1 Mbit/s connection, the maximum transfer rate is 125 KB/s. If all this traffic ends up in the cache, in one hour it would add up to 450 MB and, assuming that all this traffic is generated in only eight working hours, it would reach 3.6 GB in one day. Because the connection is normally not used to its upper volume limit, it can be assumed that the total data volume handled by the cache is approximately 2 GB. This is why 2 GB of disk space is required in the example for Squid to keep one day's worth of browsed data cached.

29.2.3 RAM

The amount of memory (RAM) required by Squid directly correlates to the number of objects in the cache. Squid also stores cache object references and frequently requested objects in the main memory to speed up retrieval of this data. Random access memory is much faster than a hard disk.

In addition to that, there is other data that Squid needs to keep in memory, such as a table with all the IP addresses handled, an exact domain name cache, the most frequently requested objects, access control lists, buffers, and more.

It is very important to have sufficient memory for the Squid process, because system performance is dramatically reduced if it must be swapped to disk. The `cachemgr.cgi` tool can be used for the cache memory management. This tool is introduced in [Section 29.6, “cachemgr.cgi” \(page 536\)](#). Sites with huge network traffic should consider using an AMD64 or Intel EM64T system with more than 4 GB of memory.

29.2.4 CPU

Squid is not a program that requires intensive CPU usage. The load of the processor is only increased while the contents of the cache are loaded or checked. Using a multiprocessor machine does not increase the performance of the system. To increase efficiency, it is better to buy faster disks or add more memory.

29.3 Starting Squid

Squid is already preconfigured in SUSE Linux, you can start it right after the installation. To ensure a smooth start-up, the network should be configured in a way that at least one name server and the Internet can be reached. Problems can arise if a dial-up connection is used with a dynamic DNS configuration. In this case, at least the name server should be entered, because Squid does not start if it does not detect a DNS server in `/etc/resolv.conf`.

29.3.1 Commands for Starting and Stopping Squid

To start Squid, enter `rcsquid start` at the command line as `root`. In the initial start-up, the directory structure of the cache must first be defined in `/var/cache/squid`. This is done automatically by the start script `/etc/init.d/squid` and can take a few seconds or even minutes. If `done` appears to the right in green, Squid has been successfully loaded. To test the functionality of Squid on the local system, enter `localhost` as the proxy and `3128` as the port in the browser.

To allow users from the local system and other systems to access Squid and the Internet, change the entry in the configuration files `/etc/squid/squid.conf` from `http_access deny all` to `http_access allow all`. However, in doing

so, consider that Squid is made completely accessible to anyone by this action. Therefore, define ACLs that control access to the proxy. More information about this is available in [Section 29.4.2, “Options for Access Controls”](#) (page 530).

After modifying the configuration file `/etc/squid/squid.conf`, Squid must reload the configuration file. Do this with `rcsquid reload`. Alternatively, completely restart Squid with `rcsquid restart`.

The command `rcsquid status` can be used to check if the proxy is running. The command `rcsquid stop` causes Squid to shut down. This can take a while, because Squid waits up to half a minute (`shutdown_lifetime` option in `/etc/squid/squid.conf`) before dropping the connections to the clients and writing its data to the disk.

WARNING: Terminating Squid

Terminating Squid with `kill` or `killall` can damage the cache. To be able to restart Squid, the damaged cache must be deleted.

If Squid dies after a short period of time even though it was started successfully, check whether there is a faulty name server entry or whether the `/etc/resolv.conf` file is missing. Squid logs the cause of a start-up failure in the file `/var/log/squid/cache.log`. If Squid should be loaded automatically when the system boots, use the YaST runlevel editor to activate Squid for the desired runlevels. See Section “System Services (Runlevel)” (Chapter 2, *System Configuration with YaST*, ↑Start-Up).

An uninstall of Squid does not remove the cache hierarchy or the log files. To remove these, delete the `/var/cache/squid` directory manually.

29.3.2 Local DNS Server

Setting up a local DNS server makes sense even if it does not manage its own domain. It then simply acts as a caching-only name server and is also able to resolve DNS requests via the root name servers without requiring any special configuration (see [Section 20.3, “Starting the Name Server BIND”](#) (page 376)). How this can be done depends on whether you chose dynamic DNS during the configuration of the Internet connection.

Dynamic DNS

Normally, with dynamic DNS, the DNS server is set by the provider during the establishment of the Internet connection and the local file `/etc/resolv.conf` is adjusted automatically. This behavior is controlled in the file `/etc/sysconfig/network/config` with the sysconfig variable `MODIFY_RESOLV_CONF_DYNAMICALLY`, which is set to "yes". Set this variable to "no" with the YaST sysconfig editor (see [Section 8.3.1, “Changing the System Configuration Using the YaST sysconfig Editor”](#) (page 190)). Then enter the local DNS server in the file `/etc/resolv.conf` with the IP address `127.0.0.1` for `localhost`. This way Squid can always find the local name server when it starts.

To make the provider's name server accessible, enter it in the configuration file `/etc/named.conf` under `forwarders` along with its IP address. With dynamic DNS, this can be achieved automatically during connection establishment by setting the sysconfig variable `MODIFY_NAMED_CONF_DYNAMICALLY` to YES.

Static DNS

With static DNS, no automatic DNS adjustments take place while establishing a connection, so there is no need to change any sysconfig variables. You must, however, enter the local DNS server in the file `/etc/resolv.conf` as described above. Additionally, the providers static name server must be entered manually in the file `/etc/named.conf` under `forwarders` along with its IP address.

TIP: DNS and Firewall

If you have a firewall running, make sure DNS requests can pass it.

29.4 The Configuration File `/etc/squid/squid.conf`

All Squid proxy server settings are made in the `/etc/squid/squid.conf` file. To start Squid for the first time, no changes are necessary in this file, but external clients are initially denied access. The proxy is available for `localhost`. The default port is 3128. The preinstalled configuration file `/etc/squid/squid.conf` provides detailed information about the options and many examples. Nearly all entries begin with # (the lines are commented) and the relevant specifications can be found at the

end of the line. The given values almost always correlate with the default values, so removing the comment signs without changing any of the parameters actually has little effect in most cases. If possible, leave the sample as it is and insert the options along with the modified parameters in the line below. This way, the default values may easily be recovered and compared with the changes.

TIP: Adapting the Configuration File after an Update

If you have updated from an earlier Squid version, it is recommended to edit the new `/etc/squid/squid.conf` and only apply the changes made in the previous file. If you try to use the old `squid.conf`, risk that the configuration no longer works, because options are sometimes modified and new changes added.

29.4.1 General Configuration Options (Selection)

`http_port 3128`

This is the port on which Squid listens for client requests. The default port is 3128, but 8080 is also common. If desired, specify several port numbers separated by blank spaces.

`cache_peer hostname type proxy-port icp-port`

Here, enter a parent proxy, for example, if you want to use the proxy of your ISP. As `hostname`, enter the name and IP address of the proxy to use and, as `type`, enter `parent`. For `proxy-port`, enter the port number that is also given by the operator of the parent for use in the browser, usually 8080. Set the `icp-port` to 7 or 0 if the ICP port of the parent is not known and its use is irrelevant to the provider. In addition, `default` and `no-query` may be specified after the port numbers to prohibit the use of the ICP protocol. Squid then behaves like a normal browser as far as the provider's proxy is concerned.

`cache_mem 8 MB`

This entry defines the amount of memory Squid can use for very popular replies. The default is 8 MB. This does not specify the memory usage of Squid and may be exceeded.

cache_dir ufs /var/cache/squid/ 100 16 256

The entry *cache_dir* defines the directory where all the objects are stored on disk. The numbers at the end indicate the maximum disk space in MB to use and the number of directories in the first and second level. The *ufs* parameter should be left alone. The default is 100 MB occupied disk space in the */var/cache/squid* directory and creation of 16 subdirectories inside it, each containing 256 more subdirectories. When specifying the disk space to use, leave sufficient reserve disk space. Values from a minimum of 50% to a maximum of 80% of the available disk space make the most sense here. The last two numbers for the directories should only be increased with caution, because too many directories can also lead to performance problems. If you have several disks that share the cache, enter several *cache_dir* lines.

cache_access_log /var/log/squid/access.log, **cache_log** /var/log/squid/cache.log,
cache_store_log /var/log/squid/store.log

These three entries specify the paths where Squid logs all its actions. Normally, nothing is changed here. If Squid is experiencing a heavy usage burden, it might make sense to distribute the cache and the log files over several disks.

emulate_httpd_log off

If the entry is set to *on*, obtain readable log files. Some evaluation programs cannot interpret this, however.

client_netmask 255.255.255.255

With this entry, mask IP addresses of clients in the log files. The last digit of the IP address is set to zero if you enter 255.255.255.0 here. You may protect the privacy of your clients that way.

ftp_user Squid@

With this, set the password Squid should use for the anonymous FTP login. It can make sense to specify a valid e-mail address here, because some FTP servers check these for validity.

cache_mgr webmaster

An e-mail address to which Squid sends a message if it unexpectedly crashes. The default is *webmaster*.

logfile_rotate 0

If you run *squid -k rotate*, Squid can rotate secured log files. The files are numbered in this process and, after reaching the specified value, the oldest file is

overwritten. The default value is 0 because archiving and deleting log files in SUSE Linux is carried out by a cron job set in the configuration file /etc/logrotate/squid.

append_domain <domain>

With *append_domain*, specify which domain to append automatically when none is given. Usually, your own domain is entered here, so entering *www* in the browser accesses your own Web server.

forwarded_for on

If you set the entry to *off*, Squid removes the IP address and the system name of the client from HTTP requests. Otherwise it adds a line to the header like

```
X-Forwarded-For: 192.168.0.0
```

negative_ttl 5 minutes; negative_dns_ttl 5 minutes

Normally, you do not need to change these values. If you have a dial-up connection, however, the Internet may, at times, not be accessible. Squid makes a note of the failed requests then refuses to issue new ones, although the Internet connection has been reestablished. In a case such as this, change the *minutes* to *seconds* then, after clicking *Reload* in the browser, the dial-up process should be reengaged after a few seconds.

never_direct allow *acl_name*

To prevent Squid from taking requests directly from the Internet, use the above command to force connection to another proxy. This must have previously been entered in *cache_peer*. If *all* is specified as the *acl_name*, force all requests to be forwarded directly to the *parent*. This might be necessary, for example, if you are using a provider that strictly stipulates the use of its proxies or denies its firewall direct Internet access.

29.4.2 Options for Access Controls

Squid provides a detailed system for controlling the access to the proxy. By implementing ACLs, it can be configured easily and comprehensively. This involves lists with rules that are processed sequentially. ACLs must be defined before they can be used. Some default ACLs, such as *all* and *localhost*, already exist. However, the mere definition of an ACL does not mean that it is actually applied. This only happens in conjunction with *http_access* rules.

```
acl <acl_name> <type> <data>
```

An ACL requires at least three specifications to define it. The name *<acl_name>* can be chosen arbitrarily. For *<type>*, select from a variety of different options, which can be found in the *ACCESS CONTROLS* section in the */etc/squid/squid.conf* file. The specification for *<data>* depends on the individual ACL type and can also be read from a file, for example, via hostnames, IP addresses, or URLs. The following are some simple examples:

```
acl mysurfers srcdomain .my-domain.com  
acl teachers src 192.168.1.0/255.255.255.0  
acl students src 192.168.7.0-192.168.9.0/255.255.255.0  
acl lunch time MTWHF 12:00-15:00
```

```
http_access allow <acl_name>
```

http_access defines who is allowed to use the proxy and who can access what on the Internet. For this, ACLs must be given. *localhost* and *all* have already been defined above, which can deny or allow access via *deny* or *allow*. A list containing any number of *http_access* entries can be created, processed from top to bottom, and, depending on which occurs first, access is allowed or denied to the respective URL. The last entry should always be *http_access deny all*. In the following example, the *localhost* has free access to everything while all other hosts are denied access completely.

```
http_access allow localhost  
http_access deny all
```

In another example using these rules, the group *teachers* always has access to the Internet. The group *students* only gets access Monday to Friday during lunch time.

```
http_access deny localhost  
http_access allow teachers  
http_access allow students lunch time  
http_access deny all
```

The list with the *http_access* entries should only be entered, for the sake of readability, at the designated position in the */etc/squid/squid.conf* file. That is, between the text

```
# INSERT YOUR OWN RULE(S) HERE TO ALLOW ACCESS FROM YOUR  
# CLIENTS
```

and the last

```
http_access deny all
```

redirect_program /usr/bin/squidGuard

With this option, specify a redirector such as squidGuard, which allows blocking unwanted URLs. Internet access can be individually controlled for various user groups with the help of proxy authentication and the appropriate ACLs. squidGuard is a separate package that can be installed and configured.

auth_param basic program /usr/sbin/pam_auth

If users must be authenticated on the proxy, set a corresponding program, such as pam_auth. When accessing pam_auth for the first time, the user sees a login window in which to enter the username and password. In addition, an ACL is still required, so only clients with a valid login can use the Internet:

```
acl password proxy_auth REQUIRED  
  
http_access allow password  
http_access deny all
```

The *REQUIRED* after *proxy_auth* can be replaced with a list of permitted usernames or with the path to such a list.

ident_lookup_access allow <acl_name>

With this, have an ident request run for all ACL-defined clients to find each user's identity. If you apply *all* to the <*acl_name*>, this is valid for all clients. Also, an ident daemon must be running on all clients. For Linux, install the pidnrd package for this purpose. For Microsoft Windows, free software is available for download from the Internet. To ensure that only clients with a successful ident lookup are permitted, define a corresponding ACL here:

```
acl identhosts ident REQUIRED  
  
http_access allow identhosts  
http_access deny all
```

Here, too, replace *REQUIRED* with a list of permitted usernames. Using *ident* can slow down the access time quite a bit, because ident lookups are repeated for each request.

29.5 Configuring a Transparent Proxy

The usual way of working with proxy servers is the following: the Web browser sends requests to a certain port in the proxy server and the proxy provides these required objects, whether they are in its cache or not. When working in a network, several situations may arise:

- For security reasons, it is recommended that all clients use a proxy to surf the Internet.
- All clients must use a proxy, regardless of whether they are aware of it.
- The proxy in a network is moved, but the existing clients should retain their old configuration.

In all these cases, a transparent proxy may be used. The principle is very easy: the proxy intercepts and answers the requests of the Web browser, so the Web browser receives the requested pages without knowing from where they are coming. As the name indicates, the entire process is done transparently.

29.5.1 Configuration Options in /etc/squid/squid.conf

The options to activate in the `/etc/squid/squid.conf` file to get the transparent proxy up and running are:

- `httpd_accel_host virtual`
- `httpd_accel_port 80`

The port number where the actual HTTP server is located

- `httpd_accel_with_proxy on`
- `httpd_accel_uses_host_header on`

29.5.2 Firewall Configuration with SuSEfirewall2

Now redirect all incoming requests via the firewall with help of a port forwarding rule to the Squid port. To do this, use the enclosed tool SuSEfirewall2, described in [Section “Configuring with YaST” \(page 106\)](#). Its configuration file can be found in `/etc/sysconfig/SuSEfirewall2`. The configuration file consists of well-documented entries. To set a transparent proxy, you must configure several firewall options:

- Device pointing to the Internet: `FW_DEV_EXT="eth1"`
- Device pointing to the network: `FW_DEV_INT="eth0"`

Define ports and services (see `/etc/services`) on the firewall that are accessed from untrusted (external) networks such as the Internet. In this example, only Web services are offered to the outside:

```
FW_SERVICES_EXT_TCP="www"
```

Define ports or services (see `/etc/services`) on the firewall that are accessed from the secure (internal) network, both via TCP and UDP:

```
FW_SERVICES_INT_TCP="domain www 3128"  
FW_SERVICES_INT_UDP="domain"
```

This allows accessing Web services and Squid (whose default port is 3128). The service “domain” stands for DNS (domain name service). This service is commonly used. Otherwise, simply take it out of the above entries and set the following option to no:

```
FW_SERVICE_DNS="yes"
```

The most important option is option number 15:

Example 29.1 Firewall Configuration: Option 15

```
# 15.)
# Which accesses to services should be redirected to a local port
# on the firewall machine?
#
# This can be used to force all internal users to surf via your
# Squid proxy, or transparently redirect incoming Web traffic to
# a secure Web server.
#
# Choice: leave empty or use the following explained syntax of
# redirecting rules, separated with spaces.
# A redirecting rule consists of 1) source IP/net,
# 2) destination IP/net, 3) original destination port and
# 4) local port to redirect the traffic to, separated by a colon,
# e.g. "10.0.0.0/8,0/0,80,3128 0/0,172.20.1.1,80,8080"
```

The comments above show the syntax to follow. First, enter the IP address and the netmask of the internal networks accessing the proxy firewall. Second, enter the IP address and the netmask to which these clients send their requests. In the case of Web browsers, specify the networks 0/0, a wild card that means “to everywhere.” After that, enter the original port to which these requests are sent and, finally, the port to which all these requests are redirected. Because Squid supports protocols other than HTTP, redirect requests from other ports to the proxy, such as FTP (port 21), HTTPS, or SSL (port 443). In this example, Web services (port 80) are redirected to the proxy port (port 3128). If there are more networks or services to add, they must be separated by a blank space in the respective entry.

```
FW_REDIRECT_TCP="192.168.0.0/16,0/0,80,3128 192.168.0.0/16,0/0,21,3128"
FW_REDIRECT_UDP="192.168.0.0/16,0/0,80,3128 192.168.0.0/16,0/0,21,3128"
```

To start the firewall and the new configuration with it, change an entry in the /etc/sysconfig/SuSEfirewall2 file. The entry START_FW must be set to "yes".

Start Squid as shown in [Section 29.3, “Starting Squid”](#) (page 525). To check if everything is working properly, check the Squid logs in /var/log/squid/access.log. To verify that all ports are correctly configured, perform a port scan on the machine from any computer outside your network. Only the Web services (port 80) should be open. To scan the ports with nmap, the command syntax is nmap -O IP_address.

29.6 cachemgr.cgi

The cache manager (cachemgr.cgi) is a CGI utility for displaying statistics about the memory usage of a running Squid process. It is also a more convenient way to manage the cache and view statistics without logging the server.

29.6.1 Setup

First, a running Web server on your system is required. Configure Apache as described in [Chapter 26, The Apache HTTP Server](#) (page 447). To check if Apache is already running, as root enter the command `rcapache status`. If a message like this appears:

```
Checking for service httpd: OK
Server uptime: 1 day 18 hours 29 minutes 39 seconds
```

Apache is running on the machine. Otherwise, enter `rcapache start` to start Apache with the SUSE Linux default settings. The last step to set it up is to copy the file `cachemgr.cgi` to the Apache directory `cgi-bin`:

```
cp /usr/share/doc/packages/squid/scripts/cachemgr.cgi /srv/www/cgi-bin/
```

29.6.2 Cache Manager ACLs in `/etc/squid/squid.conf`

There are some default settings in the original file required for the cache manager. First, two ACLs are defined then `http_access` options use these ACLs to grant access from the CGI script to Squid. The first ACL is the most important, because the cache manager tries to communicate with Squid over the `cache_object` protocol.

```
acl manager proto cache_object
acl localhost src 127.0.0.1/255.255.255.255
```

The following rules give Apache the access rights to Squid:

```
http_access allow manager localhost
http_access deny manager
```

These rules assume that the Web server and Squid are running on the same machine. If the communication between the cache manager and Squid originates at the Web

server on another computer, include an extra ACL as in [Example 29.2, “Access Rules”](#) (page 537).

Example 29.2 Access Rules

```
acl manager proto cache_object  
acl localhost src 127.0.0.1/255.255.255.255  
acl webserver src 192.168.1.7/255.255.255.255 # webserver IP
```

Then add the rules in [Example 29.3, “Access Rules”](#) (page 537) to permit access from the Web server.

Example 29.3 Access Rules

```
http_access allow manager localhost  
http_access allow manager webserver  
http_access deny manager
```

Configure a password for the manager for access to more options, like closing the cache remotely or viewing more information about the cache. For this, configure the entry `cachemgr_passwd` with a password for the manager and the list of options to view. This list appears as a part of the entry comments in `/etc/squid/squid.conf`.

Restart Squid every time the configuration file is changed. Do this easily with `rcsquid reload`.

29.6.3 Viewing the Statistics

Go to the corresponding Web site—<http://webserver.example.org/cgi-bin/cachemgr.cgi>. Press *continue* and browse through the different statistics. More details for each entry shown by the cache manager is in the Squid FAQ at <http://www.squid-cache.org/Doc/FAQ/FAQ-9.html>.

29.7 squidGuard

This section is not intended to explain an extensive configuration of squidGuard, only to introduce it and give some advice for using it. For more in-depth configuration issues, refer to the squidGuard Web site at <http://www.squidguard.org>.

squidGuard is a free (GPL), flexible, and fast filter, redirector, and access controller plug-in for Squid. It lets you define multiple access rules with different restrictions for different user groups on a Squid cache. squidGuard uses Squid's standard redirector interface. squidGuard can do the following:

- Limit the Web access for some users to a list of accepted or well-known Web servers or URLs.
- Block access to some listed or blacklisted Web servers or URLs for some users.
- Block access to URLs matching a list of regular expressions or words for some users.
- Redirect blocked URLs to an “intelligent” CGI-based information page.
- Redirect unregistered users to a registration form.
- Redirect banners to an empty GIF.
- Use different access rules based on time of day, day of the week, date, etc.
- Use different rules for different user groups.

squidGuard and Squid cannot be used to:

- Edit, filter, or censor text inside documents.
- Edit, filter, or censor HTML-embedded script languages, such as JavaScript or VBscript.

Before it can be used, install `squidGuard`. Provide a minimal configuration file as `/etc/squidguard.conf`. Find configuration examples in <http://www.squidguard.org/config/>. Experiment later with more complicated configuration settings.

Next, create a dummy “access denied” page or a more or less complex CGI page to redirect Squid if the client requests a blacklisted Web site. Using Apache is strongly recommended.

Now, configure Squid to use squidGuard. Use the following entry in the `/etc/squid/squid.conf` file:

```
redirect_program /usr/bin/squidGuard
```

Another option called `redirect_children` configures the number of “redirect” (in this case `squidGuard`) processes running on the machine. `squidGuard` is fast enough to handle many requests: on a 500 MHz Pentium with 5,900 domains and 7,880 URLs (totalling 13,780), 100,000 requests can be processed within 10 seconds. Therefore, it is not recommended to set more than four processes, because the allocation of these processes would consume an excessive amount of memory

```
redirect_children 4
```

Last, have Squid load the new configuration by running `rcsquid reload`. Now, test your settings with a browser.

29.8 Cache Report Generation with Calamaris

Calamaris is a Perl script used to generate reports of cache activity in ASCII or HTML format. It works with native Squid access log files. The Calamaris home page is located at <http://Calamaris.Cord.de/>. The program is quite easy to use.

Log in as root then enter `cat access.log.files | calamaris options > reportfile`. It is important when piping more than one log file that the log files are chronologically ordered with older files first. These are some options of the program:

- a output all available reports
- w output as HTML report
- l include a message or logo in report header

More information about the various options can be found in the program's manual page with `man calamaris`.

A typical example is:

```
cat access.log.2 access.log.1 access.log | calamaris -a -w \  
> /usr/local/httpd/htdocs/Squid/squidreport.html
```

This puts the report in the directory of the Web server. Apache is required to view the reports.

Another powerful cache report generator tool is SARG (Squid Analysis Report Generator). More information about this is available at: <http://sarg.sourceforge.net/>.

29.9 For More Information

Visit the home page of Squid at <http://www.squid-cache.org/>. Here, find the “Squid User Guide” and a very extensive collection of FAQs on Squid.

Following the installation, a small howto about transparent proxies is available in `howtoen.html` as `/usr/share/doc/howto/en/txt/TransparentProxy.gz`. In addition, mailing lists are available for Squid at squid-users@squid-cache.org. The archive for this is located at <http://www.squid-cache.org/mail-archive/squid-users/>.

Part V. Mobility

30

Mobile Computing with Linux

Mobile computing is mostly associated with laptops, PDAs, and cellular phones and the data exchange between them. Mobile hardware components, such as external hard disks, flash drives, or digital cameras, can be connected to laptops or desktop systems. A number of software components are involved in mobile computing scenarios and some applications are tailor-made for mobile use.

30.1 Laptops

The hardware of laptops differs from that of a normal desktop system. This is because criteria like exchangeability, occupied space, and power consumption are relevant properties. The manufacturers of mobile hardware have developed the PCMCIA (Personal Computer Memory Card International Association) standard. This standard covers memory cards, network interface cards, ISDN and modem cards, and external hard disks. How the support for such hardware is implemented in Linux, what needs to be taken into account during configuration, what software is available for the control of PCMCIA, and how to troubleshoot any possible problems is described in [Chapter 31, PCMCIA](#) (page 555).

30.1.1 Power Conservation

The inclusion of energy-optimized system components when manufacturing laptops contributes to their suitability for use without access to the electrical power grid. Their contribution towards conservation of power is at least as important as that of the operating system. SUSE Linux supports various methods that influence the power consumption of a laptop and have varying effects on the operating time under battery power. The following list is in descending order of contribution towards power conservation:

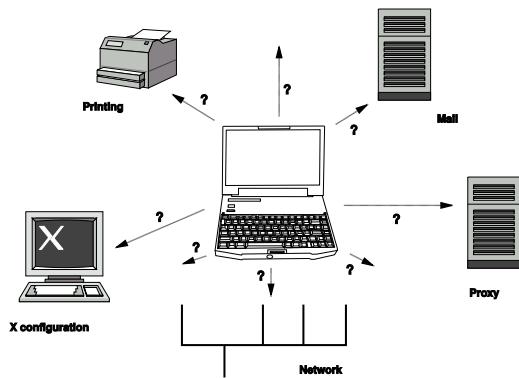
- Throttling the CPU speed
- Switching off the display illumination during pauses
- Manually adjusting the display illumination
- Disconnecting unused, hotplug-enabled accessories (USB CD-ROM, external mouse, unused PCMCIA cards, etc.)
- Spinning down the hard disk when idling

Detailed background information about power management in SUSE Linux and about operating the YaST power management module is provided in [Chapter 33, Power Management](#) (page 577).

30.1.2 Integration in Changing Operating Environments

Your system needs to adapt to changing operating environments when used for mobile computing. A lot of services depend on the environment and the underlying clients must be reconfigured. SUSE Linux handles this task for you.

Figure 30.1 Integrating a Laptop in a Network



The services affected in the case of a laptop commuting back and forth between a small home network and an office network are:

Network

This includes IP address assignment, name resolution, Internet connectivity, and connectivity to other networks.

Printing

A current database of available printers and an available print server must be present, depending on the network.

E-Mail and Proxies

As with printing, the list of the corresponding servers must be current.

X

If your laptop is temporarily connected to a beamer or an external monitor, the different display configurations must be available.

SUSE Linux offers several ways of integrating a laptop into existing operating environments:

SCPM

SCPM (system configuration profile management) allows storage of arbitrary configuration states of a system into a kind of “snapshot” called a *profile*. Profiles can be created for different situations. They are useful when a system is operated in changing environments (home network, office network). It is always possible to switch between profiles. Find information about SCPM in [Chapter 32, System Configuration Profile Management](#) (page 563). You can use the kicker applet Profile Chooser in KDE to switch between profiles. The application requires the root password before switching.

NetworkManager

NetworkManager is especially tailored for mobile networking on laptops. It provides a means to easily and automatically switch between network environments or different types of networks, such as wireless LAN and ethernet. NetworkManager supports WEP and WPA-PSK encryption in wireless LANs, dial-up connections (with smpppd). Both desktop environments (GNOME and KDE) include a front-end to NetworkManager.

Table 30.1 Use Cases for NetworkManager

My computer...	Use NetworkManager
is a laptop	Yes
is sometimes attached to different networks	Yes
provides network services (such as DNS or DHCP)	No
only uses a static IP address	No

Use the YaST tools to configure networking whenever NetworkManager should not handle network configuration.

For more information about NetworkManager and its applets refer to [Section 18.5, “Managing Network Connections with NetworkManager”](#) (page 345).

SLP

The service location protocol (SLP) simplifies the connection of a laptop to an existing network. Without SLP, the administrator of a laptop usually requires detailed knowledge of the services available in a network. SLP broadcasts the availability of a certain type of service to all clients in a local network. Applications that support SLP can process the information dispatched by SLP and be configured automatically. SLP can even be used for the installation of a system, sparing the effort of searching for a suitable installation source. Find detailed information about SLP in [Chapter 19, *SLP Services in the Network*](#) (page 363).

The emphasis of SCPM lies on enabling and maintaining reproducible system conditions. SLP makes configuration of a networked computer a lot easier by automating much of it.

30.1.3 Software Options

There are various special task areas in mobile use that are covered by dedicated software: system monitoring (especially the battery charge), data synchronization, and wireless communication with peripherals and the Internet. The following sections cover the most important applications that SUSE Linux provides for each task.

System Monitoring

Two KDE system monitoring tools are provided by SUSE Linux:

KPowersave

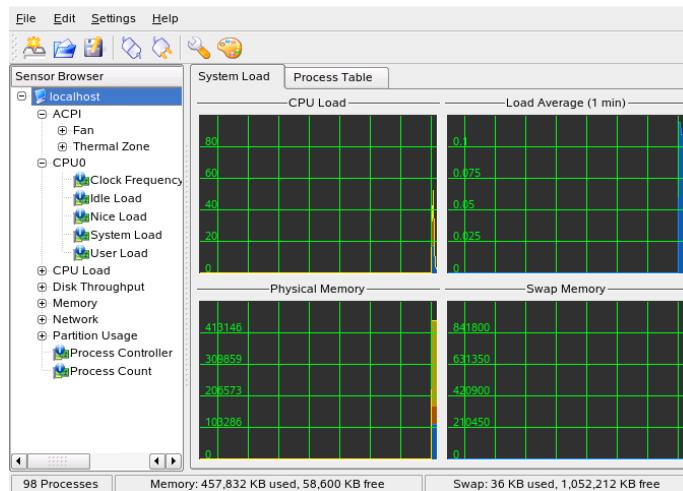
KPowersave is an applet that displays the state of the rechargeable battery in the control panel. The icon adjusts to represent the type of power supply. When working on AC power, a small plug icon is displayed. When working on batteries, the icon changes to a battery. The corresponding menu opens the YaST module for power management after requesting the root password. This allows setting the behavior of the system under different types of power supply. Find information about power management and about the corresponding YaST module in [Chapter 33, *Power Management*](#) (page 577).

KSysguard

KSysguard is an independent application that gathers all measurable parameters of the system into one monitoring environment. KSysguard has monitors for ACPI (battery status), CPU load, network, partitioning, and memory usage. It can also

watch and display all system processes. The presentation and filtering of the collected data can be customized. It is possible to monitor different system parameters in various data pages or collect the data of various machines in parallel over the network. KSysguard can also run as a daemon on machines without a KDE environment. Find more information about this program in its integrated help function or in the SUSE help pages.

Figure 30.2 Monitoring the Battery State with KSysguard



In the GNOME desktop, use the panel applet GNOME ACPI and the application System Monitor.

Synchronizing Data

When switching between working on a mobile machine disconnected from the network and working at a networked workstation in an office, it is necessary to keep processed data synchronized across all instances. This could include e-mail folders, directories, and individual files that need to be present for work on the road as well as at the office. The solution in both cases is as follows:

Synchronizing E-Mail

Use an IMAP account for storing your e-mails in the office network. Then access the e-mails from the workstation using any disconnected IMAP-enabled e-mail client, like Mozilla Thunderbird Mail, Evolution, or KMail as described in *Applications*. The e-mail client must be configured so that the same folder is always accessed for Sent messages. This ensures that all messages are available along with their status information after the synchronization process has completed. Use an SMTP server implemented in the mail client for sending messages instead of the systemwide MTA postfix or sendmail to receive reliable feedback about unsent mail.

Synchronizing Files and Directories

There are several utilities suitable for synchronizing data between a laptop and a workstation. For detailed information, refer to [Chapter 27, File Synchronization](#) (page 487).

Wireless Communication

As well as connecting to a home or office network with a cable, a laptop can also wirelessly connect to other computers, peripherals, cellular phones, or PDAs. Linux supports three types of wireless communication:

WLAN

With the largest range of these wireless technologies, WLAN is the only one suitable for the operation of large and sometimes even spatially disjointed networks. Single machines can connect with each other to form an independent wireless network or access the Internet. Devices called access points act as base stations for WLAN-enabled devices and act as intermediaries for access to the Internet. A mobile user can switch among access points depending on location and which access point is offering the best connection. Like in cellular telephony, a large network is available to WLAN users without binding them to a specific location for accessing it. Find details about WLAN in [Section 34.1, “Wireless LAN”](#) (page 603).

Bluetooth

Bluetooth has the broadest application spectrum of all wireless technologies. It can be used for communication between computers (laptops) and PDAs or cellular phones, as can IrDA. It can also be used to connect various computers within visible range. Bluetooth is also used to connect wireless system components, like a keyboard or mouse. The range of this technology is, however, not sufficient to connect remote systems to a network. WLAN is the technology of choice for communicating

through physical obstacles like walls. Find more information about Bluetooth, its applications, and configuration in [Section 34.2, “Bluetooth”](#) (page 614).

IrDA

IrDA is the wireless technology with the shortest range. Both communication parties must be within viewing distance of each other. Obstacles like walls cannot be overcome. One possible application of IrDA is the transmission of a file from a laptop to a cellular phone. The short path from the laptop to the cellular phone is then covered using IrDA. The long range transport of the file to the recipient of the file is handled by the mobile network. Another application of IrDA is the wireless transmission of printing jobs in the office. Find more information about IrDA in [Section 34.3, “Infrared Data Transmission”](#) (page 625).

30.1.4 Data Security

Ideally, you protect data on your laptop against unauthorized access in multiple ways. Possible security measures can be taken in the following areas:

Protection against Theft

Always physically secure your system against theft whenever possible. Various securing tools, like chains, are available in retail stores.

Securing Data on the System

Important data should not only be encrypted during transmission, but also on the hard disk. This ensures its safety in case of theft. The creation of an encrypted partition with SUSE Linux is described in [Section 4.3, “Encrypting Partitions and Files”](#) (page 116).

IMPORTANT: Data Security and Suspend to Disk

Encrypted partitions are not unmounted during a suspend to disk event. Thus, all data on these partitions is available to any party who manages to steal the hardware and issue a resume of the hard disk.

Network Security

Any transfer of data should be secured, no matter how it takes place. Find general security issues regarding Linux and networks in [Section 4.5, “Security and Confidentiality”](#) (page 128). Security measures related to wireless networking are provided in [Chapter 34, *Wireless Communication*](#) (page 603).

30.2 Mobile Hardware

SUSE Linux supports the automatic detection of mobile storage devices over firewire (IEEE 1394) or USB. The term *mobile storage device* applies to any kind of firewire or USB hard disk, USB flash drive, or digital camera. These devices are automatically detected and configured via hotplug as soon as they are connected with the system over the corresponding interface. subfs and submount ensure that the devices are mounted to the corresponding locations in the file system. The user is completely spared the manual mounting and unmounting that was found in previous versions of SUSE Linux. A device can simply be disconnected as soon as no program accesses it.

External Hard Disks (USB and Firewire)

As soon as an external hard disk has been correctly recognized by the system, its icon appears in *My Computer* (KDE) or *Computer* (GNOME) in the list of mounted drives. Clicking the icon displays the contents of the drive. It is possible to create folders and files here and edit or delete them. To rename a hard disk from the name it had been given by the system, select the corresponding menu item from the menu that opens when the icon is right-clicked. This name change is limited to display in the file manager. The descriptor by which the device is mounted in `/media` remains unaffected by this.

USB Flash Drives

These devices are handled by the system just like external hard disks. It is similarly possible to rename the entries in the file manager.

Digital Cameras (USB and Firewire)

Digital cameras recognized by the system also appear as external drives in the overview of the file manager. KDE allows reading and accessing the pictures at the URL `camera:/`. The images can then be processed using Digikam or The GIMP. When using GNOME, Nautilus displays the pictures in their own folder. A simple image processing and management utility is f-spot. Advanced photo processing is done with The GIMP. For more details on digital cameras and image management, refer to Chapter *Digital Cameras and Linux* (\uparrow Applications).

30.3 Cellular Phones and PDAs

A desktop system or a laptop can communicate with a cellular phone via Bluetooth or IrDA. Some models support both protocols and some only one of the two. The usage areas for the two protocols and the corresponding extended documentation has already been mentioned in [Section “Wireless Communication”](#) (page 549). The configuration of these protocols on the cellular phones themselves is described in their manuals. The configuration of the Linux side is described in [Section 34.2, “Bluetooth”](#) (page 614) and [Section 34.3, “Infrared Data Transmission”](#) (page 625).

The support for synchronizing with handheld devices manufactured by Palm, Inc., is already built into Evolution and Kontact. Initial connection with the device is, in both cases, easily performed with the assistance of a wizard. Once the support for Palm Pilots is configured, it is necessary to determine which type of data should be synchronized (addresses, appointments, etc.). Both groupware applications are described in *Applications*.

The program KPilot as integrated in Kontact is also available as an independent utility. It is described in *Applications*. The program KitchenSync is also available for synchronizing address data.

30.4 For More Information

The central point of reference for all questions regarding mobile devices and Linux is <http://tuxmobil.org/>. Various sections of that Web site deal with the hardware and software aspects of laptops, PDAs, cellular phones, and other mobile hardware.

A similar approach to that of <http://tuxmobil.org/> is made by <http://www.linux-on-laptops.com/>. Information about laptops and handhelds can be found here.

SUSE maintains a mailing list in German dedicated to the subject of laptops. See <http://lists.suse.com/archive/suse-laptop/>. On this list, users and developers discuss all aspects of mobile computing with SUSE Linux. Postings in English are answered, but the majority of the archived information is only available in German.

In the case of problems with power management with SUSE Linux on laptops, it is advisable to read the file `README` in `/usr/share/doc/packages/powersave`. This directory often contains last minute feedback by testers and developers, so provides valuable hints for the solution of problems.

31

PCMCIA

PCMCIA is often used to refer the hardware itself, although it originates from the organization that standardized all possible types of PC cards, the *PC Memory Card International Association*. In the beginning, PCMCIA only included PC cards (using a 16-bit bus like ISA cards), but later on CardBus cards (using a 32-bit bus) were included. A wide range of PCMCIA hardware is supported in Linux. Linux additionally includes tools for managing PCMCIA.

PCMCIA cards are mainly used in mobile computing for different purposes. Examples include:

- Ethernet and wireless LAN adapters
- Bluetooth cards
- Memory cards (Flash, SRAM, and others)
- Memory card adapters (SD, MMC, SmartMedia, CompactFlash, MemoryStick)
- Modems

Most of the card management is silently handled by udev and hotplug. When user interaction is required, you use `pccardctl` command. For PCMCIA background information, refer to [Section 31.2, “PCMCIA in Detail”](#) (page 556). For details on `pccardctl`, refer to [Section 31.1, “Controlling PCMCIA Cards Using pccardctl”](#) (page 556).

31.1 Controlling PCMCIA Cards Using pccardctl

Card management is normally handled by udev and hotplug without requiring any user interaction at all. pccardctl offers manual control of the card in case the automated process does not work flawlessly.

The following is a list of the most important pccardctl commands. All commands must be executed as root:

`pccardctl insert`

If the card has not been detected automatically, notify the client drivers that the card has just been inserted.

`pccardctl eject`

Eject the card manually and notify the client drivers that it will be ejected. Cut power to the socket. This option is especially useful if you noticed problems with suspend and resume as described in [Section 31.3.2, “General Suspend Issues with PCMCIA”](#) (page 561).

`pccardctl suspend`

Shut down and disable power for a socket, but do not eject the card (unbind the appropriate modules).

`pccardctl resume`

After a `pccardctl resume`, bring up power for the socket and restore the configuration from before the suspend event.

For further information, refer to the manual page of `pccardctl`.

31.2 PCMCIA in Detail

The following sections outlines what happens in your Linux system when a PCMCIA device is plugged into your machine. Components interact with each other and many requirements need to be met to support a PCMCIA device.

The following is a very rough outline of the PCMCIA initialization process in Linux:

1. The PCMCIA bridge (or socket) must be set up properly as described in [Section 31.2.1, “Bridge Initialization”](#) (page 557). Prerequisites are:
 - an appropriate driver for the bridge
 - additional I/O and memory ranges for PC cards
2. After the bridge is properly set up, the bridge driver detects the presence of a card and triggers its initialization as described in [Section 31.2.2, “Card Initialization”](#) (page 558):
 - a. Determine the card type.
 - b. Supply the proper voltage.
 - c. Assign I/O and memory ranges and IRQ lines to the card.
 - d. Trigger the card or device initialization by binding the appropriate card driver.
 - e. For some cards, the Card Information Structure (CIS) needs to be uploaded.
3. Finally, the interface itself is set up and ready for use. See [Section 31.2.3, “Interface Setup”](#) (page 559) for details on this.

31.2.1 Bridge Initialization

Most PCMCIA bridges are PCI devices and are treated as such. The bridge initialization process can be summarized as follows:

1. Hotplug creates a PCI event.
2. udev calls `/sbin/hwup` to load the driver. `/sbin/hwup` checks `/etc/sysconfig/hardware` for an existing device configuration. If an appropriate configuration is found, that configuration is used. Otherwise `/sbin/hwup` calls `modprobe` with the `modalias` string provided by the kernel to load the driver module.
3. New hotplug events are sent (one per PCMCIA socket).
4. The following steps are omitted if only CardBus cards are used:

- a. The `pcmcia_socket` events trigger udev to call `/sbin/hwup` and load the `pcmcia` kernel module.
- b. All I/O and memory ranges specified in `/etc/pcmcia/config.opts` are added to the socket.
- c. The card services in the kernel check these ranges. If the memory ranges in `/etc/pcmcia/config.opts` are wrong, this step may crash your machine. See [Section 31.3.1, “Machine Crashes on PCMCIA”](#) (page 559) for information about how to debug and fix this issue.

After these steps have been successfully completed, the bridge is fully initialized. After this, the card itself is initialized as described in the following section.

31.2.2 Card Initialization

The events caused by plugging in a PCMCIA card can be summarized as follows:

1. A hotplug event occurs. For PC cards, this is a `pcmcia` event. For CardBus cards, this is a `pci` event.
2. For any events, udev calls `/sbin/hwup` to load a driver module. The module name is either specified in a `hwcfg*` file under `/etc/sysconfig/hardware` or via `modprobe` `modalias`.
3. If needed, device initialization triggers a firmware hotplug event. This searches for firmware and loads it.
4. The device driver registers the interfaces.

After these steps have been completed, the system proceeds with interface setup as described in the next section.

If your card is a PC card, you might need some of the following parameters in `/etc/sysconfig/pcmcia` to get it fully supported and working flawlessly:

`PCMCIA_LOAD_CIS`

A PC card's firmware is referred to as *CIS* (Card Information Structure). It provides additional implementation details of the card. `hwup` checks the integrity of the card's built-in CIS and tries to load another CIS from disk if the card's CIS proves

to be defective. The default setting is `yes`. To disable CIS loading from disk, set this variable to `no`.

`PCMCIA_ALLOW_FUNC_MATCH`

Linux device drivers contain a device ID table that tells drivers which devices to handle. This means that only those devices whose IDs are known to the kernel are supported. To support those cards whose ID is not listed, you can use function matching. This means that the driver is not selected by ID, but by the function of the card (such as a network card) and would be responsible for any PC card inserted with that function (such as network cards). The default setting is `yes`. To disable function matching, set this variable to `no`.

`PCMCIA_COLDPLUG_REINSERT`

Cards that have been inserted before booting sometimes fail to be detected. To prevent that, cause a soft eject and a soft insert of the card by setting `PCMCIA_COLDPLUG_REINSERT` to `yes`. The default setting is `no`.

31.2.3 Interface Setup

Depending on the card type, different interfaces are registered after initialization has been successfully completed. Interface registration is handled by udev's hotplug. For details on udev and hotplug, refer to [Chapter 12, Dynamic Kernel Device Management with udev](#) (page 251).

31.3 Troubleshooting

The following is a list of the most prominent issues that are occasionally encountered with PCMCIA. More information about this is available in the PCMCIA README (`/usr/share/doc/packages/pcmciautils/README.SuSE`).

31.3.1 Machine Crashes on PCMCIA

Your machine crashes when PCMCIA is started on boot. To find out what caused your machine to crash, set it up manually as described below. In carefully setting up PCMCIA manually, you can clearly identify the step or component that crashed your machine.

Once the culprit has been identified, you can circumvent the problematic step or component.

To manually set up PCMCIA, proceed as follows:

- 1 Prevent PCMCIA from being started on system boot and enable SysRq for easier debugging by appending the following options to the boot prompt:

```
init=3 pcmcia=off sysrq=1
```

For more information about SysRq, refer to `/usr/src/linux/Documentation/sysrq.txt`.

- 2 Boot the system into a text-based environment and log in as root.

- 3 Add the appropriate PCMCIA modules to the kernel:

```
/sbin/modprobe yenta_socket  
/sbin/modprobe pcmcia
```

- 4 Start the PCMCIA socket:

```
/sbin/pcmcia-socket-startupN
```

Replace *N* with the number of the socket. Repeat this step for each socket.

- 5 If the previous step crashed your machine, this might have been caused by wrong I/O or memory ranges specified in `/etc/pcmcia/config.opts`. To prevent this, do one of the following:

- Exclude ranges in `/etc/pcmcia/config.opts` and retry the socket setup.
- Add the ranges manually as described below.

After you successfully added the appropriate ranges manually, set them permanently by including them in `/etc/pcmcia/config.opts`.

- 6 After the socket setup has been successfully completed, card initialization and interface setup work as described in [Section 31.2.2, “Card Initialization”](#) (page 558) and [Section 31.2.3, “Interface Setup”](#) (page 559).

To manually add I/O ranges, proceed as follows (for each socket):

- 1 Change into the directory that holds the range configurations (in this case, `pcmcia_socket0`, adapt for other socket numbers):

```
cd /sys/class/pcmcia_socket/pcmcia_socket0
```

- 2 Execute the following command:

```
echo begin - end > available_resources_io
```

Replace `begin` and `end` with the addresses where the new range should start and end. The correct values can only be determined by trial and error.

Manually adding the following ranges:

```
echo 0x800 - 0x8ff > available_resources_io
echo 0xc00 - 0cff > available_resources_io
```

equals the following line from `/etc/pcmcia/config.opts`:

```
include port 0x800-0x8ff, port 0xc00 0cff
```

The same procedure applies for the memory ranges under `available_resources_mem`.

IMPORTANT: Identifying Faulty Default Settings

If you find a faulty range in the default configuration file (`/etc/pcmcia/config.opts`) shipped with this product, file a bug against it in <http://bugzilla.novell.com>, so that developers can look into this issue.

31.3.2 General Suspend Issues with PCMCIA

Whenever suspending your system (suspend to disk, suspend to RAM, or standby), do not plug or unplug any hardware items while the system is in suspend mode. Otherwise, the system might not resume properly.

To automatically eject PCMCIA cards on suspend, proceed as follows:

- 1 Log in as root.
- 2 Open the file `/etc/powersave/sleep`.

3 Set the following variables:

```
SUSPEND2DISK_EJECT_PCMCIA="yes"  
SUSPEND2RAM_EJECT_PCMCIA="yes"  
STANDBY_EJECT_PCMCIA="yes"
```

4 Save the file to apply your settings.

If additional modules need to be ejected on suspend, proceed as above and add the module names to the following variables:

```
UNLOAD_MODULES_BEFORE_SUSPEND2DISK=""  
UNLOAD_MODULES_BEFORE_SUSPEND2RAM=""  
UNLOAD_MODULES_BEFORE_STANDBY=""
```

For general information about the powersave daemon, refer to [Section 33.5, “The powersave Package”](#) (page 589).

31.3.3 For More Information

Find the latest up-to-date information about PCMCIA in `/usr/share/doc/packages/pcmciautils/README.SuSE`. For a comprehensive overview of PCMCIA hardware and its fields of use, turn to the official PCMCIA Web site (<http://www.pcmcia.org/pccard.htm>). To check whether a certain card or device is generally supported by Linux, refer to the *Linux PCMCIA/CF/CardBus Card Survey* at http://tuxmobil.org/pcmcia_linux.html.

System Configuration Profile Management

With the help of SCPM (system configuration profile management), adapt the configuration of your computer to different operating environments or hardware configurations. SCPM manages a set of system profiles for the different scenarios. SCPM enables easy switching between system profiles, eliminating the need for manually reconfiguring the system.

Some situations require a modified system configuration. This would mostly be the case for mobile computers that are operated in varying locations. If a desktop system should be operated temporarily using other hardware components than usual, SCPM comes in handy. Restoring the original system configuration should be easy and the modification of the system configuration can be reproduced. With SCPM, any part of the system configuration can be kept in a customized profile.

The main field of application of SCPM is network configuration on laptops. Different network configurations often require different settings of other services, such as e-mail or proxies. Then other elements follow, like different printers at home and at the office, a customized X server configuration for the multimedia projector at conferences, special power-saving settings for the road, or a different time zone at an overseas subsidiary.

32.1 Terminology

The following are some terms used in SCPM documentation and in the YaST module.

system configuration

Refers to the complete configuration of the computer. It covers all fundamental settings, such as the use of hard disk partitions, network settings, time zone selection, and keyboard mappings.

profile or system profile

A state that has been preserved and can be restored at any time.

active profile

Refers to the profile last selected. This does not mean that the current system configuration corresponds exactly to this profile, because the configuration can be modified at any time.

resource

An element that contributes to the system configuration. This can be a file or a softlink including metadata (like the user), permissions, or access time. This can also be a system service that runs in this profile, but is deactivated in another one.

resource group

Every resource belongs to a certain *resource group*. These groups contain all resources that logically belong together—most groups would contain both a service and its configuration files. It is very easy to assemble resources managed by SCPM because this does not require any knowledge of the configuration files of the desired service. SCPM ships with a selection of preconfigured resource groups that should be sufficient for most scenarios.

32.2 Setting up SCPM

The following sections introduce you to SCPM configuration by means of a real life example: a mobile computer that is run in several different networks. The major challenges faced in this scenario are:

- Varying network environments, like wireless LAN at home and an ethernet at work
- Different printer configuration at home and at work

To get SCPM up and running and have it manage your changing system configuration, proceed as follows:

- 1 Add the Profile Chooser applet to your panel and configure it to allow user switching as described in [Section 32.3.1, “Configuring the Profile Chooser Panel Applet”](#) (page 565).
- 2 Configure SCPM using the YaST Profile Manager module as described in [Section 32.3.2, “Configuring Basic SCPM Settings”](#) (page 566).
- 3 Create a profile for each of the different setups using SUMF (SCPM Unified Management Front-End) as described in [Section 32.3.3, “Creating a New Profile”](#) (page 568).
- 4 Switch to the profile appropriate for your current situation as described in [Section 32.3.4, “Switching Profiles”](#) (page 569).

If you prefer to control SCPM with its command line interface, refer to [Section 32.4, “Configuring SCPM Using the Command Line”](#) (page 572) for details.

32.3 Configuring SCPM Using a Graphical User Interface

The following sections introduce the graphical tools used for controlling your profile settings.

32.3.1 Configuring the Profile Chooser Panel Applet

Before you can use Profile Chooser to control your system configuration, configure it to be started automatically on login:

- In GNOME, right-click the panel and select Profile Chooser from the list of available applets.
- In KDE, select *System → Desktop Applet → Profile Chooser* to add Profile Chooser to your panel.

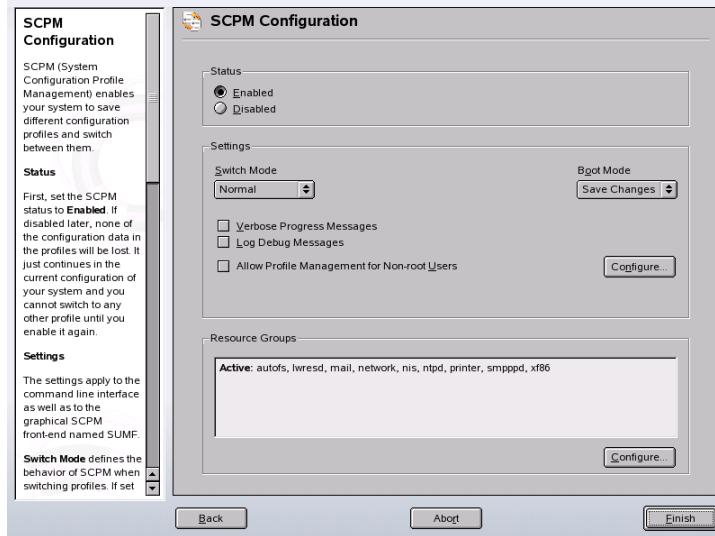
32.3.2 Configuring Basic SCPM Settings

Configure the basic behavior of SCPM through YaST.

- 1 Start YaST from the main menu and select the YaST Profile Manager.
- 2 In *System Configuration Profile Management*, click *Options* and select *Enabled*.
- 3 Determine how verbose SCPM should be by selecting any or both of *Verbose Progress Messages* and *Log Debug Messages*.
- 4 Determine the appropriate switch mode for your setup:
 - Should SCPM list any changed resource when switching to another profile and save these changes to the active profile? Select *Normal* or *Save Changes*.
 - Should SCPM drop any changed resource configuration when switching? Select *Drop Changes*.
- 5 Set the boot mode and determine whether changes to the current profile should be saved or whether they should be discarded with profile switching triggered at boot time.
- 6 Make sure that all resource groups you need are covered by the active selection, displayed in the *Resource Groups* section. If you need additional resource groups, adjust the resources with *Configure Resources*. For details, refer to [Section 32.3.6, “Configuring Resource Groups” \(page 570\)](#).

For the example scenario, you do not need to configure additional resources, because printer and network resources are included by default.

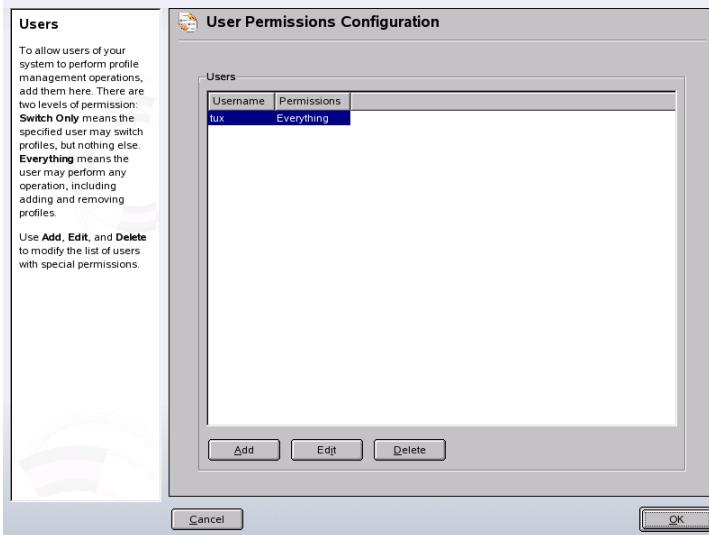
Figure 32.1 YaST: Basic SCPM Configuration



To allow users other than root to manage profiles, proceed as follows:

- 1 Start YaST from the main menu and select the YaST Profile Manager.
- 2 Check *Permit non-root Users to Manage Profiles*. See [Figure 32.2, “YaST: Configure SCPM Users”](#) (page 568).
- 3 Click *Configure Users*.
- 4 Click *Add* to add any user who should be able to manage profiles.
- 5 For each user, specify whether he should have switch permissions only or whether this user should be allowed to switch, modify, and create profiles.
- 6 Click *Accept* to apply your settings and close YaST.

Figure 32.2 YaST: Configure SCPM Users



32.3.3 Creating a New Profile

After you have enabled SCPM, you have a profile named `default` that contains your current system configuration. Create another profile that matches the requirements of the other setup.

To add a new profile based on the current system configuration, proceed as follows:

- 1 Right-click the Profile Chooser and select *Run Profile Manager (SUMF)*.
- 2 Select *Profiles* → *Add*.
- 3 Enter the name of the new profile and click *OK*.
- 4 Determine whether the new profile should be the active profile.

If you selected *Yes*, SCPM switches to the new profile immediately after it has been created.

For this example, do the following:

- 1 In your home setup, enable SCPM.
- 2 Rename the `default` profile to a more descriptive name by starting SUMF and selecting *Profiles* → *Edit* and entering the new name.
- 3 In your setup at work, start SUMF and create the profile for your system environment at work.

Once you have all profiles you need, you are ready to switch to them whenever a different system setup is required. Switching profiles is described in [Section 32.3.4, “Switching Profiles”](#) (page 569).

32.3.4 Switching Profiles

There are two ways to switch profiles. You can either select a new profile at boot or switch profiles in the running system.

To select a profile at boot, proceed as follows:

- 1 In the boot screen, press `F2` to enter the *Other Options* menu.
- 2 Press `F3` to access the list of profiles available.
- 3 Use the arrow keys to select the appropriate profile and hit `Enter`.

The system boots into the configuration selected.

To switch profiles in a running system, proceed as follows:

- 1 Make sure that you are allowed to switch profiles as a non-root user. If you are not allowed to do so, refer to [Section 32.3.2, “Configuring Basic SCPM Settings”](#) (page 566).
- 2 Left-click the Profile Chooser panel applet.
- 3 Select the profile you need in the menu that opens using the arrow keys and hit `Enter`.

SCPM runs a check for modified resources and prompts you for a confirmation of the switch. If changes have been made to the system configuration before the

switch, SCPM asks you to either keep them or discard them when switching to another profile.

32.3.5 Editing a Profile

To adjust existing profiles to a changed environment, for example, if you want to change the printer configuration of your home network, proceed as follows:

- 1 Switch to the profile to adjust as described in [Section 32.3.4, “Switching Profiles”](#) (page 569). In this example, you would choose the `home` profile.
- 2 Change the resources that need to be adjusted using the appropriate YaST module. In this example, run the YaST printer configuration.
- 3 After the configuration changes have been applied, SCPM asks whether these changes should be permanently applied to your active profile on the next profile switch.

TIP: Forcing a Profile Update

If you want to force an update of the active profile, click the profile in the profile selection menu of the Profile Chooser panel applet. This triggers a reload of your profile and you are asked whether to apply the configuration changes or discard them.

32.3.6 Configuring Resource Groups

SCPM comes with a set of predefined resource groups that are included in any profile by default. However, some scenarios require the inclusion of additional resources and resource groups.

To change the resource (group) configuration, proceed as follows:

- 1 Start YaST from the main menu and start the YaST Profile Manager module.
- 2 In the *System Configuration Profile Management* dialog, click *Configure Resources*.

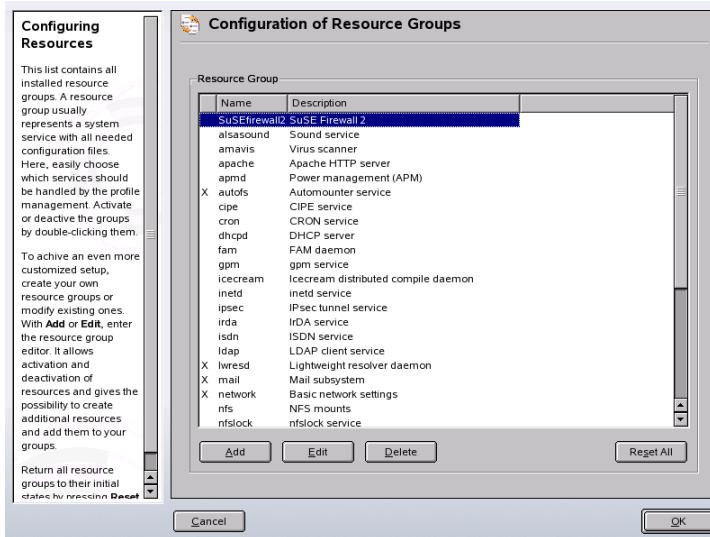
All resource groups available on your system are listed as shown in [Figure 32.3, “Configuring Resource Groups”](#) (page 571).

3 To add or edit a resource group:

- a** Set or edit *Resource Group* and *Description*.
- b** Enter the appropriate resources (resources, services, or both) and delete those that are not needed. To reset the status of the selected resources—discard any changes made to them and return to the initial configuration values—choose *Reset Group*.
- c** Click *OK* to leave the resource configuration.

4 Click *OK* to save your changes to the active profile.

Figure 32.3 Configuring Resource Groups



32.4 Configuring SCPM Using the Command Line

This section introduces the command line configuration of SCPM. Learn how to start it, configure it, and work with profiles.

32.4.1 Starting SCPM and Defining Resource Groups

SCPM must be activated before use. Activate SCPM with `scpm enable`. When run for the first time, SCPM is initialized, which takes a few seconds. Deactivate SCPM with `scpm disable` at any time to prevent the unintentional switching of profiles. A subsequent reactivation simply resumes the initialization.

By default, SCPM handles network and printer settings as well as the X.Org configuration. To manage special services or configuration files, activate the respective resource groups. To list the predefined resource groups, use `scpm list_groups`. To see only the groups already activated, use `scpm list_groups -a`. Issue these commands as `root` on the command line.

```
scpm list_groups -a

nis          Network Information Service client
mail         Mail subsystem
ntpd         Network Time Protocol daemon
xf86         X Server settings
autofs       Automounter service
network      Basic network settings
printer      Printer settings
```

Activate or deactivate a group with `scpm activate_group NAME` or `scpm deactivate_group NAME`. Replace NAME with the relevant group name.

32.4.2 Creating and Managing Profiles

A profile named `default` already exists after SCPM has been activated. Get a list of all available profiles with `scpm list`. This one existing profile is also the active one, which can be verified with `scpm active`. The profile `default` is a basic configuration from which the other profiles are derived. For this reason, all settings that should be identical in all profiles should be made first. Then store these modifications in the active profile with `scpm reload`. The `default` profile can be copied and renamed as the basis for new profiles.

There are two ways to add a new profile. If the new profile (named `work` here) should be based on the profile `default`, create it with `scpm copy default work`. The command `scpm switch work` changes into the new profile, which can then be modified. You may want to modify the system configuration for special purposes and save the changes to a new profile. The command `scpm add work` creates a new profile by saving the current system configuration in the profile `work` and marking it as active. Running `scpm reload` then saves changes to the profile `work`.

Profiles can be renamed or deleted with the commands `scpm rename x y` and `scpm delete z`. For example, to rename `work` to `project`, enter `scpm rename work project`. To delete `project`, enter `scpm delete project`. The active profile cannot be deleted.

32.4.3 Switching Configuration Profiles

The command `scpm switch work` switches to another profile (the profile `work`, in this case). Switch to the active profile to include modified settings of the system configuration in the profile. This corresponds to the command `scpm reload`.

When switching profiles, SCPM first checks which resources of the active profile have been modified. It then queries whether the modification of each resource should be added to the active profile or dropped. If you prefer a separate listing of the resources (as in former versions of SCPM), use the `switch` command with the `-r` parameter:
`scpm switch -r work.`

```
scpm switch -r work
```

```
Checking for modified resources
Checking for Resources to be started/shut down
```

```
Checking for dependencies
Restoring profile default
```

SCPM then compares the current system configuration with the profile to which to switch. In this phase, SCPM evaluates which system services need to be stopped or restarted due to mutual dependencies or to reflect the changes in configuration. This is like a partial system reboot that concerns only a small part of the system while the rest continues operating without change. It is only at this point that the system services are stopped, all modified resources, such as configuration files, are written, and the system services are restarted.

32.4.4 Advanced Profile Settings

You can enter a description for every profile that is displayed with `scpm list`. For the active profile, set it with `scpm set description "text"`. Provide the name of the profile for inactive profiles, for example, `scpm set description "text" work`. Sometimes it might be desirable to perform additional actions not provided by SCPM while switching profiles. Attach up to four executables for each profile. They are invoked at different stages of the switching process. These stages are referred to as:

`prestop`

 prior to stopping services when leaving the profile

`poststop`

 after stopping services when leaving the profile

`prestart`

 prior to starting services when activating the profile

`poststart`

 after starting services when activating the profiles

Insert these actions with the command `set` by entering `scpm set prestop filename`, `scpm set poststop filename`, `scpm set prestart filename`, or `scpm set poststart filename`. The scripts must be executable and refer to the correct interpreter.

WARNING: Integrating a Custom Script

Additional scripts to be executed by SCPM must be made readable and executable for the superuser (`root`). The access to these files must be blocked for all other users. Enter the commands `chmod 700 filename` and `chown root:root filename` to give `root` exclusive permissions to the files.

Query all additional settings entered with `set` with `get`. The command `scpm get poststart`, for example, returns the name of the poststart call or simply nothing if nothing has been attached. Reset such settings by overwriting with " ". The command `scpm set prestop " "` removes the attached prestop program.

All `set` and `get` commands can be applied to an arbitrary profile in the same manner as comments are added. For example, `scpm get prestop filename work` or `scpm get prestop work`.

32.5 Troubleshooting

This section covers frequent problems encountered with SCPM. Learn how they can arise and how you can solve these issues.

32.5.1 SCPM and NetworkManager

NetworkManager and SCPM share functionality. Both integrate a machine into an existing network, hiding this transaction from the user. NetworkManager works dynamically and adapts to any new environment. SCPM is used to restore defined system setups.

Using NetworkManager and SCPM in parallel does not work properly, because NetworkManager does not provide configurations that can be restored by SCPM. SCPM works exceedingly well for anyone who needs reproducible setups. Any private user constantly switching (wireless) networks should consider using NetworkManager if network setup is the only component that needs to be adjusted.

If you want to use SCPM for network configuration management, disable NetworkManager.

32.5.2 Termination During the Switch Process

Sometimes SCPM stops working during a switch procedure. This may be caused by some outside effect, such as a user abort, a power failure, or even an error in SCPM itself. If this happens, an error message stating SCPM is locked appears the next time you start SCPM. This is for system safety, because the data stored in its database may differ from the state of the system. To resolve this issue, run `scpm recover`. SCPM performs all missing operations of the previous run. You can also run `scpm recover -b`, which tries to undo all already performed operations of the previous run. If you are using the YaST profile manager, get a recover dialog on start-up that offers to perform the commands described above.

32.6 For More Information

The latest documentation is available in the SCPM info pages (`info scpm`). Information for developers is available in `/usr/share/doc/packages/scpm`.

Power Management

Power management is especially important on laptop computers, but is also useful on other systems. Two technologies are available: APM (advanced power management) and ACPI (advanced configuration and power interface). In addition to these, it is also possible to control CPU frequency scaling to save power or decrease noise. These options can be configured manually or using a special YaST module.

Unlike APM, which was previously used on laptops for power management only, the hardware information and configuration tool ACPI is available on all modern computers (laptops, desktops, and servers). All power management technologies require suitable hardware and BIOS routines. Most laptops and many modern desktops and servers meet these requirements.

APM had been used in many older computers. Because APM largely consists of a function set implemented in the BIOS, the level of APM support may vary depending on the hardware. This is even more true of ACPI, which is even more complex. For this reason, it is virtually impossible to recommend one over the other. Simply test the various procedures on your hardware then select the technology that is best supported.

IMPORTANT: Power Management for AMD64 Processors

AMD64 processors with a 64-bit kernel only support ACPI.

33.1 Power Saving Functions

Power saving functions are not only significant for the mobile use of laptops, but also for desktop systems. The main functions and their use in the power management systems APM and ACPI are:

Standby

This operating mode turns off the display. On some computers, the processor performance is throttled. This function is not available in all APM implementations. This function corresponds to the ACPI state S1 or S2.

Suspend (to memory)

This mode writes the entire system state to the RAM. Subsequently, the entire system except the RAM is put to sleep. In this state, the computer consumes very little power. The advantage of this state is the possibility of resuming work at the same point within a few seconds without having to boot and restart applications. Devices using APM can usually be suspended by closing the lid and activated by opening it. This function corresponds to the ACPI state S3. The support of this state is still under development and therefore largely depends on the hardware.

Hibernation (suspend to disk)

In this operating mode, the entire system state is written to the hard disk and the system is powered off. There must be a swap partition at least as big as the RAM to write all the active data. Reactivation from this state takes about 30 to 90 seconds. The state prior to the suspend is restored. Some manufacturers offer useful hybrid variants of this mode, such as RediSafe in IBM Thinkpads. The corresponding ACPI state is S4. In Linux, suspend to disk is performed by kernel routines that are independent from APM and ACPI.

Battery Monitor

ACPI and APM check the battery charge status and provide information about it. Additionally, both systems coordinate actions to perform when a critical charge status is reached.

Automatic Power-Off

Following a shutdown, the computer is powered off. This is especially important when an automatic shutdown is performed shortly before the battery is empty.

Shutdown of System Components

Switching off the hard disk is the greatest single aspect of the power saving potential of the overall system. Depending on the reliability of the overall system, the hard disk can be put to sleep for some time. However, the risk of losing data increases with the duration of the sleep periods. Other components, like PCI devices that can be put into a special power saving mode, can be deactivated with ACPI (at least theoretically) or permanently disabled in the BIOS setup.

Processor Speed Control

In connection with the CPU, energy can be saved in three different ways: frequency and voltage scaling (also known as PowerNow! or Speedstep), throttling, and putting the processor to sleep (C states). Depending on the operating mode of the computer, these methods can also be combined.

33.2 APM

Some of the power saving functions are performed by the APM BIOS itself. On many laptops, standby and suspend states can be activated with key combinations or by closing the lid without any special operating system function. However, to activate these modes with a command, certain actions must be triggered before the system is suspended. To view the battery charge level, you need special program packages and a suitable kernel.

SUSE Linux kernels have built-in APM support. However, APM is only activated if ACPI is not implemented in the BIOS and an APM BIOS is detected. To activate APM support, ACPI must be disabled with `acpi=off` at the boot prompt. Enter `cat /proc/apm` to check if APM is active. An output consisting of various numbers indicates that everything is OK. You should now be able to shut down the computer with the command `shutdown -h`.

BIOS implementations that are not fully standard-compliant can cause problems with APM. Some problems can be circumvented with special boot parameters. All parameters are entered at the boot prompt in the form of `apm=parameter`. `parameter` is one of:

on or off

Enable or disable APM support.

(no-)allow-ints

Allow interrupts during the execution of BIOS functions.

(no-)broken-psr

The “GetPowerStatus” function of the BIOS does not work properly.

(no-)realmode-power-off

Reset processor to real mode prior to shutdown.

(no-)debug

Log APM events in system log.

(no-)power-off

Power system off after shutdown.

bounce-interval=*n*

Time in hundredths of a second after a suspend event during which additional suspend events are ignored.

idle-threshold=*n*

System inactivity percentage from which the BIOS function `idle` is executed (0=always, 100=never).

idle-period=*n*

Time in hundredths of a second after which the system activity is measured.

The APM daemon (`apmd`) is no longer used. Its functionality is now handled by the new powersaved, which also supports ACPI and provides many other features.

33.3 ACPI

ACPI (advanced configuration and power interface) was designed to enable the operating system to set up and control the individual hardware components. ACPI supersedes both PnP and APM. It delivers information about the battery, AC adapter, temperature, fan, and system events, like “close lid” or “battery low.”

The BIOS provides tables containing information about the individual components and hardware access methods. The operating system uses this information for tasks like assigning interrupts or activating and deactivating components. Because the operating

system executes commands stored in the BIOS, the functionality depends on the BIOS implementation. The tables ACPI can detect and load are reported in `/var/log/boot.msg`. See [Section 33.3.4, “Troubleshooting”](#) (page 586) for more information about troubleshooting ACPI problems.

33.3.1 ACPI in Action

If the kernel detects an ACPI BIOS when the system is booted, ACPI is activated automatically and APM is deactivated. The boot parameter `acpi=force` may be necessary for some older machines. The computer must support ACPI 2.0 or later. Check the kernel boot messages in `/var/log/boot.msg` to see if ACPI was activated.

Subsequently, a number of modules must be loaded. This is done by the start script of `acpid`. If any of these modules cause problems, the respective module can be excluded from loading or unloading in `/etc/sysconfig/powersave/common`. The system log (`/var/log/messages`) contains the messages of the modules, enabling you to see which components were detected.

`/proc/acpi` now contains a number of files that provide information about the system state or can be used to change some of the states. Some features do not work yet because they are still under development and the support of some functions largely depends on the implementation of the manufacturer.

All files (except `dsdt` and `fadt`) can be read with `cat`. In some files, settings can be modified with `echo`, for example, `echo X > file` to specify suitable values for `X`. One possibility for easy access to those values is the `powersave` command, which acts as a front-end for the Powersave daemon. The following describes the most important files:

`/proc/acpi/info`

General information about ACPI.

`/proc/acpi/alarm`

Here, specify when the system should wake from a sleep state. Currently, this feature is not fully supported.

`/proc/acpi/sleep`

Provides information about possible sleep states.

/proc/acpi/event

All events are reported here and processed by the Powersave daemon (`powersaved`). If no daemon accesses this file, events, such as a brief click on the power button or closing the lid, can be read with `cat /proc/acpi/event` (terminate with **[Ctrl]** + **[C]**).

/proc/acpi/dsdt and /proc/acpi/fadt

These files contain the ACPI tables DSDT (differentiated system description table) and FADT (fixed ACPI description table). They can be read with `acpidmp`, `acpidisasm`, and `dmdecode`. These programs and their documentation are located in the package `pmtools`. For example, `acpidmp DSDT | acpidisasm`.

/proc/acpi/ac_adapter/AC/state

Shows whether the AC adapter is connected.

/proc/acpi/battery/BAT*/{alarm,info,state}

Detailed information about the battery state. The charge level is read by comparing the last full capacity from `info` with the remaining capacity from `state`. A more comfortable way to do this is to use one of the special programs introduced in [Section 33.3.3, “ACPI Tools”](#) (page 586). The charge level at which a battery event (such as warning, low and critical) is triggered can be specified in `alarm`.

/proc/acpi/button

This directory contains information about various switches, like the laptop lid and buttons.

/proc/acpi/fan/FAN/state

Shows if the fan is currently active. Activate or deactivate the fan manually by writing 0 (on) or 3 (off) into this file. However, both the ACPI code in the kernel and the hardware (or the BIOS) overwrite this setting when the system gets too warm.

/proc/acpi/processor/*

A separate subdirectory is kept for each CPU included in your system.

/proc/acpi/processor/*/info

Information about the energy saving options of the processor.

/proc/acpi/processor/*/*power

Information about the current processor state. An asterisk next to C2 indicates that the processor is idle. This is the most frequent state, as can be seen from the usage value.

/proc/acpi/processor/*/*throttling

Can be used to set the throttling of the processor clock. Usually, throttling is possible in eight levels. This is independent of the frequency control of the CPU.

/proc/acpi/processor/*/*limit

If the performance (outdated) and the throttling are automatically controlled by a daemon, the maximum limits can be specified here. Some of the limits are determined by the system. Some can be adjusted by the user.

/proc/acpi/thermal_zone/

A separate subdirectory exists for every thermal zone. A thermal zone is an area with similar thermal properties whose number and names are designated by the hardware manufacturer. However, many of the possibilities offered by ACPI are rarely implemented. Instead, the temperature control is handled conventionally by the BIOS. The operating system is not given much opportunity to intervene, because the life span of the hardware is at stake. Therefore, some of the files only have a theoretical value.

/proc/acpi/thermal_zone/*/*temperature

Current temperature of the thermal zone.

/proc/acpi/thermal_zone/*/*state

The state indicates if everything is ok or if ACPI applies active or passive cooling. In the case of ACPI-independent fan control, this state is always ok.

/proc/acpi/thermal_zone/*/*cooling_mode

Select the cooling method controlled by ACPI. Choose from passive (less performance, economical) or active cooling mode (full performance, fan noise).

/proc/acpi/thermal_zone/*/*trip_points

Enables the determination of temperature limits for triggering specific actions, like passive or active cooling, suspension (hot), or a shutdown (critical). The possible actions are defined in the DSDT (device-dependent). The trip points determined in the ACPI specification are critical, hot, passive, active1, and active2. Even if not all of them are implemented, they must always be entered

in this file in this order. For example, the entry `echo 90:0:70:0:0 > trip_points` sets the temperature for `critical` to 90 and the temperature for `passive` to 70 (all temperatures measured in degrees Celsius).

`/proc/acpi/thermal_zone/*/polling_frequency`

If the value in `temperature` is not updated automatically when the temperature changes, toggle the polling mode here. The command `echo X > /proc/acpi/thermal_zone/*/polling_frequency` causes the temperature to be queried every X seconds. Set `X=0` to disable polling.

None of these settings, information, and events need to be edited manually. This can be done with the Powersave daemon (`powersaved`) and its various front-ends, like `powersave`, `kpowersave`, and `wmpowersave`. See [Section 33.3.3, “ACPI Tools”](#) (page 586).

33.3.2 Controlling the CPU Performance

The CPU can save energy in three ways. Depending on the operating mode of the computer, these methods can be combined. Saving energy also means that the system heats up less and the fans are activated less frequently.

Frequency and Voltage Scaling

PowerNow! and Speedstep are the designations AMD and Intel use for this technology. However, this technology is also applied in processors of other manufacturers. The clock frequency of the CPU and its core voltage are reduced at the same time, resulting in more than linear energy savings. This means that when the frequency is halved (half performance), far less than half of the energy is consumed. This technology is independent from APM or ACPI. There are two main approaches to performing CPU frequency scaling—by the kernel itself or by a userspace application. Therefore, there are different kernel governors that can be set below `/sys/devices/system/cpu/cpu*/cpufreq/`.

userspace governor

If the userspace governor is set, the kernel gives the control of CPU frequency scaling to a userspace application, usually a daemon. In SUSE Linux distributions, this daemon is the `powersaved` package. When this implementation is used, the CPU frequency is adjusted in regard to the current system load. By default, one of the kernel implementations is used. However, on some

hardware or in regard to specific processors or drivers, the userspace implementation is still the only working solution.

ondemand governor

This is the kernel implementation of a dynamic CPU frequency policy and should work on most systems. As soon as there is a high system load, the CPU frequency is immediately increased. It is lowered on a low system load.

conservative governor

This governor is similar to the ondemand implementation, except that a more conservative policy is used. The load of the system must be high for a specific amount of time before the CPU frequency is increased.

powersave governor

The cpu frequency is statically set to the lowest possible.

performance governor

The cpu frequency is statically set to the highest possible.

Throttling the Clock Frequency

This technology omits a certain percentage of the clock signal impulses for the CPU. At 25% throttling, every fourth impulse is omitted. At 87.5%, only every eighth impulse reaches the processor. However, the energy savings are a little less than linear. Normally, throttling is only used if frequency scaling is not available or to maximize power savings. This technology, too, must be controlled by a special process. The system interface is `/proc/acpi/processor/*/throttling`.

Putting the Processor to Sleep

The operating system puts the processor to sleep whenever there is nothing to do. In this case, the operating system sends the CPU a `halt` command. There are three states: C1, C2, and C3. In the most economic state, C3, even the synchronization of the processor cache with the main memory is halted. Therefore, this state can only be applied if no other device modifies the contents of the main memory via bus master activity. Some drivers prevent the use of C3. The current state is displayed in `/proc/acpi/processor/*/power`.

Frequency scaling and throttling are only relevant if the processor is busy, because the most economic C state is applied anyway when the processor is idle. If the CPU is busy, frequency scaling is the recommended power saving method. Often the processor only works with a partial load. In this case, it can be run with a lower frequency. Usually, dynamic frequency scaling controlled by the kernel ondemand governor or a daemon,

such as powersaved, is the best approach. A static setting to a low frequency is useful for battery operation or if you want the computer to be cool or quiet.

Throttling should be used as the last resort, for example, to extend the battery operation time despite a high system load. However, some systems do not run smoothly when they are throttled too much. Moreover, CPU throttling does not make sense if the CPU has little to do.

In SUSE Linux these technologies are controlled by the powersave daemon. The configuration is explained in [Section 33.5, “The powersave Package”](#) (page 589).

33.3.3 ACPI Tools

The range of more or less comprehensive ACPI utilities includes tools that merely display information, like the battery charge level and the temperature (acpi, klaptopdaemon, wmacpimon, etc.), tools that facilitate the access to the structures in `/proc/acpi` or that assist in monitoring changes (akpi, acpiw, gtkacpiw), and tools for editing the ACPI tables in the BIOS (package `pmtools`).

33.3.4 Troubleshooting

There are two different types of problems. On one hand, the ACPI code of the kernel may contain bugs that were not detected in time. In this case, a solution will be made available for download. More often, however, the problems are caused by the BIOS. Sometimes, deviations from the ACPI specification are purposely integrated in the BIOS to circumvent errors in the ACPI implementation in other widespread operating systems. Hardware components that have serious errors in the ACPI implementation are recorded in a blacklist that prevents the Linux kernel from using ACPI for these components.

The first thing to do when problems are encountered is to update the BIOS. If the computer does not boot at all, one of the following boot parameters may be helpful:

`pci=noacpi`

Do not use ACPI for configuring the PCI devices.

`acpi=oldboot`

Only perform a simple resource configuration. Do not use ACPI for other purposes.

acpi=off

Disable ACPI.

WARNING: Problems Booting without ACPI

Some newer machines (especially SMP systems and AMD64 systems) need ACPI for configuring the hardware correctly. On these machines, disabling ACPI can cause problems.

Monitor the boot messages of the system with the command `dmesg | grep -2i acpi` (or all messages, because the problem may not be caused by ACPI) after booting. If an error occurs while parsing an ACPI table, the most important table—the DS-DT—can be replaced with an improved version. In this case, the faulty DSDT of the BIOS is ignored. The procedure is described in [Section 33.5.4, “Troubleshooting”](#) (page 595).

In the kernel configuration, there is a switch for activating ACPI debug messages. If a kernel with ACPI debugging is compiled and installed, experts searching for an error can be supported with detailed information.

If you experience BIOS or hardware problems, it is always advisable to contact the manufacturers. Especially if they do not always provide assistance for Linux, they should be confronted with the problems. Manufacturers will only take the issue seriously if they realize that an adequate number of their customers use Linux.

For More Information

Additional documentation and help on ACPI:

- <http://www.cpqlinux.com/acpi-howto.html> (detailed ACPI HOWTO, contains DSDT patches)
- <http://www.intel.com/technology/iapc/acpi/faq.htm> (ACPI FAQ @Intel)
- <http://acpi.sourceforge.net/> (the ACPI4Linux project at Sourceforge)
- <http://www.poupinou.org/acpi/> (DSDT patches by Bruno Ducrot)

33.4 Rest for the Hard Disk

In Linux, the hard disk can be put to sleep entirely if it is not needed or it can be run in a more economic or quieter mode. On modern laptops, you do not need to switch off the hard disks manually, because they automatically enter an economic operating mode whenever they are not needed. However, if you want to maximize power savings, test some of the following methods. Most of the functions can be controlled with powersaved and the YaST power management module, which is discussed in further detail in [Section 33.6, “The YaST Power Management Module”](#) (page 597).

The hdparm application can be used to modify various hard disk settings. The option `-y` instantly switches the hard disk to the standby mode. `-Y` puts it to sleep. `hdparm -S x` causes the hard disk to be spun down after a certain period of inactivity. Replace `x` as follows: 0 disables this mechanism, causing the hard disk to run continuously. Values from 1 to 240 are multiplied by 5 seconds. Values from 241 to 251 correspond to 1 to 11 times 30 minutes.

Internal power saving options of the hard disk can be controlled with the option `-B`. Select a value from 0 to 255 for maximum saving to maximum throughput. The result depends on the hard disk used and is difficult to assess. To make a hard disk quieter, use the option `-M`. Select a value from 128 to 254 for quiet to fast.

Often, it is not so easy to put the hard disk to sleep. In Linux, numerous processes write to the hard disk, waking it up repeatedly. Therefore, it is important to understand how Linux handles data that needs to be written to the hard disk. First, all data is buffered in the RAM. This buffer is monitored by the kernel update daemon (`kupdated`). When the data reaches a certain age limit or when the buffer is filled to a certain degree, the buffer content is flushed to the hard disk. The buffer size is dynamic and depends on the size of the memory and the system load. By default, `kupdated` is set to short intervals to achieve maximum data integrity. It checks the buffer every 5 seconds and notifies the `bdflush` daemon when data is older than 30 seconds or the buffer reaches a fill level of 30%. The `bdflush` daemon then writes the data to the hard disk. It also writes independently from `kupdated` if, for instance, the buffer is full.

WARNING: Impairment of the Data Integrity

Changes to the kernel update daemon settings endanger the data integrity.

Apart from these processes, journaling file systems, like ReiserFS and Ext3, write their metadata independently from bdflush, which also prevents the hard disk from spinning down. To avoid this, a special kernel extension has been developed for mobile devices. See `/usr/src/linux/Documentation/laptop-mode.txt` for details.

Another important factor is the way active programs behave. For example, good editors regularly write hidden backups of the currently modified file to the hard disk, causing the disk to wake up. Features like this can be disabled at the expense of data integrity.

In this connection, the mail daemon postfix makes use of the variable `POSTFIX_LAPTOP`. If this variable is set to `yes`, postfix accesses the hard disk far less frequently. However, this is irrelevant if the interval for `kupdated` was increased.

33.5 The powersave Package

The `powersave` package cares about all the previously-mentioned power saving functions. Due to the increasing demand for lower energy consumption in general, some of its features are also important on workstations and servers, such as suspend, standby, or cpu frequency scaling.

This package contains all power management features of your computer. It supports hardware using ACPI, APM, IDE hard disks, and PowerNow! or SpeedStep technologies. The functionalities from the packages `apmd`, `acpid`, `ospmd`, and `cpufreqd` (now `cpuspeed`) have been consolidated in the `powersave` package. Daemons from these packages, except `acpid` that acts as a multiplexer for acpi events, should not be run concurrently with the `powersave` daemon.

Even if your system does not contain all the hardware elements listed above, use the `powersave` daemon for controlling the power saving function. Because ACPI and APM are mutually exclusive, you can only use one of these systems on your computer. The daemon automatically detects any changes in the hardware configuration.

33.5.1 Configuring the powersave Package

The configuration of `powersave` is distributed to several files. Every configuration option listed there contains additional documentation about its functionality.

/etc/sysconfig/powersave/common

This file contains general settings for the powersave daemon. For example, the amount of debug messages in /var/log/messages can be increased by increasing the value of the variable DEBUG.

/etc/sysconfig/powersave/events

The powersave daemon needs this file for processing system events. An event can be assigned external actions or actions performed by the daemon itself. For external actions, the daemon tries to run an executable file (usually a Bash script) in /usr/lib/powersave/scripts/. Predefined internal actions are:

- ignore
- throttle
- dethrottle
- suspend_to_disk
- suspend_to_ram
- standby
- do_suspend_to_disk
- do_suspend_to_ram
- do_standby
- notify
- screen_saver
- reread_cpu_capabilities

throttle slows down the processor by the value defined in MAX_THROTTLING. This value depends on the current scheme. dethrottle sets the processor to full performance. suspend_to_disk, suspend_to_ram, and standby trigger the system event for a sleep mode. These three actions are generally responsible for triggering the sleep mode, but they should always be associated with specific system events.

The directory `/usr/lib/powersave/scripts` contains scripts for processing events:

`switch_vt`

Useful if the screen is displaced after a suspend or standby.

`wm_logout`

Saves the settings and logs out from GNOME, KDE, or other window managers.

`wm_shutdown`

Saves the GNOME or KDE settings and shuts down the system.

`set_disk_settings`

Executes the disk settings made in `/etc/sysconfig/powersave/disk`.

If, for example, the variable

`EVENT_GLOBAL_SUSPEND2DISK="prepare_suspend_to_disk do_suspend_to_disk"` is set, the two scripts or actions are processed in the specified order as soon as the user gives powersaved the command for the sleep mode `suspend to disk`. The daemon runs the external script `/usr/lib/powersave/scripts/prepare_suspend_to_disk`. After this script has been processed successfully, the daemon runs the internal action `do_suspend_to_disk` and sets the computer to the sleep mode after the script has unloaded critical modules and stopped services.

The actions for the event of a sleep button could be modified as in

`EVENT_BUTTON_SLEEP="notify suspend_to_disk"`. In this case, the user is informed about the suspend by a pop-up window in X or a message on the console. Subsequently, the event `EVENT_GLOBAL_SUSPEND2DISK` is generated, resulting in the execution of the mentioned actions and a secure system suspend mode. The internal action `notify` can be customized using the variable `NOTIFY_METHOD` in `/etc/sysconfig/powersave/common`.

`/etc/sysconfig/powersave/cpufreq`

Contains variables for optimizing the dynamic CPU frequency settings and whether the userspace or the kernel implementation should be used.

`/etc/sysconfig/powersave/battery`

Contains battery limits and other battery-specific settings.

`/etc/sysconfig/powersave/sleep`

In this file, activate the sleep modes and determine which critical modules should be unloaded and which services should be stopped prior to a suspend or standby event. When the system is resumed, these modules are reloaded and the services are restarted. You can even delay a triggered sleep mode, for example, to save files. The default settings mainly concern USB and PCMCIA modules. A failure of suspend or standby is usually caused by certain modules. See [Section 33.5.4, “Troubleshooting”](#) (page 595) for more information about identifying the error.

`/etc/sysconfig/powersave/thermal`

Activates cooling and thermal control. Details about this subject are available in the file `/usr/share/doc/packages/powersave/README.thermal`.

`/etc/sysconfig/powersave/disk`

This configuration file controls the actions and settings made regarding the hard disk.

`/etc/sysconfig/powersave/scheme_*`

These are the various schemes that adapt the power consumption to certain deployment scenarios. A number of schemes are preconfigured and can be used as they are. Custom schemes can be saved here.

33.5.2 Configuring APM and ACPI

Suspend and Standby

There are three basic ACPI sleep modes and two APM sleep modes:

Suspend to Disk (ACPI S4, APM suspend)

Saves the entire memory content to the hard disk. The computer is switched off completely and does not consume any power. This sleep mode is enabled by default and should work on all systems.

Suspend to RAM (ACPI S3, APM suspend)

Saves the states of all devices to the main memory. Only the main memory continues consuming power. This sleep mode is disabled by default because still causes problems on some systems. However, support has been extended greatly.

Standby (ACPI S1, APM standby)

Switches some devices off (manufacturer-dependent).

Make sure that the following default options are set in the file `/etc/sysconfig/powersave/events` for the correct processing of suspend, standby, and resume (default settings following the installation of SUSE Linux):

```
EVENT_GLOBAL_SUSPEND2DISK=
    "prepare_suspend_to_disk screen_saver do_suspend_to_disk"
EVENT_GLOBAL_SUSPEND2RAM=
    "prepare_suspend_to_ram screen_saver do_suspend_to_ram"
EVENT_GLOBAL_STANDBY=
    "prepare_standby screen_saver do_standby"
EVENT_GLOBAL_RESUME_SUSPEND2DISK=
    "restore_after_suspend_to_disk"
EVENT_GLOBAL_RESUME_SUSPEND2RAM=
    "restore_after_suspend_to_ram"
EVENT_GLOBAL_RESUME_STANDBY=
    "restore_after_standby"
```

Custom Battery States

In the file `/etc/sysconfig/powersave/battery`, define three battery charge levels (in percent) that trigger system alerts or specific actions when they are reached.

```
BATTERY_WARNING=12
BATTERY_LOW=7
BATTERY_CRITICAL=2
```

The actions or scripts to execute when the charge levels drop under the specified limits are defined in the configuration file `/etc/sysconfig/powersave/events`.

The standard actions for buttons can be modified as described in [Section 33.5.1, “Configuring the powersave Package”](#) (page 589).

```
EVENT_BATTERY_NORMAL="ignore"
EVENT_BATTERY_WARNING="notify"
EVENT_BATTERY_LOW="notify"
EVENT_BATTERY_CRITICAL="wm_shutdown"
```

Adapting Power Consumption to Various Conditions

The system behavior can be adapted to the type of power supply. The power consumption of the system should be reduced when the system is disconnected from the AC power supply and operated with the battery. Similarly, the performance should automatically

increase as soon as the system is connected to the AC power supply. The CPU frequency, the power saving function of IDE, and a number of other parameters can be modified.

The actions to execute when the computer is disconnected from or connected to the AC power supply are defined in `/etc/sysconfig/powersave/events`. Select the schemes to use in `/etc/sysconfig/powersave/common`:

```
AC_SCHEME="performance"
BATTERY_SCHEME="powersave"
```

The schemes are stored in files in `/etc/sysconfig/powersave`. The filenames are in the format `scheme_name-of-the-scheme`. The example refers to two schemes: `scheme_performance` and `scheme_powersave`. `performance`, `powersave`, `presentation`, and `acoustic` are preconfigured. Existing schemes can be edited, created, deleted, or associated with different power supply states with the help of the YaST power management module described in [Section 33.6, “The YaST Power Management Module”](#) (page 597).

33.5.3 Additional ACPI Features

If you use ACPI, you can control the response of your system to *ACPI buttons* (power, sleep, lid open, and lid closed). Configure execution of the actions in `/etc/sysconfig/powersave/events`. Refer to this configuration file for an explanation of the individual options.

`EVENT_BUTTON_POWER="wm_shutdown"`

When the power button is pressed, the system responds by shutting down the respective window manager (KDE, GNOME, fvwm, etc.).

`EVENT_BUTTON_SLEEP="suspend_to_disk"`

When the sleep button is pressed, the system is set to the suspend-to-disk mode.

`EVENT_BUTTON_LID_OPEN="ignore"`

Nothing happens when the lid is opened.

`EVENT_BUTTON_LID_CLOSED="screen_saver"`

When the lid is closed, the screen saver is activated.

`EVENT_OTHER="ignore"`

This event happens if an unknown event is encountered by the daemon. Unknown events include ACPI hotkeys on some machines.

Further throttling of the CPU performance is possible if the CPU load does not exceed a specified limit for a specified time. Specify the load limit in PROCESSOR_IDLE_LIMIT and the time-out in CPU_IDLE_TIMEOUT. If the CPU load stays below the limit longer than the time-out, the event configured in EVENT_PROCESSOR_IDLE is activated. If the CPU is busy again, EVENT_PROCESSOR_BUSY is executed.

33.5.4 Troubleshooting

All error messages and alerts are logged in the file /var/log/messages. If you cannot find the needed information, increase the verbosity of the messages of powersave using DEBUG in the file /etc/sysconfig/powersave/common. Increase the value of the variable to 7 or even 15 and restart the daemon. The more detailed error messages in /var/log/messages should help you to find the error. The following sections cover the most common problems with powersave.

ACPI Activated with Hardware Support but Functions Do Not Work

If you experience problems with ACPI, use the command dmesg | grep -i acpi to search the output of dmesg for ACPI-specific messages. A BIOS update may be required to resolve the problem. Go to the home page of your laptop manufacturer, look for an updated BIOS version, and install it. Ask the manufacturer to comply with the latest ACPI specification. If the errors persist after the BIOS update, proceed as follows to replace the faulty DSDT table in your BIOS with an updated DS DT:

- 1 Download the DSDT for your system from <http://acpi.sourceforge.net/dsdt/tables>. Check if the file is decompressed and compiled as shown by the file extension .aml (ACPI machine language). If this is the case, continue with step 3.
- 2 If the file extension of the downloaded table is .asl (ACPI source language), compile it with iasl (package pmtools). Enter the command iasl -sa file.asl. The latest version of iasl (Intel ACPI compiler) is available at <http://developer.intel.com/technology/iapc/acpi/downloads.htm>.

- 3** Copy the file `DSDT.aml` to any location (`/etc/DSDT.aml` is recommended). Edit `/etc/sysconfig/kernel` and adapt the path to the DSDT file accordingly. Start `mkinitrd` (package `mkinitrd`). Whenever you install the kernel and use `mkinitrd` to create an `initrd`, the modified DSDT is integrated and loaded when the system is booted.

CPU Frequency Does Not Work

Refer to the kernel sources (`kernel-source`) to see if your processor is supported. You may need a special kernel module or module option to activate CPU frequency control. This information is available in `/usr/src/linux/Documentation/cpu-freq/*`. If a special module or module option is needed, configure it in the file `/etc/sysconfig/powersave/cpufreq` by means of the variables `CPUFREQD_MODULE` and `CPUFREQD_MODULE_OPTS`.

Suspend and Standby Do Not Work

There are several kernel-related problems that prevent the use of suspend and standby on ACPI systems:

- Currently, systems with more than 1 GB RAM do not support suspend.
- Currently, multiprocessor systems and systems with a P4 processor (with hyper-threading) do not support suspend.

The error may also be due to a faulty DSDT implementation (BIOS). If this is the case, install a new DSDT.

On ACPI and APM systems: When the system tries to unload faulty modules, the system is arrested or the suspend event is not triggered. The same can also happen if you do not unload modules or stop services that prevent a successful suspend. In both cases, try to identify the faulty module that prevented the sleep mode. The log files generated by the powersave daemon in `/var/log/suspend2ram.log` and `/var/log/suspend2disk.log` are very helpful in this regard. If the computer does not enter the sleep mode, the cause lies in the last module unloaded. Manipulate the following settings in `/etc/sysconfig/powersave/sleep` to unload problematic modules prior to a suspend or standby.

```
UNLOAD_MODULES_BEFORE_SUSPEND2DISK=""  
UNLOAD_MODULES_BEFORE_SUSPEND2RAM=""
```

```
UNLOAD_MODULES_BEFORE_STANDBY=""  
SUSPEND2DISK_RESTART_SERVICES=""  
SUSPEND2RAM_RESTART_SERVICES=""  
STANDBY_RESTART_SERVICES=""
```

If you use suspend or standby in changing network environments or in connection with remotely mounted file systems, such as Samba and NIS, use automounter to mount them or add the respective services, for example, `smbfs` or `nfs`, in the above-mentioned variable. If an application accesses the remotely mounted file system prior to a suspend or standby, the service cannot be stopped correctly and the file system cannot be unmounted properly. After resuming the system, the file system may be corrupt and must be remounted.

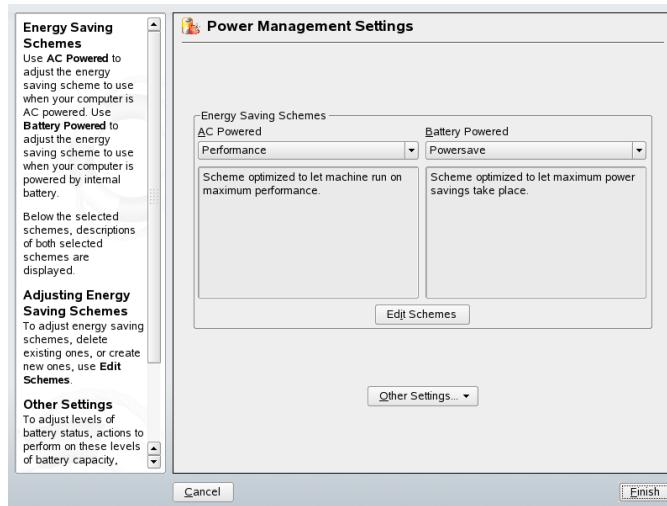
33.5.5 For More Information

- `/usr/share/doc/packages/powersave`—Local Powersave daemon documentation
- <http://powersave.sourceforge.net>—Most recent Powersave daemon documentation
- http://www.opensuse.org/Projects_Powersave—Project page in the openSUSE wiki

33.6 The YaST Power Management Module

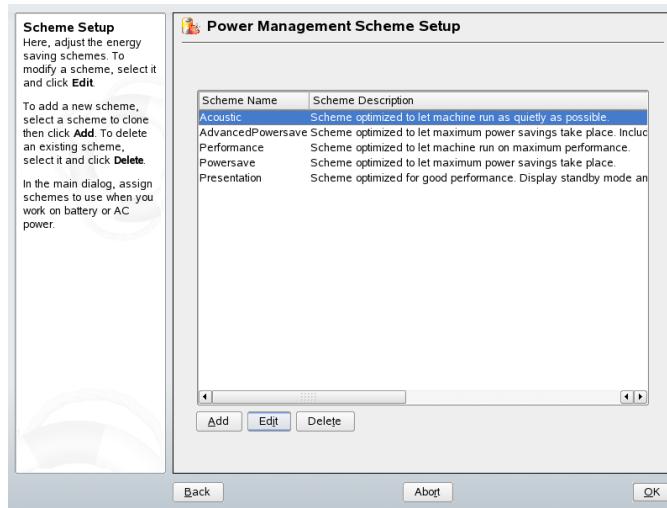
The YaST power management module can configure all power management settings already described. When started from the YaST Control Center with *System → Power Management*, the first dialog of the module opens (see [Figure 33.1, “Scheme Selection”](#) (page 598)).

Figure 33.1 Scheme Selection



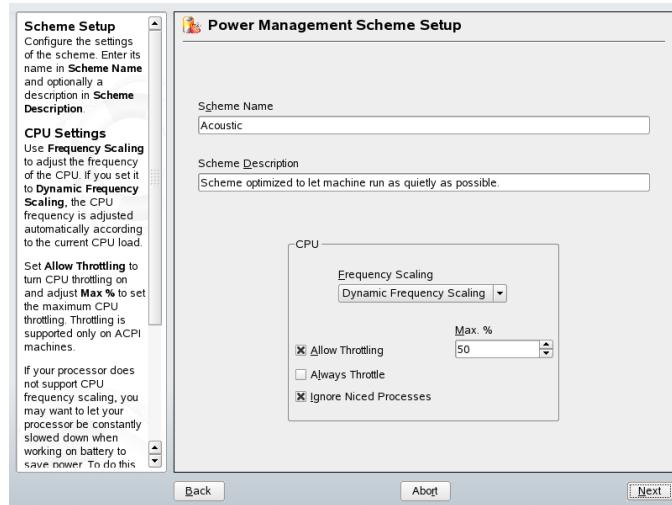
In this dialog, select the schemes to use for battery operation and AC operation. To add or modify the schemes, click *Edit Schemes*, which opens an overview of the existing schemes like that shown in [Figure 33.2, “Overview of Existing Schemes”](#) (page 598).

Figure 33.2 Overview of Existing Schemes



In the scheme overview, select the scheme to modify then click *Edit*. To create a new scheme, click *Add*. The dialog that opens is the same in both cases and is shown in Figure 33.3, “Configuring a Scheme” (page 599).

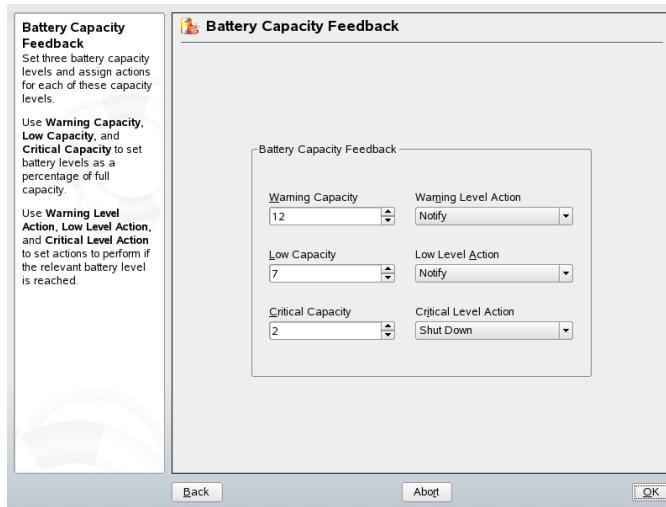
Figure 33.3 Configuring a Scheme



First, enter a suitable name and description for the new or edited scheme. Determine if and how the CPU performance should be controlled for this scheme. Decide if and to what extent frequency scaling and throttling should be used. In the following dialog for the hard disk, define a *Standby Policy* for maximum performance or for energy saving. The *Acoustic Policy* controls the noise level of the hard disk (supported by few hard disks). The *Cooling Policy* determines the cooling method to use. Unfortunately, this type of thermal control is rarely supported by the BIOS. Read `/usr/share/doc/packages/powersave/powersave_manual.html#Thermal` to learn how you can use the fan and passive cooling methods.

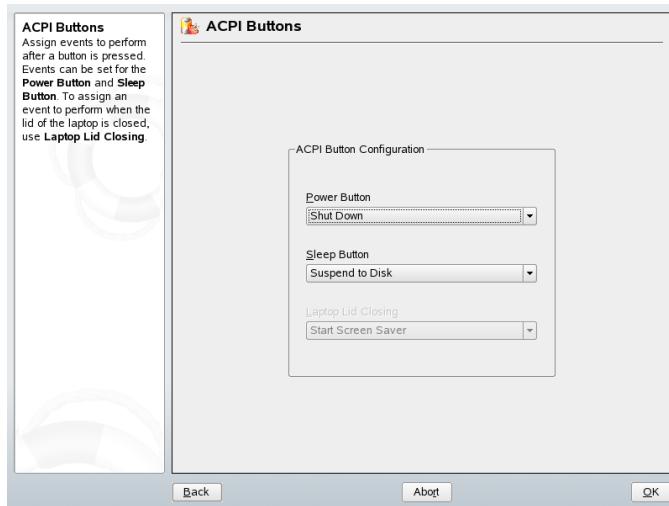
Global power management settings can also be made from the initial dialog using *Battery Warnings*, *ACPI Settings*, or *Enable Suspend*. Click *Battery Warnings* to access the dialog for the battery charge level, shown in Figure 33.4, “Battery Charge Level” (page 600).

Figure 33.4 Battery Charge Level



The BIOS of your system notifies the operating system whenever the charge level drops under certain configurable limits. In this dialog, define three limits: *Warning Capacity*, *Low Capacity*, and *Critical Capacity*. Specific actions are triggered when the charge level drops under these limits. Usually, the first two states merely trigger a notification to the user. The third critical level triggers a shutdown, because the remaining energy is not sufficient for continued system operation. Select suitable charge levels and the desired actions then click *OK* to return to the start dialog.

Figure 33.5 *ACPI Settings*



Access the dialog for configuring the ACPI buttons using *ACPI Settings*. It is shown in [Figure 33.5, “ACPI Settings”](#) (page 601). The settings for the ACPI buttons determine how the system should respond to certain switches. Configure the system response to pressing the power button, pressing the sleep button, and closing the laptop lid. Click *OK* to complete the configuration and return to the start dialog.

Click *Enable Suspend* to enter a dialog in which to determine if and how users of this system may use the suspend or standby functionality. Click *OK* to return to the main dialog. Click *OK* again to exit the module and confirm your power management settings.

34

Wireless Communication

There are several possibilities for using your Linux system to communicate with other computers, cellular phones, or peripheral devices. WLAN (wireless LAN) can be used to network laptops. Bluetooth can be used to connect individual system components (mouse, keyboard), peripheral devices, cellular phones, PDAs, and individual computers with each other. IrDA is mostly used for communication with PDAs or cellular phones. This chapter introduces all three technologies and their configuration.

34.1 Wireless LAN

Wireless LANs have become an indispensable aspect of mobile computing. Today, most laptops have built-in WLAN cards. The 802.11 standard for the wireless communication of WLAN cards was prepared by the IEEE organization. Originally, this standard provided for a maximum transmission rate of 2 MBit/s. Meanwhile, several supplements have been added to increase the data rate. These supplements define details such as the modulation, transmission output, and transmission rates:

Table 34.1 Overview of Various WLAN Standards

Name	Band (GHz)	Maximum Transmission Rate (MBit/s)	Note
802.11	2.4	2	Outdated; virtually no end devices available
802.11b	2.4	11	Widespread
802.11a	5	54	Less common
802.11g	2.4	54	Backward-compatible with 11b

Additionally, there are proprietary standards, like the 802.11b variation of Texas Instruments with a maximum transmission rate of 22 MBit/s (sometimes referred to as 802.11b+). However, the popularity of cards using this standard is limited.

34.1.1 Hardware

802.11 cards are not supported by SUSE Linux. Most cards using 802.11a, 802.11b, and 802.11g are supported. New cards usually comply with the 802.11g standard, but cards using 802.11b are still available. Normally, cards with the following chips are supported:

- Aironet 4500, 4800
- Atheros 5210, 5211, 5212
- Atmel at76c502, at76c503, at76c504, at76c506
- Intel PRO/Wireless 2100, 2200BG, 2915ABG
- Intersil Prism2/2.5/3
- Intersil PrismGT
- Lucent/Agere Hermes

- Ralink RT2400, RT2500
- Texas Instruments ACX100, ACX111
- ZyDAS zd1201

A number of older cards that are rarely used and no longer available are also supported. An extensive list of WLAN cards and the chips they use is available at the Web site of *AbsoluteValue Systems* at http://www.linux-wlan.org/docs/wlan_adapters.html.gz. <http://wiki.uni-konstanz.de/wiki/bin/view/Wireless/ListeChipsatz> provides an overview of the various WLAN chips.

Some cards need a firmware image that must be loaded into the card when the driver is initialized. This is the case with Intersil PrismGT, Atmel, and TI ACX100 and ACX111. The firmware can easily be installed with the YaST Online Update. The firmware for Intel PRO/Wireless cards ships with SUSE Linux and is automatically installed by YaST as soon as a card of this type is detected. More information about this subject is available in the installed system in `/usr/share/doc/packages/wireless-tools/README.firmware`.

34.1.2 Function

In wireless networking, various techniques and configurations are used to ensure fast, high-quality, and secure connections. Different operating types suit different setups. It can be difficult to choose the right authentication method. The available encryption methods have different advantages and pitfalls.

Operating Mode

Basically, wireless networks can be classified as managed networks and ad-hoc networks. Managed networks have a managing element: the access point. In this mode (also referred to as infrastructure mode), all connections of the WLAN stations in the network run over the access point, which may also serve as a connection to an ethernet. Ad-hoc networks do not have an access point. The stations communicate directly with each other. The transmission range and number of participating stations are greatly limited in ad-hoc networks. Therefore, an access point is usually more efficient. It is even possible to use a WLAN card as an access point. Most cards support this functionality.

Because a wireless network is much easier to intercept and compromise than a wired network, the various standards include authentication and encryption methods. In the original version of the IEEE 802.11 standard, these are described under the term WEP. However, because WEP has proven to be insecure (see [Section “Security”](#) (page 612)), the WLAN industry (joined under the name *Wi-Fi Alliance*) has defined a new extension called WPA, which is supposed to eliminate the weaknesses of WEP. The later IEEE 802.11i standard (also referred to as WPA2, because WPA is based on a draft version 802.11i) includes WPA and some other authentication and encryption methods.

Authentication

To make sure that only authorized stations can connect, various authentication mechanisms are used in managed networks:

Open

An open system is a system that does not require authentication. Any station can join the network. Nevertheless, WEP encryption (see [Section “Encryption”](#) (page 607)) can be used.

Shared Key (according to IEEE 802.11)

In this procedure, the WEP key is used for the authentication. However, this procedure is not recommended, because it makes the WEP key more susceptible to attacks. All an attacker needs to do is to listen long enough to the communication between the station and the access point. During the authentication process, both sides exchange the same information, once in encrypted form and once in unencrypted form. This makes it possible for the key to be reconstructed with suitable tools. Because this method makes use of the WEP key for the authentication and for the encryption, it does not enhance the security of the network. A station that has the correct WEP key can authenticate, encrypt, and decrypt. A station that does not have the key cannot decrypt received packets. Accordingly, it cannot communicate, regardless of whether it had to authenticate itself.

WPA-PSK (according to IEEE 802.1x)

WPA-PSK (PSK stands for preshared key) works similarly to the Shared Key procedure. All participating stations as well as the access point need the same key. The key is 256 bits in length and is usually entered as a passphrase. This system does not need a complex key management like WPA-EAP and is more suitable for private use. Therefore, WPA-PSK is sometimes referred to as WPA “Home”.

WPA-EAP (according to IEEE 802.1x)

Actually, WPA-EAP is not an authentication system but a protocol for transporting authentication information. WPA-EAP is used to protect wireless networks in enterprises. In private networks, it is scarcely used. For this reason, WPA-EAP is sometimes referred to as WPA “Enterprise”.

WPA-EAP needs a Radius server to authenticate users. EAP offers three different methods for connecting and authenticating to the server: TLS (Transport Layer Security), TTLS (Tunneled Transport Layer Security), and PEAP (Protected Extensible Authentication Protocol). In a nutshell, these options work as follows:

EAP-TLS

TLS authentication relies on the mutual exchange of certificates both for server and client. First, the server presents its certificate to the client where it is evaluated. If the certificate is considered valid, the client in turn presents its certificate to the server. While TLS is secure, it requires a working certification management infrastructure in your network. This infrastructure is rarely found in private networks.

EAP-TTLS and PEAP

Both TTLS and PEAP are two-stage protocols. In the first stage, a secure connection is established and in the second one the client authentication data is exchanged. They require far less certification management overhead than TLS, if any.

Encryption

There are various encryption methods to ensure that no unauthorized person can read the data packets that are exchanged in a wireless network or gain access to the network:

WEP (defined in IEEE 802.11)

This standard makes use of the RC4 encryption algorithm, originally with a key length of 40 bits, later also with 104 bits. Often, the length is declared as 64 bits or 128 bits, depending on whether the 24 bits of the initialization vector are included. However, this standard has some weaknesses. Attacks against the keys generated by this system may be successful. Nevertheless, it is better to use WEP than not encrypt the network at all.

TKIP (defined in WPA/IEEE 802.11i)

This key management protocol defined in the WPA standard uses the same encryption algorithm as WEP, but eliminates its weakness. Because a new key is generated

for every data packet, attacks against these keys are in vain. TKIP is used together with WPA-PSK.

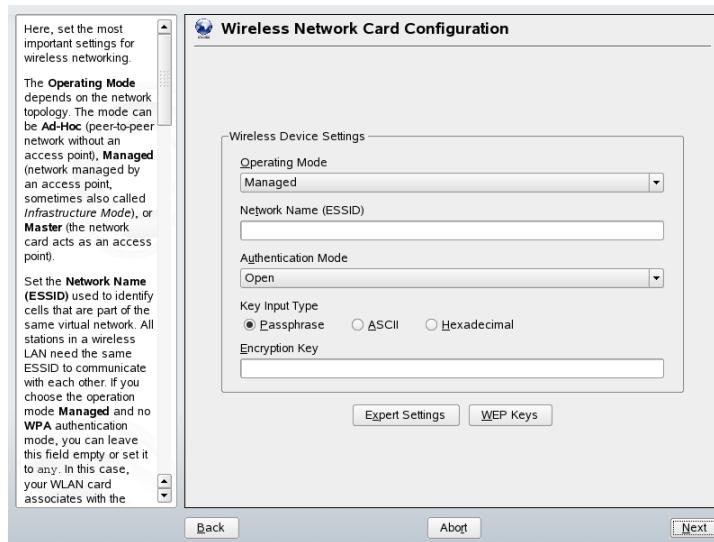
CCMP (defined in IEEE 802.11i)

CCMP describes the key management. Usually, it is used in connection with WPA-EAP, but it can also be used with WPA-PSK. The encryption takes place according to AES and is stronger than the RC4 encryption of the WEP standard.

34.1.3 Configuration with YaST

To configure your wireless network card, start the YaST *Network Card* module. Here you can also choose whether to use YaST or NetworkManager for managing your network card. If you select YaST, select the device type *Wireless* in *Network Address Setup* and click *Next*. In *Wireless Network Card Configuration*, shown in [Figure 34.1, “YaST: Configuring the Wireless Network Card”](#) (page 608), make the basic settings for the WLAN operation:

Figure 34.1 YaST: Configuring the Wireless Network Card



Operating Mode

A station can be integrated in a WLAN in three different modes. The suitable mode depends on the network in which to communicate: *Ad-hoc* (peer-to-peer network without access point), *Managed* (network is managed by an access point), or *Master* (your network card should be used as the access point). To use any of the WPA-PSK or WPA-EAP modes, the operating mode must be set to *managed*.

Network Name (ESSID)

All stations in a wireless network need the same ESSID for communicating with each other. If nothing is specified, the card automatically selects an access point, which may not be the one you intended to use.

Authentication Mode

Select a suitable authentication method for your network: *Open*, *Shared Key*, *WPA-PSK*, or *WPA-EAP*. If you select WPA authentication, a network name must be set.

Expert Settings

This button opens a dialog for the detailed configuration of your WLAN connection. A detailed description of this dialog is provided later.

After completing the basic settings, your station is ready for deployment in the WLAN.

IMPORTANT: Security in Wireless Networks

Be sure to use one of the supported authentication and encryption methods to protect your network traffic. Unencrypted WLAN connections allow third parties to intercept all network data. Even a weak encryption (WEP) is better than none at all. Refer to [Section “Encryption”](#) (page 607) and [Section “Security”](#) (page 612) for information.

Depending on the selected authentication method, YaST prompts you to fine-tune the settings in another dialog. For *Open*, there is nothing to configure, because this setting implements unencrypted operation without authentication.

WEP Keys

Set a key input type. Choose from *Passphrase*, *ASCII*, or *Hexadecimal*. You may keep up to four different keys to encrypt the transmitted data. Click *Multiple Keys* to enter the key configuration dialog. Set the length of the key: *128 bit* or *64 bit*. The default setting is *128 bit*. In the list area at the bottom of the dialog, up to four

different keys can be specified for your station to use for the encryption. Press *Set as Default* to define one of them as the default key. Unless you change this, YaST uses the first entered key as the default key. If the standard key is deleted, one of the other keys must be marked manually as the default key. Click *Edit* to modify existing list entries or create new keys. In this case, a pop-up window prompts you to select an input type (*Passphrase*, *ASCII*, or *Hexadecimal*). If you select *Passphrase*, enter a word or a character string from which a key is generated according to the length previously specified. *ASCII* requests an input of 5 characters for a 64-bit key and 13 characters for a 128-bit key. For *Hexadecimal*, enter 10 characters for a 64-bit key or 26 characters for a 128-bit key in hexadecimal notation.

WPA-PSK

To enter a key for WPA-PSK, select the input method *Passphrase* or *Hexadecimal*. In the *Passphrase* mode, the input must be 8 to 63 characters. In the *Hexadecimal* mode, enter 64 characters.

WPA-EAP

Enter the credentials you have been given by your network administrator. For TLS, provide the *Client Certificate* and *Server Certificate*. TTLS and PEAP require *Identity* and *Password*. *Server Certificate* is optional. YaST searches for any certificate under `/etc/cert`, so save the certificates given to you to this location and restrict access to these files to 0600 (owner read and write).

Click *Advanced Settings* to enter the advanced authentication dialog for your WPA-EAP setup. Select the authentication method for the second stage of EAP-TTLS or EAP-PEAP communication. If you selected TTLS in the previous dialog, choose `auto`, MD5, GTC, CHAP, PAP, MSCHAPv1, or MSCHAPv2. If you selected PEAP, choose `auto`, MD5, GTC, or MSCHAPv2. *PEAP version* can be used to force the use of a certain PEAP implementation if the automatically-determined setting does not work for you. Leave this dialog with *OK*.

Click *Expert Settings* to leave the dialog for the basic configuration of the WLAN connection and enter the expert configuration. The following options are available in this dialog:

Channel

The specification of a channel on which the WLAN station should work is only needed in *Ad-hoc* and *Master* modes. In *Managed* mode, the card automatically searches the available channels for access points. In *Ad-hoc* mode, select one of the 12 offered channels for the communication of your station with the other stations.

In *Master* mode, determine on which channel your card should offer access point functionality. The default setting for this option is *Auto*.

Bit Rate

Depending on the performance of your network, you may want to set a certain bit rate for the transmission from one point to another. In the default setting *Auto*, the system tries to use the highest possible data transmission rate. Some WLAN cards do not support the setting of bit rates.

Access Point

In an environment with several access points, one of them can be preselected by specifying the MAC address.

Use Power Management

When you are on the road, use power saving technologies to maximize the operating time of your battery. More information about power management is available in [Chapter 33, *Power Management* \(page 577\)](#).

34.1.4 Utilities

`hostap` (package `hostap`) is used to run a WLAN card as an access point. More information about this package is available at the project home page (<http://hostap.epitest.fi/>).

`kismet` (package `kismet`) is a network diagnosis tool with which to listen to the WLAN packet traffic. In this way, you can also detect any intrusion attempts in your network. More information is available at <http://www.kismetwireless.net/> and in the manual page.

34.1.5 Tips and Tricks for Setting Up a WLAN

These tips can help tweak speed and stability as well as security aspects of your WLAN.

Stability and Speed

The performance and reliability of a wireless network mainly depend on whether the participating stations receive a clean signal from the other stations. Obstructions like walls greatly weaken the signal. The more the signal strength sinks, the more the transmission slows down. During operation, check the signal strength with the iwconfig utility on the command line (Link Quality field) or with KInternet in KDE. If you have problems with the signal quality, try to set up the devices somewhere else or adjust the position of the antennas of your access points. Auxiliary antennas that substantially improve the reception are available for a number of PCMCIA WLAN cards. The rate specified by the manufacturer, such as 54 MBit/s, is a nominal value that represents the theoretical maximum. In practice, the maximum data throughput is no more than half this value.

Security

If you want to set up a wireless network, remember that anybody within the transmission range can easily access it if no security measures are implemented. Therefore, be sure to activate an encryption method. All WLAN cards and access points support WEP encryption. Although this is not entirely safe, it does present an obstacle for a potential attacker. WEP is usually adequate for private use. WPA-PSK would be even better, but it is not implemented in older access points or routers with WLAN functionality. On some devices, WPA can be implemented by means of a firmware update. Furthermore, Linux does not support WPA on all hardware components. When this documentation was prepared, WPA only worked with cards using Atheros, Intel PRO/Wireless, or Prism2/2.5/3 chips. On Prism2/2.5/3, WPA only works if the hostap driver is used (see [Section “Problems with Prism2 Cards”](#) (page 613)). If WPA is not available, WEP is better than no encryption. In enterprises with advanced security requirements, wireless networks should only be operated with WPA.

34.1.6 Troubleshooting

If your WLAN card fails to respond, check if you have downloaded the needed firmware. Refer to [Section 34.1.1, “Hardware”](#) (page 604). The following paragraphs cover some known problems.

Multiple Network Devices

Modern laptops usually have a network card and a WLAN card. If you configured both devices with DHCP (automatic address assignment), you may encounter problems with the name resolution and the default gateway. This is evident from the fact that you can ping the router but cannot surf the Internet. The Support Database at <http://portal.suse.com> features an article on this subject. To find the article, enter “DHCP” in the search dialog.

Problems with Prism2 Cards

Several drivers are available for devices with Prism2 chips. The various cards work more or less smoothly with the various drivers. With these cards, WPA is only possible with the hostap driver. If such a card does not work properly or not at all or you want to use WPA, read `/usr/share/doc/packages/wireless-tools/README.prism2`.

WPA

WPA support is quite new in SUSE Linux and still under development. Thus, YaST does not support the configuration of all WPA authentication methods. Not all wireless LAN cards and drivers support WPA. Some cards need a firmware update to enable WPA. If you want to use WPA, read `/usr/share/doc/packages/wireless-tools/README.wpa`.

34.1.7 For More Information

The Internet pages of Jean Tourrilhes, who developed the *Wireless Tools* for Linux, present a wealth of useful information about wireless networks. See http://www.hpl.hp.com/personal/Jean_Tourrilhes/Linux/Wireless.html.

34.2 Bluetooth

Bluetooth is a wireless technology for connecting various devices, such as cellular phones, PDAs, peripheral devices, laptops, or system components like the keyboard or mouse. The name is derived from the Danish king Harold Bluetooth, who united various warring factions in Scandinavia. The Bluetooth logo is based on the runes for “H” (resembles a star) and “B”.

A number of important aspects distinguish Bluetooth from IrDA. First, the individual devices do not need to “see” each other directly and, second, several devices can be connected in a network. However, the maximum data rate is 720 Kbps (in the current version 1.2). Theoretically, Bluetooth can even communicate through walls. In practice, however, this depends on the properties of the wall and the device class. There are three device classes with transmission ranges between ten and a hundred meters.

34.2.1 Basics

The following sections outline the basic principles of how Bluetooth works. Learn which software requirements need to be met, how Bluetooth interacts with your system, and how Bluetooth profiles work.

Software

To be able to use Bluetooth, you need a Bluetooth adapter (either a built-in adapter or an external device), drivers, and a Bluetooth protocol stack. The Linux kernel already contains the basic drivers for using Bluetooth. The Bluez system is used as protocol stack. To make sure that the applications work with Bluetooth, the base packages `bluez-libs` and `bluez-utils` must be installed. These packages provide a number of needed services and utilities. Additionally, some adapters, such as Broadcom or AVM BlueFritz!, require the `bluez-firmware` package to be installed. The `bluez-cups` package enables printing over Bluetooth connections. If you need to debug problems with Bluetooth connections, install the package `bluez-hcidump`.

General Interaction

A Bluetooth system consists of four interlocked layers that provide the desired functionality:

Hardware

The adapter and a suitable driver for support by the Linux kernel.

Configuration Files

Used for controlling the Bluetooth system.

Daemons

Services that are controlled by the configuration files and provide the functionality.

Applications

The applications allow the functionality provided by the daemons to be used and controlled by the user.

When inserting a Bluetooth adapter, its driver is loaded by the hotplug system. After the driver is loaded, the system checks the configuration files to see if Bluetooth should be started. If this is the case, it determines the services to start. Based on this information, the respective daemons are started. Bluetooth adapters are probed upon installation. If one or more are found, Bluetooth is enabled. Otherwise the Bluetooth system is deactivated. Any Bluetooth device added later must be enabled manually.

Profiles

In Bluetooth, services are defined by means of profiles, such as the file transfer profile, the basic printing profile, and the personal area network profile. To enable a device to use the services of another device, both must understand the same profile—a piece of information that is often missing in the device package and manual. Unfortunately, some manufacturers do not comply strictly with the definitions of the individual profiles. Despite this, communication between the devices usually works smoothly.

In the following text, local devices are those physically connected to the computer. All other devices that can only be accessed over wireless connections are referred to as remote devices.

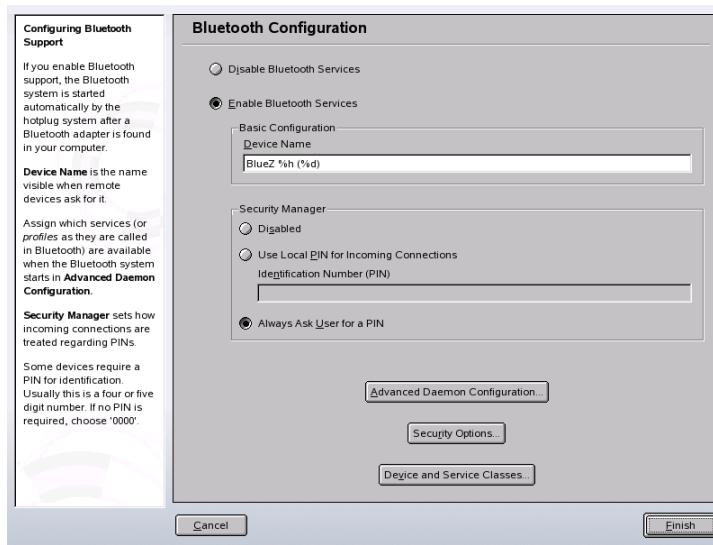
34.2.2 Configuration

This section introduces Bluetooth configuration. Learn which configuration files are involved, which tools are needed, and how to configure Bluetooth with YaST or manually.

Configuring Bluetooth with YaST

Use the YaST Bluetooth module, shown in [Figure 34.2, “YaST Bluetooth Configuration”](#) (page 616), to configure Bluetooth support on your system. As soon as hotplug detects a Bluetooth adapter on your system (for example, during booting or when you plug in an adapter), Bluetooth is automatically started with the settings configured in this module.

Figure 34.2 YaST Bluetooth Configuration



In the first step of the configuration, determine whether Bluetooth services should be started on your system. If you have enabled the Bluetooth services, two things can be configured. First, the *Device Name*. This is the name other devices display when your computer has been discovered. There are two placeholders available—%h stands for the hostname of the system (useful, for example, if it is assigned dynamically by DHCP) and %d inserts the interface number (only useful if you have more than one Bluetooth adapter in your computer). For example, if you enter Laptop %h in the field and DHCP assigns the name unit123 to your computer, other remote devices would know your computer as Laptop unit123.

The *Security Manager* parameter is related to the behavior of the local system when a remote device tries to connect. The difference is in the handling of the PIN number. Either allow any device to connect without a PIN or determine how the correct PIN is chosen if one is needed. You can enter a PIN (stored in a configuration file) in the appropriate input field. If a device tries to connect, it first uses this PIN. If it fails, it falls back to using no PIN. For maximum security, it is best to choose *Always Ask User for PIN*. This option allows you to use different PINs for different (remote) devices.

Click *Advanced Daemon Configuration* to enter the dialog for selecting and configuring the available services (called *profiles* in Bluetooth). All available services are displayed in a list and can be enabled or disabled by clicking *Activate* or *Deactivate*. Click *Edit* to open a dialog in which to specify additional arguments for the selected service (daemon). Do not change anything unless you are familiar with the service. After completing the configuration of the daemons, exit this dialog by clicking *OK*.

Back in the main dialog, click *Security Options* to enter the security dialog and specify encryption, authentication, and scan settings. Then exit the security dialog to return to the main dialog. After you close the main dialog with *Finish*, your Bluetooth system is ready for use.

From the main dialog, you can reach the *Device and Service Classes* dialog, too. Bluetooth devices are grouped into various device classes. In this dialog, choose the correct one for your computer, such as *Desktop* or *Laptop*. The device class is not very important, unlike the service class, also set here. Sometimes remote Bluetooth devices, like cell phones, only allow certain functions if they can detect the correct service class set on your system. This is often the case for cell phones that expect a class called *Object Transfer* before they allow the transfer of files from or to the computer. You can choose multiple classes. It is not useful to select all classes “just in case.” The default selection should be appropriate in most cases.

To use Bluetooth to set up a network, activate *PAND* in the *Advanced Daemon Configuration* dialog and set the mode of the daemon with *Edit*. For a functional Bluetooth network connection, one pand must operate in the *Listen* mode and the peer in the *Search* mode. By default, the *Listen* mode is preset. Adapt the behavior of your local pand. Additionally, configure the bnepx interface (X stands for the device number in the system) in the YaST *Network Card* module.

Configuring Bluetooth Manually

The configuration files for the individual components of the Bluez system are located in the directory `/etc/bluetooth`. The only exception is the file `/etc/sysconfig/bluetooth` for starting the components, which is modified by the YaST module.

The configuration files described below can only be modified by the user `root`. Currently, there is no graphical user interface to change all settings. The most important ones can be set using the YaST Bluetooth module, described in [Section “Configuring Bluetooth with YaST” \(page 616\)](#). All other settings are only of interest for experienced users with special cases. Usually, the default settings should be adequate.

A PIN number provides basic protection against unwanted connections. Mobile phones usually query the PIN when establishing the first contact (or when setting up a device contact on the phone). For two devices to be able to communicate, both must identify themselves with the same PIN. On the computer, the PIN is located in the file `/etc/bluetooth/pin`.

IMPORTANT: Security of Bluetooth Connections

Despite the PINs, the transmission between two devices may not be fully secure. By default, the authentication and encryption of Bluetooth connections is deactivated. Activating authentication and encryption may result in communication problems with some Bluetooth devices.

Various settings, such as the device names and the security mode, can be changed in the configuration file `/etc/bluetooth/hcid.conf`. Usually, the default settings should be adequate. The file contains comments describing the options for the various settings.

Two sections in the included file are designated as `options` and `device`. The first contains general information that `hcid` uses for starting. The latter contains settings for the individual local Bluetooth devices.

One of the most important settings of the `options` section is `security auto;`. If set to `auto`, `hcid` tries to use the local PIN for incoming connections. If it fails, it switches to `none` and establishes the connection anyway. For increased security, this default setting should be set to `user` to make sure that the user is requested to enter a PIN every time a connection is established.

Set the name under which the computer is displayed on the other side in the device section. The device class, such as Desktop, Laptop, or Server, is defined in this section. Authentication and encryption are also enabled or disabled here.

34.2.3 System Components and Utilities

The operability of Bluetooth depends on the interaction of various services. At least two background daemons are needed: hcid (host controller interface daemon), which serves as an interface for the Bluetooth device and controls it, and sdpd (service discovery protocol daemon), by means of which a device can find out which services the host makes available. If they are not activated automatically when the system is started, both hcid and sdpd can be activated with the command `rcbluetooth start`. This command must be executed as `root`.

The following paragraphs briefly describe the most important shell tools that can be used for working with Bluetooth. Although various graphical components are now available for controlling Bluetooth, it can be worthwhile to check these programs.

Some of the commands can only be executed as `root`. This includes the command `l2ping device_address` for testing the connection to a remote device.

hcitool

hcitool can be used to determine whether local and remote devices are detected. The command `hcitool dev` lists the local devices. The output generates a line in the form `interface_name device_address` for every detected local device.

Search for remote devices with the command `hcitool inq`. Three values are returned for every detected device: the device address, the clock offset, and the device class. The device address is important, because other commands use it for identifying the target device. The clock offset mainly serves a technical purpose. The class specifies the device type and the service type as a hexadecimal value.

The command `hcitool name device-address` can be used to determine the device name of a remote device. In the case of a remote computer, the class and the device name correspond to the information in its `/etc/bluetooth/hcid.conf`. Local device addresses generate an error output.

hciconfig

The command `/usr/sbin/hciconfig` delivers further information about the local device. If `hciconfig` is executed without any arguments, the output shows device information, such as the device name (`hciX`), the physical device address (a 12-digit number in the form `00:12:34:56:78`), and information about the amount of transmitted data.

`hciconfig hci0 name` displays the name that is returned by your computer when it receives requests from remote devices. As well as querying the settings of the local device, `hciconfig` can be used for modifying these settings. For example, `hciconfig hci0 name TEST` sets the name to `TEST`.

sdptool

The program `sdptool` can be used to check which services are made available by a specific device. The command `sdptool browse device_address` returns all services of a device. Use the command `sdptool search service_code` to search for a specific service. This command scans all accessible devices for the requested service. If one of the devices offers the service, the program prints the full service name returned by the device together with a brief description. View a list of all possible service codes by entering `sdptool` without any parameters.

34.2.4 Graphical Applications

In Konqueror, enter the URL `bluetooth:/` to list local and remote Bluetooth devices. Double-click a device for an overview of the services provided by the device. If you move across one of the specified services with the mouse, the browser's status bar shows which profile is used for the service. If you click a service, a dialog opens, asking what to do: save, use the service (an application must be started to do this), or cancel the action. Mark a check box if you do not want the dialog to be displayed again but always want the selected action to be performed. For some services, support is not yet available. For others, additional packages may need to be installed.

34.2.5 Examples

This section features two typical examples of possible Bluetooth scenarios. The first shows how a network connection between two hosts can be established via Bluetooth. The second features a connection between a computer and a mobile phone.

Network Connection between Two Hosts

In the first example, a network connection is established between the hosts *H1* and *H2*. These two hosts have the Bluetooth device addresses *baddr1* and *baddr2* (determined on both hosts with the command `hcitool dev` as described above). The hosts should be identified with the IP addresses 192.168.1.3 (*H1*) and 192.168.1.4 (*H2*).

The Bluetooth connection is established with the help of `pand` (personal area networking daemon). The following commands must be executed by the user `root`. The description focuses on the Bluetooth-specific actions and does not provide a detailed explanation of the network command `ip`.

Enter `pand -s` to start `pand` on the host *H1*. Subsequently, a connection can be established on the host *H2* with `pand -c baddr1`. If you enter `ip link show` on one of the hosts to list the available network interfaces, the output should contain an entry like the following:

```
bnep0: <BROADCAST,MULTICAST> mtu 1500 qdisc noop qlen 1000  
link/ether 00:12:34:56:89:90 brd ff:ff:ff:ff:ff:ff
```

Instead of `00:12:34:56:89:90`, the output should contain the local device address *baddr1* or *baddr2*. Now this interface must be assigned an IP address and activated. On *H1*, this can be done with the following two commands:

```
ip addr add 192.168.1.3/24 dev bnep0  
ip link set bnep0 up
```

On *H2*:

```
ip addr add 192.168.1.4/24 dev bnep0  
ip link set bnep0 up
```

Now *H1* can be accessed from *H2* under the IP 192.168.1.3. Use the command `ssh 192.168.1.4` to access *H2* from *H1*, assuming *H2* runs an `sshd`, which is activated by default in SUSE Linux. The command `ssh 192.168.1.4` can also be run as a normal user.

Data Transfer from a Mobile Phone to the Computer

The second example shows how to transfer a photograph created with a mobile phone with a built-in digital camera to a computer (without incurring additional costs for the transmission of a multimedia message). Although the menu structure may differ on various mobile phones, the procedure is usually quite similar. Refer to the manual of your phone, if necessary. This example describes the transfer of a photograph from a Sony Ericsson mobile phone to a laptop. The service Obex-Push must be available on the computer and the computer must grant the mobile phone access. In the first step, the service is made available on the laptop. You need a special service daemon running on the laptop to get the data from the phone. If the package `kbluetooth` is installed, you do not need to start a special daemon. If `kbluetooth` is not installed, use the `opd` daemon from the `bluez-utils` package. Start the daemon with the following command:

```
opd --mode OBEX --channel 10 --daemonize --path /tmp --sdp
```

Two important parameters are used: `--sdp` registers the service with `sdpd` and `--path /tmp` instructs the program where to save the received data—in this case to `/tmp`. You can also specify any other directory to which you have write access.

If you use `kbluetooth`, you are prompted for a directory when the photograph is received on the laptop.

Now the mobile phone must get to know the computer. To do this, open the *Connect* menu on the phone and select *Bluetooth*. If necessary, click *Turn On* before selecting *My devices*. Select *New device* and let your phone search for the laptop. If a device is detected, its name appears in the display. Select the device associated with the laptop. If you encounter a PIN query, enter the PIN specified in `/etc/bluetooth/pin`. Now your phone recognizes the laptop and is able to exchange data with the laptop. Exit the current menu and go to the image menu. Select the image to transfer and press *More*. In the next menu, press *Send* to select a transmission mode. Select *Via Bluetooth*. The laptop should be listed as a target device. Select the laptop to start the transmission. The image is then saved to the directory specified with the `opd` command. Audio tracks can be transferred to the laptop in the same way.

34.2.6 Troubleshooting

If you have difficulties establishing a connection, proceed according to the following list. Remember that the error can be on either side of a connection or even on both sides. If possible, reconstruct the problem with another Bluetooth device to verify that the device is not defective.

Is the local device listed in the output of `hcitool dev`?

If the local device is not listed in this output, hcid is not started or the device is not recognized as a Bluetooth device. This can have various causes. The device may be defective or the correct driver may be missing. Laptops with built-in Bluetooth often have an on and off switch for wireless devices, like WLAN and Bluetooth. Check the manual of your laptop to see if your device has such a switch. Restart the Bluetooth system with the command `rcbluetooth restart` and check if any errors are reported in `/var/log/messages`.

Does your Bluetooth adapter need a firmware file?

If it does, install `bluez-bluefw` and restart the Bluetooth system with `rcbluetooth restart`.

Does the output of `hcitool inq` return other devices?

Test this command more than once. The connection may have interferences, because the frequency band of Bluetooth is also used by other devices.

Do the PINs match?

Check if the PIN number of the computer (in `/etc/bluetooth/pin`) matches that of the target device.

Can the remote device “see” your computer?

Try to establish the connection from the remote device. Check if this device sees the computer.

Can a network connection be established (see [Section “Network Connection between Two Hosts”](#) (page 621))?

The setup described in [Section “Network Connection between Two Hosts”](#) (page 621) may not work for several reasons. For example, one of the two computers may not support the ssh protocol. Try `ping 192.168.1.3` or `ping 192.168.1.4`. If this works, check if `sshd` is active. Another problem could be that one of the two devices already has network settings that conflict with the address `192.168.1.X`

in the example. If this is the case, try different addresses, such as 10.123.1.2 and 10.123.1.3.

Does the laptop appear as a target device (see [Section “Data Transfer from a Mobile Phone to the Computer”](#) (page 622))? Does the mobile device recognize the Obex-Push service on the laptop?

In *My devices*, select the respective device and view the list of *Services*. If Obex-Push is not displayed (even after the list is updated), the problem is caused by opd on the laptop. Is opd active? Do you have write access to the specified directory?

Does the scenario described in [Section “Data Transfer from a Mobile Phone to the Computer”](#) (page 622) work the other way around?

If the obexftp package is installed, the command `obexftp -b device_address -B 10 -p image` can be used on some devices. Several Siemens and Sony Ericsson models have been tested and found to be functional. Refer to the documentation in `/usr/share/doc/packages/obexftp`.

If you have installed the bluez-hcidump package, you can use `hcidump -X` to check what is sent between the devices. Sometimes the output helps give a hint where the problem is, but be aware of the fact that it is only partly in “clear text.”

34.2.7 For More Information

Some additional (last-minute) documentation can be found in `/usr/share/doc/packages/bluez-utils/` (German and English version available).

An extensive overview of various instructions for the use and configuration of Bluetooth is available at <http://www.holtmann.org/linux/bluetooth/>. Other useful information and instructions:

- Official howto of the Bluetooth protocol stack integrated in the kernel: <http://bluez.sourceforge.net/howto/index.html>
- Connection to PalmOS PDA: <http://www.cs.ucl.ac.uk/staff/s.zachariadis/btpalmlinux.html>

34.3 Infrared Data Transmission

IrDA (Infrared Data Association) is an industry standard for wireless communication with infrared light. Many laptops sold today are equipped with an IrDA-compatible transceiver that enables communication with other devices, such as printers, modems, LANs, or other laptops. The transfer speed ranges from 2400 bps to 4 Mbps.

There are two IrDA operation modes. The standard mode, SIR, accesses the infrared port through a serial interface. This mode works on almost all systems and is sufficient for most requirements. The faster mode, FIR, requires a special driver for the IrDA chip. Not all chip types are supported in FIR mode because of a lack of appropriate drivers. Set the desired IrDA mode in the BIOS of your computer. The BIOS also shows which serial interface is used in SIR mode.

Information about IrDA can be found in the IrDA how-to by Werner Heuser at <http://tuxmobil.org/Infrared-HOWTO/Infrared-HOWTO.html>. Additionally refer to the Web site of the Linux IrDA Project at <http://irda.sourceforge.net/>.

34.3.1 Software

The necessary kernel modules are included in the kernel package. The package `irda` provides the necessary helper applications for supporting the infrared interface. The documentation can be found at `/usr/share/doc/packages/irda/README` after the installation of the package.

34.3.2 Configuration

The IrDA system service is not started automatically when the system is booted. Use the YaST IrDA module for the activation. Only one setting can be modified in this module: the serial interface of the infrared device. The test window shows two outputs. One is the output of `irdadump`, which logs all sent and received IrDA packets. This output should contain the name of the computer and the names of all infrared devices in transmission range. An example for these messages is shown in [Section 34.3.4, “Troubleshooting”](#) (page 627). All devices to which an IrDA connection exists are listed in the lower part of the window.

IrDA consumes a considerable amount of battery power, because a discovery packet is sent every few seconds to detect other peripheral devices. Therefore, IrDA should only be started when necessary if you depend on battery power. Enter the command `rcirda start` to activate it or `rcirda stop` to deactivate it. All needed kernel modules are loaded automatically when the interface is activated.

Manual configuration can be performed in the file `/etc/sysconfig/irda`. This file contains only one variable, `IRDA_PORT`, which determines the interface to use in SIR mode.

34.3.3 Usage

Data can be sent to the device file `/dev/irlpt0` for printing. The device file `/dev/irlpt0` acts just like the normal `/dev/lp0` cabled interface, except the printing data is sent wirelessly with infrared light. For printing, make sure that the printer is in visual range of the computer's infrared interface and the infrared support is started.

A printer that is operated over the infrared interface can be configured with the YaST Printer module. Because it is not detected automatically, configure it manually by clicking *Other (not detected)*. In the following dialog, select *IrDA printer*. Usually, `irlpt0` is the right connection. Details about operating printers in Linux are available in [Chapter 11, *Printer Operation*](#) (page 227).

Communication with other hosts and with mobile phones or other similar devices is conducted through the device file `/dev/ircomm0`. The Siemens S25 and Nokia 6210 mobile phones, for example, can dial and connect to the Internet with the wvdial application using the infrared interface. Synchronizing data with a Palm Pilot is also possible, provided the device setting of the corresponding application has been set to `/dev/ircomm0`.

If you want, you can address only devices that support the printer or IrCOMM protocols. Devices that support the IROBEX protocol, such as the 3Com Palm Pilot, can be accessed with special applications, like `irobexpalm` and `irobexreceive`. Refer to the *IR-HOWTO* (<http://tldp.org/HOWTO/Infrared-HOWTO/>) for information. The protocols supported by the device are listed in brackets after the name of the device in the output of `irdadump`. IrLAN protocol support is still a “work in progress.”

34.3.4 Troubleshooting

If devices connected to the infrared port do not respond, use the command `irdadump` (as `root`) to check if the other device is recognized by the computer. Something similar to [Example 34.1, “Output of irdadump”](#) (page 627) appears regularly when a Canon BJC-80 printer is in visible range of the computer:

Example 34.1 Output of irdadump

```
21:41:38.435239 xid:cmd 5b62bed5 > ffffffff S=6 s=0 (14)
21:41:38.525167 xid:cmd 5b62bed5 > ffffffff S=6 s=1 (14)
21:41:38.615159 xid:cmd 5b62bed5 > ffffffff S=6 s=2 (14)
21:41:38.705178 xid:cmd 5b62bed5 > ffffffff S=6 s=3 (14)
21:41:38.795198 xid:cmd 5b62bed5 > ffffffff S=6 s=4 (14)
21:41:38.885163 xid:cmd 5b62bed5 > ffffffff S=6 s=5 (14)
21:41:38.965133 xid:rsp 5b62bed5 < 6cac38dc S=6 s=5 BJC-80
                           hint=8804 [Printer IrCOMM ] (23)
21:41:38.975176 xid:cmd 5b62bed5 > ffffffff S=6 s=*> earth
                           hint=0500 [ PnP Computer ] (21)
```

Check the configuration of the interface if there is no output or the other device does not reply. Verify that the correct interface is used. The infrared interface is sometimes located at `/dev/ttys2` or at `/dev/ttys3` and an interrupt other than IRQ 3 is sometimes used. These settings can be checked and modified in the BIOS setup menu of almost every laptop.

A simple video camera can also help in determining whether the infrared LED lights up at all. Most video cameras can see infrared light; the human eye cannot.

Index

Symbols

64-bit Linux, 173
kernel specifications, 176
runtime support, 173
software development, 174

A

ACLs, 141–152
access, 143, 146
check algorithm, 151
default, 144, 149
definitions, 143
effects, 149
handling, 144
masks, 148
permission bits, 145
structure, 144
support, 152
addresses
 IP, 321
Apache, 447–485
 CGI scripts, 472
 configuring, 449
 files, 449
 manually, 449–456
 virtual host, 452
 YaST, 457–463
 YaST HTTP Server Configuration, 461
 YaST HTTP wizard, 457
installing, 448
modules, 465–472
 available, 466
 building modules, 471
 external modules, 470
 installing, 466

 multiprocessing modules, 469
quick start, 447
security, 480
Squid, 536
SSL, 475–480
 configuring Apache with, 479
 creating an SSL certificate, 475
starting, 463
stopping, 463
troubleshooting, 482
authentication
 Kerberos, 80
 PAM, 301–308

B

Bash
 .bashrc, 214
 .profile, 214
 profile, 213
BIND, 376–386
Bluetooth, 549, 614
 hciconfig, 620
 hcitool, 619
 network, 617
 opd, 622
 pand, 621
 sdptool, 620
booting, 177
 boot sectors, 193–194
 configuring
 YaST, 203–208
 graphic, 209
 GRUB, 193–212
 initramfs, 179
 initrd, 179

C

cards
 graphics, 278

network, 334
cellular phones, 552
chown, 71
CJK, 221
commands
 chown, 71
 fonts-config, 279
 free, 218
 getfacl, 147
 grub, 194
 head, 71
 ldapadd, 433
 ldapdelete, 436
 ldapmodify, 435
 ldapsearch, 435
 lp, 237
 nice, 71
 rpm, 86
 rpmbuild, 86
 scp, 112
 setfacl, 147
 sftp, 112
 slptool, 364
 smbpasswd, 518
 sort, 71
 ssh, 111
 ssh-agent, 115
 ssh-keygen, 114
 tail, 71
configuration files, 351
 .bashrc, 214, 217
 .emacs, 219
 .mailsync, 504
 .profile, 214
 .xsession, 115
 acpi, 581
 crontab, 214
 csh.cshrc, 223
 dhclient.conf, 410
 dhcp, 351
 dhcpd.conf, 411
 exports, 403–404
 group, 66
 grub.conf, 201
 gshadow, 72
 host.conf, 354
 HOSTNAME, 358
 hosts, 333, 353
 ifcfg-*, 351
 inittab, 181, 183, 220
 inputrc, 221
 irda, 626
 kernel, 179
 language, 221, 223
 logrotate.conf, 215
 menu.lst, 196
 modprobe.conf, 68
 modules.conf, 68
 named.conf, 377–386, 526
 network, 351
 networks, 354
 nscd.conf, 357
 nsswitch.conf, 355, 437
 pam_unix2.conf, 436
 passwd, 66
 permissions, 138
 powersave, 581
 powersave.conf, 77
 profile, 213, 217, 223
 resolv.conf, 218, 352, 377, 525
 routes, 352
 samba, 512
 services, 512, 534
 slapd.conf, 427
 smb.conf, 513, 519
 smppd.conf, 359
 smpppd-c.conf, 360
 squid.conf, 525, 527, 531, 533, 536, 538
 squidguard.conf, 538

ssh_config, 115
suseconfig, 192
sysconfig, 189–192
termcap, 221
wireless, 351
XF86Config, 81
xorg.conf, 81, 273
 Device, 277
 Monitor, 278
 Screen, 275
configuring, 189
 cable modem, 343
 DNS, 337, 367
 DSL, 343
 GRUB, 194, 201
 IPv6, 331
 IrDA, 625
 ISDN, 339
 modems, 337
 networks, 334
 manually, 348–358
 PAM, 81
 printing, 230–233
 routing, 337, 352
 Samba, 511–516
 clients, 517
 Squid, 527
 SSH, 110
 T-DSL, 345
consoles
 assigning, 220
 graphical, 209
 switching, 220
core files, 217
cpuspeed, 589
cron, 214
CVS, 488, 497–499

D

data security, 550
deltarpm, 90
DHCP, 405–414
 configuring with YaST, 406
 dhcpd, 411–412
 packages, 410
 server, 411–412
 static address assignment, 412
digital cameras, 551
disks
 boot, 208
DNS, 333
 BIND, 376–386
 configuring, 337, 367
 domains, 353
 forwarding, 377
 logging, 381
 mail exchanger, 333
 multicast, 71
 name servers, 353
 NIC, 333
 options, 379
 reverse lookup, 385
 security and, 137
 Squid and, 526
 starting, 377
 terminology, 367
 top level domain, 333
 troubleshooting, 377
 zones
 files, 382
domain name system (see DNS)
domains
 .local, 71
DOS
 sharing files, 509

E

e-mail
 synchronizing, 489, 549
 mailsync, 504–507
editors
 Emacs, 219–220
Emacs, 219–220
 .emacs, 219
 default.el, 219
encoding
 ISO-8859-1, 223
 UTF-8, 71
encrypting, 116–119
 creating partitions, 117
 files, 118–119
 files with vi, 118
 partitions, 117–118
 removable media, 119
 YaST, with, 117
Evolution, 552

F

file systems, 259–269
 ACLs, 141–152
 cryptofs, 116
 encrypting, 116
 Ext2, 261–262
 Ext3, 262–264
 LFS, 267
 limitations, 267
 Reiser4, 264
 ReiserFS, 260–261
 selecting, 260
 supported, 266–267
 terms, 259
 XFS, 265
files
 encrypting, 118
 finding, 216

synchronizing, 487–507
 CVS, 488, 497–499
 iFolder, 490
 mailsync, 489, 504–507
 rsync, 489
 Subversion, 489
 Unison, 488, 495–496

Firefox

 URL open command, 86
firewalls, 101
 packet filters, 101, 104
 Squid and, 534
 SuSEfirewall2, 101, 105
Firewire (IEEE1394)
 hard disks, 551
flash drives, 551
fonts, 279
 CID-keyed, 284
 TrueType, 279
 X11 core, 280
 Xft, 281
FreeNX, 289–300

G

graphics
 3D, 284–287
 3Ddiag, 286
 diagnosis, 285
 drivers, 284
 installation support for, 286
 SaX, 285
 support for, 284
 testing, 286
 troubleshooting, 286
 cards
 3D, 284–287
 drivers, 278
 GLIDE, 284–287
 OpenGL, 284–287

drivers, 284
testing, 286
GRUB, 193–212
 boot menu, 196
 boot password, 202
 boot sectors, 194
 booting, 194
 commands, 194–203
 device names, 197
 device.map, 195, 200
 GRUB Geom Error, 210
 grub.conf, 195, 201
 JFS and GRUB, 210
 limitations, 194
 Master Boot Record (MBR), 193
 menu editor, 199
 menu.lst, 195–196
 partition names, 197
 troubleshooting, 210
 uninstalling, 208

H

hardware
 ISDN, 339
hciconfig, 620
hcitool, 619
head, 71
help
 info pages, 219
 man pages, 219
X, 278

I

I18N, 221
iFolder, 490
info pages, 219
init, 181
 adding scripts, 186
inittab, 181

scripts, 184–187
installation support
 3D graphics cards and, 286
installing
 GRUB, 194
 manually, 80
 packages, 87
internationalization, 221
Internet
 cinternet, 360
 dial-up, 359–361
 DSL, 343
 ISDN, 339
 KInternet, 360
 qinternet, 360
 smpppd, 359–361
 TDSL, 345
IP addresses
 classes, 321
 dynamic assignment, 405
IPv6, 323
 configuring, 331
 masquerading, 103
 private, 323
IrDA, 550, 625–627
 configuring, 625
 starting, 625
 stopping, 625
 troubleshooting, 627

K

kernels
 caches, 218
 limits, 268
 modules
 modprobe.conf, 68
 version 2.6, 68
keyboard
 Asian characters, 221

layout, 220
mapping, 220
 compose, 221
 multikey, 221
X Keyboard Extension, 221
XKB, 221
Kontact, 552
KPilot, 552
KPowersave, 547
KSysguard, 547

L

L10N, 221
laptops, 543–550
 hardware, 543
 IrDA, 625–627
 NetworkManager, 546
 PCMCIA, 543
 power management, 544, 577–589
 SCPM, 544, 563
 SLP, 547
LDAP, 421–446
 access control, 430
 ACLs, 428
 adding data, 432
 administering groups, 444
 administering users, 444
 deleting data, 436
 directory tree, 424
 ldapadd, 432
 ldapdelete, 436
 ldapmodify, 434
 ldapsearch, 435
 modifying data, 434
 searching data, 435
 server configuration, 427
YaST
 modules, 437
 templates, 437

YaST LDAP client, 436
LFS, 267
Lightweight Directory Access Protocol
(see LDAP)
Linux
 networks and, 317
 sharing files with another OS, 509
 uninstalling, 208
linuxrc
 manual installation, 80
linuxthreads, 69
locale
 UTF-8, 71
localization, 221
locate, 216
log files, 215
 boot.msg, 581
 messages, 110, 377
 Squid, 526, 529, 535
 Unison, 496
 XFree86, 286
Logical Volume Manager (see LVM)
logrotate, 215
LSB
 installing packages, 86
LVM
 YaST, 53

M

man pages, 219
masquerading, 103
 configuring with SuSEfirewall2, 105
Master Boot Record (see MBR)
MBR, 193–194
memory
 RAM, 218
mobility, 543–553
 cellular phones, 552
 data security, 550

digital cameras, 551
external hard disks, 551
Firewire (IEEE1394), 551
laptops, 543
PDAs, 552
USB, 551
modems
 cable, 343
 YaST, 337
mountd, 404

N

name servers (see DNS)
NAT (see masquerading)
NetBIOS, 509
Network File System (see NFS)
Network Information Service (see NIS)
NetworkManager, 546
networks, 317

 base network address, 322
 Bluetooth, 549, 617
 broadcast address, 323
 configuration files, 351–358
 configuring, 334–345, 348–358
 IPv6, 331
 DHCP, 405
 DNS, 333
 IrDA, 550
 localhost, 323
 netmasks, 321
 routing, 321, 337
 SLP, 363
 TCP/IP, 317
 wireless, 549
WLAN, 549
YaST, 334

NFS, 399
 clients, 399
 exporting, 402

importing, 400
mounting, 400
permissions, 403
servers, 401
nfsd, 404
NGPT, 69
nice, 71
NIS, 391–398
 clients, 397
 masters, 391–397
 slaves, 391–397
notebooks (see laptops)
Novell iFolder, 490
NPTL, 69–70
NSS, 355
 databases, 356

O

opd, 622
OpenLDAP (see LDAP)
OpenSSH (see SSH)
OS/2
 sharing files, 509

P

packages
 compiling, 94
 compiling with build, 96
 installing, 87
 LSB, 86
 package manager, 86
 RPMs, 86
 uninstalling, 87
 verifying, 87
packet filters (see firewalls)
PAM, 301–308
 configuring, 81
pand, 621
partitions

encrypting, 117
partition table, 193

PCMCIA, 543, 555
IrDA, 625–627

PDAs, 552

permissions
 ACLs, 141–152
 file permissions, 216

Pluggable Authentication Modules (see PAM)

ports
 53, 379
 scanning, 535

PostgreSQL
 updating, 66

power management, 544, 577–597
 ACPI, 577, 580–587, 592
 APM, 577, 579–580, 592
 battery monitor, 578
 charge level, 593
 cpufreq, 589
 cpuspeed, 589
 hibernation, 578
 powersave, 589
 standby, 578
 suspend, 578
 YaST, 597

powersave, 589
 configuring, 589

printing, 227, 230–233
 applications, from, 237
 command line, 237
 configuring with YaST, 230
 connection, 231
 CUPS, 237
 drivers, 232
 GDI printers, 243
 Ghostscript driver, 232
 IrDA, 626
 kprinter, 237

network, 245
port, 231
PPD file, 232
queues, 232
Samba, 510
test page, 232

troubleshooting
 network, 245

xpp, 237

private branch exchange, 341

protocols
 CIFS, 509
 IPv6, 323
 LDAP, 421
 SLP, 363
 SMB, 509

proxies (see Squid)
 advantages, 521
 caches, 521
 transparent, 533

R

RAID
 YaST, 60

remote computing
 FreeNX, 289–300

removable media
 subfs, 74

RFCs, 317

routing, 321, 337, 352
 masquerading, 103
 netmasks, 321
 routes, 352
 static, 352

RPM, 86–97
 database
 rebuilding, 89, 94
 deltarpm, 90
 dependencies, 87

patches, 89
queries, 91
rpmnew, 87
rpmod, 87
rpmsave, 87
security, 139
SRPMS, 95
tools, 97
uninstalling, 89
updating, 88
verify, 93
verifying, 87
rpmbuild, 86
rsync, 489, 502
runlevels, 181–183
 changing, 183
 editing in YaST, 188

S

Samba, 509–519
 CIFS, 509
 clients, 510, 516–517
 configuring, 511–516
 installing, 511
 login, 517
 names, 509
 permissions, 516
 printers, 510
 printing, 517
 security, 516
 server, 510
 servers, 511–516
 shares, 510, 514
 SMB, 509
 starting, 511
 stopping, 511
 swat, 512
 TCP/IP and, 509
SCPM, 563

advanced settings, 574
laptops, 544
managing profiles, 573
resource groups, 572
starting, 572
switching profiles, 573

screen
 resolution, 277

scripts
 init.d, 181, 184–187, 358
 boot, 185
 boot.local, 185
 boot.setup, 186
 halt, 186
 network, 358
 nfsserver, 358, 403
 portmap, 358, 403
 rc, 183–184, 186
 sendmail, 358
 squid, 525
 xinetd, 358
 ypbind, 358
 ypserv, 358

irda, 626
mkinitrd, 179
modify_resolvconf, 218, 353
SuSEconfig, 189–192
 disabling, 192

sdptool, 620

security, 128–140
 attacks, 136–137
 booting, 129, 131
 bugs and, 132, 135
 DNS, 137
 encrypted file system, 550
 engineering, 129
 firewalls, 101
 intrusion detection, 81
 local, 130–134
 network, 134–137

passwords, 130–131
permissions, 131–132
reporting problems, 140
RPM signatures, 139
Samba, 516
serial terminals, 129–130
Squid, 522
SSH, 110–116
tcpd, 139
telnet, 110
tips and tricks, 137
viruses, 133
worms, 137
X and, 134

Service Location Protocol (see SLP)

SGML
 directories, 74

SLP, 363, 547
 browser, 365
 Konqueror, 365
 registering services, 363
 slptool, 364

SMB (see Samba)

smpd, 509

soft RAID (see RAID)

software
 compiling, 94

sort, 71

sound
 mixers, 79

source
 compiling, 94

spm, 94

Squid, 521
 access controls, 536
 ACLs, 530
 Apache, 536
 cachemgr.cgi, 536–537
 caches, 521–522
 damaged, 526

size, 524
Calamaris, 539–540
configuring, 527
CPU and, 525
directories, 525
DNS, 526
features, 521
firewalls and, 534
log files, 526, 529, 535
object status, 523
permissions, 525, 530
RAM and, 524
reports, 539–540
security, 522
squidGuard, 537
starting, 525
statistics, 536–537
stopping, 526
system requirements, 523
transparent proxies, 533, 535
troubleshooting, 526
uninstalling, 526

SSH, 110–116
 authentication mechanisms, 114
 daemon, 113
 key pairs, 113–114
 scp, 112
 sftp, 112
 ssh, 111
 ssh-agent, 115
 ssh-keygen, 114
 sshd, 113
 X and, 115

subfs
 removable media, 74

Subversion, 489, 499

SVN (Subversion), 489

synchronizing data, 549
 e-mail, 549
 Evolution, 552

- Kontact, 552
KPilot, 552
- system
 limiting resource use, 217
 localizing, 221
 updating, 65–68
- system monitoring, 547
 KPowerSave, 547
 KSysGuard, 547
- T**
- tail, 71
- TCP/IP, 317
 ICMP, 318
 IGMP, 318
 layer model, 318
 packets, 319–320
 TCP, 318
 UDP, 318
- TEI XSL stylesheets
 new location, 84
- thread packages
 NPTL, 70
- Tripwire
 replaced by AIDE, 81
- U**
- ulimit, 217
 options, 217
- uninstalling
 GRUB, 208
 Linux, 208
- updating, 65–68
 passwd and group, 66
 problems, 66
 sound mixers, 79
 YaST, 66
- USB
 flash drives, 551
- hard disks, 551
- users
 /etc/passwd, 304, 437
- UTF-8
 encoding, 71
- V**
- variables
 environment, 222
- W**
- whois, 334
- Windows
 sharing files, 509
- wireless connections
 Bluetooth, 614
- WLAN, 549
- X**
- X
 character sets, 279
 CID-keyed fonts, 284
 drivers, 278
 font systems, 279
 fonts, 279
 help, 278
 optimizing, 273–279
 SaX2, 273
 security, 134
 SSH and, 115
 TrueType fonts, 279
 virtual screen, 277
 X11 core fonts, 280
 xf86config, 273
 xft, 279
 Xft, 281
- X Keyboard Extension (see keyboard, XKB)
- X Window System (see X)

X.Org, 273
Xen, 309
 overview, 309
Xft, 281
XKB (see keyboard, XKB)
XML
 directories, 74
xorg.conf
 color depth, 276
 Depth, 276
 Device, 276
 Display, 276
 Files, 274
 InputDevice, 274
 Modeline, 276
 modelines, 274
 Modes, 274, 276
 Monitor, 274, 276
 ServerFlags, 274

LVM, 53
modems, 337
network card, 334
NIS clients, 397
power management, 597
printing, 230–233
RAID, 60
runlevels, 188
Samba
 clients, 517
SLP browser, 365
sysconfig editor, 190
T-DSL, 345
updating, 66
YP (see NIS)

Y

YaST
 3D, 285
 boot configuration, 203
 default system, 206
 security, 207
 time-out, 206
 boot loader
 disk order, 207
 location, 205
 password, 207
 type, 204
 cable modem, 343
 DHCP, 406
 DSL, 343
 GRUB, 204
 ISDN, 339
 LDAP client, 436
 LILO, 204